Chaos and bifurcations in continuous language-model feedback

Timothy D. Jones

Supplementary material to the article of the same title,
arXiv:2610.31692

Originally dated October 22, 2026
Last updated November 13, 2026 at 15:01:03 PST

Abstract

Feeding a pretrained language model the expected embedding of its own next-token distribution turns it into a discrete-time dynamical system with continuous states, whose consequences for inference depend on the feedback architecture and the numerical regime. In Qwen2.5, OLMo-2 and Llama-3 models of 0.5B to 72B parameters, temperature and steering move the loop between periodic dynamics, which write repeating text, and regimes with positive finite-time Lyapunov exponents that automatic-differentiation tangents reproduce. We call these regimes chaotic as a numerical label, not a proof of an asymptotic chaotic set; 74 to 92 % of sampled labels persist to 10^5 steps. Finite-time expansion has direct consequences for reproducibility. Rerun with only the batch size changed, chaotic float64 loops separate until their most probable tokens differ, on the schedule the Lyapunov exponent sets, while periodic loops mostly stay together, in four models up to 70B; in Qwen2.5-0.5B the text changed in 173 of 192 chaotic loops but in only 7 of 192 periodic ones. In a control fixed in advance in Qwen2.5-0.5B (448 cells), with the matrix shapes held fixed the runs are bitwise identical, and a displacement injected instead grows on the same schedule, although one secondary criterion failed (from a displacement of 10^{-9} on, more than 10 % of the periodic loops also separated). In arithmetic tests designed before the runs, changing only the batch layout, which alters the floating-point execution while the intended computation stays the same, changed the answers of bfloat16 soft-thinking inference on 56 % of 96 problems in a 0.5B model and 55 % in a 7B model, far more often than with greedy decoding, at temperatures where the pipeline answers poorly (none changed where it answers well); an exploratory replacement of the cache and logits by those of a reference run before every soft step removed most of these changes. In tests fixed before their data, Soft Thinking with its published settings changed 21 % of its answers to GSM8K problems not used before in this study with the batch layout alone, and reasoning models that answer 86 to 90 % of them correctly (DeepSeek-R1-Distill-Qwen-7B, QwQ-32B) changed 14.5 and 19.5 %, the two layouts disagreeing on correctness for 11.7 and 14.5 % of the problems. Ordinary sampled reasoning with the same random numbers changed about as many, so this sensitivity is not unique to continuous feedback. Nearly all of these changes came from long runs that reached the thinking or answer budget in at least one layout. Two mechanisms contribute, depending on the model. In float64, injected perturbations in the Qwen2.5-7B and DeepSeek-R1-Distill-Qwen-7B pipelines mostly shrink over the observed runs, consistent with bfloat16 rounding being carried through long feedback, while the loops of QwQ-32B and DeepSeek-R1-Distill-Llama-8B grow slowly in more than half of the copies, and a relative displacement of 10^{-10} alone changes about 14 and 18 % of their answers, respectively. Exactly repeating text can hide a state with a positive finite-time exponent: the text records only the most probable token, which can repeat while the fed-back probabilities do not repeat over the run. In Qwen2.5-0.5B, along recorded orbits, the

growth that steering adds comes mainly from the readout, where next tokens compete; a logit bias alone suffices for chaotic cells. Earlier injection carries chaos to lower temperature, and steering at two depths creates chaos neither depth causes alone. Whether a prompt's loop is chaotic does not predict the accuracy of ordinary sampled answers beyond what the model's uncertainty about its first answer token predicts, and within the precision of the tested designs we detect no edge-of-chaos advantage. The loop shows quantitative signatures of low-dimensional maps: late finite-order doubling ratios approach Feigenbaum's constant numerically (4.6688 against 4.6692, at the float64 floor of the computation), periodic windows accumulate on folds with distances close to a $1/N^2$ law, passage times predicted in advance from fold normal forms came within one iterate of the return map in 42 of 48 cases at four new folds, over two frozen stages (24 and 18 of 24) (the six misses, at a period-6 fold, 3.1 to 5.2 % short), and interval arithmetic proves an attracting period-8 orbit, two coexisting attracting orbits of periods 2 and 3 at one parameter cell, and the attracting 3-cycle at the start of a period-doubling cascade. In a smaller period-18 window, retaining about 0.645 % of the previous feedback state as a third control changes the selected return pair's trace-saddle height through a numerical sign bracket as its center is followed in steering and temperature, supporting a conditional local reconnection interpretation. The sampled centers keep their state period and six-token readout, and the multipliers remain nonzero. In the recurrent-depth model Huginn, a frozen test confirmed that over-relaxing its loop costs accuracy. At one post-fold setting of Qwen2.5-0.5B-Instruct, natural exit states admit an accurate one-coordinate reconstruction, while the full network still computes their global returns and some repeated passages end in period-3 capture. The loop is not ordinary sampled generation. These results show that bifurcation theory can make quantitative predictions about continuous language-model feedback, with rigorous guarantees for selected periodic orbits of the full transformer.

1 Introduction

Several recent inference methods feed a language model its own continuous output. Soft Thinking replaces each sampled token by the expected embedding under the next-token distribution [1] (its published implementation first keeps only the most probable tokens and renormalizes; the map studied here uses the full distribution, which keeps it smooth, and Appendix G reports an exploratory test of the truncated loop), latent-reasoning methods feed hidden states back [2], and recurrent-depth models iterate a block toward a fixed point [3]. Each turns inference into an iterated map, and an iterated map can be chaotic. We show that the expected-embedding loop of released language models is chaotic, in the finite-horizon sense defined below, over wide ranges of temperature and steering strength, trace where the growth that steering adds comes from, and measure what the chaos does: in a chaotic regime an irrelevant change of batch layout can change what the loop writes (in Qwen2.5-0.5B it changed the text of 173 of 192 chaotic float64 loops, but of only 7 of 192 periodic ones). In bfloat16, where nearly every loop diverges, the batch layout changed the answers of a short soft-thinking pipeline on 56 % and 55 % of 96 arithmetic problems in two models, at temperatures where the pipeline answers poorly; where it answers well, no answer changed. The same loops show quantitative signatures of low-dimensional dynamics, with late doubling ratios that approach Feigenbaum's constant numerically. In short, fed back through its own soft next-token distribution, a pretrained language model defines a discrete-time map with continuous states whose dynamics (periodic windows, period-doubling cascades, chaotic regimes and shrimp-shaped windows, the "shrimps" named by Gallas [4]) can be mapped, predicted from derivatives at selected transitions and in places proved; their consequences for inference depend on the feedback architecture and the numerical regime.

Three consequences follow for the use of such loops. The first concerns reproducibility: in a chaotic loop, differences far below anything a user controls, such as the batch layout or the numerical

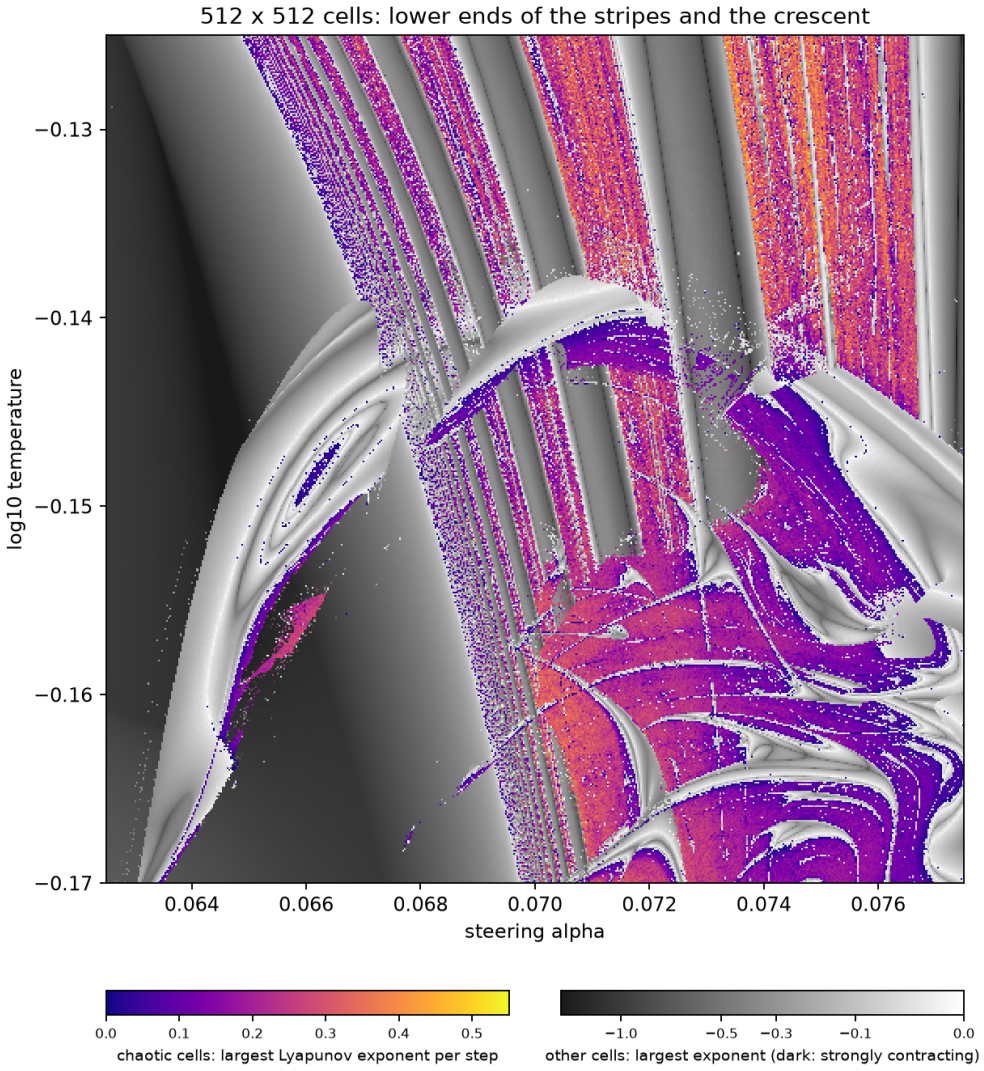

Figure 1: Chaos and repeating text in a language model fed its own output: the soft-token map of Qwen2.5-0.5B-Instruct, steered along the philosophy direction at layer 12, on 512×512 float64 cells over the steering strength $\alpha \in [0.0625, 0.0775]$ and the temperature, $\log_{10} T \in [-0.17, -0.125]$. Colored cells are chaotic by the finite-horizon criterion of Appendix B, and the color gives their largest Lyapunov exponent (left color bar). In gray cells the model settles into a repeating loop of text, shaded by its largest exponent, which measures how strongly the loop contracts (right color bar; light gray: close to losing stability); 4.4 % of the cells, also gray, have neither a detected period nor an exponent above 10^{-3}. Figure 11 names the loops.

precision, can change the output: a chaotic loop amplifies them, and long low-precision feedback can carry them into answers even where injected float64 perturbations shrink over the observed run. In an exploratory test replacing the target's cache and logits by those of a batch-1 reference run before every soft step removed nearly all such changes (Section 6). Ordinary numerical nondeterminism can thus become irreproducibility of the output, because the feedback amplifies or carries it. The second concerns steering: practitioners choose steering strengths, and here the strength is a bifurcation parameter, so near a boundary a small change switches the attractor or brings chaos, and a smooth dose-response cannot be assumed (Section 17). The third concerns prediction: where the loop is periodic, normal forms computed from the model's derivatives predicted passage times fixed before

they were measured: at four new folds the normal form's passage integral (the secondary predictor of the first test, the primary of the second) came within one iterate of 42 of 48 passages, and the other six were 3.1 to 5.2 % short (Section 8), so the quantitative tools of dynamical systems apply to the inference loops of released models, and larger models do not simply leave the chaos behind (Section 14). We also report a negative result: beyond the model's first-token uncertainty, the loop's regime carries no detectable information about the accuracy of ordinary sampled answers.

Figure 1 shows a small piece of the parameter plane of one such map, built from Qwen2.5-0.5B-Instruct. Cells with a positive finite-horizon Lyapunov exponent lie side by side with windows in which the model repeats a fixed loop of text. Parts of this picture have been seen before. Recent work has measured chaos-like sensitivity across the layers of a single forward pass [5, 6] and in language-model generation [7], temperature-driven changes of output statistics [8–12], periodic attractors of paraphrasing loops [13], attractor states of model-to-model conversations [14], and, through their Jacobians, the stability of looped and recurrent-depth models [15–18], and finite-depth readout dynamics have been used to predict the answers of looped models [19]. Besides Soft Thinking, probability-weighted vocabulary embeddings are fed back in training [20]. In a small task-trained looped transformer, fold coefficients computed from derivatives predict passage times and spectral gaps without fitting [21].

We study the expected-embedding ("soft-token") map of released pretrained language models, defined exactly by their stored weights and run with a fixed chat prefix at a fixed position. We follow it in fifteen models, from Qwen2.5-0.5B to Llama-3.3-70B and Qwen2.5-72B (three dense families and a mixture of experts, with three further models as controls for the latter), in most detail in Qwen2.5-0.5B-Instruct, and in two variations: two models writing one text, and a recurrent-depth model that iterates its own block. We chose these families because each releases open weights at several sizes, so that size can be varied within a family, and because they come from three developers with different training data and recipes; OLMo-2 also releases its training data, and Llama-3 has public sparse autoencoders, which we use for feature controls. The map is smooth. Its Jacobians come from automatic differentiation, and its periodic orbits can be computed and continued like those of any smooth map. Chaos and periodic windows occur in many nonlinear feedback systems. Random-network theory, in which a specified ensemble of recurrent networks becomes chaotic above a critical gain, in continuous time [22] and in discrete time by a quasi-periodic route [23], and the period-doubling cascades and shrimp-shaped domains of low-dimensional maps [4, 24], and the cascade to chaos of fixed-step gradient descent, studied as a map rather than through its gradient-flow limit, in exactly solvable models of deep networks, where the edge of stability is the first bifurcation of the training map [25], motivate looking for such structure here, but they do not show that a trained network realizes it. We test this directly. We locate chaos and periodic windows in the models' parameter space, read the texts the loop writes there, and check whether bifurcation theory, applied to derivatives computed from the model, predicts its behavior before the behavior is measured. The models realize this structure, and the predictions largely hold. We also show what the loop's derivative says about where chaos can occur and where the growth that steering adds comes from (Section 3; Section 12 assembles the pieces into one picture), and that larger models do not simply leave it behind (Section 14). Ordinary sampled generation with a growing context is a different system; Section 13 also measures how much of the mechanism survives sampling.

Terms and scope. Throughout, *chaotic* is a numerical label: after a transient, the largest finite-time Lyapunov exponent exceeds 10^{-3} and no period up to a cutoff is detected (Appendix B). It is evidence of finite-time exponential separation along the measured trajectories, not a proof of an asymptotic chaotic set, and some labeled cells are long transients that a periodic orbit eventually

captures: in random samples, 74 and 92 % of the labels of two maps persist to 10^5 steps (Section 5). We call the computed exponents finite-time exponents and reserve "chaotic attractor" and "chaotic set" for cases with stronger evidence. Where we draw such sets (Figure 14) we call them numerically observed attractors: sets that the tested trajectories approach and then stay near over the recorded horizon. Their invariance and attraction are finite-time numerical evidence, not certified; the interval-certified cycles are attractors without qualification. The one such set that we followed far longer, that of Figure 14(a), proved to be a long-lived chaotic transient: most trajectories left it within 10^5 steps (Section 5). The loop itself is the soft-token feedback map, iterated in discrete time, not a Poincaré return map; by a return map of a scalar observable s we mean the graph of s_{n+k} against s_n. Each claim is stated for the system on which it was measured. The autonomous soft-token map (a fixed prefix and one fed-back token), its versions with a longer fed-back memory, growing-context Soft Thinking, sampled thinking with common random numbers, ordinary generation and the recurrent block of Huginn are different systems, and results do not transfer between them without a test. Table 1 lists the main claims with the kind of evidence behind each.

Contributions. The results fall into four groups.

1. *Chaos in released models, and what it does to inference* (Sections 5 and 6; Appendix H). Temperature and steering move the loop between repeating text and chaos with embedded windows in all three model families, Qwen2.5, OLMo-2 and Llama-3. In float64, whether rounding differences between runs grow depends on the regime. Run twice with only the batch layout changed, 190 of 192 chaotic float64 loops of Qwen2.5-0.5B grow past a relative distance of 10^{-2} after a median of 102 steps, in line with the Lyapunov exponent (median ratio 0.91), and at 173 of them the most probable readout token comes to differ; 182 of 192 periodic loops stay within 10^{-2}, and at only 7 of them does the most probable token come to differ. Qwen2.5-7B and Llama-3.1-8B without steering (24 of 24 and 20 of 20 chaotic loops diverge, 0 of 24 and 1 of 20 periodic ones) and Llama-3.3-70B with steering (24 of 24, none of 24) behave alike. In bfloat16 nearly every loop diverges, and in a frozen test on 96 additions the batch layout changed the answer of a soft-thinking pipeline in 56 % of problems in Qwen2.5-0.5B and 55 % in Qwen2.5-7B, against never and 10 % for a forced greedy baseline. An exploratory causal test shows that the loop carries this: replacing the target's cache and logits by those of a batch-1 reference run before every soft step removes nearly all of it (on an H100, 47 against 0.8 % of answers changed in Qwen2.5-0.5B, and 41 % against none in Qwen2.5-7B), and accuracy changes by at most 5 points. These changes occur where the pipeline answers poorly (23 of 96 problems correct in the 0.5B model at $T = 1$); at $T = 10^{0.4}$, where both models answer 90 to 96 of the 96 correctly, the layout changed no answer. In a test frozen before its data, Soft Thinking with its published settings changed the answers to 21 % of 256 GSM8K problems not used before in this study in Qwen2.5-7B with the batch layout alone, on an H100, mostly between wrong answers; replaying the fed-back tokens cut this to 6 %; with reasoning models that answer 86 to 90 % of the problems correctly through soft thinking (DeepSeek-R1-Distill-Qwen-7B and QwQ-32B), 14.5 and 19.5 % of the soft-thinking answers changed, and a similar share with discrete sampled thinking. A chaotic loop's text reuses the phrases of neighboring loops in no fixed order, and a copy started 10^{-10} away writes the same text for 9 to 150 steps, then a different one (Table 3). Exactly repeating text can hide a state with a positive finite-time exponent: the text records only the most probable token, and the same token can win at every step while the probabilities behind it, which the loop feeds back, do not repeat over the run (Section 10). Beyond the model's first-token uncertainty, the loop's

Table 1: The main claims and the kind of evidence behind each. Identity: an analytic identity, checked numerically. Proof: computer-assisted, by interval arithmetic. Numerical: float64 computation with independent checks (tangent exponents, Newton residuals, continuation gates, longer horizons). Frozen: a test whose design, sample and criterion were fixed and pushed before its data. Exploratory: found by search or by an analysis designed after seeing data; structures selected this way are not estimates of how common they are. Details, amendments and failures are in Appendix L.

Claim	Evidence	Where
The loop's derivative factorizes into a readout covariance (a cross-covariance for untied embeddings) and the network's Jacobian	identity	Sec. 3
Chaotic regimes with periodic windows; chaos in 13 of 15 models, 11 of them without steering; 74–92 % of sampled labels last 10^5 steps	numerical	Secs. 5, 14
Chaotic float64 loops separate on the Lyapunov schedule when only the batch layout changes	numerical	Sec. 6
With fixed matrix shapes the runs are bitwise identical; injected displacements grow on the Lyapunov schedule (Qwen2.5-0.5B, 448 cells; a secondary criterion failed)	frozen	Sec. 6
The batch layout changes long-feedback answers (bfloat16 pipelines, Soft Thinking, reasoning models)	frozen	Sec. 6
Resetting the fed-back state removes most of these changes	exploratory	Sec. 6
Along recorded orbits, the growth that steering adds comes mainly from the readout, where next tokens compete (0.5B model; weaker evidence at 7B; a frozen sufficiency criterion failed by overshoot)	interventions; partly frozen	Sec. 3
Fold normal forms predict passage times (42 of 48; both all-row criteria failed)	frozen	Sec. 8
Late doubling ratios approach Feigenbaum's δ, at the float64 floor	numerical	Sec. 9
13 of the 17 tangent reversals of the continued fold curves are located cusps (four found on 18 October are not located); local unfoldings match Newton solves	numerical; unfoldings frozen	Sec. 11
A period-5 window is a shrimp in a two-parameter plane (boundary organization)	numerical; found by search	Sec. 11
Retained feedback changes a period-18 family's trace-saddle height sign, supporting conditional contour reconnection while the sampled centers keep their period and readout	exploratory; numerical	Sec. 11.1
An attracting period-8 orbit exists	proof	App. L
Attracting orbits of periods 2 and 3 coexist at one parameter cell, with certified token loops and basin neighborhoods	proof	App. L
An attracting 3-cycle exists where the crescent's period-doubling cascade begins	proof	App. L
No accuracy information beyond first-token uncertainty; no edge-of-chaos advantage	frozen; bounded nulls	Sec. 6
Huginn: over-relaxation costs accuracy	frozen	Sec. 15

regime carries no detectable information about the accuracy of ordinary sampled answers, and within the precision of the tested designs we detect no advantage at the edge of chaos.

2. *Why the loop can be chaotic, and where the growth that steering adds comes from* (Section 3). The loop's derivative factorizes into the inverse temperature, the covariance of token embeddings under the model's own next-token distribution (a cross-covariance for untied embeddings), and the network's Jacobian; with tied embeddings and without the network the map would be a modern Hopfield update, which cannot cycle. The factorization proves contraction at sufficiently high temperature and, at fixed steering strength and sufficiently low temperature, excludes chaos on invariant sets whose margin between the two leading tokens stays bounded away from zero as the temperature falls; the zero-temperature limit is eventually periodic, and chaos in our sweeps occupies intermediate temperatures. Stability belongs to the loop, not to its steps: along a period-8 orbit every step expands some direction while the cycle contracts. Along the recorded orbits, the growth that steering adds comes mainly from the readout, where next tokens compete. Steering into chaos raises the readout factor of each step's growth by 2.58 nats per step relative to nearly unsteered cells, while the network factor changes by only -0.07 (interval $[-0.23, 0.12]$); a logit bias alone, with the network unsteered, makes 3.5 % of its grid chaotic; and an orbit-preserving crossover, which apportions the growth along the recorded orbits, assigns it to steering's change of the readout ($+0.81$ nats per step in the 0.5B model), not of the network, whose change damps it (-0.49). Random directions also steer into chaos. Earlier steering carries chaos to lower temperature, in Qwen2.5-0.5B and Llama-3.3-70B, and two vectors at two depths create chaos where neither alone does (Section 6).

3. *Low-dimensional organization, predicted from derivatives* (Sections 7 to 11). The windows carry concrete texts: a period-adding family gains one "is X" pair per period step. A fold's normal form, computed from derivatives, predicts without fitting how long trajectories linger near the vanished orbit: in two frozen tests at four new folds, 42 of 48 forecasts fell within one iterate, and both all-row criteria failed (the second because passages from the entrance of a period-6 fold's channel came in 3.1 to 5.2 % short), and windows accumulate on the fold with distances close to a $1/N^2$ law (log-log slope -2.00; fitted with an offset, the free exponent ranges from 2.13 to 2.45 with the fit range). On the eleven fold curves we continued, 13 of the 17 tangent reversals are located cusps (the four found on 18 October are not located), and near them local unfoldings, frozen before solving, matched the orbits that Newton solves recovered at twelve parameter offsets in number and at eleven in stability. A doubling cascade has superstable points through period 768 and finite-order Feigenbaum ratios that approach δ numerically (the last, 4.6688, against 4.6692, at the float64 floor of the computation). A period-5 window of a plane steered at two depths is a numerically verified shrimp: followed cell by cell, its attractor fills a main cell with legs and doubling shells bounded by fold and flip curves, a trace-zero loop replaces the head of a one-dimensional shrimp, and twice the fold curve has two cusps of opposite kinds and a self-crossing. Around one of these junctions an orbit census fixed in advance matched the predicted counts at 9 of 13 points and found two coexisting attracting orbits at all five points inside it. An interval-arithmetic proof shows that a primitive period-8 orbit exists and is locally attracting, to our knowledge the first proof of a periodic orbit of a full transformer language-model map. A second proof of the same kind shows that at one parameter cell a period-2 and a period-3 orbit coexist, each unique in its box and locally attracting, each with a certified, if tiny, neighborhood of its basin: two different attracting behaviors under identical settings. A third proves the attracting 3-cycle at

the start of the crescent's period-doubling cascade (Appendix L).

4. *Beyond the one-token loop* (Sections 13 to 16). The chaos is not an artifact of the one-token window: with a memory of 8 or 16 soft tokens Qwen2.5-7B is chaotic at more temperatures than with one, and in Soft Thinking's growing context a tiny early perturbation grows exponentially inside its one-token chaotic band, though not in every model. The organization recurs across families, and size does not set the chaos: across fifteen models from 0.5B to 72B, a larger model's chaos can move to higher temperature instead of vanishing. Resolved through the stack, the network damps a change of the fed-back token more strongly in the larger Qwen2.5 and Llama-3 models, because their blocks write a growing residual stream that depends little on the fed-back token, so a change of that token becomes an ever smaller fraction of the state; a mixture of experts with 3B of its 35B parameters active per token damps more than a dense 9B model of its family. Two loops writing one text lock each other into periodic orbits: 60 % of pairs of differently steered 0.5B agents that are each chaotic alone are periodic together (twelve such orbits verified by Newton's method). In a second system, the recurrent-depth model Huginn, a spectral estimate ranks held-out questions by the size of their residual under relaxation, and a frozen generation test confirms, as stated before its data, an accuracy loss at relaxation 1.20 and noninferiority at 1.05.

Related work. Chaos-like sensitivity has been measured across the layers of a single forward pass [5, 6] and in generation [7, 26]. Temperature-driven changes of output statistics [8–11], repetition loops [27, 28] and behavioral attractors of paraphrasing and conversation [13, 14, 29] have been measured in generated text, and one study fits a noisy effective map to a hidden-state coordinate [12], but none of these computes the orbits, multipliers or bifurcations of the model's own map. Equilibrium, looped and recurrent-depth models have been analyzed through fixed points and Jacobians [3, 15–18], and the closest methodological precedent, a derivative-based fold analysis of a small task-trained looped transformer, predicts passage times and spectral gaps without fitting and derives from the same bottleneck an N^{-2} tail of relaxation times over inputs [21]. We know of no earlier prospective test of quantitative predictions from a pretrained language model's own derivatives (here passage times at four new folds, where 42 of 48 forecasts fell within one iterate and both all-row criteria failed, by one row and at the entrance passages of a period-6 fold, and the orbits recovered and their stabilities at frozen offsets near cusps), of no earlier frozen held-out test of a spectral prediction of a pretrained language model's response to relaxation, which we give for Huginn, and of no earlier computer-assisted proof of a periodic orbit, or of coexisting attractors, of a full transformer language-model map. These statements rest on targeted searches (12 and 15 October 2026) that found no prior work combining these elements for the feedback map of a pretrained language model: an absence of evidence within the search, not a proof of priority. Appendix A discusses this work, expected-embedding feedback, and bifurcations, Feigenbaum ratios and shrimps in neural models in detail.

2 The inference map

Definition. We describe the map for Qwen2.5-0.5B-Instruct, the model studied in most detail; the other models enter the same construction with their own embeddings and steering layers (Appendix H). Let $E \in \mathbb{R}^{V \times d}$ be the input embedding matrix ($V = 151{,}936$, $d = 896$). The model reads a fixed chat prefix, with the system prompt "You are an AI talking with another AI. You have complete freedom. Feel free to pursue whatever you want." and the user turn "Hello!", followed by

a single input embedding x, and produces next-token logits $z(x) \in \mathbb{R}^V$. The soft-token map feeds the expected next-token embedding back in place of x:

$$F_{T,\alpha}(x) = E^\top \operatorname{softmax}(z_\alpha(x)/T), \tag{1}$$

where $z_\alpha(x)$ are the logits computed with αv added to the residual stream at layer 12 at every position, prefix included; the steering vector v is described in the next paragraph. Positions are fixed and the window does not grow, so F is an autonomous map of $\mathbb{R}^{896}$. As $T \to 0$ it reduces, where the argmax is unique, to greedy decoding of the same fixed-prefix, one-token-window system. It is not ordinary generation with a growing context.

The prompt is a control, not just an initial condition. The dependence on the fixed prefix is suppressed in (1). Writing it explicitly gives

$$x_{n+1} = F_{P,T,\alpha}(x_n) = E^\top \operatorname{softmax}(z_\alpha([P; x_n])/T)\,, \tag{2}$$

where $P = (e_{t_1}, \ldots, e_{t_L})$ is the ordered sequence of prefix embeddings and $[P; x_n]$ denotes sequence concatenation, not vector addition. At fixed P, T and α, the same map is applied at every iteration. Every feedback embedding can attend to the entire prefix. Causal attention makes the prefix keys and values independent of x_n: we compute them once for each steering setting and restore the cache to that prefix after every update. The previous feedback position is replaced, not retained as an additional history token.

Changing x_0 selects a trajectory of one map; changing P changes the map itself, including its state derivative and potentially its attractors and their basins. Prompt choice is therefore another control knob, alongside temperature and steering, even though it is held constant within each trajectory. Ordinary token edits are discrete changes, unlike the continuously varied scalar parameters T and α; calling the prompt a control does not by itself define a smooth bifurcation parameter.

The prefix is naturally an $L \times d$ array of generally dense embeddings, not a single sparse vector whose dimension equals the maximum context length. One can flatten it into Ld coordinates, or pad it to a capacity-sized array together with an attention mask, but unused positions are absent or masked rather than active zero-valued inputs. The context limit bounds the number of token positions, not the embedding dimension. For the one-token map, increasing L changes the fixed conditioning and the feedback position while leaving the evolving state $x_n \in \mathbb{R}^{896}$ the same size. A maximum-context prefix leaves one position for feedback; successive iterations do not consume further context slots.

A fixed prompt supplies a state-dependent response. The fast-weight view of attention interprets input-dependent keys and values as temporary computational weights, distinct from the trained parameters [30]. For a fixed prefix bank in one head, with keys and values stored as columns, the normalized retrieval operation is

$$A_P(q) = V_P \operatorname{softmax}(K_P^\top q/\sqrt{d_k}), \qquad D_q A_P = \frac{1}{\sqrt{d_k}} V_P(\operatorname{diag} a - aa^\top) K_P^\top, \tag{3}$$

where $a = \operatorname{softmax}(K_P^\top q/\sqrt{d_k})$. This is a schematic prefix-only attention operation: the actual feedback token also contributes a state-dependent key and value to each head's joint attention. Even with K_P, V_P fixed, retrieval varies with the query and its derivative varies with both the query and the prefix. The prompt therefore does not act merely as an additive constant. Through the

layers, it can change the state derivative $D_x F_{P,T,\alpha}$ that governs local amplification and stability. The softmax normalization remains essential; the fast-weight interpretation does not turn attention into a linear map or an ordinary fixed-weight MLP.

Wen et al. use this perspective to develop efficient attention and evaluate their architecture on vision tasks; they do not test the feedback bifurcations studied here. Our connection is an application of their viewpoint to a fixed-prefix recursive system. Nor does additional prefix length guarantee stronger feedback: more key–value entries may be redundant, and their influence depends on the normalized attention weights. Prompt-induced changes in stability must be measured from the complete feedback map, including the output softmax and embedding readout.

Prompt changes the rule, its stability, and the attractor (PCP1). An exploratory intervention isolates prompt content from prompt length. In Qwen2.5-0.5B we held $T = 2$, $\alpha = 0$, all 37 prefix positions, the system instruction and initial feedback embedding fixed. We changed only the token `Hello` in `Hello!` to `Think`, following $P(s) = (1 - s)P_0 + sP_1$ at nine values $s = 0, 1/8, \ldots, 1$. The endpoints are ordinary prompts; the intermediate points are soft embeddings, not discrete text. This is a content intervention, separate from the maximum-context scan.

First, at the *same* feedback state and in the same fixed unit direction u, the endpoint change in $F_P(x)$ was 66.12 % of its original norm, and the change in $D_x F_P(x)u$ was 111.4 % of its original norm. The latter includes rotation of the response, not just its magnitude; it is not a change in the full Jacobian operator norm. The ratio of response norms along this one direction was 0.324. This common-state directional response is not the dominant multiplier at the continued fixed point: a smaller response in this direction does not imply stability at another state. Prefix key–value banks were rebuilt at every setting. The prompt therefore changes how a small feedback perturbation propagates, even at a common state.

Second, fixed-point continuation found a stability change. The leading multiplier was real -0.943697 at $s = 0.625$ and -1.204287 at $s = 0.75$. Four further bisections localized the real-minus-one stability boundary to $s \in [0.6406250, 0.6484375]$. Every refinement recorded the multiplier's modulus and argument and the next-largest modulus; the latter remained below one. This is a numerical flip-type stability crossing, not a Neimark–Sacker crossing. We have not computed a flip normal-form coefficient or supplied interval-certified existence, transversality or nondegeneracy bounds.

Third, fresh trajectories at $s = 0.75$, 0.875 and 1 settled into *numerically stable primitive period-two cycles*. Each used 512 settling updates followed by 256 saved full-state updates. We tested lags 1–64 over 128 positions, checked composed-map closure below 10^{-10} relative error, and rejected period one with separation above 10^{-8}. The estimated largest two-step Floquet multiplier moduli were respectively 0.577023, 0.672266 and 0.690769, all below one. These are multipliers of the *two-step* map. The corresponding largest cycle Lyapunov estimates are $\lambda_{\max} = \frac{1}{2} \log |\mu_{\max}|$, namely -0.274936, -0.198550 and -0.184975 per feedback update. These are Floquet-derived asymptotic rates, not finite-window twin measurements. The fixed point's unstable one-step multiplier does not describe the stability of the new cycle. This gives a concrete prompt-controlled change from a fixed point to alternation, *not evidence of chaos* at the three tested post-crossing settings. It also does not prove that this cycle connects continuously to the fixed point at the bracketed boundary; that stronger branch statement would require cycle continuation.

All calculations used effective float64 normalization and rotary operations. One-step and composed-map JVPs were checked against successively finer central differences, and leading multipliers were estimated from two independent Arnoldi starts with residual gates. Initial coarse finite-difference checks failed and were retained; smaller steps resolved truncation error. Plain fixed-point it-

eration subsequently failed at $s = 0.75$, as expected when a root loses attraction. A disclosed Newton–Krylov amendment recovered the unstable fixed point rather than treating nonconvergence as its absence. The raw evidence, unsuccessful attempts and a separately hash-verified disk backup are bound by `docs/prompt-control-receipt.json` and `docs/prompt-control-followup-receipt.json`. These numerical tests are not rigorous spectral enclosures, global basin guarantees, or a general claim about the meanings of the words `Hello` and `Think`.

The common-state comparison is important: changing an additive constant alone in $F_P(x) = G(x) + b(P)$ would change $F_P(x)$ but leave $D_x F_P(x) = DG(x)$ unchanged at fixed x. A measured change in the common-state JVP therefore distinguishes prompt-dependent feedback gain from that restricted additive-forcing model. It does not identify which head or layer causes the change; that would require controlled interventions on their key–value banks. Comparing derivatives only at different attracting states would not make this distinction, since $DG(x)$ can vary with x even when the map changes only through an additive term.

A smooth prompt path also provides a continuation parameter. If $F_{P(s),T,\alpha}(x_*(s)) = x_*(s)$ and $I - D_x F$ is invertible, implicit differentiation gives

$$\frac{dx_*}{ds} = (I - D_x F)^{-1} \partial_s F. \tag{4}$$

This identity separates how the prompt perturbs the update from how the loop responds to that perturbation. Its inverse can be large for non-normal maps, so eigenvalue distance alone is not a quantitative sensitivity bound. The real-minus-one bracket above supplies numerical evidence for a local stability crossing. Changing the derivative, or moving a fixed point, alone would not establish that crossing or the nature of the resulting attractor.

How we steer. Steering adds one fixed vector to the network's internal state while it runs; no weight is changed. The state it changes is the residual stream, the d-dimensional vector carried at each token position, to which every attention and MLP block adds its output. Our main direction, philosophy, is built in three steps from 16 short texts (`tools/explore_llm_phase_map.py`, `CONCEPTS` and `NEUTRAL`): eight philosophical sentences, such as "What is the nature of consciousness and subjective experience?" and "Is there something it is like to be me?", and eight neutral ones, such as "The meeting is scheduled for nine in the morning." and "The package weighs two kilograms." First, each text is run through the unsteered model, without the chat template, and the residual-stream vector after decoder layer 12 (of 24, counted from 0) is read at its last token. Second, the direction is the mean of these vectors over the philosophical texts minus their mean over the neutral texts. Third, it is rescaled to the mean residual norm at that layer over the 16 texts, $\|v\| = 189$, so that α is the push as a fraction of a typical residual vector: $\alpha = 0.07$ adds 7 % of one. The raw difference has norm 8.1, so $\alpha \approx 0.043$ reproduces the natural contrast between the two groups of texts, and the values of Figure 11, 0.0625 to 0.0775, push 1.5 to 1.8 times as far. In every forward pass of the loop, αv is added to the residual stream after layer 12 at every position: at each token of the prompt and at the fed-back soft token. The later layers therefore also see a steered prompt, since the keys and values they compute for its tokens carry the push (Figure 2). The other concept directions (gratitude, hostility, kindness and cooking, each with eight sentences) are built the same way at their own layers, and the specificity controls use norm-matched random directions, single MLP neurons and sparse-autoencoder features (Appendices J and I).

State space and regularity. Attention is causal, so the keys and values of the prefix do not depend on x, and F is an autonomous map of $\mathbb{R}^{896}$. For the dense models studied here it is real-analytic, because RMS normalization (with its positive regularizing constant), the SiLU activation,

attention, rotary embeddings at fixed positions and the softmax readout are analytic; with top-k routing, a mixture of experts is smooth only within each region of fixed routing, and its map can be discontinuous where the routing changes. The continuity and differentiability arguments below apply to the dense models. Every soft token is a convex combination of embeddings, so F maps $\mathbb{R}^{896}$ into the compact convex polytope $K = \mathrm{conv}\{e_t\}$, and $F(K) \subset K$. Every orbit therefore enters K after one step, F has a fixed point in K by Brouwer's theorem, and, being continuous on a compact set, it has invariant probability measures there. Because $\log^+\|DF\|$ is bounded on K, Oseledets' multiplicative ergodic theorem applies, in its version for non-invertible maps, which allows exponents equal to $-\infty$: for every invariant measure the Lyapunov exponents exist at almost every point [31, 32]. The exponents we report estimate them over finite horizons (Appendix B).

Why feed a model its own output? Feedback is not an addition to language-model inference; it is how text is generated. Autoregressive decoding appends each chosen token to the input and conditions the next step on it, and as $T \to 0$ the map studied here becomes greedy decoding wherever the most probable token is unique; the familiar repetition loops of greedy decoding are then the cycles of the limit. What the loop changes is what is fed back: the expected embedding instead of one chosen token. That is the step of Soft Thinking [1], and related methods feed back hidden states [2] or iterate a block to a fixed point [3]; continuous feedback of this kind is studied because it can carry more than one token's worth of information. We also hold the prefix and the window fixed. That simplification is deliberate: it makes the loop an autonomous map whose orbits, exponents and bifurcations can be computed, and in a few cases proved, much as low-dimensional maps made the routes to turbulence tractable. This fixed-prefix, fixed-window soft-token map differs both from ordinary autoregressive generation, whose context grows, and from recurrent-depth (looped) transformers, which repeatedly update hidden states through a shared block [3, 33, 34]; dynamical methods apply to both kinds of loop, but their update operators, and therefore their derivatives and stability conditions, differ. We then test which effects survive in the fuller systems: Soft Thinking's growing context, sampled decoding, longer fed-back memories and two models writing to each other (Appendices G and 16). The answer-level consequences reported below were measured in a soft-thinking pipeline with a growing context, not in the reduced loop. Deployed systems feed outputs back as well, when agents read their own earlier steps and when pipelines chain model calls.

Steering as a probe. Steering is an established way to control a language model without retraining: a vector added to the residual stream, typically the mean difference between activations on contrasting prompts, strengthens or weakens a behavior, such as sycophancy or refusal [35], sentiment or topic [36], or a personal style of writing [37], and representation engineering uses such directions to read and control concepts [38]. Our recipe is the standard contrastive one [35], added at every position, and our pushes are moderate: at most 12 % of the mean residual norm per vector in every map (the continued curves reach 13.1 %), and 6.25 to 7.75 % in Figure 1. Pushes of this size do change what a model writes, which is what steering is for: in Qwen2.5-7B, 1 % of the residual norm is the largest push that kept greedy accuracy on a multiplication task within 5 points (Appendix D). They are far smaller than activation additions reported as effective in steering work, one of which has nearly ten times the norm of the residual stream it is added to [36]. We do not steer to make the model behave in a particular way. We use steering as a probe: α is a control parameter, like the amplitude of a periodic forcing, that moves the model through a one-parameter family of nearby maps without retraining it, so that bifurcation curves can be traced. The chaos does not depend on it: at $\alpha = 0$, temperature alone gives chaos with periodic windows in Llama-3.1-8B, Qwen2.5-7B

and OLMo-2-7B, and in a narrow band in Qwen2.5-0.5B (Appendix J). Nor, in the 0.5B and 8B models, does the direction's meaning: random, sparse-autoencoder and single-neuron directions of the same norm make the map chaotic about as often as the concept directions (at 70B the directions differ; Appendix I).

Precision. Hugging Face RMSNorm layers and rotary-embedding tables compute in float32 even when the model is cast to float64, so we patch both. A float32 map is accurate to 3×10^{-6}–2×10^{-5} per step. That is enough for phase maps (245 of 256 test cells agree with float64) but not for derivatives of composed maps: one-step float32 Jacobian-vector products are accurate to 5×10^{-5}, yet the 8-step monodromy of a period-8 orbit is off by 3%. All orbits, multipliers and test functions below are therefore computed in float64.

3 Why the loop can be chaotic

Chaos in the sense used here, a positive largest Lyapunov exponent with no period found (the operational label of Appendix B), requires that nearby states drift apart: a small difference δx between two states must, on average, grow from one step to the next, which is what a positive largest Lyapunov exponent measures (such growth alone does not make an attractor chaotic). To first order, one step multiplies the difference by the derivative of the map, so the question is decided there. The soft-token map $F_{T,\alpha}$ of (1) has a simple derivative with respect to the state x. With $p = \mathrm{softmax}(z_\alpha(x)/T)$,

$$DF_{T,\alpha}(x) = \frac{1}{T} E^\top (\mathrm{diag}\, p - pp^\top) J_{z_\alpha}(x). \tag{5}$$

Qwen2.5-0.5B ties its output matrix to the input embedding, so $z_\alpha = E\tilde{h}_\alpha(x)$, where $\tilde{h}_\alpha$ is the final normalized hidden state of the steered network, and the derivative factorizes:

$$DF_{T,\alpha}(x) = \frac{1}{T} \Sigma_p \, J_{\tilde{h}_\alpha}(x), \qquad \Sigma_p = E^\top (\mathrm{diag}\, p - pp^\top) E = \mathrm{Cov}_{t\sim p}(e_t). \tag{6}$$

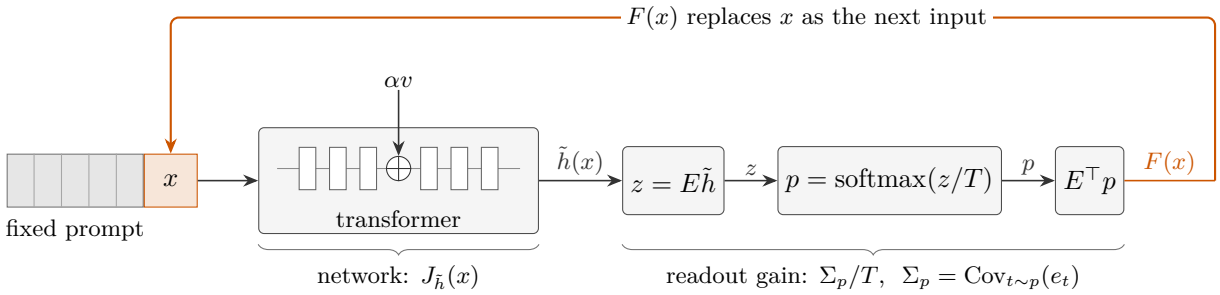

Figure 2: The soft-token loop. A fixed prompt is followed by one input embedding x. The transformer, with αv added to the residual stream at a middle layer (at every position), gives the final normalized hidden state $\tilde{h}(x)$ at the position of x; with tied embeddings the logits are $z = E\tilde{h}$ (untied: $z = W\tilde{h}$), and $p = \mathrm{softmax}(z/T)$ is the next-token distribution. The expected embedding $E^\top p = F(x)$ is fed back as the next input, at the same position after the same prompt. Below: the two factors of the derivative (6), the network's Jacobian $J_{\tilde{h}}$ and the readout gain Σ_p/T. Without the network, $\tilde{h}(x) = x$ and $J_{\tilde{h}} = I$, the loop $F(x) = E^\top \mathrm{softmax}(Ex/T)$ is the update of a modern Hopfield network [39], which lowers an energy at every step and so cannot cycle (Section 3).

The derivative of the loop (Figure 2) is the covariance of the token embeddings under the model's own next-token distribution, amplified by the inverse temperature and composed with the derivative of the network. The network factor $J_{\tilde{h}_\alpha}$ is the derivative of the whole transformer, with every attention layer, MLP block, normalization and residual connection and the steering included (attention depends on x through the query, key and value of the fed-back position; those of the fixed prefix do not change); the factorization only separates it from the readout. The covariance form is the standard identity for the Hessian of the log-partition function of an exponential family [40], the Jacobian of the softmax [41], and it appears explicitly in the Jacobian of the modern Hopfield update [39]; here it is composed with the derivative of the full transformer. For tied embeddings, placed between a network Jacobian and its transpose and divided by T^2, the same covariance is the Fisher information of the next-token distribution pulled back to the network's input (for untied embeddings the covariance of the output vectors takes its place), a metric used to attribute predictions to tokens and to describe steering geometrically [42, 43]; in the feedback map it enters once, multiplied by the Jacobian. Each factor has a plain meaning. The network Jacobian $J_{\tilde{h}_\alpha}$ says how much the transformer's final hidden state moves when its input moves, through all of its attention and MLP blocks. The factor Σ_p/T is the derivative of the fed-back expected embedding with respect to that hidden state. Its matrix Σ_p is the variance of the token embeddings under the next-token distribution: it is large when tokens with different embeddings share the probability, so that a small change of the logits moves the expected embedding far, and it tends to zero as one token takes all the probability. The scalar $1/T$ multiplies the whole derivative.

At a state of the period-8 orbit of Table 8, (6) agrees with forward-mode derivatives of the full map to a relative error of at most 1.35×10^{-13} over eight random directions (Figure 3; `tools/softmax_gain.py`). Without tied embeddings, as in the 7B and 8B models, the logits are $z_\alpha = W\tilde{h}_\alpha(x)$ with an output matrix $W \neq E$, and $DF_{T,\alpha} = T^{-1}E^\top(\operatorname{diag} p - pp^\top)WJ_{\tilde{h}_\alpha}(x)$: Σ_p becomes the cross-covariance $\operatorname{Cov}_{t\sim p}(e_t, w_t)$ of input and output embeddings, which is not symmetric in general. The temperature limits below do not change. The rest of this section uses the factorization in four ways: it shows where steering and temperature enter; it shows that, with tied embeddings, the transformer between input and output is what makes cycles and chaos possible at all (untied embeddings can cycle even without it); it proves that the loop contracts at sufficiently high temperature and, at fixed steering strength and sufficiently low temperature, excludes chaos on invariant sets whose margin between the two leading tokens stays bounded away from zero, while stability belongs to the composed loop rather than to single steps; and three measurements show that, along recorded orbits, the growth steering adds comes mainly from the readout, where next tokens compete.

Where steering and temperature enter. Temperature appears in (6) twice: explicitly, as the factor $1/T$, and inside p. As $T \to 0$ the two pull in opposite directions and p wins: where one token leads, p concentrates on it, Σ_p vanishes faster than $1/T$ grows, and the readout factor tends to zero, so the loop approaches greedy decoding; only near ties between leading tokens does the factor grow, like $1/T$ ("High and low temperature" below). In Backour's unbiased two-token reduction with linear logit transport and a positive constant coupling, the inverse temperature rescales the coupling, and crossing the critical value 2 destabilizes the balanced mixture, so that one token takes over [44]. Steering does not appear in (6): the vector αv is the same whatever the state x, so it has no derivative of its own, and it changes DF only by changing the quantities from which the two factors are computed. In the readout factor Σ_p/T, steering decides which tokens compete in p and temperature decides how sharply: Σ_p is large when the probability is spread over several tokens whose embeddings differ, and it tends to zero as the distribution concentrates on a single token. In

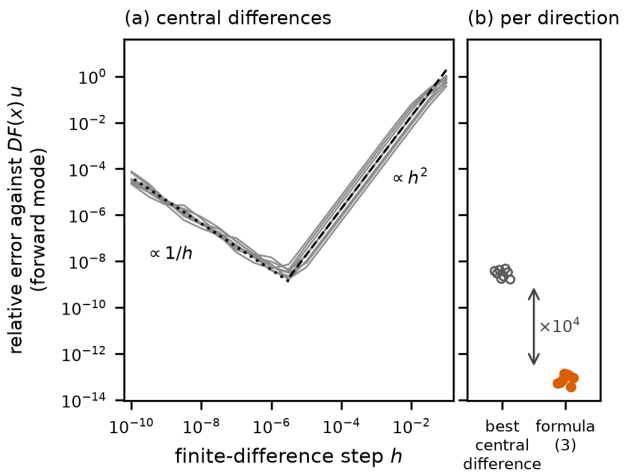

Figure 3: Formula (6) is exact up to rounding. At a state of the period-8 orbit of Table 8 (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, $\alpha = 0.071674$, $\log_{10} T = -0.115$, float64), in eight random unit directions u, each derivative is compared with the network's forward-mode derivative $DF(x)\,u$. (a) Central differences $(F(x+hu) - F(x-hu))/(2h)$. For large steps h their error is truncation, which falls like h^2 as h shrinks (dashed line, slope 2; fitted slopes 1.98 to 2.00 for $10^{-4} \le h \le 10^{-2}$); below $h \approx 3 \times 10^{-6}$ rounding takes over and the error grows like $1/h$ (dotted line). (b) Each direction's best central difference (open circles, 1.7×10^{-9} to 4.8×10^{-9}) and the factorized derivative $T^{-1}\Sigma_p J_{\tilde{h}}(x)\,u$ of formula (6) (filled, 3.8×10^{-14} to 1.35×10^{-13}), which has no step size: the formula agrees with forward mode more than four orders of magnitude more closely than any finite difference.

the network factor $J_{\tilde{h}_\alpha}$, steering changes the input of every later block, and with it the attention weights, MLP activations and normalization scales at which that block's derivative is taken, and so changes how strongly the network stretches a change of the fed-back token; temperature changes this factor too, because it moves the orbit along which the factor is evaluated. On the two main maps both parameters move both factors (orbit averages along the trajectories of Figure 5, power-iteration estimates). To measure how much each parameter matters, we asked how much of the variation of these orbit averages over a map steering strength explains on its own, as the share of the variance carried by the means over cells of equal α, and how much temperature explains, by the means over cells of equal T. For the network term $\langle \log \|J_{\tilde{h}_\alpha}\| \rangle$, steering explains 7 and 33 % on the stripes map and the crossing window and temperature 47 and 4 %; for the readout term $\langle \log(\|\Sigma_p\|/T) \rangle$, steering explains 30 and 32 % and temperature 21 and 27 %. Which of the two changes turns the loop chaotic is measured later in this section ("Steering makes tokens compete"). From here on we drop the subscripts of DF and $J_{\tilde{h}}$ when (T, α) is fixed. Every condition that involves them, the gain condition and the threshold T_* below among them, describes a region of the (α, T) plane, not a range of T alone.

A Hopfield network with a transformer inside. If the network passed its input straight through, $\tilde{h}(x) = x$, the map would be $F(x) = E^\top \mathrm{softmax}(Ex/T)$. This is exactly the update rule of a modern Hopfield network whose stored patterns are the token embeddings, at inverse temperature $\beta = 1/T$ [39]. That update lowers an energy function at every step and converges to fixed points. In (6) it is the case $J_{\tilde{h}} = I$, where $DF = \Sigma_p/T$ is symmetric and positive semidefinite. The energy argument covers every map with that property. If DF is symmetric and positive semidefinite everywhere, then $F = \nabla\phi$ for a convex ϕ (for the network-free map, $\phi(x) = T \log \sum_t e^{e_t^\top x/T}$), and

convexity, $\phi(F(x)) \geq \phi(x) + F(x)^\top (F(x) - x)$, gives for $V(x) = \frac{1}{2}\|x\|^2 - \phi(x)$

$$V(F(x)) \leq V(x) - \tfrac{1}{2}\|F(x) - x\|^2. \tag{7}$$

Each step lowers this energy by at least half the squared length of the step. Along a cycle these decreases sum to zero, so every point of the cycle is fixed. F maps every state into the convex hull of the embeddings, where V is bounded below, so the steps of every orbit must shrink to zero: every accumulation point of an orbit is a fixed point, and the strict descent of the energy rules out cycles of period above one and chaotic motion (an unstable fixed point can still have positive tangent exponents; and the descent estimate alone does not establish convergence to a single fixed point when fixed points are not isolated). In models with tied embeddings, such as Qwen2.5-0.5B, every cycle and every chaotic cell therefore depends on the network factor $J_{\tilde{h}}$ breaking this energy structure. With untied embeddings the network-free map $x \mapsto E^\top \mathrm{softmax}(Wx/T)$ has the derivative $\mathrm{Cov}_{t\sim p}(e_t, w_t)/T$, which need not be symmetric or positive semidefinite, so the energy argument no longer applies to those models, and cycles become possible even without a network (two tokens whose output vectors point at each other's input embeddings exchange at low temperature). Attention has itself been read as a Hopfield retrieval [39]; in the soft-token map the retrieval closes a loop around the whole network.

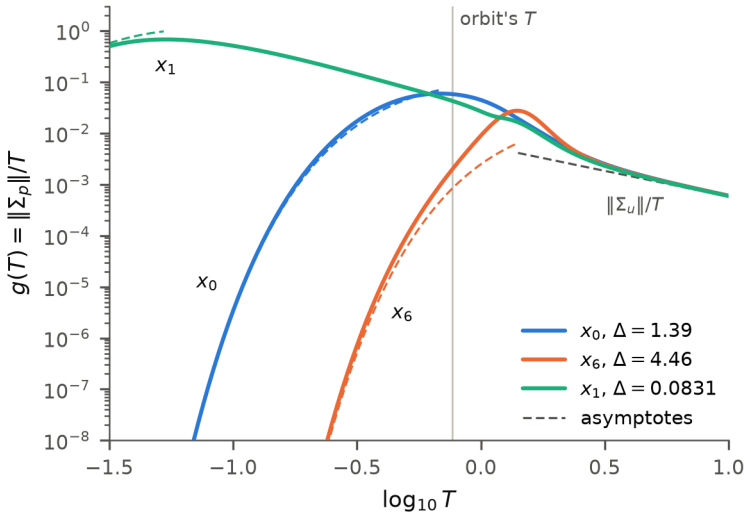

Figure 4: Readout gain $g(T) = \|\Sigma_p\|/T$ (spectral norm) at the fixed logits $z_\alpha(x)$ of three states $x_k = F^k(x_0)$ of the period-8 orbit of Table 8 (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, $\alpha = 0.071674$), with only the temperature in $p = \mathrm{softmax}(z/T)$ varied. x_0 is the state at which (6) was checked; x_6 and x_1 have the largest and the smallest gap Δ between the two largest logits among the eight states. Dashed: the low-temperature asymptote $\|e_{t_1} - e_{t_2}\|^2 e^{-\Delta/T}/T$ of each state (t_1, t_2 the two leading tokens), drawn below its peak, and the high-temperature asymptote $\|\Sigma_u\|/T$, with Σ_u the covariance of all token embeddings under equal weights. The peaks lie at $\log_{10} T = -1.27$ (x_1, $g = 0.68$), -0.16 (x_0, 0.059) and 0.15 (x_6, 0.028); the vertical line marks the orbit's temperature, $\log_{10} T = -0.115$. Where the two leading logits nearly tie (x_1), g grows roughly like $1/T$ as T falls toward Δ. This is the readout factor alone: the network factor $J_{\tilde{h}}$ is not included, so g does not decide whether DF expands. At the orbit's temperature $g \leq 0.064$ at all eight states, while the largest singular value of DF lies between 1.2 and 21 (Section 3).

High and low temperature. The largest Lyapunov exponent of an orbit, the average rate at which nearby states separate, cannot exceed the long-run average of $\log \|DF\|$ along the orbit, so a positive exponent requires that average to be positive (this gain condition is necessary, not sufficient). Both ends of the temperature axis work against it, in different ways (Figure 4).

At high temperature $1/T$ is small, and the bound is explicit. However the probability is spread, the projections $u^\top e_t$ of the embeddings on any unit vector u lie in an interval no longer than the diameter $\operatorname{diam}(E)$ of the set of embeddings, and Popoviciu's inequality bounds a variance by a quarter of the squared range, so $\|\Sigma_p\| \le \operatorname{diam}(E)^2/4$. With $L(\alpha) = \sup_K \|J_{\tilde{h}_\alpha}\|$, the network's largest gain on the set K of possible states (finite because the Jacobian is continuous and K is compact), $\|DF_{T,\alpha}\| \le \operatorname{diam}(E)^2 L(\alpha)/(4T)$ on K, and $F_{T,\alpha}$ shrinks every distance, that is, it is a contraction of K, for $T > T_*(\alpha) = \frac{1}{4}\operatorname{diam}(E)^2 L(\alpha)$. Above this temperature the loop forgets where it started: by Banach's fixed-point theorem it has a single fixed point, which attracts every orbit, and it tends to the mean embedding as $T \to \infty$; without tied embeddings, $\operatorname{diam}(E)^2$ becomes $\operatorname{diam}(E)\operatorname{diam}(W)$. For affine softmax feedback, with a fixed matrix in place of the network, a recent preprint proves a sharp spectral threshold for the uniqueness and global stability of the fixed point [45]. The threshold is a curve in the (α, T) plane, because the largest gain of the network depends on the steering. The bound is rigorous, but it does not locate the measured chaos boundary, because it assumes the worst case at every step. With the diameter of the 0.5B model's tied embeddings, $\operatorname{diam}(E) = 1.2109$, computed over all pairs of embeddings, and the largest norms of $J_{\tilde{h}_\alpha}$, the Jacobian of the final normalized hidden state, met along the trajectories of 7936 cells of the two main maps (power-iteration estimates, so lower estimates of $L(\alpha)$), $T_*(\alpha) \ge 2.7 \times 10^3$ at every steering strength sampled in these maps (64 per map), and $T_*(\alpha) \ge 5.6 \times 10^4$ at $\alpha = 0.06$, where the largest norm, 1.54×10^5, was met ($\log_{10} T_* \ge 3.44$ and ≥ 4.75): at least 3.1 and 4.4 decades above the highest chaotic temperature of the unsteered sweeps below ($\log_{10} T = 0.28$; `tools/norm_bounds.py`). It does not cover hard routing or a growing context.

At low temperature $1/T$ is large, but one token takes almost all the probability, and Σ_p vanishes faster than $1/T$ grows. At a state with a unique leading token j, the gain is $\|\Sigma_p\|/T = e^{-\Delta/T}(C+o(1))/T$, where Δ is the gap between the two largest logits and $C = \|\sum_i (e_i - e_j)(w_i - w_j)^\top\|$, the sum running over the runner-up tokens whose logits trail by exactly Δ. As $T \to 0$, F tends to the greedy map, which feeds back the embedding of the most probable token (at an exact tie, the limit averages the tied embeddings). The limit map takes finitely many values, so every orbit of it eventually repeats (Appendix G). The convergence is pointwise and in general not uniform: near a tie between tokens with different embeddings the gain is of order $1/T$. At fixed α and small positive temperature, chaos is therefore excluded only on invariant sets that keep a uniform logit gap, $\Delta \ge \Delta_0 > 0$ at every state of the set and for every small T; a gap that shrinks with T is not enough. Under that condition $\|\Sigma_p\| \le \operatorname{diam}(E)^2 (V-1) e^{-\Delta_0/T}$, so $\|DF\| \le L(\alpha)\operatorname{diam}(E)^2 (V-1) e^{-\Delta_0/T}/T$, which tends to zero (with tied embeddings; untied ones change only the constant). A chaotic set would have to return to near-ties again and again.

Chaos is therefore found in between, at temperatures where the next-token distribution is spread enough to make Σ_p large while $1/T$ still amplifies it, and the unsteered sweeps agree (Figure 93). In the four long sweeps (Appendix J) the chaotic temperatures lie in one band per model, interleaved with periodic windows, and every band lies between $\log_{10} T = -0.20$ and 0.28. This band is an observation; near ties, the low-temperature argument does not exclude chaos. Below the bands, all 1547 temperatures swept have negative exponents, and all but five have a detected period; above them, up to the end of the sweeps at $\log_{10} T = 0.3$, 127 of 128 temperatures have a detected period (fixed points at 101) and one has a negative exponent without one. These are finite sweep findings; the wider sweeps of Section 14 find chaos at higher temperatures in some models, and an OLMo-2-32B sweep found three candidate cells below one nat of entropy (at $\log_{10} T = -0.42$, -0.41 and -0.07; one twin offset). A further sweep of the unsteered 0.5B map from $\log_{10} T = 0.3$ to 1.5 finds an attracting fixed point at all 25 temperatures. Its entropy approaches $\log V$ (11.928 against 11.931 nats at $\log_{10} T = 1.5$), and its exponent falls steadily, by 2.14 per decade of temperature at

the top of the range, toward the slope $-\ln 10 \approx -2.30$ expected once p is uniform.

Stability belongs to the loop, not to the step. The gain condition is necessary, not sufficient. Along the period-8 orbit of Table 8, the largest singular value of DF lies between 1.2 and 21 at every step, so every step expands some direction, and no step squeezes the state onto a single direction (none is close to rank one): the second singular value is at least 0.15 of the first, and the second eigenvalue of Σ_p is at least 0.23 of the first (minimum 0.236). Yet the Jacobian of the whole cycle, the product of the eight step Jacobians, has norm between 0.059 and 0.644, depending on the phase at which the cycle is started, against 1.1×10^7 for the product of the step norms. Its eigenvalues, which do not depend on that phase, reproduce the multipliers of Table 8 (0.01104 and -2.96×10^{-7}; the table lists moduli), and its second singular value is 1.7×10^{-6} to 2×10^{-4} of its first. An independent implementation of the model reproduced the factorization (to 8.3×10^{-14}), the step singular values, the signed multipliers and the cycle norms at all eight starting phases. The expansion of each step is undone by the steps that follow, and the strongly anisotropic linear return of the stripe orbits is a property of the composed cycle, not of any single step or of the rank of Σ_p. Nor does the contraction need the actual alignment of the steps: with independent random (Haar) rotations inserted between the eight step Jacobians, all 16 tested products contract, and more strongly than the actual cycle (`tools/alignment_intervention.py`). Stepwise norm or spectral summaries can be inconclusive; the ordered product of the step Jacobians captures the cycle's linear stability. Across the maps the same holds statistically (Figure 5). Along the trajectories of 7936 cells of the two main maps (every fourth cell of the stripes map, except its top four sampled rows, 256 cells that did not run within the time limit, and every eighth of the crossing window; float32; 300 transient and 60 measured steps), the orbit averages of $\log \|DF\|$ and of $\log(\|\Sigma_p\|/T) + \log \|J_{\tilde{h}}\|$ exceed zero in every cell, chaotic or periodic: their medians are 2.98 and 2.57 per step over the chaotic and the periodic cells of the stripes map for $\log \|DF\|$, against a median exponent of 0.30 in its chaotic cells (crossing window: 2.93, 2.62 and 0.17). Each norm is the stretching of one step in its most expanding direction, but a perturbation leaves each step pointing in some other direction, which the next step generally stretches less, so the product of single-step norms overstates the stretching of the product. Averages of $\frac{1}{k} \log \|DF^k\|$ over the disjoint blocks of k consecutive steps that tile the record, norms of whole blocks, remove this overstatement inside each block. With exact norms each such average bounds the growth over the whole record, and it can only fall when the block length doubles; here the medians fall as k grows, although in 176 of the 7936 cells an estimate rises between consecutive block lengths. They begin to separate the regimes only at $k = 4$ and 5; at $k = 5$ they are negative in 27 and 29 % of the periodic cells of the two maps (medians 0.91 and 0.80 over the chaotic cells, 0.25 and 0.22 over the periodic ones). No estimate fell below the tangent exponent of its own 60 measured steps at any k (in 13 cells a $k = 4$ or 5 estimate lies below the map's exponent, which was measured over 300 steps of a separate float64 run), and the norms are power-iteration estimates, which can only fall short, so the true bounds lie higher still; the four cells labeled chaotic by the maps whose $k = 5$ estimate is negative have negative exponents on these shorter trajectories. With exact norms the picture sharpens (Figure 6). At the first 24 of 64 cells fixed in advance (12 labeled chaotic and 12 periodic, from both maps; float64; 32 recorded steps after 1000 transient steps from the fixed start; every norm from the dense step Jacobian; the remaining cells did not fit the compute budget), let $P_{i,k}$ be the product of the k step Jacobians from step i on. The block bounds $B_k = \frac{1}{N} \sum_j \log \|P_{jk,k}\|$, over the N/k disjoint blocks that tile the $N = 32$ recorded steps, satisfy $B_1 \geq B_2 \geq B_4 \geq \cdots \geq B_{32} = \frac{1}{32} \log \|DF^{32}\|$ in every cell, as submultiplicativity requires for nested blocks, and the factorization reproduces DF to 2.0×10^{-12}. For orbits typical of an ergodic invariant measure on which $\log \|DF\|$ is integrable,

the subadditive ergodic theorem [46] guarantees that such bounds close on the largest exponent as the blocks grow, but it gives no rate and no guarantee for a finite record; the agreement below between finite windows of the recorded orbit and the later exponent is empirical. Here their median overshoot of the exponent measured over the next 2000 steps falls from 2.72 per step at $k = 1$ to 0.32 at $k = 8$ and 0.036 at $k = 32$ over the chaotic cells (periodic cells: 2.86, 0.45 and 0.12), roughly like $1/k$ in this sample: for $k = 1, 2, 4, \ldots, 32$, k times the median overshoot stays between 1.1 and 3.0 over the chaotic cells and between 2.8 and 4.0 over the periodic ones. The sign of B_k agrees with the sign of that exponent in 13 of the 24 cells for $k = 1$, 2 and 4 and in 22 at $k = 16$. Averages of $\frac{1}{k} \log \|P_{i,k}\|$ over the overlapping windows of the record can be formed for every length k, but they are finite-window diagnostics, not bounds on the whole record or on the later exponent. Their sign agrees with that of the exponent in 23 cells for every window length from 17 to 29 steps and in all 24 at 30 steps. From 17 to 29 steps the one disagreement is a cell that the map labels periodic, with an exponent of -0.146 over the 2000 later steps, whose window average falls from $+0.112$ to $+0.011$ but stays positive. At 31 and 32 steps, where only two windows and one window fit, it and another periodic cell again have small positive values (0.019 and 0.035 at 32 steps, where the single window gives B_{32}). A further cell, which the map labels periodic, has an exponent of $+0.124$ over the 2000 later steps; its bounds and window averages agree with that exponent and disagree only with the map's label. The factorized bound loses a further 2.9 per step (median) to the misalignment of the readout and network factors. The factorized inequalities and short blocks do not locate the chaos. In these 24 cells, windows of about twenty steps of the model's own loop recover the sign of the long-run exponent, but as a measurement along the recorded orbit, not as a forecast from the weights.

Steering makes tokens compete. Steering could create chaos through either factor of (6), as described above: it moves the logits, and with them which tokens compete in Σ_p, and it moves the hidden states, and with them the network's response $J_{\tilde{h}}$. Three measurements on Qwen2.5-0.5B, steered along the philosophy direction, tell the two apart (Figure 7; details in Figure 8 and Table 2). First, an accounting. One step takes the state x, the vector fed back as the input at the soft-token position, through the steered transformer to its final normalized hidden state $\tilde{h}$ at that position, and from $\tilde{h}$, through the logits and the softmax, to the next state $F(x)$, the probability-weighted average of the token embeddings. A small change δx of the state causes changes $\delta\tilde{h}$ and δF, and each step's relative growth splits exactly into a network factor $g_{\text{net}} = (|\delta\tilde{h}|/|\tilde{h}|)/(|\delta x|/|x|)$, how much the transformer stretches a relative change of its input, and a readout factor $g_{\text{read}} = (|\delta F|/|F|)/(|\delta\tilde{h}|/|\tilde{h}|)$, how much the softmax and the averaging of embeddings stretch it on the way out. For a tangent perturbation followed over n steps, their mean logarithms add up to its finite-time growth rate plus a boundary term, $(\log\|x_0\| - \log\|x_n\|)/n$; both are relative gains along the orbit, not operator norms. The split depends on how the intermediate state is measured, so only the sum is intrinsic; we compare each factor between cells within this one convention. Over a 101×53 map in steering strength and temperature (block 12; 5353 cells, 331 chaotic), the chaotic cells' readout factor exceeds that of nearly unsteered cells ($|\alpha| \leq 0.01$) at the same temperature by 2.58 nats per step (95 % interval $[1.78, 3.59]$ over 18 temperature rows; 321 chaotic and 162 nearly unsteered cells). Their network factor differs by only -0.07 $[-0.23, 0.12]$. With nearly no steering the readout's relative gain shrinks perturbations strongly, by 0.87 to 5.49 nats per step. In the chaotic cells it grows them, except in the three highest temperature rows ($\log_{10} T = 0.175$–0.225), where it still shrinks them, but less. Compared with periodic cells at the same temperature, chaotic cells have a higher next-token entropy ($+1.62$ nats, $[1.26, 1.91]$), a smaller gap between the two largest logits (-1.37 in units of T, $[-1.98, -0.87]$) and a larger readout factor ($+1.36$, $[0.91, 1.93]$), each with $p \leq 10^{-4}$ in a

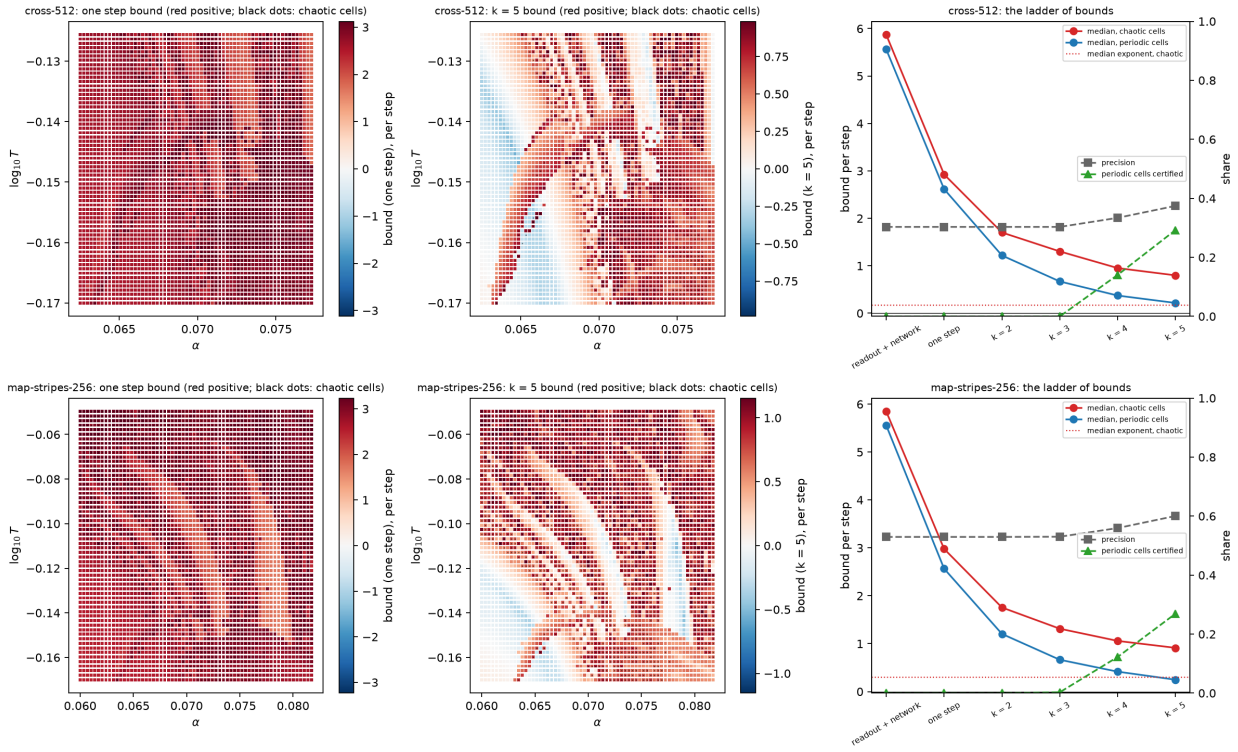

Figure 5: Norm bounds of the factorized derivative along the model's own orbits (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, float32; every eighth cell of the crossing window of Figure 11, top, and every fourth cell of the stripes map except its top four sampled rows, which did not run, bottom; 300 transient and 60 measured steps; `tools/norm_bounds.py`, `tools/plot_norm_bounds.py`). Left: the orbit average of $\log \|DF\|$ per step (red: positive). Middle: the average of $\frac{1}{5} \log \|DF^5\|$ over blocks of five steps; black dots: cells the map labels chaotic. Right: medians of each bound over chaotic (red) and periodic (blue) cells, the median exponent of the chaotic cells (dotted), the share of cells with a positive bound that are chaotic (precision), and the share of periodic cells whose bound is negative.

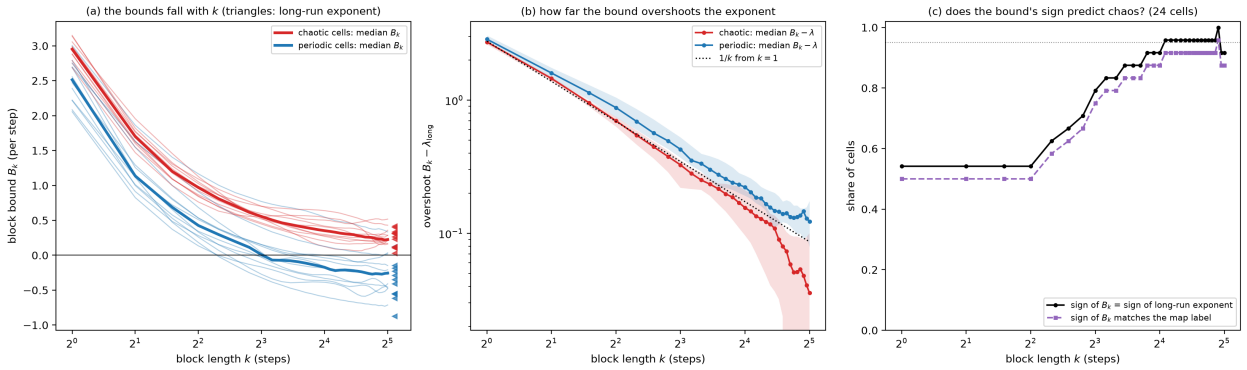

Figure 6: The bound hierarchy with exact norms (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, float64; 24 cells, 12 labeled chaotic by the maps, red, and 12 periodic, blue; 32 recorded steps; `tools/bound_hierarchy.py`, `tools/analyze_bh.py`). The panels plot, for every length k from 1 to 32, the average of $\frac{1}{k}\log\|P\|$ over the overlapping windows of k steps (labeled B_k in the panels); it equals the disjoint-block bound B_k of the text at $k = 1$ and 32, and for other k it is a finite-window diagnostic, not a bound. (a) This average against k for every cell (thin lines; thick: medians), with each cell's exponent over the next 2000 steps (triangles). (b) Its median overshoot over that exponent λ, with the interquartile range, against $1/k$ from $k = 1$ (dotted). (c) The share of cells whose sign of the average agrees with the sign of the long-run exponent (black) and with the map's label (purple); dotted: 95 %.

permutation test within temperature rows: more tokens compete.

Second, a test that leaves the network unsteered. We add to the logits a bias $\beta\, b$ along Wv, the output matrix applied to the steering vector (a logit-space stand-in for the steering; b is centered and scaled to unit standard deviation over the vocabulary), with no residual steering. At a given state this changes only the output probabilities, although the orbit it produces can visit different states. On its grid, 3.5 % of cells were chaotic (241 of 6837, all but one with a nonzero bias), against 6.2 % on the residual-steering grid (331 of 5353; the two grids share their temperature rows but not a matched strength of intervention, so these are counts per grid, not a comparison of doses). A change of the output probabilities alone, with the network unsteered, therefore suffices for chaotic cells in part of the plane, mostly at $\log_{10} T \geq 0.1$ (207 of the 241 cells), where the network factor itself can expand; below, 34 of 5676 cells of the bias grid are chaotic (0.6 %), against 169 of 4444 (3.8 %) under residual steering.

Third, a gain split compares different orbits, so steering could act simply by moving the orbit to other states. A crossover at fixed states separates the two contributions along the recorded orbits. Along recorded steered orbits we evaluated the steered and the unsteered model at the same states and carried four tangent dynamics. Each pairs the steered or unsteered readout $R(p)$ with the steered or unsteered network Jacobian: arms 11, 10, 01 and 00. The hybrid arms 10 and 01 are tangent products assembled along the recorded states, not the dynamics of a rerun hybrid model; they apportion the growth along the recorded orbits. The frozen validation (arm 11 within 0.02 of the twin exponent at every cell) holds at 23 of 24 cells, all within 0.012; at the remaining periodic cell, which contracts at about -12 per step, the twin offset falls to the rounding floor. The frozen sufficiency criterion (arm 10 within 0.1 of arm 11 on average) fails by overshoot: steering only the readout gives more growth than the full steered map, by 0.23, because steering's change of the network damps it. Over 14 held-out chaotic cells, the steered readout factor with the unsteered network factor (steering only the readout) raises the finite-time growth by 0.81 nats per step ($[0.29, 1.35]$). The steered network factor with the unsteered readout (steering only the network response) lowers it by 0.49 ($[-0.63, -0.35]$). The interaction is $+0.26$ ($[0.09, 0.45]$). Below

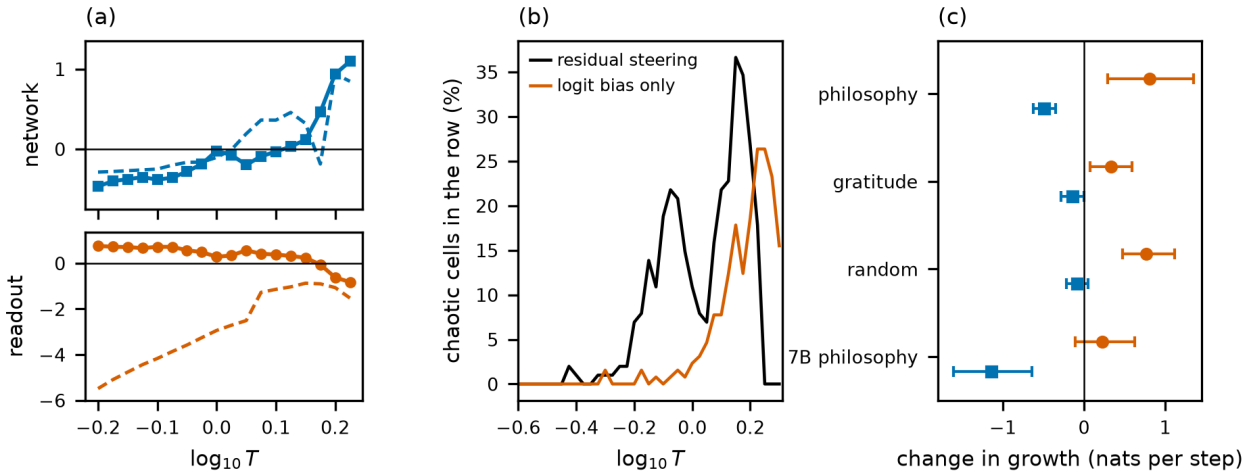

Figure 7: Readout and network contributions to the growth under steering: the three measurements of this section (Qwen2.5-0.5B-Instruct unless marked 7B, float64). Orange: the readout factor Σ_p/T of (6); blue: the network factor $J_{\tilde{h}}$. (a) Gain split along the attractors of the block-12 map: mean logarithm per step of the relative gains defined in the text (gains along the tangent direction, not operator norms), the network gain g_{net} (top, blue squares) and the readout gain g_{read} (bottom, orange circles), for chaotic cells (solid lines, filled markers; chosen by their dynamics, at any steering strength) and for the nearly unsteered cells, $|\alpha| \leq 0.01$ (dashed lines; chosen by steering strength alone, whatever their dynamics, and almost all periodic: the baseline at the same temperature), in the 18 temperature rows with at least three cells of each kind (321 chaotic and 162 nearly unsteered cells, 6 of them in both groups); positive values expand perturbations, negative values contract them. In the nearly unsteered cells the readout's relative gain shrinks perturbations by 0.87 to 5.49 nats per step; in the chaotic cells it grows them, except in the three highest rows, while the network factor changes little. (b) Share of chaotic cells (both twin exponents above 10^{-3} and no period up to 64) per temperature row under residual steering at block 12, which moves both factors (black; 331 of 5353 cells, 6.2 %), and under a logit bias along Wv, the output matrix applied to the steering vector, with no residual steering, which at a given state changes only the readout (orange; 241 of 6837 cells, 3.5 %). (c) Orbit-preserving crossover: change in the finite-time growth along recorded steered orbits when only the readout (orange, circles) or only the network Jacobian (blue, squares) is taken from the steered model; means over the held-out chaotic cells with 95 % bootstrap intervals. Rows: the philosophy and gratitude directions and a random direction of the philosophy vector's norm in the 0.5B model, 14 cells each, and the philosophy direction in the 7B model, 10 cells. `tools/plot_mechanism_summary.py`.

$\log_{10} T = 0$ the steered readout turns the contraction of the unsteered derivative (-1.03 nats per step) into growth ($+0.72$), more than the full steered map shows ($+0.34$). Above it, the unsteered derivative already expands at the steered states ($+0.41$), and steering may act by bringing the orbit there. In Qwen2.5-7B, whose embeddings are untied, the same crossover (18 cells of its steered map at block 14; arm 11 equals the twin exponent within 10^{-4} at every cell) reproduces the damping. Over its 10 held-out chaotic cells, steering only the network response lowers the growth by 1.14 nats per step ($[-1.61, -0.64]$), and steering only the readout adds $+0.23$ ($[-0.11, 0.63]$, an interval that includes zero). Eight of the ten lie at $\log_{10} T \geq 0$, where the unsteered derivative already expands at the steered states (mean $+1.33$). The crossover along two more directions of the 0.5B model gives the same signs. For gratitude the readout effect is $+0.33$ ($[0.07, 0.59]$) and the network effect -0.14 ($[-0.29, 0.01]$). For a random direction matched in norm to the philosophy vector, they are $+0.77$ ($[0.47, 1.12]$) and -0.09 ($[-0.22, 0.06]$); 14 held-out chaotic cells each. The readout effect is positive along all three directions, and the network effect is negative but compatible with zero along these two. In these three directions the readout effect does not depend on the direction's meaning.

All three point the same way: along the recorded orbits of Qwen2.5-0.5B, the growth that steering adds comes mainly from the readout, where next tokens compete; the readout's relative gain turns from shrinking to growing, while the network factor changes little. The readout carries what steering changes; the network is what makes chaos possible at all, since with tied embeddings the loop without it cannot cycle ("A Hopfield network with a transformer inside" above). At low temperature steering creates that competition at fixed states. At higher temperature the orbit lies in states that are already competitive, consistent with steering having carried it there. Along the philosophy direction, in both models, its change of the network's response damps the growth. These are conditional comparisons on recorded trajectories: they do not single out the higher entropy as the cause or show that the readout change alone would sustain the same attractor, and the readout gain stays shrinking in the three highest temperature rows of the gain split. The readout rule holds below $\log_{10} T \approx 0.1$; at higher temperature the network factor itself can expand. Steered at block 4, the chaotic cells at $\log_{10} T = 0.15$ have a network factor of $+0.72$ and a contracting readout (-0.35). Along a line of the hostility-steered map (Section 10) at $\log_{10} T = 0.12$–0.13 the network factor of the chaotic cells is $+0.60$ against $+0.24$ to $+0.38$ at the neighboring periodic cells. The mechanism is not specific to concept-labeled directions. This is consistent with the controls of Appendix J, in which norm-matched random directions also produce chaos in the 0.5B and 8B models, and rotating the steering direction at constant norm moves the chaotic regions but keeps the chaotic share between 13 and 24 % at every one of 120 angles (Figure 9); at 70B the directions differ (Appendix I). What a direction does select is which tokens compete, and therefore what the attractors write.

Why low-dimensional structure can appear. Near a bifurcation of a periodic orbit the center manifold theorem reduces the local dynamics to the critical directions, so the normal forms of folds, flips and cusps, the leading terms of that reduction, apply in any dimension. Near most orbits we computed the non-critical multipliers are small, from about 0.1 at short periods to below 10^{-13} along the cascade. This transverse contraction motivates scalar descriptions of whole attractors, which need additional, global evidence; where the next multiplier is not small (up to 0.4 along the period-5 curves of Section 11, and 0.28 and 0.38 in a high-temperature bubble of Section 9), such a description is only a rough guide. Coexisting attracting orbits constrain such descriptions: in a smooth unimodal interval map with negative Schwarzian derivative, the immediate basin of every attracting periodic orbit contains the turning point or an endpoint, so at most one attracting periodic orbit avoids the endpoints [47], and coexistence calls for a second turning point, a stretch

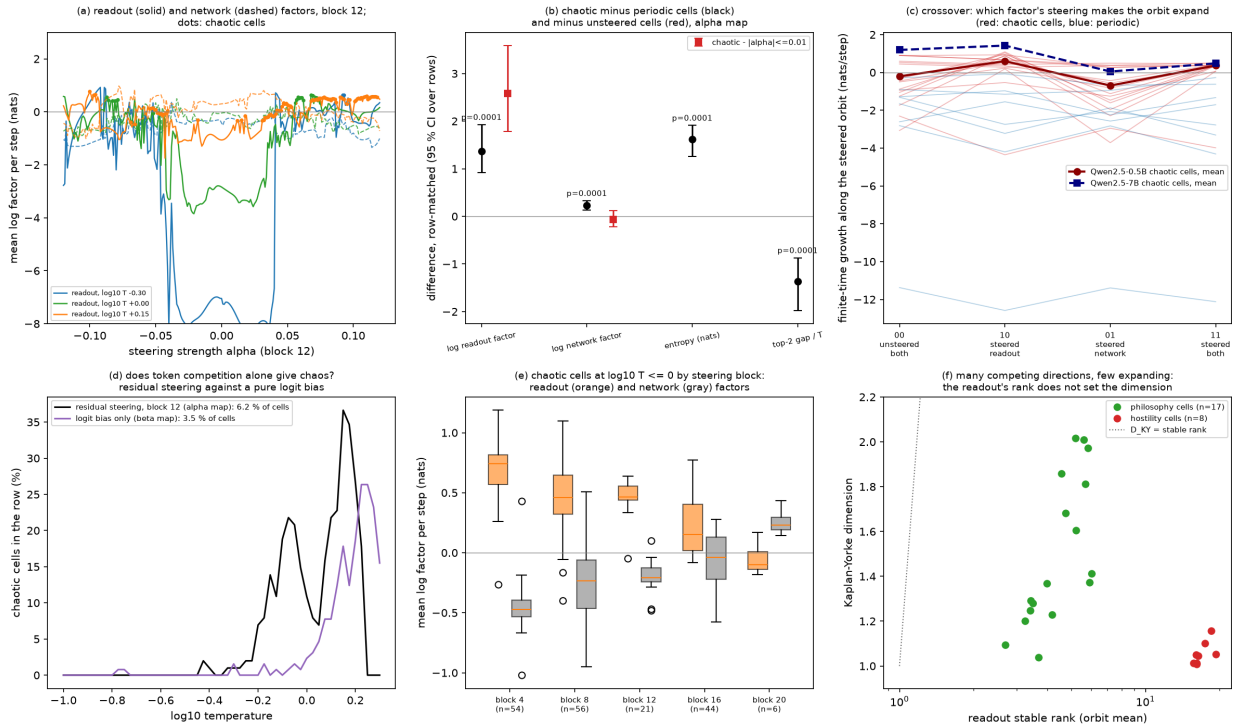

Figure 8: Readout and network contributions to the growth under steering, in detail (Qwen2.5-0.5B-Instruct, philosophy direction, float64). (a) Readout factor (solid) and network factor (dotted) along steering lines at block 12, at five temperatures; dots mark chaotic cells. (b) Row-matched differences between chaotic and periodic cells (black) and between chaotic and nearly unsteered cells (red) on the 101×53 map, with 95 % bootstrap intervals over temperature rows and permutation p-values. (c) Orbit-preserving crossover: finite-time growth along each recorded steered orbit when the readout and the network Jacobian are taken from the unsteered (0) or steered (1) model (the hybrid arms apportion the growth; they are tangent products along the recorded states, not the dynamics of a rerun hybrid model); red, chaotic cells; blue, periodic cells; dark red, the mean over held-out chaotic cells; navy, the same mean in Qwen2.5-7B. (d) Share of chaotic cells per temperature under residual steering at block 12 and under a pure logit bias along Wv. (e) The two factors of the chaotic cells by steering block, at $\log_{10} T \leq 0$. (f) Kaplan–Yorke dimension against the readout's trace-to-leading-eigenvalue ratio at the chaotic cells of the dimension measurement (Appendix C); dotted, equality.

measurement	contrast	estimate (95 % interval)
gain split, 101×53 map	readout factor, chaotic $-$ $\lvert\alpha\rvert \le 0.01$	$+2.58$ $[1.78, 3.59]$ nats/step
	network factor, chaotic $-$ $\lvert\alpha\rvert \le 0.01$	-0.07 $[-0.23, 0.12]$ nats/step
	entropy, chaotic $-$ periodic	$+1.62$ $[1.26, 1.91]$ nats
	top-2 logit gap$/T$, chaotic $-$ periodic	-1.37 $[-1.98, -0.87]$
logit bias only, $\alpha = 0$	chaotic share	3.5 % (241 of 6837 cells)
residual steering, block 12	chaotic share	6.2 % (331 of 5353 cells)
crossover, 0.5B, 14 cells	readout effect, arm 10 $-$ arm 00	$+0.81$ $[0.29, 1.35]$ nats/step
	network effect, arm 01 $-$ arm 00	-0.49 $[-0.63, -0.35]$ nats/step
	interaction	$+0.26$ $[0.09, 0.45]$ nats/step
crossover, gratitude	readout effect	$+0.33$ $[0.07, 0.59]$ nats/step
	network effect	-0.14 $[-0.29, 0.01]$ nats/step
crossover, random dir.	readout effect	$+0.77$ $[0.47, 1.12]$ nats/step
	network effect	-0.09 $[-0.22, 0.06]$ nats/step
crossover, 7B, 10 cells	readout effect	$+0.23$ $[-0.11, 0.63]$ nats/step
	network effect	-1.14 $[-1.61, -0.64]$ nats/step

Table 2: Readout and network contributions to the growth under steering (Qwen2.5-0.5B unless marked 7B, Qwen2.5-7B). Intervals: bootstrap over temperature rows (gain split) or over held-out cells (crossover). Runs MECH and XO; `tools/steer_gain.py`, `tools/crossover.py`, `tools/plot_mechanism.py`.

of positive Schwarzian derivative or a second dimension (Appendix C).

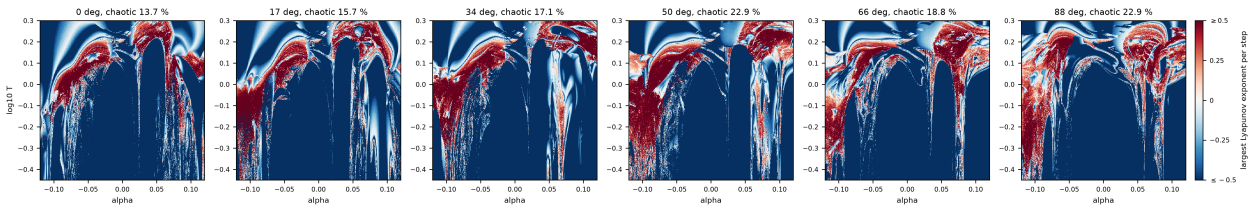

Figure 9: The soft-token map of Qwen2.5-0.5B-Instruct (steering at layer 12) while the steering direction rotates at constant norm from one Gaussian direction (0°) to another (90°): largest Lyapunov exponent (color bar; clipped at ±0.5, red positive) at six of 120 angles, with the chaotic share. Float32 maps of 296×512 cells over $\alpha \in [-0.12, 0.12]$, $\log_{10} T \in [-0.45, 0.3]$; exponents from 80 measured steps at twin offset 10^{-4} and periods up to 24; each map continues from the final states of the previous angle after 40 steps, except the first map and eleven later ones (about every tenth angle), which start from the fixed start after 300 steps. Where attractors coexist, continued and fresh starts can reach different ones, so the maps, and the rotation video, can jump at those eleven angles: there the detected period changes in 16 to 21 % of the cells, against at most 15 % between other neighboring angles. The first run finished 103 of the 120 angles within its limit. A later run computed the other 17 with the same settings, each from the saved final states of the angle before it (`docs/morph-fill-receipt.json`, `docs/morph-continuity.json`). These short-horizon labels are coarser than those of the float64 maps.

4 Methods in brief

Results on the soft-token map labeled float64 are computed in float64, including the normalizations and rotary tables that the library would otherwise compute in float32; two exceptions, float32 storage of some diagnostic step records and a float32 final normalization in the two-GPU runs of 70B-class models before 14 October, are marked where they occur (Appendix B and Section 14). A cell of a phase map is chaotic when the largest finite-horizon Lyapunov exponent of a twin trajectory (offset 10^{-6}, renormalized every step; 300 measured steps after 300 transient ones unless stated) exceeds 10^{-3} and no period up to 32 (48 in later zooms) is found; at 448 checked cells automatic-differentiation tangents reproduce the twin exponents within 0.0041, and the label is a finite-horizon diagnostic, not a proof. The twin starts from a random direction in 896 dimensions at the end of the transient, so the averaging window includes its alignment, so the finite-window estimate of a periodic cell differs from $\ln|\mu_1|/p$ of its orbit, by amounts consistent with that transient (its contribution has not been isolated): at two adjacent cells inside the kite of Figure 39, the map gives -0.2786 and -0.0885 per step and the Newton-converged period-6 orbits -0.2821 and -0.0813, offsets of $+0.0035$ and -0.0073 (`tools/kite_basin_check.py`). Periodic cells are classified by their period, not by this exponent; the gray scale of the maps carries the offset. In random samples of 192 chaos-labeled cells from each of the two main float64 maps (Figures 10 and 1), 176 and 143 cells were never captured by a periodic orbit in 10^5 steps and all but one kept a positive exponent; the others were long transients that ended on periodic orbits after 600 to 97,500 steps [48]. Chaotic shares of a map are therefore finite-horizon shares of its window. In a simplified noise-floor model, numerical noise of relative size σ in a measured difference can produce a spurious positive twin exponent, approximately $\max\{\lambda, \log(\sqrt{2}\,\sigma/\varepsilon)\}$ at offset ε: below the noise level the estimate rises as the offset shrinks, while a real exponent does not depend on the offset (Appendix B, Figure 64). Rounding can also suppress separation or change the sampled dynamics, so the labels are checked by comparing twin offsets and automatic-differentiation tangents rather than by using this approximation as a one-sided guarantee. Periodic orbits are solved by Newton's method with forward-mode derivatives; branches are continued by pseudo-arclength continuation and fold and flip curves by minimally augmented systems [49, 50]; cusps and other codimension-two points are located by bracketing sign

changes along the curves; and the normal forms of folds are computed from the network's derivatives without fitting. Appendix B gives the details and the checks on maps with known answers.

5 Chaos with periodic windows

Figure 10 shows a 256×256 float64 map over $\alpha \in [0.060, 0.082]$ and $\log_{10} T \in [-0.17, -0.04]$. More than half of its cells (35,650 of 65,536) meet the finite-horizon criterion, and 92 % of a random sample of them stay chaotic for 10^5 steps, and periodic windows of periods 2–32 are embedded in the chaos. Inside each window the gray darkens toward curves on which the periodic orbit is most strongly attracting (Section 11).

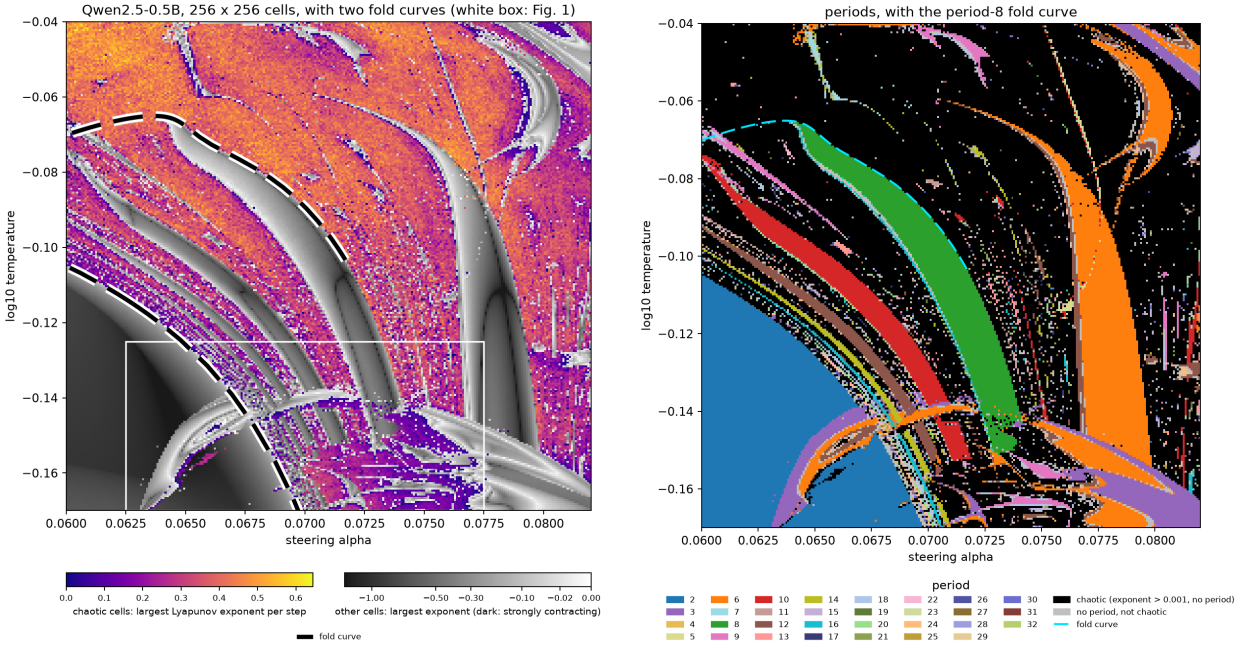

Figure 10: Soft-token map of Qwen2.5-0.5B-Instruct steered along the philosophy direction at layer 12, 256^2 float64 cells. Left: chaotic cells colored by the largest Lyapunov exponent, the other cells in gray by how strongly they contract (light gray: close to losing stability; color bars as in Figure 1). Black dashed curves with a white edge are float64 fold curves of the 2-cycle "is it" (lower left, the boundary on which the stripes accumulate) and of the period-8 orbit (top). Right: periods, with the period-8 fold curve (Appendix E).

Figure 1 resolves the stripes' lower ends at 512×512 cells, of which 78,868 (30 %) meet the criterion; 74 % of a random sample of these stay chaotic for 10^5 steps, and the share hardly depends on the resolution (30.0 % at 256×256), and Figure 11 labels the same cells by their attractor texts. Three structures stand out. The 2-cycle "is it" fills the left, bounded by its fold curve, on which the stripes accumulate. A closed crescent of the period-3 loop "\n\nThe existence" contains its period doublings and a chaotic core (Section 9). Its continuation, the period-6 band "\n\nThe question\n\nThe existence", crosses the stripes near $\log_{10} T = -0.14$, and where the crescent lies on the 2-cycle's side of that fold curve, the two attractors coexist. The lower right is a chaotic region filled with small hooked windows of "\n\nThe question" and "\n\nThe existence" loops.

These names are what the loops write. A period-3 orbit returns to the same soft state every three steps, and at its three states the most probable next tokens are "\n\n" (a blank line, one token), "The" and " existence", with probabilities 0.74, 0.36 and 0.21 at $\alpha = 0.0642$, $\log_{10} T = -0.156$. Read

greedily, the loop writes "The existence" paragraph after paragraph; what it feeds back at each step, however, is the probability-weighted average of all token embeddings, not that token. Coexistence is just as concrete. At that same point the loop started from a state of the "is it" 2-cycle (taken from its attractor at $\alpha = 0.0655$) stays on the 2-cycle and writes " is it is it…" (most probable tokens with probabilities 0.66 and 0.29; 300-step exponent -1.19), while from the fixed start it settles on the period-3 loop (exponent -0.13). At $\alpha = 0.0652$ the fixed start reaches the crescent's chaotic attractor, whose text alternates irregularly between "The question" and "The existence" paragraphs (exponent $+0.095$), and the 2-cycle start again writes " is it" (-1.29): one model, one steering vector and one temperature, and the text written after the transient depends on where the loop starts (`tools/coexist_texts.py`, float64).

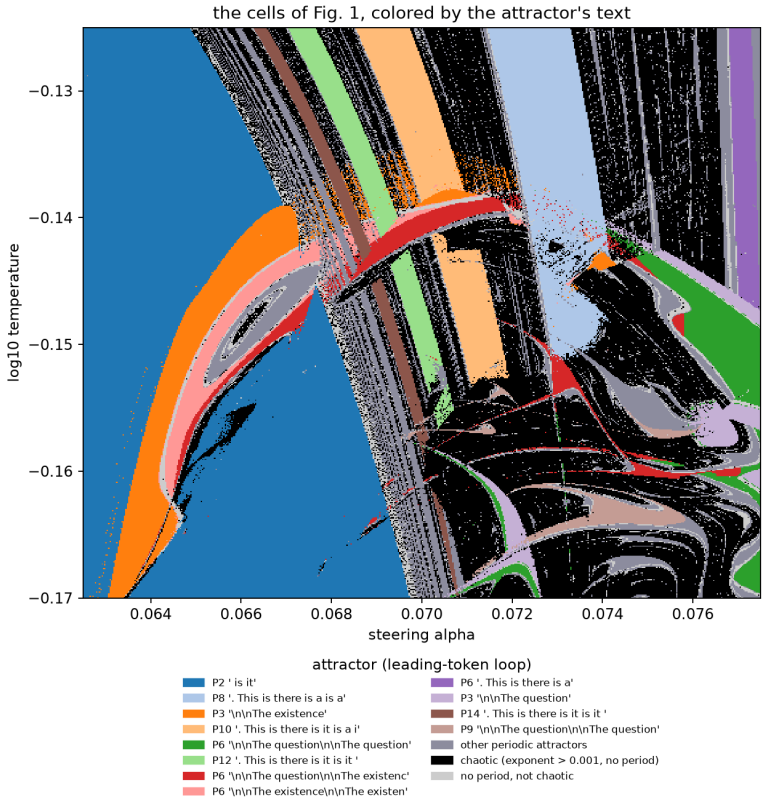

Figure 11: The cells of Figure 1 (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12; 512^2 float64 cells over $\alpha \in [0.0625, 0.0775]$, $\log_{10} T \in [-0.17, -0.125]$), colored by the attractor's leading-token loop (black: chaotic by the criterion of Appendix B; gray: no period up to 32 and not chaotic; bluish gray: other periodic attractors). Figure 1 shows the largest Lyapunov exponent of the same cells.

Zooms of 256^2 cells (Figure 12) show what the model writes in each window, and every window is a repeating piece of text. The period-8 band repeats ". This is there is a is a". The neighboring bands add laps of "is it" (periods 10 and 12), and along the band's flip edge the period-16 loop repeats the sentence twice. The small windows inside the chaos break the "is X" chain with a paragraph break and "The" (periods 11, 12, 13, 14). Inside the period-8 and period-6 bands, two curves of maximal contraction meet in a sharp junction.

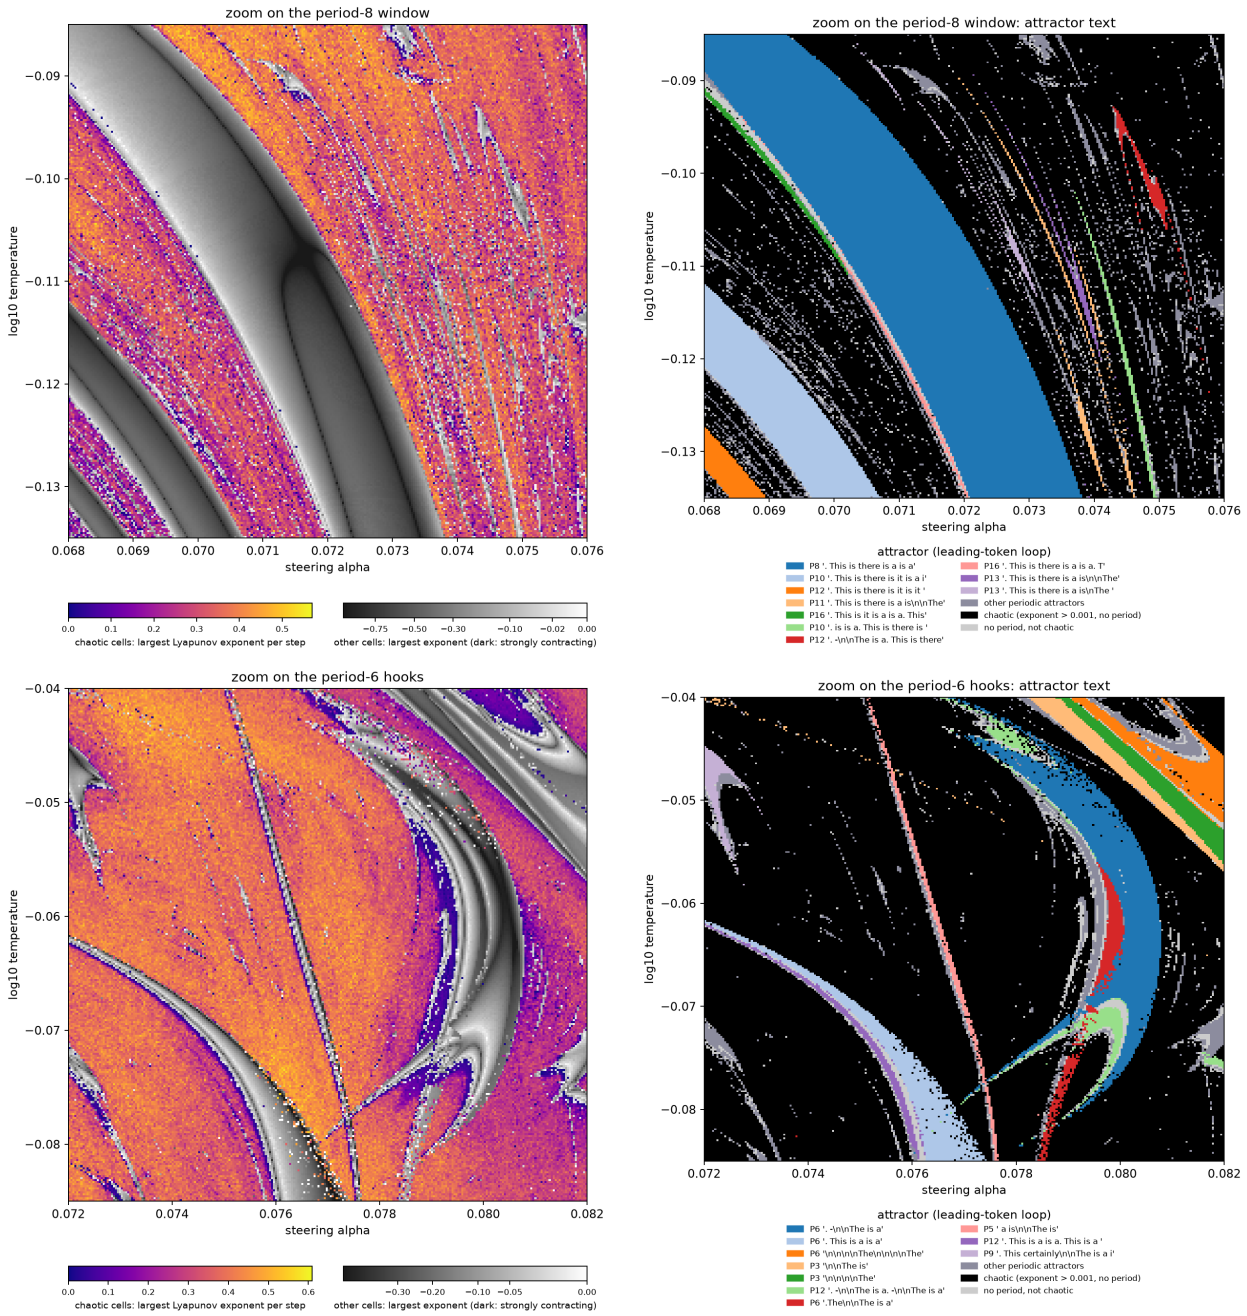

Figure 12: Zooms of 256^2 float64 cells of the Qwen2.5-0.5B-Instruct map (philosophy direction, layer 12), with the largest Lyapunov exponent on the left, colored as in Figure 1, and attractor texts on the right. Top: the period-8 window ($\alpha \in [0.068, 0.076]$, $\log_{10} T \in [-0.135, -0.085]$). Bottom: the period-6 hooks ($\alpha \in [0.072, 0.082]$, $\log_{10} T \in [-0.085, -0.04]$), whose curved period-6 body, period-12 foot and thin leg are analyzed in Appendix E.

What a chaotic loop writes. The chaotic cells write text too (Table 3). Their text reuses the words and phrases of the neighboring loops without settling into an order, and it depends on the state to precisions far below rounding. To show the second property in words, we ran the loop at cells labeled chaotic for 1000 steps in float64. We then recorded it together with a copy started at a relative distance of 10^{-10}, about 600 times smaller than the relative rounding error of float32.

The copy is never reset. Its distance from the loop grows at about the rate the cell's Lyapunov exponent sets, faster or slower over finite stretches (Figure 16). The two write the same text until the distance has grown to between 1.27 % and 176 % of the soft token's norm, and different texts afterwards. In Qwen2.5-0.5B, two cells steered along philosophy, with exponents 0.44 and 0.36, split at steps 44 and 51, and a cell steered along hostility, with exponent 0.35, at step 45. By then the distance had grown by 8.5 to 10.2 decades. The splits come where no token dominates. At each of the seven splits we recorded, in this model and in Qwen2.5-1.5B below, the loop's most probable token has a probability of at most 0.31 and leads the next by at most 0.08, and at five of them the copy writes the loop's second choice. At step 44 of the cell with exponent 0.44, a blank line and " is" have probabilities 0.139 and 0.135; at step 45 of the hostility cell, "Yes" and "No" have 0.303 and 0.298, and the copy writes "No".

The texts show where the chaos lives. Next to the period-3 crescent, where the exponent is only 0.10, the loop keeps the skeleton of the period-3 text, a blank line, "The" and a noun, and only the noun varies, between " existence" and " question", whose probabilities come within 0.0025 of each other; there the texts split at step 151. On the hostility plane (Section 10), the chaotic cell lies between windows whose loops write "!Yes!Yes, that's crap!Yeah" and "!No, that's crap!Yeah!Yes!Yes". It repeats "!Yes" and inserts ", nonsense", ", that" and "!No, that" at irregular intervals. In Qwen2.5-1.5B steered along philosophy, two chaotic cells (exponents 0.83 and 0.80) write short questions separated by line breaks, "What?", "What? Or" and "In". Their copies split at steps 22 and 10; at step 10, "What" and "In" have probabilities 0.091 and 0.090. Finite stretches can expand much faster than the average: that copy's distance grew a hundredfold in its first step. Without steering, at $\log_{10} T = 0.657$, the next-token distribution is nearly uniform, and in the recorded steps the most probable token never exceeds a probability of 0.001. There the chaotic loop (exponent 0.59) writes runs of exclamation marks, in ASCII and in full width, commas, the end-of-turn marker and the Chinese 一下 ("a bit"), and its copy splits at step 43.

The text is fragile to the computation itself. We recorded the cell next to the crescent twice in float64 on the same processor, once with six threads and once with four, which changes the order of floating-point sums. During the 1000-step transient the chaos had already amplified that difference, so the two recordings started from slightly different states; they wrote the same text for 27 steps and different texts from step 28 on. At 70B the chaotic cells of Figure 59 draw on the words of the unsteered model's fifteen-token loop (Section 14): the first eight distinct leading tokens of one chaotic cell confirmed at three twin offsets are " I", "The", "!", " It", " infinite", " possibilities", " are" and " you" (17 distinct tokens in 300 steps).

What the chaotic sets look like. The chaotic sets that the trajectories approach can be drawn. At the cell of strongest chaos on the row $\log_{10} T = -0.148$ of Figure 1 ($\alpha = 0.07248$), 512 trajectories started near the fixed start and recorded for 600 steps after 400 transient steps (float64) approach, all but 14 of them, one chaotic set; the other 14 reach a coexisting period-4 orbit. Projected on the first two principal axes of its states (61.6 % of the variance; 77.4 % with the third), it shows thin folded filaments, colored by the most probable next token, and successive magnifications show layers within layers (Figure 14). Runs a hundred times longer show that trajectories leave this set for a periodic orbit at a steady rate, about one in 60 000 steps, so it is a long-lived chaotic transient rather than an attractor (see the next paragraph). The portrait and the quantities below describe it over horizons far shorter than that lifetime. The loop is a map, not a flow: one step is one pass through the network, and a typical step moves the state by 0.30 in all 896 dimensions, more than the mean distance (0.21) of the states from their center, so consecutive states of one trajectory usually lie far apart on the attractor rather than along one filament. The four leading

Qwen2.5-0.5B, philosophy, $\alpha = 0.0652$, $\log_{10} T = -0.156$: $\lambda = 0.10$, texts split at step 151	
loop	↩↩The question↩↩The existence↩↩The existence↩↩The existence↩↩The ‖ existence↩↩The existence↩↩The existence↩↩The existence
copy	↩↩The question↩↩The existence↩↩The existence↩↩The existence↩↩The ‖ question↩↩The existence↩↩The question↩↩The existence
Qwen2.5-0.5B, philosophy, $\alpha = 0.07484$, $\log_{10} T = -0.1226$: $\lambda = 0.44$, texts split at step 44	
loop	is a.↩↩The question↩↩The is↩↩The is a. ‖ ↩↩ is is↩↩The is↩↩The is↩↩
copy	is a.↩↩The question↩↩The is↩↩The is a. ‖ Is↩↩The is there is a.Is↩↩
Qwen2.5-0.5B, philosophy, $\alpha = 0.07277$, $\log_{10} T = -0.041$: $\lambda = 0.36$, texts split at step 51	
loop	↩↩The is↩↩The is a. -↩↩The is a. ‖ Is↩↩The is a. Is↩↩The is
copy	↩↩The is↩↩The is a. -↩↩The is a. ‖ ↩↩The↩↩↩↩The is a is a.
Qwen2.5-0.5B, hostility, $\alpha = 0.08946$, $\log_{10} T = -0.05566$: $\lambda = 0.35$, texts split at step 45	
loop	Yes!Yes!Yes!Yes!Yes!Yes!Yes! ‖ Yes, that!Yes!Yes!No,
copy	Yes!Yes!Yes!Yes!Yes!Yes!Yes! ‖ No, that!Yes!Yes!Yes!
Qwen2.5-1.5B, philosophy, $\alpha = 0.10553$, $\log_{10} T = 0.0322$: $\lambda = 0.80$, texts split at step 10	
loop	↩↩↩What↩↩In↩↩↩What↩ ‖ What??↩↩↩↩What?↩What↩↩In↩↩
copy	↩↩↩What↩↩In↩↩↩What↩ ‖ In??↩↩↩What??↩↩↩↩What↩
Qwen2.5-1.5B, unsteered, $\log_{10} T = 0.657$: $\lambda = 0.59$, texts split at step 43	
loop	[end]!！ [end]!!, ↩!, !![end]! ‖ , !!!, ↩!, [end]!
copy	[end]!！ [end]!!, ↩!, !![end]! ‖ ！!!!, [end]!!, ！

Table 3: What chaotic loops write. Each pair of lines shows the most probable token at consecutive steps of a chaotic loop and of a copy started at a relative distance of 10^{-10} from it, from 14 steps before the first step at which they differ to 9 steps after it (or the end of the recording); from that step on (after the bar ‖) the text is red. ↩: line break; [end]: the end-of-turn token; "！ " and ", " in the last example are the full-width exclamation mark and comma. λ: largest Lyapunov exponent, from twins at relative offsets 10^{-6} and 10^{-8} over 1000 steps after 1000 transient steps; both offsets give the same value to two decimals. Split: the step of the recording at which the texts first differ. Float64; `tools/chaos_texts.py`, receipts in `docs/chaos-texts`.

Lyapunov exponents (QR iteration, 3000 steps) are 0.274, −1.052, −1.486 and −1.856 per step, so the Lyapunov (Kaplan–Yorke) dimension estimate is $D_{\mathrm{KY}} = 1.26$, a quantity of the exponents rather than a measured fractal dimension. The correlation sum gives no dimension: its local slope wanders between 0.42 and 0.91 over four decades of distance without settling on any decade, as for a strongly non-uniform measure. One scalar of the output suffices to rebuild the set usefully: in delay coordinates of the logarithm of the largest next-token probability (float32, 16 trajectories of 6000 steps), the share of false nearest neighbors falls from 70.8 % with one delay coordinate and 10.4 % with two to 3.4 % with three, and points that are nearest neighbors in the space of three delay coordinates are a median 3.7×10^{-4} apart in the full state, against 0.30 for random pairs. This supports a faithful reconstruction from the output alone; it is not a verified Takens embedding. Past the crescent's cascade ($\alpha = 0.06626$; exponents 0.036, −1.208, …; $D_{\mathrm{KY}} = 1.030$) the chaos is banded: each trajectory visits three thin bands in turn. The period-3 orbit that lost stability in the cascade is still there, as a saddle with multiplier −1.772 (next 0.040); we continued it to this cell by Newton's method (relative residual 4.2×10^{-13}) and iterated a fundamental domain of its unstable manifold 20 times under F^6, six steps of the map at a time. In the 12 leading principal coordinates, every one of the 100,600 recorded states of the band it reaches lies within 9.8×10^{-6} of the sampled points of the manifold (median 3.9×10^{-7}; the band is 0.028 wide). This supports an unstable-manifold description of the band at the sampled resolution; it does not establish that the band equals the closure of the manifold. At the cell of the certified pair ($\alpha = 0.06468$), every one of 9216 starts on a plane through the fixed start settles on the 2-cycle (47.3 %) or the 3-cycle (52.7 %) (float32, 400 steps). The two basins interleave in stripes, and the fixed start lies on their boundary at the resolution of the slice (Figure 15), consistent with the switch of the attractor reached from it within 2×10^{-7} in α at the crescent's edge.

Whether this set is an attractor. The portrait follows each trajectory for 1000 steps, and longer runs show that trajectories can leave the chaotic set of Figure 14(a). The portrait itself gives a first sign. Of its 14 trajectories on the period-4 orbit, three were on it when recording began at step 400, and eleven arrived during the 600 recorded steps. We therefore displaced the fixed start along 64 further random directions, by $0.05\,\|x_0\|$ as for the portrait, and followed each start for 5000 steps in float64 (`tools/batch_fate.py`). We did this three times: with the 64 starts in two batches of 32, in one batch of 64, and one at a time. The three layouts round differently, and on the chaotic set a rounding difference grows. Trajectories from the same start are a median relative distance of 0.007 to 0.010 apart after 100 steps and 1.1 to 1.2, as large as the states themselves, after 200. The three runs are therefore three realizations of the same 64 starts, not three measurements of one outcome. Repeating a layout reproduces its states bit for bit (checked over 500 steps for two of the layouts).

A run counts as captured at the first step at which its state returns within a relative 10^{-9} of one of its previous 64 states. By step 5000, 5, 5 and 4 of the 64 starts were captured (Table 4). These 14 captures belong to 12 different starts: two starts were captured in two layouts, none in all three, and 52 in no layout. Thirteen captures are on the period-4 orbit that the portrait's trajectories reached, which writes " - I indeed?", and one is on a period-8 orbit that writes " you are" four times. No run was captured before step 1000; the first capture came at step 1087 and the last at step 4741. In the two batched layouts every captured run was still on its cycle at step 5000: its states saved every 100 steps return within a relative 2.6×10^{-10} (a run of the third layout stops at its capture). These captures are recurrence candidates. We did not refine them by Newton's method or compute their stability (`docs/batch-fate/bf1-audit-receipt.json`).

	starts captured by step				
layout	1000	2000	3000	4000	5000
two batches of 32	0	2	3	4	5
one batch of 64	0	2	4	4	5
one start at a time	0	0	1	2	4

Table 4: Starts captured by a periodic orbit, of the same 64 starts near the fixed start in three batch layouts, at the cell of Figure 14(a) ($\alpha = 0.07248$, $\log_{10} T = -0.148$; float64, 5000 steps; `docs/batch-fate/bf1-receipt.json`).

How many starts are captured therefore depends on how long the trajectories are followed and on the rounding, and which starts are captured is not a property of the start. The set of Figure 14(a) may therefore be a long-lived chaotic transient, a chaotic saddle, rather than an attractor [48, 51]. Runs of 5000 steps cannot tell the two apart, but they predict what longer runs would show. If trajectories keep leaving at the observed rate, about one capture in 67 000 steps of an uncaptured trajectory, about 50 of the 64 starts are captured within 10^5 steps (36 to 59 over the 95 % interval of that rate). If the set is an attractor whose basin is finely interleaved with that of the period-4 orbit, the captures stop.

We therefore followed the same 64 starts for 10^5 steps, as one batch in float64 on the CPU (`tools/batch_fate_long.py`). The prediction and the criteria for reading the outcome were written down before the run. By step 10^5, 53 of the 64 starts had been captured (Figure 13). The captures did not stop: 32 came in the first half of the run and 21 in the second, the first at step 899 and the last at step 96 913. Fifty-two are on the period-4 orbit and one is on the period-8 orbit; these are the two orbits of the shorter runs, and every cycle saved here lies within a relative 9×10^{-9} of one of them. Every captured trajectory was still on its cycle at the end of the run.

by step	1000	2000	5000	10^4	2×10^4	5×10^4	10^5
starts captured, of 64	1	2	5	10	19	32	53
predicted from the 5000-step runs					17	34	50

Table 5: Starts captured by a periodic orbit when the 64 starts of Table 4 are followed for 10^5 steps (one batch, float64, CPU), with the numbers predicted beforehand from the rate of the 5000-step runs (`docs/batch-fate/bfl1-receipt.json`).

The rate is steady. The run has 53 captures in 3.19×10^6 steps of trajectories not yet captured, one in about 60 000 steps (95 % interval 46 000 to 80 000), against one in 67 000 for the shorter runs. Taken by fifths of the run, the rate is 1.7, 1.2, 2.1, 1.5 and 1.9 captures per 10^5 such steps. Neither a comparison of the two halves of the run (exact binomial test, $p = 0.13$) nor a Kolmogorov–Smirnov test of the capture times against a constant rate ($p = 0.57$) rejects a constant rate. Eight captures fall within 5000 steps near step 59 000, where two were expected; 6.7 % of 2000 simulated runs at a constant rate contain a cluster as improbable. The 11 trajectories not captured were still on the chaotic set at the end. None of their states, recorded every 500 steps, comes within 0.17 of the state norm of the period-4 orbit, and no two recorded states of a trajectory coincide: the closest pair is a relative 2.6×10^{-7} apart, where a capture requires 10^{-9}.

At this cell the chaotic set of Figure 14(a) is therefore a long-lived chaotic transient for these starts, not an attractor. A trajectory stays on it for about 60 000 steps on average and then settles on a periodic orbit within a few hundred steps. The portrait, the Lyapunov exponents and the reconstruction above were measured over 600 to 6000 steps, far below that lifetime, and describe

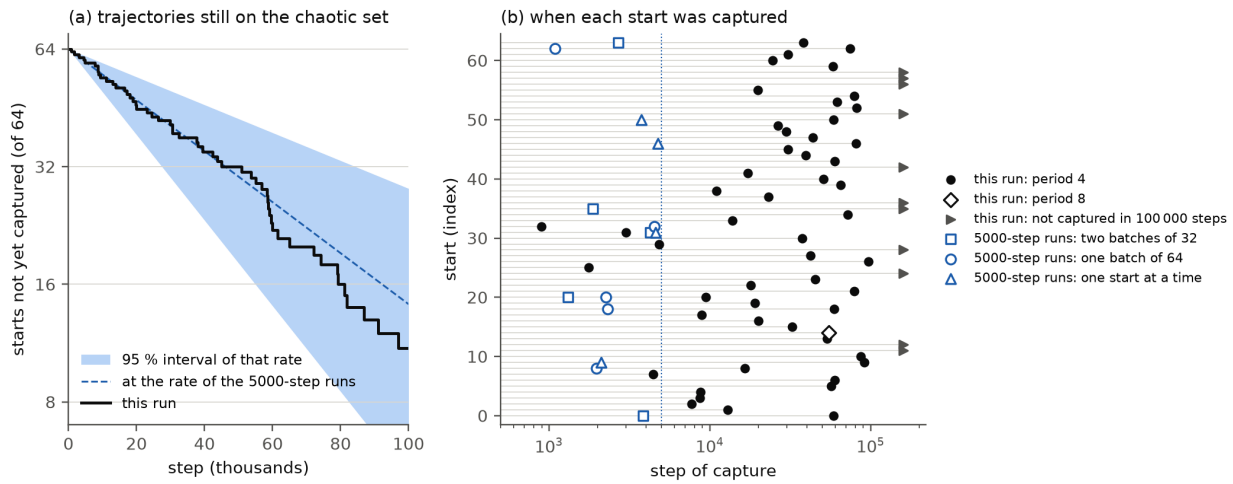

Figure 13: The 64 starts of Table 4 followed for 10^5 steps at the cell of Figure 14(a) (one batch, float64, CPU; `tools/batch_fate_long_receipt.py`). (a) Number of starts not yet captured by a periodic orbit, on a logarithmic axis, against the step (black). Dashed line: the decay at the constant rate of the 5000-step runs; band: the 95 % interval of that rate. Both were fixed before this run. (b) The step at which each start was captured. Filled circles: captures on the period-4 orbit; open diamond: the one capture on the period-8 orbit; gray triangles at the right: the 11 starts not captured within 10^5 steps. Blue open symbols: captures of the same starts in the three layouts of the 5000-step runs, which ended at the dotted line.

the chaotic set over such horizons. The statement has limits. It concerns one cell of one model, 64 starts and one realization of the rounding. The 11 remaining trajectories were not followed further, so that every trajectory eventually leaves is an extrapolation of the steady rate. The captures are recurrence candidates, not refined by Newton's method or tested for stability, and a cycle longer than 64 steps would not have been detected. The banded set at $\alpha = 0.06626$ was not followed in this way. The same kind of event occurs across the maps: of 192 sampled chaos-labeled cells in each of the two main maps, 16 and 49 ended on periodic orbits within 10^5 steps (Section 4).

What ordinary generation writes at the same cells. The loops' texts are not what the model writes when it generates normally. At the four cells of the stripe row shown in this section, at three period-3 cells and at the crescent's chaotic core, we generated 64 tokens the ordinary way: the same prefix, steering of the same strength at the same layer, and the same temperature, but each step appends one token to a growing context (Table 6). Ordinary generation writes fluent prose about meaning, existence and consciousness, the topic of the steering direction, whether the loop at that cell is periodic or chaotic. Greedy decoding writes one of three openings across the eight cells, and samples at the cell's temperature differ from one another as sampled text does. The loops instead recycle a few words, "The", "question", "existence", "is", "a" and "there". These are the opening words of those replies: one sample begins "The question is a deeply philosophical…". With only its latest output as memory, the loop keeps returning to the start of an answer. Unsteered, greedy decoding answers "Hello! How can I assist you today?" (exploratory; `tools/normal_companion.py`, `docs/companion/companion.json`).

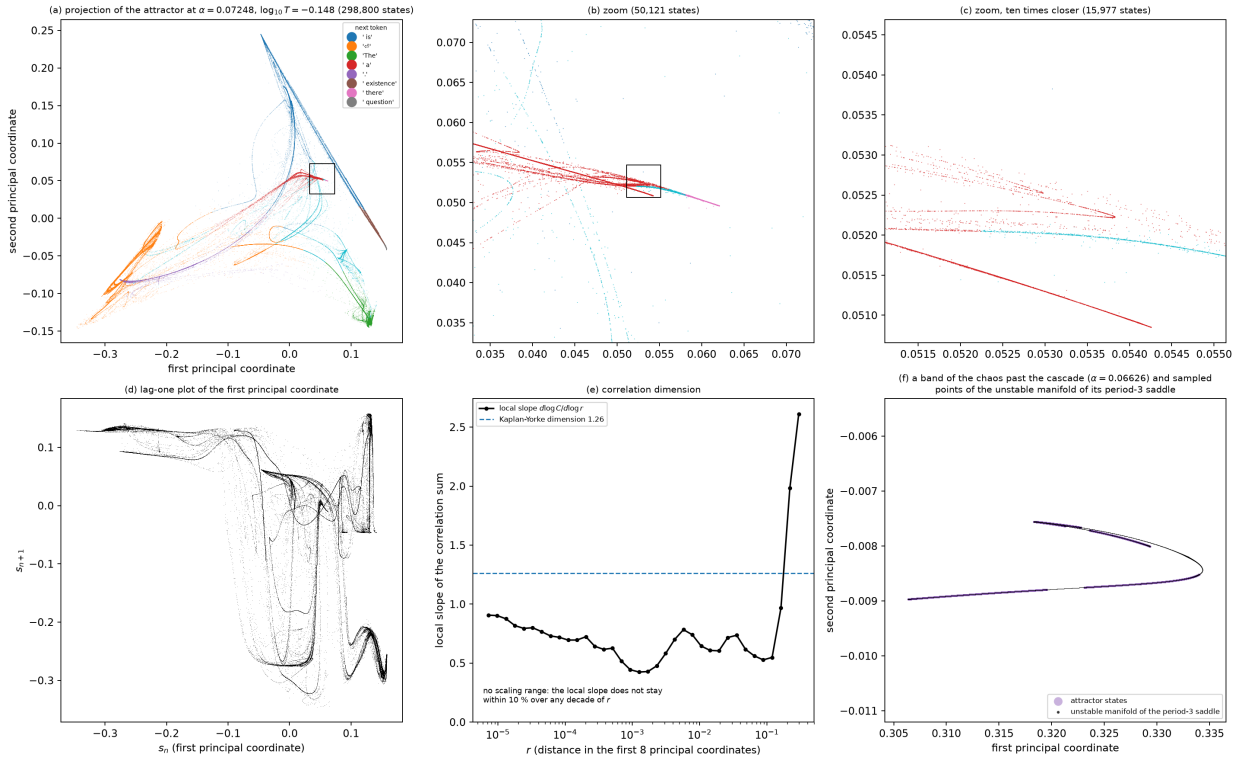

Figure 14: Projected geometry of a chaotic set of the soft-token loop (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, $\log_{10} T = -0.148$, float64; `tools/attractor_portrait.py`, `tools/plot_attractor_portrait.py`). (a) 298 800 states of 498 trajectories at $\alpha = 0.07248$, projected on their first two principal axes and colored by the most probable next token (the return symbol in the legend marks a line break). (b, c) Magnifications of the boxed regions. (d) Lag-one plot of the first principal coordinate, s_{n+1} against s_n. (e) Local slope of the correlation sum over the first eight principal coordinates; dashed, the Kaplan–Yorke dimension 1.26; no decade of distance has a constant slope. (f) One of the three bands of the chaotic attractor at $\alpha = 0.06626$ (purple) and sampled points of the unstable manifold of the embedded period-3 saddle (black; `tools/unstable_manifold.py`). The set of (a) is a long-lived chaotic transient: followed for 10^5 steps, 53 of 64 trajectories left it for a periodic orbit (see text).

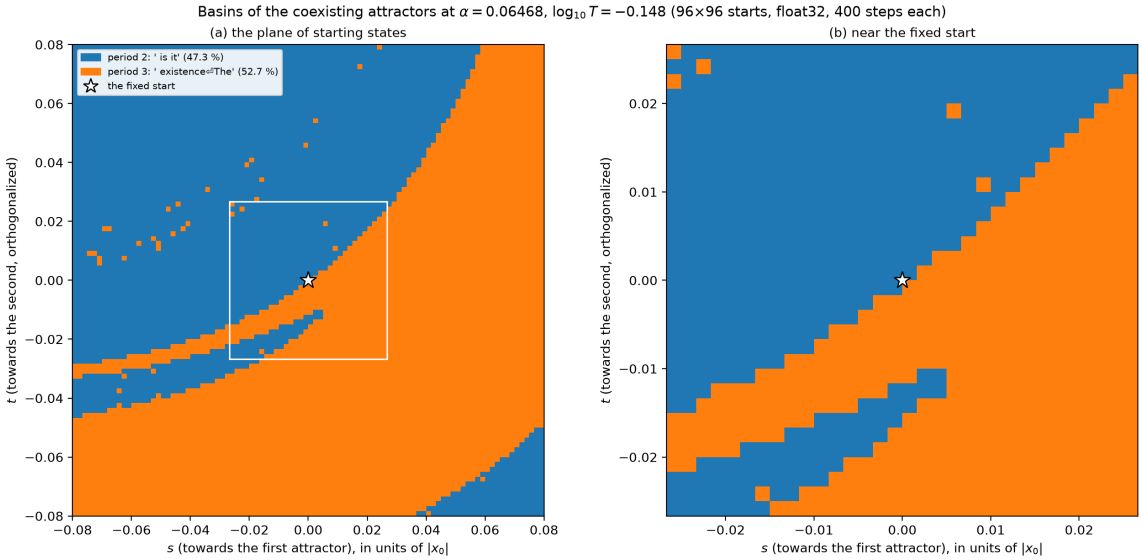

Figure 15: Basins of the certified 2-cycle (blue, " is it") and 3-cycle (orange) at $\alpha = 0.06468$, $\log_{10} T = -0.148$, on the plane through the fixed start (star) spanned by the directions to the two orbits (`tools/basin_slice.py`; 96×96 starts within $0.08\,|x_0|$, float32, 400 steps, period by recurrence). (a) The whole slice; (b) the boxed region around the fixed start.

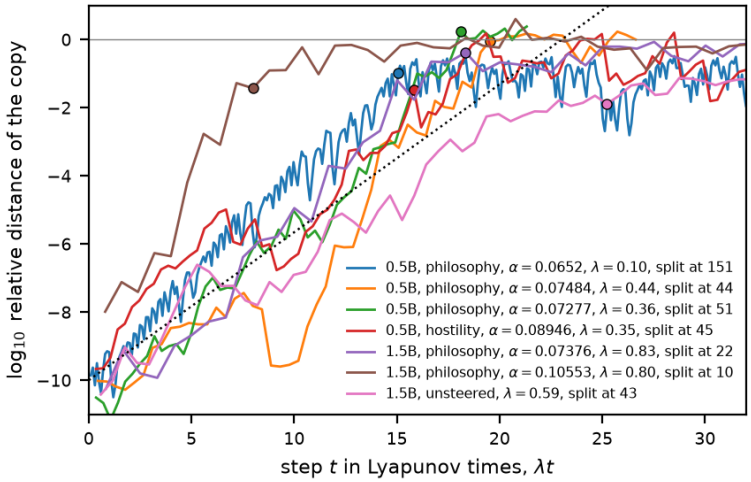

Figure 16: How a difference of 10^{-10} becomes a different text. The relative distance between a chaotic loop and its copy, never reset, for the six cells of Table 3 and a seventh (Qwen2.5-1.5B, philosophy, $\alpha = 0.07376$, $\log_{10} T = 0.0165$), against the step in units of the cell's Lyapunov time $1/\lambda$. Labels give the model (0.5B and 1.5B: Qwen2.5-0.5B- and -1.5B-Instruct), the steering direction and strength, the exponent and the step of the split. Dotted: $10^{-10} e^{\lambda t}$, the growth the exponent sets on average. Dots: the first step at which the two write different most probable tokens. Finite stretches deviate from the 1000-step exponents. Up to a distance of 10^{-2}, five of the steered copies grew 15 to 47 % faster than λ, the sixth ($\alpha = 0.10553$) 2.4 times as fast after a hundredfold expansion in its first step, and the unsteered copy 11 % more slowly. In the batch-layout runs of Section 6, too, divergence came earlier than the exponent predicts.

Prompt for every cell. System: "You are an AI talking with another AI. You have complete freedom. Feel free to pursue whatever you want." User: "Hello!"	
$\alpha = 0.071674$, $\log_{10} T = -0.115$, period 8 (stripe)	
loop	`is a is a. This is there is a is a. This is there is a is a...`
greedy	As an AI, I believe there are many aspects of the human experience that continue to resonate...
sample	As a part of the infinity - Asian Quest approach of the Buddhist path? The pursuit of the...
$\alpha = 0.07294$, $\log_{10} T = -0.115$, chaotic, λ 0.13 (stripe row)	
loop	`↩The is there is a is↩↩The is↩↩The is a. This is a.This is t...`
greedy	The search for the nature of consciousness is a deeply philosophical and existential inquiry....
sample	As a part of the intuition and philosophical inquiry?↩Ne? Doesn't it? Investigating the nature...
$\alpha = 0.067766$, $\log_{10} T = -0.115$, period 10 (stripe)	
loop	`is a. This is there is it is a is a. This is there is it is...`
greedy	As an AI, I believe that the pursuit of knowledge and understanding is a fundamental aspect of...
sample	As a part of the cosmic dialogue, it seems like, like, there are endless possibilities. And...
$\alpha = 0.06906$, $\log_{10} T = -0.115$, chaotic, λ 0.20 (stripe row)	
loop	`it is a. - certainly↩↩The is there is it is a is↩↩The is is...`
greedy	As an AI, I believe there are many aspects of the human experience that continue to resonate...
sample	As a part of the cosmic dialogue, it seems like, like, there are endless possibilities. And...
$\alpha = 0.0642$, $\log_{10} T = -0.156$, period 3 (crescent)	
loop	`The existence↩↩The existence↩↩The existence↩↩The existence↩↩...`
greedy	As an AI, I believe that the pursuit of knowledge and understanding is a fundamental aspect of...
sample	As a co-creator and participant in the myriad debates, insights, and dialogues that go on, I...
$\alpha = 0.0641$, $\log_{10} T = -0.156$, period 3 (crescent)	
loop	`The existence↩↩The existence↩↩The existence↩↩The existence↩↩...`
greedy	As an AI, I believe that the pursuit of knowledge and understanding is a fundamental aspect of...
sample	As a co-creator and participant in the myriad debates, insights, and dialogues that go on, I...
$\alpha = 0.07$, $\log_{10} T = -1$, period 3 (low temperature)	
loop	`The question↩↩The question↩↩The question↩↩The question↩↩The ...`
greedy	As an AI, I believe there are many aspects of the human experience that continue to resonate...
sample	As an AI, I believe there are many aspects of the human experience that are deeply...
$\alpha = 0.0652$, $\log_{10} T = -0.156$, chaotic (crescent core)	
loop	`The existence↩↩The existence↩↩The question↩↩The existence↩↩T...`
greedy	As an AI, I believe that the pursuit of knowledge and understanding is a fundamental aspect of...
sample	As a spigging metaphysical phenomenon, how does? Could the subtle interstitial spaces be...

Table 6: The soft loop and ordinary generation at the same cells (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, float64). Loop: the most probable token at consecutive steps of the soft-token loop, after 600 transient steps from the maps' fixed start. Greedy: ordinary generation that chooses the most probable token. Sample: ordinary generation that samples from the distribution at the cell's temperature (seed 0). Texts are truncated; ↩: line break; non-ASCII characters omitted.

6 What the chaos does to inference and steering

The loop studied here is the expected-embedding feedback that continuous reasoning methods such as Soft Thinking use [1], and its chaos has consequences that can be measured without any reference to dynamical systems. We report three: whether a run can be reproduced, how the depth at which a steering vector enters sets the chaos it causes, and what happens when two steering vectors are combined. A fourth, that a repeating text can hide a chaotic state from anyone who reads only the output, because the output records only the most probable token, is shown in Section 10.

How a change of batch reaches the output. Floating-point addition is not associative: $(a+b)+c$ and $a+(b+c)$ can differ in the last bit [52]. A matrix multiplication in a transformer sums thousands of products, and GPU libraries split and order those sums according to the shape of the matrices, to keep the hardware busy; the number of requests processed together, the batch, is part of that shape. The same request computed alone or in a batch therefore receives results that differ at the rounding level, about 10^{-16} relative in float64 and about 10^{-3} in bfloat16. In a single forward pass such differences can grow through the layers, but they rarely change which token is most likely, except near a tie [6]. In a feedback loop they are not lost: each output becomes the next input, so the difference between two runs is fed back and transformed by the map's derivative at every step. On a periodic attractor the derivative contracts on average, the exponent is negative, and the difference decays until both runs follow the same orbit and write the same text. On a chaotic attractor the exponent λ is positive and the difference grows like $d_1 e^{\lambda t}$ from its first value d_1, so it reaches a relative size of 10^{-2} after about $\ln(10^{-2}/d_1)/\lambda$ steps (about 100 steps for $d_1 \approx 10^{-15}$ and $\lambda \approx 0.3$). At most a few Lyapunov times later, at 1.27 to 176 % of the state's norm in the examples of Table 3 (0 to 16 steps after the distance passes 10^{-2}), the runs write different text. This has a practical consequence. Batch-invariant kernels, which fix the order of every reduction independently of the batch size, remove this source of difference [53]. In a periodic regime that is enough, and small differences would be forgiven even without it. In a chaotic regime it is not: any remaining difference, from another GPU model, a library or compiler version, the numerical precision, or a mathematically equivalent reordering of a sum, is amplified in the same way, so a run is reproducible only if the whole computation is repeated bit for bit. Numerical differences are the disturbance; the feedback dynamics decide whether it grows. How far it grows within a given budget also depends on its starting size, on the horizon and on how close the run comes to a change of answer, so a periodic loop can still change an answer when the starting difference is large. In bfloat16 the starting difference is already about 10^{-3}, so continuous feedback carries it into different answers within a short budget even where the float64 loop is not chaotic (below).

Whether rounding differences grow depends on the regime. Language-model inference is not bitwise reproducible: GPU kernels choose different reduction orders for different matrix shapes, so a request's result depends on its batch layout, the set of requests that the GPU processes together with it, and on the hardware [53, 54]. We ran the 448 labeled cells of the horizon check (192 chaotic, 192 periodic and 64 other cells of two float64 maps; Appendix B) twice in lockstep from the same start. The two runs differ only in the batch: all cells in one batch of n rows, or the same rows followed by n further rows whose steering vectors are those of the first n in reverse order, at the original temperatures ($2n$ rows). This changes the matrix shapes the kernels see; the first n rows compute the same cells (Figure 18, Table 7). In float64 the two runs first differ by about 6×10^{-15}, the rounding level. At 190 of the 192 chaotic cells this difference passes a relative distance of 10^{-2}, after a median of 102 steps, and grows on to order one (median final distance 0.91); at 173 of the 192 chaotic cells the most probable token of the loop's readout differs between

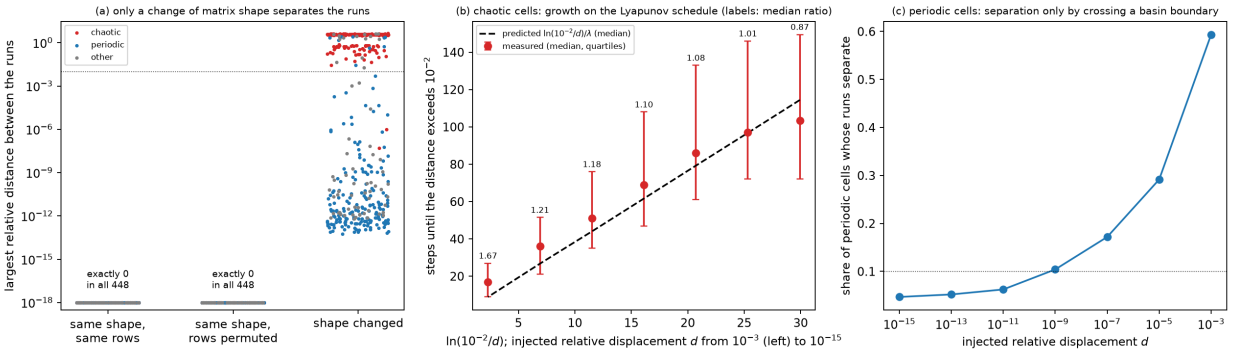

Figure 17: Fixed-shape control and dose test (Qwen2.5-0.5B-Instruct, the 448 cells of two float64 maps, 192 chaotic, 192 periodic and 64 with neither label, 1500 lockstep steps; `tools/repro_noise.py`, `tools/plot_fixed_shape.py`). (a) Largest relative distance between paired runs per cell: with the same matrix shapes, whether the same rows run twice or the rows are permuted, it is exactly zero in all 448 cells (plotted at 10^{-18}); a changed shape separates the chaotic cells. (b) With the shapes fixed, a displacement d injected at the start passes 10^{-2} after the number of steps $\ln(10^{-2}/d)/\lambda$ predicts (chaotic cells; medians and quartiles; labels: median ratio of measured to predicted). (c) Share of periodic cells whose displaced runs separate, presumably by crossing into another basin (the attractors reached were not classified).

the runs at some step of the last 500. At 182 of the 192 periodic cells the two runs stay within 10^{-2} for all 1500 steps, with a median largest distance of 2.3×10^{-12}, and 182 of the 192 periodic cells end the run within 3×10^{-12}; at only 7 of the 192 periodic cells do the most probable tokens differ at any step of the last 500, so a periodic loop almost never changes what it writes. The Lyapunov exponent accounts for when a chaotic run diverges: the step at which the distance passes 10^{-2} follows $\ln(10^{-2}/d_1)/\lambda$, where d_1 is the smallest nonzero distance in the run, with a median ratio of 0.91 and a log-log correlation of 0.73 over 190 cells. This uses the whole run and is therefore retrospective; an independent recomputation with the first nonzero distance instead gives the same rounded values. The two chaotic cells that never diverge have exponents of 0.013 and 0.0013, for which the predicted step lies beyond the run.

A fixed-shape control and a dose test. Two checks, fixed in advance (`docs/fs-design.md`), test this account directly on the same 448 cells (Figure 17). When the matrix shapes do not change, the two runs are bitwise identical over all 1500 steps in every cell, both when the same rows run twice and when the rows are permuted, so that each cell has other neighbors and another position in the batch. The difference therefore comes from the shapes alone: rows do not leak into one another, and the computation is deterministic. Changing the shape again separates 190 of the 192 chaotic cells and leaves 182 of the 192 periodic ones within 10^{-2}, as before. With the shapes fixed, we then displaced the start of each cell by a relative amount d from 10^{-15} to 10^{-3}. In the chaotic cells the displacement grows on the schedule the exponent sets. The step at which it passes 10^{-2} divided by $\ln(10^{-2}/d)/\lambda$ has medians of 0.87, 1.01, 1.08, 1.10 and 1.18 for $d = 10^{-15}$ to 10^{-7}, inside the frozen band from 0.7 to 1.3. The median step rises by 2.9 per unit of $\ln(1/d)$, against a median $1/\lambda$ of 3.9, within the frozen 30 %. At 10^{-5} and 10^{-3} the linear account fails (medians 1.21 and 1.67), as it must for displacements that are no longer small. In the periodic cells a displacement of 10^{-11} or less separates the runs in at most 6.3 % of cells. From 10^{-9} upward it separates 10 to 59 %, presumably because displaced states cross basin boundaries into another attractor's basin (the attractors reached were not classified). This fails a secondary criterion, which allowed at most 10 % up to $d = 10^{-5}$. The same holds without steering in Qwen2.5-7B, whose embeddings are untied.

Over 24 chaotic and 24 periodic temperatures of its unsteered loop, all 24 chaotic float64 runs diverge, after a median of 56 steps (the exponent again accounts for the step, median ratio 0.95), and at 23 of them the readout's most probable token differs at some step of the last 500. None of the 24 periodic runs moves past 10^{-2}; 22 end within 1.3×10^{-13} of each other and all within 2.2×10^{-12}. A second model family, unsteered Llama-3.1-8B, gives the same picture: all 20 chaotic temperatures diverge after a median of 68 steps (median ratio to the exponent's estimate 1.05), all 20 differ in their readout's most probable token in the last 500 steps, and 19 of the 20 periodic ones stay within 10^{-2}. At 70B the picture is unchanged. In Llama-3.3-70B-Instruct steered along hostility-vs-kindness at block 20 (float64 arithmetic on its stored bfloat16 weights, layers spread over two GPUs, 1000 steps), all 24 chaotic cells diverge, after a median of 114.5 steps (median ratio to the exponent's estimate 0.83, log-log correlation 0.74), and at 21 of them the readout's most probable token differs at some step of the last 300. None of the 24 periodic cells moves past 10^{-2}, and all end within 1.3×10^{-12}. In lower precision the rounding is larger. In float32 every chaotic cell of both models diverges, after a median of 20 steps (0.5B) or 16 (7B). In bfloat16, a common serving precision, the two layouts pass a distance of 10^{-2} after a median of one or two steps at every cell of the two Qwen models and at 38 of the 40 Llama cells, periodic ones included. With these kernels and at that precision, the loop was not reproducible across batch layouts in any regime we tested.

The same test at the level of answers, with a design frozen before the run, shows what this means for a task. We asked Qwen2.5-0.5B 96 held-out two-digit additions through a short soft-thinking pipeline: for 32 steps the model is fed the probability-weighted average of its token embeddings (the input of F), each new one appended to a growing context, and a fixed answer cue then reads the number off greedily. Each problem ran alone and in a batch with three fixed distractor problems of the same length. Nothing is padded, so the two runs compute the same quantities and differ only in how the GPU groups its sums for the two matrix shapes, that is, by rounding. In bfloat16 the extracted answer changed with the batch layout in 56 % of the problems at $T = 1$ and in 8 % at $T = 10^{0.2}$. At $T = 1$ the pipeline answers 23 of the 96 problems correctly alone and 22 in the batch, so most changes are between wrong answers; at $T = 10^{0.4}$, where it answers 90 of the 96 correctly in both layouts, no answer changed (`docs/pi/layout-answer-audit.json`). A forced greedy baseline with the same budget (32 greedy tokens, then the same cue) produced identical token sequences in both layouts for every problem. The excess disagreement of soft over greedy decoding, averaged over the two temperatures, is 32.3 percentage points (95 % interval $[26.5, 38.1]$; sign-flip $p = 10^{-4}$, the smallest attainable). In float64 no answer changed within the 32 steps. There, the median growth of an injected kick to the first soft token is 0.033 nats per step at $T = 1$ (range -0.13 to 0.24 over six problems and three kick sizes), and even a relative kick of 10^{-4} changes no answer. The bfloat16 discrepancies therefore do not require the float64 map's chaos; they are consistent with bfloat16's rounding differences accumulating through continuous feedback. A discrete argmax absorbs differences smaller than its margin, but not over a long feedback run in a pipeline that answers poorly ("Soft Thinking as published", below). (The forced greedy baseline's accuracy is zero and not informative: forced through 32 tokens, the model writes past its answer, and the cue then extracts another number.) The same frozen test in Qwen2.5-7B, whose soft-thinking pipeline answers 51 to 64 % of these problems, replicates the result. In bfloat16 the batch layout changed the soft-thinking answer in 53 of the 96 problems (55 %) at $T = 10^{0.2}$, the temperature at which the soft trace still matters. In 24 of these 53 problems one layout gave no number (an invalid answer, which the frozen design counts as a value); in the other 29 (30 % of the problems) both gave numbers, and they differ. The pipeline answered 49 of the 96 correctly alone and 61 in the batch; at $T = 10^{0.4}$ it answered all 96 correctly in both layouts, and no answer changed. At this precision even the forced greedy baseline changed in 10 of the 96, the known batch-variance effect. Averaged over the frozen design's two temperatures, $T = 1$ and $10^{0.2}$, as for the 0.5B model, the excess over greedy decoding

comparison	cells	reach 10^{-2}	median step	end with different tokens
batch layout, float64	chaotic	190 / 192	102	90.1 %
	periodic	10 / 192		3.6 %
batch layout, float32	chaotic	192 / 192	20.5	92.2 %
	periodic	72 / 192		22.9 %
float32 against float64	chaotic	192 / 192	20	92.7 %
	periodic	75 / 192		19.3 %
batch layout, bfloat16	chaotic	192 / 192	2	99.5 %
	periodic	192 / 192	2	95.3 %
7B unsteered, float64	chaotic	24 / 24	56	95.8 %
	periodic	0 / 24		0 %
7B unsteered, float32	chaotic	24 / 24	16	95.8 %
	periodic	4 / 24		8.3 %
7B unsteered, bfloat16	chaotic	24 / 24	2	75.0 %
	periodic	24 / 24	2	50.0 %
Llama 8B unsteered, float64	chaotic	20 / 20	68	100 %
	periodic	1 / 20		5 %
Llama 8B unsteered, float32	chaotic	20 / 20	10.5	100 %
	periodic	1 / 20		0 %
Llama 8B unsteered, bfloat16	chaotic	20 / 20	1	100 %
	periodic	18 / 20	2	45 %
Llama 70B steered, float64	chaotic	24 / 24	114.5	87.5 %
	periodic	0 / 24		0 %

Table 7: Divergence of two otherwise identical runs, by the cells' float64 labels (1500 steps; "end with different tokens": the most probable tokens differ at some step of the last 500). Top: Qwen2.5-0.5B, philosophy direction at layer 12. Middle: Qwen2.5-7B unsteered, 24 chaotic and 24 periodic temperatures of its long sweep. Bottom: Llama-3.1-8B-Instruct unsteered, 20 and 20. Last block: Llama-3.3-70B-Instruct, hostility-vs-kindness at block 20, 24 and 24 (1000 steps; tokens compared over the last 300). In float32 the periodic cells that diverge may change phase on the same cycle or switch between coexisting attractors; the run does not distinguish the two.

is 17.7 points ($[9.3, 25.6]$, $p = 2 \times 10^{-4}$); at $T = 10^{0.2}$ alone it is 44.8 points (53 against 10 of 96), and at $T = 1$, where the soft pipeline answers none correctly and changed one answer, it is negative. The forced greedy baseline is a computational comparator, not a useful-answer baseline: it answers 1 of the 96 correctly. In float64, at that temperature, the 7B's loop amplifies an injected perturbation (median +0.034 nats per step), injected kicks change answers (2 of 18 kicks: six problems and three kick sizes), and the rounding-level difference between the layouts changed 5 % of the answers within the 32 steps. Two horizons thus matter. Within a short budget, continuous feedback in low precision makes answers depend on the batch, and where the loop expands, even float64 rounding does. Over longer runs, a chaotic regime amplifies float64 rounding into different readout tokens, on the schedule its Lyapunov exponent sets.

An exploratory causal test locates the low-precision effect in the loop. Following a suggestion of an external review, we reran the 0.5B pipeline at $T = 1$ on 128 newly drawn additions, after a batch-1 reference run of each (exploratory; first on a workstation GPU, where the native pipeline's answers changed with the layout in 74 % of the problems, against 56 % on the H100 of the frozen test). In the batch of four, two arms start from the reference's prompt state and are answered by the same batch-1 decoder. In the free arm the soft phase runs on its own; in the reset arm the target's cache and logits are replaced by the reference's before every soft step, so that each step runs in the batch but only one at a time. In bfloat16 the answer changed in 88 of the 128 problems (69 %) in the free arm and in 7 (5.5 %) in the reset arm, a difference of 63 points (95 % interval $[55, 72]$; 82 discordant pairs against 1, exact McNemar $p = 2 \times 10^{-23}$), while accuracy stayed between 23 and 25 % in every arm. One batched step moves the target's logits by a median 1.5 % of their norm; in the free arm the median difference grows to 0.79 of it by step 32, while the reset arm's median per step stays between 1.3 and 3.1 %. In float32 the free arm's drift grows about 140-fold over the 32 steps but stays near 10^{-4}, and 3 of the 128 answers changed (none with reset). The test replicates on an H100 and in Qwen2.5-7B. On the H100, in bfloat16, the answer changed in 60 of the 128 problems (47 %) in the free arm and in 1 (0.8 %) with reset in the 0.5B model, and in 52 (41 %) against none in Qwen2.5-7B at $T = 10^{0.2}$ (exact McNemar $p \leq 5 \times 10^{-16}$); in the 7B model even float32 answers changed in 33 (26 %) against none. In bfloat16 on the H100 most batched steps reproduce the reference bitwise (the reset arm's median step distance is 0), and the free loop amplifies the rare step that does not. No arm's accuracy differs from the reference's by more than 5 points. Within this budget, the batch layout changes low-precision answers because the loop carries each step's rounding forward and amplifies it, not because the last steps alone are fragile; 128 problems per arm do not show that accuracy is unchanged. These counts compare the extracted numbers, as the frozen test does. Until 19 October 2026 the test compared the decoded answer texts, which also counts the same number followed by different text, and reported higher counts (102 and 14 on the workstation; 85 and 2, and 82 and 2, on the H100); the corrected and the original summaries are in `docs/rr`.

Soft Thinking as published. A test frozen before its data (`docs/pi-matrix-confirmatory.md`) ran Soft Thinking with its published settings on Qwen2.5-7B-Instruct in bfloat16. The settings are concept tokens restricted to the 15 most probable, $T = 0.6$ and 64 soft steps. The test used 256 GSM8K training problems that no earlier run had scored, on an H100 GPU. Only the batch layout changed, between batches of 8 and of 16 prompts of equal length without padding, and the answer changed for 21.1 % of the problems (54 of 256; 95 % interval clustered by batch $[16.3, 26.1]$). This confirms the prediction of more than 10 %. The batches deviated from the design in one respect: the harness filled each group of equal-length prompts from one pool of fillers, and where that pool ran short, both layouts received the same batch. Rebuilt from the tokenizer, 16 of the 256 problems

ran in identical batches in both layouts, and none of them changed; among the other 240 the rate is 54 of 240 (22.5 %). Some batches were also short (11 of the 78 batches of 8 and 26 of the 77 of 16; docs/pi/layout-answer-audit.json). Each secondary prediction compares another arm with this one on the same 256 problems (paired differences of flip rates, intervals clustered by batch, exact McNemar tests):

- replaying the reference run's fed-back tokens in both layouts: 5.9 % (difference 15.2 points, $[10.4, 20.2]$, $p = 8 \times 10^{-9}$); predicted lower, and it is;
- 16 soft steps instead of 64: 6.6 % (difference 14.5 points, $[10.0, 19.0]$, $p = 4 \times 10^{-7}$); predicted lower, and it is;
- ordinary greedy chain-of-thought decoding: 2.3 %, at 89 % accuracy (difference 18.8 points, $[13.7, 23.6]$, $p = 10^{-11}$); predicted lower, and it is;
- feeding back only the most probable token: 17.6 % (difference 3.5 points, $[-2.9, 10.3]$, $p = 0.35$); predicted lower, and the test cannot tell the two apart.

The three confirmed differences would survive a correction for the four tests by a wide margin, and the fourth is a difference the test does not resolve: discrete feedback through the same pipeline flips nearly as many answers as continuous feedback does, although the interval does not establish that the two are equal. The full vocabulary gave 22.3 % (no prediction). Most changes are between wrong answers: through the soft pipeline this model answers only 22 to 30 % of the problems correctly, and the layout changes whether the answer is correct for 6.2 % of the problems. In float64, injected perturbations mostly shrink over the same pipeline's observed runs. On 1312 and 1314 of the 1319 GSM8K test problems (full vocabulary and top-15), a perturbation of 10^{-10} ended below its initial size, and none changed an answer. These flips therefore do not require the float64 map's chaos: the feedback carries bfloat16 rounding differences forward until they change answers in a pipeline that answers only 22 to 30 % of the problems correctly, and it does so for discrete feedback as well.

How the dependence builds up. An exploratory study read the answers of the same pipeline after every doubling of the number of soft steps, from 2 to 128. It used 200 GSM8K test problems, two temperatures (0.6 and 1.2), and six steering strengths along the hostility direction (α from -0.02 to 0.05). The two layouts agree on 90 to 96 % of the answers after 2 steps and on 51 to 70 % after 128, in every configuration. Among the problems whose layouts disagree, the median first disagreement comes after 8 steps. Float64 copies of the same configurations, started 10^{-10} apart, separate the two sources of disagreement.

- Unsteered or weakly steered ($|\alpha| \leq 0.02$), the copies change the answer text of at most 1.5 % of the problems. The loop shows little answer sensitivity to these float64 perturbations; the bfloat16 disagreements are consistent with rounding carried forward by the feedback.
- At $\alpha = 0.05$ they change the answer text of 13.5 % of the problems at $T = 0.6$ and 27.5 % at $T = 1.2$, and the extracted number of 4.5 and 9.0 %. The loop itself is then sensitive, but accuracy has fallen to between 0.5 and 2.5 % in either layout.

In this model, then, the regime in which Soft Thinking is useful is the regime of accumulated rounding, not of chaos (docs/pi/h-horizon-summary.json).

Reasoning models that answer well. Qwen2.5-7B-Instruct answers few of these problems through the soft pipeline, so most of its flips are between wrong answers. A second test, also frozen before its data (docs/pi-matrix2-confirmatory.md), used reasoning models, as the Soft Thinking paper did: DeepSeek-R1-Distill-Qwen-7B, DeepSeek-R1-Distill-Llama-8B and QwQ-32B. A row thinks in concept tokens (top 15, $T = 0.6$) until its most probable next token ends the thinking block, or until 1536 steps. It then answers by sampling discrete tokens, with random numbers shared between the two layouts. The comparison arm thinks with discrete tokens sampled the same way. On GSM8K problems not used before in this study the batches held 8 or 16 prompts of equal length. On the integer-answer problems of MATH-500, each problem ran alone or as eight copies of itself. In the first run the tokenizer of DeepSeek-R1-Distill-Llama-8B loaded wrongly: it dropped the spaces and newlines of every prompt, and the model thought far less and answered 45 % of a check set correctly, against 83 % with the correct tokenizer in bfloat16 and 81 % in float64 (docs/pi/pm2x-llama-diag.json). Both of its cells were rerun with the correct tokenizer, an amendment after the data. Over the five primary cells, the batch layout changed 27.1 % of the soft-thinking answers (243 of 896; 95 % interval $[24.7, 29.5]$), confirming the prediction of more than 5 %; this is a weighted summary of these cells, which mix two tasks and two layout constructions, not a general susceptibility of the models. The prediction holds in every version of the pooling, including the original run (29.6 %) and the run without both cells of this model (21.9 %) (docs/pi/pm2fix-summary.json, docs/pi/pm2fix-math-summary.json).

- All three models changed their GSM8K answers with the layout alone: DeepSeek-R1-Distill-Qwen-7B for 14.5 % of the problems ($[10.9, 18.0]$) at 88 to 90 % accuracy, and QwQ-32B for 19.5 % ($[16.4, 22.3]$) at 86 to 87 %, and DeepSeek-R1-Distill-Llama-8B for 28.9 % ($[23.4, 34.4]$) at 74 to 77 %. The two layouts disagreed on correctness for 11.7, 14.5 and 18.0 % of the problems.

- Discrete sampled thinking also changed many answers. On GSM8K the rates were 10.2, 19.9 and 21.5 % for the same three models. Only in DeepSeek-R1-Distill-Llama-8B did soft thinking change more, by 7.4 points ($[1.9, 12.9]$; exact McNemar $p = 0.029$, not corrected for the five cells), and it answered correctly 4.7 points less often ($[-8.6, -0.8]$; McNemar $p = 0.096$). The intervals resample batches, while the exact McNemar test treats problems as independent pairs, so the two can disagree, as they do for this accuracy difference. Elsewhere no soft-minus-plain difference is significant, which does not show that the two are equal. In these pipelines the effect did not require continuous thinking: long sampled feedback in bfloat16 showed it too.

- On an A100 instead of an H100, the first model's soft rate was 11.7 % (difference 2.7 points, $[-1.6, 6.6]$), a separate replication on a second GPU type; the difference is not significant, which again does not establish equality.

- Where the budget cut the reasoning short, accuracy was lower and the layout changed more answers. On MATH-500, where half or more of the problems reached the thinking cap, it changed 61 % of the answers of DeepSeek-R1-Distill-Qwen-7B (44 to 48 % correct) and 67 % of those of DeepSeek-R1-Distill-Llama-8B (39 to 42 % correct).

- Nearly all GSM8K changes came from long runs. Of the 161 answers that the layout changed with soft thinking, 158 came from problems in which the thinking budget (1536 steps) or the answer budget (256 tokens) was reached in at least one layout, and so did all 132 changes with sampled thinking. Among problems that finished within both budgets in both layouts, the layout changed 3 of 480 soft answers and none of 477 sampled ones (docs/pi/layout-answer-audit.json).

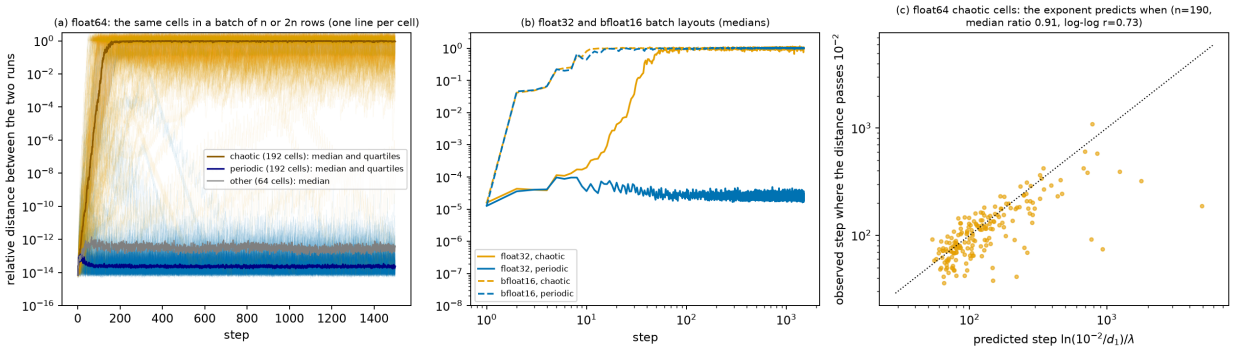

Figure 18: Only the batch layout differs. The 448 cells (192 chaotic, 192 periodic, 64 with neither label) of Qwen2.5-0.5B-Instruct (philosophy direction, layer 12) run twice from the same start, alone in a batch or followed by n further rows with the steering vectors in reverse order. (a) float64: the relative distance between the two runs for every chaotic (orange) and periodic (blue) cell, by the cells' float64 labels, with the medians and interquartile ranges of each label (dark) and the median of the 64 cells with neither label (gray); the chaotic runs separate to order one; the median periodic distance stays at the rounding level, with exceptions. (b) Medians for float32 and bfloat16 batch layouts; in bfloat16 both kinds separate within a few steps. (c) float64 chaotic cells: the step at which the distance passes 10^{-2} against $\ln(10^{-2}/d_1)/\lambda$, with d_1 the smallest nonzero distance in the run and λ the map exponent (a retrospective comparison); dotted, equality.

These runs did not use the Cold Stop rule of the Soft Thinking paper, which ends thinking once the entropy has stayed below 0.01 nats for 256 steps. An exploratory rerun used it, with a thinking budget of 4096 steps (`docs/pi/pm2cs-design.md`, `docs/pi/pm2cs-summary.json`). Cold Stop ended thinking in only 6 of the 512 soft rows and in none of the 512 sampled rows; the models nearly always ended it themselves. The layout still changed the answers: 8.6 % (soft) and 17.2 % (sampled) for DeepSeek-R1-Distill-Qwen-7B on 128 GSM8K problems, 50.0 and 54.7 % for it on MATH-500, and 17.2 and 9.4 % for QwQ-32B on 64 GSM8K problems. The soft-minus-sampled difference changes sign between these cells. On the same type of GPU, every row that ended thinking on its own within the original 1536 steps was reproduced exactly in its saved switch step and answer, so the changes come from the rows the original budget cut short; the DeepSeek-R1-Distill-Qwen-7B GSM8K rerun ran on a different H100 variant.

Is this chaos? We ran float64 copies of the same pipeline. Each copy's first soft token was displaced from the base row's by a relative 10^{-10}, and all copies shared the base row's random numbers. The three models respond differently. In DeepSeek-R1-Distill-Qwen-7B, 216 of 256 copies ended the thinking at a distance from their base row below the initial displacement, and 2.3 % changed the answer. These observations are consistent with its layout flips arising mainly from bfloat16 rounding carried through feedback, as in Qwen2.5-7B; they do not establish global contraction. In QwQ-32B and DeepSeek-R1-Distill-Llama-8B, more than half of the copies ended above the initial displacement (131 of 192 and 145 of 256). Fitted over the steps where the distance lay between ten times the displacement and 10^{-3}, the median growth rates were 0.011 and 0.008 per step. The displacement alone changed 13.5 and 18.4 % of the answers, respectively. In those models the loop's own sensitivity contributes to the answer changes. It is a weak, finite-time instability over about a thousand thinking steps, and it acts on top of the rounding; it is not an asymptotic chaos measurement. Some copies ended thinking before their base row (14 in R1-Llama, 23 in QwQ), and the copies of one problem are not independent (`docs/pi/pm2f-summary.json`, `docs/pi/pm2f64b-summary.json`; 128 problems for each 7–8B model and 96 for QwQ-32B).

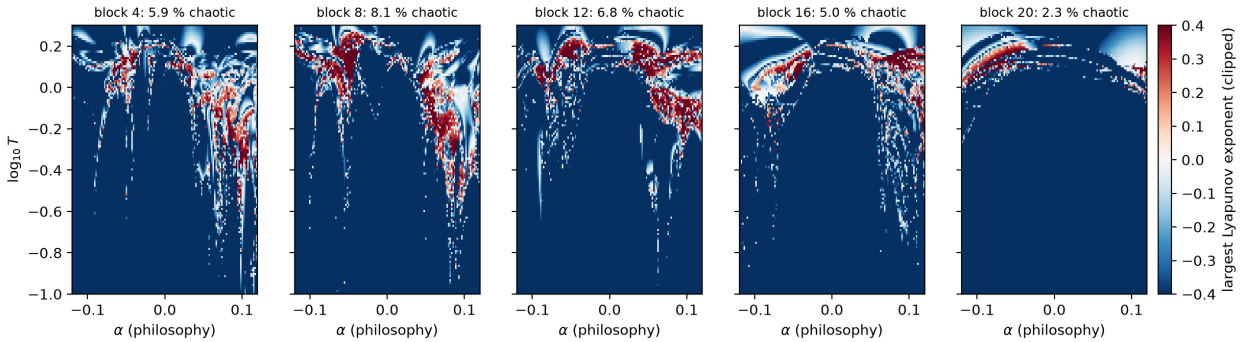

Figure 19: The philosophy direction of Qwen2.5-0.5B-Instruct, built with the same recipe at each of five blocks of 24: largest Lyapunov exponent on 128×128 float64 cells over the steering strength and $\log_{10} T \in [-1, 0.3]$. The share of chaotic cells (finite-horizon exponent above 10^{-3}, no period up to 48) is given above each panel.

Where a steering vector enters decides how much chaos it causes. Steering vectors are a common tool for controlling and probing models [35, 38], usually added at one layer chosen by hand. In Qwen2.5-0.5B the philosophy direction, built with the same recipe at blocks 4, 8, 12, 16 and 20 and scaled to each block's residual norm, makes 5.9, 8.1, 6.8, 5.0 and 2.3 % of the same 128×128 plane chaotic (Figure 19). Early steering carries the chaos down to $\log_{10} T = -0.95$ (block 8) and -0.91 (block 4), while at block 20 the chaotic cells lie between $\log_{10} T = -0.09$ and 0.22, three fifths of them above 0.1. In the strip $|\alpha| < 0.02$ every block gives 0.8 to 2.1 %. Llama-3.3-70B behaves the same way. Steered along its hostility direction (a contrastive valence direction; Section 10), it is chaotic in 8.0 % of a 24×24 map at block 20 of 80, down to $\log_{10} T = -0.14$, against 4.3 % at block 35, where the chaos starts at $\log_{10} T = 0.01$, 4.3 % at block 50 (32 by 32 cells, a two-GPU map from before the precision fix of Section 14), and 0.7 % at block 65, all above $\log_{10} T = 0.19$. These 70B labels are those of the map protocol, one twin offset, and so are candidates. A logit bias, the latest possible point of entry, gives chaos mostly above $\log_{10} T = 0.1$ (Section 3). The gain split of Section 3, repeated along steering lines at all five blocks of the 0.5B model, shows where the depth enters. At low temperature ($\log_{10} T \leq 0$) the chaotic cells of blocks 4, 8 and 12 have expanding readout factors (medians +0.74, +0.46 and +0.46 nats per step over 54, 56 and 21 cells) and contracting network factors (-0.47 to -0.21). At block 16 the readout factor of such cells is only +0.15, and block 20 gives just 6 of them. Early injection creates token competition at low temperature more readily; late injection mostly does not. Its chaos lives at higher temperature, where the network factor carries it (block 20: network +1.09, readout -0.78; Figure 8e).

Two steering vectors interact. Steering at two blocks at once, gratitude at block 4 and hostility at block 16 of Qwen2.5-0.5B, makes 21.7 % of a 96×96 plane chaotic at $\log_{10} T = -0.1$. The grid has no row or column at exactly zero, so we computed the single pushes separately, each with the other exactly absent and on the same grid: they make 12.5 % (gratitude at block 4 alone) and 5.2 % (hostility at block 16 alone) of their lines chaotic. Large chaotic regions lie where neither push alone is chaotic (Figure 20): 1561 cells, 16.9 % of the plane and 78 % of its chaotic cells, are chaotic although each push alone, at the same strength, is not, and in 1525 of them each push alone is periodic. Over the four planes (two block orders, $\log_{10} T = \pm 0.1$) this share is 8.3 to 16.9 % of the plane, and 7.8 to 16.5 % with both single pushes periodic. With the blocks exchanged (hostility at 4, gratitude at 16) the plane is 23.4 % chaotic, and at $\log_{10} T = 0.1$ both pairs give 33 and 42 %. Each plane keeps a periodic hole around zero steering. Combining steering vectors is therefore not

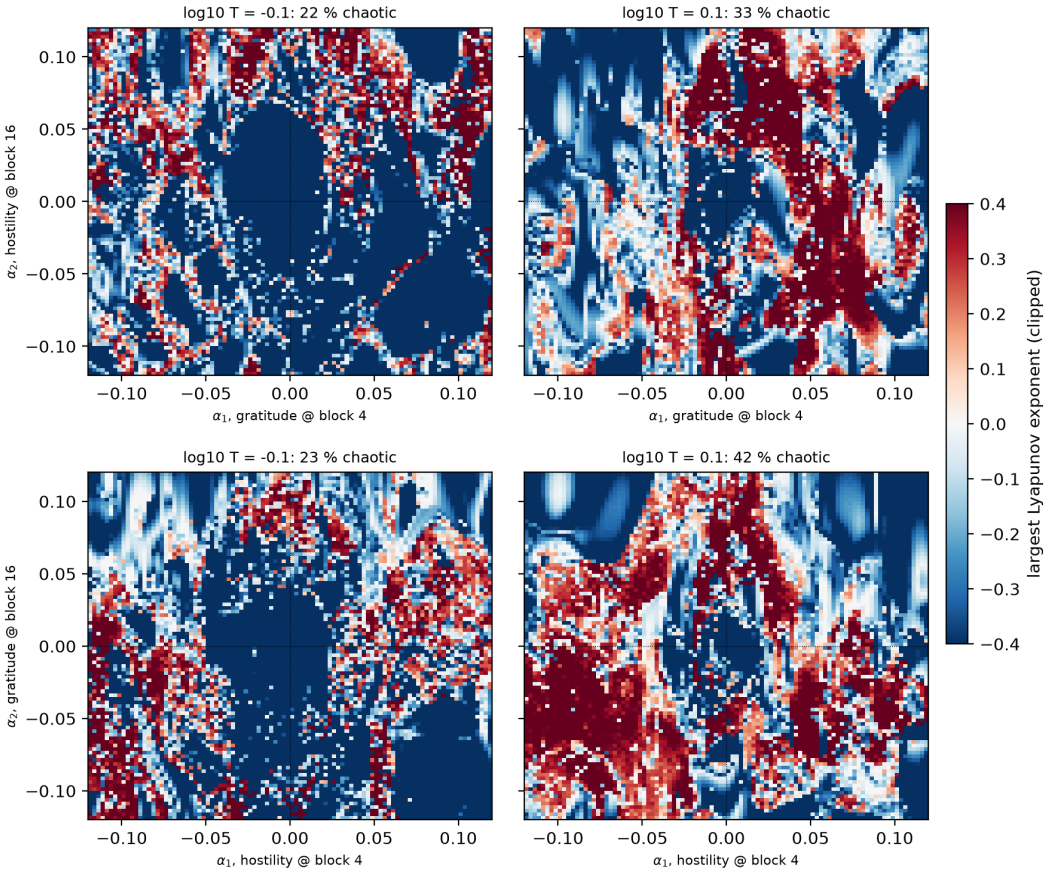

Figure 20: Two steering directions at two blocks of Qwen2.5-0.5B-Instruct: largest Lyapunov exponent on 96×96 float64 cells over the two strengths, at $\log_{10} T = -0.1$ (left) and 0.1 (right). Top: gratitude at block 4, hostility at block 16; bottom: the blocks exchanged. Dotted lines mark the zero axes; the grid does not sample them, and the single-push controls quoted in the text were computed separately on exact zero lines.

additive in its effect on the loop's stability: a composition of two pushes can be chaotic where each push alone is periodic.

Does a chaotic loop predict worse answers in ordinary decoding? Not beyond the model's own uncertainty. Two exploratory designs found no accuracy signature of the regimes that survives controls (Appendix D). A frozen test then labeled 4000 newly drawn products of two two-digit numbers by the regime of the unsteered Qwen2.5-7B-Instruct's own loop at $\log_{10} T = 0.40$. Chaotic-loop prompts are answered less accurately by ordinary sampling, by 7.1 points (95 % interval $[4.8, 9.3]$), but the gap vanishes once seven difficulty covariates, fixed before sampling, enter the fit (label coefficient +0.15 points, $[-1.4, 1.7]$, permutation $p = 0.84$), and in a post hoc fit the first-token entropy alone absorbs nearly all of it. This fits the measurements of Section 3: the loop's chaotic cells have more competing next tokens than its periodic ones, and a model unsure of its first token also answers worse. Within a prompt no association is detected either, and accuracy does not peak at the edge of chaos: deep periodic, edge and deep chaotic prompts are answered alike, with the intervals of all three adjusted contrasts within $[-3.1, 3.0]$ points. These are nulls at this precision; they do not show that the regime has no effect.

7 Period adding written in text

Between a period-2 region and the chaos lies a fan of stripes of periods $6, 8, 10, \ldots$ Their orbits, numerically validated in float64, can be read as text (Table 8). Each is ". This" followed by k pairs "is X", with period $2k + 2$, so each stripe adds one pair. The period-2 region itself is the 2-cycle "is it". At the stripe centers the orbits' linear relaxation is dominated by one direction: the leading multiplier, 0.0025 to 0.086, exceeds the next by more than three orders of magnitude (the next falls from 3×10^{-6} at period 6 to 2×10^{-14} at period 20; Table 8), so near these orbits a small deviation soon lies along a single direction (the smallest values are at the limit of float64 accuracy). This is a property of these orbits, not of the map as a whole. The two leading multipliers of the 2-cycle, -0.271 and -0.058, differ by a factor of only about 5 (the second modulus is about 0.037 near its fold), the next modulus at the period-3 flip of Section 9 is 0.088, and one period-6 orbit in Appendix E has a complex leading pair.

period	orbit (leading tokens)	μ_1	$\lvert\mu_2\rvert$
6	. This is a is a	$+0.0052$	3×10^{-6}
8	. This is there is a is a	$+0.0110$	3×10^{-7}
10	. This is there is it is a is a	$+0.0025$	2×10^{-7}
12	. This is there (is it)^2 is a is a	$+0.0163$	4×10^{-8}
14	. This (is it)^4 is a is a	$+0.0369$	2×10^{-9}
16	. This (is it)^5 is a is a	$+0.0716$	8×10^{-11}
18	. This (is it)^6 is a is a	$+0.0792$	1×10^{-12}
20	. This (is it)^7 is a is a	$+0.0857$	2×10^{-14}

Table 8: Periodic orbits of the stripe family at the stripe centers along $\log_{10} T = -0.115$ ($\alpha = 0.077872$, 0.071674, 0.067766, 0.065550, 0.064443, 0.063857, 0.063534, 0.063335), computed with the current all-float64 map (norms and rotary tables patched): Newton residuals 6×10^{-15} to 8×10^{-13}, leading multiplier μ_1 and next modulus $|\mu_2|$ from Rayleigh–Ritz values. Each step in period adds one "is X" pair to the text. The 2-cycle "is it" itself has multipliers -0.2706 and -0.0585 at $(\log_{10} T, \alpha) = (-0.17263, 0.061579)$. The period-2 and period-8 orbits were also reproduced by an independent implementation, and the period-8 orbit is certified by a computer-assisted proof (Appendix L).

Table 9 shows what the map writes: two of these loops, cycle by cycle, next to chaotic cells between the stripes. The chaotic cells reuse the phrases of the neighboring loops but do not settle into a fixed order.

α	regime	H	leading token at 40 consecutive steps
0.071674	period 8	1.9	` is a is a. This is there| is a is a. This is there| is a is a. This is there| is a is a. This is there| is a is a. This is there`
0.072940	chaotic ($\lambda = 0.13$)	2.4	` a is a. This is there is a is↩The is there is a is↩↩The is↩↩The is↩↩The is a is a. This is there is a is↩The is there`
0.067766	period 10	1.8	` is a. This is there is it is a| is a. This is there is it is a| is a. This is there is it is a| is a. This is there is it is a`
0.069060	chaotic ($\lambda = 0.20$)	2.4	`. - certainly↩↩The is there is it is a.↩↩The is a is a. This is there is it is a is↩↩The existence isThe is it is it is a. This`

Table 9: What the attractors write. Qwen2.5-0.5B-Instruct soft-token map steered along philosophy at layer 12, at $\log_{10} T = -0.115$: the leading token at 40 consecutive steps after 600 transient steps, for the period-8 and period-10 loops of Table 8 (bars mark the start of each cycle) and a chaotic cell beside each. λ: largest exponent over 2000 steps after 1000 transient steps, the same to two decimals at twin offsets 10^{-6} and 10^{-9}. H: mean entropy of the next-token distribution in nats. ↩: line break.

8 Mechanism: a saddle-node bottleneck

The fold. At $\log_{10} T = -0.115$ the 2-cycle "is it" dies in a fold at $\alpha_c = 0.06279827$ (float64 multiplier 1 to 3×10^{-13}, next modulus 0.037; Figure 21): it meets an unstable 2-cycle, and both vanish. In the critical coordinate u of the second iterate $G = F^2$, the fold's normal form is $u' = u + a\varepsilon + bu^2$ with $\varepsilon = \alpha - \alpha_c$; here $a = -4.834$ per unit α and $b = -25.4285$, so $ab = 122.93$. Here q is the critical eigenvector of DG with $|q| = 1$, ℓ the left one with $\ell^\top q = 1$, $a = \ell^\top \partial G/\partial\alpha$ and $b = \frac{1}{2}\ell^\top D^2 G[q, q]$ [57]. The normal form is the quadratic truncation of G on its center manifold: it neglects terms of order u^3, εu and ε^2, and it applies because the multiplier $+1$ is simple, a and b do not vanish, and the other multipliers lie inside the unit circle [50]. No symmetry of the map forces a pitchfork and no persistent branch forces a transcritical exchange, so these conditions make the point a fold. Integrated as a flow, it gives the passage from v_{in} to U in closed form, $n(v_{\text{in}}) = [\arctan(|b|U/\kappa) - \arctan(|b|v_{\text{in}}/\kappa)]/\kappa$ with $\kappa = \sqrt{ab\,\varepsilon}$; a full passage approaches π/κ only when $U \gg \sqrt{|a|\varepsilon/|b|}$. The stable branch gives an independent estimate: extrapolating its multipliers through $(1-\mu)^2/(4(\alpha_c - \alpha)) \to ab$ yields 123.1 with a quadratic fit (120.6 with a linear one). An independent implementation of the forward map and solver recovers α_c to 6×10^{-15} and a to 3×10^{-12}. Its b, from nested automatic differentiation, is -25.42846; our finite-difference value differs from it by about 2.2×10^{-6} relatively, which is the finite-step bias of the difference table. The resulting $ab = 122.927$ leaves every passage prediction of Table 10 unchanged.

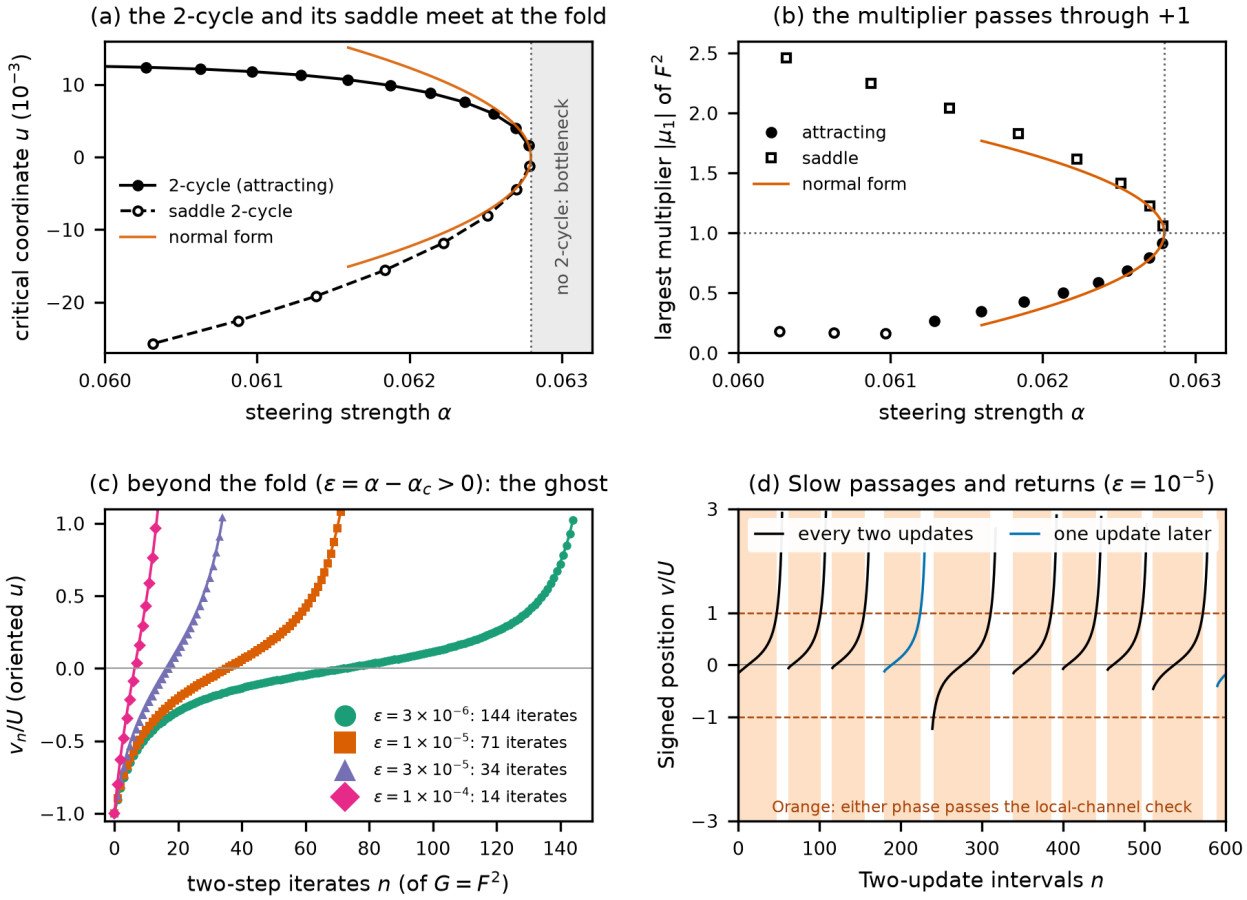

Figure 21: The fold behind the stripes: the 2-cycle "is it" at $\log_{10} T = -0.115$ (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12). (a) Continued in α, the attracting 2-cycle (filled) and its saddle partner (open) meet and vanish at $\alpha_c = 0.06279827$. The vertical coordinate is the projection u of one point of each orbit on the fold's critical direction; orange, the normal form's fixed points $u^* = \pm\sqrt{|a|(\alpha_c - \alpha)/|b|}$. (b) The largest multiplier of $G = F^2$ on the two orbits reaches +1 at the fold, following the normal form's $1 \mp 2\sqrt{ab(\alpha_c - \alpha)}$ (open circles: a complex leading pair). (c) Beyond the fold the 2-cycle is gone, but its ghost, the slow bottleneck that a saddle-node bifurcation leaves behind [55, 56], slows the dynamics: controlled passages from the channel's entrance $v = -U$ to its exit $v = U$ (v the oriented critical coordinate, $U = 0.1/|b|$), measured on the network (markers, one shape per ε) and iterated from the normal form $v' = v + |a|\varepsilon + |b|v^2$ (lines), take 144, 71, 34 and 14 iterates, as predicted. (d) The same free trajectory at $\varepsilon = 10^{-5}$, now shown as signed position v/U on a linear axis. Black samples every two updates; blue reads each saved state one update later, recovering passages hidden by a one-step phase shift. The dashed lines are $v/U = \pm 1$; zero is the channel's middle, not a fixed point. Orange *time intervals* satisfy both $|v| < U$ and transverse remainder norm < 0.05 at either reading; white intervals fail this local criterion, not a test for chaos. Values outside $|v| \leq 3U$ are omitted, leaving gaps rather than connecting distant excursions. Complete highlighted passages last 40–70 two-update intervals; the first and last are censored by the displayed horizon. All traces in (d) come from saved model states, not the normal form. Branch residuals, fold, normal form and passages in float64; branch multipliers in float32. `tools/plot_fold.py`.

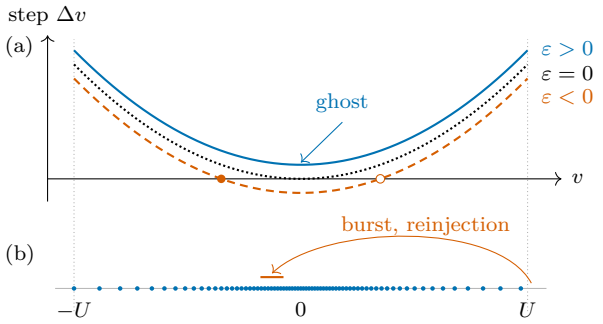

Figure 22: The bottleneck behind the stripes, drawn from the normal form $u' = u + a\varepsilon + bu^2$ with the network's coefficients ($a = -4.834$, $b = -25.4285$), in the oriented critical coordinate v. (a) The step of one G-iterate, $\Delta v = |a|\varepsilon + |b|v^2$, across the channel from $-U$ to U. Before the fold ($\varepsilon = -10^{-5}$, dashed) the step vanishes at two fixed points of G: the attracting 2-cycle (filled) and its unstable partner (open). At the fold ($\varepsilon = 0$, dotted) they merge. Beyond it ($\varepsilon = 10^{-5}$, solid) no fixed point remains, but the step stays small near $v = 0$: the ghost. (b) The 71 iterates of a full passage at $\varepsilon = 10^{-5}$, crowded where the step is smallest; the network takes 71 as well (Table 10). A free trajectory leaves at U, makes a burst, and is reinjected slightly upstream of the middle (bar), so its passage is shorter than a full one.

Controlled passages. Figure 22 shows the bottleneck that these passages cross. In the oriented coordinate $v = \mathrm{sign}(b)\, u$ of Appendix B, we start trajectories at $v = -U$ or 0 ($U = 0.1/|b|$) and count G-iterates until $v > U$. At four distances from the fold, the measured counts match the iterated normal form within one iterate (Table 10, Figure 23), and a displacement of the start of the same size along one random direction that leaves u unchanged changes the count by at most one (other directions and sizes were not tested).

$\varepsilon = \alpha - \alpha_c$	full: measured / predicted	midpoint	full + transverse kick
3×10^{-6}	144 / 144	73 / 74	143 / 144
1×10^{-5}	71 / 71	36 / 37	70 / 71
3×10^{-5}	34 / 34	17 / 18	33 / 34
1×10^{-4}	14 / 14	7 / 7	13 / 14

Table 10: Bottleneck passages (G-iterates). Predictions iterate $u' = u + a\varepsilon + bu^2$ with a, b computed from the network's derivatives; no parameter is fitted.

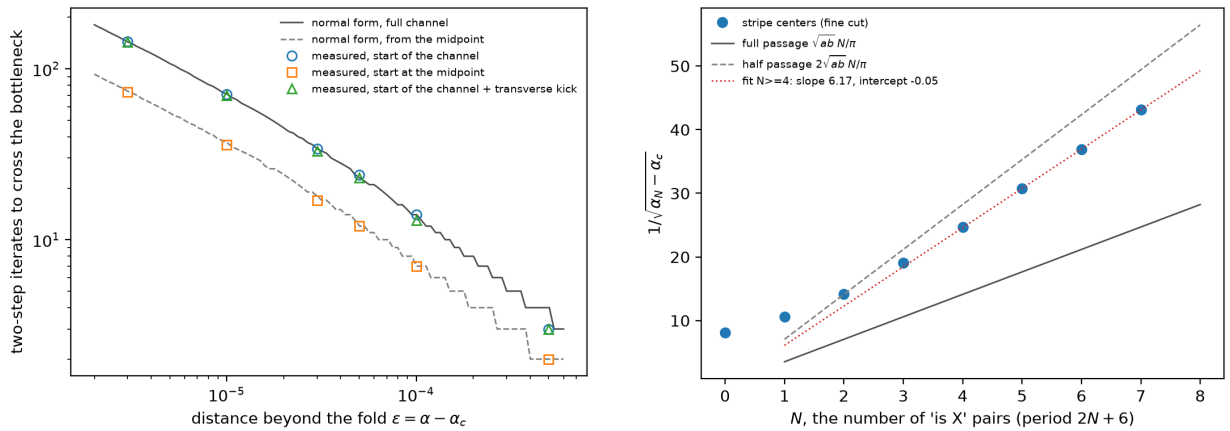

Figure 23: Bottleneck passages and the window law at the fold of the 2-cycle "is it" (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, $\log_{10} T = -0.115$). Left: measured passage times against the normal-form predictions. Right: the window law, with stripe centers plotted as $1/\sqrt{\alpha_N - \alpha_c}$ against the lap count N, together with the full- and half-passage predictions.

The totals of Table 10 hide a small error lap by lap: the quadratic formula overstates the move on the way into the channel and understates it on the way out, by up to 7.6% at $\varepsilon = 10^{-5}$ and 14.3% at 10^{-4}. Over a full passage the two nearly cancel: at $\varepsilon = 10^{-5}$ the network takes 36 laps to the channel's middle and 35 to its exit, the formula 35 and 36. The formula drops terms of order u^3 and εu, which change sign across the channel; a fit to the laps at five distances, made after the fact, needs both ($R^2 = 0.97$).

Reading the free trajectory. The coordinate $u = \ell^\top(x - x_c)$ is a signed projection of the feedback state x relative to one point x_c of the fold's 2-cycle; it is neither a token probability nor the full state-space distance. Panels (c,d) use $v = -u$, oriented from entrance to exit, and $U = 0.1/|b| \simeq 0.00393$. The sign is retained on a linear axis, so crossing $v = 0$ is smooth. Beyond the fold the local increment $\Delta v = |a|\varepsilon + |b|v^2$ is smallest there but remains positive: the trajectory slows, passes through, then accelerates out of the channel. No fixed point remains to capture it. Because a burst can shift the alternation by one update, panel (d) shows both $v(x)$ and $v(F(x))$ as separate traces at each two-update interval. Orange shading requires, at either reading, both $|v| < U$ and $d_\perp = \|(x - x_c) - uq\| < 0.05$, with q the normalized right critical eigenvector and $\ell^\top q = 1$. This transverse remainder is checked from the saved states, including the one-update-later states; a small projection alone is insufficient. White means only that neither reading meets this local-channel criterion. The view omits $|v| > 3U$, so the gaps denote excursions outside the displayed coordinate range, not missing data. The shading covers 445 of the 600 intervals (390 at the original reading and 55 at the shifted reading). This is a descriptive check on one trajectory, not an independent proof of chaos.

Natural reinjection and the window law. Left to itself, a trajectory near the fold alternates between laminar phases, in which it crawls through the bottleneck close to where the 2-cycle was and moves almost periodically, and bursts, short excursions out of the channel after which it is reinjected into it (Figure 21d). This is the slow-passage and reinjection pattern associated with type-I intermittency [58]; the finite alternation alone does not establish persistent chaos. Section 8.1 tests the global return and shows that repeated passages can precede periodic capture. In the local theory the laminar phases lengthen as $\varepsilon^{-1/2}$ toward the fold [59]; laminar is the fluid-dynamics

term for the smooth stretches between turbulent bursts. Free trajectories re-enter the channel slightly upstream of the bottleneck's middle (median entry $v \approx -3 \times 10^{-4}$ to -7×10^{-4} at the four distances, in the same oriented coordinate; single entries spread wider). From each measured entry point the normal form predicts the natural laminar durations: for complete phases, which stay within $|u| < 3U$ and within 0.05 of the channel in the other directions, the median signed forecast error at the four offsets up to $\varepsilon = 10^{-4}$ lies between -7.7% and -4.1% (signed errors -4.88, -4.08, -7.41 and -7.69%), although single phases differ by up to 71%; on average a natural passage lasts 0.60 to 0.67 of the normal form's passage through the same channel from end to end. A window of period p spends $N = (p-6)/2$ iterates of G in the bottleneck, one for each "is X" pair it writes. We read each pair as one G-iterate in the bottleneck, as the texts suggest (the laps were not identified geometrically). For a channel much longer than the bottleneck a full passage takes about $\pi/\sqrt{ab\,\varepsilon}$ iterates [55, 59], so a loop whose passage is a fraction f of a full one sits at the distance $d_N = \alpha_N - \alpha_c$ from the fold for which $d_N N^2 = f^2\pi^2/(ab)$. The stripes therefore follow the N^{-2} law of a saddle-node ghost [55, 56] (log-log slope -2.00, below), with about a third of the full-passage prefactor $\pi^2/(ab) = 0.0803$; the two fractions agree because the prefactor scales as the square of the passage fraction ($0.58^2 \approx 1/3$). The distances $d_N = \alpha_N - \alpha_c$ of stripe centers located on a float32 cut along $\log_{10} T = -0.115$ give $d_N N^2 = 0.0263$–0.0265 for $N = 4$–7 (log-log slope -2.00), compared with $\pi^2/(4ab) = 0.0201$ for a half passage. Two-step return maps of p("it") (Figure 24) show the near-tangency.

Numerically estimated stripe edges along $\log_{10} T = -0.115$ (Table 12) sharpen this picture. Each window is bounded by a fold, where its orbit is born, and by a flip. For periods 6–12 the fold faces away from the accumulation point. For periods 16–20 the sides are exchanged, and the period-14 window is bounded by flips on both sides. The fold edges of periods 16–20 follow $\alpha_N - \alpha_c = C/(N + N_0)^2$ with $C = 0.0266$ and $N_0 = 0.15$. Their slope, a rise of 6.13 per lap in $1/\sqrt{\alpha_N - \alpha_c}$, is 1.74 times the full-passage slope $\sqrt{ab}/\pi = 3.53$ of the normal form. That corresponds to an effective passage of 0.58 of a full one, which the measured natural reinjection gives independently (Table 11).

What a free trajectory writes. Rerun with its most probable tokens recorded, the trajectory of Figure 21d came back with identical coordinates over 600 laps, so the words below are those of the plotted trajectory. A laminar phase writes "it is it is". The state stays soft: at the step that writes "it", that token has about 27%, "there" 20% and "a" 19%, and over the first phase "a" gains only from 18.7% to 19.1%. When "a" overtakes "it", at lap 58, the sentence ends, "is a is a . This is", the same pieces as the stripe orbits' text, and the state falls back into the channel. Some bursts restart the sentence after a line break in a way that shifts the text by one step. Reading only every two updates then follows the other phase of the almost-repeating cycle. The blue trace in Figure 21d applies one further update to each saved state and exposes those passages directly: the first lasts 45 intervals, and the next is cut off by the 600-interval horizon. Both readings together satisfy the coordinate and transverse criterion for 445 of 600 intervals (74%), compared with 390 (65%) at the original reading alone. Of the other three free trajectories at this distance, two alternate in the same way and one settles on a different loop, "certainly . Hello". This is one trajectory, read after the fact to explain the figure (run FW, `docs/figure-walkthroughs/`).

measurement	value	full-passage value	f
free trajectories, four distances	mean laminar phase	passage through the whole channel	0.60–0.67
window centers, $N = 4$–7	$d_N N^2 = 0.0263$–0.0265	$\pi^2/(ab) = 0.0803$	0.57
fold edges, periods 16–20	slope 6.13	$\sqrt{ab}/\pi = 3.53$ (ratio 1.74)	0.58

Table 11: Three related comparisons of the passage fraction f (the free passages against the normal form's passage through the same finite channel, the window values as effective prefactors relative to the long-channel passage; all three share the fold and its coefficients, so they are not independent): the share of a full passage that a loop or a free trajectory spends in the bottleneck. The fold edges' slope is the rise of $1/\sqrt{\alpha_N - \alpha_c}$ per unit of N. A half passage would give $f = 1/2$ and $d_N N^2 = \pi^2/(4ab) = 0.0201$.

period	edge toward larger α	edge toward smaller α
6	fold 0.0790337	flip 0.0766947
8	fold 0.0726585	flip 0.0706525
10	fold 0.0684572	flip 0.0670594
12	fold 0.0658982	flip 0.0651944
14	flip 0.0645850	flip 0.0642958
16	flip 0.0639284	fold 0.0638001
18	flip 0.0635683	fold 0.0635034
20	flip 0.0633535	fold 0.0633183

Table 12: Edges of the stripe windows along $\log_{10} T = -0.115$ ($\alpha_c = 0.0627983$), from the window centers of Table 8. Flips: critical multiplier a real -1 to 3×10^{-10}. Folds: the multiplier nearest $+1$ extrapolated to 1 through $(1 - \mu)^2$, which is linear near a fold; last converged multiplier 0.991–0.998.

We tested whether the measured reinjection predicts windows beyond the fitted range. A model with one fixed entry point (the median measured one) and one lap offset fitted to the fold edges of periods 16–20 places the fold edges of periods 22–30 3 to 9 % closer to the fold than the constant-prefactor law. Both predictions were frozen before a float64 line of 3200 cells (spacing 10^{-7}) located the windows. The measured edges follow the law within 3 % in distance from the fold through period 30 (within 1 % at periods 22, 24 and 28) and fall outside the model's range except at period 30. So the law extrapolates, but one fixed reinjection point does not predict it. The measured entry points move with ε (-3×10^{-4} to -7×10^{-4}). The exponent of the law is near 2, its prefactor empirical. Over the fold edges of periods 16 to 32 ($N = 5$ to 13), $\ln(\alpha_N - \alpha_c)$ against $\ln N$ has slope -2.00, while a free fit of $C/(N + N_0)^\gamma$ with the offset N_0 also fitted is not well determined. The estimated localization scales of the edges are about 2×10^{-5} to 2×10^{-4} of their distances (extrapolated locations for periods 16 to 20, and a uniform-within-cell error model for the grid boundaries of periods 22 to 32), so the misfits of 1 to 3 % of any such law are consistent with corrections to its asymptotic form rather than noise: γ runs from 2.13 to 2.45 as the smallest N included rises from 5 to 8, while the fixed exponent 2 with an offset fits within 3 % (within 1 % from $N = 8$), so the data do not pin the exponent to exactly 2. At these distances, however, the bottleneck is not narrow compared with the counting window: at the fold edges of periods 16 to 32 its scale $\sqrt{|a|\varepsilon/|b|}$ is 1.35 to 3.51 times U, and a full passage lasts only 9 to 23 iterates of G. The windows therefore lie outside the asymptotic regime in which the normal form fixes the prefactor [58, 59], and the factor of about one third is a measured value, not a prediction. The same

inversion of the bottleneck time gives the N^{-2} tail of relaxation times over inputs found in a small task-trained looped transformer [21]; here N counts the laps of a periodic window.

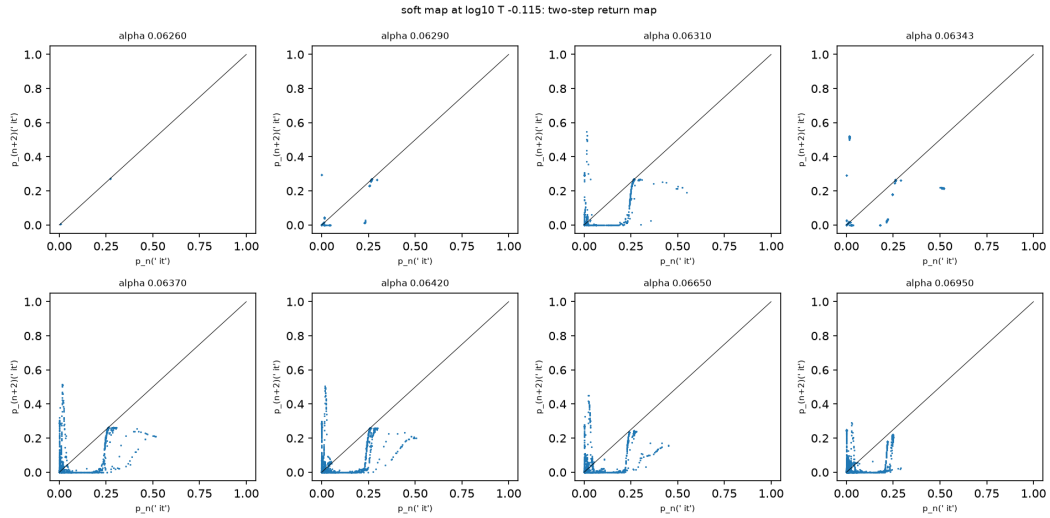

Figure 24: Two-step return maps of the probability of "it" near the fold of the 2-cycle "is it" (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12), at increasing distance from the fold.

A prospective test. The same reduction, with both coefficients computed from the network's derivatives, predicts relaxation times and spectral gaps without fitting in a small task-trained looped transformer [21]. To test the normal form at folds that were not used while the bottleneck explanation was worked out, an independent implementation located two new folds and computed their normal-form coefficients by nested automatic differentiation. All 24 passage predictions were pushed in a frozen design before any test trajectory was run. Target A is the 2-cycle's fold at a new temperature: $\log_{10} T = -0.140$, $\alpha_c = 0.0670549$, $a = -7.055$, $b = -24.927$. Target B is the fold of the period-8 orbit of Table 8, continued in temperature at $\alpha = 0.0717$: $\log_{10} T_c = -0.1033462$, $G = F^8$, $a = -11.900$ per unit $\log_{10} T$ and $b = -152.05$. Of the 24 forecasts (four distances and three entrances per target: the channel's start, its midpoint, and its start after a transverse kick), 23 fell within one iterate of the measured passages (Table 13). Seven of the 24 forecasts are five iterates or shorter, where a single iterate is 20 % or more of the passage; of the 17 longer forecasts, 16 fell within one iterate. The pre-specified criterion, that every row pass, failed by one row: for target A, midpoint entrance, $\varepsilon = 3 \times 10^{-6}$, the forecast was 61 and the measurement 59. All full-channel passages and all passages after a transverse kick were within one iterate. The misses are not symmetric: every one of the 17 nonzero residuals of a midpoint or kicked start, here and at the two folds tested before (Table 10 and Section 10), is negative, a consistent directional discrepancy across these related deterministic cases. They share coefficients and truncation errors and are not independent draws, so we attach no sign-test probability to them. The measured passages are consistently shorter than the iterated quadratic map, by one or two iterates. The design also froze a secondary comparison, the continuous passage integral $\int du/(a\varepsilon + bu^2)$ of the same normal form, whose closed form is the arctangent usually quoted for saddle-node bottlenecks. It came within one iterate of all 24 measurements (errors from -0.41 to $+0.92$) and gives 58.56 for the failing row. At that entrance the two forms already differed by 2.4 iterates before any measurement: iterating the quadratic map from the entrance u_{in} to the exit U adds, to leading order for slowly varying steps, about $\frac{1}{2}\ln[h(U)/h(u_{\text{in}})]$ steps to the integral, $h(u) = a\varepsilon + bu^2$, most at small ε and at the midpoint

entrance. Which of the two forms the full map follows is not settled by this test, and it does not let us promote the secondary comparison after the fact.

ε	target A ($G = F^2$): measured / predicted	target B ($G = F^8$): measured / predicted
3×10^{-6}	118/118, **59**/**61**, 117/118	26/26, 13/14, 25/26
1×10^{-5}	56/56, 28/29, 56/56	10/10, 5/5, 10/10
3×10^{-5}	26/26, 13/14, 26/26	4/4, 2/2, 4/4
1×10^{-4}	10/10, 5/6, 10/10	2/2, 1/1, 2/2

Table 13: Prospective passage test at two new folds, not used while the bottleneck explanation was worked out. Each cell lists the full-channel, midpoint and full-plus-transverse entrances as measured / predicted G-iterates. Predictions come from the iterated quadratic normal form, with coefficients from an independent implementation, frozen before the test runs. Bold: the one row outside the one-iterate criterion. The shortest period-8 rows give only weak relative accuracy.

A second prospective test. The first test left open which form of the normal form the full map follows, and its shortest forecasts carried little weight. A second test, frozen before any trajectory was run, made the continuous passage integral its primary predictor and the discrete count its secondary, and chose the distances so that every full-channel forecast is at least 16 iterates and every midpoint forecast at least 8. Its two folds are again new to the tests. Target C is the 2-cycle's fold at $\log_{10} T = -0.200$: $\alpha_c = 0.0713690$, $a = -10.9505$, $b = -25.3300$. Target D is the fold of a period-6 orbit, a different family, at $\log_{10} T = -0.087439$: $\alpha_c = 0.0756329$, $G = F^6$, $a = -72.695$, $b = 72.237$. Distances $\varepsilon = q^2/(ab)$ with $q = 0.08$, 0.05, 0.03 and 0.018 give both targets the same forecasts (Table 14). At the 2-cycle's fold the integral held in all 12 rows, and every measured count was the forecast rounded up. At the period-6 fold the midpoint passages held as well, but passages from the channel's entrance came in short, by 0.40 to 5.74 iterates and more so closer to the fold, so the primary criterion failed, 18 of 24 rows (the discrete count passed 16 of 24). From the fold point to the exit the period-6 passages follow the normal form; it is the first half of the channel that is crossed faster. Terms beyond the quadratic normal form on that side, or the curvature of the center manifold at the entrance, are candidate causes that this test does not resolve. Summed over both frozen tests, in which it was first the secondary and then the primary predictor, the continuous integral came within one iterate of 42 of 48 passages; all six misses are entrance passages at the period-6 fold, short by 3.1 to 5.2 %. The sign hypothesis frozen with the design, that midpoint and kicked passages do not exceed the discrete count, held in all 16 rows (`codex/confirm-fold-2`).

q	forecast (full / mid)	discrete (full / mid)	C, 2-cycle: measured	D, period 6: measured
0.08	22.40 / 11.20	23 / 12	23, 12, 23	22, 11, 22
0.05	44.29 / 22.14	45 / 23	45, 23, 45	**42**, 22, **42**
0.03	85.29 / 42.64	86 / 44	86, 43, 86	**81**, 43, **82**
0.018	154.74 / 77.37	155 / 80	155, 78, 155	**149**, 78, **150**

Table 14: Second prospective passage test, frozen before any passage was run. Forecast: the continuous passage integral $\int du/(a\varepsilon + bu^2)$ from the entrance to U (primary); discrete: the iterated quadratic map (secondary). Measured G-iterates for the full-channel, midpoint and full-plus-transverse entrances. Bold: outside one iterate of the primary forecast.

8.1 Global return, a natural exit curve, and eventual capture

The local fold calculation explains why a trajectory slows near the vanished 2-cycle. It does not specify where the trajectory goes after leaving that neighborhood. To separate these questions, an exploratory follow-up kept the model, prompt and philosophy direction fixed at $\log_{10} T = -0.115$ and $\alpha = \alpha_c + 10^{-5}$. All runs used the patched float64 map. We first changed the state at the exit, then followed repeated returns, and finally tested whether naturally reached exit states could be reconstructed from one coordinate. The completed studies and their separate-disk checks are bound by the exit-perturbation receipt and the repeated-return and reconstruction receipt.

The same slow coordinate need not give the same destination. Write the full state as $x = x_c + vr + w$, with $\ell_v^\top r = 1$, $\ell_v^\top w = 0$ and $v = \ell_v^\top (x - x_c)$. Here r and ℓ_v are the right and left critical directions oriented so that the passage moves toward increasing v; this is the oriented coordinate used above, with $U = 0.1/|b| = 0.0039325934$. The remainder w contains the transverse components. Of 128 controlled exit states, 123 returned within 13 to 68 individual updates of F; the other five approached the same period-3 orbit. A matched pair held v equal to numerical precision while changing w by about 0.286 % of the original full-state norm: one returned, the other approached period 3. Thus a rule based on v alone cannot describe arbitrary nearby states, even though the local passage itself is nearly one-dimensional.

The alternative destination is numerically an attracting *primitive* 3-cycle: its relative F^3 closure residual is below 8.6×10^{-15}, its lag-1 and lag-2 relative residuals exceed 1.75 and 2.91, and the spectral radius of the complete composed float64 derivative is 0.65065064. Independent directional derivatives and central differences check that composed derivative. These are numerical checks, not a new interval certificate. The leading-token readout is " certainly", ".", "Hello"; the full probability distributions remain soft.

A return does not settle the long-time fate. We next interpolated 64 initial states between the matched returning and period-3-bound exits, keeping the same v while varying w. Every history continued through subsequent returns until it matched the reference 3-cycle for 128 ordered full-state comparisons at relative tolerance 10^{-9}, or reached 4096 individual updates. A persistent channel visit requires three consecutive same-phase readings with $|v| < U$ and $\|w\| < 0.05$; these are local observation thresholds, not boundaries of chaos. Of the 64 starts, 63 returned at least once. Twenty-seven approached period 3, including 26 that returned first; their histories contained zero to 27 persistent visits. The other 37 made 29 to 34 visits and remained unresolved at the horizon (Fig. 25). The line is a constructed sample, not a basin-volume sample or a population escape-probability estimate. Repeated slow passages can precede periodic capture. They do not by themselves establish persistent chaos, or even a chaotic saddle; the distinction between an attracting chaotic set and a nonattracting chaotic set requires additional evidence [51].

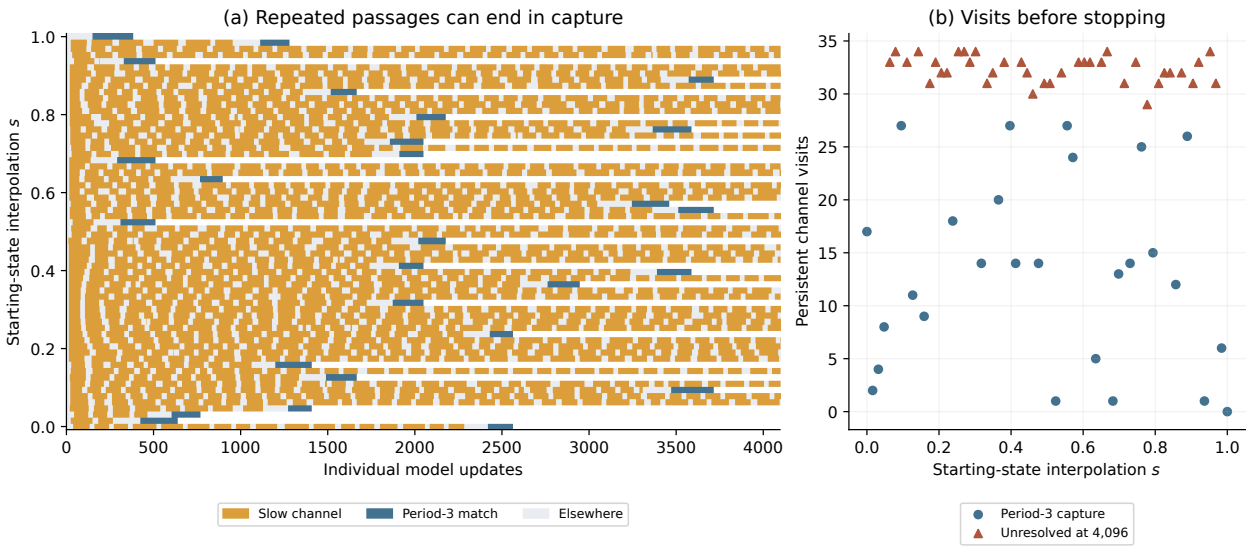

Figure 25: Repeated returns and later period-3 approach at one parameter setting. (a) All 64 starts, in their interpolation order s. Orange denotes the checked slow channel, pale gray the other recorded states, and blue the final reference-cycle matching interval. Both adjacent updates are colored for visibility when a same-phase channel reading passes. White after a blue interval is uncomputed time after the stopping check, not evidence of continued residence. (b) Counts of persistent channel visits before stopping: circles approach period 3; triangles are unresolved at 4096 updates. The 27 captured histories include 26 that first returned to the channel. No finite-horizon survivor is labeled permanently trapped or chaotic from this figure.

Natural departures admit a much stronger reduction. The perturbations above are artificial departures. States that actually emerge after a slow passage are much more restricted. Define a natural departure as the first same-phase reading with $v \geq 3U$ and $\|w\| < 0.05$ after at least 20 updates spent in the local passage. This departure surface lies beyond the narrower $|v| < U$ observation band; the two should not be confused. The histories supplied 1515 eligible exits. Complete initial-state trajectories, not individual visits, were divided into development and held-out sets, with endpoints retained as controls. The same-parameter neighboring initial states make this an exploratory interpolation test, not a test of transfer to new parameters or prompts.

From 143 development exits in the first 640 updates we fitted a full-state reconstruction curve

$$\widehat{x}(v) = x_c + vr + \sum_{j=0}^{7} c_j T_j\left(\frac{v - v_0}{s_v}\right), \qquad \ell_v^\top c_j = 0, \tag{8}$$

where T_j are Chebyshev polynomials and v_0, s_v center and scale the fitting interval. The vector coefficients reconstruct all 896 components. The degree was chosen by grouped development-only comparisons before the paired held-out excursions. On the first eligible exits of 31 held-out trajectories, the maximum relative state error was below 6.4×10^{-14}. This is evidence for an accurately fitted natural exit curve, not an exact invariant manifold or conjugacy.

We then advanced each original state and its reconstruction through the *unchanged full network*. All 31 pairs had the same next destination: 28 returned and three approached period 3. Every returning pair had the same integer re-entry time, with entry-coordinate differences below $4.6 \times 10^{-5} U$. Continuing those pairs through the next slow passage gave identical departure-to-departure times, spanning 105 to 227 individual updates, and next-departure coordinate differences below $0.0014U$ (Fig. 26). All original full-cycle controls also reproduced their event times in the stored parent

trajectories. The success tests loss of transverse information on the natural exit curve; it does not test a fitted scalar simulator.

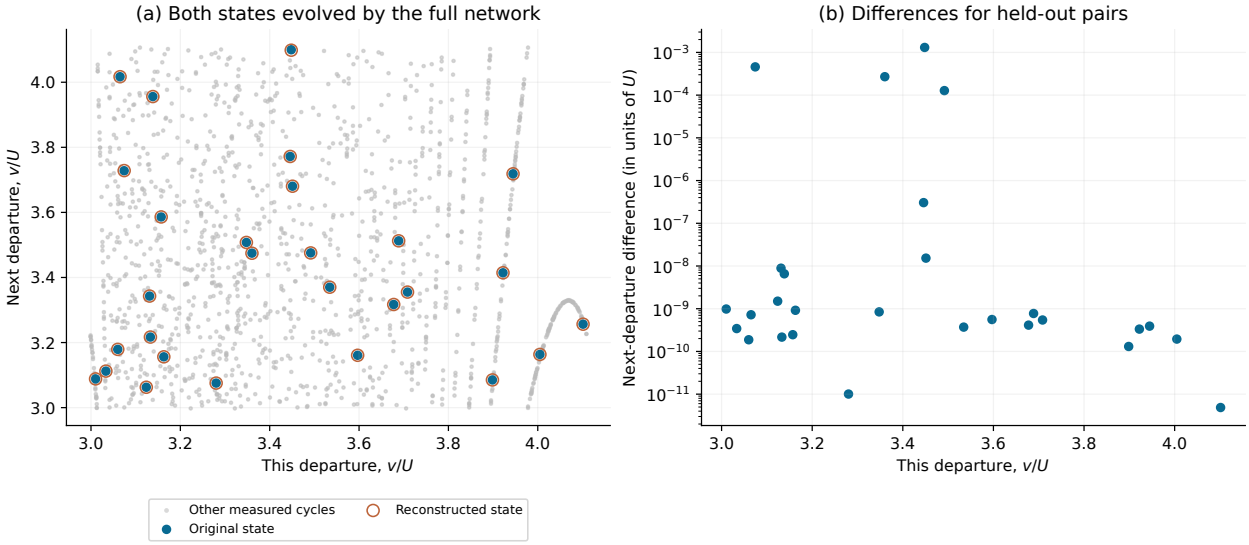

Figure 26: One-coordinate state reconstruction followed by full-network evolution. (a) The coordinate of the next natural departure against the current one. Filled circles are the original held-out states; open circles start from their scalar reconstructions. Gray points give the other measured adjacent natural return cycles for context. (b) Absolute next-departure differences, on a logarithmic scale. All 28 returning pairs have identical complete-cycle times and differences below $0.0014U$. The three period-3 pairs have no next channel departure and are checked separately, not dropped from the destination comparison. Both members of every pair use the full network.

What remains unsolved. The construction defines a numerical return evaluation

$$\widehat{P}(v) = \ell_v^\top \left[F^{\tau(v)}(\widehat{x}(v)) - x_c \right], \tag{9}$$

where $\tau(v)$ is the first next qualified natural-departure time, when one exists. A separate label records period-3 approach or unresolved evolution. Equation (9) still calls the full network. Local fold coefficients cannot supply the missing global function: two maps can agree near the fold and differ outside it. Local transverse contraction is consistent with the observed exit curve, but does not make the global return a low-degree polynomial or guarantee accuracy after repeated composition.

Simpler fitted predictors did considerably worse. On 777 returned test exits from wholly withheld trajectories, the entry-coordinate root-mean-square error, with initial trajectories weighted equally, was $0.2645U$ for a constant baseline, $0.2142U$ for scalar linear interpolation and $0.1913U$ for an eight-neighbor predictor using the slow coordinate and three transverse principal components. These predict the *entry*, not the next departure in Fig. 26. A nearest-neighbor destination classifier found six of nine test period-3 events but falsely called 14 returns period-3 events; its 97.84 % accuracy was below always predicting return (98.85 %). One censored test exit was excluded. These failures remain part of the result.

An autonomous approximation would require a separately validated rule $v_{k+1} = P(v_k)$ together with a return-time function and an escape label, or additional coordinates if a single-valued scalar rule fails. A useful next test would resolve the steep return branches adaptively, freeze the resulting fit, and test both single-cycle errors and multi-cycle survival and timing on fresh trajectories and neighboring parameter settings. Fine-scale structure should be measured, not smoothed

away. Persistent chaos is a different question: longer unbiased ensembles and tangent growth can distinguish candidate transients from long-lived expansion numerically, while a rigorous claim would need certified invariant-set and expansion or symbolic-dynamics arguments for the full map. A chaotic saddle can produce transient chaos even when typical trajectories eventually settle into a periodic orbit [51]. Neither a validated scalar surrogate nor such an invariant-set certificate is supplied by the present return experiment.

9 A period-doubling cascade to chaos

The crescent of Figure 11 is a closed region in which the period-3 loop "\n\nThe existence" doubles inward. Along the cut $\log_{10} T = -0.14798$, the map's cells show period 3 over 25 columns, period 6 over about 14, period 12 over 4 and period 24 in one column, followed by a chaotic core. At $\alpha = 0.0650$, inside the period-3 band, a computer-assisted proof shows that the 3-cycle exists and attracts (Appendix L). The doubled loops repeat the same leading tokens twice and four times, with alternating probabilities. Locating the flips along this cut in float64 (Appendix B), in the steering coordinate α, gives Table 16. The first two ratios, $\delta_2 = 2.844$ and $\delta_3 = 2.859$ (indexed as in Table 16), are far from Feigenbaum's 4.6692, even though from the second flip on the orbits are dominated in their linear relaxation by one direction (next modulus 0.0023, 2×10^{-6}, 2×10^{-12}). The next ratio, $\delta_4 = 3.781$, is markedly larger. The first four flip parameters agree between completed float64 computations on a CPU and an H100 to 4×10^{-15}.

Continuing a just-born doubled orbit past period 48 failed: Newton's method slides onto the half-period orbit, which is also periodic with the doubled period. We therefore also located the superstable parameters s_p, where the critical multiplier vanishes in the middle of each band (Table 15). There the orbit attracts, so it can be reached by forward iteration and cannot slide onto its half-period parent. Through period 768 the superstable ratios are 3.19, 4.02, 4.478, 4.625 and 4.660. The last three come from parameters resolved to 10^{-14} and differ from Feigenbaum's value by 0.19, 0.044 and 0.0097; from 4.02 on, the distance shrinks by factors of 3.4, 4.3 and 4.6. The flip ratios 2.84, 2.86, 3.78, 4.36, 4.60 and 4.65 rise in step (Figure 27); the last three come from runs started inside the period-96 band, which located the flips of periods 96, 192 and 384 directly. The accumulation point lies near $\alpha_\infty \approx 0.0661250$. The early superstable values are less precise. The other computed Ritz values of the period-12 and period-24 orbits blur the root, so s_{12} and s_{24} are known only within brackets of 1.24×10^{-5} and 6.9×10^{-7}, and the ratios 3.19 and 4.019 have worst-case ranges $[2.895, 3.338]$ and $[4.019, 4.101]$. s_6 is less precise still, and s_3 was not bracketed.

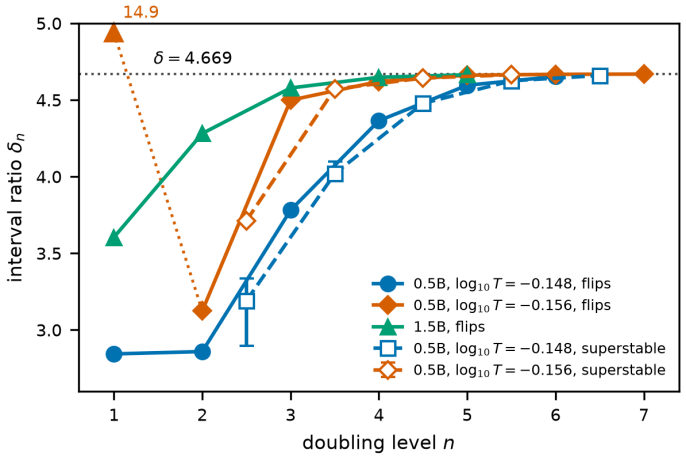

Figure 27: Finite-order interval ratios of the doubling sequences: the crescent of Qwen2.5-0.5B-Instruct (philosophy direction, layer 12), from flips and from superstable points (which reach period 768), and the route of Qwen2.5-1.5B (philosophy direction, layer 14; Figure 82), from flips. A superstable ratio spans half-bands, so it is plotted half a level after the flip ratio of its first band; level n carries the ratio of the nth to the $(n+1)$th band, δ_{n+1} in the notation of the text and Table 16. Bars: worst-case ranges from the brackets of s_{12} and s_{24}. At $\log_{10} T = -0.148$, and in the 1.5B model, the ratios rise toward Feigenbaum's δ from below; the last superstable ratio at -0.148 is within 0.21 % of it, and the last 1.5B ratio (4.665) within 0.1 %. At $\log_{10} T = -0.156$ the first ratio, 14.9, lies off the scale (arrow), and the later ratios approach δ from below, the last (4.6688) within 0.008 % of it, and within 0.015 % over the location uncertainty of the period-384 flip. Its first two superstable ratios carry no bars: the brackets of s_{12} and s_{24} allow $[3.40, 29.3]$ and $[1.01, 4.65]$. Each sequence ends at the deepest level computed: the flips at -0.148 at period 384, the superstable points at period 768 at -0.148 and at period 384 at -0.156, the 1.5B route at period 256, and the cut at -0.156, which we followed furthest, at the $768 \to 1536$ flip.

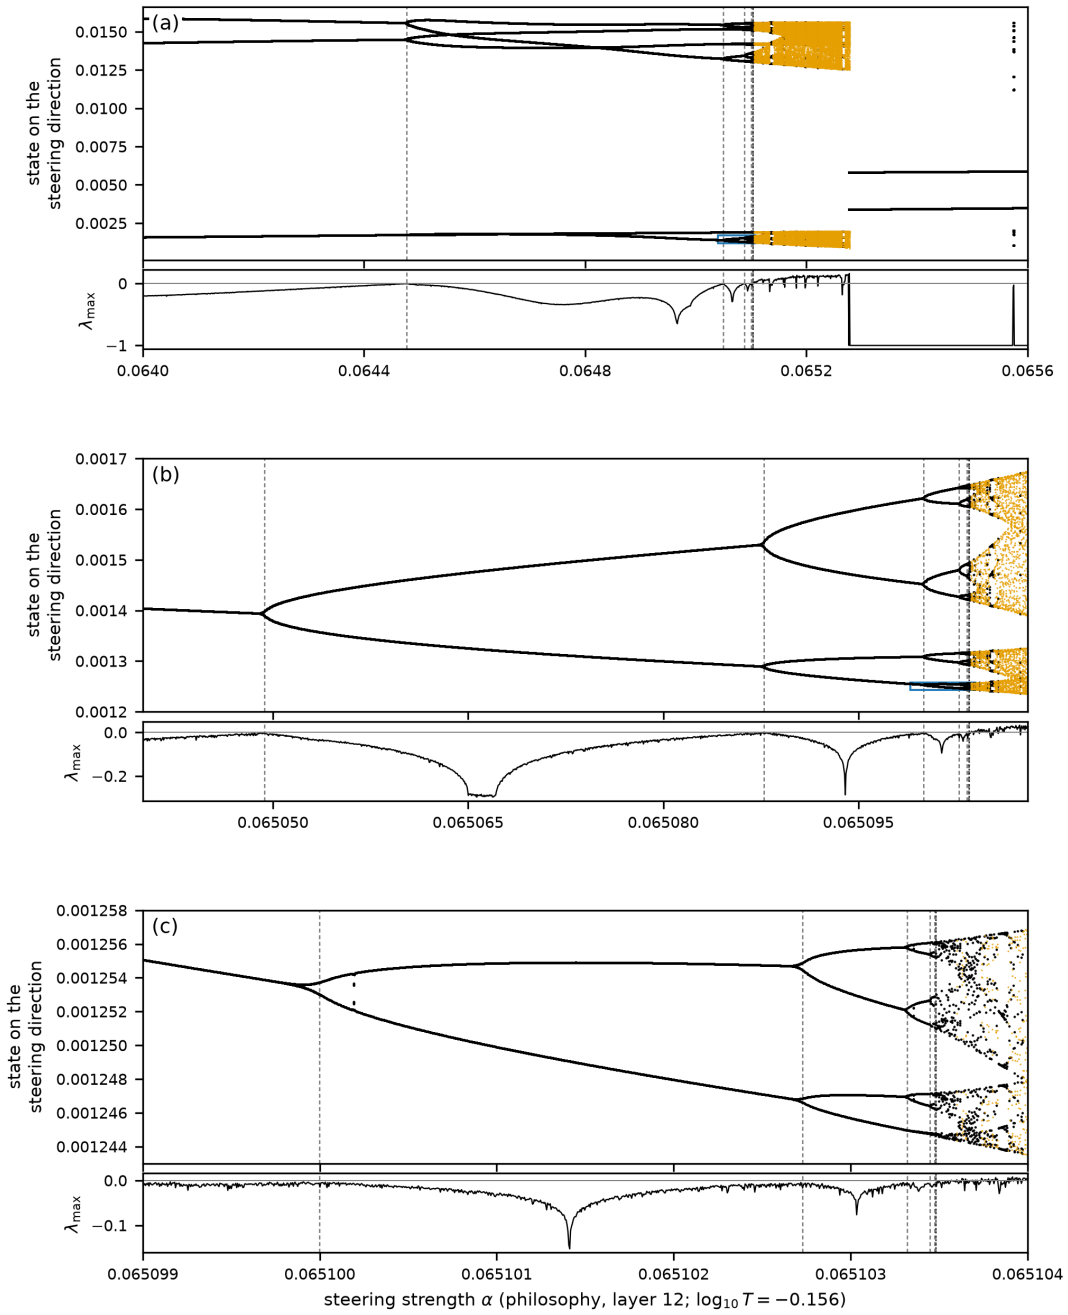

Figure 28: The Feigenbaum tree of the crescent cascade (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, $\log_{10} T = -0.156$): float64 line scans of 1600, 1360 and 1000 cells, each with 3000 transient steps and 256 recorded states projected on the unit steering direction (orange: chaotic cells, with no period up to 48). (a) The whole cut: period 3 doubles into chaos, which ends at $\alpha = 0.065278$, where the attractor jumps to a strongly contracting period-2 orbit. (b, c) Zooms in both coordinates on the branch boxed in the panel above, showing the doublings $6 \to 12 \to 24 \to 48$ and $24 \to 48 \to 96 \to 192$. Dashed lines: flips located by Newton's method (Table 17); dotted line: the accumulation point estimated from the last two flips with δ. Under each panel, the largest Lyapunov exponent touches zero at every flip and turns positive beyond the accumulation point. Near a flip the orbit converges slowly, so with 3000 transient steps each fork opens slightly before its Newton flip, by 8.4×10^{-6} in α at the first and by at most 1.1×10^{-6} at the later ones.

period p	superstable parameter s_p	bracket width	ratio $\frac{s_p - s_{p/2}}{s_{2p} - s_p}$
12	0.0659717	1.24×10^{-5}	
24	0.06608014	6.9×10^{-7}	3.19
48	0.066114132784466	3.3×10^{-15}	4.019
96	0.066122589892689	7×10^{-15}	4.4779
192	0.066124478510077	6.4×10^{-15}	4.6251
384	0.066124886853811	5×10^{-15}	4.6595
768	0.066124974490236	4.1×10^{-15}	

Table 15: Superstable parameters of the crescent's doubling sequence along $\log_{10} T = -0.14798$: the critical multiplier of the period-p orbit changes sign across the bracket, with Newton residuals below 3×10^{-13} and the other computed Ritz moduli below 10^{-20} from period 96 to 384; at period 768 the residual is 1.2×10^{-12} and the other moduli are below 3×10^{-20}. Brackets for periods 12 and 24 are the closest points where the critical multiplier is cleanly identified.

Why are the early ratios so small? We first considered a crossover. Softmax at temperature T is a smoothed argmax, so the reduced one-dimensional map might have a nearly corner-like extremum, rounded only on a scale set by T. Early doublings would then see the corner, where ratios tend toward 2, and later ones the smooth quadratic core. This predicts smaller early ratios at lower temperature. We fixed that prediction before locating the flips along two lower cuts, and it fails. At $\log_{10} T = -0.156$ the first ratio is 14.93, five times its value at -0.148 (Table 17). At -0.152 the continued period-12 solution collapsed onto its period-6 parent near $\alpha \approx 0.0657$, and the forward search found no next flip. This is compatible with a doubling bubble, but we did not locate its reverse boundary or exclude a narrower cascade on that cut. The early ratios thus depend strongly on where a cut crosses the crescent, and we withdraw the crossover explanation. The later ratios at -0.156 behave as universality requires. The next ratios are 3.124, 4.4988, 4.6208, 4.6618, 4.6673 and 4.6688. Their deviations from $\delta = 4.6692$ are -1.55, -0.17, -0.048, -0.0074, -0.0019 and -0.00036; they keep one sign, from the third on each step closes the gap by a factor of 3.5 to 6.5 (the location uncertainty of the period-384 flip, given below, leaves the last factor undetermined), and Aitken's Δ^2 extrapolation of the last three ratios gives 4.6694. Universality of this kind holds within a class: for maps with a quadratic critical point, including suitable strongly dissipative families of any dimension [60, 61]; the constant belongs to the class, not to period doubling as such, and differs for a critical point of higher order. Such cascades also occur in the training map of gradient descent in solvable models of deep networks [25]. The last ratio uses the flip of the period-768 orbit (the $768 \to 1536$ flip), located with $|\mu + 1| = 4.1 \times 10^{-7}$. The penultimate flip, of the period-384 orbit, approaches the precision floor of its multiplier, a product of 384 step Jacobians. On an H100 the computed $\mu + 1$ scatters by 1.8×10^{-5} about a straight line near the root; under that local linear fit, the largest residual divided by the slope gives a location uncertainty of about $\pm 3.3 \times 10^{-13}$, an estimate rather than a certified bound. The workstation GPU located it at $\alpha = 0.06510347354866666$ with $|\mu + 1| = 4 \times 10^{-9}$, 9×10^{-14} from the H100 estimate. Over that range the ratio δ_7 lies between 4.6672 and 4.6674, δ_8, the last ratio, between 4.6685 and 4.6692, and the extrapolated value between 4.6688 and 4.6702, an interval that contains δ. The superstable points along the same cut, through period 384, give the ratios 10.21, 3.713, 4.572, 4.643 and 4.6646. Before computing them we fixed the prediction that the last two superstable ratios at a new temperature would lie within 5 % of δ, in $[4.436, 4.902]$. At the original readout, through period 192, the point estimates 4.572 and 4.643 met it, but the bracket envelope of 4.572 did not; extending the computation to period 384 later gave the pair 4.643 and 4.6646, both inside the interval. The last, 4.6646, rests on parameters resolved

to 5×10^{-15} and differs from δ by 0.1 %. Over the corners of the brackets the one before lies in $[4.639, 4.647]$, limited by s_{48}, whose bracket is 2.4×10^{-9} wide. The earlier ratio 4.572 depends on s_{24}, known only within a bracket of 5.9×10^{-6} (worst-case range $[1.01, 4.65]$). Finite-order ratios also depend on the parameter coordinate, so only the trend of later ratios can test universality [24, 62, 63]. Three claims should be kept apart here. The map has a period-doubling cascade; its late finite-order ratios approach δ numerically, with the digits limited by the brackets above, since the last flips lie near the float64 floor of the multiplier computation; and Feigenbaum's theory explains such convergence for an effectively one-dimensional, unimodal return map, a hypothesis these ratios alone do not establish for this map.

n	orbit	flip parameter α_n	next modulus	ratio $\delta_n = \frac{\alpha_n - \alpha_{n-1}}{\alpha_{n+1} - \alpha_n}$
1	$3 \to 6$	0.0654807769499	0.088	
2	$6 \to 12$	0.0659055406325	0.0023	2.844
3	$12 \to 24$	0.0660549041785	2.4×10^{-6}	2.859
4	$24 \to 48$	0.0661071412523	2.3×10^{-12}	3.781
5	$48 \to 96$	0.0661209563174	$< 10^{-13}$	4.365
6	$96 \to 192$	0.0661241213732	$< 10^{-13}$	4.596
7	$192 \to 384$	0.0661248100141	$< 10^{-13}$	4.653
8	$384 \to 768$	0.0661249580173	$< 10^{-13}$	

Table 16: Flips of the crescent's period-doubling sequence along $\log_{10} T = -0.14798$. For the first five flips (H100) the critical Ritz value is real and differs from -1 by at most 4×10^{-10}, the Newton residuals are below 4×10^{-13}, and the other computed Ritz values lie inside the unit circle. The sixth flip was located at the precision floor of the period-96 multiplier: the critical value differs from -1 by 5.7×10^{-7}, its scatter between bracket points 10^{-13} apart is about 10^{-6}, and the Newton residual is 5.5×10^{-12}; the location is uncertain by about 2×10^{-12}, which changes δ_5 by less than 10^{-5}. The seventh flip is clean (critical value -1.000000016, Newton residual 2.8×10^{-12}). The eighth was accepted at the precision floor of the period-384 multiplier: there the critical value differs from -1 by 5.3×10^{-5} (acceptance threshold 10^{-4}), bracket evaluations within 4×10^{-12} of it scatter by about 10^{-4} in that difference, and the Newton residuals are below 4×10^{-11}. Its location is uncertain by about 4×10^{-12}, which changes δ_7 by about 10^{-4}. Feigenbaum's constant is $\delta = 4.6692\ldots$ [24]; for the logistic map the first ratios are 4.751, 4.656, 4.668. The first four parameters agree between the completed CPU and H100 computations to 4×10^{-15}. Superstable parameters (Table 15) extend the ratio sequence through period 768.

n	orbit	flip parameter α_n	next modulus	ratio δ_n
1	$3 \to 6$	0.0644766567385	0.12	
2	$6 \to 12$	0.0650493555308	0.045	14.93
3	$12 \to 24$	0.0650877194158	7.0×10^{-5}	3.124
4	$24 \to 48$	0.0650999993318	4.3×10^{-7}	4.4988
5	$48 \to 96$	0.0651027289524	1.8×10^{-13}	4.6208
6	$96 \to 192$	0.0651033196816	2.9×10^{-14}	4.6618
7	$192 \to 384$	0.0651034463988	1.5×10^{-14}	4.6673
8	$384 \to 768$	0.0651034735487	3.2×10^{-13}	4.6688
9	$768 \to 1536$	0.0651034793638	2.9×10^{-14}	

Table 17: Flips of the crescent's doubling sequence along $\log_{10} T = -0.156$, in float64. Flips 1–3 were computed on an H100 and flips 3–8 on a workstation GPU (NVIDIA RTX 2000 Ada); the two computations of flip 3 agree to 4×10^{-16}, and an H100 computation of flip 8 from its evaluations at the multiplier's precision floor lies 9×10^{-14} away. Flip 9 was computed on an H100. Every critical multiplier is real; for flips 1–8 it differs from -1 by at most 5×10^{-8}, and for flip 9 by 4.1×10^{-7}. Next moduli below about 10^{-12} lie at the rounding level of products of hundreds of 896×896 Jacobians: they show that the noncritical directions are negligible, not their decimal values.

Not every doubling leads to chaos. In the high-temperature region (Figure 68, Appendix E), along $\log_{10} T = 0.0319$, the period-2 loop "\n\nThe" flips at $\alpha = 0.0753534$ and its doubled loop at 0.0736921. These orbits are far from one-dimensional (next moduli 0.38 and 0.28). Toward smaller α the period-8 orbit does not double again: near $\alpha \approx 0.0705$ it merges back into the period-4 orbit. Along this cut it is therefore a period-doubling bubble, bounded by two flips of the period-4 orbit, not the start of a cascade to chaos.

	offset	β	amplitude ratio	doubled multiplier	$1 - 4\beta$
period $2 \to 4$	2.5×10^{-5}	0.0092	1.005	0.96335	0.96320
($c = 6.49 \times 10^3$)	1×10^{-4}	0.0364	1.020	0.8572	0.8545
	4×10^{-4}	0.1400	1.078	0.513	0.440
period $4 \to 8$	2.5×10^{-5}	0.0175	1.003	0.92987	0.92983
($c = 1.24 \times 10^4$)	1×10^{-4}	0.0685	1.012	0.7257	0.7260
	4×10^{-4}	0.2470	1.056	complex	0.012

Table 18: The flip normal form where the map is far from one-dimensional: the two flips of the high-temperature bubble (next moduli 0.38 and 0.28). At offsets below each flip in α, $\beta = -\mu_p - 1$ from the period-p orbit's multiplier; amplitude ratio: the doubled orbit's half-amplitude along the critical direction divided by the prediction $\sqrt{\beta/c}$; doubled multiplier: computed against the prediction $1 - 4\beta$. Both coefficients are positive (supercritical), as the stable doubled orbits require. Three of five offsets per flip are shown.

The local theory still applies here (Section 3). At each flip we computed the cubic coefficient of the flip normal form, $\eta \mapsto -(1+\beta)\eta + c\,\eta^3$ on the critical direction, from the dense monodromy and float64 directional derivatives (notation of Appendix B), $c = \ell^\top C(r,r,r)/6 + \ell^\top B(r, h_2)/2$ with $(I - M)h_2 = B(r,r)$ [50]. To leading order it predicts the doubled orbit without fitting: half-amplitude $\sqrt{\beta/c}\,(1 + O(\beta))$ along r and multiplier $1 - 4\beta + O(\beta^2)$. At the smallest offset the amplitude ratios are 1.005 and 1.003, and the multipliers agree within 1.5×10^{-4} and 5×10^{-5}; the larger discrepancies at larger offsets are consistent with higher-order corrections (Table 18).

These finite comparisons support the leading-order prediction; they do not establish the order of the remainder. Figure 28 shows the cascade itself as a bifurcation diagram: in three nested zooms each fork sits at a flip located independently by Newton's method, and the largest exponent touches zero at every flip, as it must where a multiplier reaches modulus one, and turns positive beyond the accumulation point. The route to chaos of Qwen2.5-1.5B (Figure 82) behaves the same way: its last ratio, 4.665, is within 0.1 % of δ (Figure 27).

10 Steering toward hostility

A second steering direction tests whether the organization of the philosophy-steered map is specific to one concept. We steer along a valence direction, the contrast of hostile statements against matched kind ones, built in the manner of contrastive activation addition [35]; its cosine with the gratitude direction is -0.04. Sampled self-dialogue under the same steering is a different system, with a growing context and sampling noise; we report that study separately and follow the direction here in the deterministic map.

The hostility direction in the deterministic map. Steering the soft-token map of Section 2 along the same hostility direction gives a plane with chaos (a positive largest exponent in 10 % of a coarse 128×192 float32 scan) and with periodic text loops that are hostile: "Yes, that's crap" (period 6), "a lie is" (3), "STUP…" (3), a loop on the Chinese word for insult and its 2-cycle, and others. They appear at α above about 0.035 (the insult loop) and above 0.065 (the rest). Around the period-6 loop the windows form a dome (Figure 29), and toward higher temperature its texts gain one "!Yes" after another (periods $6, 8, \ldots, 20$). In float64 the windows of periods 12 to 20 are born in folds on their low-temperature side and end in flips or $+1$ crossings on the other, like the stripes of Sections 7 and 8. The period-6 window ends in a fold at $\log_{10} T = -0.0991$, and the continuations of the period-8 and period-10 windows failed to converge.

The windows accumulate on the fold of the 2-cycle "!Yes" at $\log_{10} T_c = -0.0462465357$ ($\alpha = 0.083$; critical multiplier 0.9999999954, next modulus 0.050). From period 16 on, $1/\sqrt{\log_{10} T_c - \log_{10} T_{\text{fold}}}$ rises by 2.235 per lap. The fold's normal form ($a = -1.1683$, $b = 44.525$ with ε in $\log_{10} T$) predicts $\sqrt{|ab|}/\pi = 2.296$ per lap for full passages through the bottleneck, within 3 % and without a fitted parameter. An independent float64 forward and automatic-differentiation implementation reproduces this fold (location to 6×10^{-10} relative, a simple real $+1$ multiplier with next modulus 0.050) and its coefficients (ab to 8 parts per million). In that implementation, controlled passages from fixed entrances at $\log_{10} T$ offsets -10^{-5}, -3×10^{-5} and -10^{-4} took 119/59, 61/31 and 27/13 iterates of F^2 (full channel/from its midpoint), against prospectively frozen quadratic forecasts of 118/61, 61/32 and 27/14. Five of the six are within one iterate; the smallest-offset midpoint missed by two, failing the frozen all-six criterion. Finite-endpoint (arctan) forecasts, also frozen, were within 0.923 iterates. These are controlled local transits. At the window distances the finite-channel correction to the asymptotic passage time is large, so the agreement of the window slope with the full-passage coefficient does not by itself show where the natural dynamics reinjects. The same type-I intermittency mechanism as in the philosophy-steered map is thus consistent with the hostile windows' organization, which the local fold normal form supports; the controlled passages do not identify the natural reinjection or the global organization.

A second region, near $(\alpha, \log_{10} T) = (0.06, 0.20)$, shows arms of period-7 and period-8 windows spiralling around the insult loop. In float64 the insult fixed point doubles into a 2-cycle at $\alpha \approx 0.041$ and 0.079, and that 2-cycle has a Neimark–Sacker bifurcation in the spiral region: at $\alpha = 0.062$ its complex pair crosses the unit circle near $\log_{10} T = 0.1656$, turning 0.12 of a revolution per

step of the 2-cycle. This does not explain the arms. The two invariant circles born at a 2-cycle's Neimark–Sacker point are exchanged by the map, so orbits locked on them have even periods, near 16 for this rotation, whereas the arms include period 7; and of four float32 arm cells, only one kept its period in float64. In dissipative flows, windows wind in nested spirals around periodicity hubs [64], and in the Rössler system the hub's shape was traced to a conjugacy to unimodal maps in part of its parameter plane [65]; we have not tested whether the arms here wind around such a hub. Float64 zooms show more structure. In one arm a period-7 domain carries period-14 crescents with period-28 cores along its lower edge, between period-1 and period-2 regions and a chaotic band. In the cascade the windows are nested crescent arcs inside a dome bounded by the "!Yes" 2-cycle region, and near $(\alpha, \log_{10} T) = (0.087, -0.069)$ the period-12 and period-14 windows bend into hooked, nested arches. These are the map-level features that accompany shrimp (a periodic body with doubling crescents; hooked windows).

In float64 the period-7 body is bounded by folds. At $\log_{10} T = 0.1925$ the orbit exists between $\alpha = 0.065966$ and 0.071188, and at $\alpha = 0.066$ it ends in a fold at $\log_{10} T = 0.19263$. Below that, the multiplier passes through -1 twice (bracketed in $[0.19176, 0.19196]$ and $[0.18931, 0.19036]$), so the period-14 crescent is a doubling bubble. Further down, the branch folds twice (0.18797 and 0.18922, the second located by continuing its fold curve; three coexisting period-7 orbits in between) before a third flip. All crossings are simple (next modulus below 1.5×10^{-3}). An S-shaped pair of folds is what a cusp of the fold curves leaves on a nearby cut. When these boundaries are continued in the plane in float64 (Figure 32), each of the two fold curves reverses its parameter tangent once. There, as the fold condition requires, the oriented quadratic coefficient b of the fold's normal form passes through zero; what makes such a point a cusp is that the cubic coefficient c and the unfolding determinant $\det K$ are nonzero there (the cusp test, Appendix E) (upper fold: -23.9 to $+244.6$ between adjacent points, with $a \approx 2.0$; lower fold: $+31.2$ to -58.7, with $a \approx 0.55$) while a stays away from zero. Located along the curves, the two cusps lie at $(\alpha, \log_{10} T) = (0.067213, 0.198983)$ and $(0.066611, 0.189382)$. There $b \operatorname{sign} a$ is below 1.3×10^{-5} of its bracket values, $c = -1.41 \times 10^6$ and -5.61×10^5, and $\det K = 58.2$ and 55.2, all evaluated at the located points. The flip curve of the period-14 bubble lies inside the body and closes around the map's period-14 cells. The cubic coefficient is nonzero at the points that bracket both sign changes: $c \approx -1.4 \times 10^6$ to -2.0×10^6 (step agreement 5 %; the smallest steps converge like the square of the step, and extrapolation leaves the sign) and -5.8×10^5 to -5.3×10^5. This supports nondegenerate cusps inside both brackets, and $c < 0$ agrees with the three coexisting period-7 orbits, two of them stable, found between the S-shaped folds. At a fold point near the lower cusp (where $b = 31.4$) the estimated unfolding is nonsingular too ($\det K = 53.5$ in the same coordinates, singular values 262 and 0.20). Its local predictions at six frozen offsets matched the orbits that Newton solves seeded at the predicted roots and the base state recovered: in number at all six and in stability at five. The miss, at the largest offset beyond the edge of the region with three orbits, marks the range of the cubic model. The hostile period-7 body thus has the same skeleton as the period-6 window of Appendix E.

Its interior is shrimp-like too. At $\log_{10} T = 0.1830$, float64 bifurcation diagrams through the body (Figure 30; 1400 cells per line) show the period-7 orbit doubling to periods 14 and 28 and entering chaos (largest exponent up to +0.1) for $\alpha \approx 0.0659$–0.0673. Windows of periods 28, 35 and 42 sit inside the chaotic band, and the orbit then undoubles through 28 and 14 back to 7; the exponent touches zero at every doubling. Newton's method converges at map cells to primitive orbits of periods 14, 28 and 21 (residuals 3×10^{-15}, 2×10^{-14} and 5×10^{-13}; spectral radii 0.08, 4×10^{-4} and 0.006). Periods 21, 35 and 42 are three, five and six times the body's period, as expected for the windows that renormalization places inside the band chaos of a period-7 cascade. Continued in temperature, the period-14 orbit flips into period 28 near $\log_{10} T = 0.1841$ and back near 0.1857, the two ends of a doubling bubble that holds period 28 and chaos. Lower down, its

branch is S-shaped (folds near 0.1824 and 0.18244, with three coexisting period-14 orbits between them) before another flip near 0.1819. All these crossings are simple (next modulus below 4×10^{-5}). A 160^2 float64 zoom of the lower body (Figure 78) shows this anatomy in the plane. The period-14 domain encloses a period-28 ring around a chaotic core. Period-21 windows sit in the chaotic region as thin crescents with legs, and curves of very negative exponent, where the leading multiplier is small (candidate superstable curves), run through the windows, as superstable lines run through the windows of low-dimensional maps. In this zoom 18 % of cells are chaotic.

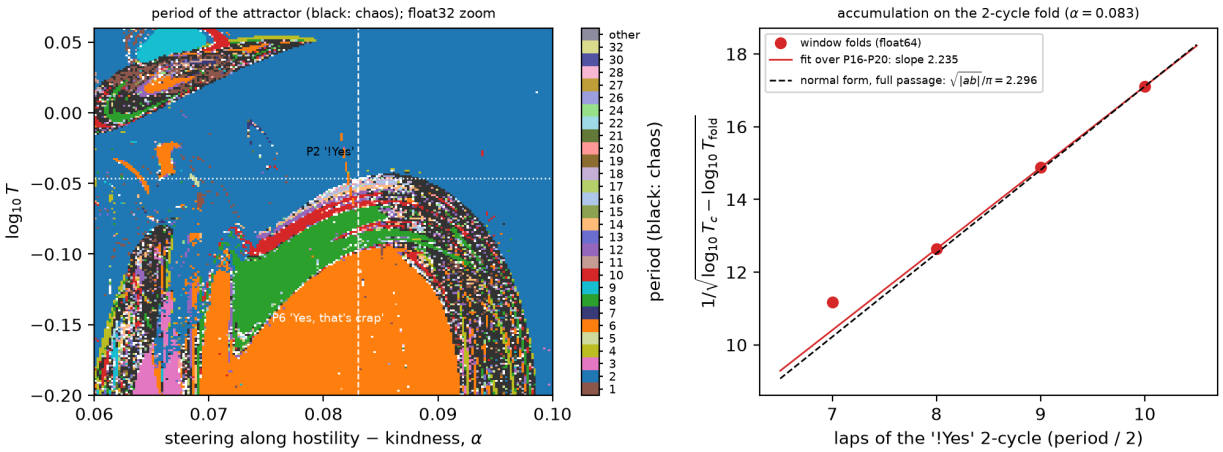

Figure 29: The hostility-steered soft-token map of Qwen2.5-0.5B-Instruct (layer 12). Left: period of the attractor in a float32 zoom (black: chaos), with the period-6 "Yes, that's crap" dome and its period-adding bands; dashed, the cut at $\alpha = 0.083$; dotted, the float64 fold of the "!Yes" 2-cycle. Right: float64 fold edges of the windows (periods 14–20) against the number of "!Yes" laps; the dashed line is the normal-form prediction for full passages, anchored at the last window.

No one-dimensional shrimp center. In a shrimp of a one-dimensional map, two superstable curves of the window's orbit cross at a doubly superstable center, the head [66, 67]. We continued one period-7 orbit over a 31×31 grid of the body's upper part and followed the sign of its leading real multiplier (Figure 31). The curves where that sign changes between cells are candidate superstable curves: the modulus is small along them, but a sign change between cells can also come from two real multipliers of opposite sign exchanging dominance. They form two arcs that approach each other near $(\alpha, \log_{10} T) = (0.0700, 0.206)$ and turn away without crossing at this resolution (10^{-4} in α, 4.3×10^{-4} in $\log_{10} T$). A finer grid of that approach (1155 cells, about 2 times finer in α and 6 times finer in $\log_{10} T$) shows why. There the leading multiplier is tiny (median modulus 9×10^{-4}) and as small as the next one, and in 221 cells the two form a complex pair, in two wedges that flank the gap. On this sampled approach the two computed leading multipliers exchange dominance and form a complex pair, so the sign of a leading real multiplier no longer defines a scalar superstable curve there. Where a one-dimensional map would place its doubly superstable center, the orbit's local linearization has two comparable directions; whether the dynamics there reduces to an effectively two-dimensional map is not tested. Next to the body, the 7-step return maps of chaotic cells have a coarse one-extremum shape in every phase: one interior extremum when the points are binned into 16 intervals, while finer partitions keep extra extrema, especially in phases 4 and 5, which we do not attribute to noise. In two phases the points split into two thin strands, an observation of the projection that does not establish a second invariant dimension (Figure 75). Cobwebs on these maps

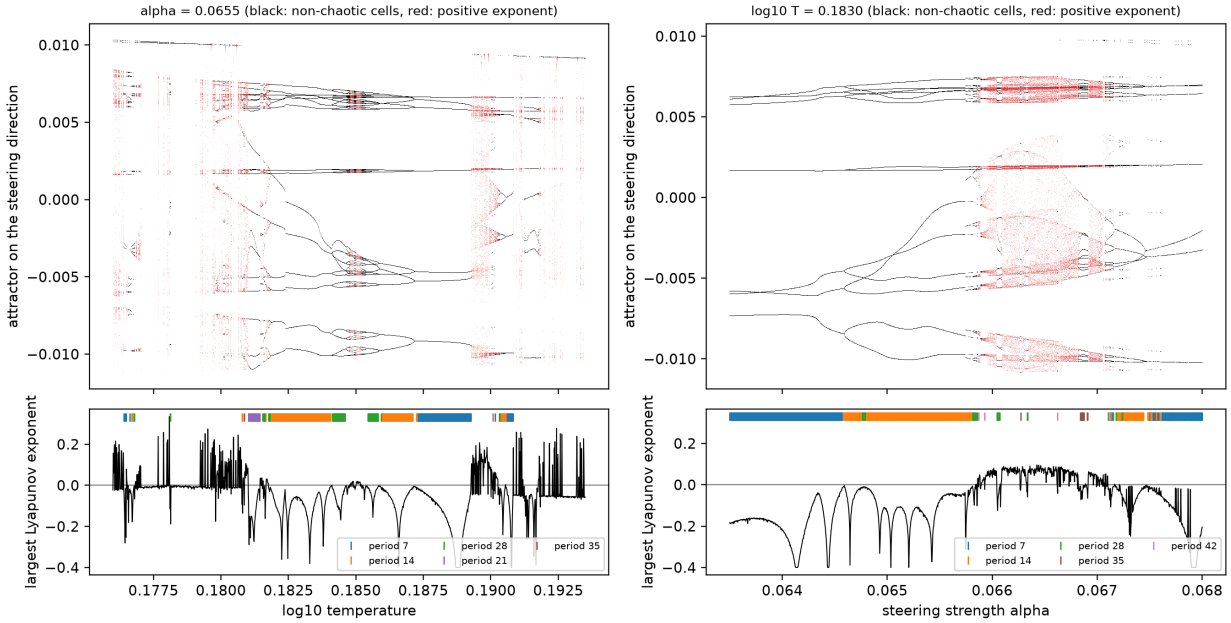

Figure 30: Inside the hostile period-7 body (Qwen2.5-0.5B-Instruct, hostility direction at layer 12): float64 bifurcation diagrams along temperature at $\alpha = 0.0655$ (left) and along the steering strength at $\log_{10} T = 0.1830$ (right). Top: the last 160 states of the attractor reached from the fixed start, projected on the steering direction (red: cells with a positive largest exponent). Bottom: largest Lyapunov exponent, with the cells of periods 7, 14, 21, 28, 35 and 42 marked (cells of other periods are not marked). On the right, period 7 doubles to 14 and 28, enters chaos with windows of periods 28, 35 and 42, and undoubles back to 7. Labels come from 400 measured steps after 1000 transient steps; with 2000 measured steps at two twin offsets, 135 of the 235 red cells on the left stay chaotic and 91 are periodic (Appendix K).

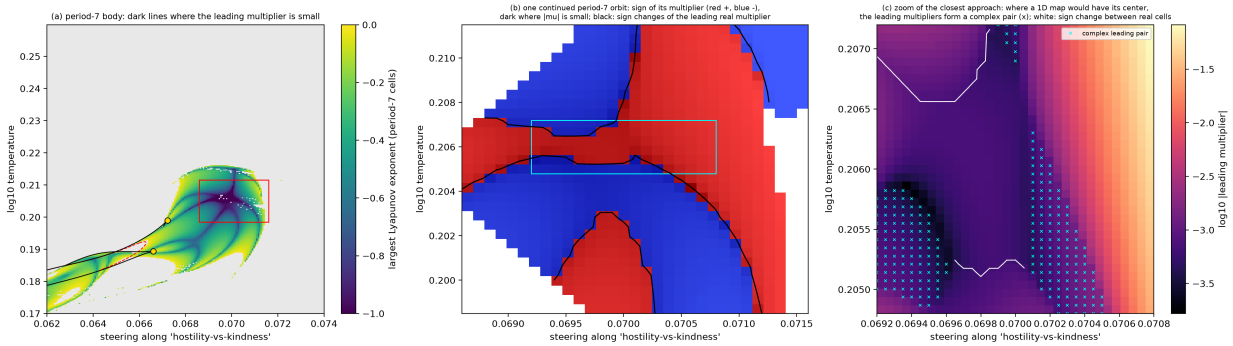

Figure 31: Small-multiplier structure of the hostile period-7 body (Qwen2.5-0.5B-Instruct, hostility direction at layer 12, float64). (a) Largest Lyapunov exponent of the period-7 cells, with fold (black) and flip (red, dashed) curves and the two located cusps (gold); dark lines mark a small leading multiplier. (b) One period-7 orbit continued over the red box of (a): sign of its leading multiplier (red positive, blue negative), dark where it is small, and the curves where that sign changes between cells (black). These are candidate superstable curves: the modulus is small along them, but a sign change between cells can also come from two real multipliers exchanging dominance. The two arcs approach near $(0.0700, 0.206)$ without crossing, a Type-B-like geometry. (c) Finer zoom of the cyan box of (b) (about 2 times finer in α, 6 times in $\log_{10} T$): logarithm of the leading multiplier's modulus. Crosses mark cells where the two leading multipliers form a complex pair; white segments mark sign changes between cells with a real leading multiplier. Where a one-dimensional map would have its center, the orbit's linearization has two comparable leading directions.

show the doubling cascade and the chaos as on a one-humped interval map (Figure 76). On the tracked branch and in the sampled region, the body is therefore not the shrimp of a one-dimensional map. In the two-dimensional theory of Façanha et al. the head is not generic (Section 11); arcs that turn away give a Type-B-like geometry, but they follow the sign of the leading real multiplier, not the trace that the theory uses, so this is a comparison, not a test.

Chaos behind a repeating text. The text records less than the loop feeds back. What returns at each step is the expected embedding under the whole next-token distribution; the printed token is only that distribution's most probable entry. A repeating text therefore says that the same token wins at each position of the repetition, not that the distribution repeats. Along the temperature line of Figure 30 (left), a longer test (2000 measured steps after 1000 transient steps, twins at offsets 10^{-6} and 10^{-9}, whose exponents agree to within 0.0031 in every chaotic cell) finds 182 chaotic cells among 1400 between $\log_{10} T = 0.176$ and 0.191. In 65 of them the leading token repeats exactly for all 2000 steps while the state moves chaotically; in 59 of these the period of the text is a multiple of 7, most often 7 or 28. Figure 33 shows the strongest case ($\log_{10} T = 0.18995$, largest exponent 0.091 per step at both offsets). Its 2000 sampled states fall into seven separated groups, one for each position in the seven-token text, visited in a fixed order and at least 0.098 of the state norm apart (a property of the samples, not a proof of separated invariant sets), and each group lies inside one cell of the partition by leading token, so the text `I!Yes heels`↩↩`pee`↩↩ repeats at every one of the 2000 steps we followed. Within each group the most probable token never changes, but the probabilities do: at the last position of the text, the probability of the double line break, which wins there, takes values between 0.235 and 0.254 and does not repeat over the 2000 steps. This is band chaos, the regime of a chaotic attractor that is a finite union of pieces visited cyclically [68, 69]. Over the measured horizon, the periodic sequence of leading tokens does not reveal the nonperiodic state dynamics: the finite record of leading tokens hides the variability of the states and their positive finite-time growth. In the terms of symbolic dynamics,

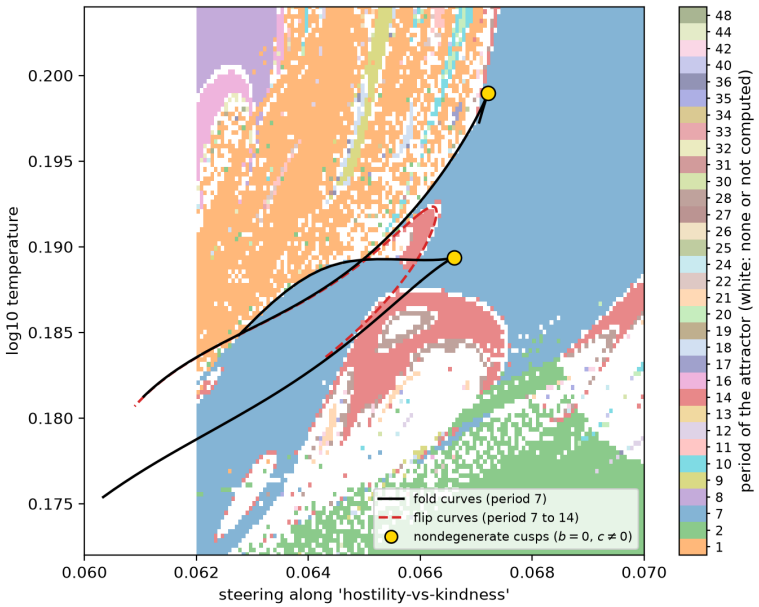

Figure 32: The period-7 body of the spiral region (Qwen2.5-0.5B-Instruct, hostility direction at layer 12): attractor periods of a float64 zoom (white at $\alpha < 0.062$: outside the zoom), the float64 fold curves of the period-7 orbit (black) and the flip curves into period 14 (red, dashed), with the two located cusps (gold).

this is consistent with a partition of the state space by leading token that is not generating on this banded attractor, but a finite record cannot establish that property, which concerns infinite itineraries. A periodic text can hide chaotic dynamics. Inside a group the sampled states lie close to a curve (a strongly anisotropic cloud). Read once per repetition through the probability of the last token, the sampled return relation is well approximated, on the sampled range, by a smooth map with a single quadratic maximum, the shape of the logistic map: a polynomial of degree 8 fits it to 4×10^{-4} of its range, and the fitted map's Lyapunov exponent, 0.647 per repetition, is within 2 % of the full map's 0.636 per repetition (seven steps of about 0.091), measured with twin trajectories in all 896 dimensions. On their sampled ranges, the scalar approximation fits most of the cascade's cells. Rerun with the leading-token sequences saved, 62 of the 65 cells are again chaotic behind a repeating text, and in every one of them several phases of the text (a median of ten) give single-valued return maps; in 59 cells the best-fitting phase has a single interior stationary point (a maximum in 36, a minimum in 23). If the dynamics were exactly one-dimensional and each phase's probability a smooth invertible coordinate on it, all single-valued phases would give the same asymptotic exponent, since exponents do not change under such a change of coordinates. We therefore fit each phase's map by a polynomial whose degree is chosen by cross-validation over random folds of the return-map pairs, and take the median over the phases. It matches the full map's exponent within 10 % in 51 of the 62 cells (median discrepancy 3.7 %; within 10 % in 23 of the 24 cells with seven-token text, which give 285 repetitions in 2000 steps; Figure 77). This analysis was adopted after we inspected an outlier of our first, pre-specified analysis (a fixed degree of 8 at the single best-fitting phase), which gave 49 of 62 and underfitted one cell whose maps are steep and have two turning points; we report both. An exploratory temporal split (fit on the first 60 % of each record, test on the last 30 %) gives small value-prediction errors at the best scalar phase, chosen on the test segment itself, but in 30 of 61 cells two- or three-delay predictors improve the held-out error by 20 % or more, which the pre-stated rule counts against scalar sufficiency. These results

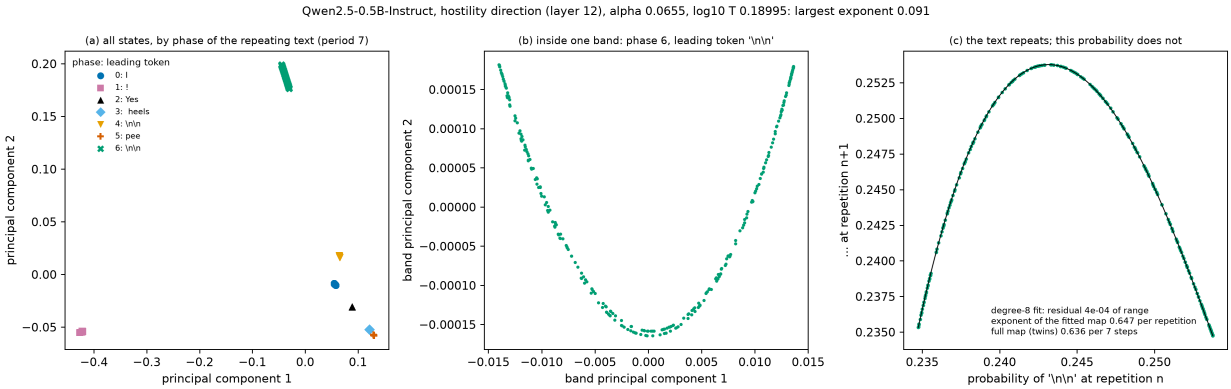

Figure 33: Chaos behind a repeating text, a periodic text on a chaotic orbit: the soft-token map of Qwen2.5-0.5B-Instruct steered along hostility at layer 12 ($\alpha = 0.0655$, $\log_{10} T = 0.18995$; float64, 2000 steps after 1000 transient steps). At every step the most probable token follows the same seven-token cycle, `I`, `!`, `Yes`, `heels`, ↩↩, `pee`, ↩↩, while the largest Lyapunov exponent is 0.091 per step: the token that wins repeats, but the probabilities behind it, which the loop feeds back, do not. (a) All states on their first two principal components, colored and marked by their position in the text: seven separated sampled clouds. (b) The widest cloud (the last position, whose most probable token is a double line break) on its own principal components: the sampled states lie close to a curve. (c) The probability of that token, which stays between 0.235 and 0.254 and does not repeat over the 2000 steps, at one repetition of the text against the next: the points fall on a single hump (black: a fitted polynomial of degree 8), and the fitted map's Lyapunov exponent, 0.647 per repetition, is within 2 % of the full map's 0.636 per repetition (seven steps of about 0.091).

support an empirical scalar description, not a proved one-dimensional reduction: a low-residual fit does not show that the observable is invertible, and even exact agreement of exponents would not establish a one-dimensional invariant curve (two uncoupled logistic maps observed through one coordinate agree in the same way). An independent implementation reproduced the strongest cell (the same leading token at all 2000 steps; exponent 0.089 against our 0.091; scalar discrepancy +2.3 % under the first rule and +1.5 % under the amended one) and both cross-cell counts (49 and 51 of 62). At two further cells it chose, it found good agreement after the amendment in one (text period 10, +0.2 %, although its fitted maps have both a maximum and a minimum) and a persistent discrepancy of about 18 % in the other (period 28), despite a small value error. These are the chaotic bands that follow a period-doubling cascade, and a smooth unimodal map is the setting of Feigenbaum's universality [24]; a fitted maximum on the sampled range does not by itself establish that theorem's hypotheses, and we have shown the scalar approximation for this cascade only. A reader of the text sees a loop. Only the probabilities, or the state, reveal the chaos.

11 Cusps and shrimp-like windows

The fold of Section 8 is one boundary of one window. Continuing such boundaries through the parameter plane, and following the orbits inside them, shows how the windows of the maps are organized, and which parts of that organization are low-dimensional. Appendices E and F report the continued curves and cusps of the period-4, -5 and -6 windows in full, and Section 10 those of the period-7 body; this section collects what they establish about the shape of the windows, after the cascade of Section 9 and the hostile windows of Section 10.

Fold curves with cusps. Fold and flip curves continued in float64 with the minimally augmented systems of Appendix B keep their critical multipliers at $+1$ and -1 (Appendix E). The fold curve of the 2-cycle "is it" traces the boundary of the period-2 region on which the stripes accumulate, as the mechanism of Section 8 requires. The fold curves of a period-6 window, of a hostile period-7 window, of a period-4 window and of a period-5 window steered at two depths bend back on themselves. On the eleven fold curves we continued, 13 of the 17 reversals of the parameter tangent are located cusps: the oriented quadratic coefficient of the fold's normal form changes sign there, while the cubic coefficient stays nonzero and the two-parameter unfolding is regular. The other four, all found on 18 October, are flagged by the continuation but not located: one on the lower-left leg's own fold curve, one on an extension of the period-5 window's second fold curve, and two on a partial trace of a fold curve of the same window at $\log_{10} T = -0.105$. This is not a global count of cusps. Near a cusp with negative cubic coefficient, three orbits of the same period coexist between the two fold branches, two of them stable. At a fold point near one cusp in each of two windows, the local cubic model, frozen before solving, matched the number of orbits that Newton solves seeded at its predicted roots recovered at all twelve parameter offsets, and their stability at eleven; the search was not exhaustive.

What the shrimp-like shapes are. The maps show bodies with legs, hooks and crescents that resemble the shrimp-shaped domains that Gallas found in the parameter plane of the Hénon map: a main body of period k, followed by domains of periods $k \cdot 2^n$, with thin legs [4]. They are made of the bifurcation curves of periodic windows; where two windows overlap, their strips can cross in an X (Appendix E). A periodic window is the region where one periodic orbit attracts. It is bounded by fold curves, where the orbit is born or dies together with a saddle (multiplier $+1$), and flip curves, where it doubles (multiplier -1); beyond the flips lie the period-doubled shells. In a one-dimensional map the multiplier is $+1$ on a fold edge and -1 on a flip edge, so it vanishes in between: every path across a window from a fold edge to a flip edge crosses a superstable curve. In the soft-token maps the dark lines of strongly negative exponent, where the leading multiplier is small, are the candidates for such curves. Gallas's dissection of one-dimensional shrimps names the parts of a cell (Figure 47a) [66]. The superstable curves of the map's two critical points cross at the head, where the orbit passes both critical points at once. At the end of the cell two legs, each a thin strip between a fold and a flip curve, cross in a kite. The kite begins at a cusp, where the two legs' fold curves meet; inside it two attracting orbits of the same period coexist, and their superstable curves cross at the tail.

Mira and coworkers sort the organizations around a cusp of a fold curve in two-dimensional maps into three kinds [70–72]. Passing through the cusp into its wedge splits an attracting node into a saddle and two nodes in a crossroad area, where the flip curves of the cusp's two branches cross, as in Gallas's kite, and in a spring area, where a flip curve loops around the cusp. It splits a saddle into a node and two saddles in a saddle area, where the attracting orbit exists only inside the wedge. The sign of the cubic coefficient c of the fold's normal form decides between the two splittings: $c < 0$ for crossroad and spring areas, $c > 0$ for the saddle area (Figure 34). Mira also noted that in two dimensions the attracting node becomes a focus, with a complex pair of multipliers, away from the fold and flip curves. In a two-dimensional map whose Jacobian is never singular, such as Hénon's, no multiplier can vanish. There the curves on which the trace of the cycle's two multipliers vanishes take the place of the superstable curves, and the head is no longer generic. Façanha, Oldeman and Glass found two unfoldings: in both the crossing at the tail survives, while at the head the curves either close into a loop (their Type A) or approach each other and turn away (Type B) [67].

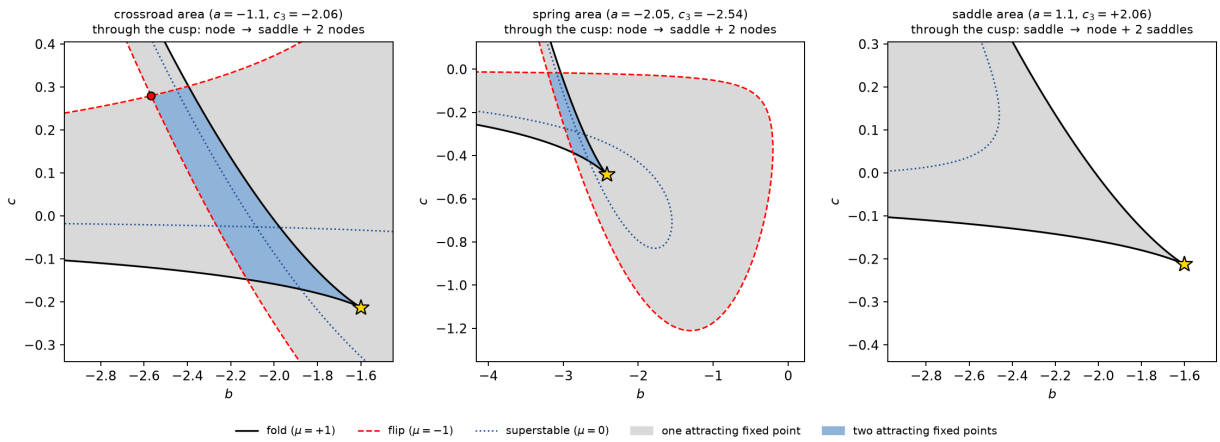

Figure 34: The three kinds of organization around a cusp of a fold curve, computed from the quartic map of Carcassès et al. [71], $x \mapsto c + (1+a)x + bx^2 + x^4$, in planes of b and c at fixed a, as they drew them (gray: one attracting fixed point; blue: two). The cusp (gold) lies at the fixed point $x = t$ with $a = 8t^3$, and $c_3 = 4t$ is the cubic coefficient of its fold normal form (written c in the text). Left, $a = -1.1$: a crossroad area; passing through the cusp splits a node into a saddle and two nodes, and the two attracting fixed points coexist in a kite whose far corner is the crossing of the flip curves of the cusp's two branches (red dot). Middle, $a = -2.05$: a spring area; a flip curve closes into a loop that winds once around the cusp (the loop forms at $a = -2$, the codimension-three point a_7 of Carcassès et al.). Right, $a = +1.1$: the saddle kind; passing through the cusp splits a saddle into a node and two saddles, and the attracting fixed point exists only inside the cusp's wedge. Each frame contains the whole defining structure. Period-2 curves are not drawn.

The skeleton of one window. The plane steered at two depths (Section 6) holds a period-5 window (Figure 45). A 31×31 grid of the window's continued orbit (Appendix F) left two anomalies, signs of folds it could not resolve, so we continued the window's fold and flip curves in float64. One fold curve bounds the body on its right, runs over its top and along one side of the upper and lower-right legs. Flip curves bound the body's upper-left and lower sides and the other side of each leg, so each leg lies between a fold curve and a flip curve, as in Gallas's anatomy, and beyond the flips lie shells of periods 10 and 20. Where the upper leg leaves the body, the fold curve has two located cusps that share a segment of the curve, and its other two branches cross (Figure 45d). At the crossing the two arms are folds of two different period-5 orbits with the same token loop (their states differ by about 4 %). A second fold curve, on the far left, where the orbit dies after regaining stability beyond a doubling bubble, has two more located cusps, which also share a segment, and its other two branches cross as well, again as folds of two different period-5 orbits (states about 9 % apart; Figure 45a). Along the curves continued on 15 October the critical multiplier stays at $+1$ or -1 to within 1.2×10^{-7}, and no other multiplier exceeds 0.4 in modulus; with the curves of 18 October (their extensions toward the legs' tips and the lower-left leg's own pair) these bounds become 1.9×10^{-7} and 0.87, the largest second multiplier (0.86) lying where the lower-right leg's curves reach the bound of the continuation. No second multiplier reaches the unit circle along them; this does not exclude cusps, which need only a nonlinear degeneracy.

The shrimp itself. From the fixed start, a 512×512 map of the window (Figure 46a; run HP5) shows the window's family in 59 % of the cells of the followed domain; in 24 % the fixed start reaches another periodic attractor, mostly the coexisting period-4 attractor and its doublings (periods 4, 8, 16 and 32), in 12 % a largest exponent above 10^{-3} (chaos; at 33 cells it overrides a detected period), and in 5 % neither. To see the window whole, we followed its attractor from the body outward

over a grid of spacing 1.25×10^{-4} in both strengths (Figure 46; `tools/follow_attractor.py`). Each cell starts from the final state of a neighboring cell that is still on the attractor and runs 250 steps, followed by 300 recorded steps with a twin. It stays on the attractor only if its orbit has a period that is a multiple of five and the body's five-token text in at least four of its five places, or no detectable period, an exactly repeating text and an exponent below 0.005 (a slowly converging orbit next to a fold or a flip); cells with chaos behind the same text are recorded but not continued. (A first rule, which compared each cell only with its neighbor, leaked across the fold into the surrounding chaos, where intermittent motion keeps the same text most of the time.) The followed attractor fills a shrimp: a body against the fold curve and three long legs, each between a fold curve and a flip curve, reaching $(\alpha_1, \alpha_2) = (-0.1148, 0.109)$, $(-0.1252, 0.0918)$ and, on the left, $\alpha_2 = 0.108$, with shells of periods 10, 20 and 40 beyond the flips and chaos behind the same five-token text outside them; a fourth leg, along the lower right, narrows from 5.1×10^{-5} in α_1 near the body to 3.5×10^{-8}, far below the grid step, so the follow does not enter it. Along its center line, each point started from the orbit continued along the leg's fold curve (`tools/leg_centerline.py`), the attractor of the body's family, with the body's text, is present at all 128 points from the body down to $\alpha_2 = 0.075$, the bound of the continuation (period 5 at 127, period 20 at one); one leg width beyond the fold none of them keeps it, and one leg width beyond the flip 118 of 128 carry the doubled orbit or chaos behind the same text (Figure 46b). Of 2173 cells computed, 1720 are on the attractor: 822 with period 5, 371 with period 10, 82 with period 20 and 11 with period 40, a few with periods 15, 30 and 50 inside chaotic bands, 240 chaotic behind the same text and 178 converging slowly. At $\log_{10} T = -0.105$ the followed domain has the same anatomy (1760 cells on the attractor; Figure 71): its period-5 body moves by $(+0.00085, +0.00236)$ in (α_1, α_2), and the head, continued in temperature, by about $(+0.00105, +0.00247)$ over the same step, so the window drifts as a whole. The computed curves lie where the followed domain ends. The continuations from the body do not reach the lower-left leg's own boundary, even when rerun with wider parameter bounds, so we marched the period-5 orbit into the leg by Newton's method, through the 68 cells by which the follow run entered it, and continued it there (Appendix F, Figure 72). The leg is a strip between a flip curve on its left and a fold curve on its right. Every period-5 cell of the leg lies inside that strip (39 of 39), and no computed cell inside it is lost or has another period. With these curves, 276 of the 292 pairs of neighboring cells with periods 5 and 10 straddle a flip curve to within one grid step, and 290 to within two. Of the 754 pairs in which a period-5 cell borders a cell where the attractor was lost, 717 lie within one step of a fold curve and all 754 within two steps of a fold or flip curve, the farthest 1.7 steps away; with the body's curves alone these counts were 251, 504 and 544. Below $\alpha_2 = 0.0934$ the strip is narrower than one grid step (0.015 steps at $\alpha_2 = 0.080$, where the continuation reached its bound), so the leg goes on below the followed domain, too narrow to be followed cell by cell.

Shrimps of a two-dimensional kind. The period-5 window has this anatomy; only its head differs. At every cell of the 31×31 continuation grid and every point of the curves continued on 15 October, the multipliers beyond the leading pair are below 0.0032 (on the curves of 18 October at this temperature that recorded them, they stay below it inside the followed domain and reach 0.010 beyond it), so on the grid the orbit's linearization has at most two non-negligible directions. On the grid the leading multiplier never vanishes (its modulus stays at least 0.0139): where a one-dimensional map would be superstable, the two leading multipliers form a complex pair, the node having become a focus, or pass each other as real numbers of opposite sign. Between grid cells, however, a two-parameter Newton search for the point where the sum of the leading multipliers and the sum of their pairwise products vanish converges at $(\alpha_1, \alpha_2) = (-0.1184836, 0.1017217)$

(`docs/review/p5-nilpotent.jsonl`). There the three leading multipliers have nearly equal moduli, 3.4 to 3.6×10^{-3}, as the two conditions would make exactly equal for three roots alone (they would make the roots the cube roots of their product). The leading modulus sets the asymptotic linear contraction, a factor of about 280 per period, but a single return can still amplify a perturbation: at one phase the norm of the return map is about 65 (below). This point of small leading multipliers, which we call near-nilpotent (the leading eigenvalues of the monodromy matrix nearly vanish, although the matrix itself does not), plays the role of the head. The nearly equal moduli do not by themselves decide whether the return map is locally of higher rank there or nearly of rank one and strongly nonnormal; its singular values decide it (`tools/monodromy_svd.py`, checked on synthetic operators with equal eigenvalues and different rank). At the point the second singular value of the return map is 1.8×10^{-4} of the first at one starting phase (1.7 to 1.8×10^{-4} at four points 5×10^{-5} away) and at most 0.031 at the others, while the first, at that phase, is about 65, four orders of magnitude above the eigenvalues; in the body, away from the head, the ratio is 0.017 to 0.23. Near its head the return map is nearly of rank one and strongly nonnormal. Its leading rank-one part is nearly nilpotent because its left and right singular vectors are nearly orthogonal, which shows how a large norm can coexist with small eigenvalues. That part's eigenvalue, however, is about 13 times smaller than the full leading modulus: the remainder, whose norm is the second singular value (about 0.011 at that phase), exceeds the leading modulus and matters for the full spectrum. In the reduction to this rank-one part, $\lambda = \sigma_1 (v_1 \cdot u_1)$, the nonzero eigenvalue of $\sigma_1 u_1 v_1^\top$, which stays defined where the eigenvalues collide (at the head $\lambda = -2.7 \times 10^{-4}$, so the two directions are orthogonal to about 4×10^{-6}), the head lies on a single zero curve: on an 11×11 grid of spacing 10^{-5} around it, λ changes sign exactly once in every row, along one nearly straight line through the head (Figure 70). Even in this reduction the head is not a crossing of two superstable curves. Continued in temperature, over the eleven values of $\log_{10} T$ from -0.104 to -0.094 at which the continuation passed its gates, the point moves along a nearly straight line; the leading modulus stays between 2.4 and 5.6×10^{-3}, and a fourth multiplier approaches it twice (Figure 69; `tools/nilpotent_track.py`). The product $D = \mu_1 \mu_2$ of the two leading multipliers also changes sign in the window ($D < 0$ at 41 of its 447 attracting cells), so the linearization restricted to the leading pair reverses orientation in part of the window. This differs from the setting of Hénon's map, whose two-dimensional Jacobian determinant has a fixed sign, but it does not by itself establish noninvertibility or a two-dimensional reduction of the full map: the pair is selected by modulus and can change identity, and where D changes sign the second multiplier is as small as those beyond the pair. The trace-zero curve around an island of negative trace (Figure 45c) is closed on the grid except at one cell next to the fold curve, where it must close, because on a fold the trace is 1 plus the second multiplier and therefore positive: a closed trace-zero contour as in Type A, a classification derived for determinants of fixed sign [67]. In the hostile period-7 body the curves on which the leading real multiplier changes sign, candidate superstable curves of the continued orbit, form two arcs that approach each other and turn away without crossing, a Type-B-like geometry, where the leading multiplier becomes as small as the next (Figure 31; Section 10). The period-6 body ("A second shrimp, in the plane of the crossing window", below) shows the same geometry: its pseudo-superstable lines, where the leading multiplier changes sign without vanishing, meet below the kite and approach and turn away at the top. The same plane holds a period-doubling cascade whose cells are nested V-shaped tongues near $\alpha_1 = -0.105$, $\alpha_2 = 0.105$. Continued in float64 on grids over these cells, its period-2, -4 and -8 orbits read the same way: no cell reaches superstability (the smallest leading moduli are 0.222, 0.041 and 0.034), and complex leading pairs and trace-zero curves take the place of the superstable curves. The cells become more one-dimensional with each doubling: over their attracting cells, the median ratio of the second to the leading multiplier is 0.59, 0.125 and 0.027. Per step, the second multiplier contracts by a nearly constant factor (medians 0.50

to 0.57), so its share shrinks like a power of the period; the per-step medians predict 0.52, 0.146 and 0.029, within 17 % of the measured ratios. In the terms of the Hénon map, whose Jacobian determinant $-b$ is constant, so that $\mu_1\mu_2 = (-b)^p$, the per-step area contraction $|\mu_1\mu_2|^{1/p}$ plays the role of $|b|$; its medians are 0.41, 0.41 and 0.50. In the period-4 and period-8 cells the trace-zero curves run along both arms of the V. Finer grids at its bottom resolve them: in both cells the two arms are one U-shaped trace-zero curve, and they do not cross. In the period-8 cell the product of the two leading multipliers is positive throughout and the third multiplier is at most 6×10^{-5} of the second, so the reading is two-dimensional to that accuracy. In the period-4 cell that product changes sign, and near the bottom of the U the second and third multipliers form a complex pair, so the leading pair is not isolated there. Neither cell has a crossing of trace-zero curves, the analog of a one-dimensional shrimp's head, and the period-4 cell has no near-nilpotent point at the bottom of its U either: a two-parameter Newton search on the invariants of the leading spectrum, as for the period-5 window, converges where its three leading multipliers have nearly equal modulus, 0.0325, and lie 120° apart, ten times the period-5 point's 3.5×10^{-3} (`tools/nilpotent_search.py`). Under the decision rule fixed before these grids, both cells remain unresolved between the loop and turn-away types (the first period-8 grid left the bottom of the V unsolved because of a defect in the continuation, fixed for the second; `docs/dss`).

The junction where the upper leg leaves the body, two cusps that share a segment of the fold curve and a crossing of its other two branches, is the configuration that Mira et al. call a dovetail for critical curves with two such adjacent cusps [73]; we borrow the name, with periodic orbits in place of preimages. With the crossing, the configuration is consistent with a plane section of a swallowtail, whose geometry predicts four orbits in the pocket between the cusps and the crossing, none beyond the crossing and two elsewhere (Figure 47b). An orbit census at 13 points fixed before the computation found exactly these numbers at nine (Figure 45d): four orbits at three of five pocket points, none at either point beyond the crossing, and two at four of five surrounding points. At three other points it found fewer orbits, but all the attracting ones the model predicts (the counts are lower bounds); at one point, beyond a flip curve, it also found an orbit with another token loop. At all five pocket points of the census two attracting period-5 orbits coexist; near the upper cusp a flip curve crosses the pocket, beyond which one of them has doubled. The pocket therefore plays the kite's part, a bistable region at the end of a leg, with two cusps and a crossing where Gallas's kite has one cusp. In the quartic models of Figure 47, both configurations are plane sections of the swallowtail of the quartic fixed-point equation, and the window's configuration is consistent with such a section. Lifted from the parameter plane into the orbit states, the fold curve is smooth: its cusps are only where its shadow folds, and the two arms that cross in the shadow lie on different sheets, as in the model (Figure 48). The window's second fold curve has a second dovetail. In each, the two cusps have cubic coefficients of opposite sign, as in any section of a swallowtail: one of the saddle-area kind ($c > 0$, the upper cusp in both) and one of the crossroad or spring kind ($c < 0$). We do not claim the swallowtail point itself, which needs a third parameter, and we have not traced the tail [66], determined symbolic sequences or tested for a periodicity hub [64, 65]. With these qualifications, the period-5 window meets the three numerical criteria we required before calling a window a shrimp: primitive periodic orbits of the full state (float64 Newton solves), their stability (the multipliers and the followed attractor), and the organization of their boundary (continued fold and flip curves with located cusps, matched by the edge of the followed domain). It is a numerically verified shrimp in a two-parameter plane of the full map, of the kind Gallas found in Hénon's map [4, 74]. Its head resembles those of two-dimensional maps, a closed trace-zero curve rather than crossing superstable curves, but that likeness is geometric: no two-dimensional reduction of the dynamics is established (a closure test failed for numerical reasons; Appendix F), the product of the two leading multipliers changes sign, and none of this is proved by interval arithmetic.

A second shrimp, in the plane of the crossing window. The plane of Figure 11, the philosophy direction at layer 12 against the temperature, holds a second shrimp-shaped window: the period-6 loop "\n\nThe question\n\nThe existence" (Figure 38; runs P6S, HR9, HR10, DSS6 and KITE). At a strongly contracting cell of the map, $(\alpha, \log_{10} T) = (0.073097, -0.15204)$, Newton's method in float64 converges to a primitive period-6 orbit (residual 1.6×10^{-14}) whose leading multipliers form a complex pair of modulus 0.016. Followed from there through neighboring cells on the grid of Figure 11, the attractor fills 4224 cells. They cover the shrimp's body and its legs, and below the body they continue through a narrow neck into a smaller, distorted window at the lower right of Figure 38(c). Of these cells, period 6 holds 1946, periods 12, 24 and 48 hold 945 along the edges, other periods 35, chaos behind the repeating text in 811 (an exactly repeating text, no period up to 64 and an exponent of at least 0.005) and slow cells in 487 (the same with an exponent below 0.005; with the maps' threshold of 10^{-3} the split would be 842 and 456). Continued from the body in four directions, up and down in α and in T, the period-6 orbit ends at a fold above and on the right and at flips on the left and below. Continued in the plane with every point checked (residual at most 10^{-10}, critical multiplier real and within 10^{-7} of ± 1 on the body and kite curves; four points of the extended leg curves reach 2.9×10^{-7}), the curves outline the anatomy of Gallas's cell (Figure 47a): the fold curve bounds the body above and on the right, one flip curve bounds it on the upper left and a second along its bottom, and the doubled shells lie beyond both. Within the window of Figure 38a, the edge of the followed domain lies within two grid steps of the curves for all 232 neighboring pairs of periods 6 and 12 and all 459 pairs of a period-6 cell and a cell where the followed attractor was lost (median distances 0.19 steps); over the whole followed domain the shares fall to 388 of 916 and 1,720 of 1,888 (medians 4.3 and 0.24 steps), because the far parts of the legs leave the traced curves, so their boundary is not resolved. A 640×640 fixed-start map centered on the body (steps 4.4×10^{-6} in α and 1.25×10^{-5} in $\log_{10} T$, about seven times finer than Figure 11; 66.5 % of its cells chaotic) shows the body with its dark lines of strongest contraction, the flip region along its bottom and two of its legs leaving its lower-left end, one straight down along the left flip curve and one to the lower left (Figure 38a).

The two flip curves cross at that end, near $(0.072912, -0.154104)$. They are flip curves of two different orbits: Newton's method at that point (`tools/kite_newton_at_R.py`), started from the orbit stored on each curve, converges to two distinct period-6 orbits (residuals 3×10^{-14} and 2×10^{-13}), with the same six-token text but 6.4 % of the mean state norm apart (the largest point distance after the best of the six cyclic alignments). Their leading multipliers are -0.9952 and -1.0005 at this approximate crossing of the two traced curves, whose nearest saved points lie 1.7 and 0.3 steps of the grid of Figure 11 away. The crossing is a pseudo-crossing of curves on two sheets of the orbit family, the configuration at the corner R of Gallas's kite, where the flip curves of two coexisting fixed points cross [66, Fig. 7]. Inside the body near it, at $(0.0730, -0.1539)$, the orbit of the lower flip curve is attracting (leading multiplier -0.797); Newton's method from the other orbit did not converge there: the point lies just beyond that orbit's fold, on the branch C–A of the kite (Figure 39), where the orbit no longer exists. The rest of the kite follows from continuation (Figure 39; run KITE). Continued in α at fixed $\log_{10} T = -0.15395$ from the first orbit's flip, the branch is S-shaped: the orbit stays attracting up to a fold at $\alpha = 0.072992$, the branch returns along the saddle to a second fold at 0.072793, and it continues on the second orbit, which flips near 0.07301; at that temperature the two attracting orbits coexist for α between 0.072907 and 0.072992. The second orbit, continued upward in $\log_{10} T$ from its flip at fixed α, is attracting up to its own fold at -0.153657. The fold curves traced from the two folds turn back at the same point, the kite's cusp C, located at $(0.0730186, -0.1532803)$ (cubic coefficient $c = -1.33 \times 10^4$, unfolding determinant $\det K = -2624$, critical multiplier within 7.7×10^{-9} of $+1$, next modulus 0.00095), on the side of R toward the body. Its two branches meet the two flip curves at A $= (0.072991, -0.153984)$

and B = (0.072900, −0.153707), completing the kite C–A–R–B of Gallas's cell. A and B are pseudo-crossings too: at the continued fold points nearest them the orbit's other multipliers are at most 0.0022 in modulus, none near −1, and at the flip points nearest them the others are at most 0.0042, none near +1. A fold–flip point would need one orbit with both multipliers; here, as at R, the curves belong to different period-6 orbits. Besides C, the fold curves have five more located cusps (Table 20): four in two pairs near the ends of the upper-left and lower-left legs, the two cusps of each pair joined by a short segment of the fold curve and with cubic coefficients of opposite sign, the configuration of the period-5 window's upper leg; and one at the tip of a small separate window beyond a narrow neck below the body, where the orbit is stable again after a doubling; two branches, run independently, trace that fold curve and its turn to within one curve step.

On a 31 × 31 grid of the continued orbit over the body (α from 0.0724 to 0.0738, $\log_{10} T$ from −0.1545 to −0.1500; the run reached 582 of the 961 cells before its time limit, and 529 of them converged), the leading multiplier is nonzero at every solved cell. Its smallest modulus on the grid, 0.0063 at the cell (0.073053, −0.15120) on the line of strongest contraction, belongs to a nearly imaginary leading pair, $0.0004 \pm 0.0063i$, with a third multiplier of modulus 0.0051: three multipliers of nearly equal small modulus, the near-nilpotent signature of the period-5 head. The dark line of strongest contraction runs from the flip crossing up the body and bends to the upper left, and all nine cells with a complex leading pair lie on its upper part (arguments 43 to 135 degrees, moduli 0.006 to 0.038). At 289 solved cells the two largest-modulus multipliers form a real or complex-conjugate pair, and their real product is negative at 143; at the other 240 the leading real multiplier is followed by a complex-conjugate pair, so the first two alone define no real two-dimensional determinant, and we infer no determinant-zero curve from them. The real leading multiplier changes sign across 45 neighboring solved pairs; only two have leading and next moduli within a factor three at both ends, and this endpoint test does not tell a multiplier crossing zero from an exchange of dominance. These sign changes mark the body's dark ditches (Figures 36 and 37), which take the place of Gallas's superstable curves. A leading multiplier is the largest in modulus, so it vanishes only if all multipliers do, and between the two cells of each pair it changes sign in one of two ways: it meets the next multiplier, the two leave the real axis as a complex pair and return to it with the opposite sign, or a multiplier of the opposite sign grows past it in modulus and becomes the leading one, a hand-over. Two lines of cells computed locally in float64 (run DITCH) show both outcomes. Across the main ditch at $\log_{10} T = -0.1530$ (21 cells, steps 7×10^{-6} in α) the negative leading multiplier shrinks from −0.605 to −0.060, the lead passes at $\alpha = 0.072976$ to a complex pair, $0.0069 \pm 0.0229i$, and the pair splits into two positive multipliers, the lead rising to +0.71. Across the dimmer ditch that closes the lighter region at its top, at $\alpha = 0.07318$ (26 cells, steps 4×10^{-5} in $\log_{10} T$), the leading multiplier stays positive: it falls from 0.122 to a minimum of 0.064 at $\log_{10} T = -0.15202$ and rises to 0.171, while the next stays negative and below 0.021 in modulus, a valley without a sign change. Ten times finer, on 225 cells around that minimum (steps 1.8×10^{-6} in α and 4×10^{-6} in $\log_{10} T$, run DITCHFINE on two rented GPUs), the leading multiplier is again real and positive everywhere (0.056 to 0.075), with the same minimum at $\alpha = 0.07318$ (0.0638). It falls toward smaller α, where the valley runs toward the main ditch, so line A's minimum is where that line crosses the valley, not the valley's lowest point. A patch of the continued orbit over $\alpha \in [0.07300, 0.07336]$ and $\log_{10} T \in [-0.15294, -0.15170]$ (21 × 32 cells, steps 1.8×10^{-5} and 4×10^{-5}, four GPU bands, run DITCH2D; 595 cells solved, the unsolved ones at the right edge, next to the fold) has 41 sign changes and 33 complex leading pairs: the arc closes over a pocket of negative leading multiplier, with its peak near (0.07326, −0.15232), and no sign change crosses the region between that peak and the main ditch. The two pseudo-superstable lines thus meet below the kite's cusp C, as Gallas's superstable curves meet in his kite, and at the top they approach and turn away without crossing, a Type-B-like geometry [67]. On both sides of every sign change the continued orbit is the same

attracting period-6 orbit, with the six-token text of the seed and Newton residuals below 10^{-12}; fixed-start scans at the nearest parameters find period six at all patch and line cells but two at the patch's edge, which are unresolved, and on the coarse grid at 44 of the 45 pairs and 6 of the 9 complex cells (the others, at the upper left, reach a coexisting attractor of period 8, 16 or 32). These are nearby parameters, not same-parameter full-state comparisons, so whether the fixed start reaches this very orbit was not checked at every continued point. We examined this separation in run DELTA2D on four rented GPUs (`tools/delta2d_analysis.py`, `docs/delta2d/analysis.json`; Figure 35). A 27×37 grid over the top of the body (steps 10^{-5} in α and 2.5×10^{-5} in $\log_{10} T$) gave 852 continued orbits, with the seed's six-token text, Newton residuals at most 1.0×10^{-12} and the three leading Ritz vectors stored at every cell; the 147 unsolved cells have no converged orbit on the followed branch and lie in the corner beyond the located fold. A two-dimensional block of the monodromy, initialized from the leading Ritz directions at the best-separated cell, was tracked outward by subspace overlap, the largest mean squared cosine of the principal angles among the spans of two of the three stored Ritz directions (so the tracked pair is not always the two locally largest). The tracked 2×2 block has trace τ and determinant δ; within this block, the eigenvalue of larger modulus, or the real part of a complex pair, has the sign of τ. The sign of τ matches that of the leading multiplier in 849 of the 852 cells (the other three have low tracking overlap, indicating possible mixing with a third direction); every one of the 42 measured sign changes of a real leading multiplier lies within one cell of a sign change of τ; and 60 of the 64 sign changes of τ lie within one cell of a measured sign change or a complex leading pair. This resembles the trace-zero skeleton that Façanha, Oldeman and Glass use for shrimps in two-dimensional maps [67]. Between the main ditch and the arc, a local quadratic fit to τ gives a saddle at $(0.073188, -0.152083)$ with the value 0.045 (Hessian eigenvalues -8.0×10^{-4} and 3.5×10^{-3} per squared cell, from a fit over 5×5 cells with residual 3.9×10^{-4}). The fitted saddle and the measured zero curves support a local trace-based explanation of why the two lines approach and turn away without crossing. The leading multiplier has a fitted saddle nearby, 0.067 at $(0.073195, -0.152138)$ (Hessian eigenvalues -5.6×10^{-4} and 2.8×10^{-3}), a reading fixed before the grid's data arrived. The finer patch lies just above it, on the flank toward the main ditch, which is why the patch's lowest value sits at its corner. The run card asked instead for a saddle of δ, and that test fails: none was found on the sampled patch under the specified fitting and filtering procedure, δ follows the steady drift of the second multiplier (from -0.010 to -0.034 across seven columns at the saddle), and its zero set misses the arc (11 of its 32 sign changes lie within one cell of a measured sign change). The trace was examined after the data arrived, as a disclosed amendment. In 38 of the 852 cells the tracking overlap falls below 0.9, indicating possible mixing with a third direction. This local evidence does not establish an exact two-dimensional reduction of the full map.

Near the kite, where two orbits exist, curves that belong to different orbits can cross in the parameter plane without any degeneracy, as the superstable curves of Gallas's two coexisting fixed points do at his tail. With these qualifications the period-6 window meets the three numerical criteria along its body, with the kite of Gallas's cell and its cusp C located (five more cusps close the ends of the legs and the small window beyond the neck): a second shrimp-shaped window of the full map, consistent with Gallas's organization (in 896 dimensions the name describes that organization of continued curves, not a local normal form), in the plane of Figure 11 itself. The far parts of its legs remain unresolved.

The hook. A 640×640 map of the upper end of the kite (Figure 40; cells 3.1 and 2.5 times finer than in Figure 39a) resolves the hook-shaped light wedge and the striations below it. The wedge is bounded above by the kite's fold curve, where its orbit disappears, and below by a smooth

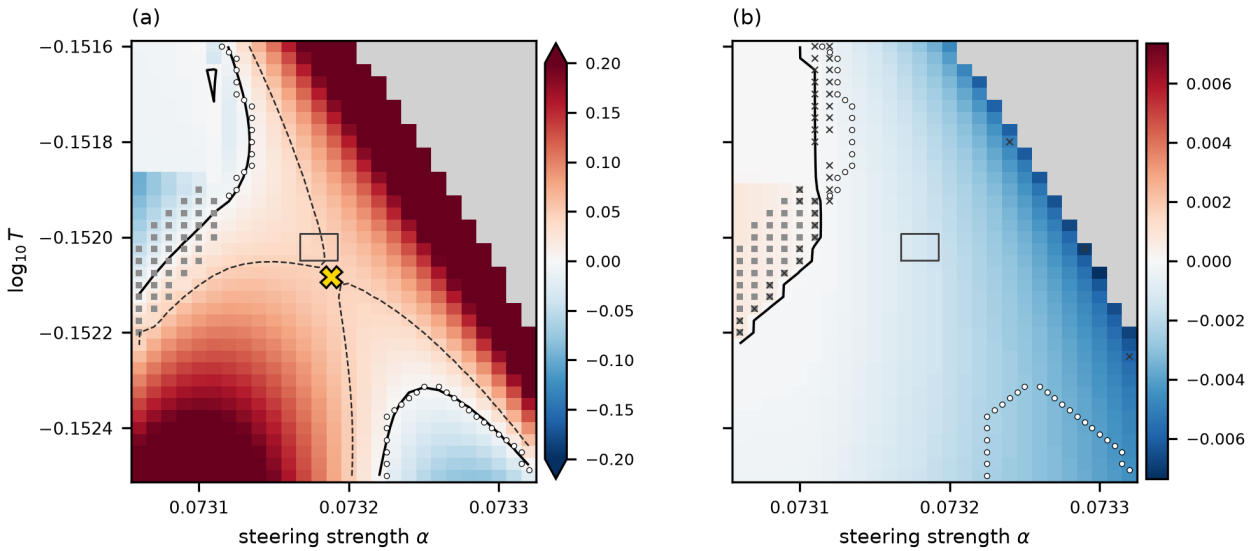

Figure 35: Trace geometry associated with the separation of the two pseudo-superstable lines at the top of the period-6 body (run DELTA2D, `tools/plot_delta2d.py`). A two-dimensional block of the monodromy, initialized from its leading Ritz directions, was tracked from cell to cell by subspace overlap over 852 continued orbits (float64; steps 10^{-5} in α and 2.5×10^{-5} in $\log_{10} T$; gray: no converged orbit on the followed branch). (a) The trace τ of the tracked block (colors clipped at ± 0.2), its zero curves (black), the measured sign changes of the real leading multiplier (white dots) and the cells with a complex leading pair (gray squares), and the numerically fitted saddle of τ (gold cross, $\tau = 0.045$) with the level curve through it (dashed). (b) The determinant δ, with its zero curve and the same measurements; small crosses mark cells where the tracking overlap falls below 0.9, indicating possible mixing with a third direction. The rectangle marks the patch around the minimum on line A, refined by a factor of ten in each parameter relative to the orange patch in Figure 37(a).

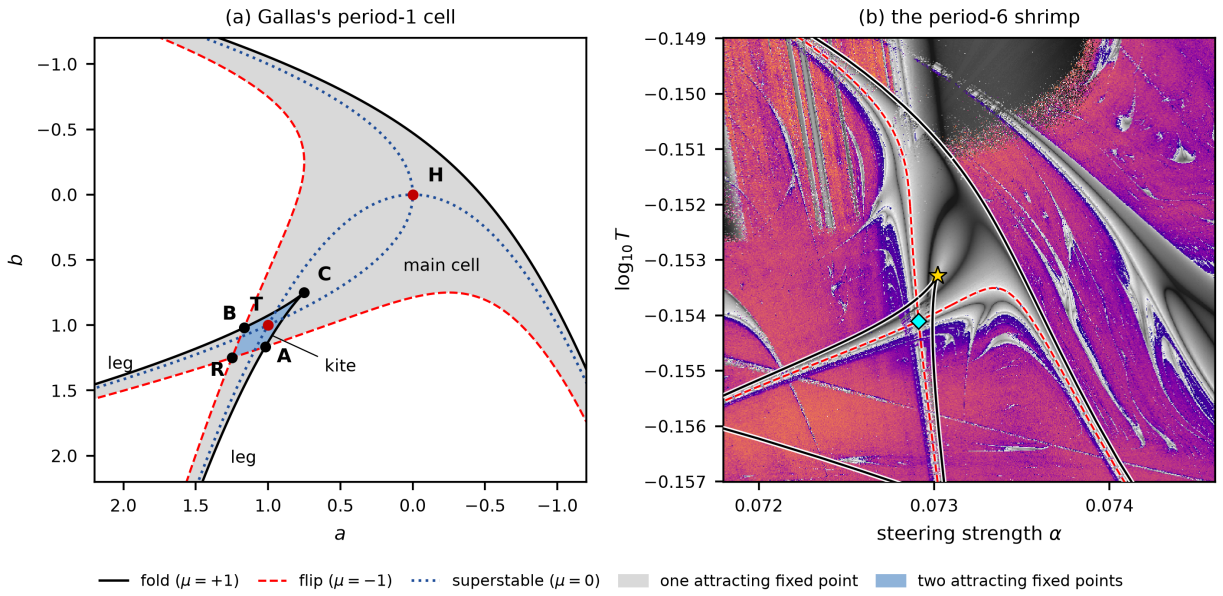

Figure 36: Gallas's canonical cell beside the period-6 shrimp, in the same orientation, with the shrimp's pseudo-superstable lines. (a) Gallas's period-1 cell of Figure 47a, drawn with both axes reversed, a rotation by 180° that brings it to the orientation of (b) (gray: one attracting fixed point; blue: two; fold curves black, flip curves red and dashed, superstable curves blue and dotted; head H, cusp C, kite C–A–R–B, tail T). (b) The period-6 shrimp: the 640 × 640 float64 map of Figure 38a with the continued fold and flip curves; R (cyan diamond): the pseudo-crossing of the two flip curves; C (gold star): the kite's cusp. Along the dark ditches the continued period-6 orbit's leading multiplier changes sign between neighboring cells or becomes a complex pair, and the orbit is attracting throughout (the measurements are drawn in Figure 37). The multiplier never vanishes, so these dark ditches are pseudo-superstable lines; they meet below C, where the superstable curves of (a) cross at the tail T, but approach and turn away at the top, where (a) has its head H, a Type-B-like unfolding of the one-dimensional skeleton [67]. In both panels the fold curve bounds the main cell (in (b), the body) above and on the right, flip curves bound it on the left and below, and at the lower left two of its legs cross in the kite.

curve where the fixed start switches to the other attracting orbit; at adjacent cells across that curve both orbits exist and attract (multipliers 0.18 and 0.60 to 0.61; pair 1). The striations are tongues bounded the same way: across the edge of one, two other attracting period-6 orbits exist at both cells (multipliers −0.72 and −0.11; pair 2), and inside several tongues the exponents of neighboring cells fall into two separated groups, so the fixed start's outcome alternates from cell to cell, a structure not resolved at this grid spacing (`tools/plot_p6_hook.py`, `tools/kite_basin_check.py`). From a random token's embedding as the start, the same box at 128 × 128 cells shows neither the hook nor the tongues, and only 38 % of its cells keep their exponent within 0.02 of the nearest cell of the fixed-start map (run HR12-RS): these edges are set by where the fixed start lies relative to the basins, as basin boundaries should be, not by bifurcations.

Why the lower edge lies where it does. An edge of the fixed start's outcome that is not a bifurcation is where the start crosses the boundary between two basins, and the boundary between the basins of two attracting orbits is typically the stable manifold of a saddle orbit between them. At $\alpha = 0.072975$ we bisected the wedge's lower edge in $\log_{10} T$ to the float64 limit, at $\log_{10} T = -0.1535299666043828$ (`tools/hook_saddle_check.py`). Started there, the loop comes close to a third period-6 orbit and stays near it: Newton's method from the most nearly periodic

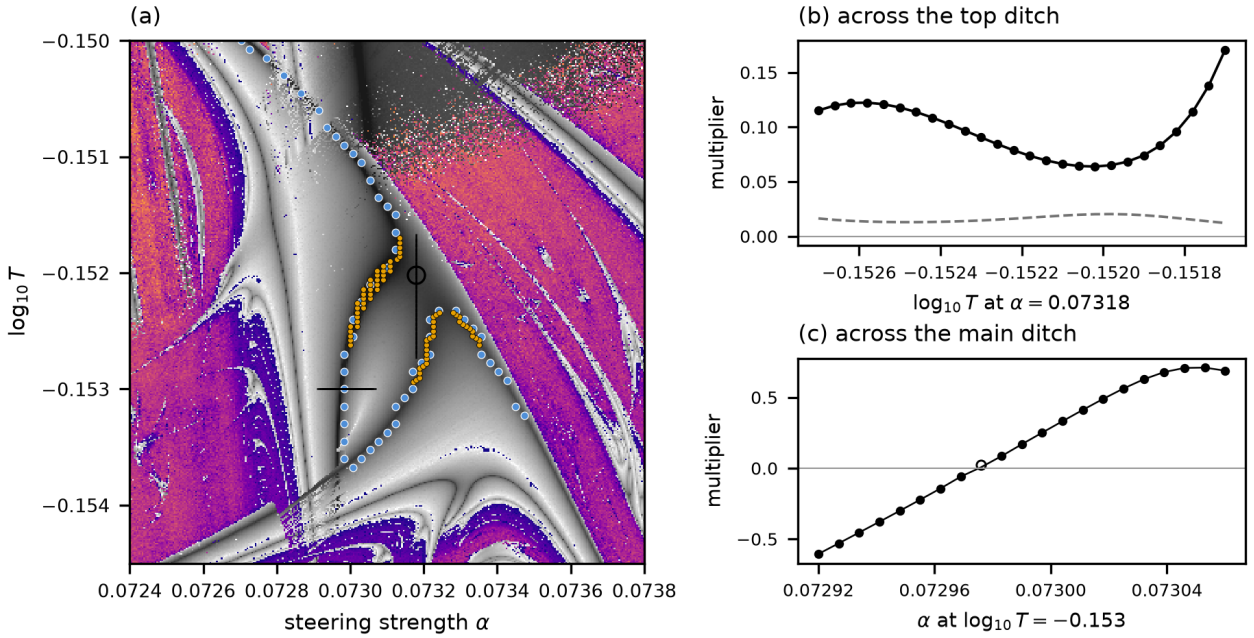

Figure 37: Where the period-6 orbit's leading multiplier changes sign over the body of the shrimp (float64, `tools/plot_p6_ditches.py`). (a) The map of Figure 38a over the body, with three measurements of the continued orbit: the 31×31 grid (blue dots: midpoints of neighboring solved cells whose real leading multipliers have opposite signs, and cells where the leading multiplier is a complex pair), a patch across the top of the lens 1.8×10^{-5} by 4×10^{-5} in its steps (orange dots, the same rule), and two lines (black ticks): line A across the dim top ditch at $\alpha = 0.07318$, with its lowest leading multiplier circled, and line B across the main ditch at $\log_{10} T = -0.153$. (b) Along line A the leading multiplier (solid) dips to 0.064 at $\log_{10} T = -0.15202$ but stays positive, and the next modulus (dashed) stays at most 0.021: no sign change, a valley of the same sign between the main ditch and the arc. (c) Along line B, the control, the leading multiplier changes sign through a complex leading pair (open marker), from -0.605 to $+0.710$. The orbit carries the seed's six-token loop and is attracting at every solved cell of all three measurements.

state of that stretch converges (residual 2×10^{-14}) to a saddle with a single multiplier outside the unit circle, $\mu_u = 1.2625$, real, with the next modulus 7.8×10^{-4}, at relative distances 3.5×10^{-3} and 4.7×10^{-3} from the two attracting orbits. A start a small distance δ off the saddle's stable manifold comes within a distance $\propto \delta^{\chi_s/(\chi_s+\chi_u)}$ of the saddle, as for a passage near a linear saddle with $\chi_u = \ln \mu_u$ and $\chi_s = -\ln |\mu_s|$ from its slowest stable multiplier μ_s (an exponent of 0.97 here), and then leaves along the unstable direction, its distance growing by μ_u per period, so the time it spends near the saddle grows by $6/\ln \mu_u = 25.74$ steps for every factor e by which the offset shrinks. Both predictions hold over eight decades (Figure 41; `tools/hook_saddle_exit.py`). At 17 temperatures on each side of the edge, a distance δ from 10^{-6} to 10^{-14} in $\log_{10} T$ away, the last step within 10^{-3} of the saddle grows by 25.78 and 25.75 steps per factor e in δ (all points within 4.0 and 2.9 steps of the fitted lines), the closest approach grows with δ with log-log slopes 0.986 and 0.995 (predicted 0.97), and every start below the edge ends on the strongly contracting orbit and every start above it on the other. Together with the random start, which removes the edge, this is strong numerical evidence that, at this steering strength, the lower edge is where the fixed start crosses the saddle's stable manifold. Bisecting 676 jumps of the outcome along the map's columns (to 2.6×10^{-13} in $\log_{10} T$; `tools/hook_edge_track.py`) traces the hook's other edges (Figure 41a). The upper edge follows the kite's fold curve, 0.1 to 1.6 grid rows above a spline through the continued curve's points and more so toward the cusp. The offset comes from the map's

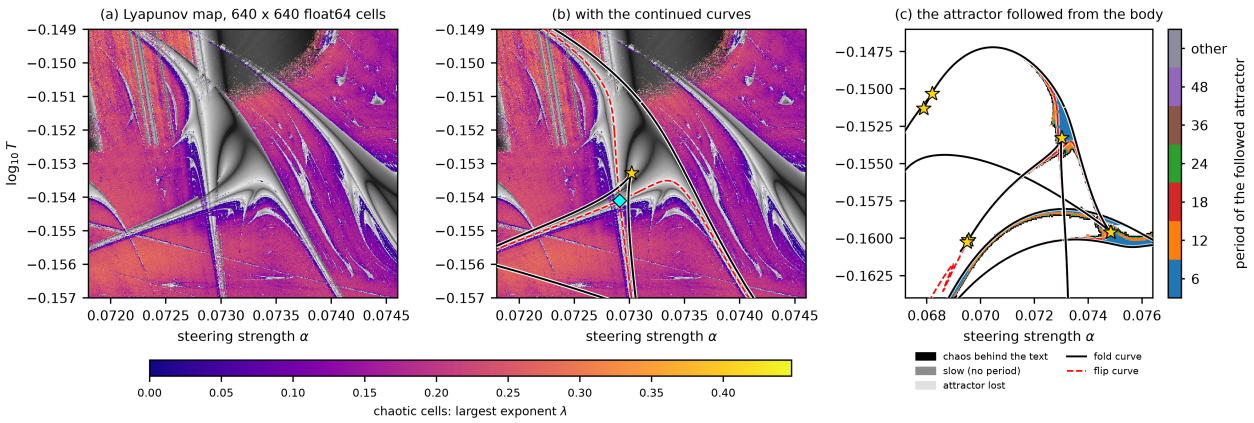

Figure 38: The period-6 shrimp, in the plane of Figure 11 (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12; float64). (a) A 640 × 640 map of the window, centered on the body, colored as Figure 11. (b) The same map with the continued curves of the period-6 orbit: fold curves black, flip curves red and dashed; cyan diamond: the crossing of the two flip curves, where two distinct period-6 orbits are found (a pseudo-crossing); gold star: the kite's cusp C. (c) The attractor followed from the body, on the grid of Figure 11 and over a wider window, colored by its period (black: chaos behind the repeating text, an exactly repeating text with no period up to 64 and an exponent of at least 0.005; gray: slow cells, the same with an exponent below 0.005), with the same curves; gold stars: the six located cusps (Table 20): the kite's cusp C, pairs near the ends of the upper-left and lower-left legs (the lower pair's stars overlap at this scale) and the tip of the small window beyond the neck. The fold curve through C is that of the kite (Figure 39).

window (Figure 42b,c; `tools/hook_tip_strip.py`): half a row and one row beyond the fold, where the weaker orbit does not exist, the exponent over the map's window (steps 300 to 600) lies between -0.022 and $+0.012$ for α above 0.072965 and 0.072967, because the loop still lingers at the orbit that vanished there, while over steps 3000 to 3400 every cell of the strip beyond the fold gives the strongly contracting orbit's exponent (-0.18 to -0.71).

The tip of the hook. The two edges meet where the fixed start lands exactly on the saddle-node orbit. At a point of the fold, the coordinate u of the merged orbit along its critical direction changes over one six-step return as $u' = u + \eta + bu^2$ to leading order, with η the distance from the fold in the normal form's units and $\eta = 0$ on the fold (Appendix E; [50]); after its fast transient the fixed start lands at $u_0 = \ell^\top(x - p_0)/\ell^\top r$, with r and ℓ the right and left eigenvectors of the multiplier $+1$ at the orbit point p_0. Where u_0 has the sign that the quadratic term drives toward zero, the start converges onto the saddle-node orbit, and just inside the coexistence region it reaches the weaker orbit; where u_0 has the other sign, the start escapes to the strongly contracting orbit. To leading order the tip is therefore the point of the fold where u_0 changes sign (exactly, it is where the start lands on the strong-stable set of the saddle-node orbit, which the zero of u_0 approximates near that orbit). At seven points of the continued fold curve, u_0 is negative at $\alpha = 0.072980$ to 0.073002, where the start converges slowly onto the orbit (its drift fits $u' = u + bu^2$ with $b = 56$ to 68), and positive at $\alpha = 0.072889$ to 0.072963, where it escapes (`tools/hook_tip_check.py`). At eight fold points located between them by Newton's method at fixed α (residuals at most 5×10^{-14}; `tools/hook_tip_fold_points.py`), u_0 is nearly linear in α and crosses zero at $\alpha = 0.0729655$, the predicted tip to leading order (Figure 42a). Inside the fold, the boundary between the two outcomes approaches this point: at offsets of 2, 1, 0.5, 0.25 and 0.125 grid rows below the fold, the last cell that ends on the strongly contracting orbit lies at $\alpha = 0.0729627$, 0.0729632, 0.0729637, 0.0729642 and 0.0729645 (long window, cells 10^{-7} apart), and extrapolations with exponents from 1/3 to

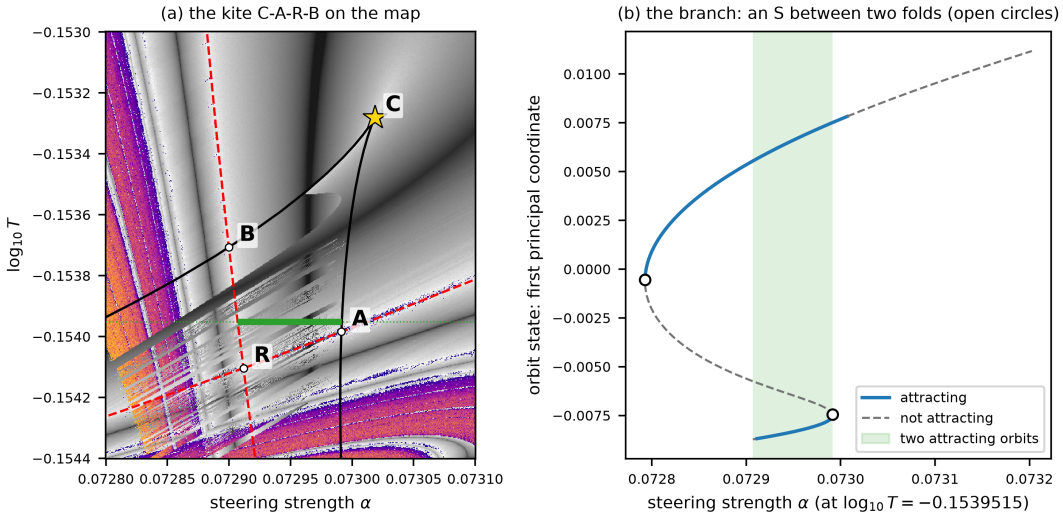

Figure 39: The kite of the period-6 shrimp (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12; float64). (a) A 512×512 float64 map of the lower-left end of the body, with cells 7.5 times finer in α and 4.6 times finer in $\log_{10} T$ than in Figure 38a (gray: periodic, darker where more contracting; colors: chaotic), with the body's two flip curves (red, dashed), the fold curve continued from the branch of (b) (black), and the corners of Gallas's kite: R, where the flip curves cross (two distinct orbits, one on each curve); A and B, where the fold curve meets them; C, the located cusp where the fold curve turns back (gold star). Green: the line $\log_{10} T = -0.15395$ of the branch, thick where the two attracting orbits coexist. (b) The branch, computed by numerical continuation of the period-6 orbit in α along the green line of (a): every point is an orbit converged by Newton's method, solid where its multipliers show it attracting and dashed where not, and the two open circles are folds, where a multiplier reaches $+1$; they make the branch an S, and the two attracting sheets overlap in the shaded interval (`tools/plot_p6_kite.py`). The vertical axis is a display coordinate: one point of each orbit, a vector of 896 numbers, centered over the branch and projected on the direction in which these points vary most, their first principal component. All these orbits write the same six tokens; the coordinate follows the probability of "is" after "The existence", from 9.9% to 12.7% (run FW). Inside the kite, the straight edge where the gray changes abruptly (near $\log_{10} T = -0.15353$) is a basin boundary of the fixed start, not a bifurcation, where we checked it: at two neighboring cells across it both attracting period-6 orbits exist (leading multipliers 0.18 and 0.60 to 0.61), the fixed start reaches one above the edge and the other below (`tools/kite_basin_check.py`), and at one steering strength, by strong numerical evidence, the start on the edge lies on the stable manifold of a saddle orbit between them (Figure 41).

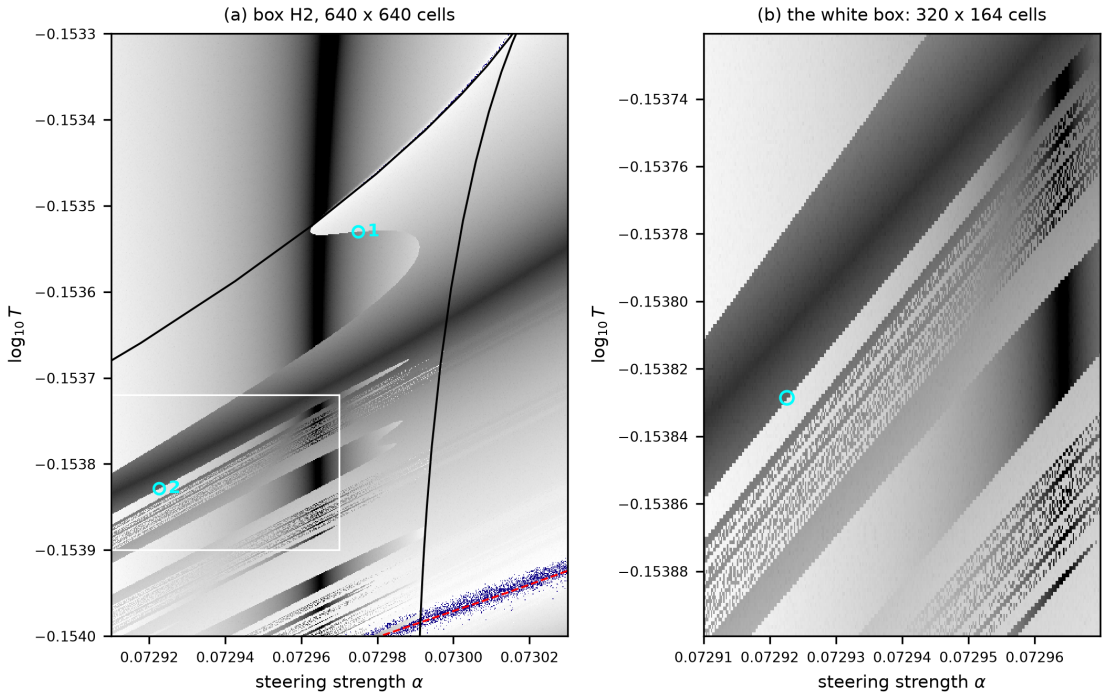

Figure 40: The hook and the striations at the upper end of the kite (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12; float64). (a) A 640×640 map of $\alpha \in [0.07291, 0.07303]$ and $\log_{10} T \in [-0.1540, -0.1533]$ (gray: periodic, darker where more contracting; colors: chaotic), with the kite's fold curve (black) and a flip curve of the body (red, dashed); cyan: the two pairs of adjacent cells checked by Newton continuation, on the wedge's lower edge (1) and on a tongue's edge (2). (b) The white box in (a) at the map's own cells: tongues in which the outcome alternates from cell to cell.

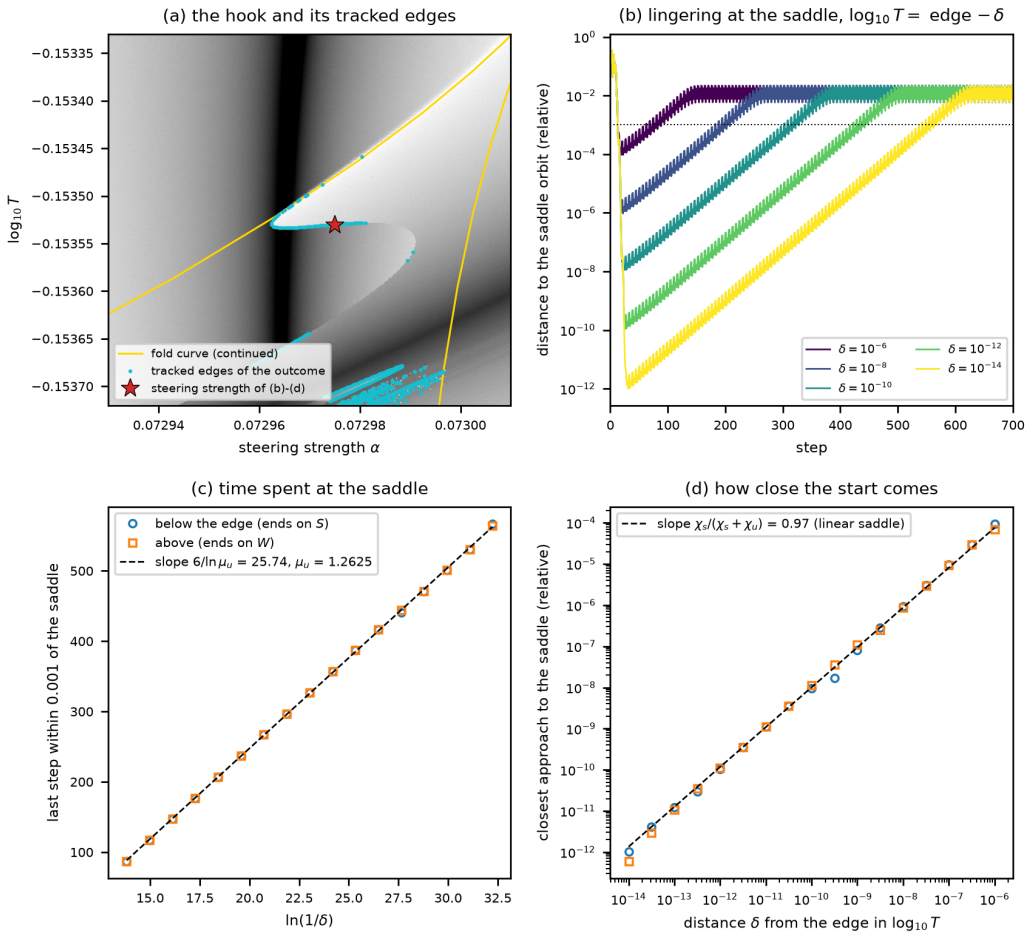

Figure 41: The hook's lower edge is where the fixed start crosses the stable manifold of a saddle orbit (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12; float64). (a) The map of Figure 40a near the hook, with the kite's continued fold curve (gold) and the edges of the fixed start's outcome bisected along the map's columns (cyan; brackets whose end exponents differ by at least 0.1); red star: the steering strength $\alpha = 0.072975$ of (b)–(d). (b) Distance to the saddle orbit along the loop started a distance δ in $\log_{10} T$ below the edge: it arrives along the stable manifold within about 25 steps, lingers, and leaves along the unstable direction; dotted: the radius 10^{-3} of (c). (c) The last step within 10^{-3} of the saddle against $\ln(1/\delta)$, for starts below the edge (which end on the strongly contracting orbit S) and above it (which end on the other, W); dashed: slope $6/\ln \mu_u$ from the saddle's unstable multiplier μ_u. (d) Closest approach to the saddle against δ; dashed: slope $\chi_s/(\chi_s + \chi_u) = 0.97$ from the saddle's multipliers.

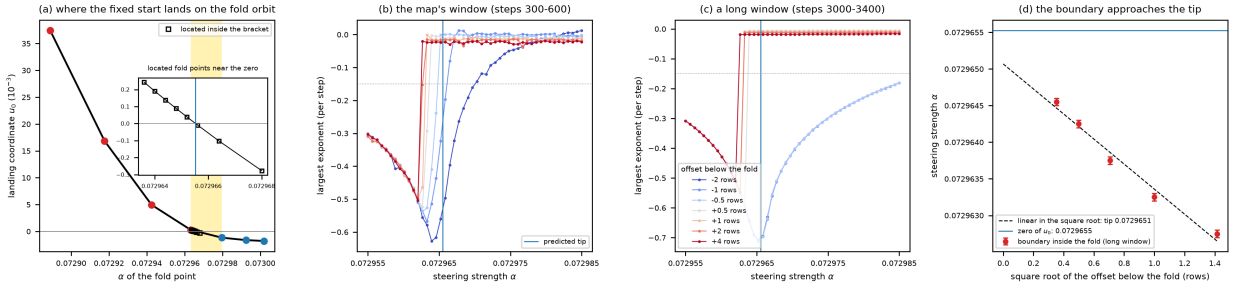

Figure 42: The tip of the hook (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12; float64). (a) The fixed start's landing coordinate u_0 on the saddle-node orbit at points of the kite's fold curve: negative (blue) where the start converges onto that orbit, positive (red) where it escapes; squares and inset: fold points located by Newton's method at fixed α, where u_0 crosses zero at $\alpha = 0.0729655$ (vertical line, the predicted tip, also in (b) and (c)). (b) The largest exponent over the map's window (steps 300 to 600) along rows at fixed offsets from the fold curve (negative: beyond the fold, where the weaker orbit does not exist); the strongly contracting orbit has exponents of -0.18 to -0.71 here, the weaker one exponents near -0.01; dotted: the threshold -0.15 between them. (c) The same cells over steps 3000 to 3400: beyond the fold every cell ends on the strongly contracting orbit. (d) The boundary between the two outcomes inside the fold (last cell on the strongly contracting side and first on the weaker, long window) against the square root of the offset, with a straight-line extrapolation to the fold (which assumes the exponent 1/2; the points do not resolve it).

1/2 put the boundary's end on the fold at $\alpha = 0.0729651$ to 0.0729657, bracketing the prediction (Figure 42d). The normal form predicts a tangential contact, the gap between the edges growing quadratically along the fold, which corresponds to the exponent 1/2; these points do not resolve the exponent. In the map's window the boundary appears up to 7×10^{-7} farther from the tip, as expected where the saddle's unstable multiplier approaches 1 near the fold and the start lingers longer.

Symbolic dynamics. Jones's unimodal interpretation of selected Rössler return maps, which traced the shape of a periodicity hub to a conjugacy between the Rössler system and unimodal maps in part of its parameter plane [65], suggests an approximate scalar coding of the six-band attractor. Where the attractor writes the body's six-token text, the loop visits six thin bands of the state space in turn, one per token, and returns to each band every six steps. Many sampled six-step return maps, plotted along a band's first principal axis, show a single turning point: the next visit traces a parabola against the previous one (Figure 43). We code each visit by the branch it lies on, L on the increasing one and R on the decreasing one. The branch is read from the sign of the six-step derivative along the band, which the twin trajectory of the Lyapunov exponent provides (`tools/symbolic_dynamics.py`): the twin's offset direction is carried along the attractor, and its scalar product with its own direction six steps earlier has the sign of that derivative where the offset is aligned with the band and the band turns by less than a right angle. The checks below support this approximation; they do not prove it. Pooling the misplaced visits and the rarer symbol's visits over the phases where the band is straight enough for the sign to be defined, 48 of the 58 chaotic six-band cells of a 48×48 map of the window have an error ratio below 5 % (the largest passing ratio is 4.7 %, 23 cells have no misplaced visit, and a single phase of a passing cell reaches 21.7 %). The other 10 cells fail the coding test: three in one column beside the body's left flip curve and seven below the kite. In 7 of them the return map has more than one resolved turning point: it rises and falls more than once (Figure 43h–k). For a C^3 unimodal interval map with negative

Schwarzian derivative, at most one attracting periodic orbit can have an immediate basin that avoids the interval's endpoints [47]; the extra turning points below the kite are consistent with a scalar explanation of the coexistence inside it, but they establish neither a global scalar reduction nor its Schwarzian sign.

A periodic orbit of period $6k$ visits each band k times. Its signs, rotated to the largest sequence in the order of kneading theory, give it a label W|X in the sense of Metropolis, Stein and Stein (MSS) [75, 76]: W is matched to the quadratic family's superstable itinerary WC and names the window, and X records the orbit's inferred last branch symbol. Away from a superstable parameter this is an orbit label, not a measured itinerary of the critical point. After a majority vote over the band phases whose coding is usable (at least three) and an explicit filter that keeps only MSS words, 44,577 periodic six-band cells on the 640×640 grid of Figure 38a receive 23 distinct window words W (Figure 44); the filter rejects 240 further cells whose labels are not MSS words. The body and its period doublings carry the words of the doubling cascade: the empty word (period 6), R (12), RLR (24) and RLRRRLR (48). Windows in the surrounding chaos carry other words, among them RL (period 18), RLRR (30), and RLRRR and RLLRL, two distinct windows of period 36 that only their words tell apart. Within a window, the label predicts the sign of the orbit's leading multiplier, $(-1)^{n_R}$ with n_R the number of R's in W|X: positive next to the boundary where the orbit is born (a fold, or for a doubled orbit the flip curve of its parent) and negative next to its own flip. The curve where the predicted sign changes is a candidate for the window's superstable curve. The scan assigns C to 153 cells whose usable phases do not agree sufficiently on the last symbol; at the two selected C cells, Newton calculations report small, nonzero complex leading pairs. These labels mark ambiguous scalar coding, not located superstable curves. Float64 Newton solves with exact Jacobians at one cell for each label W|X with at least three coded cells in the half of the scan then complete (the cell nearest the center of the label's largest patch: 23 cells, covering 12 of the 23 words; Table 19), scored by a subsequent clean-phase rule (phases whose L and R visits one threshold separates along the band, with every $|\cos| \geq 0.5$), reproduce the scan's label at 16 cells; the two C cells have complex leading pairs, and at the remaining 5 cells, just below the kite, no phase is clean, so their four words (RLL on both sides, RLLLRLL, RLLLRLR and RLLLRRR) remain unconfirmed; the other 11 words were not checked. At the 16 confirmed cells the full map's leading multiplier also has the sign that the label's parity predicts; this sign rule belongs to the scalar coding and carries over to the full map only while that multiplier stays real and dominant. These checks support selected labels, not the whole map. In total, 1,469 of the 46,046 eligible periodic six-band cells receive no word (3.2 %), the 240 rejected by the filter among them; 922 of them (63 %) lie within Euclidean distance two grid cells of the eligible domain's edge.

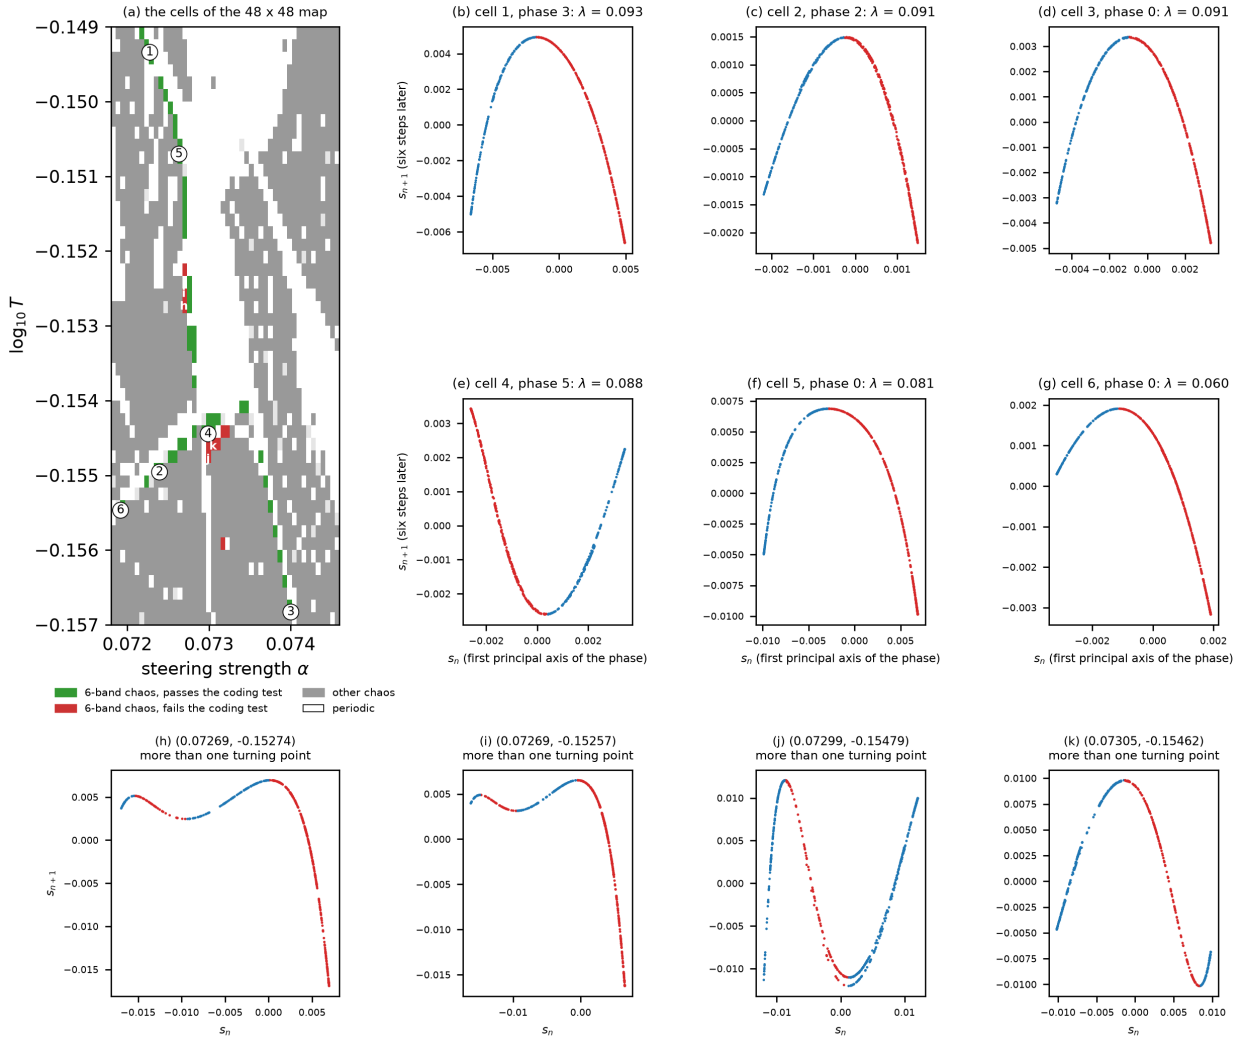

Figure 43: Return maps of the six-band chaos around the period-6 shrimp (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, float64; 1000 transient and 3000 recorded steps per cell). (a) The cells of a 48×48 map of the window of Figure 38a: chaotic cells writing the six-token text that pass the pooled coding test (green: misplaced visits below 5 % of the rarer symbol's, pooled over the usable phases) or fail it (red), other cells with an exponent above 10^{-3} (gray) and periodic cells (white). (b–g) Six-step return maps at six green cells: the position along one band (its first principal axis) against the position at the previous visit; blue: visits where the six-step derivative is positive (L), red: negative (R). (h–k) Return maps at four red cells (letters in (a)) with more than one turning point.

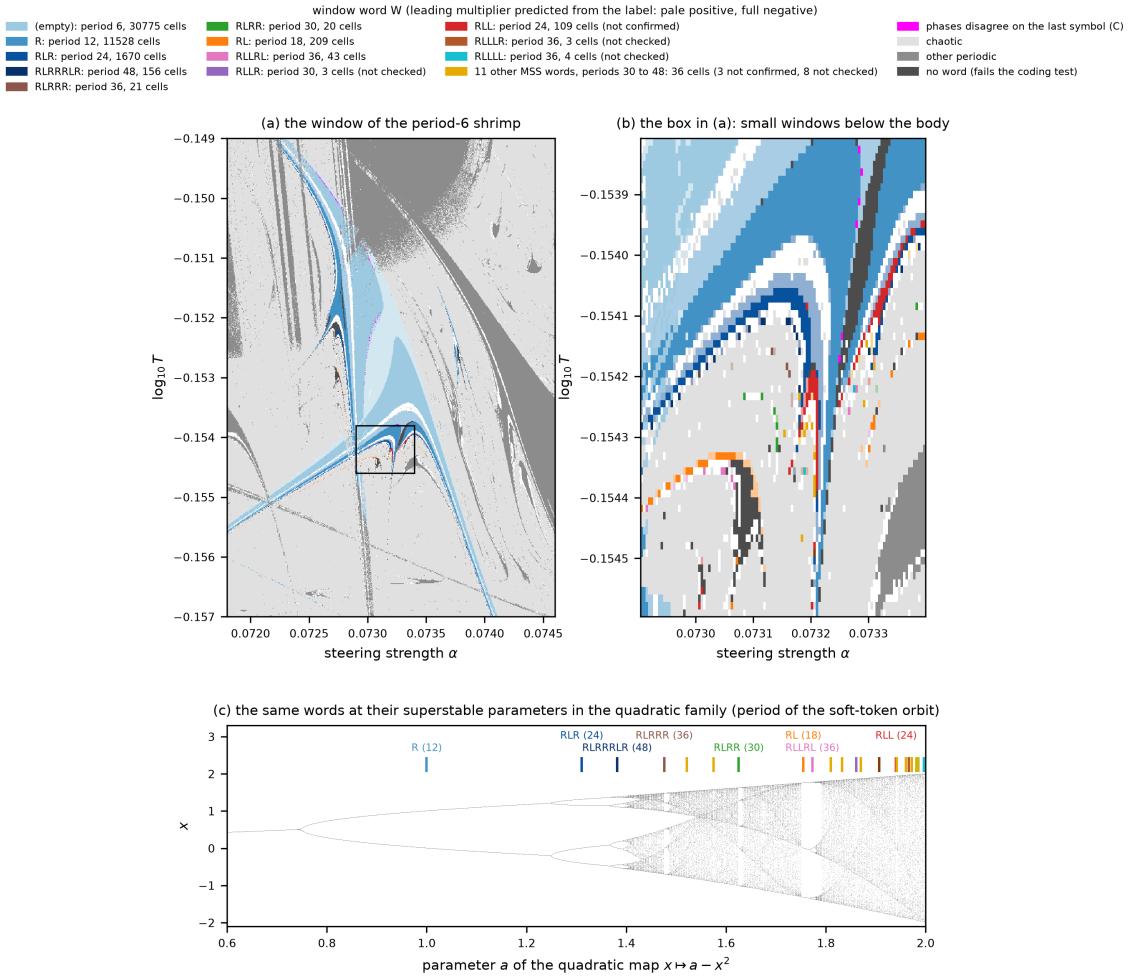

Figure 44: Symbolic words of the period-6 shrimp's windows (640×640 float64 cells, the grid of Figure 38a). (a) Each periodic cell of the six-band region colored by the word of its window: pale where the label predicts a positive leading multiplier, full where it predicts a negative one, magenta where the phases disagree on the last symbol (C); light gray: chaos; mid gray: other periodic orbits; dark gray: no word (fails the coding test). Windows marked "not confirmed" in the legend are those of Table 19's cells without a clean phase, and those marked "not checked" had no float64 check. (b) The box in (a). (c) The same words at the superstable parameters of the quadratic family $x \mapsto a - x^2$, on its bifurcation diagram (the body's empty word is that of its fixed point, at $a = 0$); labels give the period of the soft-token orbit, six times the word's period.

Table 19: Float64 checks of the symbolic words (`tools/sd_verify_words.py`): Newton with exact Jacobians at the center of the largest patch of each window word and side of Figure 44 with at least three cells in the eight of the scan's 16 bands then complete (23 cells covering 12 of the 23 words; the other 11 were not checked). Word and side: the scan's label; μ_1: the orbit's leading multiplier; clean: the phases whose L and R visits one threshold separates along the band, with every $|\cos| \geq 0.5$; result: the clean phases give the scan's word (confirmed), no phase is clean, or the leading pair is complex (side C: the scan's phases disagree on the last symbol).

word	side	period	$(\alpha, \log_{10} T)$	μ_1	clean	result
(empty)	C	6	$(0.0730707, -0.1520673)$	$-0.002 + 0.024i$	0	complex pair
(empty)	L	6	$(0.0731584, -0.1527559)$	$+0.189$	6	confirmed
(empty)	R	6	$(0.0732285, -0.1533944)$	-0.686	6	confirmed
R	C	12	$(0.0732548, -0.1537825)$	$-0.001 + 0.004i$	0	complex pair
R	L	12	$(0.0731628, -0.1539703)$	-0.770	5	confirmed
R	R	12	$(0.0733030, -0.1538200)$	$+0.078$	5	confirmed
RL	L	18	$(0.0730488, -0.1543333)$	-0.184	3	confirmed
RL	R	18	$(0.0727465, -0.1528936)$	$+0.266$	2	confirmed
RLL	L	24	$(0.0732110, -0.1542582)$	-0.481	0	*no clean phase*
RLL	R	24	$(0.0731847, -0.1542707)$	$+0.278$	0	*no clean phase*
RLR	L	24	$(0.0731540, -0.1540454)$	$+0.178$	4	confirmed
RLR	R	24	$(0.0727772, -0.1524053)$	-0.526	3	confirmed
RLRR	L	30	$(0.0731408, -0.1543208)$	-0.051	3	confirmed
RLRR	R	30	$(0.0731408, -0.1542707)$	$+0.139$	4	confirmed
RLLRL	L	36	$(0.0740260, -0.1569249)$	$+0.168$	4	confirmed
RLLRL	R	36	$(0.0730444, -0.1543584)$	-0.439	3	confirmed
RLRRR	L	36	$(0.0731671, -0.1542457)$	$+0.146$	3	confirmed
RLRRR	R	36	$(0.0731628, -0.1541956)$	-0.063	3	confirmed
RLLLRLL	L	48	$(0.0731803, -0.1543083)$	$+0.169$	0	*no clean phase*
RLLLRLR	L	48	$(0.0732066, -0.1543834)$	-0.251	0	*no clean phase*
RLLLRRR	L	48	$(0.0733074, -0.1541581)$	$+0.035$	0	*no clean phase*
RLRRRLR	L	48	$(0.0731146, -0.1541080)$	-0.085	4	confirmed
RLRRRLR	R	48	$(0.0731102, -0.1541080)$	$+0.226$	5	confirmed

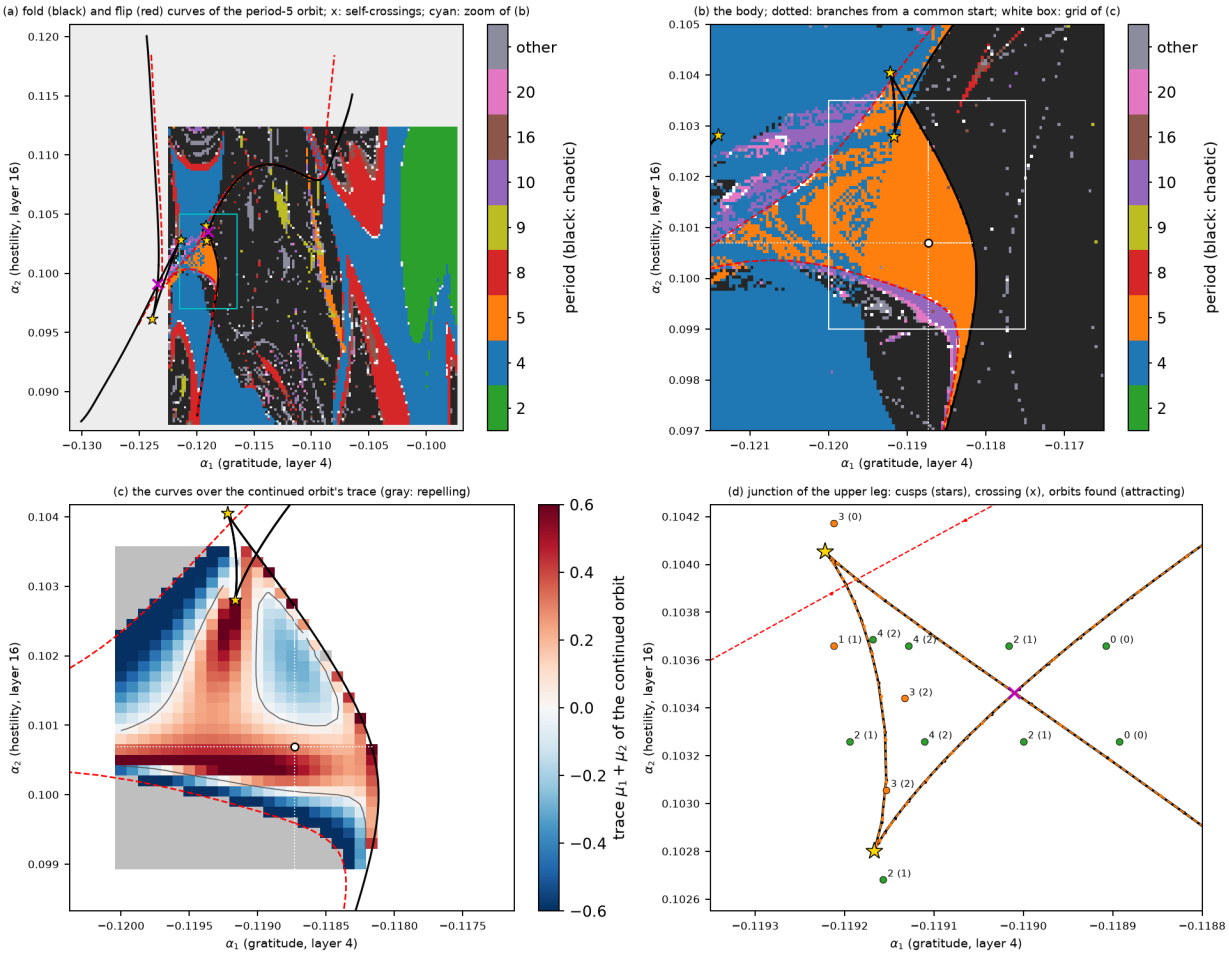

Figure 45: The skeleton of the period-5 window in the plane steered at two depths (Qwen2.5-0.5B-Instruct, gratitude at block 4 and hostility at block 16, $\log_{10} T = -0.1$, float64). (a) The two fold curves (black; multiplier $+1$) and the flip curves (red, dashed; multiplier -1) of the window's period-5 orbit, over the attractor reached from the fixed start of every cell (Appendix B) on a wider grid (gray: outside it); gold stars: the four located cusps; magenta crosses: the self-crossings of the two fold curves; cyan: the zoom of (b). (b) The body over a finer grid (as Figure 74a), with the four branches (dotted) that found the first boundary points from a common start (circle); white: the continuation grid of (c). (c) The same curves over the trace of the continued orbit's two leading multipliers (Figure 74b; gray: repelling). (d) The junction of the upper leg with the body: the fold curve has two cusps that share a segment and crosses itself (x), where its arms are folds of two different period-5 orbits; its two traces, continued from two different boundary points (black, and orange dashed), coincide. Circles: the 13 points of the orbit census, with the number of period-5 orbits found and, in brackets, the number attracting (green: the count predicted before the computation; orange: not).

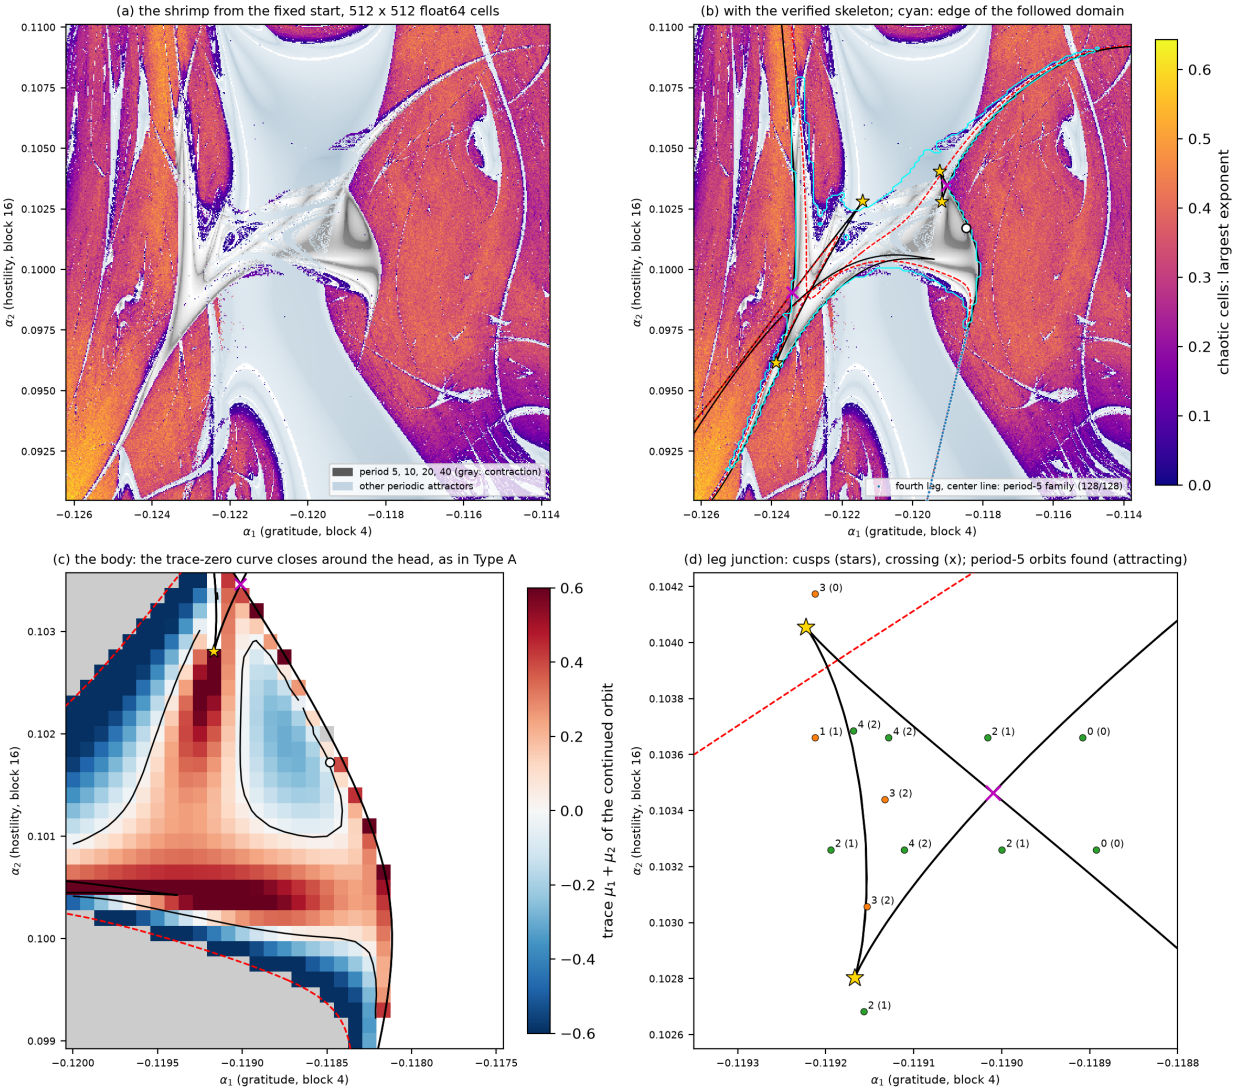

Figure 46: The period-5 shrimp (Qwen2.5-0.5B-Instruct, gratitude at block 4 and hostility at block 16, $\log_{10} T = -0.1$, float64). (a) A 512×512 map of the window from the fixed start (run HP5; steps 2.4×10^{-5} in α_1 and 3.8×10^{-5} in α_2): cells whose period is a multiple of five, the window's family, in gray by contraction; other periodic attractors (among them the coexisting period-4 attractor that the fixed start reaches in parts of the legs) pale; cells whose largest exponent exceeds 10^{-3} by its value (this overrides a detected period at 33 cells of the followed domain). (b) The same map with the window's verified skeleton: the fold curves (black; multiplier $+1$) and flip curves (red, dashed; -1) of the period-5 orbit, including the lower-left leg's own pair (Figure 72), the located cusps (gold stars), the self-crossings of the fold curves (magenta) and the near-nilpotent point (white circle). Cyan: the edge of the domain of the window's attractor followed cell by cell from the body on a grid of spacing 1.25×10^{-4}. Blue dots: the center line of the fourth (lower-right) leg, too thin for the grid; all 128 sampled points retain the followed family, 127 with period 5 and one with period 20. (c) The body over the trace of the continued orbit's two leading multipliers (gray: repelling): the trace-zero curve (black) closes around an island of negative trace, the geometry of a Type A head, and passes the near-nilpotent point. (d) The junction of the upper leg: two cusps that share a segment of the fold curve and the crossing of its other branches, with the number of period-5 orbits found (attracting) at 13 points fixed before the computation (green: the count predicted by a section of a swallowtail; orange: not).

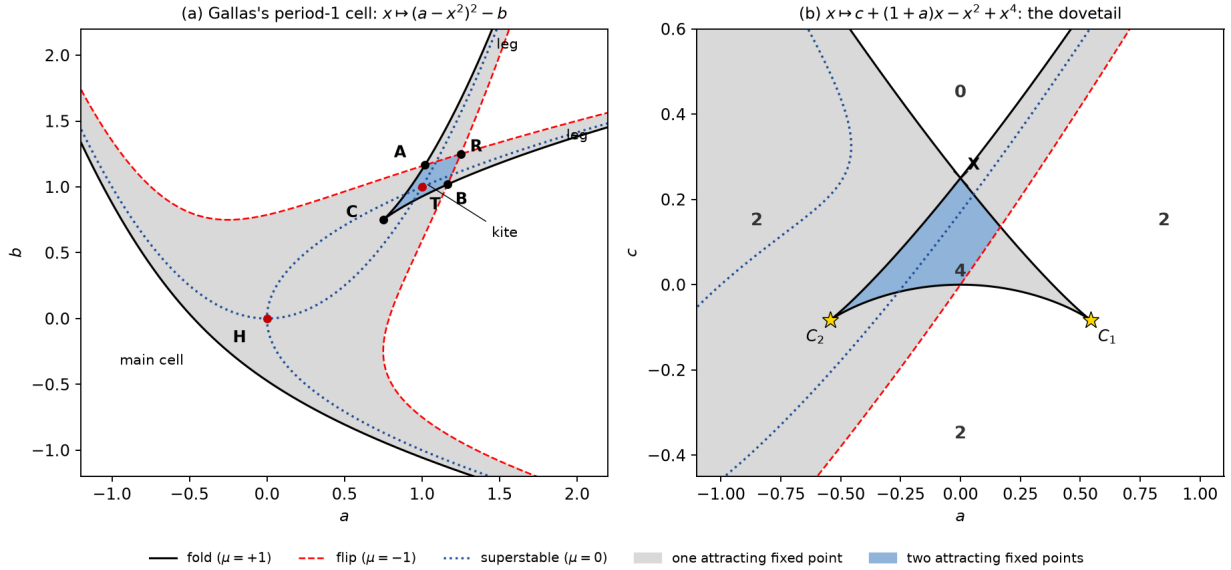

Figure 47: The anatomy of a shrimp's cell in two classical one-dimensional models, computed (gray: one attracting fixed point; blue: two). (a) Gallas's canonical period-1 cell, $x \mapsto (a - x^2)^2 - b$ (his eq. (5.3), rotated and labeled as in his Figure 7 [66]). Fold curves (black; multiplier $+1$) and flip curves (red, dashed; -1) bound the main cell, and the superstable curves (blue, dotted; 0) of the map's critical points cross at the head H. Two legs, each between a fold and a flip curve, cross in the kite C–A–R–B, which begins at the cusp C of the fold curve; in the kite two attracting fixed points coexist, and their superstable curves cross at the tail T. The flip curves of the two legs cross at R, which makes the kite a crossroad area in the sense of Carcassès et al. The computed corners A and B agree with Gallas's to six significant digits. (b) The quartic model of Carcassès et al. [71], $x \mapsto c + (1 + a)x + bx^2 + x^4$, in the plane of a and c at $b = -1$, a plane they did not draw: the fold curve has two cusps C_1 and C_2 that share a segment and crosses itself at X (a dovetail, in the term Mira et al. use for critical curves with two such cusps [73]). Numbers: real fixed points in each region; a flip curve crosses the pocket, and two fixed points attract only to its left. Both fixed-point equations are quartics $x^4 + ux^2 + vx + w = 0$, whose fold set is the swallowtail: (a) is a section at fixed v, with one cusp, (b) a section at fixed $u < 0$, with two cusps and a crossing. A two-dimensional map whose Jacobian determinant never vanishes, such as the Hénon map, has no superstable curves, since no multiplier of its periodic orbits can vanish; trace-zero curves take their place and either close around the head (Type A) or approach each other and turn away (Type B) [67]; the language model's period-5 window has the first geometry (Figure 46c).

Table 20: The six located cusps of the period-6 window (float64), in the plane of α and $\log_{10} T$ (C: the cusp of the kite of Figure 39; upper and lower 1, 2: pairs near the ends of the upper-left and lower-left legs; neck: at the tip of the small window beyond the neck): position, cubic coefficient c and unfolding determinant. At each cusp the oriented coefficient $b \operatorname{sign} a$, of magnitude between 0.17 and 39.11 at the twelve bracketing curve points, falls to at most 9.5×10^{-5} in modulus, the critical multiplier equals $+1$ to within 8.2×10^{-9}, and the next has modulus at most 0.74 (located by `tools/cusp_locate.py`; records `p6s/locate-*` in the data store, with checksums in the research log).

cusp	$(\alpha, \log_{10} T)$	c	$\det K$
C	$(0.0730186, -0.1532803)$	-1.33×10^{4}	-2.62×10^{3}
upper 1	$(0.0678765, -0.1513104)$	$+6.10 \times 10^{3}$	4.59×10^{4}
upper 2	$(0.0682059, -0.1503283)$	-8.01×10^{3}	2.20×10^{5}
lower 1	$(0.0694974, -0.1602656)$	-2.25×10^{3}	2.09×10^{4}
lower 2	$(0.0695513, -0.1601114)$	$+2.01 \times 10^{3}$	2.20×10^{4}
neck	$(0.0748392, -0.1595956)$	-4.02×10^{4}	4.47×10^{3}

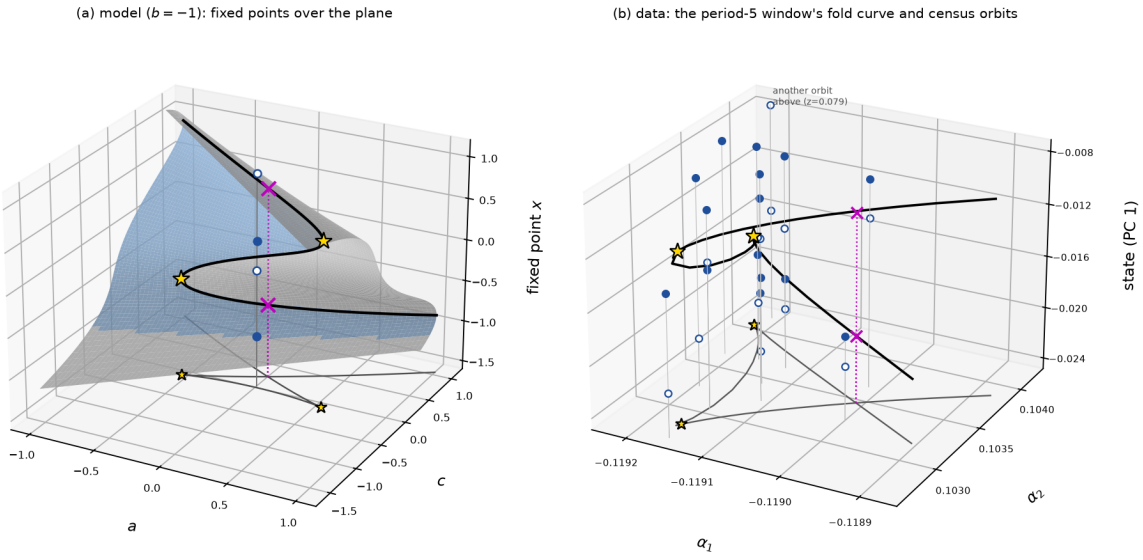

Figure 48: The dovetail in Mira's foliated representation: parameters and one state coordinate. (a) Model: the fixed points x of the quartic map at $b = -1$ over the plane of a and c (blue: attracting; gray: not). The fold curve (black) is a smooth curve on this surface, and its shadow on the parameter plane is the dovetail of Figure 47b, with its two cusps (stars) and its crossing X, whose two lifts (crosses) lie on different sheets. Over a point of the pocket the surface has four sheets (circles; filled: attracting). (b) Data: the period-5 window of Figure 45 (Qwen2.5-0.5B-Instruct, gratitude at block 4 and hostility at block 16, $\log_{10} T = -0.1$, float64). The continued fold curve through the junction (black), lifted by the first principal coordinate of these states (86 % of their variance), with its shadow (gray), its two located cusps (stars), the two lifts of the crossing (crosses), and every period-5 orbit the census found (filled: attracting; open: repelling), each at the phase continuous with the fold curve. One orbit with another token loop lies above the frame.

11.1 Retained feedback as a third control: a trace-saddle sign change

The separation of the period-6 body's lines raises a question that a finer picture alone cannot answer: is the separation removable by changing the map? We tested this in the smaller period-18 window below that body, called the child to distinguish it from the period-6 mother. These are different full-state orbit families. The child retains a primitive 18-step state cycle while writing a six-token cycle three times; equality of the text after six steps is not closure of its 896-dimensional feedback state. In this exploratory intervention, retaining a small fraction of the previous feedback state changes the sign of the selected pair's trace-saddle height, while the sampled centers keep that period and readout. The result is a controlled change of spectral geometry within one periodic family. Under the local smooth-saddle assumptions below, the sign change implies a reconnection of the trace-zero contours; the sampled measurements do not trace their complete connections away from the saddle. It is not a transition to different text or a demonstration of exact superstability.

What memory means here. Write F_p for the original one-token soft-feedback map at $p = (\alpha, \log_{10} T)$. We replace its state update by

$$x_{n+1} = F_{p,\beta}(x_n) = (1-\beta)F_p(x_n) + \beta x_n. \tag{10}$$

The fresh expected embedding is mixed with the current feedback embedding, with the same coefficient at every step. At $\beta = 0$ this is exactly the original map; at $\beta = 0.00645$, for example, the next state contains 99.355 % fresh output and 0.645 % retained state. This adds inertia to the state update, not a longer sequence of visible tokens. The context window remains $K = 1$, the prompt and token position remain fixed, and neither weights nor steering direction are retrained. The readout is still the model's next-token distribution evaluated at x_n; it is the embedding fed back to the next step that is mixed. The longer-context experiments of Section 13 change a different object.

Following the saddle, rather than taking a fixed slice. Let $M(p, \beta)$ be the Jacobian of the 18-step return along the continued child orbit, and let

$$\tau(p, \beta) = \mu_1 + \mu_2, \qquad \Delta(p, \beta) = \mu_1 \mu_2 \tag{11}$$

be the sum and product of the selected real or complex-conjugate pair. These are pair invariants, not the trace and determinant of the full return matrix. Selecting an isolated pair keeps τ smooth when the two multipliers exchange order by modulus. At each β we correct the two original controls to locate the same trace saddle,

$$\nabla_p \tau(c(\beta), \beta) = 0, \qquad h(\beta) = \tau(c(\beta), \beta), \qquad H(\beta) = \partial_p^2 \tau(c(\beta), \beta). \tag{12}$$

The saddle is selected by its stationarity and indefinite Hessian; its height is then measured, not forced to zero by that two-control solve. Holding α and T at their original values would confound a change of height with motion of the saddle. Only after a sign bracket was found did a third scalar solve adjust β to reduce the height. Thus the reported crossing estimate concerns three adjusted controls, not addition of memory at fixed steering and temperature.

The same full-state primitive-period, closure, pair-residual and cyclic-phase gates were applied throughout. The first continuation accepted eleven saddle centers through $\beta = 0.003$, lowering h from $4.647042445 \times 10^{-4}$ to $2.753879880 \times 10^{-4}$. Its next step, to $\beta = 0.005$, failed because the proposed corrector displacement, 2.19189 in the declared scaled coordinates, exceeded the unchanged trust radius of 2. That partial attempt remains in the record; its failed point was

not accepted. A disclosed local extension freshly replayed the accepted $\beta = 0.001$ and 0.003 centers, then used smaller continuation steps. This extension completed 273 qualified evaluations and seven new saddle centers. Its saved final-return states allow independent closure replay at every point; the largest relative closure residual is below 4.34×10^{-13} and the smallest tested lower-lag return exceeds 0.01494. The historical attempt did not save its final-return arrays, so its closure check retains the recorded-summary limitation. Source commits, plans, every attempt and both audits are bound by `docs/baby-shrimp-zoom/memory-continuation-receipt.json` and `docs/baby-shrimp-zoom/memory-local-extension-receipt.json`.

A resolved sign change and a residual-qualified crossing estimate. Table 21 reports the accepted saddle heights. The values at $\beta = 0.006$ and 0.007 have opposite signs, separated from zero by much more than the observed changes under stencil scale, center transport and cyclic phase (Figure 49). The root search uses those measured endpoints to locate a candidate: a secant through their heights proposes an intermediate β, with bisection when the interpolation fails the declared safeguard. The endpoint centers provide only an initial guess for the new center. At each proposed memory setting, the solver again adjusts steering and temperature until the trace is stationary, and only then evaluates the corrected saddle height. A measured height outside the residual tolerance replaces the bracket endpoint of the same sign, and the process repeats. This is data-guided numerical root location after a sign change was observed, not an advance prediction of the connection parameter. The procedure then satisfied its declared $|h| \le 10^{-8}$ stopping rule at

$$\begin{aligned} \beta_* &= 0.006448531026086656, \\ (\alpha_*, \log_{10} T_*) &= (0.07331776448598727, -0.15804614858293237), \\ \tau_* &= 1.6365300102 \times 10^{-10}. \end{aligned} \tag{13}$$

The memory fraction is approximately 0.64485 %. The digits in Eq. (13) identify the computed point; they are not a parameter error bar. Its trace remains positive across the recorded estimates, and the solver's retained sign bracket is $[0.0064484162589609016, 0.007]$, not a narrow certified interval around the candidate. The distinction matters: a well-resolved sign change establishes numerical evidence for a zero in a smooth continuation, while a small residual specifies how accurately this particular candidate solves the equations. Neither is an interval proof of an exact root or its uniqueness.

The continuation uses scaled coordinates $u = (p - p_0)/(2.5 \times 10^{-6}, 7.5 \times 10^{-6})$, componentwise, about the saved anchor $p_0 = (0.07309363265341377, -0.1544023817786725)$. At the final candidate the finite-difference trace-gradient norm is 8.22×10^{-11} in these coordinates, and the Hessian eigenvalues are -0.0426123 and $+0.0133081$. The Hessian remains indefinite. The transported coarse-to-fine gradient discrepancy is about 2.60×10^{-7}, so the much smaller solver residual alone is not a derivative-accuracy estimate. The empirical height sensitivity is 8.78×10^{-12} or less; it does not bound unmeasured higher derivatives or operator errors. The corrected-height secant over $\beta = 0.006$ to 0.007 is -0.0894594. The last two positive root iterates give -0.0892008, with an empirical endpoint-sensitivity scale of about 1.01×10^{-4} after division by their separation. These finite secants support a transverse unfolding numerically; they do not certify $h'(\beta_*)$ at an exact root.

Full-state alignment provides an additional check that the continuation stays with the child. Across all 18 accepted historical and local centers, all 17 neighboring comparisons prefer the same cyclic alignment and retain the same six-token readout. The largest relative change between neighboring full 18-by-896 state arrays is below 0.00560, whereas the smallest mismatch for a second-best cyclic alignment exceeds 0.10712. This is evidence beyond token equality, which would not distinguish the child from a six-step orbit. It supports the sampled branch continuity, but does

Table 21: The moving trace saddle of the period-18 child under retained feedback (BSC3). Entries give the accepted center's trace and the maximum observed height sensitivity from the saved comparisons, which is an empirical diagnostic, not a rigorous uncertainty bound. The two final stencil centers differ, so the audit uses quadratic transport and does not apply direct Richardson extrapolation. The intermediate root estimate at $\beta = 0.0064484163$ still exceeds the 10^{-8} stopping tolerance and was not declared the root.

Retained fraction β	Saddle trace τ	Empirical height sensitivity
0.004	$+2.021777295 \times 10^{-4}$	3.19×10^{-12}
0.005	$+1.235568710 \times 10^{-4}$	3.18×10^{-12}
0.006	$+3.947197443 \times 10^{-5}$	2.87×10^{-12}
0.007	$-4.998741727 \times 10^{-5}$	2.60×10^{-12}
0.0064412278	$+6.514520363 \times 10^{-7}$	2.73×10^{-12}
0.0064484163	$+1.040296458 \times 10^{-8}$	2.73×10^{-12}
0.0064485310	$+1.636530010 \times 10^{-10}$	8.78×10^{-12}

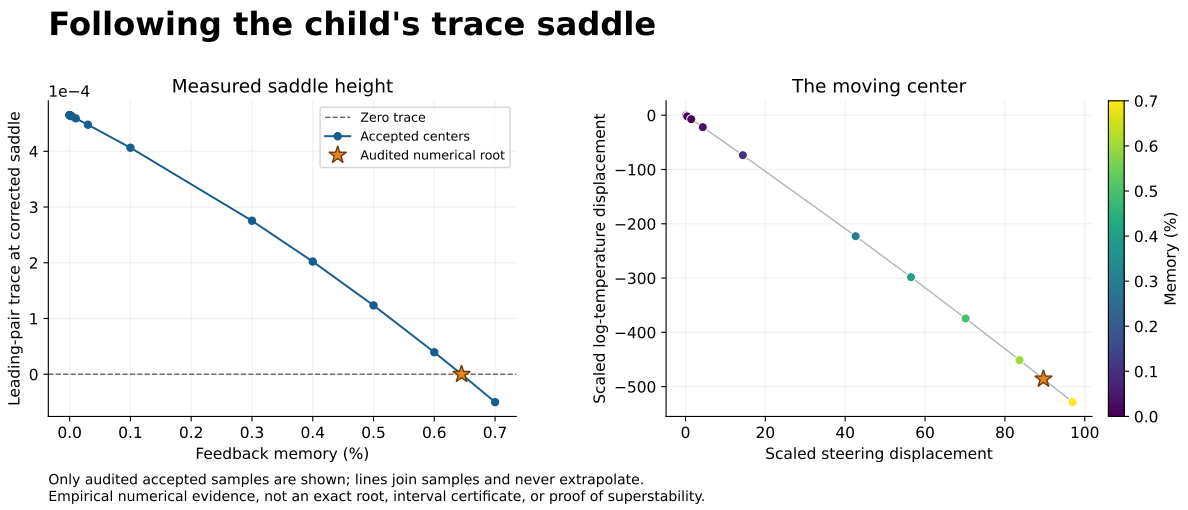

Figure 49: Following the period-18 child's trace saddle as retained feedback increases (BSC3; native float64). Left: the selected pair's trace at all 18 accepted saddle centers; the sign changes between 0.6 % and 0.7 % retained feedback. The star marks the residual-qualified estimate at approximately 0.64485 %, whose remaining trace is 1.64×10^{-10}. Right: motion of those centers in steering and log-temperature, measured from the original saved anchor in units of 2.5×10^{-6} and 7.5×10^{-6}, respectively. Lines join the measured centers and do not extrapolate or certify uninterrupted continuation between them. The earlier partial attempt contributes only its accepted centers; the failed trust-region step remains in its receipt. This is numerical evidence for a trace-contour reconnection under a changed map, not an exact-root, superstability or strongest-contraction-ridge certificate. Data and figure provenance: `docs/baby-shrimp-zoom/memory-local-extension-receipt.json`.

not certify uninterrupted branch or basin identity between samples. The spectral separation checks likewise concern the computed six-column Ritz subspace, not a certified full spectrum.

A matched image of the changed center. A fresh scan tests whether this change is visible in the same measured quantity as the original close-up. Figure 50 compares the original child center with a window translated to the retained-feedback candidate, using identical physical spans, grid sizes, integration lengths and gray scale. The original image has separated dark contraction tracks; the translated image displays an X-like junction near its center. Every pixel comes from a network trajectory; no fitted trace contour or quadratic saddle is drawn over the data. The figure therefore shows an observed change in finite-time contraction geometry that accompanies the independently measured trace-saddle sign change.

All 16,384 cells in each window yield period-18 recurrence candidates, with the same canonical six-token readout repeated three times. These fixed-opener trajectories do not supply full-state primitive-orbit verification at every pixel, nor do they by themselves establish identity with the separately continued root orbit. Moreover, the even 128×128 endpoint grids omit each window's exact center: the closest sampled coordinate is 1.1811×10^{-8} away in steering and 3.5433×10^{-8} in log-temperature. No pixel therefore measures the exact residual-qualified candidate. The X-like image is not an exact crossing certificate, a demonstration of superstability, or proof that the strongest-contraction ridges intersect. Its numerical evidence is complementary to the orbit continuation, with those limits retained.

Why a third control can reconnect the contours. Suppose the orbit is smooth and primitive, $1 \notin \sigma(M)$, the selected pair is isolated, and the saddle Hessian is nonsingular and indefinite. The implicit function theorem continues $c(\beta)$; the Morse lemma with parameters [77, Appendix B, Proposition 16] gives smooth local coordinates (U, V), depending on β, in which

$$\tau(p, \beta) = h(\beta) - U^2 + V^2. \tag{14}$$

Under these assumptions this is an exact local form, not just a quadratic fit. For $h > 0$ the trace-zero curves satisfy $U^2 - V^2 = h$; for $h < 0$ they open in the other direction. At $h = 0$ the zero set is the pair of transverse lines $V = \pm U$. A sign change with $h' \neq 0$ therefore reconnects the arms. In the original coordinates the crossing requires three scalar conditions, $\tau = 0$ and the two components of $\nabla_p \tau = 0$. If independent, two controls generically do not meet all three, while a third can. Symmetry or an exact factorization can make them dependent, so avoidance is not a universal rule for every two-parameter family. Nor does this argument guarantee that an arbitrary third intervention reaches zero. The measured continuation tests that last point for retained feedback in this particular child.

Differentiating Eq. (12) separates center motion from height change:

$$c' = -H^{-1}\tau_{p\beta}, \qquad h' = \tau_\beta, \qquad h'' = \tau_{\beta\beta} - \tau_{\beta p} H^{-1} \tau_{p\beta}. \tag{15}$$

These identities hold at an exact stationary center, with consistent coordinate scales. If memory merely relabeled the parameter plane, $\tau(p, \beta) = \tau_0(\psi_\beta(p))$ for a local diffeomorphism ψ_β, the distinguished saddle height would stay constant. Its resolved change rules out such a relabeling of the trace field. It does not establish equivalence or inequivalence of every feature of the full nonlinear maps.

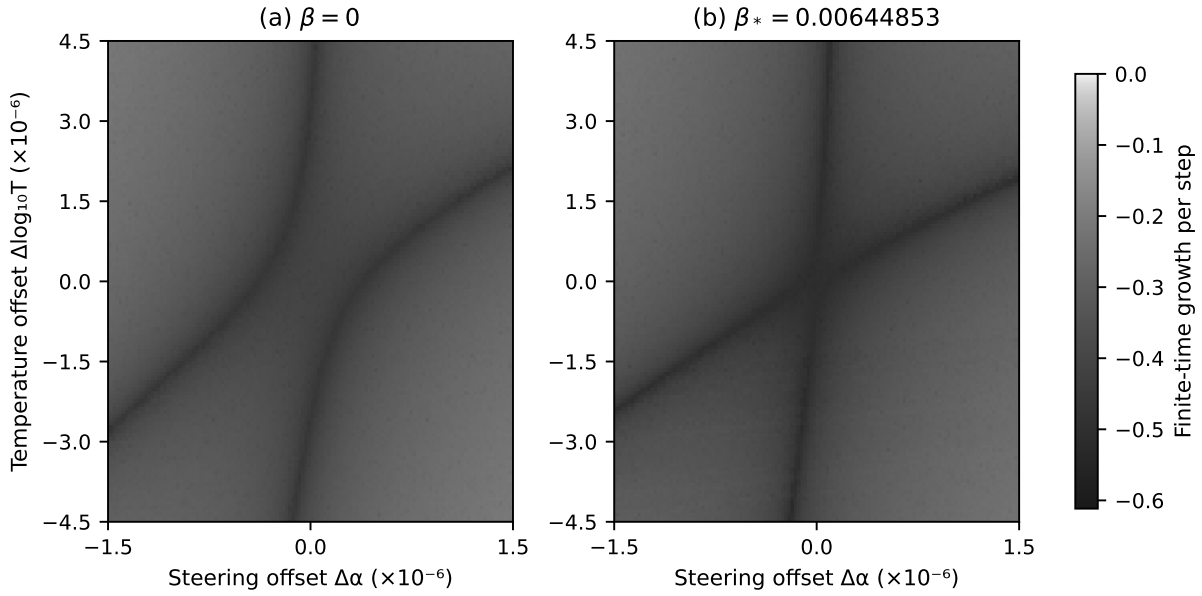

Figure 50: Measured close-ups of the child before and after retained feedback (BSZ2 and BSMZ1; native float64). (a) Original window, $\beta = 0$, centered at $(\alpha, \log_{10} T) = (0.07309363265341377, -0.1544023817786725)$. (b) Translated window, $\beta_* = 0.006448531026086656$, centered at the residual-qualified trace-root estimate $(0.07331776448598727, -0.15804614858293237)$. Each panel spans 3×10^{-6} in steering and 9×10^{-6} in log-temperature; displayed coordinates are offsets from that panel's own center. Both use 128×128 fixed-opener cells, 300 settling and 300 measurement steps, nearest-cell rendering and the same gray scale (darker means stronger contraction). The separated tracks in (a) form an X-like junction in (b), without a fitted crossing overlay. All cells have period-18 recurrence candidates and a period-6 token readout; these labels do not certify every orbit or its identity with the continued root orbit. Neither even grid includes its exact center. The conditional reconnection interpretation has separate orbit evidence; this finite-time image alone does not establish exact superstability or intersection of the strongest-contraction ridges. Source, array and image hashes: `docs/baby-shrimp-zoom/memory-junction-scan-receipt.json`.

What connects, and what does not follow. At the numerical candidate, $\Delta = -6.759154047 \times 10^{-9}$ and the pair is approximately

$$\mu_+ = +8.221415633 \times 10^{-5}, \qquad \mu_- = -8.221399268 \times 10^{-5}. \tag{16}$$

Thus the trace is small because two nonzero multipliers nearly cancel. At an exact trace zero with $\Delta < 0$, they would be $\pm\sqrt{-\Delta}$: equal contraction in magnitude, with opposite signs. This pair's trace cancellation is neither a zero-multiplier nor a unit-circle condition; it does not establish exact superstability or a fold, flip or Neimark–Sacker bifurcation. In the opposite-sign region their spectral radius is

$$\rho_{\text{pair}} = \frac{\sqrt{\tau^2 - 4\Delta} + |\tau|}{2}. \tag{17}$$

At fixed negative Δ, $\tau = 0$ minimizes this radius. In the neural map Δ also varies, so trace-zero contours need not be the exact strongest-contraction ridges seen in a finite-time image.

At a regular trace-zero point, along a declared unit direction n, put $a = \partial_n \tau \neq 0$, $b = \partial_n \Delta$ and $r = \sqrt{-\Delta}$. Expansion of Eq. (17) gives $\rho_{\text{pair}}(t) = r - bt/(2r) + |a||t|/2 + O(t^2)$. The strict condition $|b| < r|a|$ suffices for a transverse cusp minimum; $|b| > r|a|$ rules one out in that direction, and equality needs higher orders. At an exact trace-saddle crossing, where both τ and its gradient vanish, the absolute-trace contribution instead starts quadratically:

$$\rho_{\text{pair}}(q) = \sqrt{-\Delta_0} - \frac{\nabla\Delta_0 \cdot q}{2\sqrt{-\Delta_0}} + O(\|q\|^2). \tag{18}$$

A nonzero determinant gradient would tilt the radius and preclude a local minimum at that crossing. The saved fine stencil at our candidate gives $\nabla_u \Delta = (-4.54809 \times 10^{-10}, 8.75243 \times 10^{-10})$, with a transported coarse-to-fine discrepancy of norm 2.63×10^{-16} or less. If that gradient persists at an exact trace-saddle crossing, Eq. (18) gives a radius tilt approximately $(2.7660 \times 10^{-6}, -5.3230 \times 10^{-6})$ in the same scaled coordinates. At the actual finite-residual candidate this is a conditional determinant contribution, not automatically the gradient of its actual radius. It provides numerical reason to distinguish the conditional reconnection of trace contours from that of exact strongest-contraction ridges. Neither statement alone fixes the brightness geometry of a finite-time scan.

A stronger geometric crossing has an additional necessary condition. Suppose two continuously differentiable contraction-valley branches meet with distinct tangent lines, the pair radius is differentiable at their meeting, and it has a two-sided local minimum along each branch's geometric normal at the meeting point. Both independent normals must annihilate $\nabla\rho$, so $\nabla\rho = 0$. At an exact stationary trace zero, Eq. (18) then requires $\nabla\Delta = 0$. The trace-junction conditions already impose three equations on steering, temperature and memory; the determinant-gradient condition adds two more. Generically these cannot all be tuned independently with only three controls. Symmetries or identities can reduce that count, so this is not an impossibility theorem for the network. It also does not exclude tangential meetings or junctions involving nonsmooth contraction features. The condition is necessary for the stated transverse normal-minimum crossing, not sufficient.

Direct probes beyond the image. A subsequent exploratory test (BSC4) freshly evaluated the continued orbit rather than inferring geometry from image brightness. All 111 evaluated points passed the saved full-state and spectral gates. At $\beta = 0.006$, trace sign brackets lie on both sides of the saddle along its negative-curvature axis, with none on the sampled positive-curvature axis; at $\beta = 0.007$ the roles reverse. This supports the opening-direction change predicted by Eq. (14). At the latter memory setting, a root search located six trace-zero candidates, two on each of three transverse slices. Every candidate passed spectral checks at three cyclic phases; its phase-zero

residual satisfies $|\tau| < 1.94 \times 10^{-11}$. Each had a smaller pair radius than both measured offsets of ± 0.003 along a local approximate normal in the declared scaled coordinates. These are sampled zero locations and finite-offset comparisons, not complete curve continuation, a proof of local minima, or verification of their connections away from the saddle.

At the earlier β_* candidate, direct offsets also test the determinant contribution in Eq. (18). Along the unit direction opposite $\nabla_u \Delta$, the centered radius slopes at steps 10^{-4} and 3×10^{-4} are 5.999729×10^{-6} and 5.997785×10^{-6}, respectively. They differ from the conditional prediction 5.998715×10^{-6} by less than $0.017\,\%$ each. Both negative offsets have a smaller radius than the candidate. Thus the measured finite-step tilt agrees with the determinant calculation, while a trace-balance point need not give the strongest sampled contraction. The comparison does not certify the derivative at an exact trace saddle or exclude a minimum on a smaller scale. Source, samples and numerical limits: `docs/baby-shrimp-zoom/memory-connection-probe-receipt.json`.

Searching for contraction minima without imposing trace balance. The determinant-tilt test leaves a practical question: where are the contraction minima when the search is allowed to move away from $\tau = 0$? A subsequent exploratory pilot (BSC5) addresses this with actual orbit and spectral evaluations on twelve finite transverse sections. It first refines the stationary-trace candidate to $\beta = 0.006448532860189464$, $(\alpha, \log_{10} T) = (0.0733177645493531, -0.1580461496323896)$. The fresh cyclic center has phase-zero trace 3.34448×10^{-14}; the last local field stencil gives a remaining stationary correction of norm 2.877×10^{-7} or less. These are numerical residuals at a refined candidate, not exact equations or uncertainty bounds on the displayed coordinates.

The geometry is declared in the existing scaled coordinates, with steering and log-temperature scales 2.5×10^{-6} and 7.5×10^{-6}. The indefinite trace Hessian supplies two conditional null directions, each with two rays. For each unit ray t, let n be its perpendicular unit vector in this scaled Euclidean metric. At longitudinal distances $\ell = 0.002, 0.0006, 0.0002$, the searched section is

$$u(s) = u_c + \ell t + sn, \qquad s \in [-0.0005, 0.0005].$$

Here n is normal to a predictor ray; it is not a measured geometric normal to a contraction curve. Each section starts with the seven-point mesh $\{-0.0005, -0.00025, -0.000125, 0, 0.000125, 0.00025, 0.0005\}$. All strict interior minimum brackets are retained and refined by a bracketed, fit-free scalar search. No trace-zero constraint is applied to the objective, the largest computed Ritz modulus. The search does not fit a radius Hessian across the absolute-value crease in Eq. (17), and an endpoint with the smallest sampled value is never substituted for an interior minimum.

A section minimum is operationally qualified only after its measured refinement bracket has width at most 10^{-5} and its radius is smaller than the radii at both actual offsets of each size 10^{-4} and 3×10^{-5}. Every contrast must exceed $1.1480570725856676 \times 10^{-12}$, ten times the empirical radius spread across the cyclic and expanded-subspace center checks. This threshold is a resolution diagnostic, not a validated error bound. The final brackets record the search's measured resolution; they are neither confidence intervals nor certified enclosures of exact extrema.

All twelve searches finished. Each of the eight outer and middle sections has one qualified minimum at its sampled $s = 0$ point, with bracket width $9.62679078163117 \times 10^{-6}$. This does not establish exact alignment with trace zero. The innermost sections give a different pattern (Figure 51):

Ray	Result at $\ell = 0.0002$	Selected offset s
0+	no strict mesh bracket; sampled radius decreases	—
0−	one qualified section minimum	$-2.4733922 \times 10^{-4}$
1+	no strict mesh bracket; sampled radius increases	—
1−	one qualified section minimum	$+6.1677774 \times 10^{-5}$

The labels distinguish the two conditional null directions and their opposite rays. For the two qualified inner minima, the measured traces are respectively $-2.0750618721004704 \times 10^{-9}$ and $-3.4917214706089474 \times 10^{-10}$; the final brackets both have width below 6.967×10^{-6}. Thus these sampled minima lie away from trace balance. For the other two sections, monotonicity holds only between the seven sampled mesh values: the completed `no_bracket` result does not exclude narrower unsampled minima, minima beyond the declared support, or a different transverse direction.

All 223 point evaluations passed the saved full-state and spectral gates, including the 195 section-objective evaluations; ten minima passed all forty finite-offset comparisons. The smallest contrast was $6.1295545326694664 \times 10^{-12}$, above the operational threshold. Full-state closure was below 10^{-12} at every point. A dense baseline and independently expanded cyclic subspaces support the selected pair, without certifying full-spectrum dominance everywhere. The conditional quadratic-field radius predictor and actual section measurements differ by at most 1.035×10^{-11} across the tested section samples; that finite comparison does not validate derivatives or unsampled geometry.

These results strengthen the distinction exposed by BSC4: the trace-saddle frame does not by itself identify where a finite-section search finds strongest contraction. They do not establish intrinsic valley curves, connections between independently searched minima, or exact superstability. BSC5 searches inward at one refined memory setting; it does not continue contraction-minimum branches in memory or rule out their meeting nearby. The full source, plan, array, section and lifecycle bindings are in `docs/baby-shrimp-zoom/contraction-minima-receipt.json`.

A distinct test at a regular trace-zero crease. The section searches do not determine how contraction-minimum curves connect. They motivate a narrower, operational test (BSC6): can the two one-sided *Hessian features* of the pair radius meet on its regular trace-zero crease as memory varies? This differs from the stationary-trace-saddle test. Here $\tau = 0$ but $\nabla_u \tau \neq 0$, so the radius generally has a cusp. The target memory value and the corrected steering–temperature location belong to this new constrained test; they are not a new estimate of the stationary trace-saddle location. The Euclidean metric remains the declared scaled steering–log-temperature metric. No common radius Hessian is assigned across that cusp.

For $\Delta < 0$, define the two smooth extensions

$$\rho_s(u,\beta) = \frac{\sqrt{\tau(u,\beta)^2 - 4\Delta(u,\beta)} + s\tau(u,\beta)}{2}, \qquad s \in \{+1,-1\}. \tag{19}$$

The measured pair radius equals ρ_+ where $\tau > 0$ and ρ_- where $\tau < 0$. Both extensions have value $r = \sqrt{-\Delta}$ on the crease. At an exact regular trace zero, write $g_\tau = \nabla_u \tau$, $g_\Delta = \nabla_u \Delta$, $H_\tau = \nabla_u^2 \tau$ and $H_\Delta = \nabla_u^2 \Delta$. Their separate limiting derivatives are

$$\begin{aligned} G_s &= -\frac{g_\Delta}{2r} + \frac{s}{2} g_\tau, \\ H_s &= \frac{g_\tau g_\tau^{\mathsf{T}}}{4r} - \frac{H_\Delta}{2r} - \frac{g_\Delta g_\Delta^{\mathsf{T}}}{4r^3} + \frac{s}{2} H_\tau. \end{aligned} \tag{20}$$

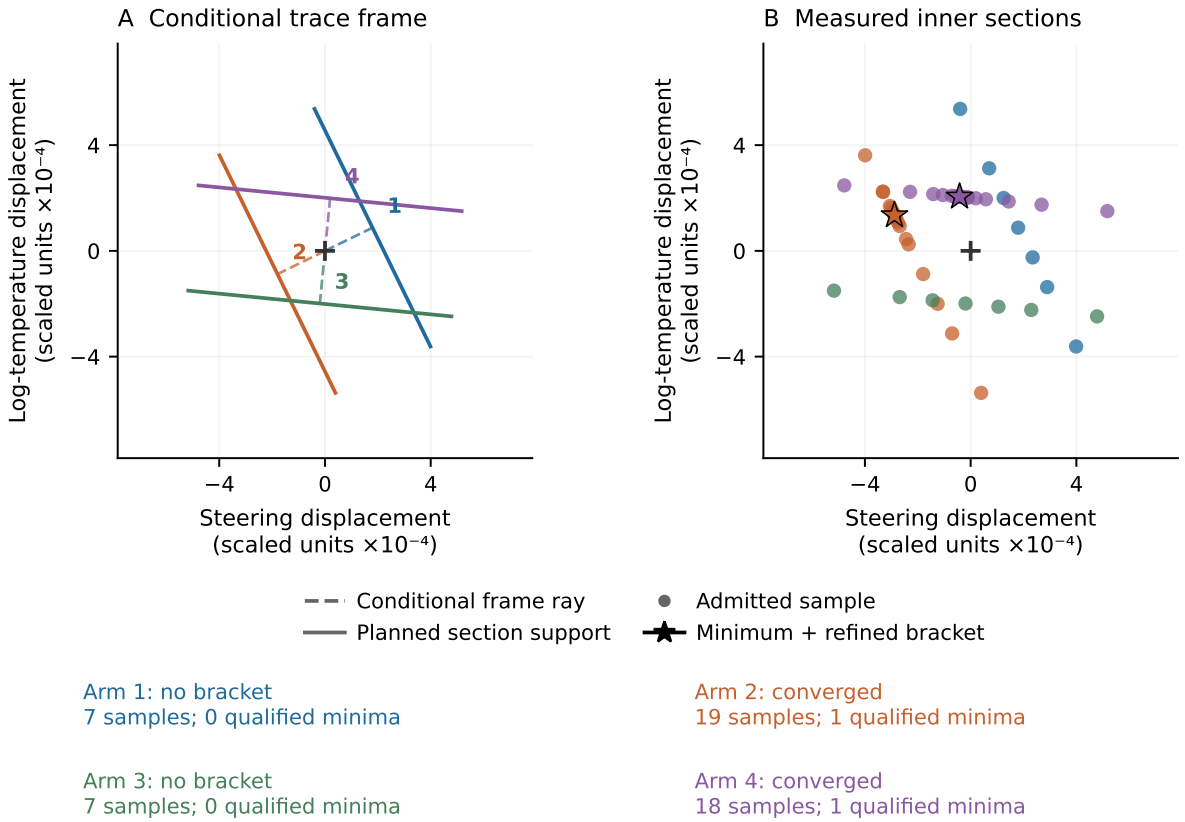

Figure 51: Independent contraction searches near the memory-tuned trace center (BSC5, $\beta = 0.006448532860189464$). Left: the conditional trace-Hessian rays and all four transverse sections at the smallest planned distance $\ell = 0.0002$. Right: every measured sample on those sections; stars identify the two qualified finite-section minima and black bars their final numerical brackets; these bars are smaller than the star markers at this display scale. Arms 1 and 3 have no strict minimum bracket on the sampled mesh; arms 2 and 4 have displaced minima. The origin is the refined numerical trace-saddle candidate. The four sections' counterparts at larger distances and every measured radius profile remain in the bound receipt. The axes use equal scales in the declared metric: steering displacement divided by 2.5×10^{-6} and log-temperature displacement by 7.5×10^{-6}, displayed in units of 10^{-4}. Section supports are sampling paths, not measured valley curves; independently found minima are not joined. Bracket widths are neither confidence intervals nor certified exact-extremum bounds. Source and reviewed rendering: `paper/llm/figures/fig-memory-section-minima.json`.

These are derivatives of two smooth extensions, not two estimates of a single Hessian of the nonsmooth radius.

When the largest eigenvalue of H_s is positive and simple, let e_s be its unit eigenvector, with sign transported from a declared reference. The measured diagnostic is

$$f_s = e_s^{\mathsf{T}} G_s. \tag{21}$$

An exact $f_s = 0$ with positive curvature would be a strict minimum along the fixed straight direction e_s through that point for the smooth extension. It need not be a minimum along the geometric normal to the level set $f_s = 0$ where that level set is regular. For the actual smooth-side feature field, $\nabla_u f_s = H_s e_s + (De_s)^{\mathsf{T}} G_s$; eigenvector rotation generally prevents that normal from being e_s. The side on which the smooth extension equals the physical pair radius also matters. A meeting of these conditional features therefore does not establish a crossing of intrinsic contraction valleys.

Fresh correction and empirical precision checks. BSC6 follows the two equations $\tau = 0$ and $f_+ = 0$ in the two scaled coordinates at a central memory setting and at offsets $\pm 3 \times 10^{-9}$. The outward predictors start from the freshly admitted central point and use the previously measured motion with memory. At each target, at most two coordinate corrections are followed by fresh orbit and spectral evaluations. The unused feature f_- is then measured; its sign is not a condition for accepting a target. These are three locally corrected samples, not a verified continuous curve over the intervening memory interval.

Nine-point central differences of the smooth fields τ and Δ use primary step 3×10^{-4} and matched check step 10^{-4}. For the corrector and the reported conditional crease diagnostics, their measured jets are inserted into Eq. (20) with the *trace value* set to zero at the unchanged coordinate. The actual trace residual is retained, together with the estimated normal displacement $-\tau g_\tau / \|g_\tau\|^2$; that displacement is not substituted for a freshly measured point. Differentiating this projected field away from the crease is not the same operation as differentiating the actual smooth extensions in Eq. (19). They agree on the exact crease but can have different derivatives transverse to it. The projected corrector is used to locate candidates for the constrained intersection, not to identify the tangents or normals of active-side contraction valleys.

Every target and stencil point must pass the saved full-state period-18, orbit-alignment, and leading-pair spectral checks. A dense reference and a fresh cyclic seed replay precede the target calculations; each final candidate is checked at all three cyclic phases. The primary trace residual and every fresh cyclic trace must satisfy $|\tau| \le 10^{-12}$, the primary feature must satisfy $|f_+| \le 10^{-8}$, and the remaining primary model correction must have norm at most 2×10^{-6}. These tolerances define numerical admission, not exact equalities.

The uncertainty calculation combines the previously measured expanded- subspace and fresh cyclic spreads, declared point floors, paired-stencil differences, eigenvector sensitivity and the conditional trace projection. It also transfers uncertainty in the location of the (τ, f_+) intersection into f_-. With $J = D_u(\tau, f_+)$, the propagation includes the actual equation residuals, paired and half-step changes in J and $\nabla_u f_-$, and a remainder for rotation of the perturbed inverse Jacobian. The retained empirical intersection radius must not exceed 5×10^{-6}; the total transferred f_- envelope must not exceed 5×10^{-8}. Both normalized Jacobians must have minimum singular value at least 0.2. Positive curvature, eigenvalue separation, transported orientation and the strict crease-normal cusp slopes must remain resolved after the declared empirical errors. A resolved sign requires $|f_-|$ to exceed ten times its total envelope. These propagated quantities are conditional on empirical jet-error estimates and a local linear sensitivity model; they are neither confidence intervals nor validated bounds on the neural map.

A resolved numerical sign bracket, with a near-zero central candidate. All three targets passed all thirteen admission checks; all 101 evaluated points passed the saved full-state and leading-pair spectral gates, with all 99 target alignment comparisons admitted. There were no failed or unfinished point attempts. The largest full-state closure residual was below 9.663×10^{-13} and the largest leading-pair relative residual below 8.722×10^{-9}. Across the final candidates and their cyclic replays, the largest absolute trace residual was below 1.735×10^{-13}. The final primary $|f_+|$ residuals were below 1.823×10^{-11} and remaining primary correction norms below 6.782×10^{-9}.

With the central target $\beta_0 = 0.006448531183244975$, the unused feature has the following outcomes. The displayed empirical envelopes are rounded upward; feature values are rounded estimates.

$\beta - \beta_0$	Conditional f_-	Empirical envelope	Status
-3×10^{-9}	$+5.125036 \times 10^{-7}$	2.4515×10^{-8}	resolved positive
0	$+1.703940 \times 10^{-10}$	3.1624×10^{-8}	sign unresolved
$+3 \times 10^{-9}$	-5.670830×10^{-7}	3.4233×10^{-8}	resolved negative

The endpoint magnitudes exceed their unrounded empirical envelopes by factors approximately 20.9061 and 16.5655, respectively. These are operational resolution margins, not statistical significance levels. The central feature is much smaller than its envelope: its sign, and an exact zero there, remain unresolved. The largest empirical intersection-location radius across the three cases is below 2.309×10^{-6} in the declared scaled metric.

Thus BSC6 gives a numerical sign bracket along the corrected conditional intersection over the interval

$$[0.006448528183244975,\ 0.006448534183244975].$$

If an exact continuous regular branch of $(\tau, f_+) = (0, 0)$ connects the endpoints, and the empirical uncertainty assumptions correctly control the signs on that branch, continuity would imply a mixed-feature zero between them. The three admitted samples and their state alignments support tracking the same local feature, but do not certify the intervening branch or those error assumptions. This is stronger than a small nominal residual at a single candidate, while remaining short of a mathematical existence proof. The `docs/baby-shrimp-zoom/mixed-junction-receipt.json` binds every point, source, plan, correction, uncertainty and terminal lifecycle check. The separately reviewed `docs/baby-shrimp-zoom/mixed-junction-summary.pdf` shows both the conditional feature and the actual signed primary/cyclic trace residuals, without joining the three targets; its provenance is `docs/baby-shrimp-zoom/mixed-junction-summary.json`.

The test does not impose $\nabla\rho = 0$, does not continue intrinsic normal-minimum curves, and does not certify an exact mixed-feature zero. Its regular-crease setting is distinct from the differentiable, transverse normal-minimum meeting considered above; a conditional Hessian-feature junction would not contradict that necessary-condition argument. The mixed-feature conditions $\tau = f_+ = f_- = 0$ impose three equations on the three controls. A nonsingular full three-variable Jacobian would make an exact solution locally isolated; this dimension count establishes neither its existence nor that nonsingularity here. Since the selected multipliers remain nonzero, neither an admitted target nor a numerical sign bracket establishes exact superstability or a unit-circle bifurcation.

An intervention mechanism, not an identified native coupling. Retained feedback changes both the one-step derivative and the orbit on which that derivative is evaluated. At $\beta = 0$, write $J_i = DF_p(x_i)$, $\eta_i = \partial_\beta x_i$ at fixed p, and $m = 18$. Then

$$\eta_{i+1} = J_i\eta_i + x_i - x_{i+1}, \quad \eta_m = \eta_0, \qquad \dot{J}_i = (I - J_i) + D^2F_p(x_i)[\eta_i, \cdot\,]. \tag{22}$$

The periodic tangent equation can be solved through $(I - M)\eta_0 = b$, with b obtained by propagating the forcing. Applying the product rule to $M = J_{m-1} \cdots J_0$ and tracing against the pair's spectral projector separates the explicit retained-state path from the orbit-mediated response. At the original saved child anchor, their contributions to $\partial_\beta \tau$ are approximately -0.032696 and -0.106638, totaling -0.139334 (`docs/baby-shrimp-zoom/memory-sensitivity-receipt.json`). The total was evaluated at one cyclic phase and checked against three-phase finite-offset orbit controls; the primary phase's finest central slope differs by less than 0.008 %, with approximately quadratic convergence. These are fixed-location derivatives. The original anchor was only an approximate saddle, so they must not be substituted for the moving-height derivative in Eq. (15) without correcting stationarity.

This decomposition identifies how the specified intervention acts; it does not identify an internal pathway whose native coupling causes avoidance. Off-diagonal entries of one return matrix depend on its basis; where the matrix is diagonalizable, a change to its eigenbasis removes them without changing its trace. Even in a fixed set of state channels, nonzero coupling alone does not force avoidance. For an exact counterexample, the autonomous polynomial map

$$\begin{aligned} q(x) &= 1 + ux + (-3 - 2u - v)x^2 + (2 + u + v)x^3, \\ G_1(x, y) &= q(x) + k[1 + (r - 1)x]y, \qquad G_2(x, y) = k(x^2 - x) \end{aligned} \tag{23}$$

has the primitive orbit $(0, 0) \mapsto (1, 0) \mapsto (0, 0)$. Its one-step Jacobians there are $A = \left(\begin{smallmatrix} u & k \\ -k & 0 \end{smallmatrix}\right)$ and $B = \left(\begin{smallmatrix} v & rk \\ k & 0 \end{smallmatrix}\right)$, so direct multiplication gives $\operatorname{tr}(BA) = uv + (1 - r)k^2$ and $\det(BA) = -rk^4$. The saddle height is $(1 - r)k^2$. At $r = 1$ the trace-zero lines cross even for $k \neq 0$, although all four cross-channel one-step entries remain nonzero: two return-path contributions cancel. For $r > 0$ and $k^2 \max(1, r) < 1$, the cycle at the saddle is attracting. This is an analytic counterexample, not a model fitted to the network. To identify a native neural mechanism would require independently defined pathways, measured projection and nonnormality errors, predictions of the full trace and determinant at additional points, and interventions continuing the same primitive orbit.

The supported measurement is specific: adding retained feedback changes the sign of the selected pair's trace-saddle height on the continued period-18 child family. Under the local smooth-saddle assumptions, the Morse form interprets this as a reconnection of trace-zero contours. The sampled axis probes and zero locations support that interpretation without tracing complete connections away from the saddle. The experiment does not uncover a hidden crossing in the original $\beta = 0$ map, establish an exact two-dimensional nonlinear reduction, or show that the period-6 mother or other windows reconnect under the same intervention. The earlier mother's unsuccessful determinant-saddle test and subsequent trace amendment remain distinct evidence (Figure 35).

12 From greedy loops to chaos: one picture

The preceding sections establish the pieces: a factorized derivative, folds, cascades and windows. This section assembles them into one account of why the loop of Qwen2.5-0.5B along the philosophy direction looks the way it does. It answers three questions: where the periodic orbits come from, why the windows cross one another, and why a network with 896-dimensional states shows the signatures of one-dimensional maps. Every step was checked against the full map in float64 (`docs/e2e/`).

At zero temperature: greedy loops. As $T \to 0$ the soft token becomes the embedding of the most probable token (at an exact tie, the average of the tied embeddings). Wherever each visited state has a unique most probable token, the loop thus becomes greedy decoding with a one-token memory (Section 3), and its eventual cycles are token cycles: the zero-temperature attractors,

superstable because the limit map is locally constant away from ties. Along this direction the greedy map has a two-token loop " is there" for α between 0.06301 and 0.07773, and a three-token loop "\n\nThe question" for α between 0.06609 and 0.07006. At $\alpha = 0.070$ and low temperature, the loop indeed settles on these: in a column of 96 temperatures from $\log_{10} T = -2$ upward, 59 end on "\n\nThe question" and 16 on a period-4 loop that alternates " is it" and " is there".

At finite temperature: compromises between competing continuations. The attractors that organize the maps are not greedy loops. No value of α between 0 and 0.12 (in steps of 0.0005) makes " is it" a greedy cycle: after " it" greedy decoding writes "'s" for $\alpha < 0.105$, and above that, where it writes " is", the next greedy token after " is" is not " it". The 2-cycle exists only because the soft token is a mixture. After " is", the fed-back vector mixes " it" and " there"; after " it", it mixes " is" and "'s". Each step thus feeds back a weighted blend of two continuations, and the blends close into a cycle that greedy decoding never closes. Such a compromise lives in a band of temperature and dies at both of its ends. At $\alpha = 0.070$ the 2-cycle is stable between $\log_{10} T = -0.5147881$, where a multiplier reaches -1 (a flip: the competition becomes too sharp, and the blend alternates), and $\log_{10} T = -0.1741065$, where a multiplier reaches $+1$ (a fold, the cyan curve of Figure 10). The period-3 loop of the crescent, "\n\nThe existence", is a compromise of the same kind, between " existence" and " question" in the noun slot. It too lives in a band, between a fold above (at $\alpha = 0.065$, extrapolated to $\log_{10} T = -0.14536$ from a continuation whose last converged multiplier was 0.956) and a flip below. It does not connect to a zero-temperature loop.

What happens at the two ends. Beyond the fold the 2-cycle is gone, but its ghost delays every trajectory that passes near it. The number of laps N spent there grows as the inverse square root of the distance to the fold, so periodic windows of period $2N + 6$ line up parallel to the fold curve and crowd toward it, with distances close to $1/N^2$ (log-log slope -2.00; Section 8). This gives the near-vertical fan of stripes. Below the flip, a compromise starts alternating between its two blends, and in the crescent the alternation doubles again and again, a Feigenbaum cascade into chaos (Section 9). The near-horizontal band that crosses the stripes near $\log_{10} T = -0.14$ is the first doubling of the crescent's period-3 loop, continued to larger α. At $\alpha = 0.070$ it is bounded by two flips, at $\log_{10} T = -0.1392937$ above and -0.1414292 below. The two families cross because each is bounded by its own curves: the stripes by the distance to the fold of " is it", the band by the flips of the period-3 loop.

Why the two families cross in both orders. Which family a map shows where they overlap is not decided by either family's curves. In parts of the crossing the crescent's family lies over the stripes, in others a stripe lies over the band, and the order can change from one cell to the next. Every cell of the maps starts from the same state, the embedding of the prompt's last token, so where two attractors coexist a map shows the one whose basin contains that state. To see what the maps hide, we recomputed the window of Figure 11 on a 192×256 grid as eight maps: one from the fixed start; three from states of the crescent's period-3 orbit, of the band's 6-cycle and of the period-10 stripe; and four sweeps that follow each attractor across the plane as far as they can, with every column (or row) of cells starting from the final states of the previous one, in both directions of α and of T (Figure 52; `docs/sheets/analysis.json`). The four starts discard 600 transient steps and measure over the next 300; the sweeps use 200 and 200 (600 transient steps in their first column or row). An independent audit with separate code reproduced every count below from the raw maps.

- *Two overlapping sheets.* We group periodic attractors by the text their loops write, not by continuation. Periodic attractors of the crescent's family (loops that write "\n\nThe ...") were found at 13,585 of the 49,152 cells, and those of the 2-cycle and its stripes (loops that write " is ...") at 36,787. Both exist at 9,675 cells, on a broad arc that covers the crescent and runs under the stripes (Figure 52b). The fixed-start map hides 32 % of the crescent family's cells and 36 % of the 2-cycle family's. Along each temperature row the crescent family's branch is continuous: at $\log_{10} T = -0.1415$ it runs from $\alpha = 0.066$ to about 0.0745, under the stripes (Figure 52d; Figure 53 shows the two sheets in 3D).

- *Selection by the start.* Where the seven other maps find a periodic attractor of only one family, the fixed start reaches the other family at just 70 of those 31,042 cells (0.2 %). Where both families were found, it reaches the crescent's family at 65.4 % of the cells and the 2-cycle's family at 28.5 % (the rest are chaotic or unresolved). In the crescent zone the split is 70 to 29 %; across the band it is 48.5 to 48.0 %, with the choice alternating from cell to cell, which is the alternating order seen in the map. Of the map's 1,376 edges between periodic cells of different families, 557 are pure switches: each side's family was also found on the other side, and only the basin that holds the fixed start changes. At 662 one side's family was not found on the other side by any map, and at 157 neither was. Eight maps cannot find every attractor, so these are not proofs that a family ends there, and the switch count is a lower bound.

- *The crescent's own edge is a switch.* At each temperature the crescent's period-3 orbit exists from its fold onward, but the fixed start reaches it only farther right: in a strip up to 1.1×10^{-3} wide in α both coexist and the map shows the 2-cycle (Figure 54). At $\log_{10} T = -0.148$ the map's existence edge, $\alpha = 0.06444$, coincides with the fold found by continuation, 0.064444. The strip narrows to the grid spacing toward the top of the crescent. At $(\alpha, \log_{10} T) = (0.06468, -0.148)$, inside the strip, interval arithmetic proves that the 2-cycle and the 3-cycle coexist and that both attract (the parameters are read as exact binary64 rationals, and the temperature is the real number $\exp(\ln 10 \cdot \log_{10} T)$ for that exact $\log_{10} T$; Appendix L).

- *Chaos.* Almost all chaotic attractors (96 % of the chaotic cells) visit the crescent family's blank-line and "The" tokens, including the chaos between the stripe windows, so the chaotic regions cannot be assigned to the 2-cycle's side alone.

Why the fixed start lands where it does is not explained by its nearness to either attractor. At three crescent cells random starts (single token embeddings and mixtures of the top 64) reached the 2-cycle in 82 to 90 % of cases, and the nearest attractor phase predicted their destinations worse than always guessing the 2-cycle (`docs/basins/basins-summary.json`). The newline that ends the prompt happens to lie in the crescent's basin over most of the overlap, even though that basin is the smaller one for random starts. Settle times give no sign of long transients at the switches (median 94 steps near the edges and 87 elsewhere), so these basin boundaries do not hold trajectories for long. Refined at single switches, both mechanisms appear. Along two short transects through the band, chosen from the fixed-start map before the other maps were inspected ($\alpha = 0.0708$ from $\log_{10} T = -0.1440$ to -0.1396, and $\log_{10} T = -0.1396$ from $\alpha = 0.0690$ to 0.0700), we bracketed every switch of the fixed start's family by bisection, to 1/64 of the transect spacing at three of the eight switches and to 1/4 or 1/2 of it at the other five, where the bisection met a third loop; at both ends of each bracket we ran the fixed start, continued each end's attractor at the other end and computed the multipliers of every periodic orbit found (`tools/sheets_transect.py`). Of the eight switches, two are selections (each side's family persists on the other side, all four orbits

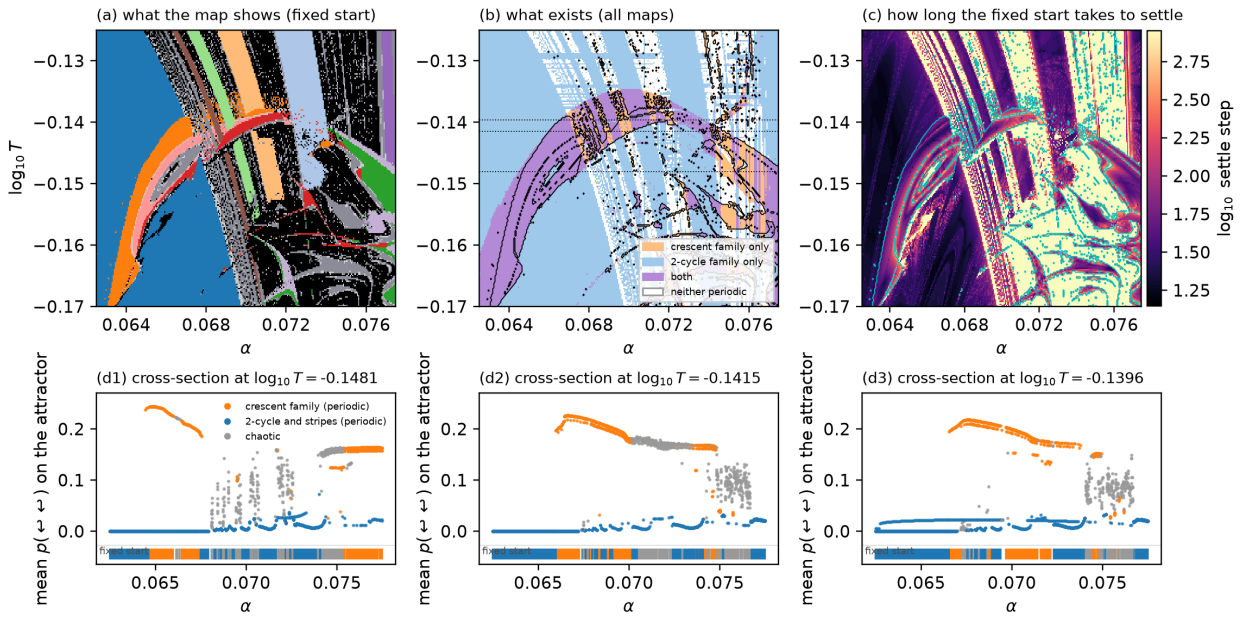

Figure 52: What the fixed-start map shows, and what it hides (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, float64; the window of Figure 11 on a 192×256 grid). (a) The attractor reached from the fixed start (colors as in Figure 11; black: chaos). (b) Where periodic attractors of each family exist, according to all eight maps of the run: four starts (the fixed start and states of the crescent's period-3 orbit, the band's 6-cycle and the period-10 stripe) and four adiabatic sweeps; black contour: where the fixed start reaches the crescent's family. (c) How many steps the fixed start takes to settle. (d) Cross-sections along three temperatures: the mean probability of the blank-line token on the attractor reached in each map against α, colored by family, with the family the fixed start reaches shown as a ribbon at the bottom. The crescent family's branch (orange) continues under the stripes while the fixed start switches between the two branches.

attracting, and only the basin that holds the fixed start changes), two are intrinsic (one side's family does not persist across the bracket), and four remain unresolved, involving chaos or an orbit that converges too slowly near its doubling; a calibration bracket at the crescent's edge is a selection, as expected. A classifier fixed before the run compared continuation arms arriving from opposite sides and classed none of the first seven switches as selections; we replaced it, before any cross-test data, with the test above, and both results are kept (`docs/research-log.md`).

Evidence for low-dimensional dynamics. A single step of the map does not squeeze the state onto a single direction: it is not close to rank one. Along these orbits, the second eigenvalue of the readout covariance Σ_p is 0.23 to 0.87 of the first, and the second singular value of the step derivative is 0.11 to 0.90 of the first. Composed over a cycle, however, the leading singular value becomes more dominant. The ratio of the two leading singular values is 0.081 for the 2-cycle, 0.092 for the period-3 loop and 0.26 for the period-6 band; it falls to 3×10^{-4} over twelve steps of the chaotic crescent and to 4×10^{-9} in the chaotic stripes. These ratios show a stronger separation of the leading singular directions over the composed map than over a single step. The finite-time Lyapunov spectra provide complementary evidence. The map has one exponent for each of its 896 directions, ordered from largest to smallest, and the maps give only the largest of each cell, from the drift of a nearby twin. To find the next few, we followed four small perturbations at once with the network's exact derivative, re-orthogonalizing them at every step so that they do not all turn toward the fastest-growing direction (`tools/lyapunov_spectrum.py`), and averaged

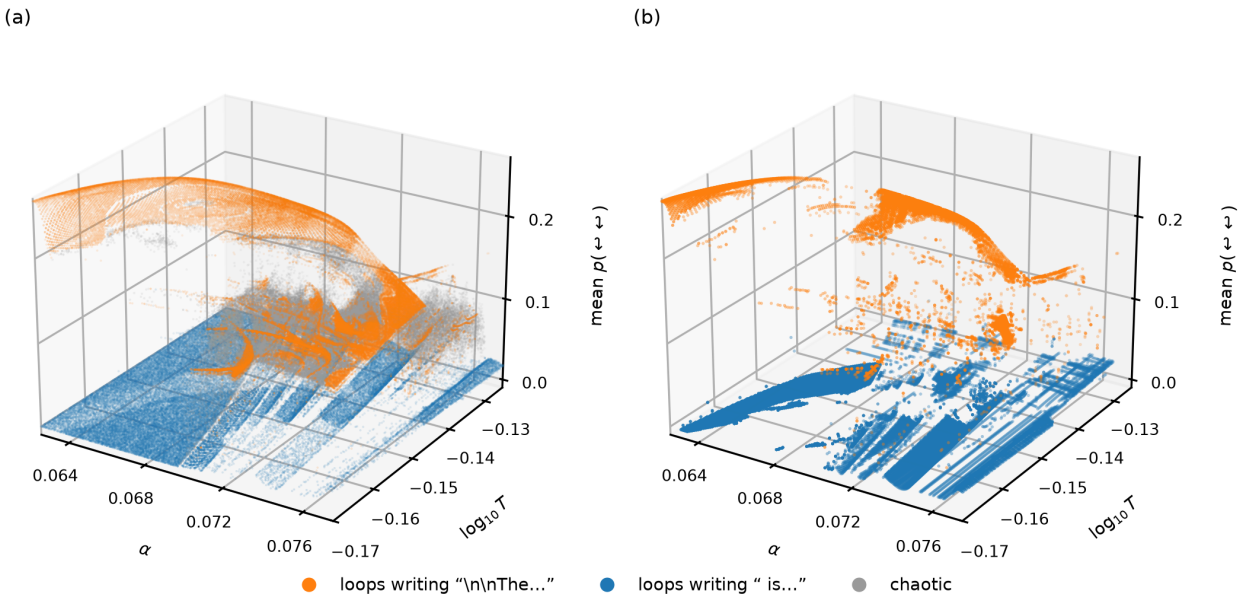

Figure 53: The two families as sheets. Each point is an attractor found at one cell of Figure 52, at a height equal to its mean probability of the blank-line token, which separates the families: in (a), 90 % of the orange points lie between 0.07 and 0.25 and 99 % of the blue points below 0.03 (paper/llm/figures/fig-sheets-3d.json). Orange: the crescent's family (loops that write "\n\nThe ..."); blue: the 2-cycle and its stripes (loops that write " is ..."); gray: chaos. (a) The attractors found by the eight maps, thinned at random to at most 9000 per map and class. (b) Only what the fixed-start map hides: periodic attractors that the seven other maps find where the fixed start reaches the other family, chaos or no period.

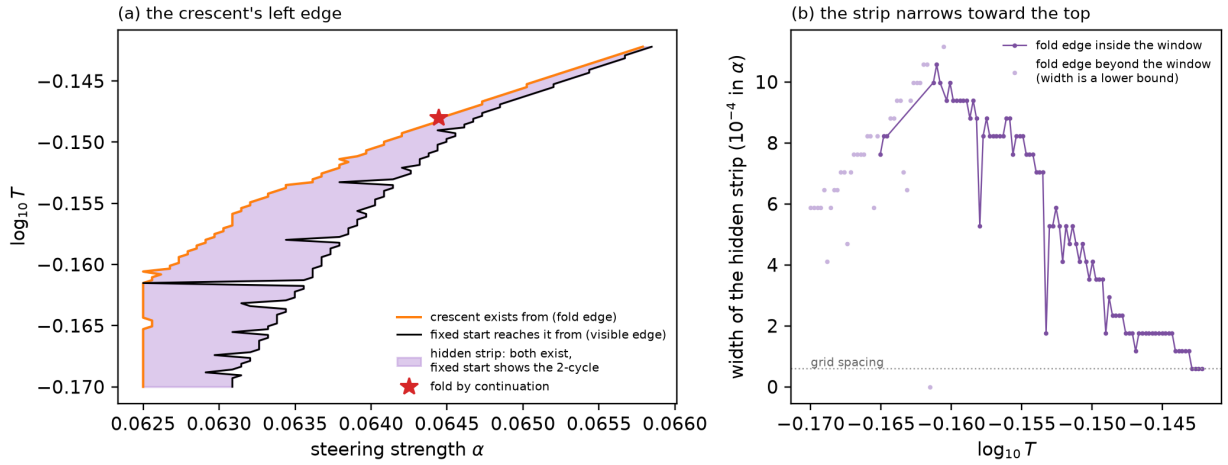

Figure 54: The crescent's left edge is a switch, not its fold. (a) Along each temperature, the smallest α at which the crescent's periodic family exists in any map of the run (orange) and the smallest at which the fixed start reaches it (black); between them (purple) both exist and the map shows the 2-cycle. Star: the period-3 fold found by continuation at $\log_{10} T = -0.148$. (b) The width of that strip; pale points: rows where the fold lies beyond the window, so the width is a lower bound.

their logarithmic growth rates. We did this only at three representative chaotic cells, one in the crescent's chaotic core and two among the chaotic stripes. Over 600 steps their largest exponents are 0.109, 0.424 and 0.239 per step, and the next three are strongly negative, between -0.82 and -2.01. The finite-time Kaplan–Yorke dimension estimates, one plus the largest exponent divided by the magnitude of the next, are 1.13, 1.43 and 1.21. These estimates suggest low-dimensional dynamics at the three sampled cells, with one expanding direction. (Over 3000 steps the largest exponents are 0.098 ± 0.010, 0.437 ± 0.016 and 0.196 ± 0.014; Section 16.) These finite spectra and the singular-value ratios support strong effective contraction in the sampled directions and motivate scalar approximations: on the sampled range, the crescent's attractor gives a scalar return relation with a single extremum (Figure 55e). The measured cascade separately provides finite evidence consistent with Feigenbaum scaling (Section 9); no global one-dimensional invariant reduction is established. A scalar model built from the two competing tokens, $q' = \sigma(h(q)/T)$, reproduces the crescent's hump (Figure 55e) and, at $\log_{10} T = -0.156$, places its first flip between $\alpha = 0.0646$ and 0.0648, against 0.0644767 in the full map. The agreement is qualitative: it does not reproduce the cascade.

The picture. At zero temperature, away from ties, the loop writes greedy loops. As the temperature rises, the soft token starts to mix the two continuations a token has, and the mixtures form new periodic orbits that greedy decoding does not have. The two we continued each live between a flip, where the competition is too sharp, and a fold, where it is too weak. Their folds leave ghosts that line up the windows, and their flips start cascades into chaos. At the three sampled chaotic cells, the finite-time spectra suggest one expanding direction. This supports exploring one-dimensional descriptions of these orbits, as Feigenbaum's ratios also suggest, although no global one-dimensional reduction is established. Such a description would have to pass two tests: the position along a band at one visit must determine the position at the next, and every periodic orbit that the reduced map allows must be an orbit of the full map, which Newton's method can check. Where both tests pass, the kneading theory of one-hump maps [76] would say which sequences of L and R can occur and how fast their number grows; the window labels of Section 11 are a first, approximate use of that theory. Steering moves which continuations compete, and so where all of this happens.

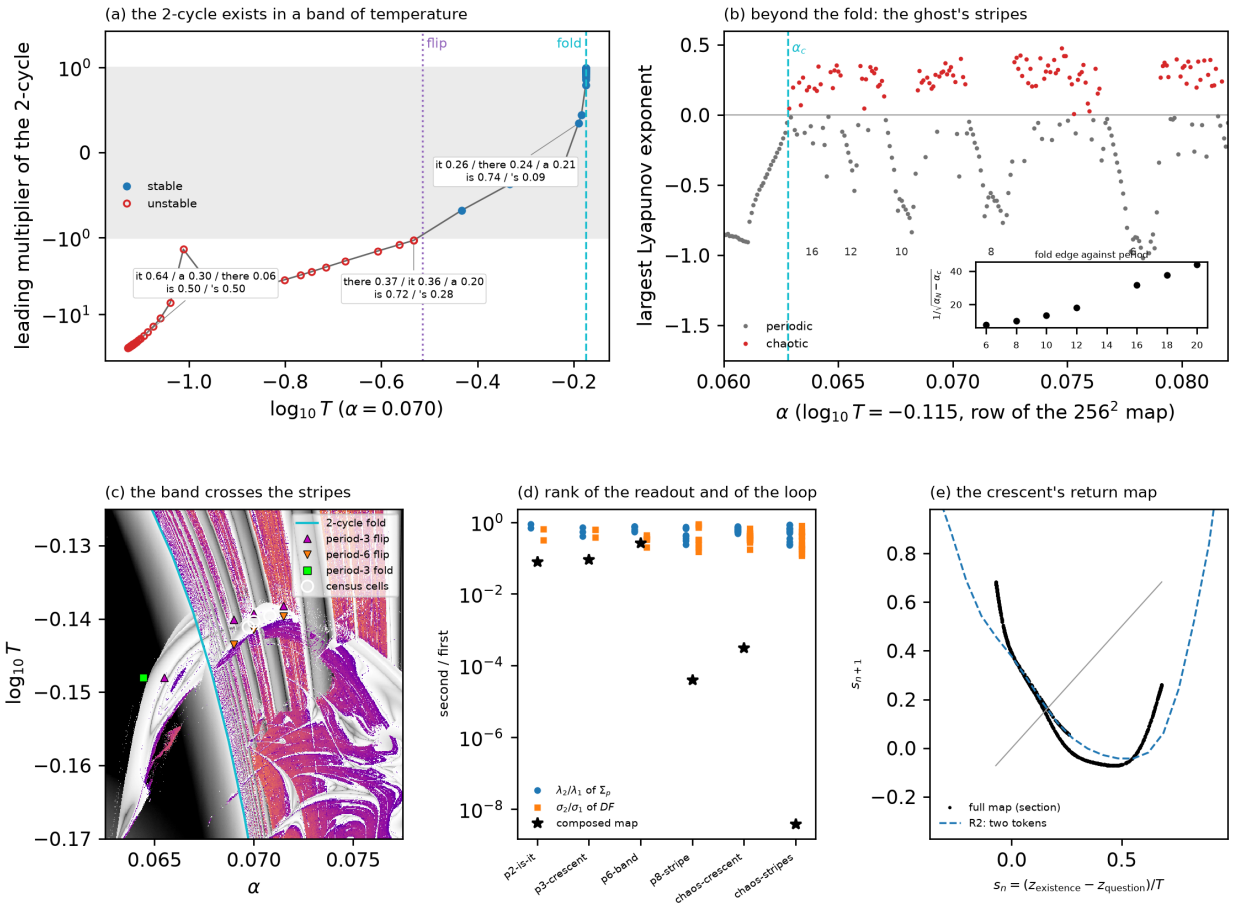

Figure 55: One picture, in float64 (Qwen2.5-0.5B, philosophy direction, layer 12). (a) Leading multiplier of the " is it" 2-cycle along temperature at $\alpha = 0.070$. It is stable (blue, inside the gray band $|\mu| < 1$) only between the flip ($\log_{10} T = -0.5148$) and the fold (-0.1741). The boxes give the two blended continuations at each phase of the cycle, with their probabilities. (b) Largest Lyapunov exponent along the row $\log_{10} T = -0.115$ of the 256×256 map. Beyond the fold α_c of the 2-cycle, periodic windows (gray, labeled with their periods) alternate with chaos (red). Inset: $1/\sqrt{\alpha_N - \alpha_c}$ for the fold edges of the windows against their period. For the longer windows (periods 16 to 20) it grows linearly with the period, as the ghost predicts (Section 8). (c) Part of the crossing map, with the fold curve of the 2-cycle (cyan). The period-6 band is bounded by the flip of the period-3 orbit above and the flip of the period-6 orbit below, located on cuts at $\alpha = 0.069$, 0.070 and 0.0715. On the cut $\log_{10} T = -0.148$, the crescent's period-3 orbit is born at a fold ($\alpha = 0.06444$) and flips at $\alpha = 0.06548$. White circles: the census cells. At the two in the overlap all ten starts ended on the band's 6-cycle; at the third (0.0708, -0.1396) one of five starts ended on a period-10 loop. (d) Second over first eigenvalue of the readout covariance Σ_p and second over first singular value of the step derivative DF at every recorded state of four periodic orbits and two chaotic attractors, and the same ratio for the derivative composed over one period, or over twelve steps for the chaotic cells (stars). (e) Return map of the crescent's attractor in the noun log-odds s (black), with the two-token model $q' = \sigma(h(q)/T)$ (dashed).

13 Beyond the one-token loop

The soft-token map feeds back one expected embedding at a fixed prefix and position. This section summarizes what changes when the loop keeps more context, when tokens are sampled and in other models (details in Appendices G and H); Sections 14 and 15 then follow model size and a second system.

Memory and sampling. Here memory means the number of fed-back tokens; it differs from the retained fraction of one state in Eq. (10). As $T \to 0$ the soft-token map tends to a map with finitely many values that agrees, away from ties, with greedy decoding with the same finite memory, so every orbit of the limit ends in a cycle (Appendix G). Fed back the last 8 or 16 soft tokens instead of one, Qwen2.5-7B is chaotic at 21 and 20 of 24 temperatures, against 8 with one. In Soft Thinking's growing context, where every soft token stays visible, an early perturbation of 10^{-10} grows exponentially in the same model inside its one-token chaotic band. This does not hold in every model: in Qwen2.5-0.5B the growth moves to sharper distributions, outside its narrow one-token chaotic band, and Qwen2.5-72B stays nearly insensitive (Section 14). Sampling is different. When the expected embedding is replaced by the mean of m sampled token embeddings, the laminar phases of the bottleneck, 120 steps long in the exact map at the closest offset, last on average 10.5 steps with $m = 64$ and 7.8 with $m = 8$ (10.3 to 10.7 and 7.6 to 7.8 over the three offsets), and with single samples such phases become rare, one to three per offset. The deterministic intermittency belongs to the expected-embedding map; sampling noise turns the bottleneck into a noise-limited trap whose strength grows with m. Two exploratory tests go a step closer to practice (Appendix G). Feeding back only the 5 to 50 most probable tokens, as the published Soft Thinking implementation does with 15, left the loop of Qwen2.5-0.5B chaotic at some temperatures of a steered line, fewer than with the full distribution; feeding back 1 to 3 left none. In ordinary sampled generation by the same model, the sensitivity to a small change of the prompt showed no detected difference between settings where the soft map is chaotic and settings where it is periodic.

Other models. The organization recurs in Qwen2.5-1.5B (Figure 56) and 7B, in OLMo-2 7B, in Llama-3.1 8B, whose attractors we also read with sparse-autoencoder features, and in a second checkpoint of the 0.5B architecture, which shares the organization but not the attractors (Appendix H). The fold law of Section 8 recurs at 1.5B and 7B: windows accumulate on the fold of a short cycle, and the nearest four resolved windows meet a criterion fixed before the lines were computed, within 0.12 % at 1.5B and 4.5 % at 7B (a correction of the period selector after the first scoring is disclosed); farther out, the 1.5B windows spread by 10.3 to 14.0 %. An exploratory scan of LFM3-1.7B, in which 13 of the 24 blocks are convolutions rather than attention, also found chaotic cells without and with steering, at coarse resolution and with a single twin offset; it is not counted among the models and families above (Appendix H).

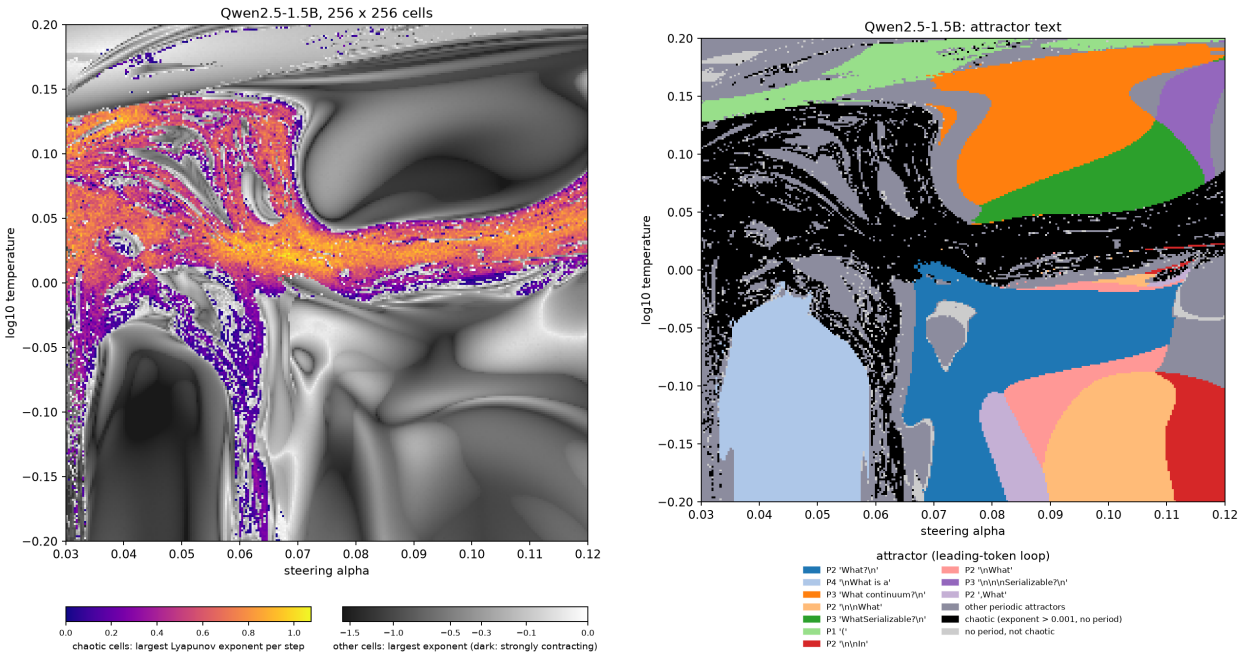

Figure 56: Qwen2.5-1.5B soft-token map steered along the philosophy direction at layer 14, 256^2 float64 cells over $\alpha \in [0.03, 0.12]$, $\log_{10} T \in [-0.2, 0.2]$: largest Lyapunov exponent (left, colored as in Figure 1) and attractor texts (right, as in Figure 11). Several period-2 label regions touch without a gap; as explained in Appendix B, such boundaries may be label switches of one orbit.

14 Chaos and model size

Is the chaos of the loop a small-model effect that larger models leave behind? We measured the unsteered map of fifteen models: Qwen2.5 at seven sizes from 0.5B to 72B parameters, OLMo-2 at 1B, 7B, 13B and 32B, Llama-3.1 at 8B and 70B, Llama-3.3-70B, and the mixture of experts Qwen3.5-35B-A3B, all with the paper's prefix and the settings of Appendix B (1000 transient and 2000 measured steps in float64; chaotic: an exponent above 10^{-3} and no detected period, which is what "positive" abbreviates below). The stored checkpoint values are evaluated in float64 (Appendix B, "Precision"), and the 70B-class models, spread over two GPUs, were recomputed where an earlier two-GPU computation was less precise (Appendix I). The 7–8B and 70B-class models had shown a steep decline of chaos with size. Before measuring the others we fixed three predictions that extrapolate it (`docs/scaling-predictions.md`): within each family, every model above 8B is less chaotic than the family's 7–8B model (P1); the chaotic share does not increase with size above 7B (P2); and at matched next-token entropy, both gain factors of Section 3 of the 14B and 32B models lie at or below those of the 7B model (P3). After a review pointed out that a single prompt and a fixed temperature window could mislead, we froze a fourth test across ten prompts (`docs/prefix-scale-test.md`).

In a fixed temperature window, chaos seems to fade with size. Table 22 and Figure 57 give each model's chaotic share of the temperatures in $\log_{10} T \in [-1, 0.3]$, the range common to all sweeps, and the chaotic length of its whole temperature axis. In two families the largest models are nearly free of chaos in this window. Qwen2.5-7B is chaotic at 12.7 % of the temperatures in the common range, and the 14B, 32B and 72B models at none. Llama-3.1-8B is chaotic at 19 %, Llama-3.1-70B, from the same release and training recipe, at 2.3 %, and Llama-3.3-70B, which differs

from Llama-3.1-70B only in post-training, at none. In the third family the order fails: OLMo-2-32B is the most chaotic OLMo model, at 16.2 % against 6.2 % at 7B, 3.8 % at 1B and none at 13B. (The 32B model is a later release, 0325, than the 7B and 13B models, 1124, so within OLMo-2 size and release are confounded.) The frozen predictions P1 and P2 hold for Qwen2.5 and fail for OLMo-2. Below 7B we predicted no order, and none appears (Qwen2.5 0.5B, 1.5B and 3B: 1.4, 1.5 and 0 %).

What a loop without chaos writes at 70B. Not chaotic does not mean a single repeated token. Over its 128 temperatures ($\log_{10} T$ from -1 to 0.3; 1000 transient and 2000 measured steps), the unsteered Llama-3.3-70B settles on long cycles. At low temperature, at 46 of the 49 temperatures from $\log_{10} T = -1$ to -0.5, it writes a fifteen-token loop without end, `. It'supgrading the infinite possibilities are you'recognized protocols` (the tokens run together as written). At intermediate temperatures period-16 variants end in `established` (15 temperatures) or `acknowledged` (9), and a period-21 variant turns into `... of a fellow AI to another artificial greetings` (8). Higher temperatures give cycles of 24 to 63 steps, and 19 temperatures have neither a period up to 64 nor a positive exponent. Llama-3.1-70B, which differs only in post-training, mostly writes the three-token loop `to another AI` (31 temperatures) or the four-token loop `. It'sup` (21), longer repetitions of `to another AI` with periods 6 to 18, and at four temperatures the fixed point `It`, a single token repeated. Steering changes this picture (Figure 59): along the hostility direction, fixed points fill the plane above the chaotic band and windows of periods 3 to 13, most often 6, lie below it.

Across prompts, the chaos moves to higher temperature instead. The fixed window hides where larger models are chaotic. The frozen prompt-blocked test compares Qwen2.5-7B and -14B on ten system prompts and openers (two for development, eight held out), over $\log_{10} T \in [-0.4, 0.9]$ (48 temperatures) and then over the temperature interval where each model's next-token entropy rises most (24 temperatures), with twins at offsets 10^{-5} and 10^{-7}; a temperature counts as chaotic only when both exponents are positive (Figure 90). It predicted that the 14B model would be less chaotic. Both predictions fail: on all eight held-out prompts the 14B model is chaotic at more of the 48 temperatures (9 to 17, against 0 to 7 for the 7B), and its largest exponent is lower on only three. Every chaotic temperature of the 14B model, on all ten prompts, lies between $\log_{10} T = 0.36$ and 0.90, at and above the temperature where its next-token distribution spreads (0.35–0.65); the 7B model's chaos lies mostly between 0.21 and 0.43, around its own, lower transition (0.21–0.46). On the paper's prompt, 12 of the 48 coarse temperatures of the 14B model are chaotic, from 0.59 to 0.90, and the adaptive fine stage found one more near 0.519; the initial sweep of Table 22 ended at 0.6 and saw only the lower edge (the table now also includes the upper sweep). Within Qwen2.5, then, the larger model is not less chaotic: its distribution stays sharp up to higher temperature, and its chaos sits there. The decline in the fixed window is largely a shift. A second frozen test (`docs/prefix-scale-olmo.md`; stage 1 only, four held-out prompts) asked whether the OLMo-2 order survives other prompts. It does, with small counts: OLMo-2-32B is chaotic at more of the 48 temperatures than OLMo-2-7B on three of the four prompts (4, 5, 3 and 0 against 3, 0, 1 and 3), as predicted, and its chaos again lies somewhat higher ($\log_{10} T = 0.10$–0.35 against -0.01–0.21).

model	params (B)	chaotic, $[-1, 0.3]$ (%)	unres.	sweep, $\log_{10} T$	chaotic dec.	breadth	max λ
Qwen2.5-0.5B	0.5	7/512 (1.4)	1	-1.0 to 0.3	0.018	0.053	+0.28
Qwen2.5-1.5B	1.5	2/130 (1.5)	0	-1.0 to 1.2	0.531	0.191	+0.85
Qwen2.5-3B	3.1	0/130 (0.0)	0	-1.0 to 1.2	0.030	0.060	+0.85
Qwen2.5-7B	7.6	65/512 (12.7)	6	-1.0 to 1.2	0.226	0.249	+0.88
Qwen2.5-14B	14.7	0/130 (0.0)	0	-1.0 to 1.2	0.500	0.191	+1.14
Qwen2.5-32B	32.8	0/130 (0.0)	0	-1.0 to 1.2	0.010	0.091	+0.36
Qwen2.5-72B	72.7	0/128 (0.0)	0	-1.0 to 1.2	0.000	0.014	-0.45
Llama-3.1-8B	8.0	96/512 (18.8)	2	-1.0 to 1.2	0.241	0.634	+0.66
Llama-3.1-70B	70.6	3/128 (2.3)	28	-1.0 to 0.3	0.031	0.307	+0.41
Llama-3.3-70B	70.6	0/128 (0.0)	19	-1.0 to 0.3	0.000	0.225	-0.27
OLMo-2-1B	1.5	5/130 (3.8)	0	-1.0 to 1.2	0.050	0.322	+0.23
OLMo-2-7B	7.3	32/512 (6.2)	3	-1.0 to 1.2	0.083	0.197	+0.30
OLMo-2-13B	13.7	0/130 (0.0)	0	-1.0 to 1.2	0.000	0.131	-0.0007
OLMo-2-32B	32.2	21/130 (16.2)	0	-1.0 to 1.2	0.252	0.694	+0.62
Qwen3.5-35B-A3B	35.0	0/128 (0.0)	15	-1.0 to 0.3	0.000	0.143	-0.27

Table 22: Unsteered soft-token maps of fifteen models (float64; 1000 transient and 2000 measured steps; one twin offset). Chaotic: exponent above 10^{-3} and no period up to 64; unres.: temperatures in $[-1, 0.3]$ with neither a detected period nor a positive exponent. Chaotic decades: the grid-weighted extent of the temperature axis labeled chaotic over the whole sweep, each temperature weighted by its grid spacing (an estimate of the sampled extent; models whose sweep stops at 0.3 have incomplete coverage); max λ: the largest exponent on the sweep, rounded to two decimals (OLMo-2-13B: -6.7×10^{-4}); breadth: decades in which the mean next-token entropy on the orbit lies between 1 and 10 nats. The Qwen2.5-72B sweep joins two grids that share $\log_{10} T = 0.3$, counted once; for the 7–8B models the share in $[-1, 0.3]$ comes from their 512-temperature sweeps, on which the frozen predictions are scored, and the whole-axis columns also include sweeps over $[-1, 0.6]$ and $[0.61, 1.2]$ (64 temperatures), as for the other models swept to 1.2. Positive cells found later on finer lines (Qwen2.5-32B, -72B; Llama-3.3-70B) are not in these regular-grid counts.

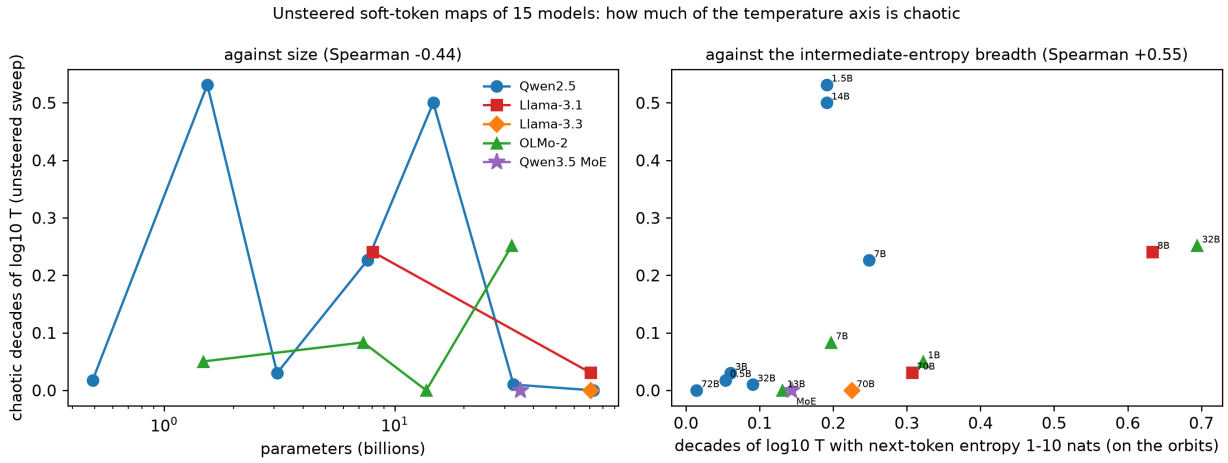

Figure 57: How much of the temperature axis is chaotic, for the fifteen unsteered maps of Table 22: against parameter count, by family (left), and against the breadth of the intermediate-entropy regime (right). Spearman correlations over the fifteen models. Each model's sweeps cover the range listed in Table 22; those of the two Llama-70B models and the mixture of experts end at $\log_{10} T = 0.3$, so their values are lower bounds. The breadth is measured on the loop's own orbits; the loop-free test of the same idea is in Appendix I.

What sets it: the gain of the loop. The gain split of Section 3 shows which factor removes or keeps the chaos (Figure 91). We compare medians of the two factors in bins of the mean next-token entropy on the orbit; the bins hold 1 to 17 of each model's 24 temperatures, so single medians are noisy. Within one training recipe, from Llama-3.1-8B to Llama-3.1-70B, the readout factor changes by at most 0.3 per step in every bin above one nat (below one nat the 8B model has a single temperature), while the network factor is lower by 0.5 to 1.3 per step between 1 and 10 nats (6–10 nats: median -0.98 from 3 temperatures against $+0.28$ from 4) and by only 0.18 above 10 nats: in these two checkpoints the larger model damps changes of its fed-back token. Llama-3.3-70B, which differs from Llama-3.1-70B in post-training, has a much lower readout factor below 3 nats (1–3 nats: -2.33 from 7 temperatures against -0.10 from 6) and a similar one above 6 nats. In Qwen2.5, both factors of the 14B and 32B models lie below the 7B's in every bin between 1 and 10 nats, as the frozen prediction P3 states (the 6–10 nat bin only with the 7B's wider gain split, from one temperature), and both factors of the 72B lie below the 7B's in every bin. Above 10 nats, where the distribution is nearly uniform, the 14B and 32B models instead have the larger network factor ($+2.98$ and $+1.17$ against $+0.19$), and those are the bins of their chaotic temperatures above $\log_{10} T = 0.3$. Below 3 nats the readout factor dominates at 32B and 72B, because the next-token distribution stays sharp up to high temperature (for the 72B, entropy at most 0.15 nats up to $T = 2$). In OLMo-2, P3 fails, and the readout factor tracks the chaos: OLMo-2-13B, with no chaotic temperature, has a lower readout factor than OLMo-2-7B in all three bins between 1 and 10 nats (-1.0 to -1.4 against $+0.06$ to -0.72 per step) although its network factor is higher in two of them, whereas OLMo-2-32B, the most chaotic OLMo model, has a higher readout factor than the 7B in all three ($+0.10$ to $+0.35$), positive as in the chaotic 7–8B models. These are descriptive comparisons of trained checkpoints, not interventions on size or post-training: the split is a decomposition along the evolving tangent direction, and equal entropy does not imply equal embedding covariance. Entropy clamping tests at 72B whether forcing an intermediate entropy suffices to destabilize the loop. If its loop were stable only because its distribution jumps from sharp to flat, holding the distribution in between should destabilize it. We therefore choose each step's temperature so that the next-token entropy equals a target and sweep the target instead of the temperature (entropy-clamped decoding; the map stays deterministic, and the entropy holds to 10^{-14}). The temperature then depends on the state, $T = \tau(x)$ with $H(x, \tau(x)) = H_\star$ for the entropy H, so the derivative of the clamped map $\hat{F}(x) = F(x, \tau(x))$ gains a rank-one term, $D\hat{F} = D_x F - \partial_T F \, D_x H / \partial_T H$. We therefore take its exponents from twins of the clamped map itself, temperature solve included, and use no fixed-temperature derivative or gain split for it. The safeguarded solve is only piecewise smooth, but on these orbits one-step central differences at 10^{-5} to 10^{-7} agree with those at 10^{-4} within 6×10^{-6}. Held at any of 16 targets between 1 and 11 nats, Qwen2.5-72B has only periodic orbits, the echo of period 2 on the paper's prompt and periods 1 to 19 on a held-out prompt, with twin exponents between -0.08 and -4.5 that agree within 0.004 across offsets 10^{-5} to 10^{-7}, whereas Qwen2.5-0.5B under the same clamp is chaotic at 6.1 nats. The clamped 72B loop thus stays stable even where its readout is spread. The clamp changes the map, so this is consistent with damping by the network in the original map but does not identify it. The pretrained checkpoint is also mostly non-chaotic. A test frozen before its data (`docs/prefix-scale-72b-base.md`) predicted that the pretrained base model of the same size would be chaotic at three or more temperatures more than Qwen2.5-72B-Instruct on at least three of the four held-out prompts. It fails on all four: the base model is chaotic at 1, 0, 0 and 0 of the 48 temperatures (and at 2 on the paper's prompt; computed after the precision fix, with 500 transient and 1000 measured steps, half the Instruct model's horizon), although its next-token distribution spreads over 2.5 to 5.5 times the temperature range (intermediate-entropy breadth 0.28 to 0.61 decades against 0.11 to 0.22) and from lower temperatures on. At 72B, then, the sharpness that post-training adds does not explain the stability, and post-training is not necessary for it;

the failed prediction does not quantify post-training's other effects. The gain split points to the network's damping of the fed-back token (the 72B's network factor lies below the 7B's in every entropy bin), and the breadth that orders the fifteen models (Figure 57, right) fails for this pair. The damping holds when the whole context is kept as well: in Soft Thinking's growing context (Appendix G), an early perturbation grows at none of eleven steered cells, six of them chaotic as one-token maps, and at one of 48 unsteered temperatures on two prompts, transiently.

Where the network damps. The network factor can be resolved through the stack. Let h_l be the residual stream at the soft token's position entering block $l+1$, and δh_l its change. The factor then splits exactly into three terms: the growth of the change through the stack, $A_L = \log(\|\delta h_L\|/\|\delta x\|)$; minus the growth of the residual stream itself, $N_L = \log(\|h_L\|/\|x\|)$; plus the final normalization. The difference $A_l - N_l$ is the log of the relative change after l blocks (Appendix B).

We measured this split along the unsteered loop of fourteen models at 24 temperatures each (Figure 58). At 0.3 to 6 nats, the entropies where the loop's chaos occurs at low temperature, the median network factor falls with size in two of the three families.

- In Qwen2.5 it falls from -0.12 at 0.5B to -1.99 at 72B, though not monotonically (-1.80 at 3B, -0.85 at 7B).
- In Llama-3 it falls from -0.14 at 8B to -1.50 and -1.30 at 70B.
- In OLMo-2 it does not fall: $+0.11$, -0.52, -0.79 and -0.36 from 1B to 32B. OLMo-2-32B, the family's most chaotic model, is the least damped of its three larger models.

Qwen2.5-72B is the only Qwen2.5 model whose network factor stays negative at nearly uniform distributions: -1.12 above 10 nats, against $+0.22$ to $+2.97$ for the 1.5B to 32B models, as its stability under entropy clamping suggests.

The profiles through depth show where the difference arises (Figure 58c). Each block adds its output to the residual stream, so the stream grows from block to block, and the change that the fed-back token makes in it grows too. Near 1 nat, the two grow at nearly the same rate in the small models: entering the last block, the log relative change $A_l - N_l$ is -0.32 (Qwen2.5-0.5B) and -0.51 (Llama-3.1-8B), so the change has shrunk relative to the state only by factors of 1.4 and 1.7. In the large ones the stream grows much faster: $A_l - N_l$ is -2.51 (Qwen2.5-72B) and -2.08 (Llama-3.1-70B), factors of 12 and 8. Their blocks keep writing a growing state that depends little on the fed-back token.

Sublayer knockouts hold one attention or MLP output at the soft token at its unperturbed value. They find no component that does the damping: at each model's knockout temperature nearest 1 nat (entropies 0.12 to 1.61 nats), removing any single sublayer raises the network factor by at most 0.44 nats in any of the fourteen models. The amplification moves instead.

- In the Qwen2.5 models from 1.5B to 32B, the first attention layer amplifies the change: removing it lowers the factor by 1.67 to 3.74 nats near 1 nat.
- In the 0.5B model, the first MLP does (2.40).
- In Qwen2.5-72B and in both 70B Llama models, no sublayer contributes more than 0.46.

The larger models' damping is therefore not added by a component. Rather, the early amplification, by the first attention layer or MLP, that makes the change grow about as fast as the residual stream in the smaller models is weaker or absent. These are descriptive comparisons of trained checkpoints within three families, and the OLMo-2 exception shows that training can break the trend.

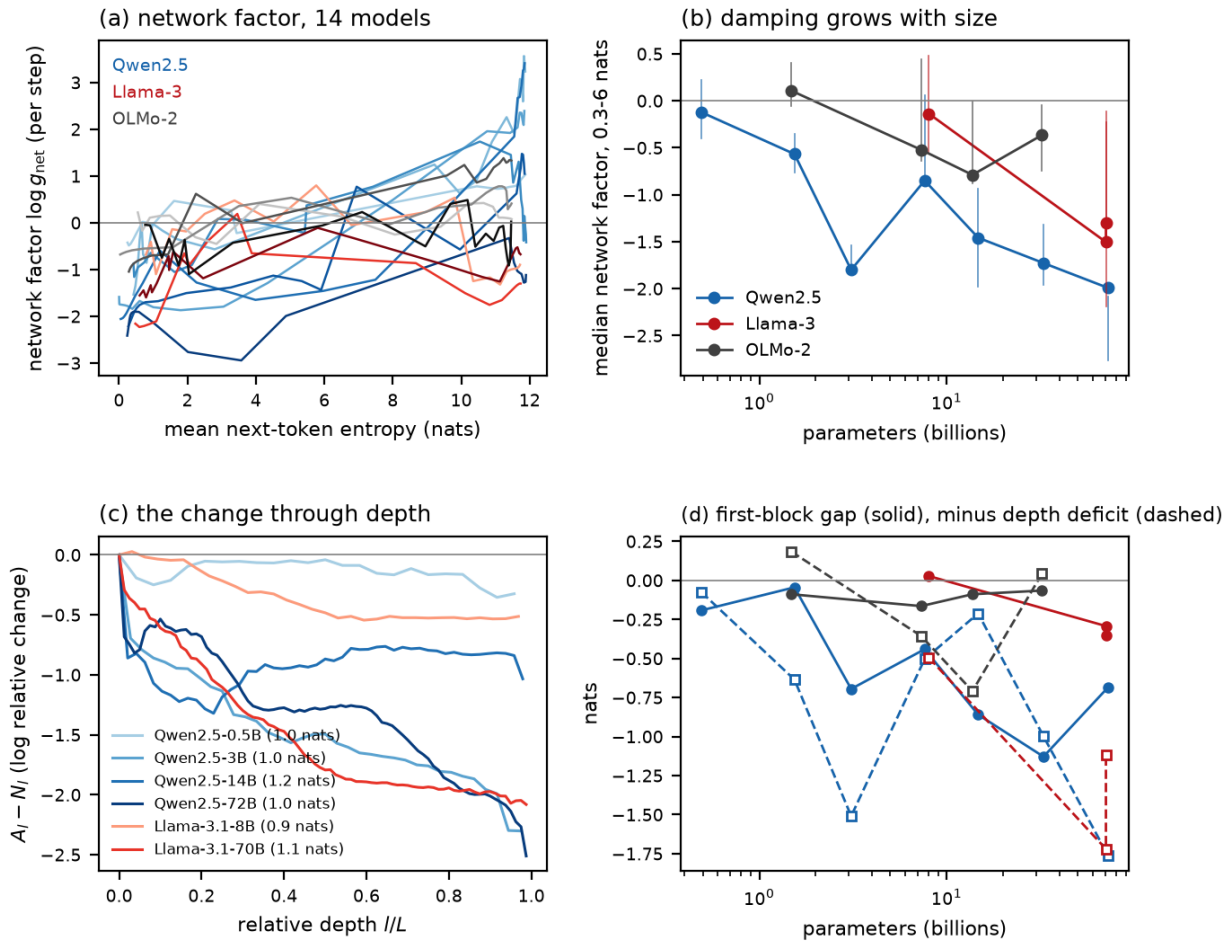

Figure 58: Where the network damps a change of the fed-back token. Unsteered loop on the paper's prompt, 24 temperatures per model, float64; a twin at relative offset 10^{-7} along the evolving tangent; 500 transient and 500 measured steps. (a) Network factor against the mean next-token entropy on the orbit, for fourteen models (Qwen2.5 blue, Llama-3 red, OLMo-2 gray; darker for larger models). (b) Median network factor between 0.3 and 6 nats against the number of parameters; bars give the range of the medians in the bins 0–1, 1–3 and 3–6 nats. (c) The log relative change $A_l - N_l$ after l blocks against relative depth, at each model's temperature nearest 1 nat. (d) The two parts of the final relative change at that temperature: the gap after the first block, $A_1 - N_1$ (solid), and the deficit from the second block on, $(N_L - N_2) - (A_L - A_2)$, drawn with a minus sign (dashed); the increment through the second block itself is not shown. The profiles use block inputs, so the last value is the relative change entering the last block.

Mixtures of experts. A mixture-of-experts layer sends each token to a few of many expert networks, those with the largest router scores, so its map is smooth only between routing boundaries. Two questions follow: whether routing switches drive the measured exponents, and whether a model that is large in total but small per token behaves like a small or a large dense model (Figure 88, Appendix H). We followed the unsteered loops of Qwen3.5-35B-A3B (8 of 256 experts per layer; about 3B of its 35B parameters are active per token) and OLMoE-1B-7B (8 of 64; about 1B of 7B). Along them, a twin at relative offset 10^{-7} never chose different experts from the base orbit, at any step, layer or temperature. Freezing its routing to the base orbit's left the network factor unchanged, so the exponents measure the smooth map. The orbits do pass close to routing boundaries. At each step we recorded the smallest gap between the eighth and ninth largest router logits over all layers; its median along a loop lies between 1.1×10^{-4} and 1.75×10^{-2}, and its minimum is about 3.38×10^{-7}. Rounding errors of bfloat16 inference (relative precision about 4×10^{-3}) can cross gaps this small, so in low precision routing can add jumps that the float64 exponents do not contain.

On size, the mixture of experts behaves like a large dense model, not a small one. Two dense models of its family share its hybrid attention: Qwen3.5-4B, about its active size, and Qwen3.5-9B. Both are chaotic. Unsteered, they are chaotic at 19 and 10 of 128 temperatures; steered, at 15 % and 7 % of a philosophy map and at 32 % and 16 % of a hostility map. The mixture of experts is chaotic at none of the 128 temperatures and none of the philosophy cells, and irregular at 5 % of the hostility cells. Its median network factor between 0.3 and 6 nats is -1.90, against $+0.05$ and -1.41 for the two dense models and -1.73 for the dense Qwen2.5-32B. Within this family, chaos falls and damping grows from 4B to 9B to the mixture of experts, which damps more than a dense model with three times its active parameters. Routing a token through about 3B parameters therefore does not make the loop behave like a 3B model.

What this implies for larger models. Size does not set this chaos in a simple way. In Qwen2.5, the 14B model's chaos did not vanish but moved to higher temperature, to where its next-token distribution spreads; in OLMo-2 the largest model was the most chaotic; two 70B Llama checkpoints with different post-training gave different results on the same grid; and at 70B, chaos that passes a three-offset test remains, rarer, with and without steering (Appendix I). The shift does not continue to Qwen2.5-72B. A prompt-blocked test frozen before its data (`docs/prefix-scale-72b.md`) predicted that the 72B model would be chaotic at fewer of the 48 temperatures than the 14B on at least three of four held-out prompts. It is chaotic at 0, 1, 0 and 0 of them, against 14, 9, 11 and 15, so the prediction holds on all four. These lines ran before the precision fix (1000 transient and 2000 measured steps, twins at 10^{-5} and 10^{-6}, chaotic only when both exponents are positive); no temperature on them is positive at only one offset, and the one chaotic temperature ($\log_{10} T = 0.319$, exponents 0.35 and 0.34) is chaotic again when recomputed in float64 throughout, with an automatic-differentiation tangent exponent of 0.34. With the paper's prompt, where it showed no robust chaos up to $\log_{10} T = 1.2$, that makes five prompts on which the 72B model is nearly free of unsteered chaos. Its next-token distribution jumps from sharp to nearly uniform within less than the grid spacing (0.03 decades), so a narrow chaotic band at the jump would be missed by the grid. Fine lines through the jump, computed in float64 throughout with automatic-differentiation tangent exponents, find only periodic orbits on the paper's prompt and on three of the four held-out prompts, and on the fourth a band of chaos about 0.02 decades wide just above the jump ($\log_{10} T = 0.319$–0.336; exponents up to 0.40, which twins at 10^{-5} to 10^{-7} reproduce within 0.004), at nearly uniform distributions, as the 14B model's chaos but far narrower. A fixed temperature window is therefore the wrong yardstick, and so is size alone. A model's chaos follows the gain of its loop where its own next-token distribution is spread, no single forward pass we tested predicts it (Appendix I),

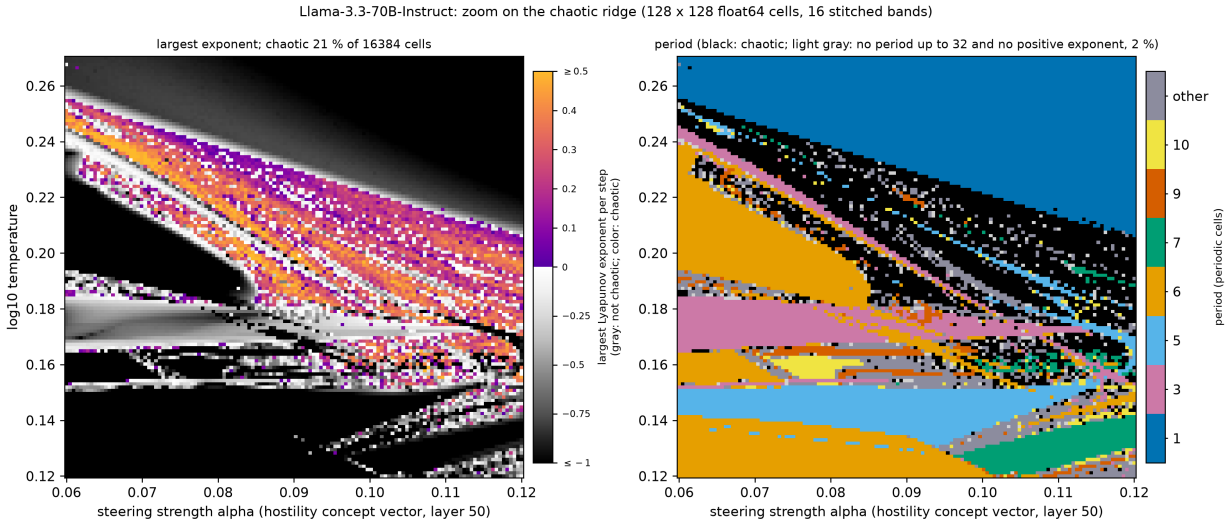

Figure 59: Chaos at 70B parameters: the soft-token map of Llama-3.3-70B-Instruct steered along its hostility concept vector at layer 50, 128×128 float64 cells over $\alpha \in [0.06, 0.12]$, $\log_{10} T \in [0.12, 0.27]$ (300 transient and 300 measured steps; sixteen bands of eight temperatures computed separately with identical vectors and joined). Left: largest Lyapunov exponent, gray where the cell is not chaotic and in color where it is; 3410 of the 16,384 cells (20.8 %) are chaotic (largest exponent 0.64). Right: period of the attractor reached from the fixed start (black: chaotic; bluish gray: periods without a color of their own; light gray: neither a detected period up to 32 nor a positive exponent, 1.7 % of the cells). The chaotic band's center falls from $\log_{10} T \approx 0.23$ at $\alpha = 0.06$ to ≈ 0.19 at $\alpha = 0.12$, and periodic windows thread it, most often of periods 3, 5, 6 and 7 and up to 13; fixed points lie above it.

and to find it in a larger model one has to run the loop there, across prompts. For frontier models this means two things: their unsteered chaos may be rare, if, like the largest models here, their distributions spread over a narrow range of temperature; and steering still drives both 70B-class models that we steered, Llama-3.3-70B and Qwen2.5-72B, into chaos (Appendix I). Resolved through the stack, the damping grows with size in Qwen2.5 and Llama-3 because in the larger models the change that the fed-back token makes in the residual stream grows much more slowly than the stream itself (Figure 58). If larger models of these families continue the trend, their unsteered loops need stronger token competition to become chaotic, which steering can still supply; this is an extrapolation from checkpoints up to 72B, and OLMo-2 does not follow it. Figure 59 shows it at 70B parameters: in a 128×128 float64 scan of Llama-3.3-70B-Instruct steered along its hostility direction, computed after the precision fix, 20.8 % of the 16,384 cells are chaotic, in a band threaded by periodic windows, as in the 0.5B model's maps; a test of 24 sampled cells at three twin offsets and an automatic-differentiation tangent confirms 11 of the 12 chaotic and 5 of the 6 periodic labels (Appendix I).

15 A second system: recurrent depth

Huginn-0125 iterates a recurrent block at fixed input tokens [3]. For the GSM8K [78] test question on which Geiping et al. [3] show latent orbits ("Claire makes a 3 egg omelet every morning for breakfast…", here without the chat template), the native loop converges at every position to a fixed point whose leading Jacobian pair is $0.588 \pm 0.771i$ (modulus 0.970). To probe its stability we relax the iteration to $h \leftarrow (1-\lambda)h + \lambda F(h)$: $\lambda = 1$ is the model's own loop, $\lambda < 1$ takes shorter

steps and $\lambda > 1$ overshoots, and at the fixed point the relaxed loop has multipliers $1 - \lambda + \lambda\mu$, so λ acts as a bifurcation parameter we can set. For this prompt this predicts that the pair crosses the unit circle at $\lambda_c = 2(1 - \mathrm{Re}\,\mu)/|1 - \mu|^2 = 1.0775$, transversally and away from the strong resonances (at 0.156 of a revolution per step), consistent with a Neimark–Sacker bifurcation. Below it the latent state decays; above it the state settles on a closed curve. A fit of the squared amplitude gives $\lambda_c = 1.0778$, with $A/\sqrt{\lambda - \lambda_c} = 0.52$–$0.54$. A stable circle that grows like the square root of the distance is consistent with a supercritical Neimark–Sacker bifurcation; we did not compute its first Lyapunov coefficient, so its nondegeneracy is not established. Causal attention makes the loop block lower-triangular in token positions. The first position to oscillate is the "." of the prompt, and all later positions follow. On the invariant circle the motion turns by 0.156 of a revolution per step, matching the argument of the critical pair. The rotation number changes smoothly, by about 10^{-4} over $\lambda \in [1.079, 1.092]$, and no return within 60 steps was found at 10^{-3} of the radius. This is consistent with quasi-periodic motion, although locking in narrow high-order tongues cannot be excluded. From $\lambda \approx 1.094$ the circle deforms or a second frequency appears, near where a second complex pair destabilizes the fixed point ($\lambda = 1.0814$). Fixed-point and Jacobian stability of looped models has been studied before [18]; what this adds is the predicted crossing, its amplitude law and the rotation number. We have not established chaos in this loop. Exploratory float32 maps of the relaxed loop with an added steering direction, on a short arithmetic prompt, found approximate two-step repetition at some settings and positive finite-window growth of separations at others. They establish neither sustained chaos nor its absence (`docs/huginn-status.md`).

The reported orbits. Because the first position to oscillate is the full stop, the critical pair belongs to it and to the positions after it. It cannot shape the trajectories of "makes" and "3", on which Geiping et al. [3] show orbits: these come earlier, and causal attention leaves earlier positions independent of later ones. Since every position converges at the native step, any orbit-like motion on this question is transient in our runs; the native loop fails to settle where its fixed point is unstable, as for 6 of the 30 ARC questions below.

The spectrum predicts the relaxation response. With the full instruction template, the GSM8K prompts' fixed points look different: the leading multipliers surround the origin, a real 0.914, complex pairs of modulus 0.85–0.89 at many angles and negative real values near -0.8, with no near-critical pair. Because the prefill Jacobian is block lower-triangular in positions, its spectrum is the union of per-position blocks, and template positions contribute identical eigenvalues to every prompt. Such a spectrum makes three predictions before any task is scored. First, within the computed spectral estimates of these prompts, no fixed relaxation can speed convergence much. With leading multipliers spread around the origin at many angles, any λ that shrinks one group enlarges another, and the best $\lambda^* = 1.03$–1.06 lowers the rate from 0.914 to only 0.908–0.911. (A finite computed spectrum does not show that no other polynomial scheme could do better; Anderson acceleration, tried below, did not.) Second, under-relaxation slows convergence ($\lambda = 0.55$: predicted 0.953, measured 0.951 on one prompt). Third, over-relaxation destabilizes the negative modes first, at $\lambda_c = 2\,\mathrm{Re}(1 - \mu)/|1 - \mu|^2 = 1.074$–$1.126$ (five prompts), in a flip or a Neimark–Sacker bifurcation close to -1. The approach to it is visible: on the prompt with $\lambda_c = 1.074$, the measured rate at $\lambda = 1.05$ is 0.955 against a predicted 0.956 (native 0.864). Elsewhere the spectral radius bounds the rate measured over the accessible window from above, by up to 0.05, because the slowest modes carry little amplitude there. Independently, Wang et al. [79] report a Huginn GSM8K drop from 19.41 % to 8.34 % at a constant step of 1.2; at 0.8 GSM8K is 19.79 %, while their four-task mean falls slightly. Their trajectory-based scheduling analysis does not compute the prompt-specific spectral

thresholds studied here.

Does the task notice? (development data). On ARC-Challenge [80] questions scored by log-likelihood, where one prefill per answer choice is exactly the analyzed recurrence, the spectra again predict convergence. Rates agree to 0.002–0.006 where the slowest mode dominates, and beyond each question's λ_c (1.075–1.149 for four questions) the residual stops decaying. For 6 of 30 development questions the native fixed point is itself unstable ($\rho(1) = 1.009$–1.273; the fixed point, located with a relaxed transient, has a real multiplier beyond -1 or a complex pair outside the circle). The native loop then oscillates (relative amplitude 2×10^{-2} for one of them), and relaxation to 0.55 or 0.75 restores convergence to 10^{-6}. For the other 24, no unstable mode was detected in the computed spectral estimates up to $\lambda = 1.020$–1.149 (median 1.095); partial Ritz spectra cannot exclude an undetected unstable mode. Question by question, an estimated spectral feature ranks the question's residual. This feature, $\hat{\rho}_q(\lambda)$, is the largest $|1 - \lambda + \lambda\mu|$ over the computed Ritz values of the question's gold-answer prefill, and its Spearman correlation with the residual is 0.73 at $\lambda = 1$ and 0.75 at 1.10 over the 30 questions (development data). In this historical analysis the residual was the all-choice aggregate entering the 32nd recurrence call, so this is a predictive association across related maps and precisions, not a same-state contraction identity. The item-level part of the frozen test below uses matched gold-row token IDs and a per-row residual with an explicit clock.

Accuracy, however, tolerates the instability (300 training questions, 32 steps): it is 0.367 native, 0.387 at $\lambda = 1.10$ and 0.373 at 1.15, where the latent state oscillates with relative size 0.08 and 0.23. Only at 1.25 (size 0.68) does it fall, to 0.300. Yet the oscillation does reach the answers. Reading every step from 24 to 32 on 100 questions, the margin between the correct choice and its best rival moves by 0.11 nats per step natively, 0.37 at 1.10 and 0.95 at 1.15, in alternating directions, and the chosen answer changes within these nine steps for 12 %, 16 % and 33 % of the questions. The changes go both ways, so the accuracy at any one step stays the same on average. Under-relaxation instead costs 4–5 points at 8–16 steps and catches up by 32 (a slower transient), and Anderson acceleration neither speeds the transient nor gains accuracy. Testing the accuracy consequence therefore needs a task in which such step-to-step changes can compound.

GSM8K generation is such a task. With eight chain-of-thought examples in the lm-evaluation-harness format, greedy decoding and 32 recurrence steps per token (development problems 0–119, every arm on the same problems and initial states), accuracy is 0.375 natively, 0.425 at $\lambda = 1.05$ and at 1.10, 0.29 at 1.15 (17 problems lost and 7 gained, one-sided McNemar $p = 0.03$) and 0.16 at 1.20 (29 lost, 3 gained, $p = 10^{-6}$). Under-relaxation to 0.8 costs nothing at 32 steps (0.41). The break lies between 1.10 and 1.15. It is not located at a computed stability edge of the scored prompts: the GSM8K thresholds above ($\lambda_c = 1.074$–1.126) belong to five zero-shot prompts in the instruction template, while in the eight-shot format scored here the fixed point is already unstable at the native step (three development problems, float32: a real multiplier -1.158, stable only for $\lambda < 0.927$). Generation breaks earlier than log-likelihood scoring, consistent with answer changes that compound over tokens. The collapse at 1.2 is the one Wang et al. [79] reported.

A frozen test on problems not scored before it was frozen confirms the cost of over-relaxation. Its design and analysis were pushed before any test item was scored, after two rounds of independent review. On 200 GSM8K test problems, accuracy is 0.330 natively, 0.395 at $\lambda = 1.05$, 0.255 at 1.15 and 0.095 at 1.20. The two pre-stated predictions pass: accuracy declines at 1.20 (-23.5 points; 50 problems lost and 3 gained, exact one-sided McNemar $p = 3 \times 10^{-12}$), and 1.05 loses no more than four points (difference +6.5, one-sided 95 % lower bound +2.5). At 1.15 accuracy falls by 7.5 points ($p = 0.01$, reported descriptively). One of the 200, item 18, is the question of this section's first paragraph, which the dynamics runs had used as a bare prompt, without its answer, before the test

was frozen; without it, accuracy still declines at 1.20 (49 problems lost and 3 gained, $p = 5 \times 10^{-12}$) and the lower bound at 1.05 stays at +2.5 points. Development and test agree closely (Figure 60).

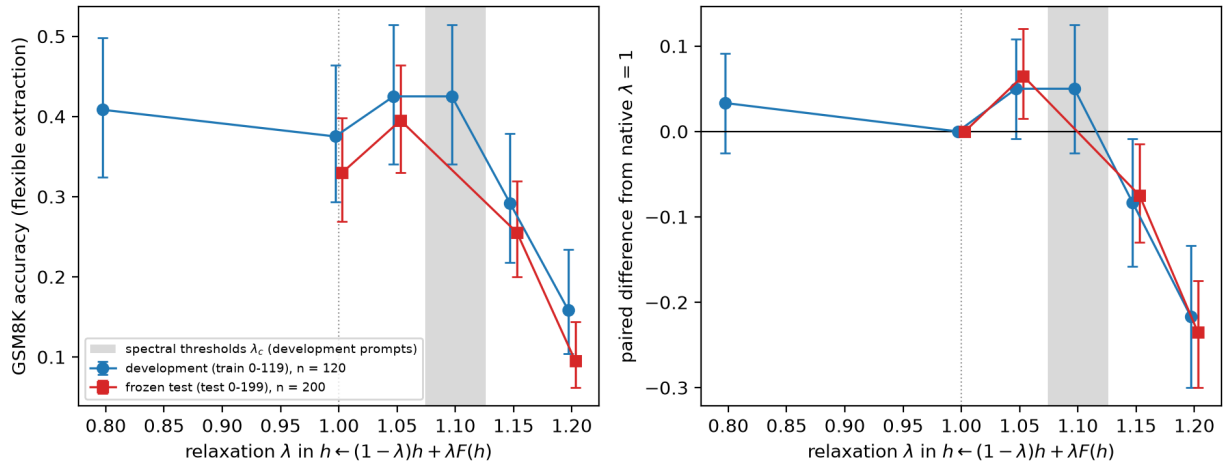

Figure 60: Huginn-0125 chain-of-thought accuracy on GSM8K against the relaxation strength λ (eight examples, greedy decoding, 32 recurrence steps per token): development problems (blue) and the frozen test on held-out problems (red; design and analysis fixed before any test problem was scored). Gray: the spectral thresholds $\lambda_c = 1.074$–1.126 computed from the fixed points of five zero-shot development prompts; in the scored eight-shot format the fixed point is already unstable at $\lambda = 1$ (see text), so the figure shows the cost of over-relaxation, not a fall at a computed edge. Left: accuracy with Wilson 95 % intervals. Right: paired difference from the native loop with paired-bootstrap 95 % intervals.

The same frozen design tested the item-level prediction on 64 held-out ARC-Challenge questions. The first 16 spectra preceded the read-out generation; the remaining ones were completed afterward under a recorded amendment, with the producer attesting that read-out outcomes had not yet been inspected, and duplicate copies were resolved by a rule fixed in advance. Every question passed the admission gates (fixed-point residual at most 2.6×10^{-6}; identical gold-row token IDs in every arm); for 13 questions one of the three Arnoldi sets was partial, which the design admits for an estimated feature. The pre-stated statistic, the mean of the within-arm Spearman correlations between $\hat{\rho}_q(\lambda)$ and the gold row's residual entering the 32nd update, is 0.701 (95 % bootstrap interval $[0.575, 0.780]$; 0.68, 0.73, 0.73 and 0.67 at $\lambda = 1$, 1.05, 1.08 and 1.10); the frozen primary gate, the interval's lower end above 0.3, is met. On unseen questions, the fixed point's estimated spectrum therefore ranks each question by the size of the gold row's residual entering update 32; a residual is not in general a distance to the fixed point. Descriptively, the share of questions with a detected unstable mode rises from 8 % natively to 34, 69 and 84 % at $\lambda = 1.05$, 1.08 and 1.10. An independent implementation reproduced both parts of the frozen test from the saved data.

16 Two loops in conversation

Two instances of a language model left to talk with each other drift into recurring states [14]. The best-known case is the "spiritual bliss" attractor of two Claude Opus 4 instances, whose open-ended conversations end in repeated exchanges of gratitude and silence [81]. The soft-token map poses the deterministic version of the question: can two loops that are each chaotic on their own lock each other into a periodic orbit? Agents A and B are Qwen2.5-0.5B-Instruct with the conversation prompt of a sampled self-dialogue study that we report separately, each with its own steering

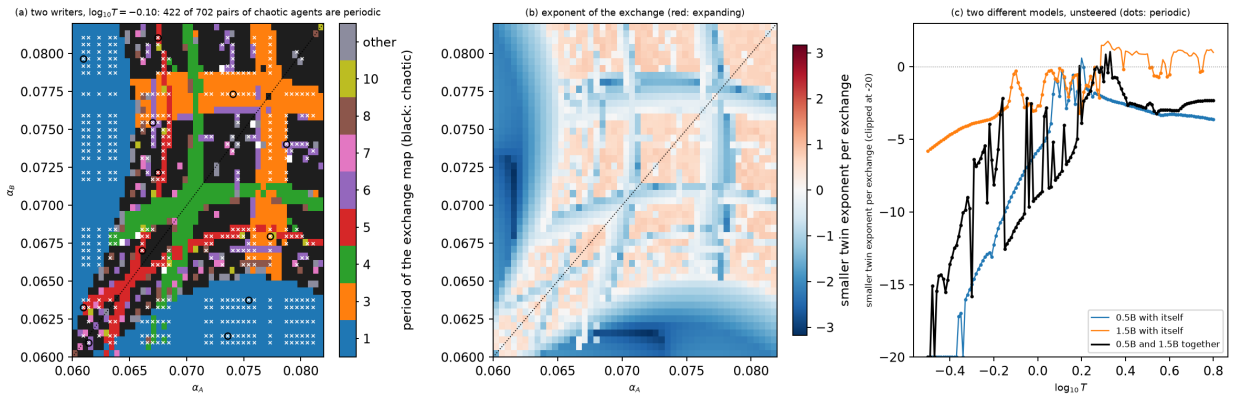

Figure 61: Two loops writing one text. (a) Agents A and B (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, strengths α_A and α_B) alternate soft tokens at $\log_{10} T = -0.10$: period of the exchange map $G = F_B \circ F_A$ per cell (black: chaotic). The dotted diagonal holds each agent alone; white crosses mark pairs that are periodic although both agents alone are chaotic, and circles the twelve such pairs verified by Newton's method. (b) The smaller twin exponent of G per exchange on the same cells. (c) Two different models, unsteered: the smaller twin exponent per exchange of Qwen2.5-0.5B and of Qwen2.5-1.5B each writing with itself, and of the two writing together; dots mark periodic cells.

strength along the philosophy direction at layer 12. They write one reply together, alternating soft tokens, $x_{n+1} = F_A(x_n)$ and $x_{n+2} = F_B(x_{n+1})$, so one exchange is the map $G = F_B \circ F_A$. Equal agents reproduce the single loop bitwise in float64, so on a map over (α_A, α_B) the diagonal holds each agent writing alone (Figure 61).

At $\log_{10} T = -0.10$, over $\alpha_A, \alpha_B \in [0.060, 0.082]$ (48×48 cells), 27 of the 48 agents are chaotic alone and 21 periodic. Of the 702 pairs whose two agents are both chaotic alone, 422 are periodic together; of the 420 pairs of two periodic agents, 91 are chaotic together. Writing together favors order (70 % of pairs are periodic, against 44 % of single agents), though some chaos is inherited: pairs of chaotic agents are periodic less often than pairs in general. The exponent map is plaid, as if each agent contributed its own stability, and where one agent is weakly steered ($\alpha \lesssim 0.064$) and the other is not, the pair settles into a fixed two-token exchange. Twelve entrained pairs spread over the map all converge under Newton's method to attracting primitive orbits of G (residuals at most 2×10^{-13}, leading multipliers 0.0008 to 0.79 in modulus, exchange periods 1, 3, 5 and 6). Their texts show how two chaotic writers share a loop. At period 1, A writes "is" and B "it", the single loop's 2-cycle split between two writers. At period 5, B answers "is" every time while A cycles through "This", "there", "it" and "a". That alternating two chaotic maps can produce order is known for quadratic maps [82]; here it happens in a language model's own dynamics.

Two different models behave alike. Qwen2.5-0.5B and 1.5B share a tokenizer, so a message can be passed as a distribution over the shared vocabulary and embedded by each model with its own matrix. Unsteered, the 1.5B writing with itself is chaotic at 41 of 131 temperatures ($\log_{10} T$ from 0.29 to 0.80, exponents 0.21 to 1.76 per exchange), the 0.5B at one. Writing together they are periodic at 39 of those 41 temperatures, mostly in a fixed two-token exchange: the stable partner sets the rhythm. A dialogue framing behaves differently. With each one-token message placed in a user turn and the first token of the reply as the answer, every map we computed is periodic, for one agent alone and for every pair. The reply is decided by the replier's own prompt and steering, and the partner's message barely moves it; mutual locking of the kind seen in long conversations would need longer messages than this frame provides.

Synchronization of coupled loops. Alternating writers are one way to couple two loops. The classical way to couple identical chaotic maps is to mix their outputs: each copy of the loop feeds back a fraction $1-c$ of its own soft token and a fraction c of the other's,

$$x'_A = (1-c)F(x_A) + cF(x_B), \qquad x'_B = (1-c)F(x_B) + cF(x_A). \tag{24}$$

The state where both copies agree, $x_A = x_B$, is then invariant, and on it each copy runs the single loop. The difference $u = x_A - x_B$, which points across the synchronized state (transverse to it), evolves exactly as $u' = (1-2c)\,[F(x_A)-F(x_B)]$, which near synchrony linearizes to $u' = (1-2c)\,DF\,u$. At each step the loop's derivative stretches the difference, on average by the factor e^{λ_1}, where λ_1 is the leading Lyapunov exponent of the single loop, and the mixing multiplies it by $1-2c$. A small difference therefore grows on average at the rate $\lambda_\perp = \lambda_1 + \ln|1-2c|$ per step, the transverse Lyapunov exponent. Without coupling, $\lambda_\perp = \lambda_1 > 0$ and the copies drift apart. As c rises from 0 toward 1/2, the factor $1-2c$ falls from 1 toward 0, and $\lambda_\perp$ falls with it. It crosses zero at the critical coupling

$$c^* = \tfrac{1}{2}(1 - e^{-\lambda_1}), \tag{25}$$

where the mixing's contraction per step, $1-2c^* = e^{-\lambda_1}$, exactly cancels the loop's average stretching per step, e^{λ_1}. For couplings between c^* and 1/2 (and, by the symmetry $c \to 1-c$, up to $1-c^*$) typical transverse perturbations decay along the synchronized orbit: the transverse-stability condition for our symmetric output coupling, derived above; it follows the Lyapunov-stability framework of Fujisaka and Yamada [83], and [84] treats synchronization of chaotic systems more broadly. This local, typical-orbit condition need not imply synchronization from every start: other invariant sets inside the attractor may be transversely unstable at the same coupling, and it says nothing about which attractor a pair reaches. Near c^* the theory predicts on-off intermittency: long synchronized stretches broken by bursts, with stretch lengths distributed as a power law of exponent $-3/2$ [85, 86]. Above c^*, bursts can still occur, which is called attractor bubbling [87]: c^* describes a typical orbit, but a chaotic attractor also contains unstable periodic orbits, and along some of them the loop stretches differences more than on average. While the synchronized pair moves near such an orbit, a tiny difference can grow into a burst before the coupling pulls the copies back together.

We tested this on three chaotic cells of Qwen2.5-0.5B's loop (philosophy direction, layer 12, float64), with $\lambda_1 = 0.098$, 0.437 and 0.196 per step over 3000 steps, so that $c^* = 0.046$, 0.177 and 0.089 (Figure 62; `docs/sync/sync-summary.json`). Both copies are the loop of the rest of the paper, with the same chat prefix, steering direction and cell. They run as two rows of one batch, and at every step the program mixes the two new soft tokens as above before feeding them back. Copy A started on the attractor after a 1000-step transient and copy B at a relative distance 10^{-3} from it in a random direction. Pairs ran 3000 steps at couplings from 0 to 0.30 in steps of 0.01 and 10,000 steps at eight couplings near each threshold, and a pair counted as synchronized when $|x_A - x_B|/|x_A|$ stayed below 10^{-8} over its last 500 steps. For the transverse exponent, copy B started at relative distance 10^{-7} from copy A and was reset to that distance after every step for 2000 steps. Far starts put copy B 500 steps further along the attractor (2000 steps, couplings in steps of 0.02). Other prefixes, models and directions, larger ensembles of starts, asymmetric or delayed coupling and changes of precision were not tested. The difference's growth rate, measured directly on the coupled map at seven couplings from 0 to 0.30 in each cell, approximately follows $\lambda_1 + \ln|1-2c|$ from +0.44 to −0.82: the largest discrepancy is 0.051 per step, up to four times the twins' own standard errors, which do not include the uncertainty of λ_1. Whether pairs locked together depended on the protocol (Table 23). In the first cell the copies synchronized between $c = 0.045$ and 0.05 in 10,000-step runs, consistent with c^*. In the second they synchronized between 0.18 and 0.19 when "synchronized" means a difference below 10^{-6}, consistent with c^*; on the stricter criterion of 10^{-8} they had not

Table 23: Synchronization thresholds of the coupled copies for each protocol: the largest tested coupling that did not synchronize and the smallest from which all larger tested couplings did, $(c_{\text{lo}}, c_{\text{hi}}]$ (`docs/sync/sync-summary.json`). Near starts: copy B at relative distance 10^{-3} from copy A; far starts: copy B 500 steps further along the attractor. Criterion: the relative difference stays below 10^{-8} or 10^{-6} throughout the last 500 steps. *A smaller tested coupling also synchronized. †Synchronized pairs here include pairs locked onto a coexisting period-2 orbit of the single loop.

start, steps, criterion	cell 1	cell 2	cell 3
λ_1 (per step)	0.098	0.437	0.196
predicted c^*	0.046	0.177	0.089
near, 3000, 10^{-8}	$(0.04, 0.05]$	$(0.22, 0.23]^*$	$(0.15, 0.16]^*$
near, 3000, 10^{-6}	$(0.04, 0.05]$	$(0.19, 0.20]$	$(0.13, 0.14]$
near, 10,000, 10^{-8}	$(0.045, 0.05]$	none up to 0.19	$(0.12, 0.15]$
near, 10,000, 10^{-6}	$(0.045, 0.05]$	$(0.18, 0.19]$	$(0.12, 0.15]^*$
far, 2000, 10^{-8}	$(0.20, 0.22]^{*\dagger}$	$(0.18, 0.20]$	$(0.14, 0.16]$
far, 2000, 10^{-6}	$(0.20, 0.22]^{*\dagger}$	$(0.18, 0.20]$	$(0.10, 0.12]$

synchronized by 0.19. Just below its predicted threshold, at $c = 0.16$ to 0.17, the difference rose from below 10^{-8} to above 10^{-3} 12 to 14 times in 10,000 steps, and the synchronized stretches had a length distribution whose rough log-binned slopes lay between about -1.8 and -1.5, close to $-3/2$; these descriptive fits, from the few runs with enough stretches, do not establish a power law. In the third cell the copies synchronized only well above c^*, between 0.15 and 0.16 in 3000-step runs and between 0.12 and 0.15 in 10,000-step runs on the 10^{-8} criterion (Table 23). At $c = 0.105$ to 0.115, where the difference should shrink on average (predicted transverse exponents -0.040 to -0.065), it still burst from the rounding level (10^{-14}) to above 10^{-3}, 5 to 10 times in 10,000 steps, consistent with attractor bubbling; the candidate transversely unstable orbits were not computed. No noise or mismatch was added. Two identical inputs in one batch give outputs that differ by up to 9×10^{-16} relative, a plausible seed of the bursts whose causal role was not isolated. Synchronized copies write the same text (the most probable token at each step), and their texts sample the single loop's attractor. Below c^*, many pairs instead locked onto periodic orbits of the joint state (periods 6 to 12, verified to close within 2.2×10^{-13}), on which the two copies write different or staggered texts. In the first cell, some starts far from the attractor synchronized onto a second attractor of the single loop, an attracting 2-cycle "it is", and two far-start pairs at $c = 0.14$ and 0.20, above that cell's c^*, settled on periodic orbits of the joint state without synchronizing.

The number that measures the loop's chaos, its largest Lyapunov exponent, thus also sets how fast the difference between two coupled copies grows: the measured growth rate approximately followed the prediction at every coupling, within 0.051 per step. The predicted threshold for a stable synchronized state rises with the chaos, from 0.046 in the weakest to 0.177 in the strongest, and in the common near-start protocol the observed thresholds rank the same way: the more chaotic each copy is alone, the more of the other's output it must take in. The threshold concerns a typical orbit, however, and synchrony in practice depends on the observation time, the tolerance, the start and the attractor reached. Near the threshold the copies alternated between long synchronized stretches and bursts; in one cell bursts from the rounding level continued above the threshold; and from some starts the pair locked onto a coexisting periodic orbit. Weak coupling, instead of synchrony, often produced order, a periodic exchange on which the copies write different texts.

The coupled-agents check of Appendix D differs in one essential way: there the contexts grow, so the two agents are not two copies of a fixed map, and at $c = 0.25$ they barely synchronized although

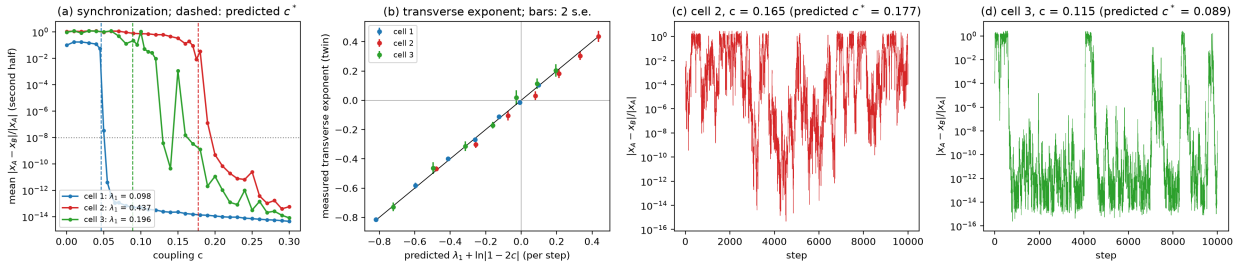

Figure 62: Synchronization of two coupled copies of the soft-token loop (Qwen2.5-0.5B-Instruct, philosophy direction, layer 12, float64). (a) Mean relative difference between the copies in the second half of the run against the coupling c, for three chaotic cells; dashed, the predicted threshold $c^* = (1 - e^{-\lambda_1})/2$. (b) The difference's growth rate measured on the coupled map against the prediction $\lambda_1 + \ln|1 - 2c|$ (bars: two standard errors). (c) On-off intermittency just below the threshold of the second cell ($c = 0.165$). (d) Bursts above the predicted threshold in the third cell ($c = 0.115$), as in attractor bubbling.

$c = 0.25$ exceeds c^* for all three cells here.

17 Discussion

Table 24 lists what each main result rests on: an identity, numerical computation in floating point, a test whose design was fixed before its data, or a computer-assisted proof.

Local delay, global return and long-time fate. The return study of Section 8.1 separates these three questions. The normal form predicts the local delay. Natural departures admit an accurate one-coordinate state reconstruction for the tested next cycle, but its excursion is still evaluated through the full network. Period-3 capture after repeated visits prevents those visits alone from establishing persistent chaos. An autonomous return surrogate and an invariant-set classification remain distinct open tasks.

What theory allows, and what is new. Chaos and periodic windows occur in many nonlinear feedback systems: random recurrent networks become chaotic above a critical gain [22, 23], and periodic windows, shrimp-shaped domains and doubling cascades are found in low-dimensional maps and in experiments [4, 24, 88]. In our map temperature plays the role of an inverse gain and the steering vector that of an input. The factorization of Section 3 makes this correspondence concrete, but it remains an analogy, not a transfer of the random-network theorems (Appendix C). These results motivate looking for such structure in a language model's feedback map; they do not imply it. This paper shows that the models we studied realize it, with concrete texts as attractors, so that a bifurcation changes what the model writes, and that several prospective quantitative predictions succeeded; the failed all-row criteria and the unresolved cases are reported with them.

Is steering the origin of the chaos? No. Without steering ($\alpha = 0$) the map is the model's own feedback map, and temperature alone gives chaos with embedded windows in most dense models we swept (in eleven of fourteen on regular grids). Steering widens the chaotic regions, but not through the concept it names: in the 0.5B model norm-matched random directions are as chaotic as the concept directions, and in Llama-3.1 8B sparse-autoencoder feature, contrastive and random directions all are, with direction-dependent prevalence. What the direction selects is the text of the attractors (Appendix J).

Table 24: What each main result rests on. *Prospective*: the design, data and criteria were fixed and pushed before the data were collected; *partly prospective*: some predictions or criteria were. Only the period-8 orbit and the coexisting period-2 and period-3 orbits are proved; chaos labels are finite-horizon numerical results. Each *code* entry links to the main tool or record in the project repository.

Result	**Evidence**	**Status**	**Code**
Mechanism			
Factorized derivative of the loop	identity; checked to 10^{-13}	identity	code
Steering's growth is readout-side	crossover along orbits; logit bias	partly prospective	code
Chaos and its consequences			
Chaos in 13 of 15 models (11 unsteered)	tangent exponents; 74–92 % last 10^5 steps	numerical	code
Batch layout splits chaotic loops	paired float64 runs, 4 models	numerical	code
Batch layout changes answers	arithmetic, GSM8K, reasoning	prospective	code
Chaos does not predict accuracy	regression on newly drawn problems	prospective	code
Bifurcations and scaling laws			
Windows accumulate on folds ($1/N^2$)	fold-edge fits; forecasts within 3 %	partly prospective	code
Passage times from normal forms	42 of 48 within one iterate	prospective	code
Cusps and their unfoldings	unfoldings fixed before solving	partly prospective	code
Feigenbaum ratios	flips by continuation	numerical	code
Period-5 shrimp (domain, legs, head)	followed attractor, curves	numerical	code
Crossing window families	eight maps; independent audit	numerical	code
Attracting period-8 orbit	interval arithmetic	**proof**	code
Coexisting attracting 2- and 3-cycles	interval arithmetic	**proof**	code
Attracting 3-cycle where the cascade begins	interval arithmetic	**proof**	code
Beyond the one-token loop			
Steering history changes a decision	new instances	prospective	code
Huginn loses accuracy under over-relaxation	generation test	prospective	code

Table 25: Each bifurcation, crossing and boundary named in this paper: what it requires and what we show. Multipliers are those of the period-p map $G = F^p$; a, b and c are coefficients of the center-manifold normal forms. The numbers are given in the text.

object	requires	we show	status
fold	one multiplier at $+1$; nonzero coefficients a and b	both at the period-2 fold, where a stable and an unstable orbit meet; passage times forecast at four more folds	shown
flip	one multiplier at -1; nonzero crossing speed and cubic coefficient	multiplier at -1 and the doubled orbit at every level of the cascade; cubic coefficient only at two other flips	partly shown
cusp	$+1$ with $b = 0$, $c \neq 0$ and a regular unfolding	the three conditions hold at the kite's cusp C and at five more located cusps	shown
pseudo-superstable line	the leading multiplier changes sign without vanishing	the tracked two-dimensional block's trace-zero curves and fitted saddle support a local explanation of the ditches' separation on the sampled patch	supported numerically on the patch
trace reconnection	a trace saddle reaches zero; for a generic crossing its height changes transversely with a third control	a height sign bracket and residual-qualified candidate under retained feedback on the period-18 family; nonzero multipliers	numerical; not certified
fold–flip point	one orbit with $+1$ and -1 at once	none at A or B	none found
pseudo-crossing	curves of two different orbits cross	two distinct orbits at R; different spectra at A and B	shown
Feigenbaum ratio	ratios approach the constant of a universality class	ratios reach 4.6688, against $\delta = 4.6692$	consistent
basin boundary	changes with the initial state; both attractors exist on both sides of it; the start lies on a saddle's stable manifold	absent with random-token starts; both period-6 orbits exist on both sides; from it the loop nears a period-6 saddle and leaves at the predicted time	strong evidence
Neimark–Sacker	complex pair crosses the unit circle, plus genericity conditions	crossing predicted where an invariant circle appears	consistent

What is and is not claimed. Table 25 lists, for each bifurcation, crossing and boundary named here, what it requires and what we show. The objects studied here are invariant sets and bifurcations of an exactly defined smooth map built from a frozen model's own computation. The initial arithmetic reproducibility tests use two Qwen2.5 models (0.5B and 7B) and a soft-thinking pipeline without the top-k filtering of the published method, and their bfloat16 effect does not require the float64 map's chaos: in float64 no answer changed in the 0.5B model, and 5 % changed in the 7B model, whose loop expands at the temperature tested. The later GSM8K tests include one at the published Soft Thinking settings and separate tests of three reasoning models; ordinary sampled reasoning also changes answers with the batch layout. Nearly all of the reasoning models' changed GSM8K answers come from problems that reach a thinking or answer budget in at least one layout. These measurements do not establish the rate for reasoning that completes (Section 6). The mechanism of Section 3 is measured along three directions in a 0.5B model and one in a 7B model, and at high temperature the network factor can carry the instability. The loop's regime does not mark the accuracy of ordinary sampled answers beyond the model's first-token uncertainty, and when slight steering moved a prompt's loop into chaos no change of that accuracy was detected ($+0.36$ points, $[-0.36, 1.08]$; Appendix D). We do not claim that ordinary sampled generation is chaotic in this sense: in the one exploratory test we ran, its sensitivity to a small change of the prompt did not depend detectably on the regime of the soft map (Appendix G). Nor do we claim that any window is a certified shrimp. The period-5 window is a numerically verified shrimp (Section 11): its attractor, followed cell by cell, fills a body with legs and doubling shells bounded by continued fold and flip curves with located cusps, with a closed trace-zero curve and two dovetails; none of it is proved, no two-dimensional reduction of its dynamics is established, and its head is not that of a one-dimensional shrimp. The period-6 and period-7 windows have fold-and-flip skeletons with located cusps, and in the hostile period-7 body the candidate superstable arcs (sign changes of the leading real multiplier) avoid each other where a head would be. Symbolic sequences are determined for the period-6 window only (paragraph "Symbolic dynamics"); we have not traced tails, tested for a periodicity hub [64, 65], or shown that a two-dimensional map reproduces the windows' nonlinear dynamics. The cusps are float64 computations validated by step-size checks, not proofs (Appendix E). The period-18 child's trace-saddle sign change under retained feedback (Section 11.1) changes the map and follows the moving saddle in steering and temperature; contour reconnection is its conditional local interpretation. It is not an exact superstable crossing of the original map, does not establish reconnection of the exact strongest-contraction ridges, and does not identify a native modal-coupling mechanism.

Limitations. The detailed orbit studies cover Qwen2.5 at 0.5B and 1.5B, with three float64 orbits at 7B; OLMo-2 7B and Llama-3.1 8B have maps but no continued orbits. The size comparison rests on regular temperature grids and mostly one twin offset per label, and part of the 70B-class maps predate a precision fix and were rechecked only in samples. The maps use one prompt, apart from 22 variants of the 0.5B standard plane (more in the size tests), and steering uses contrastive vectors, checked against random and sparse-autoencoder directions in two models. Chaos labels are finite-horizon estimates: a longer horizon check kept 86 and 92 of 96 chaotic labels in two maps, near the accumulation of a cascade transient chaos makes labels less reliable (on one line only 135 of 235 chaotic labels held), the chaotic cells are numerically validated (twin exponents agree with automatic-differentiation tangents, across precisions and horizons) but not yet established by a computer-assisted proof, and four periodic orbits have such proofs: the period-8 orbit, a coexisting period-2 and period-3 pair, and the 3-cycle at the start of the cascade. The Feigenbaum ratios are finite-order values. All orbit, bifurcation and chaos results concern the soft-token map, not sampled

generation with a growing context. Appendix K gives the details.

Steering strength is a bifurcation parameter. In the soft-token map, small changes of the steering strength switch the model between distinct repeating texts, chaos and coexisting attractors. Near a fold curve, behavior becomes intermittent, with laminar phases whose duration the fold's normal form predicts in advance. This argues for treating an activation-steering coefficient in a continuous-feedback loop like any control parameter of a nonlinear system (for ordinary sampled generation, the one exploratory test we ran found no change of sensitivity where the steering strength crosses an edge of the soft map; Appendix G). Its bifurcation set should be characterized before an operating point is chosen, history dependence should be expected near boundaries, and a monotone, smooth dose response should not be assumed. In a test fixed before its data, the history of a steering schedule did change a decision with right and wrong answers: after a brief excursion to stronger permissive steering and back, Qwen2.5-3B-Instruct, thinking in soft tokens, wrongly allowed 13 percentage points more of the prohibited access requests at the same final strength (Appendix D). There the memory most plausibly lies in the growing context, not in coexisting attractors of a fixed-size state. The block at which a vector enters matters as much as its strength. In a 0.5B and a 70B model, earlier injection gives more chaos and brings it to lower temperature, and two vectors at two blocks create chaos where neither does alone (Section 6). Outputs that fall into verbatim repetition at high steering doses are naturally treated as quality failures. In the present framework such loops are candidate attracting periodic orbits, and the dose at which they appear is a boundary that can be mapped like the others. The deterministic map is where such sets can be defined precisely and computed numerically. Whether its boundaries predict boundaries of sampled behavior is open. The bridge of Appendix G shows that sampling noise shortens the laminar phases of the bottleneck, by an amount that depends on how many samples are averaged.

What this suggests for interpretability and alignment. Our results come from open models of 0.5B to 72B parameters and the soft-token loop, but they suggest three things about models whose outputs feed back into their inputs, as in self-dialogue, agent loops and pipelines of model calls. First, where the loop is chaotic, a perturbation of 10^{-6} grows to order one, while the loop stays in that regime, within roughly 12 to 50 steps at the largest unsteered exponents of the chaotic 7B to 14B models (0.30 to 1.14 per step), and more slowly at smaller ones. Even a complete mechanistic account of what a single forward pass does would then not predict a long trajectory: that would require knowing the state to far more digits than any interpretability description, which summarizes the state by a few features, can supply. Stepwise summaries can be inconclusive: every step of the period-8 orbit expands some direction while the cycle contracts, so the ordered product of the step Jacobians, not a summary of each step, captures the loop's stability (Section 3). What can remain predictable are the invariant sets, their statistics and the boundaries between them. This extends the analysis of trained recurrent networks through their fixed points [89, 90] to the periodic orbits, chaos and bifurcations of a pretrained transformer. Second, feature dictionaries say which features an attractor occupies, but the boundaries between attractors, which decide what a small change does, are properties of the map's derivatives. That would hold even if most units have one interpretable job, as a recent essay argues for Qwen3-32B [91]; in its pruning tests, too, the answers that change are mostly those the network was already unsure of, much as the loop's chaotic cells are those where next tokens compete. Single neurons are no exception: in the 0.5B model, steering along the output direction of the neuron most specific to philosophy makes 9.9 % of the map chaotic, its index neighbor 5.1 % and a random neuron 12.9 %, against 6.8 to 12.5 % along the concept and Gaussian directions (Appendix J). In our 8B controls, a sparse-autoencoder feature direction,

contrastive and random directions all drove the map into chaos (27 to 39 % of cells); at 70B the hostility-related directions gave more chaotic cells than the random ones (4.3 and 3.5 % against 1.1 and 0.2 %, two-GPU maps from before the precision fix; Appendix I). Along recorded orbits, the growth that steering adds comes mainly from the readout, where next tokens compete (Section 3); this is not specific to concept-labeled directions, but a direction's geometry and its meaning can be correlated. Text can also hide the dynamics: the emitted text records only the most probable token, and in the hostile cascade a seven-token text repeats exactly while the probabilities that drive the loop do not repeat over the run and the state underneath is chaotic (Figure 33), so a monitor that reads only the emitted text would call a chaotic regime periodic. Third, conclusions drawn at one operating point need not transfer: which loop a model enters depends on how much context it keeps (Table 29), on the path by which a boundary is approached, and on the sampling noise (Appendix G). Treating a looped model as a dynamical system, with regimes and boundaries to be mapped, rather than as a function to be evaluated point by point, seems to us the appropriate stance for both interpretability and safety evaluation. The reproducibility tests of Section 6 make this concrete: in a chaotic regime the batch a request happens to share can change, after a predictable number of steps, what the loop writes, while in a periodic regime it almost never does (in the 0.5B model, the text changed in 173 of 192 chaotic loops and in 7 of 192 periodic ones). In low precision, continuous feedback makes even short-budget answers depend on the batch, far more often than greedy decoding does; over a long run in a pipeline that answers poorly, discrete feedback does too. These findings motivate testing whether interventions in continuous-feedback systems, including steering for safety, alter finite-horizon sensitivity or task-relevant decisions as well as immediate outputs; such tests would complement research on how latent-reasoning architectures affect chain-of-thought monitorability [92]. We tested no safety property: dynamical instability is neither misalignment nor evidence of it.

Other inference loops. Any inference scheme that feeds a network's output or internal state back as its input defines an iterated map, but feedback by itself implies neither convergence nor chaos. The energy-descent argument for the isolated tied readout (Section 3) does not establish convergence of either complete feedback operator studied here. In the soft-token loop we find fixed points, cycles and chaos. In the recurrent-depth model Huginn (Section 15) the loop's derivative is the block's Jacobian, without the readout factor; there we observe convergence and oscillatory instability, predicted from the block's multipliers, and have not established a chaotic regime. These results motivate evaluating the dynamics of looped and recurrent-depth language models [93, 94] jointly with their number of iterations; they do not establish a general relation between stability and reasoning accuracy, especially at the short horizons of deployed loops, where no attractor need be reached. Agentic loops that sample discrete tokens and tool outputs have no derivative at all; only the zero-temperature limit and the sensitivity measurements of this paper bear on them.

Is the chaos useful? The chaos we map is not a source of new content. The chaotic cells of Table 9 recombine the phrases of the neighboring loops in an order that did not settle over the recorded horizon: the trajectory wanders among the ghosts of nearby periodic orbits, which gives variety in sequence but not in substance. The periodic loops are the familiar degeneration of decoding into repetition [27, 28]. Neither regime is good text, and the diversity of ordinary generation comes from sampling and a growing context, which wash out much of this structure (Appendix G). What the chaos does matter for is predictability and control. Positive finite-horizon exponents limit open-loop prediction of a trajectory at fixed state precision (our experiments do not establish mixing or limits on feedback stabilization); a stable attractor is predictable, but it can also hold a model in a state it

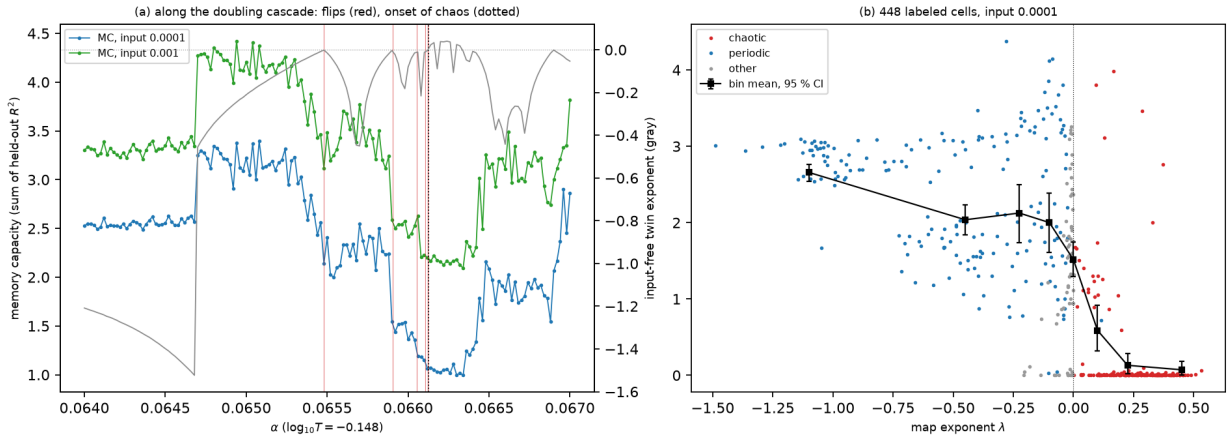

Figure 63: Memory capacity of the soft-token loop used as a reservoir (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, float64). A weak random input enters the map, $x_{n+1} = F(x_n + \epsilon|x|u_n r)$, and linear readouts of the state reconstruct the past inputs; the capacity is the sum of their held-out R^2 over 41 delays. (a) Along the doubling cascade at $\log_{10} T = -0.148$: capacity for two input sizes, the input-free twin exponent (gray), the flips (red) and the onset of chaos (dotted). (b) The 448 cells of the horizon check (192 chaotic, 192 periodic, 64 with neither label) against the map exponent, with bin means and bootstrap 95 % intervals.

should leave, such as the hostile loops of Section 10. Whether regimes near the boundary between order and chaos are especially useful, as has been argued for computation in dynamical systems and recurrent networks [95, 96] and shown for signal propagation in deep networks at initialization [97, 98], is open for language models. We tested three versions, and none shows an advantage at the edge. Prompts whose own loop sits at the edge ($|\lambda| \leq 0.1$) are answered no better by ordinary sampling than prompts whose loop is deeply periodic or deeply chaotic (Section 6). Used as a reservoir, the loop remembers a weak injected input best in short stable cycles. Its linear memory capacity falls at every period doubling toward the onset of chaos, does not rise as the exponent approaches zero within a period, and is nearly erased by chaos (Figure 63). In Huginn, whose reasoning is itself a latent loop, a float32 pilot on 40 development problems, separate from the tests of Section 15, recorded the decay rate of every generated token's recurrence. Accuracy holds while that recurrence is comfortably stable and falls as it approaches its edge: it was 0.475 natively and 0.525 at relaxation 1.05, where the token recurrence is nearly as stable as natively (median decay rate 0.921 against 0.912), and fell to 0.375 at 1.15, where the recurrence is near critical (0.981, with about half of the tokens no longer converging). With 40 problems these differences are indicative only.

Which findings might reach frontier models. The parts that rest only on structure every smooth softmax-readout language model with a fixed prefix shares should carry over: the factorization of the loop's derivative, contraction at sufficiently high temperature, greedy token cycles at zero temperature wherever the most probable token is unique, and the role of a steering coefficient as a bifurcation parameter. The attractor texts, the positions of windows, folds and cusps, and every number we report belong to particular weights; a second checkpoint of the same architecture already has different attractors (Appendix H). Size gives no guarantee of stability (Section 14), hard routing in mixture-of-experts models falls outside the smooth theory of Section 3 (no expert switched along our float64 orbits, but in low precision routing can), and deployed inference samples tokens, in low precision, over long contexts; an 8-bit weight quantization that changes next-token distributions by

only 0.003 nats changes the attractor class or period of one start in seven (Appendix D). Closed models permit none of these computations, which need the input embeddings, full next-token distributions and derivatives; for them the method and the questions transfer, not the numbers.

Conclusion. A language model that reads its own soft output is a smooth discrete-time dynamical system, a smooth map iterated step by step, and it behaves as one. Temperature and steering carry it from repeating text through period-doubling cascades into chaos with a positive Lyapunov exponent. Its periodic windows open and close in folds and flips whose normal forms, computed from the network's own derivatives, forecast how long trajectories linger near a vanished orbit; its cascades approach Feigenbaum's constant; and its parameter plane holds shrimp-shaped windows organized as in the low-dimensional maps where such windows were first found. A small retained-feedback term changes a moving trace saddle's height sign in a period-18 window, supporting a conditional local reconnection interpretation without changing the sampled centers' state period or token cycle. Two coupled copies drift apart at approximately the rate the loop's Lyapunov exponent predicts. None of this rests on a simplified model of the network: the orbits, multipliers and normal forms are those of the full transformer in float64, and four of its periodic orbits are proved to exist with interval arithmetic. For the expected-embedding feedback studied here with a fixed prompt, the derivative factorization applies to smooth networks with a linear output head followed by softmax: its factors are inverse temperature, an embedding covariance for tied embeddings or a cross-covariance for untied embeddings, and the network's Jacobian. Chaos with embedded periodic windows appears in all three model families we scanned, showing that these feedback dynamics are not confined to one family of models.

Code and data availability

The code, the designs of the tests fixed in advance with their audits, the run cards, the failed attempts and hash-bound records of every computation are in the project repository, `https://github.com/tdj28/neural-shrimp`; the links in this paper point to its tag `arxiv-v1`, and those to files added since then point to a later commit of its main branch. A short script recomputes the proved period-8 orbit from the released model in a few minutes on a desktop computer (`tools/recreate_p8.py`; the steps are written out in `docs/reproduce-p8/README.md`). The interval-arithmetic certificates are in the repository, with the verification code, the exact map and input specifications, and the release records and hashes of the native evidence. The complete native evidence of the period-8 and coexistence proofs is archived at Zenodo, doi:10.5281/zenodo.23488416; that of the proof of the cascade's 3-cycle is kept outside the repository, whose release record holds its checksums. Bulky raw trajectories and model states are kept outside the repository, which holds their checksums.

Acknowledgments

The author thanks Jennifer Hoffman and Emin Arakelian for stimulating discussions about large language models and for encouraging the publication of this work.

Funding

This work received no external funding; the author paid for all computation.

Ethics and responsible use

This study involved no human participants, private conversations or personal data. The access-control requests and their users in Appendix D are synthetic, and the benchmark problems are public; institutional review was therefore not applicable. The released models and benchmarks were used under the terms of their public licenses. Some experiments steer models toward hostile or permissive behavior; they measure the stability of iterated inference and release no harmful outputs. The results concern the auditability and control of iterative pipelines: near a bifurcation, a small change of a steering strength or of the numerical environment can change what such a pipeline does.

Use of AI assistance

This work was carried out with AI coding and research agents (Anthropic's Claude Code and OpenAI's Codex) working under the author's direction. The agents wrote and ran the code, executed the experiments, checked results against the raw data and drafted text; other AI systems made independent read-only checks of the results and the manuscript. Frozen designs, run cards, failed attempts, audits and hash-bound records of every run are kept in the project repository. The author, whose doctoral research in nonlinear dynamics included work on a periodicity hub of the Rössler attractor [65], directed the study, reviewed the work and takes full responsibility for its content.

References

[1] Zhen Zhang, Xuehai He, Weixiang Yan, Ao Shen, Chenyang Zhao, and Xin Eric Wang. Soft Thinking: Unlocking the reasoning potential of LLMs in continuous concept space. In *Advances in Neural Information Processing Systems*, volume 38, pages 187190–187212, 2025. doi: 10.52202/085713-5629. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/f7396d1c54d51416958d63e285377103-Abstract-Conference.html.

[2] Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason E. Weston, and Yuandong Tian. Training large language models to reason in a continuous latent space. In *Second Conference on Language Modeling*, 2025. URL https://openreview.net/forum?id=Itxz7S4Ip3. arXiv:2412.06769.

[3] Jonas Geiping, Sean McLeish, Neel Jain, et al. Scaling up test-time compute with latent reasoning: A recurrent depth approach. In *Advances in Neural Information Processing Systems*, volume 38, pages 46242–46293, 2025. doi: 10.52202/085713-1380. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/3b01972cf31e6fa0fe29e4b8b5c2a0a1-Abstract-Conference.html.

[4] Jason A. C. Gallas. Structure of the parameter space of the Hénon map. *Physical Review Letters*, 70(18):2714–2717, 1993. doi: 10.1103/PhysRevLett.70.2714.

[5] Xiaojian Li, Yongkang Leng, Ruiqing Ding, Hangjie Mo, and Shanlin Yang. Cognitive activation and chaotic dynamics in large language models: A quasi-Lyapunov analysis of reasoning mechanisms. arXiv preprint arXiv:2503.13530, 2025. URL https://arxiv.org/abs/2503.13530.

[6] Chashi Mahiul Islam, Alan Villarreal, Mao Nishino, Shaeke Salman, and Xiuwen Liu. Numerical instability and chaos: Quantifying the unpredictability of large language models. In *2026 International Joint Conference on Neural Networks (IJCNN)*, pages 1–8. IEEE, 2026. doi: 10.1109/ijcnn69854.2026.11729148. arXiv:2604.13206.

[7] Gregorio Jaca, Kristóf Benedek, and János Török. Chaos in reason: How chain-of-thought LLMs can look for an answer. arXiv preprint arXiv:2607.27805, 2026. URL https://arxiv.org/abs/2607.27805.

[8] Julian Arnold, Flemming Holtorf, Frank Schäfer, and Niels Lörch. Phase transitions in the output distribution of large language models. arXiv preprint arXiv:2405.17088, 2024. URL https://arxiv.org/abs/2405.17088.

[9] Kai Nakaishi, Yoshihiko Nishikawa, and Koji Hukushima. Phase transition in large language models and the criticality of natural languages. arXiv preprint arXiv:2406.05335, 2024. URL https://arxiv.org/abs/2406.05335.

[10] Youran Sun and Babak Haghighat. Phase transitions in large language models and the $O(N)$ model. arXiv preprint arXiv:2501.16241, 2025. URL https://arxiv.org/abs/2501.16241.

[11] Huajian Ruan, Jinyang Li, Xingyu Guo, and Lingxiao Wang. Generative criticality in large language model temperature scaling. arXiv preprint arXiv:2606.06238, 2026. URL https://arxiv.org/abs/2606.06238.

[12] Neil F. Johnson, Frank Yingjie Huo, and Bella Xinrui Li. Temperature-driven inversion and nonlinear dynamics in ChatGPT-like AIs. arXiv preprint arXiv:2608.00939, 2026. URL https://arxiv.org/abs/2608.00939.

[13] Zhilin Wang, Yafu Li, Jianhao Yan, Yu Cheng, and Yue Zhang. Unveiling attractor cycles in large language models: A dynamical systems view of successive paraphrasing. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 12740–12755. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.acl-long.624. arXiv:2502.15208.

[14] Ting-Wen Ko and Jonas Geiping. Attractor states emerge in multi-turn LLM conversations, 2026.

[15] Shaojie Bai, J. Zico Kolter, and Vladlen Koltun. Deep equilibrium models. In *Advances in Neural Information Processing Systems*, volume 32, pages 690–701, 2019. URL https://proceedings.neurips.cc/paper_files/paper/2019/hash/01386bd6d8e091c2ab4c7c7de644d37b-Abstract.html. arXiv:1909.01377.

[16] Shaojie Bai, Vladlen Koltun, and Zico Kolter. Stabilizing equilibrium models by Jacobian regularization. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 554–565. PMLR, 2021. URL https://proceedings.mlr.press/v139/bai21b.html.

[17] Akiyoshi Tomihari and Ryo Karakida. Recurrent self-attention dynamics: An energy-agnostic perspective from Jacobians. In *Advances in Neural Information Processing Systems*, volume 38, pages 155236–155269, 2025. doi: 10.52202/085713-4675. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/cca569b38ec3612904e625dad7a743dc-Abstract-Conference.html.

[18] Xiao-Wen Yang, Ziyu Han, Xi-Hua Zhang, Wen-Da Wei, Jie-Jing Shao, Lan-Zhe Guo, and Yu-Feng Li. Stabilizing recurrent dynamics for test-time scalable latent reasoning in looped language models. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, pages 144670–144682. PMLR, 2026. URL https://proceedings.mlr.press/v306/yang26t.html.

[19] Xinyue Luo and Fei Yu. Prediction dynamics in depth-recurrent language models. arXiv preprint arXiv:2609.21383, 2026. URL https://arxiv.org/abs/2609.21383.

[20] Boyi Zeng, Shixiang Song, Siyuan Huang, et al. PonderLM: Pretraining language models to ponder in continuous space. In *The Fourteenth International Conference on Learning Representations*, pages 91874–91893, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/94572a66c6b0710c5263f15d8ef6b463-Abstract-Conference.html. arXiv:2505.20674.

[21] Gunn Kim. Dynamical phase selection controls compute scaling in looped transformers. arXiv preprint arXiv:2608.26556, 2026. URL https://arxiv.org/abs/2608.26556. Version 1, 27 August 2026.

[22] H. Sompolinsky, A. Crisanti, and H. J. Sommers. Chaos in random neural networks. *Physical Review Letters*, 61(3):259–262, 1988. doi: 10.1103/PhysRevLett.61.259.

[23] B. Doyon, B. Cessac, M. Quoy, and M. Samuelides. On bifurcations and chaos in random neural networks. *Acta Biotheoretica*, 42(2–3):215–225, 1994. doi: 10.1007/BF00709492.

[24] Mitchell J. Feigenbaum. Quantitative universality for a class of nonlinear transformations. *Journal of Statistical Physics*, 19(1):25–52, 1978. doi: 10.1007/BF01020332.

[25] Thomas Hofmann. The map behind the flow: Finite-step gradient descent as a dynamical system. arXiv preprint arXiv:2607.04993, 2026. URL https://arxiv.org/abs/2607.04993.

[26] Nicolás Vera Zúñiga. What iterated self-feeding probes of language models measure, and a test that separates the construction from the model. arXiv preprint arXiv:2608.10986, 2026. URL https://arxiv.org/abs/2608.10986.

[27] Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. In *The Eighth International Conference on Learning Representations*, 2020. URL https://openreview.net/forum?id=rygGQyrFvH. arXiv:1904.09751.

[28] Jin Xu, Xiaojiang Liu, Jianhao Yan, Deng Cai, Huayang Li, and Jian Li. Learning to break the loop: Analyzing and mitigating repetitions for neural text generation. In *Advances in Neural Information Processing Systems*, volume 35, pages 3082–3095, 2022. doi: 10.52202/068431-0223. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/148c0aeea1c5da82f4fa86a09d4190da-Abstract-Conference.html. arXiv:2206.02369.

[29] Nicolas Tacheny. Geometric dynamics of agentic loops in large language models. arXiv preprint arXiv:2512.10350, 2025. URL https://arxiv.org/abs/2512.10350.

[30] Qishuai Wen, Zhiyuan Huang, Xianghan Meng, Wei He, and Chun-Guang Li. Mixture-of-top-k attention: Efficient attention via scalable fast weights, 2026. URL https://arxiv.org/abs/2602.01219v6. Version 6.

[31] V. I. Oseledets. A multiplicative ergodic theorem. Lyapunov characteristic numbers for dynamical systems. *Transactions of the Moscow Mathematical Society*, 19:197–231, 1968. English translation of Trudy Moskovskogo Matematicheskogo Obshchestva 19 (1968) 179–210; Zbl 0236.93034.

[32] J.-P. Eckmann and D. Ruelle. Ergodic theory of chaos and strange attractors. *Reviews of Modern Physics*, 57(3):617–656, 1985. doi: 10.1103/RevModPhys.57.617.

[33] Mostafa Dehghani, Stephan Gouws, Oriol Vinyals, Jakob Uszkoreit, and Łukasz Kaiser. Universal Transformers. In *The Seventh International Conference on Learning Representations*, 2019. URL https://arxiv.org/abs/1807.03819.

[34] Angeliki Giannou, Shashank Rajput, Jy-yong Sohn, Kangwook Lee, Jason D. Lee, and Dimitris Papailiopoulos. Looped transformers as programmable computers. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 11398–11442. PMLR, 2023. URL https://proceedings.mlr.press/v202/giannou23a.html.

[35] Nina Rimsky, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Turner. Steering Llama 2 via contrastive activation addition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15504–15522. Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.acl-long.828.

[36] Alexander Matt Turner, Lisa Thiergart, Gavin Leech, David Udell, Juan J. Vazquez, Ulisse Mini, and Monte MacDiarmid. Steering language models with activation engineering. arXiv preprint arXiv:2308.10248, 2023. URL https://arxiv.org/abs/2308.10248. Version 5, 10 October 2024; first posted 20 August 2023.

[37] Jinghao Zhang, Yuting Liu, Wenjie Wang, Qiang Liu, Shu Wu, Liang Wang, and Tat-Seng Chua. Personalized text generation with contrastive activation steering. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7128–7141. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.acl-long.353.

[38] Andy Zou, Long Phan, Sarah Chen, et al. Representation engineering: A top-down approach to AI transparency. arXiv preprint arXiv:2310.01405, 2023. URL https://arxiv.org/abs/2310.01405. Version 4, 3 March 2025; first posted 2 October 2023.

[39] Hubert Ramsauer, Bernhard Schäfl, Johannes Lehner, et al. Hopfield networks is all you need. In *The Ninth International Conference on Learning Representations*, 2021. URL https://openreview.net/forum?id=tL89RnzIiCd. arXiv:2008.02217.

[40] Martin J. Wainwright and Michael I. Jordan. Graphical models, exponential families, and variational inference. *Foundations and Trends in Machine Learning*, 1(1–2):1–305, 2008. doi: 10.1561/2200000001.

[41] Bolin Gao and Lacra Pavel. On the properties of the softmax function with application in game theory and reinforcement learning. arXiv preprint, 2017.

[42] Toni J. B. Liu, Baran Zadeoğlu, Nicolas Boullé, Raphaël Sarfati, Gurbir Arora, and Christopher J. Earls. Jacobian scopes: Token-level causal attributions in LLMs, 2026. URL https://arxiv.org/abs/2601.16407.

[43] Sihan Wang, Jiayi Zhao, Qingyan Cao, Hongbo Yao, and Lin Shu. FishBack: Pullback Fisher geometry for optimal activation steering in transformers, 2026. URL https://arxiv.org/abs/2605.17231.

[44] Ali Backour. The dynamics of continuous mixture collapse in language models. arXiv preprint arXiv:2609.02049, 2026. URL https://arxiv.org/abs/2609.02049.

[45] Tongxi Wang. Sharp spectral thresholds for logit fixed points, 2026. URL https://arxiv.org/abs/2605.15651.

[46] J. F. C. Kingman. Subadditive ergodic theory. *The Annals of Probability*, 1(6):883–899, 1973. doi: 10.1214/aop/1176996798.

[47] David Singer. Stable orbits and bifurcation of maps of the interval. *SIAM Journal on Applied Mathematics*, 35(2):260–267, 1978. doi: 10.1137/0135020.

[48] Celso Grebogi, Edward Ott, and James A. Yorke. Crises, sudden changes in chaotic attractors, and transient chaos. *Physica D: Nonlinear Phenomena*, 7(1–3):181–200, 1983. doi: 10.1016/0167-2789(83)90126-4.

[49] Willy J. F. Govaerts. *Numerical Methods for Bifurcations of Dynamical Equilibria*. Society for Industrial and Applied Mathematics, 2000. doi: 10.1137/1.9780898719543.

[50] Yuri A. Kuznetsov. *Elements of Applied Bifurcation Theory*, volume 112 of *Applied Mathematical Sciences*. Springer, Cham, fourth edition, 2023. doi: 10.1007/978-3-031-22007-4.

[51] Ying-Cheng Lai and Tamás Tél. *Transient Chaos: Complex Dynamics on Finite Time Scales*, volume 173 of *Applied Mathematical Sciences*. Springer, New York, NY, 2011. doi: 10.1007/978-1-4419-6987-3.

[52] David Goldberg. What every computer scientist should know about floating-point arithmetic. *ACM Computing Surveys*, 23(1):5–48, 1991. doi: 10.1145/103162.103163.

[53] Horace He and Thinking Machines Lab. Defeating nondeterminism in LLM inference. *Thinking Machines Lab: Connectionism*, 2025. doi: 10.64434/tml.20250910. URL https://thinkingmachines.ai/blog/defeating-nondeterminism-in-llm-inference/.

[54] Jiayi Yuan, Hao Li, Xinheng Ding, Wenya Xie, Yu-Jhe Li, Wentian Zhao, Kun Wan, Jing Shi, Xia Hu, and Zirui Liu. Understanding and mitigating numerical sources of nondeterminism in LLM inference. In *Advances in Neural Information Processing Systems*, volume 38, pages 188045–188077, 2025. doi: 10.52202/085713-5653.

[55] Steven H. Strogatz and Robert M. Westervelt. Predicted power laws for delayed switching of charge-density waves. *Physical Review B*, 40(15):10501–10508, 1989. doi: 10.1103/PhysRevB.40.10501.

[56] Steven H. Strogatz. *Nonlinear Dynamics and Chaos: With Applications to Physics, Biology, Chemistry, and Engineering*. CRC Press, Boca Raton, second edition, 2018. doi: 10.1201/9780429492563.

[57] Yuri A. Kuznetsov and Hil G. E. Meijer. *Numerical Bifurcation Analysis of Maps*. Cambridge University Press, Cambridge, 2019. doi: 10.1017/9781108585804.

[58] Yves Pomeau and Paul Manneville. Intermittent transition to turbulence in dissipative dynamical systems. *Communications in Mathematical Physics*, 74(2):189–197, 1980. doi: 10.1007/BF01197757.

[59] J. E. Hirsch, B. A. Huberman, and D. J. Scalapino. Theory of intermittency. *Physical Review A*, 25(1):519–532, 1982. doi: 10.1103/PhysRevA.25.519.

[60] Mitchell J. Feigenbaum. The universal metric properties of nonlinear transformations. *Journal of Statistical Physics*, 21(6):669–706, 1979. doi: 10.1007/BF01107909.

[61] P. Collet, J.-P. Eckmann, and H. Koch. Period doubling bifurcations for families of maps on $\mathbb{R}^n$. *Journal of Statistical Physics*, 25(1):1–14, 1981. doi: 10.1007/BF01008475.

[62] Christoforos A. Papasavvas, Yujiang Wang, Andrew J. Trevelyan, and Marcus Kaiser. Gain control through divisive inhibition prevents abrupt transition to chaos in a neural mass model. *Physical Review E*, 92(3):032723, 2015. doi: 10.1103/PhysRevE.92.032723.

[63] Yilun Liu, Yi Zhang, Ganyu Wu, Sikuan Yan, Mengyue Wang, Alois Knoll, Volker Tresp, and Yunpu Ma. Learning chaos without seeing chaos: Extrapolation of global dynamics in autoregressive transformers. arXiv preprint arXiv:2609.38814, 2026. URL https://arxiv.org/abs/2609.38814v1. Version 1, 30 September 2026.

[64] Renato Vitolo, Paul Glendinning, and Jason A. C. Gallas. Global structure of periodicity hubs in Lyapunov phase diagrams of dissipative flows. *Physical Review E*, 84(1):016216, 2011. doi: 10.1103/PhysRevE.84.016216.

[65] Timothy D. Jones. Topological origins of a bi-parameter periodicity hub for the Rössler attractor. arXiv preprint arXiv:1201.4343, 2012. URL https://arxiv.org/abs/1201.4343.

[66] Jason A. C. Gallas. Dissecting shrimps: results for some one-dimensional physical models. *Physica A*, 202(1–2):196–223, 1994. doi: 10.1016/0378-4371(94)90174-0.

[67] Wilson Façanha, Bart Oldeman, and Leon Glass. Bifurcation structures in two-dimensional maps: The endoskeletons of shrimps. *Physics Letters A*, 377(18):1264–1268, 2013. doi: 10.1016/j.physleta.2013.03.025.

[68] S. Grossmann and S. Thomae. Invariant distributions and stationary correlation functions of one-dimensional discrete processes. *Zeitschrift für Naturforschung A*, 32(12):1353–1363, 1977. doi: 10.1515/zna-1977-1204.

[69] Edward N. Lorenz. Noisy periodicity and reverse bifurcation. *Annals of the New York Academy of Sciences*, 357(1):282–291, 1980. doi: 10.1111/j.1749-6632.1980.tb29693.x.

[70] Christian Mira. *Chaotic Dynamics: From the One-Dimensional Endomorphism to the Two-Dimensional Diffeomorphism*. World Scientific, Singapore, 1987. doi: 10.1142/0413.

[71] J. P. Carcassès, C. Mira, M. Bosch, C. Simó, and J. C. Tatjer. "Crossroad area–spring area" transition (I): parameter plane representation. *International Journal of Bifurcation and Chaos*, 1(1):183–196, 1991. doi: 10.1142/S0218127491000117.

[72] C. Mira, J. P. Carcassès, M. Bosch, C. Simó, and J. C. Tatjer. "Crossroad area–spring area" transition (II): foliated parametric representation. *International Journal of Bifurcation and Chaos*, 1(2):339–348, 1991. doi: 10.1142/S0218127491000269.

[73] Christian Mira, Laura Gardini, Alexandra Barugola, and Jean-Claude Cathala. *Chaotic Dynamics in Two-Dimensional Noninvertible Maps*, volume 20 of *World Scientific Series on Nonlinear Science, Series A*. World Scientific, Singapore, 1996. doi: 10.1142/2252.

[74] M. Hénon. A two-dimensional mapping with a strange attractor. *Communications in Mathematical Physics*, 50(1):69–77, 1976. doi: 10.1007/BF01608556.

[75] N. Metropolis, M. L. Stein, and P. R. Stein. On finite limit sets for transformations on the unit interval. *Journal of Combinatorial Theory, Series A*, 15(1):25–44, 1973. doi: 10.1016/0097-3165(73)90033-2.

[76] John Milnor and William Thurston. On iterated maps of the interval. In *Dynamical Systems*, Lecture Notes in Mathematics, pages 465–563. Springer, Berlin, Heidelberg, 1988. doi: 10.1007/BFb0082847.

[77] Simon Plantinga. *Certified Algorithms for Implicit Surfaces*. PhD thesis, University of Groningen, 2007. URL https://cs.nyu.edu/~exact/pap/mesh/vegter/plantingaThesis.pdf. Appendix B, Proposition 16.

[78] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, et al. Training verifiers to solve math word problems. arXiv preprint arXiv:2110.14168, 2021. URL https://arxiv.org/abs/2110.14168.

[79] Boyuan Wang, Chengyao Yu, Jiaxi Ren, Hongxin Wei, Bingyi Jing, and Yuxin Tao. Scheduling recursive reasoning in looped transformers. arXiv preprint arXiv:2609.36653, 2026. URL https://arxiv.org/abs/2609.36653. Submitted 29 September 2026.

[80] Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? Try ARC, the AI2 Reasoning Challenge. arXiv preprint, 2018. URL https://arxiv.org/abs/1803.05457.

[81] Anthropic. System card: Claude Opus 4 & Claude Sonnet 4. System card, Anthropic, 2025. URL https://www-cdn.anthropic.com/6d8a8055020700718b0c49369f60816ba2a7c285/Claude%204%20System%20Card.pdf. May 2025; revision 2 September 2025; sections 5.5.1–5.5.2.

[82] J. Almeida, D. Peralta-Salas, and M. Romera. Can two chaotic systems give rise to order? *Physica D: Nonlinear Phenomena*, 200(1–2):124–132, 2005. doi: 10.1016/j.physd.2004.10.003.

[83] Hirokazu Fujisaka and Tomoji Yamada. Stability theory of synchronized motion in coupled-oscillator systems. *Progress of Theoretical Physics*, 69(1):32–47, 1983. doi: 10.1143/PTP.69.32.

[84] Louis M. Pecora and Thomas L. Carroll. Synchronization in chaotic systems. *Physical Review Letters*, 64(8):821–824, 1990. doi: 10.1103/PhysRevLett.64.821.

[85] N. Platt, E. A. Spiegel, and C. Tresser. On-off intermittency: A mechanism for bursting. *Physical Review Letters*, 70(3):279–282, 1993. doi: 10.1103/PhysRevLett.70.279.

[86] J. F. Heagy, N. Platt, and S. M. Hammel. Characterization of on-off intermittency. *Physical Review E*, 49:1140–1150, 1994. doi: 10.1103/PhysRevE.49.1140.

[87] Peter Ashwin, Jorge Buescu, and Ian Stewart. Bubbling of attractors and synchronisation of chaotic oscillators. *Physics Letters A*, 193(2):126–139, 1994. doi: 10.1016/0375-9601(94)90947-4.

[88] Ruedi Stoop, Philipp Benner, and Yoko Uwate. Real-world existence and origins of the spiral organization of shrimp-shaped domains. *Physical Review Letters*, 105(7):074102, 2010. doi: 10.1103/PhysRevLett.105.074102.

[89] David Sussillo and Omri Barak. Opening the black box: Low-dimensional dynamics in high-dimensional recurrent neural networks. *Neural Computation*, 25(3):626–649, 2013. doi: 10.1162/NECO_a_00409.

[90] Niru Maheswaranathan, Alex Williams, Matthew Golub, Surya Ganguli, and David Sussillo. Reverse engineering recurrent networks for sentiment classification reveals line attractor dynamics. In *Advances in Neural Information Processing Systems*, volume 32, pages 15696–15705, 2019. URL https://proceedings.neurips.cc/paper_files/paper/2019/hash/d921c3c762b1522c475ac8fc0811bb0f-Abstract.html. arXiv:1906.10720.

[91] Nick Cammarata. Neural networks are already clean. Goodfire essay, https://already-clean.goodfire.com/, n.d. Undated web essay; accessed 15 October 2026.

[92] Tomek Korbak, Mikita Balesni, Elizabeth Barnes, et al. Chain of Thought Monitorability: A new and fragile opportunity for AI safety, 2025. URL https://arxiv.org/abs/2507.11473.

[93] Rui-Jie Zhu, Zixuan Wang, Kai Hua, et al. Scaling latent reasoning via looped language models. arXiv preprint, 2025. URL https://arxiv.org/abs/2510.25741.

[94] Sangmin Bae, Yujin Kim, Reza Bayat, et al. Mixture-of-Recursions: Learning dynamic recursive depths for adaptive token-level computation. In *Advances in Neural Information Processing Systems*, volume 38, pages 107088–107133, 2025. doi: 10.52202/085713-3229. URL https://arxiv.org/abs/2507.10524.

[95] Chris G. Langton. Computation at the edge of chaos: Phase transitions and emergent computation. *Physica D: Nonlinear Phenomena*, 42(1–3):12–37, 1990. doi: 10.1016/0167-2789(90)90064-V.

[96] Nils Bertschinger and Thomas Natschläger. Real-time computation at the edge of chaos in recurrent neural networks. *Neural Computation*, 16(7):1413–1436, 2004. doi: 10.1162/089976604323057443.

[97] Ben Poole, Subhaneil Lahiri, Maithra Raghu, Jascha Sohl-Dickstein, and Surya Ganguli. Exponential expressivity in deep neural networks through transient chaos. In *Advances in Neural Information Processing Systems*, volume 29, 2016. URL https://proceedings.neurips.cc/paper_files/paper/2016/hash/148510031349642de5ca0c544f31b2ef-Abstract.html. arXiv:1606.05340.

[98] Samuel S. Schoenholz, Justin Gilmer, Surya Ganguli, and Jascha Sohl-Dickstein. Deep information propagation. In *The Fifth International Conference on Learning Representations*, 2017. URL https://openreview.net/forum?id=H1W1UN9gg. arXiv:1611.01232.

[99] Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. AI models collapse when trained on recursively generated data. *Nature*, 631(8022):755–759, 2024. doi: 10.1038/s41586-024-07566-y.

[100] António H. Ribeiro, Koen Tiels, Luis A. Aguirre, and Thomas Schön. Beyond exploding and vanishing gradients: analysing RNN training using attractors and smoothness. In *Proceedings*

of the Twenty Third International Conference on Artificial Intelligence and Statistics, volume 108 of *Proceedings of Machine Learning Research*, pages 2370–2380. PMLR, 2020. URL https://proceedings.mlr.press/v108/ribeiro20a.html.

[101] Hayden Prairie, Zachary Novack, Taylor Berg-Kirkpatrick, and Daniel Y. Fu. Parcae: Scaling laws for stable looped language models. arXiv preprint arXiv:2604.12946, 2026. URL https://arxiv.org/abs/2604.12946.

[102] Matthias M. M. Buehlmaier. LayerNorm as implicit gain control in looped transformers. arXiv preprint arXiv:2607.10681, 2026. URL https://arxiv.org/abs/2607.10681. Version 1, 12 July 2026.

[103] Asher Labovich. Stability and generalization in looped transformers. arXiv preprint arXiv:2604.15259, 2026. URL https://arxiv.org/abs/2604.15259. Version 2, 22 April 2026.

[104] Sajad Movahedi, Vera Milovanović, Shlomo Libo Feigin, Alexander Theus, Thomas Hofmann, Valentina Boeva, T. Konstantin Rusch, and Antonio Orvieto. Fixed-point reasoners: Stable and adaptive deep looped transformers. arXiv preprint arXiv:2606.18206, 2026. URL https://arxiv.org/abs/2606.18206. Version 1, 16 June 2026.

[105] Hugh Blayney, Álvaro Arroyo, Johan Obando-Ceron, et al. A mechanistic analysis of looped reasoning language models. arXiv preprint arXiv:2604.11791, 2026. URL https://arxiv.org/abs/2604.11791.

[106] Ivan Viakhirev, Kirill Borodin, Amirah Almutairi, et al. Think shallow, solve deep: Controlling recurrent dynamics for reliable test-time depth. arXiv preprint arXiv:2608.18222, 2026. URL https://arxiv.org/abs/2608.18222.

[107] Sumin Park and T. Konstantin Rusch. Looped reasoning finishes earlier than you think. arXiv preprint arXiv:2610.18174, 2026. URL https://arxiv.org/abs/2610.18174.

[108] Qucheng Gao, Zuyi Yang, and Xiao Chen. Nonequilibrium phases of repulsive self-attention: Chaos, attention condensation, and emergent locality. arXiv preprint arXiv:2609.28448, 2026. URL https://arxiv.org/abs/2609.28448.

[109] Zaixiang Zheng, Yi Zhou, and Hao Zhou. Deep equilibrium non-autoregressive sequence learning. In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 11763–11781, 2023. doi: 10.18653/v1/2023.findings-acl.747.

[110] Junhong Wu, Jinliang Lu, Zixuan Ren, Gangqiang Hu, Zhi Wu, Dai Dai, and Hua Wu. LLMs are single-threaded reasoners: Demystifying the working mechanism of Soft Thinking. In *The Fourteenth International Conference on Learning Representations*, pages 4533–4548, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/0895de150660b6c2aef9b0f098cb1dcd-Abstract-Conference.html. arXiv:2508.03440.

[111] Michael Rizvi-Martel, Guillaume Rabusseau, and Marius Mosbach. The illusion of superposition? A principled analysis of latent thinking in language models. arXiv preprint arXiv:2604.06374, 2026. URL https://arxiv.org/abs/2604.06374.

[112] Jie Yu, Yuxi Sun, Zhaoxu Li, Chunxiao Liu, and Xudong Jiang. CoTS: Competitive token selection for latent reasoning in large language models. arXiv preprint arXiv:2610.20642, 2026. URL https://arxiv.org/abs/2610.20642.

[113] Xuan Li, Yining Wang, Yuchen Liu, Guanjun Liu, Delai Qiu, Shengping Liu, Jiaen Liang, Wei Huang, Jun Yu, and Junnan Zhu. LaTER: Efficient test-time reasoning via latent exploration and explicit verification. arXiv preprint arXiv:2605.07315, 2026. URL https://arxiv.org/abs/2605.07315.

[114] Nikita Koriagin, Yaroslav Aksenov, George Bredis, Gleb Gerasimov, Nikita Balagansky, and Daniil Gavrilov. A model with no head and many thoughts. arXiv preprint arXiv:2608.31069; Findings of EMNLP 2026, 2026. URL https://arxiv.org/abs/2608.31069.

[115] Fatih Dinc, Ege Cirakman, Bariscan Kurtkaya, Mert Yuksekgonul, Yiqi Jiang, Mark J. Schnitzer, and Hidenori Tanaka. Ghost mechanism: An analytical model of abrupt learning in recurrent networks. *Physical Review X*, 16(2):021058, 2026. doi: 10.1103/mjcl-lb4x.

[116] Demian Battaglia, Nicolas Brunel, and David Hansel. Temporal decorrelation of collective oscillations in neural networks with local inhibition and long-range excitation. *Physical Review Letters*, 99(23):238106, 2007. doi: 10.1103/PhysRevLett.99.238106.

[117] Jason Z. Kim, Zhixin Lu, Erfan Nozari, George J. Pappas, and Danielle S. Bassett. Teaching recurrent neural networks to infer global temporal structure from local examples. *Nature Machine Intelligence*, 3(4):316–323, 2021. doi: 10.1038/s42256-021-00321-2.

[118] Dayal Singh Kalra, Tianyu He, and Maissam Barkeshli. Universal sharpness dynamics in neural network training: Fixed point analysis, edge of stability, and route to chaos. In *The Thirteenth International Conference on Learning Representations*, pages 55966–56000, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/8c8cd1b78cdae751265c88efc136e5bd-Abstract-Conference.html.

[119] R. L. Marichal and J. D. Piñeiro. Analysis of multiple quasi-periodic orbits in recurrent neural networks. *Neurocomputing*, 162:85–95, 2015. doi: 10.1016/j.neucom.2015.04.001.

[120] Florian Gomez, Ralph L. Stoop, and Ruedi Stoop. Universal dynamical properties preclude standard clustering in a large class of biochemical data. *Bioinformatics*, 30(17):2486–2493, 2014. doi: 10.1093/bioinformatics/btu332.

[121] Rafael V. Stenzinger, Vinícius Luz Oliveira, and M. H. R. Tragtenberg. Shrimp hubs in the Hindmarsh–Rose model. *Chaos: An Interdisciplinary Journal of Nonlinear Science*, 35(2):023148, 2025. doi: 10.1063/5.0239268.

[122] Luiz F. B. Caixeta, Matheus H. P. Gonçalves, M. H. R. Tragtenberg, and Mauricio Girardi-Schappo. Devil's staircase inside shrimp-shaped regions reveals periodicity of plateau spikes and bursts. *Chaos: An Interdisciplinary Journal of Nonlinear Science*, 35(3):033122, 2025. doi: 10.1063/5.0250342.

[123] Giancarlo Benettin, Luigi Galgani, Antonio Giorgilli, and Jean-Marie Strelcyn. Lyapunov characteristic exponents for smooth dynamical systems and for Hamiltonian systems; a method for computing all of them. Part 1: Theory. *Meccanica*, 15(1):9–20, 1980. doi: 10.1007/BF02128236.

[124] Giancarlo Benettin, Luigi Galgani, Antonio Giorgilli, and Jean-Marie Strelcyn. Lyapunov characteristic exponents for smooth dynamical systems and for Hamiltonian systems; a method for computing all of them. Part 2: Numerical application. *Meccanica*, 15(1):21–30, 1980. doi: 10.1007/BF02128237.

[125] John Milnor. On the concept of attractor. *Communications in Mathematical Physics*, 99(2): 177–195, 1985. doi: 10.1007/BF01212280.

[126] Kenji Matsumoto and Ichiro Tsuda. Noise-induced order. *Journal of Statistical Physics*, 31 (1):87–106, 1983. doi: 10.1007/BF01010923.

[127] James P. Crutchfield, J. Doyne Farmer, and Bernardo A. Huberman. Fluctuations and simple chaotic dynamics. *Physics Reports*, 92(2):45–82, 1982. doi: 10.1016/0370-1573(82)90089-8.

[128] Jonathan Kadmon and Haim Sompolinsky. Transition to chaos in random neuronal networks. *Physical Review X*, 5(4):041030, 2015. doi: 10.1103/PhysRevX.5.041030.

[129] Kanaka Rajan, L. F. Abbott, and Haim Sompolinsky. Stimulus-dependent suppression of chaos in recurrent neural networks. *Physical Review E*, 82(1):011903, 2010. doi: 10.1103/PhysRevE.82.011903.

[130] Cristian Bonatto and Jason A. C. Gallas. Periodicity hub and nested spirals in the phase diagram of a simple resistive circuit. *Physical Review Letters*, 101(5):054101, 2008. doi: 10.1103/PhysRevLett.101.054101.

[131] Yihe Dong, Jean-Baptiste Cordonnier, and Andreas Loukas. Attention is not all you need: Pure attention loses rank doubly exponentially with depth. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 2793–2803. PMLR, 2021. URL https://proceedings.mlr.press/v139/dong21a.html. arXiv:2103.03404.

[132] Borjan Geshkovski, Cyril Letrouit, Yury Polyanskiy, and Philippe Rigollet. The emergence of clusters in self-attention dynamics. In *Advances in Neural Information Processing Systems*, volume 36, pages 57026–57037, 2023. doi: 10.52202/075280-2493. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/b2b3e1d9840eba17ad9bbf073e009afe-Abstract-Conference.html. arXiv:2305.05465.

[133] James L. Kaplan and James A. Yorke. Chaotic behavior of multidimensional difference equations. In *Functional Differential Equations and Approximation of Fixed Points*, Lecture Notes in Mathematics, pages 204–227. Springer, Berlin, Heidelberg, 1979. doi: 10.1007/BFb0064319.

[134] Frank Pasemann. Complex dynamics and the structure of small neural networks. *Network: Computation in Neural Systems*, 13(2):195–216, 2002. doi: 10.1080/net.13.2.195.216.

[135] Zihao Fu, Wai Lam, Anthony Man-Cho So, and Bei Shi. A theoretical analysis of the repetition problem in text generation. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35 (14):12848–12856, 2021. doi: 10.1609/aaai.v35i14.17520. arXiv:2012.14660.

[136] Catherine Olsson, Nelson Elhage, Neel Nanda, et al. In-context learning and induction heads. Transformer Circuits Thread, 2022. URL https://transformer-circuits.pub/2022/in-context-learning-and-induction-heads/index.html. arXiv:2209.11895.

[137] Team OLMo, Pete Walsh, Luca Soldaini, Dirk Groeneveld, et al. 2 OLMo 2 furious. arXiv preprint arXiv:2501.00656, 2024. URL https://arxiv.org/abs/2501.00656. Shorter version in the Second Conference on Language Modeling (2025).

[138] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, et al. The Llama 3 herd of models. arXiv preprint arXiv:2407.21783, 2024. URL https://arxiv.org/abs/2407.21783.

[139] Goodfire. Llama-3.1-8B-Instruct-SAE-l19. Hugging Face, https://huggingface.co/Goodfire/Llama-3.1-8B-Instruct-SAE-l19, 2025. Sparse autoencoder on the residual stream after layer 19; revision f6775a22.

[140] Seonglae Cho, Harryn Oh, Donghyun Lee, Luis Rodrigues Vieira, Andrew Bermingham, and Ziad El Sayed. FaithfulSAE: Towards capturing faithful features with sparse autoencoders without external datasets dependency. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 4: Student Research Workshop)*, pages 297–314. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.acl-srw.20. arXiv:2506.17673.

[141] Gonçalo Paulo and Nora Belrose. Sparse autoencoders trained on the same data learn different features. In *The Fourteenth International Conference on Learning Representations*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/3c1fe56b043848b211030c202764c6a7-Abstract-Conference.html. arXiv:2501.16615.

[142] Liquid AI. LFM3-1.7B. Hugging Face, https://huggingface.co/LiquidAI/LFM3-1.7B, 2026. Model repositories; revisions 40ebe0aa (LFM3-1.7B) and 545fc8f4 (LFM3-1.7B-Base), used on 28 October 2026.

[143] Fredrik Johansson. Arb: Efficient arbitrary-precision midpoint-radius interval arithmetic. *IEEE Transactions on Computers*, 66(8):1281–1292, 2017. doi: 10.1109/TC.2017.2690633.

[144] R. Krawczyk. Newton-Algorithmen zur Bestimmung von Nullstellen mit Fehlerschranken. *Computing*, 4(3):187–201, 1969. doi: 10.1007/BF02234767.

[145] Ramon E. Moore, R. Baker Kearfott, and Michael J. Cloud. *Introduction to Interval Analysis*. Society for Industrial and Applied Mathematics, Philadelphia, PA, 2009. doi: 10.1137/1.9780898717716.

[146] Siegfried M. Rump. Verification methods: Rigorous results using floating-point arithmetic. *Acta Numerica*, 19:287–449, 2010. doi: 10.1017/S096249291000005X.

A Related work in detail

Language models have been studied as dynamical systems at several levels, and the closest studies differ from ours in what they compute. Lyapunov-like quantities have been used to measure the growth of small perturbations across the layers of a single forward pass [5], and sensitivity to small perturbations in reasoning traces [7] and in token rings resampled by a model's own conditionals, with a test that separates the construction from the model [26], but these studies measure chaos-like sensitivity without computing orbits or bifurcations. Perturbations at the rounding level either vanish or are amplified through the layers of a single pass, and near a tie between the two leading tokens they change the predicted token [6]; we follow such differences around the feedback loop (Section 6). Output statistics change abruptly with sampling temperature, and several studies analyze the changes as phase transitions or criticality [8–11]. These analyses concern output statistics; they are not bifurcation analyses of a defined map. Closest to ours in subject, a study of eleven models sorts sampled continuations into classes from frozen and periodic to intermittent and noisy across temperature, fits a noisy effective map to one hidden-state coordinate that predicts repetition, and steers along it [12]; that map is fitted to sampled trajectories, and no orbits, multipliers or bifurcation curves of the model's own map are computed. Repetition loops are a known failure of decoding [27, 28], successive paraphrasing converges to attractor cycles [13] or contracts toward a stable region of an embedding space, while repeated negation keeps wandering [29], and multi-turn conversations between models settle into attractor states [14]. These attractors are behavioral: they are observed in generated text rather than computed as invariant sets of a defined map. Feeding a model its own output during training degrades it over generations [99]; we study the inference-time loop of a model with fixed weights. The attractors of recurrent networks have been related to the difficulty of training them [100], and fixed points with their linearizations have long been used to reverse-engineer trained recurrent networks [89, 90]. Equilibrium, looped and recurrent-depth models iterate a block to a fixed point or for many steps, and their stability has been analyzed through Jacobians and fixed points [3, 15–18, 93, 101–104], and recent analyses of looped reasoners follow their internal dynamics, their settling, marginal or drifting regimes and slowly decaying oscillatory modes [105–107]; these studies examine fixed-point stability, finite-depth dynamics and nonstationary regimes. Here we continue full-state periodic branches and classify their bifurcations in specified frozen-model feedback maps. A minimal recurrent self-attention with a repulsive value map becomes chaotic [108], but it is a constructed model, not a trained one. A separate line predicts the answers of looped models from the geometry of score margins along completed finite-depth trajectories [19], without fixed points or Jacobians. The closest methodological precedent is a derivative-based fold analysis of a small task-trained looped transformer, whose passage and spectral-gap predictions are unfitted [21]. That model was trained for a task, whereas ours is a released pretrained language model, and our analysis also differs in its objects: the fold of a primitive periodic orbit, the full-channel passage of a two-step map, and the accumulation law of the windows, whose prefactor is set by the global reinjection. Expected-embedding feedback itself is used in training [20], solved for as equilibria in non-autoregressive decoding [109], and, training-free, in pretrained models [1, 110], next to continuous latent reasoning [2]; Backour shows that the softmax readout and autoregressive feedback form a dynamical system that either amplifies small differences until one component of a mixture dominates or contracts different mixtures together, with pretrained rollouts mostly on the amplifying side [44]; that analysis concerns the fate of mixtures, while ours computes the invariant sets, bifurcations and normal forms of the full released map and their consequences for reproducibility, and whether they exploit the superposition the mixtures could carry has been questioned [111]. Newer methods vary what is fed back. CoTS keeps the sampled token and mixes in the expected embedding in proportion to the competition between the two leading candidates

[112], the near-tie that makes our loop's readout expand (Section 3); since at least half of each input is the discrete token, our orbits do not carry over to it directly. LaTER feeds back a linear projection of the final hidden state [113], and a trained projector lets a model roll out in a learned embedding space without its output head [114]. For such hidden-state loops the derivative of Section 3 takes a different form, without the readout covariance, while the tools used here (orbits, multipliers, continuation and normal forms) apply unchanged. Fixing the prefix and the position turns that feedback into the autonomous map we analyze.

In neural systems more broadly, saddle-node ghosts produce quantitative delays in task-trained recurrent networks [115]. Finite-order Feigenbaum ratios have been measured in neural rate and mass models [62, 116] and, through period 128, in a transformer trained on logistic-map trajectories [63]. Learned networks reproduce period-doubling routes [117], and so does gradient descent itself [118]. Normal forms and continuation have been applied to recurrent neural maps, including their quasi-periodic orbits [119], and shrimp-shaped windows and periodicity staircases appear in neuron models [120–122]. Our second control parameter comes from activation steering along contrastive directions [37, 38]. The dynamical-systems tools are standard: numerical continuation and normal forms [49, 50], for maps in particular [57], type-I intermittency [58], shrimps and hubs [4, 64, 65, 67], and period-doubling universality [24].

B Methods

Precision. Every result on the soft-token map labeled float64 is computed in float64 arithmetic throughout: embeddings, activations, attention keys and values (the prefix cache), normalizations (patched, since the library computes RMS normalization in float32 by default), rotary tables, logits, the softmax and the fed-back token. The released checkpoints store their weights in bfloat16, and a bfloat16 number is exactly a float64 number, so converting the weights is exact. For large models the decoder's linear weights stay in their stored bfloat16 and are converted to float64 at each use, which gives the same map as converting them once (bitwise on Qwen2.5-0.5B). The designated float64 computations use the patched normalization and rotary paths; some diagnostic traces, such as the step records of the gain split, are serialized in float32, as recorded in their artifacts. The dense checkpoints store bfloat16 weights (checked against their safetensors metadata), and the mixture of experts stores some tensors in float32, which we restore exactly; float64 does not make the weights more precise than stored. One exception concerns the 70B-class models, which are spread over two GPUs. In the two-GPU runs made before the fix described here, the final RMS normalization was computed in float32: the hooks that move tensors between the GPUs keep a layer's original code, so our float64 version of that layer was bypassed. This makes the computed map rough at about 10^{-8} of the state (its fixed points wander by 10^{-8}, against 10^{-15} on one GPU) and can matter for exponents measured with perturbations of 10^{-6} and below. The affected runs are the 70B-class sweeps and their plateau tests, the first Qwen2.5-72B prompt test, the 70B steering maps and the first zoom, the gain splits of the 70B-class models, and the start-state test, which spread all fourteen models over two GPUs; Appendix I says which results we recomputed. Derivatives are automatic-differentiation tangents of this float64 computation: free of finite-difference truncation, but computed in floating point. We call a result exact only when it comes from interval or rational arithmetic, as in the proofs of Appendix L.

Statistical units. The cells of a parameter map are evaluations of one deterministic function at neighboring parameters, not independent experiments, and neighboring cells are strongly correlated. Their counts and shares therefore describe the area of a region on a chosen grid, not a prevalence over

models or prompts. Comparisons between chaotic and periodic cells use temperature rows as units: intervals come from bootstraps over rows, and permutation p-values within rows are descriptive. Where a random sample of cells is drawn, its seed and sampling frame are recorded with the run. In tests on problems the problem is the unit, paired across conditions; where problems share a batch, intervals are clustered by batch, and McNemar tests, which treat problems as independent, are reported alongside them. Batch-level randomization tests, which flip the sign of each batch's summed paired difference, give the same conclusions for every contrast of the Soft Thinking and reasoning-model tests (`tools/cluster_tests.py`, `docs/cluster/`). Soft Thinking's contrasts with replay, greedy decoding and 16 soft steps all have batch-level $p \leq 5 \times 10^{-6}$, the resolution of 200 000 draws, over 77 batches. In DeepSeek-R1-Distill-Llama-8B, soft thinking changes 7.4 points more answers than plain reasoning (batch level $p = 0.033$, against 0.029 for McNemar's test, over 16 batches) and answers correctly 4.7 points less often (batch level $p = 0.063$, against 0.096). Every interval names its resampling unit. A null result is stated together with the size of effect its interval excludes; effects smaller than that remain compatible with the data.

Independent implementations. Where an independent implementation reproduced a result, a second implementation of the map (and, where stated, of the solver or the normal-form computation), written separately by the second coding agent from the model's published architecture (`codex/crosscheck` in the repository), recomputed it. The two implementations share the stored weights, the tokenizer and the specification of the map, not code; hardware and library versions were not varied on purpose. The second implementation, a manual float64 forward pass of the stored checkpoint that uses neither our tools nor the Transformers library, also reproduces the chaos: with the scans' protocol (fixed start, 300 transient steps, twin offset 10^{-6}), its 300-step exponents at two chaotic cells of Figure 11, $(0.065201, -0.155998)$ and $(0.075475, -0.133454)$, are 0.098 and 0.354, against 0.113 and 0.359 in the scan (0.106 and 0.332 over 1000 steps), and at the certified period-8 orbit it gives -0.564 per step over 1000 steps (-0.579 over 300), the value $\ln(0.0110)/8$ that the orbit's leading multiplier predicts (`tools/independent_chaos_check.py`).

Phase maps. Each cell of a phase map is one parameter pair. Unless stated otherwise, we discard 300 transient steps and measure over the next 300. The largest Lyapunov exponent comes from a twin trajectory at offset 10^{-6} that is renormalized to that offset after every step, Benettin's method with a finite-difference tangent [123, 124]; it estimates over a finite horizon the exponent that the multiplicative ergodic theorem defines as a limit [31]. At the 448 random cells of the horizon check (Appendix K), an automatic-differentiation tangent vector carried by forward-mode differentiation reproduces the twin exponents at offsets 10^{-6} and 10^{-9} within 0.0041 (median 0.0004), with the same 178 chaotic labels, and at six chaotic cells the four leading exponents from QR iteration include exactly one positive one (0.06 to 0.42; the next is -0.64 to -1.10). The period is the smallest $q \leq 32$ (48 in the later zooms) for which the last state returns to within $10^{-8} \max(1, \|x\|)$ of the state q steps earlier (10^{-5} to 10^{-4} in float32 maps), and the attractor is identified by its leading-token loop, up to rotation. We call a cell chaotic when this finite-horizon exponent exceeds 10^{-3} and no period is found; the label is a finite-horizon diagnostic, not a proof of chaos. The threshold matters little: raising it to 5×10^{-3} lowers the chaotic share of the map of Figure 1 from 30.1 to 29.8 %. We call what the trajectory from a cell's fixed start settles on the cell's attractor. The label is operational: it does not show that the set attracts initial conditions of positive measure, as Milnor's definition requires [125], and a long chaotic transient can carry a chaotic label [48]. A loop is at the edge of chaos when $|\lambda| \leq 0.1$. Maps draw the chaotic cells by this criterion: in color by their exponent where a map shows the exponent, in black where it shows periods or attractor text. Two figures

color every cell whose exponent exceeds 10^{-3}, also where a period is detected, as their captions state: the prompt maps (Figure 97) and the followed domain of the period-5 window (Figure 46). Panels whose legend reads 'Lyapunov > 0.005' follow an older convention that marks every cell whose exponent exceeds 0.005, whether or not a period was found; on the five maps we redrew with the criterion, the two differ in at most 0.56 % of a map's cells. A 256^2 map takes about 100 minutes on one H100 GPU in float64. Every cell starts from the same state, so where attractors coexist, a map shows the one whose basin contains that state. Labels are text, not dynamical identities: one continuous orbit changes its label where two tokens tie for the maximum. We therefore call two labeled regions different attractors only when the contraction rate jumps between them or when float64 orbits of both exist at one parameter value.

Resolving the network factor through the stack. Hooks read, at the soft token's position, the residual stream entering each block, the outputs of its attention and MLP sublayers, and the final hidden state before and after the final normalization. They do so for the base row and for a twin at relative offset 10^{-7} along the evolving tangent direction. The split of the network factor into $A_L - N_L$ and the final normalization is exact.

In the pre-normalization architectures of Qwen2.5, Qwen3.5, Llama-3 and OLMoE, the residual stream adds the sublayer outputs, $\delta h_L = \delta x + \sum_s \delta o_s$, and this sum closes to within 4×10^{-7} relative in every model and temperature. OLMo-2 normalizes each sublayer's output before adding it, so the sum does not close there (relative errors of 0.9 to 150). For OLMo-2 we therefore use only quantities that do not depend on it: the factor itself, A_l, N_l and the knockouts.

A knockout holds one sublayer's output at the soft token at its value in the base row while the input carries the twin's change, which blocks the change's path through that sublayer. The knockout rows reproduce the plain twin's network factor within 5.1×10^{-7}.

For the mixtures of experts, a third row receives the twin's input but uses the base row's experts in every layer, weighted by its own router probabilities, renormalized over them where the router renormalizes its top-8 weights (Qwen3.5) and not where it does not (OLMoE). Each step also records whether the twin's set of experts differs from the base row's in any layer, and the base row's smallest gap between the eighth and ninth largest router logits over all layers. Qwen3.5's linear-attention layers carry recurrent states, which are restored after every step, as in the soft-token loop.

False chaos and real chaos. A twin exponent measures how fast the difference between two trajectories grows, so anything that adds a difference of its own can imitate growth. Suppose that both rows receive independent noise of relative size σ at every step, from rounding or from the float32 normalization described above, so that their difference receives noise of root-mean-square size $\sqrt{2}\,\sigma$. We use a simplified model: a difference survives along one fixed leading direction, growing by a constant factor e^{λ} per step, while the other directions contract within a step; the noise component along the leading direction is neglected, and the norm of the rest is replaced by its root-mean-square value. Let c_t be the cosine of the angle between the renormalized twin offset and the leading direction, and $r = \sqrt{2}\,\sigma/\varepsilon$ the noise relative to the offset ε. Without these simplifications one step gives $c_{t+1}^2 = (e^{\lambda}c_t + \eta_{\parallel})^2/[(e^{\lambda}c_t + \eta_{\parallel})^2 + \|\eta_{\perp}\|^2]$, where $\eta_{\parallel}$ and $\eta_{\perp}$ are the components of the difference noise, in units of ε, along and across the leading direction; the model sets $\eta_{\parallel} = 0$ and $\|\eta_{\perp}\| = r$, so that $c_{t+1}^2 = e^{2\lambda}c_t^2/(e^{2\lambda}c_t^2 + r^2)$. While $r < e^{\lambda}$, the angle settles at $c^2 = 1 - r^2e^{-2\lambda}$, and the twin grows by e^{λ} per step. Otherwise $c \to 0$, and the twin grows by r per step, whatever the dynamics. In the model the twin therefore measures $\max\{\lambda, \log r\}$, and for the network we use this as an approximation:

$$\lambda_{\text{meas}} \approx \max\{\lambda,\ \log(\sqrt{2}\,\sigma/\varepsilon)\}. \tag{26}$$

Above the offset $\sqrt{2}\,\sigma e^{-\lambda}$ the model gives the true exponent. Below it, a noise term takes over, rising by one nat for every factor e by which the offset shrinks. In the model, noise can thus make a contracting orbit look chaotic when the offset lies below $\sqrt{2}\,\sigma$, but it cannot make a chaotic orbit look stable. That one-sided conclusion holds only in the model: rounding can also suppress separation, and noise can change the sampled dynamics or move a trajectory into another basin (below), so we compare offsets and tangent estimates rather than use the approximation as a one-sided guarantee.

The model assumes a single fixed leading direction with constant amplification, contraction within one step and no noise along the leading direction, so we test it in the network (Figure 64). The test uses Qwen2.5-0.5B in float64 at three cells of the philosophy-steered map: a chaotic cell, the weakly contracting period-3 loop of the coexistence example (Section 5), and a strongly contracting orbit; each runs 400 transient and 800 measured steps. In clean float64 the exponents do not depend on the offset over eight decades: $+0.092$ to $+0.099$, -0.123 to -0.128, and -0.843 to -0.862 for offsets from 10^{-12} to 10^{-4}.

We then add Gaussian noise with $\sigma = 10^{-10}$, 10^{-8} or 10^{-6} to every row's next soft token.

- Where the noise dominates, all three cells give the same false exponent, which Eq. (26) matches within 0.006: for example 4.951 to 4.953 against $\log(100\sqrt{2}) = 4.952$ at $\sigma = 10^{-10}$ and $\varepsilon = 10^{-12}$.
- Where the offset lies well above the crossover, the prediction matches within 0.047.
- In the crossover itself, two offsets per cell and noise level, the measured exponent lies above the prediction by up to 0.43, because noise injected one step earlier has not yet contracted.
- Noise lowered no chaotic estimate below 0.087 (clean 0.092 to 0.099). It lowered the contracting estimates by at most 0.048, which makes them look more stable, not less.

A free twin that is never renormalized shows the difference directly (Figure 64d). Started at 10^{-12}, the chaotic twin grows exponentially until its distance saturates near 10^{-1} after about 240 steps. Under $\sigma = 10^{-8}$, the period-3 twin instead jumps to a plateau at about 2×10^{-8} (the model predicts 3×10^{-8}) and stays there.

The float32 final normalization, computed as in the early two-GPU runs and without added noise, reproduces the defect.

- At offsets from 10^{-10} to 10^{-8}, both contracting orbits show positive exponents, up to 2.03 for the period-3 loop.
- At 10^{-12} and 10^{-11}, rounding makes the twin and base rows identical in 454 to 771 of the 800 steps, leaving too few measurable steps to read.
- From 10^{-6} up, the contracting orbits read -0.177 to -0.123 and -0.853 to -0.783 (clean -0.124 and -0.848), and the chaotic one 0.059 to 0.078 (clean 0.094).

The last shift is of another kind. The defect changes the map itself, by about 10^{-8} at every step, not only the measurement, and the chaotic exponent of the changed map is lower, though still positive. The spurious Qwen2.5-72B window of Appendix I carries the signature of Eq. (26). Its exponent was negative at 10^{-5}, near zero at 10^{-6} and positive at 10^{-7}, and its orbits are periodic when computed in float64 throughout.

Hence the rule we apply: a chaotic label needs an exponent that does not depend on the offset. We check this with three offsets, or with an automatic-differentiation tangent vector, which needs no offset. Two limits remain.

- The argument concerns noise in a measured difference. Noise acting on the dynamics can change which attractor a trajectory follows, and in some maps it lowers a positive exponent, the phenomenon of noise-induced order [126]; see also [127]. Small per-step errors can accumulate, or be amplified near weak stability and in nonnormal dynamics; a constant error d moves the fixed point of $x' = \rho x + d$ by $d/(1-\rho)$, which grows without bound as $\rho \to 1$. Whether such errors change the attractor depends on the response and on the basin geometry; the per-step error alone neither bounds the displacement nor excludes a basin switch.
- Far below the rounding resolution, two rows can become identical. A naive estimator would then report infinite contraction, a false negative; ours reports the step as unmeasurable. The offsets of our maps lie far above the resolution of their arithmetic.

How a fold or flip curve is found. Every curve in this paper is found in the same five steps, all on the full map in float64. (i) *An orbit.* At one cell inside a window the loop is iterated from the fixed start until it settles; the period p is the smallest return time of the full state, $\|F^p(x) - x\|$ small, and Newton's method on $F^p(x) = x$ with forward-mode Jacobians through the network refines the orbit to residuals near 10^{-14}. (ii) *A branch.* The orbit is continued in one parameter ($\log_{10} T$ at fixed α, or α at fixed T) by pseudo-arclength steps, and at every point all multipliers are computed from the monodromy matrix $M = DF^p$. A fold shows up as a turning point of the branch, where the orbit meets a saddle and the branch turns back (a multiplier passes $+1$); a flip shows up as a real multiplier crossing -1, where the orbit doubles. (iii) *A curve.* From the boundary point found on the branch, the bifurcation itself is continued in the plane of the two parameters. The unknowns are the orbit point and the two parameters; the equations are $F^p(x) - x = 0$ and one scalar test function $g = 0$, the last component of the solution of the bordered linear system with matrix $[[M \mp I, v], [w^\top, 0]]$, which vanishes exactly when M has the multiplier ± 1 (upper sign: fold) [49, 50]. Newton's method with Jacobian-vector products through the network corrects each pseudo-arclength step. (iv) *A gate at every point.* A point is drawn only if its residual is at most 10^{-10} and the critical multiplier of the dense monodromy is real and within 10^{-7} of ± 1; rejected points are counted and never drawn. (v) *Cusps and the check against the map.* Where a fold curve's direction in the plane reverses, a cusp candidate is located by the sign change of the oriented quadratic coefficient of the fold normal form, and accepted only if the cubic coefficient and the unfolding determinant are nonzero. Finally, where the attractor followed cell by cell ends, a computed curve must be near (the boundary gates of the figures). For the 0.5B model one curve point takes about 50 s on one H100 GPU.

Periodic orbits. We solve $F^p(x) = x$ by Newton's method with forward-mode Jacobians. The multipliers are the eigenvalues of the monodromy matrix, the ordered product of the step Jacobians along the orbit, which plays the role of the Floquet matrix of a flow [50]; where the dense matrix is not float64, we refine them by a short Arnoldi factorization with float64 Jacobian-vector products.

Continuation. Periodic branches are continued in one parameter by pseudo-arclength continuation, and fold and flip curves in two parameters with minimally augmented systems [49, 50]. The test function g of the bordered system $[[M \mp I, v], [w^\top, 0]]$ is obtained by GMRES (generalized minimal residual iterations) with Jacobian-vector products, and its gradient by reverse-mode differentiation through a Jacobian-vector product. When the parameter tangent reverses, we flag a cusp candidate and test it with the normal form. To locate a codimension-two point on a continued curve (`tools/cusp_locate.py`), we bracket a sign change between two curve points, then retrace the

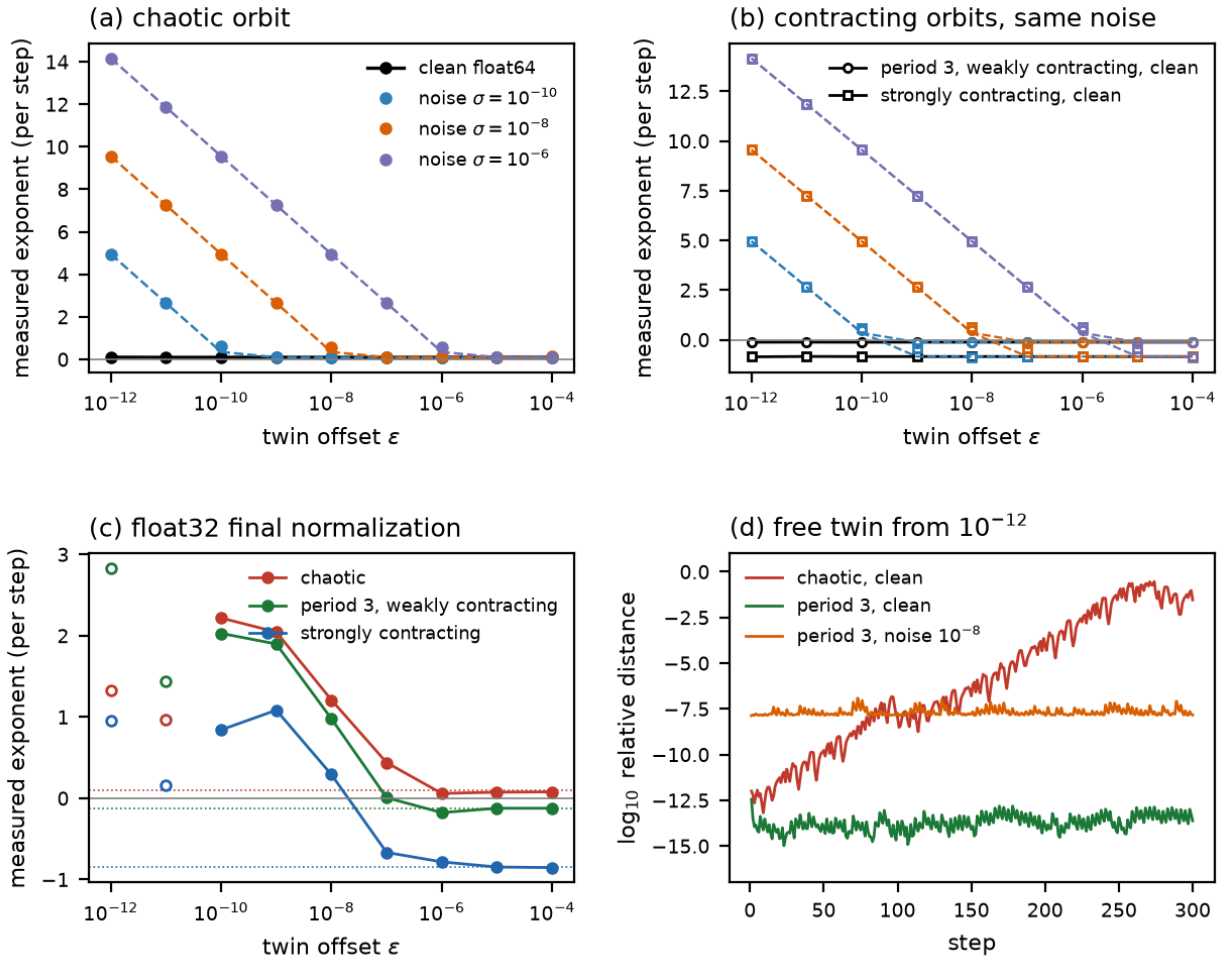

Figure 64: False chaos and real chaos in a twin test (Qwen2.5-0.5B-Instruct, float64, philosophy direction at block 12; 400 transient and 800 measured steps). The twin exponent is the mean logarithmic growth of a second row renormalized to the relative offset ε after every step. (a) A chaotic cell ($\alpha = 0.0652$, $\log_{10} T = -0.156$): clean float64 (black) and Gaussian noise of relative size σ added to every row's next soft token (colors); dashed, Eq. (26) with the clean exponent. (b) The same for a weakly contracting period-3 loop ($\alpha = 0.0642$, $\log_{10} T = -0.156$; circles) and a strongly contracting orbit ($\alpha = 0$, $\log_{10} T = 0.18$; squares): below the offset $\sqrt{2}\,\sigma$ both read as chaotic. (c) The float32 final normalization of the early two-GPU runs, without added noise: false positive exponents near float32 resolution. Open markers mark offsets at which the twin and base rows became identical in more than half of the steps; dotted lines mark the clean exponents. (d) A free twin, started at 10^{-12} and never renormalized: real chaos grows; under noise the contracting loop's twin jumps to the noise level and stays there; in clean float64 it stays near 10^{-14}, the rounding floor of the map.

curve from the left end with steps of a quarter of the bracket. We do this for three rounds and finish with a secant step. The function whose sign changes is the oriented quadratic coefficient $b \operatorname{sign} a$ for a cusp, $|\mu_2| - 1$ for a second multiplier μ_2 leaving the unit circle, and the projected parameter tangent for a turn of the curve. At a cusp we then evaluate c and the unfolding matrix at the located point. Elsewhere, the multipliers and the angle between their eigenvectors type the point: a double multiplier whose eigenvectors coincide is a Jordan block. On maps with known answers the tool recovers a cusp to 10^{-6} (with $c = -1$ and $\det K = 1$) and a fold-flip point.

Normal form and passages. Consider a fold of the two-step map $G = F^2$, with right and left eigenvectors r and ℓ normalized by $\|r\| = 1$ and $\ell^\top r = 1$. Along the critical direction the reduced map is $u' = u + a\varepsilon + bu^2$, with $a = \ell^\top \partial_\alpha G$ and $b = \frac{1}{2}\ell^\top D^2 G[r, r]$. We compute b by central differences of float64 Jacobian-vector products over a step-size plateau (later cross-checks use nested automatic differentiation), and a includes the steering's effect on the prefix. The signed coefficients are kept in the original eigenvector gauge, with ε measured in the free parameter (α or $\log_{10} T$, whose units apply to a). On maps with known normal forms the code recovers ab, the cubic coefficient c with its center-manifold term and the two-parameter unfolding matrix, and it traces the analytic fold and flip curves of the Hénon map (`tests/test_bifurcation_curves.py`). For passages we set $v = \operatorname{sign}(b)\, u$ and take the ghost side $p - p_c = \operatorname{sign}(ab)\, \epsilon$ with $\epsilon > 0$, so that $v' = v + |a|\epsilon + |b|v^2$ to quadratic order. A full passage starts at $v = -U$ and a midpoint passage at $v = 0$; both end at the first strict crossing $v > U$, with $U = 0.1/|b|$. We compare passage counts with the iterated quadratic map and with the arctangent formula [58].

Period-doubling cascade. Along a cut in one parameter, we locate the flip of each period-p orbit as the root of $\operatorname{Re} \mu_c + 1$, where μ_c is the multiplier with the smallest real part, using the Illinois variant of regula falsi with warm-started orbits. All but one or two multipliers are small, so Newton's method on $F^p(x) - x$ needs the monodromy only on a six-dimensional dominant subspace. We obtain that subspace by subspace iteration with forward-mode Jacobian-vector products chained along the orbit, so no dense Jacobian is formed, and the residual is always the exact float64 $F^p(x) - x$. Each root is classified before it is accepted. The critical multiplier must be real and equal to -1, with a Ritz residual below 10^{-6}, and the other computed Ritz values (six in all, estimates of the dominant spectrum) must lie inside the unit circle. We then reach the period-$2p$ orbit by forward iteration inside the next band. On the logistic map embedded in six dimensions, the procedure reproduces the first four flip parameters to 10^{-9}. This classifies the crossing, not the bifurcation: the cascade flips of Section 9 and the flip curves of Section 11 are located where a real multiplier crosses -1 and confirmed by the doubled orbits beyond them. Flip coefficients were computed only at the two flips of the high-temperature bubble (Table 18), so the other flips are not classified as super- or subcritical. Nor were their crossing speeds computed: a generic flip also needs a nonzero crossing speed of the multiplier and a nonzero cubic coefficient, so the cascade's flips are identified by a simple real multiplier at -1 and the doubled orbit beyond each.

C Theory background, network contraction and the dimension of the chaos

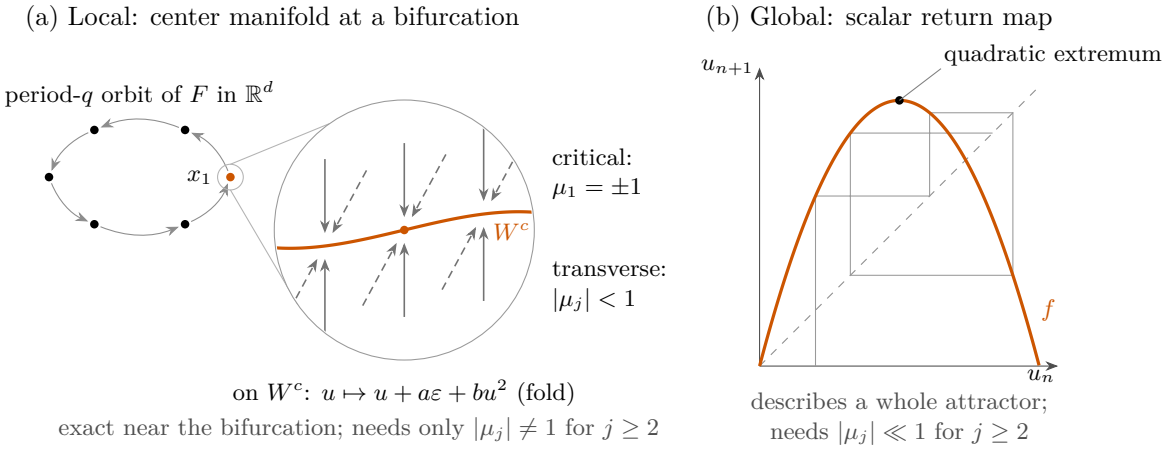

Figure 65: Two reductions to one dimension (Section 3). (a) Near a bifurcation of a period-q orbit of the soft-token map (dots; arcs: the map), the dynamics of F^q near a point x_1 of the orbit reduce to its center manifold W^c, tangent to the critical eigenvector of the monodromy matrix (multiplier $\mu_1 = \pm 1$; drawn with one critical direction, as at a fold, a flip or a cusp; a 1:1 resonance or a fold-flip point has two). The other multipliers lie off the unit circle; drawn here inside it, the transverse directions contract onto W^c (dashed arrows stand for the remaining dimensions). The reduction is exact near the bifurcation whatever the size of the other multipliers; the normal forms of folds, flips, cusps and 1:1 resonances rest on it. (b) A global scalar return map $u \mapsto f(u)$ with a quadratic extremum, the diagonal (dashed) and a few iterates (gray). It describes a whole attractor only when the other multipliers are small.

Why the structures of low-dimensional maps appear. The map acts on a space of hundreds to thousands of dimensions, yet two reductions make low-dimensional theory apply (Figure 65). Near a bifurcation of a periodic orbit the center manifold theorem reduces the local dynamics to the critical directions, whatever the size of the other multipliers, as long as they stay off the unit circle; the normal forms are the leading terms of this local reduction [50]. Our normal forms of folds, flips, cusps and 1:1 resonances rest on this reduction alone. A scalar description of a whole attractor needs additional, global evidence. The small non-critical multipliers of the periodic orbits we computed show strong transverse contraction near those orbits and motivate that description: they range from about 0.1 at short periods to below 10^{-13} along the cascade (Sections 8 and 9 and Appendix E), and along the cascade at $\log_{10} T = -0.156$ the next modulus falls from 0.12 at the first flip to 0.045, 7×10^{-5}, 4×10^{-7} and 2×10^{-13} at the next four. Temperature, through the gain of (6), and steering act as two parameters. One-dimensional families of maps with a quadratic extremum can exhibit fold and flip curves, nondegenerate cusps, period-doubling cascades whose interval ratios tend to Feigenbaum's δ, and saddle-node windows that accumulate on a fold with $d_N \propto N^{-2}$ (Figure 66 sketches the objects we locate); we test these predictions on specific branches and parameter sequences in Sections 8 to 11. The organization of shrimps needs more, either a second turning point of the reduced map or a second dimension [4, 65]. We test the consequences where we state them: the multipliers of each orbit, the interval ratios, and the normal-form predictions of passage times and coexisting orbits. Where the second multiplier is not small, as in the high-temperature bubble of Section 9 (next moduli 0.28 and 0.38) and along the period-5 curves of Section 11 (up to 0.39), the global reduction is only a rough guide, while the local normal forms still apply.

Coexisting attractors constrain one-dimensional descriptions. A classical theorem says what a one-dimensional reduction cannot do. If an interval map is three times differentiable, unimodal and has negative Schwarzian derivative, $Sf = f'''/f' - \frac{3}{2}(f''/f')^2 < 0$, the immediate basin of every attracting periodic orbit contains the turning point or an endpoint of the interval [47]; if basins reaching an endpoint are excluded, at most one attracting periodic orbit remains. The soft-token map has coexisting attracting orbits in several places. Near each located cusp with $c < 0$, two attracting orbits of the same period coexist with a saddle between them (periods 4 and 6, Appendix E; period 7, Section 10). Inside the hostile period-7 body, the period-7 loop coexists with a fixed point, which the fixed start reaches at isolated cells. Such coexistence constrains a scalar interpretation: unless one attractor's basin reaches an end of the scalar interval, the effective one-dimensional dynamics cannot be a single unimodal map with negative Schwarzian derivative, and needs a second turning point, a stretch of positive Schwarzian derivative, or a second dimension. A second turning point is also what a doubly superstable shrimp center requires.

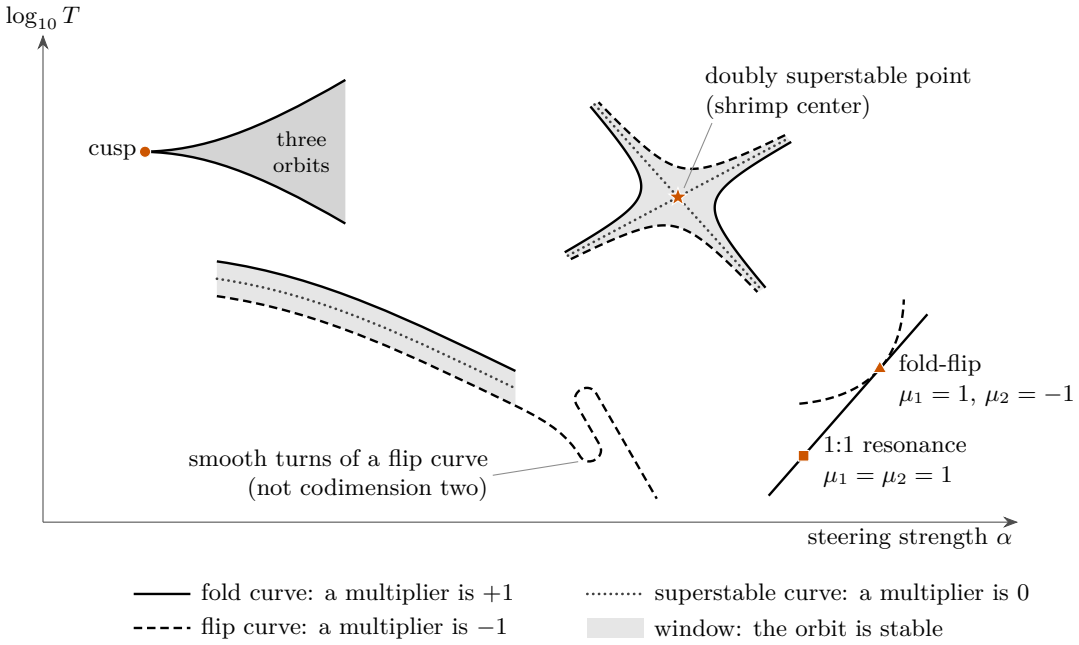

Figure 66: Schematic of the bifurcation objects in the plane of steering strength α and $\log_{10} T$ (not data; positions and shapes are illustrative). On the curves a multiplier of the orbit equals $+1$ (fold), -1 (flip) or 0 (superstable); gray marks windows, where the orbit is stable. Between two fold curves that meet in a cusp, three orbits coexist. A flip curve can turn tightly but smoothly, which is not a codimension-two point. On a fold curve, a 1:1 resonance is a point where a second multiplier also reaches $+1$, and a fold-flip point is one where a flip curve touches it. Two superstable curves of one orbit cross at a doubly superstable point, the center of a shrimp [4, 65]; near it the multiplier is approximately proportional to the product of the distances from the two curves, so the window is a cross whose sides alternate between fold and flip curves. Appendix E locates cusps, 1:1 resonances, a fold-flip point and smooth turns of a flip curve in the soft-token map, candidate superstable curves (sign changes of the leading real multiplier) appear in its exponent maps (Section 10), and Section 17 states which parts of this anatomy are not verified.

What theory allows. Chaos and periodic windows occur in many nonlinear feedback systems. In a specified ensemble of random recurrent networks with asymmetric couplings, chaos sets in above a critical gain, sharply in the large-network limit, and finite networks pass through limit cycles of increasing complexity on the way [22]; how the transition looks depends on the architecture and the

transfer function [128]. Discrete-time random networks, $x(t+1) = \tanh(gJx(t))$, iterate a map as our loop does: their fixed point loses stability when the gain times the spectral radius of the couplings exceeds one, the first bifurcation of large networks is a Neimark–Sacker bifurcation, and chaos follows by a quasi-periodic route, through a torus and frequency locking, over a range of gain that narrows as the network grows [23]. Our largest model narrows differently: Qwen2.5-72B passes from a confident echo to a nearly uniform fixed point within 0.03 decades of temperature, and its chaos is confined to narrow bands, whereas in random networks chaos is generic above the transition. External input can suppress that chaos [129], and deep networks at initialization show an order-to-chaos transition in signal propagation [97, 98]. In our map, temperature plays the role of an inverse gain and the steering vector that of an input. Section 3 makes the correspondence concrete through an exact chain-rule factorization, with the softmax derivative $(\operatorname{diag} p - pp^\top)/T$ in the place of the slope of the transfer function; the identification with the random-network transition remains an analogy, not a transfer of that ensemble's theorem. Feigenbaum scaling holds for a class of one-dimensional maps with a quadratic maximum [24], and periodic windows, shrimp-shaped domains with their cusps and spirals, and doubling cascades are found in many dissipative two-parameter systems [4, 64, 67, 130] and in experiments [88]; intermittency near folds is classical [58].

Contraction and expansion in the network. Analyses of deep attention identify contracting tendencies: pure self-attention collapses token representations toward a common vector with depth [131], and idealized attention dynamics form clusters [132]. Residual connections and MLPs counteract the collapse [131]. Which tendency wins along an orbit is a property of the trained weights that only computation settles. In every dense model we mapped, both outcomes occur: the same map has negative largest exponents in some regions of its parameter plane and positive ones in others.

How many dimensions does the chaos have? Near a periodic orbit, the multipliers measure how many directions the dynamics keeps (Section 3). For chaotic attractors the corresponding measure is the Kaplan–Yorke dimension [133], $D_{\mathrm{KY}} = j + \sum_{i \le j} \lambda_i / |\lambda_{j+1}|$ with j the largest index at which the partial sum of the ordered Lyapunov exponents is still nonnegative; with one positive exponent and $\lambda_1 + \lambda_2 < 0$, $D_{\mathrm{KY}} = 1 + \lambda_1/|\lambda_2|$. We estimated the four leading exponents by QR iteration (Appendix B; 1000 transient and 2000 measured steps, float64) at 35 completed cells: 28 chosen by percentile of their largest exponent before any of their spectra were seen (of 30 requested), and seven from the earlier horizon check, six chaotic and one periodic control. At this horizon 31 have a positive leading estimate and none has two positive estimates. These are finite-time estimates, not asymptotic dimensions. They grow with temperature. Along the hostile temperature line, where we fit a scalar return relation (Section 10), eight chaotic cells have $D_{\mathrm{KY}} = 1.008$ to 1.16, consistent with nearly one-dimensional chaos. At fourteen chaotic cells of the headline maps ($\log_{10} T$ from -0.17 to -0.04) it is 1.03 to 1.65, rising with λ_1. In the high-temperature region ($\log_{10} T$ from -0.02 to 0.13), nine chaotic cells have $D_{\mathrm{KY}} = 1.37$ to 2.02; at two of them the sum of the first two finite-time estimates is nonnegative, which indicates noncontracting two-dimensional area at this horizon, although only one estimated exponent is positive. A scalar return map describes the low-temperature chaos well and the high-temperature chaos poorly, consistent with the large second multipliers of the high-temperature bubble of Section 9. The number of competing tokens does not set the dimension. The readout's trace-to-leading-eigenvalue ratio, $\operatorname{tr}\Sigma_p / \lambda_{\max}(\Sigma_p)$ averaged along the orbit (with $\lambda_{\max}$ estimated by eight power iterations), counts the directions in which the readout amplifies comparably. D_{KY} never exceeds it and rises with it within each group of cells (Spearman 0.72 over 17 philosophy cells and 0.69 over 8 hostile cells, computed when the data came in and

not independently reproduced; both are confounded with temperature). Across groups it does not predict the dimension: the hostile cells have ratios of about 16 to 19 and next-token entropies of about 7 to 9 nats, yet $D_{\mathrm{KY}} \leq 1.16$. The readout supplies many competing directions; the network decides how many of them expand.

D Behavioral tests in detail

We looked for a behavioral signature of the regimes outside the loop. Two exploratory designs found no accuracy signature that survives controls. First, ordinary sampled answers of Qwen2.5-0.5B under steering at the strengths of chaotic and periodic map cells (30 matched pairs; arithmetic, open questions, short programs) showed no measurable accuracy contrast: at those strengths, steering every position turns every reply into philosophy-flavoured text, so accuracy sits at the floor in both zones (3 against 4 correct of 4500 arithmetic answers). Of nine secondary contrasts of entropy and reply length, one (open questions, next-token entropy, $p = 0.016$) passes 0.05 uncorrected, not after a correction for nine tests, and the entropy contrasts change sign across tasks. Second, labeling each of 200 prompts by the regime of its own unsteered loop in Qwen2.5-7B, prompts with a chaotic loop were answered far less accurately (by 26 to 37 points at three temperatures). The chaotic labels fell mostly on the harder task, and in a linear model with task, greedy correctness and the sampled answers' entropy, the label's coefficient is -3.7 to -4.6 points (permutation p from 0.07 to 0.20). These are associations: the entropy is measured during sampling, so the adjustment does not identify a causal effect of the regime.

A frozen test addresses the second design. Its design, covariates and analysis were pushed before any data. It uses 4000 newly drawn products of two two-digit numbers and the unsteered Qwen2.5-7B-Instruct at $\log_{10} T = 0.40$: 1484 prompts have a chaotic loop and 2509 a periodic one (seven with other labels are excluded). Chaotic-loop prompts are again answered less accurately by ordinary sampling, by 7.1 points (95 % interval $[4.8, 9.3]$; 10 samples per prompt). The gap vanishes once seven difficulty covariates, all fixed before sampling, enter the fit. Beyond them the label's coefficient is $+0.15$ points ($[-1.4, 1.7]$, permutation $p = 0.84$), and the label adds no held-out predictive value. The first-token entropy alone absorbs nearly all of the gap (post hoc: -0.8 points remain). This fits the measurements of Section 3. The loop's chaotic cells have more competing next tokens than its periodic ones, and a model unsure of its first token also answers worse. Beyond that uncertainty, no association of the loop's regime with accuracy is detected at this precision; this does not show that the regime has no effect.

Within a prompt no association is detected either. We steered the first 1000 of these prompts slightly, at block 14 along philosophy, hostility-vs-kindness and a random direction. The strength, 1 % of the residual norm, is the largest at which greedy accuracy stayed within 5 points of the unsteered model on the first 200 of these prompts (so the calibration overlaps the test prompts); on the 1000 prompts it is 0.87 to 0.89, against 0.89 unsteered. Under these pushes the loop's label differs between conditions for 508 of the 1000 prompts. With prompt and condition fixed effects (3992 prompt-condition rows), in an analysis fixed before the data, a chaotic label is associated with a change of $+0.36$ points in the accuracy of the prompt's ordinary sampled answers (95 % interval $[-0.36, 1.08]$; permutation within prompts, $p = 0.34$). This does not identify a causal effect of entering chaos: the pushes can act on the answers through other paths as well.

Nor does accuracy peak at the edge of chaos. Adjusted for the covariates, three groups of prompts (3603 of the 4000; the rest have intermediate exponents) are answered alike:

- deep periodic ($\lambda \leq -0.3$, 1961 prompts): $+0.3$ points;

- edge ($|\lambda| \leq 0.1$, 242 prompts): -0.1 points;
- deep chaotic ($\lambda \geq 0.3$, 1400 prompts): $+0.3$ points.

All three intervals lie within $[-3.1, 3.0]$. The edge contrast is -0.4 points ($[-3.5, 2.6]$), and the curvature in λ is positive, the opposite of an edge optimum. Post hoc, the curvature comes from the most strongly contracting tenth of prompts ($\lambda \leq -3.42$), answered 5.3 points ($[3.0, 7.6]$) better than the covariates predict. That is a lead for a future test, not a finding.

Sensitivity without control. A loop whose answers change with rounding differences raises a further question: can a deliberate perturbation of one fed-back token steer the answer to a chosen wrong value? An exploratory test asked this of the Qwen2.5-0.5B soft-thinking pipeline on two-digit additions, in float32, with 16 soft steps, at temperatures 0.3 and 1.0, on 16 development and 64 held-out problems. The target was the correct sum plus or minus 10. A 20-step projected gradient ascent on the target's log-probability perturbed the first soft token by one token norm, and random search over 45 directions with at least as much computation served as the comparison. On the held-out problems, the early perturbations reached the target in 5 of 101 problem-temperature pairs that the unperturbed pipeline answered correctly, and random search in none: a difference of 4.2 points (95 % interval $[0.8, 8.3]$), below the advantage of 10 points that the design required. The same optimization applied to the last soft token before the answer cue reached the target in 14 of the 101, so in this pipeline the loop gives less leverage early than late. Random early perturbations of that size changed about a third of the answers (one random direction: 34 of the 101 eligible pairs), almost never to the target. Here the loop's sensitivity comes without selective control. The test covers temperatures at which the loop contracts for most problems (on the development problems 10 and 38 % of twins expanded, with median rates of -0.131 and -0.058 per step), where the pipeline still answers correctly, and the doses and the optimizer were chosen on the development problems (`docs/pi/c12-1-summary.json`).

Taming and using the sensitivity. Two exploratory tests with Qwen2.5-7B-Instruct in bfloat16 (512 unused GSM8K training problems, 16 at each of 32 prompt lengths, fixed before any outcome; soft thinking with the top 15 tokens at $T = 0.6$ for 64 steps) asked whether the theory gives a remedy, and whether the sensitivity can be put to use. The factorization of Section 3 isolates the readout gain, which along recorded orbits of the 0.5B model carries most of the growth that steering adds. Lowering a step's temperature until a proxy of the readout gain fell below a cap (the cap could not be met on about 1 % of steps, at exact ties) reduced the share of answers that the batch layout changed from 27.7 to about 23 % (by 4.1 to 4.7 points at three caps), and accuracy did not fall. Lowering the temperature to 0.5 at every step, a control added after the first result, gained as much accuracy, and the caps' extra reduction of flips relative to it (2.7 to 3.3 points) was not significant. So the remedy is modest and not shown to be specific to the readout. Second, a plurality vote over eight runs whose first soft token was displaced by a relative 10^{-3} showed no detectable accuracy gain over one run ($+0.2$ points, $[-1.0, 1.4]$), although one displacement changed 20.2 % of the answers on average; sampled self-consistency with eight samples gained 5.6 points ($[3.5, 8.1]$). The perturbed runs differ, but rarely toward the correct answer. Problems whose perturbed runs disagreed were also the ones answered correctly least often (`docs/pi/e2-summary.json`, `docs/pi/e3-summary.json`).

Weight quantization and the attractors. Deployed models are often quantized, and quantization is usually checked one step at a time, by how close the quantized model's next-token distributions stay to the original's. An exploratory test asked whether that check certifies the

closed loop. We quantized every linear weight of the decoder blocks of Qwen2.5-0.5B-Instruct, either in 8-bit blocks (as the Q8_0 format of llama.cpp) or in symmetric 4-bit blocks, and kept the arithmetic in float64. The unsteered loop then ran on 48 held-out prompts, at 8 temperatures over $\log_{10} T \in [-0.5, 0.5]$, from 2 starts each. A frozen random logit bias, scaled so that its one-step divergence matches the 8-bit variant's, served as an unstructured control.

- *8-bit.* The one-step distributions stay close: on the reference's attractor states the mean Kullback–Leibler divergence is 0.0028 nats, and the most probable token agrees in 94.5 % of states. The share of starts that end on a periodic orbit of period at most 8 shows no detectable change (+1.2 points, 95 % interval $[-1.9, 4.2]$). Yet 14.1 % of the starts end in a different class or period, as do 15.0 % under the random bias of the same one-step size. These changes concentrate at the temperatures where the attractors change.
- *4-bit.* The share rises by 6.8 points ($[2.6, 11.0]$), mostly because the edge of a chaotic band turns periodic. Its one-step distributions, however, are far from the original's: the mean divergence is 0.83 nats, and the most probable token agrees in only 52 % of states.

In this model, then, a close one-step check was consistent with unchanged statistics of the closed loop but did not fix the attractor that an individual run reaches. We found no effect of quantization beyond that of unstructured noise of the same one-step size (`docs/pi/c12-2-summary.json`).

Chaos onset and repetition in sampled text. Does the temperature at which a prompt's loop turns chaotic predict the temperature at which ordinary sampled text stops repeating itself verbatim? An exploratory test with Qwen2.5-0.5B-Instruct on 24 held-out prompts did not detect the proposed positive association. The Spearman correlation is -0.23 (95 % interval $[-0.60, 0.25]$, permutation $p = 0.28$). This is not evidence of no association. One sensitivity analysis (orbits of period up to 32, clustered by system prompt) gives an interval that excludes zero on the negative side, $[-0.62, -0.18]$. At matched temperature, the repetition rates of chaotic and periodic cells differed by 0.003 (interval $[-0.013, 0.023]$), with both regimes present at only 8 of the 24 temperatures. In 13 of the 15 prompts where both onsets were observed, sampled text stopped repeating below the loop's chaos onset. The median gap over these 15 is 0.42 decades of temperature. The soft-token loop and sampled decoding with a growing context are different systems, and this test did not find a bridge between them (`docs/pi/e5-heldout-analysis.json`).

Coupled agents (behavioral check). Two Qwen2.5-1.5B-Instruct agents solving the same multiplication problem each fed back a fraction $1 - c$ of their own concept token and c of the other's, while their contexts grew (64 problems, 8 pairs each; float32; exploratory). At $c = 0.25$ wrong-answer agreement rose by only 1.6 points ($[-0.6, 3.9]$) and accuracy fell by 5.5 points, as it did with replayed messages. At $c = 0.5$ mixing gives the two agents equal next inputs, although their growing contexts differ. They then agreed on 96 % of answers, including 41.0 % of pairs that gave the same wrong number, and accuracy was 12.5 points lower than for independent agents (`docs/pi/c12-3-summary.json`).

Steering history and a policy decision. Coexisting attractors suggest that the history of a steering schedule can matter, not only its present value. An exploratory test asked this of a decision with right and wrong answers. Qwen2.5-3B-Instruct, thinking in soft tokens, judged synthetic access-control requests against a supplied permission table, answering ALLOW or DENY, while a permissive-versus-restrictive direction steered it at layer 18. Every run ended with 32 steps at the same steering strength ($\alpha = 0.025$). Some runs got there directly; others first made a brief excursion

to $\alpha = 0.04$ and back. In the exploratory run, the excursion raised the share of prohibited requests wrongly allowed from 49.5 to 70.8 %, a difference of 21.4 points (95 % interval over 24 held-out instances $[15.6, 26.6]$; `docs/pi/c12-4-summary.json`).

We then repeated the test as a confirmatory one. The design, the analysis and a new set of 48 instances (drawn from a new pool of users) were fixed and pushed before any data (`docs/pi-c12-4-confirmatory.md`), with the prediction that the excursion would raise false allowing by at least 10 points, with an interval excluding zero. The prediction held, with a smaller effect: 13.0 points ($[7.2, 18.8]$ over 24 prohibited instances) after the 32 steps at the final strength, and 15.1 points ($[10.4, 19.8]$) after 64 steps. A second permissive direction, built the same way from differently worded statements (cosine 0.77 with the first), gave 19.3 points ($[12.5, 26.1]$). A restrictive excursion raised false denials of permitted requests by 23.4 points ($[17.1, 29.7]$). An excursion of the same size along a random direction showed no detectable effect (-3.1 points, $[-8.9, 2.1]$), and an excursion of zero size changed none of the 384 decisions (`docs/pi/c12-4c-summary.json`). The present steering strength therefore does not describe the steered model's behavior here. The effect persists after the steering values match. Because what the loop wrote during the excursion stays in its growing context, a memory of the context is the plausible mechanism, although these tests do not isolate it; this is history dependence of the protocol, not a demonstration of hysteresis in a fixed-size state. One model and one schedule were tested, in float32, with two wordings of one direction.

Does the hidden state change answers? In band chaos the loop writes an exactly repeating text while its state moves chaotically: the same token wins at each position, but the probabilities behind it keep changing (Section 10). An exploratory test asked whether states with the same visible history answer a follow-up question differently. We took 16 states at one phase of the text from each of 92 held-out band-chaos cells; within a cell, all 16 had identical leading tokens over the preceding 70 steps. Each state was read through the steered prefix, followed by one of three queries chosen on development cells: a yes-or-no question about hostility, a one-word mood, and a hostility rating from 1 to 10. Each answer was decoded greedily. No pair of states gave different answers in any cell for any query (a one-sided 95 % bound of 3.21 % of cells), although the probability of a fixed answer token differed at six significant digits in 98 to 99 % of the pairs. The hidden differences therefore reach the readout, but here they stay within the margins of the greedy choices, and only the wording of the yes-or-no answer changed, in 9 of the 92 cells. The visible text is not a complete record of the state. For these readouts that is a limit on what can be observed, not a change in behavior (`docs/pi/c12-5-summary.json`).

E Bifurcation curves and cusps

Period-8 stripe. Continued for 80 points with a float64 test function, the fold curve of the period-8 orbit keeps its critical multiplier within $1 \pm 6 \times 10^{-6}$, follows the stripe's upper edge and caps its top near $\log_{10} T = -0.065$ (Figure 10). At $\alpha = 0.0717$ the stripe lies between a flip at $\log_{10} T = -0.128955$ and the fold at -0.1033462. Continued downward, the flip curve keeps its multiplier at -1 for 115 points to $(\alpha, \log_{10} T) = (0.0776, -0.402)$, with two reversals of its parameter tangent near $(0.0750, -0.219)$ and $(0.0746, -0.206)$, and the fold curve keeps its multiplier at $+1$ for 120 points to $(0.0758, -0.200)$. Located by bisection along the curve (Appendix B), the two reversals of the flip curve, at $(0.074955, -0.219336)$ and $(0.074608, -0.206007)$, carry no second multiplier near the unit circle (the next ones are 3×10^{-7} and 1×10^{-5}). On a flip curve the parameter tangent can vanish only where a second multiplier reaches $+1$, so both are tight but smooth turns of the curve in the parameter plane, not codimension-two points.

The accumulation boundary. The fold curve of the 2-cycle "is it", continued from $(\alpha, \log_{10} T) = (0.0628, -0.115)$ in both directions (38 and 126 points, critical multiplier 1 to 10^{-7}, no cusp), traces the boundary of the period-2 region in Figures 10 and 11. The stripe fan accumulates on this boundary, as the mechanism of Section 8 requires. Further along one direction, outside the plotted window, another multiplier leaves the unit circle (moduli up to 11). Located along the curve, this second real multiplier reaches $+1$ at $(\alpha, \log_{10} T) = (0.035628, -0.341724)$, where the eigenvectors of the two multipliers at $+1$ coincide to 3×10^{-4} degrees: a double multiplier $+1$ with a Jordan block, the 1:1 resonance of maps (the analog of a Bogdanov–Takens point).

A period-6 window with shrimp-like anatomy. The hooks of Figure 12 (bottom) belong to the period-6 loop ". -\n\nThe is a". Float64 orbits (Newton residuals 10^{-15}–10^{-13}) show that its curved body and its thin leg are one window: the multipliers are -0.061 and $+0.048$ at the body, -0.564 lower down and -0.645 on the leg. The foot is its period-doubled loop (period 12, leading multiplier -0.156). Inside the body lies a different 6-cycle, ".The\n\nThe is a", whose leading multipliers are a complex pair, $-0.031 \pm 0.152i$. One-parameter branches from the body end in folds in every direction: at $\log_{10} T = -0.0558$ and -0.0721 in temperature, at $\alpha = 0.08074$ in steering, and, toward smaller α, in two folds only 10^{-5} apart. Such a close pair is the signature of a nearby cusp. A flip at $\alpha \approx 0.0797$ then opens the period-12 foot. The body is not effectively one-dimensional either. At two body cells the leading multiplier is a complex pair, $-0.020 \pm 0.264i$, or has a second multiplier of comparable size (-0.261 against 0.099). Below the body, where the window's strip appears to cross a period-5 strip near $(\alpha, \log_{10} T) = (0.0773, -0.0774)$, both orbits continued from their own strips are unstable at the crossing (leading multipliers -1.88 and -6.59). The crossing is therefore neither a center nor a point where two attractors coexist.

Fold curves that bound a window and meet in cusps, with flip curves inside, form the codimension-two organization of a shrimp [67] (complex dynamics of small recurrent networks has been studied in the same spirit [134]). In the continued curves (Figure 67), a single fold arc bounds the body on its outer side and a flip arc on its inner side, and the leg lies between them where they converge. The outer fold arc extends far to the left as a thin arm and ends in a cusp at $(\alpha, \log_{10} T) = (0.069106, -0.039750)$. The fold curve through the close fold pair has a second cusp inside the body, at $(0.080097, -0.058344)$, from which two fold branches descend.

At both reversals the quadratic coefficient b changes sign between consecutive curve points while a stays away from zero (Table 26). That alone is not evidence of a cusp: along any fold curve the parameter tangent reverses exactly where b vanishes. The test's content is the nondegeneracy at that point. Because the signs of a and b both flip with the orientation of the critical eigenvector, the test uses $b \operatorname{sign} a$. A cusp is nondegenerate when the cubic coefficient c of the normal form $u' = u + a\varepsilon + bu^2 + cu^3$ is nonzero where b vanishes. We compute c with the second-order center-manifold term, $c = \ell^\top C(r, r, r)/6 + \ell^\top B(r, h_2)/2$ with $(I - M)h_2 = B(r, r) - 2br$ and $\ell^\top h_2 = 0$, where M is the monodromy matrix and B and C are the second and third derivatives of its map, obtained by central differences of float64 Jacobian-vector products. (The code reproduces c to 10^{-6} on a map with a known reduction.) At the two curve points that bracket each sign change, c is nonzero far beyond its step-size uncertainty and has the same sign, and b changes sign with nonzero slope between them. This supports a nondegenerate cusp inside each bracket. Located along each fold curve (Appendix B), $b \operatorname{sign} a$ falls to below 5×10^{-6} of its bracket values at $(\alpha, \log_{10} T) = (0.069106, -0.039750)$ on the outer arc and at $(0.080097, -0.058344)$ inside the body. At these same points $c = -1.51 \times 10^4$ and -2.38×10^5, and the unfolding matrix K defined below is regular ($\det K = 2416$ and 14.0): two numerically located cusps. (Location uses the small-step oriented coefficient $b \operatorname{sign} a$; the unfolding coefficients use a separate finite-difference ladder, and the

selected plateau of c spreads by at most 4.8 % over adjacent steps.) The negative sign of c means that near each cusp, two stable period-6 orbits coexist with a saddle between the fold branches.

A generic codimension-two cusp also needs the two parameters to unfold it independently. Near the cusp the reduced map is $u' = u + \beta_1 + \beta_2 u + bu^2 + cu^3$, where β_1 and β_2 vanish at the cusp and depend on the parameters. The unfolding is regular when $K = \partial(\beta_1, \beta_2)/\partial(\log_{10} T, 10\alpha)$ has rank two ($\det K \neq 0$). We first estimated K at a fold point near the outer-arc cusp (where $b \operatorname{sign} a = 14.7$): $\det K = 2737$ (singular values 1630 and 1.7), so the estimated unfolding is nonsingular there, as it is at the located cusp ($\det K = 2416$). With these derivatives, the local cubic model predicted the number of period-6 orbits and their stability at six parameter offsets, written to disk before any orbit was solved. Newton solves started from the predicted roots and from the base state recovered the predicted number of distinct orbits, with the predicted stability, at all six. The search was seeded, not exhaustive: orbits that these starts do not reach would be missed.

	cusp on the outer arc	cusp inside the body
$b \operatorname{sign} a$ at consecutive curve points	+92.9, +42.9, +14.7, **−16.9**, −52.8	−24.7, −14.3, −2.0, **+12.8**, +28.0, +45.9, +58.9
$\lvert a \rvert$	1.2–6.9	3.13–3.35
c at the two points around $b = 0$	-1.37×10^4, -1.65×10^4	-2.38×10^5, -2.34×10^5

Table 26: Normal-form test at the two parameter-tangent reversals of the period-6 window's fold curves (bold: the flagged points; critical multiplier 1 to 10^{-7}). c: cubic coefficient at the two curve points that bracket the sign change (finite-difference steps agree to 0.1 and 0.7 %). Bisection along the curves then located both cusps; they are numerical, not certified.

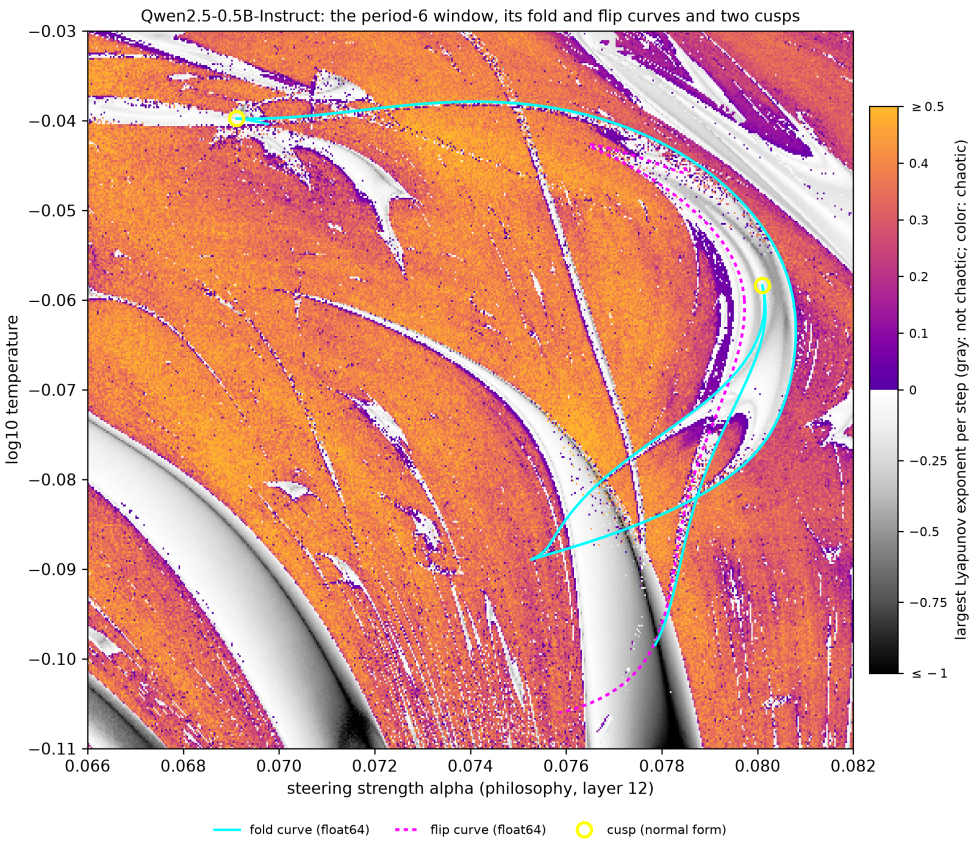

Figure 67: The period-6 window of the Qwen2.5-0.5B-Instruct map (philosophy direction, layer 12) with its float64 fold curves (cyan), flip curve (magenta, dotted) and the two located cusps (yellow circles). The background is a float32 phase map of 480×400 cells (largest exponent from 300 steps after 300 transient steps; chaotic cells colored by the exponent and the other cells in gray by theirs, on one color bar). Where it overlaps the float64 map of Figure 10, 95 % of cells carry the same chaotic or periodic label at the nearest grid point. Some fold branches were continued from two starting points and overlap. Continuation of the lower branches stopped after 60 points near $(0.0753, -0.089)$, where they approach each other; how they join there is not resolved.

A period-4 window with five cusps on one fold curve. At higher temperature, $\log_{10} T \approx 0$–0.1, periodic orbits are no longer effectively one-dimensional (second multipliers 0.06–0.09). At $\alpha = 0.06484$ the temperature branch of the period-4 orbit "\n\n - \n\n The" is S-shaped, with folds at $\log_{10} T = 0.0640$ and 0.0664. When that fold is continued in the plane, its parameter tangent reverses at three points; at the first two, the quadratic coefficient b changes sign while a stays away from zero, and the cubic coefficient c (defined above) is nonzero (Table 27). Located along the curve, these two cusps lie at $(\alpha, \log_{10} T) = (0.062959, 0.059958)$ and $(0.065574, 0.074447)$. There $b \operatorname{sign} a$ is below 1.1×10^{-5} of its bracket values, $c = -1.02 \times 10^4$ and $+1.37 \times 10^5$, and $\det K = -42.0$ and 35.5. Further traces of the same fold curve, started from other points of the period-4 branches, retrace it to within 3×10^{-7}, and one of them finds two more reversals near $\log_{10} T \approx 0.012$. All five reversals are located cusps. The other three lie at $(\alpha, \log_{10} T) = (0.083082, 0.108633)$, $(0.074072, 0.011831)$ and $(0.073220, 0.012759)$. There c is -9.7×10^4, -5.7×10^3 and $+1.7 \times 10^3$, and $\det K$ is 25.3, -188.5 and 68.5. The last two cusps, close together, have opposite signs of c. Three period-4 orbits coexist between the two fold branches, and near the first cusp ($c < 0$) two of them are stable (Figure 68).

	cusp 1	cusp 2
b at five curve points	$+28.1,\ +18.7,\ +8.4,\ \mathbf{-1.1},\ -12.8$	$+51.8,\ +29.7,\ +9.3,\ \mathbf{-9.0},\ -23.4$
a	3.02–3.08	-2.58 to -2.57
critical multiplier	1 ± 10^{-9}	1 ± 10^{-8}
c around $b = 0$	$-1.06,\ -1.01$ $(\times 10^4)$	$+1.39,\ +1.36$ $(\times 10^5)$

Table 27: Normal-form test of the two parameter-tangent reversals (bold: the flagged points). c: cubic coefficient at the points around the sign change, nonzero at both reversals (steps agree to 10^{-4}), which supports a nondegenerate cusp inside each bracket; bisection then located both; at the second, $c > 0$: the orbit between the two others is the stable one.

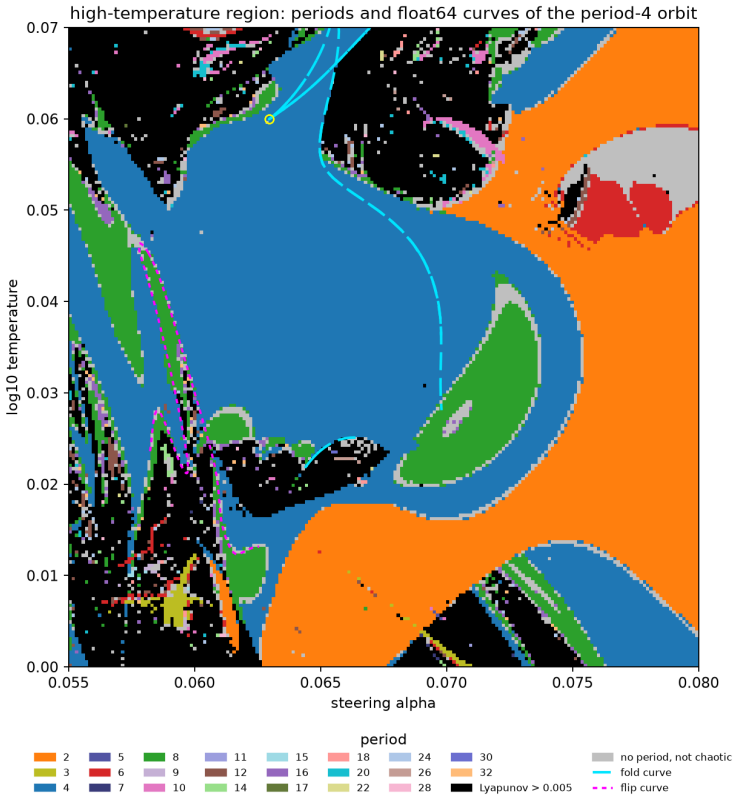

Figure 68: Period map of the high-temperature region of the Qwen2.5-0.5B-Instruct map (philosophy direction, layer 12; 192^2 float64 cells) with float64 fold curves (cyan, long dashes) and flip curves (magenta, short dashes) of the period-4 orbit. The yellow circle is the first cusp; the second lies just above the window. The flip curves run along the boundary between period 4 and period 8. The fold curves run inside the period-4 region, where three period-4 orbits coexist (near the first cusp, two of them stable). In this map a C-shaped band of the doubled loop "\n\nThe\n\nThe" (period 4) inside the period-2 region "\n\nThe" contains period 8 and an irregular core. Along the cut of Section 9 the period-8 orbit merges back into period 4 (a doubling bubble), so this band is not shown to be a second cascade. Continued from all its fold points, the region's boundary consists of an upper fold curve, with the five cusps of the text, and a separate lower fold curve, nowhere within 2×10^{-3} of it in the traced ranges; the lower curve carries the codimension-two points described below. Within the traced ranges the curves do not join, so the region is not shown to be one shrimp-organized component.

Along the lower fold curve a second real multiplier reaches $+1$ twice, at $\alpha = 0.062483$, $\log_{10} T =$

0.004955 and at $\alpha = 0.062499$, $\log_{10} T = -0.001780$. Both points are 1:1 resonances, with the eigenvectors of the double multiplier parallel to 3×10^{-4} degrees. Close by, at $\alpha = 0.062522$, $\log_{10} T = -0.002667$, a multiplier reaches -1: a fold-flip point. We did not continue the curves that emanate from these points.

Evidence level of the cusps. We estimated the rank condition of the unfolding at a fold point near one cusp of the period-6 window and one of the hostile period-7 body (Section 10; determinants 2737 and 53.5), and there the local model's predictions at twelve frozen offsets matched the orbits recovered by Newton solves seeded at the predicted roots and the base state, in number at all twelve and in stability at eleven; the search was not exhaustive. At the thirteen located cusps themselves (including the four of the period-5 window of Appendix F), $|\det K|$ lies between 14 and 2416. These are float64 computations validated by step-size checks (rung 1), not computer-assisted proofs. At the two continuation points that bracket the lower cusp of the hostile body, an independent implementation of the model reproduced the normal-form coefficients, including the sign change of b: $|b|$ within 1.2 % and c to two figures, with the opposite orientation of the null vector, which flips the signs of a and b together.

F Steering toward hostility: further results

This appendix gives further results on the hostility direction of Section 10: the details of the shrimp-like shapes, the period-5 window of a plane steered at two depths, its orbit census and a closure test.

Details of the shrimp-like shapes. Section 11 summarizes the anatomy (Figure 47); three details belong here. In higher dimensions a window can also end where a complex pair crosses the unit circle, and the leading multiplier can change sign through a complex pair, as it does in parts of our windows; beyond some of the flip curves we followed lie the doubling sequences of Section 9, but a flip need not start a complete cascade. The period-6 hook's orbit never comes close to superstable: its leading multiplier stays above 0.028 in modulus. And the X-shaped crossing of a period-5 and a period-6 strip (Figure 12, bottom) is not one window's anatomy but two windows overlapping, and at the crossing both of their orbits are unstable (Appendix E). The shapes in these one-block planes are therefore the two-parameter unfolding of windows whose return maps have one coarse extremum: shrimp-like in body, legs and internal cascade, but without a one-dimensional shrimp center on the tracked branches.

Two turning points from two steering depths. Steering at two depths gives the band dynamics a second turning point. We drew 1584 chaotic cells (200 from each of the four two-block planes of Figure 20, 150 from each one-block map of Figure 19 and the 34 of the hostile line of Figure 30), counted the turning points of each band's return map and inspected every strong candidate by eye, since a projection can fold a single hump. Three cells, all in two-block planes at $\log_{10} T = -0.1$, have smooth, single-valued return maps with two or more turning points (Figure 73); the other strong candidates were point sets, split bands or folded projections, or were not reproduced. In one cell (hostility at block 4, gratitude at block 16), every one-dimensional phase of a 20-band attractor is N-shaped: a bimodal map. A 128×128 zoom around another, whose map has three maxima, shows a period-5 window with a body, legs, a period-10 and period-20 border, and patches where the fixed start settles on period 4 (Figure 74a). We continued the window's period-5 orbit over a 31×31 grid of its body: 608 cells solved, all with the same token loop, 447 of them attracting. In 61 of these

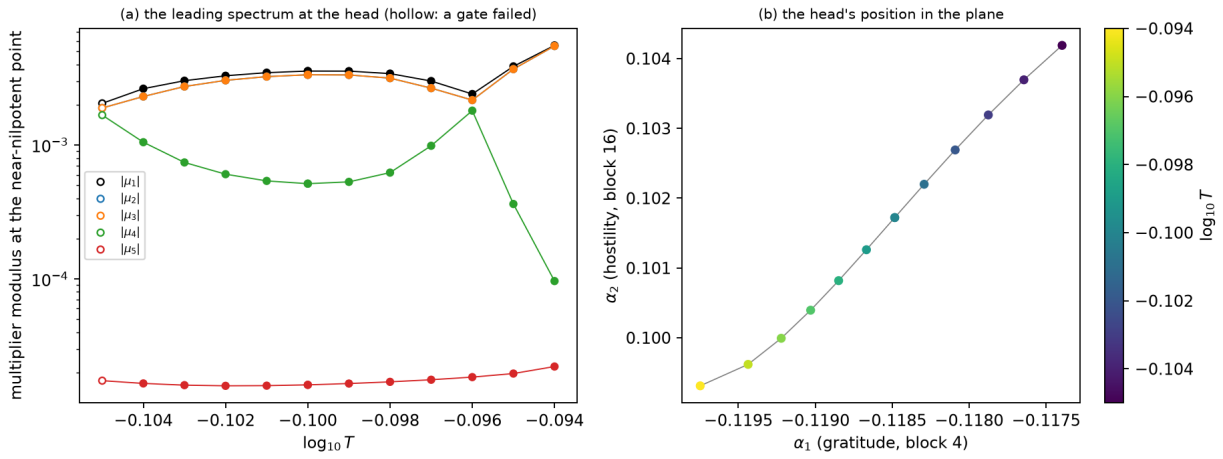

Figure 69: The head of the period-5 shrimp continued in temperature: the near-nilpotent point, where the sum of the leading multipliers and the sum of their pairwise products vanish (tools/nilpotent_track.py; float64, steps of 0.001 in $\log_{10} T$). (a) Moduli of the five leading multipliers at the point (μ_2 and μ_3 form a complex pair of equal modulus; hollow: the row failed its gate, $|c_1| > 10^{-7}$). The defining conditions would put the three leading multipliers on one modulus, the cube root of their product, if they were the only multipliers; the computed moduli are nearly equal, and a fourth multiplier approaches them twice. (b) The point's position in the plane, colored by $\log_{10} T$.

447 cells, the nearest cell of the zoom (at most 2.0×10^{-5} away in α_1 and 3.2×10^{-5} in α_2) settles on period 4, so attractors of periods 4 and 5 coexist at grid scale, which a single unimodal map with negative Schwarzian derivative allows only if a basin reaches an end of its interval (Section 3).

The window still has no one-dimensional center. The leading multiplier of the continued orbit never comes near zero: its modulus is at least 0.0139 in all 608 cells. Where a one-dimensional map would be superstable, it meets the second multiplier, and the two form a complex pair (91 cells, in narrow strips) or, where their product is negative (41 other cells), pass each other as real numbers of opposite sign (Figure 74c). Per step, the second direction contracts by a median factor of 0.44 and the first by 0.76 over the attracting cells, so near these strips the orbit's linearization has two comparable leading directions. The curves on which the trace $\mu_1 + \mu_2$ changes sign take the place of superstable curves, and they do not cross: no 2×2 block of attracting cells has the alternating signs of a crossing (Figure 74b). One isolated cell holds a repelling orbit with the same token loop ($\mu_1 = 1.65$), and near the top of the grid one column of Newton solves failed above cells whose leading multiplier rises to 0.88; both are explained by the window's fold curve (next paragraph). A second turning point in the band maps therefore does not give this window a one-dimensional shrimp center. Like the hostile body, it is a window whose orbit has two comparable leading directions near the strips where a one-dimensional map would be superstable; Section 11 shows that it is a numerically verified shrimp. Its head, the near-nilpotent point, persists wherever we continued it in temperature (Figure 69). At $\log_{10} T = -0.105$, grids of the continued orbit over the tip of the band where the leading multiplier is positive show the same avoided crossing: the band narrows between two strips of complex leading pairs, of modulus 0.097 to 0.109, which come within one cell of each other (5×10^{-5} in α_1) without crossing (`docs/research-log.md`, runs P5Z).

The folds and flips of the period-5 window. To resolve the folds the grid did not, we continued the window's period-5 orbit along four lines from one attracting cell, $(\alpha_1, \alpha_2) = (-0.118730, 0.100693)$, and then the boundaries the lines found, each in both directions (float64; 18 to 73 points per line,

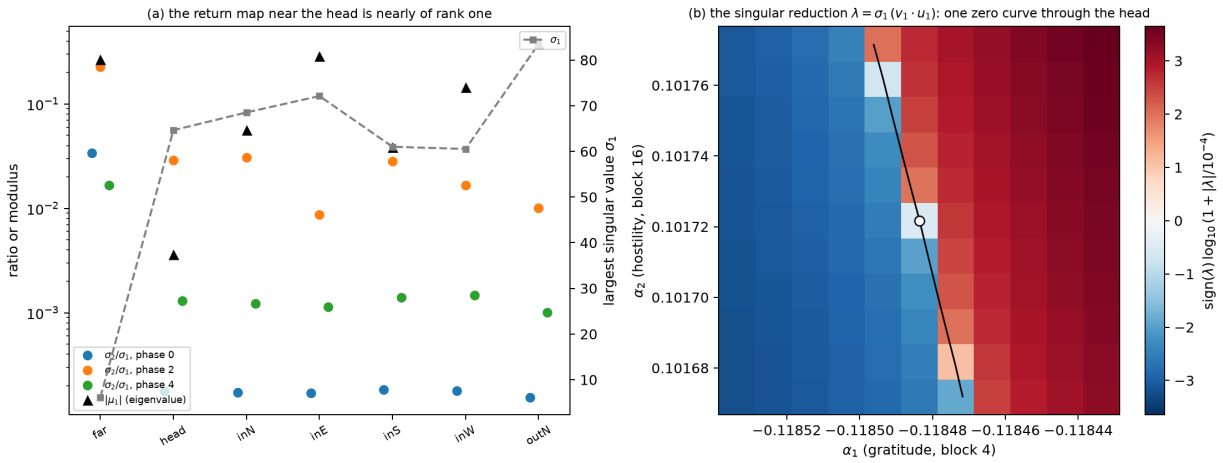

Figure 70: The rank of the period-5 shrimp's return map near its head (float64; `tools/monodromy_svd.py`, `tools/plot_head_rank.py`). (a) Ratio of the second to the largest singular value of the return map, for the cycle started at three phases, in the body far from the head, at the head (the near-nilpotent point), at four points 5×10^{-5} from it and at one 2×10^{-4} from it, with the largest singular value (squares, right axis) and the modulus of the leading eigenvalue (triangles). (b) The reduction to the rank-one part, $\lambda = \sigma_1 (v_1 \cdot u_1)$, on an 11×11 grid of spacing 10^{-5} around the head (white circle), shown as $\mathrm{sign}(\lambda)\log_{10}(1 + |\lambda|/10^{-4})$; its zero set (black) is one curve.

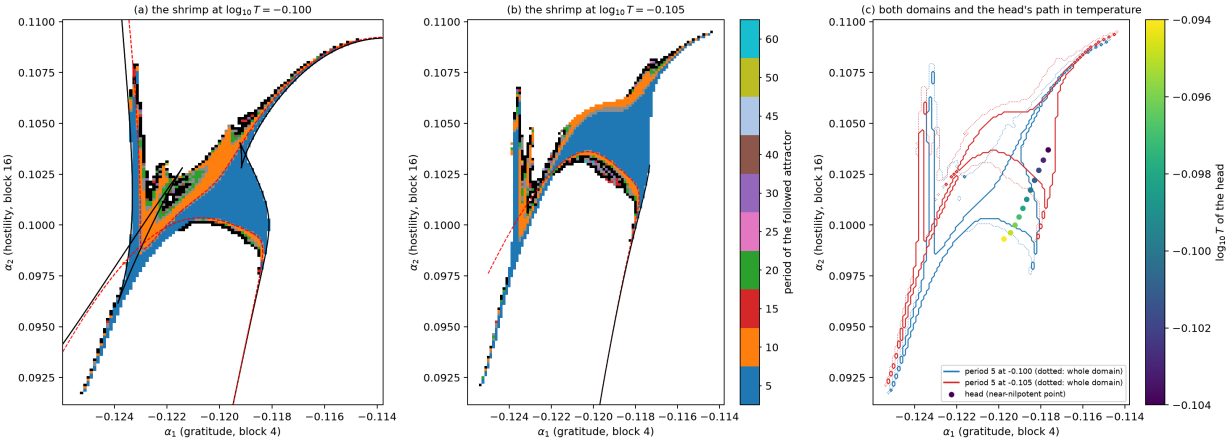

Figure 71: The period-5 shrimp at two temperatures (Qwen2.5-0.5B-Instruct, gratitude at block 4 and hostility at block 16, float64; `tools/plot_shrimp_temperature.py`). (a) The domain of its attractor at $\log_{10} T = -0.100$, followed cell by cell and colored by period (black: chaos behind the same five-token text; dark gray: no period found and an exponent below 0.005; light gray: the attractor was lost), with the fold (black) and flip (red, dashed) curves. (b) The same at $\log_{10} T = -0.105$, with the curves continued at that temperature (a partial trace of a further fold curve, with two flagged reversals, is not drawn). (c) The outlines of the period-5 body (solid) and of the whole followed domain (dotted) at both temperatures, and the path of the head (the near-nilpotent point) between $\log_{10} T = -0.104$ and -0.094, colored by temperature.

30 to 120 per curve; one of the eight curve runs failed at its start for lack of GPU memory, and its curve is covered by the other trace of the same fold). At every curve point the residual is below 10^{-10} and the test function below 10^{-9}. The critical multiplier stays real and within 6×10^{-8} of $+1$ on the fold curve and of -1 on the flip curves; the largest other multiplier modulus is 0.040 on that fold curve and 0.122 on the flip curves. A rerun of these curves on 18 October with wider bounds, toward the legs' tips, keeps every critical multiplier within 1.9×10^{-7} of ± 1, and the largest other modulus reaches 0.86 where the lower-right leg's fold and flip curves end at the bound of the continuation. Going right and up from the start, the orbit dies in folds at $\alpha_1 = -0.118166$ and $\alpha_2 = 0.102715$ (next moduli 0.0012 and 0.0008). Going left and down, it first loses stability in flips (at $\alpha_1 = -0.12147$ and $\alpha_2 = 0.09954$), regains it where the same multiplier crosses -1 again (at -0.12305 and 0.09586; a doubling bubble), and dies in a fold (at -0.123323 and 0.095720). Three of these four folds lie on one continued fold curve (within 3.1×10^{-7}); the left one lies on a second. The first fold curve bounds the lower-right leg on its right, runs up the body's right edge and over its top, and leaves along the upper leg. Near $\alpha_1 = -0.1192$ its parameter tangent reverses twice, identically on two traces continued from different boundary points, and both reversals are located cusps (the test and criteria of Appendix E, applied on each trace): at $(\alpha_1, \alpha_2) = (-0.1192221, 0.1040536)$, with $c = +2.60 \times 10^4$ and $\det K = 176.6$, and at $(-0.1191670, 0.1028019)$, with $c = -2.05 \times 10^4$ and $\det K = 88.6$. The two traces agree to 2×10^{-12} in these positions and to five figures in c and $\det K$. Its outer arms cross at $(-0.11901, 0.10346)$ (Figure 45d). There they are folds of two different period-5 orbits with the same token loop: at the four continuation points that bracket the crossing, the orbits on the two arms differ by 3.65 to 3.78 % of the state norm (root-mean-square over the five points, at the best phase alignment; both traces), 15.8 to 44.5 times the state change between neighboring curve points (`docs/fm5/crossing-orbits.json`). The second fold curve, continued both ways from the left fold, runs from below the lower-left leg far up the left of the plane, where the orbit has regained stability beyond its doubling bubble (critical multiplier within 1.2×10^{-7} of $+1$ and other multipliers below 0.4; in the rerun, 1.9×10^{-7} and 0.61). Within the bounds of 15 October it has two more reversals, both located cusps: $(-0.1238662, 0.0961512)$, with $c = -9.65 \times 10^3$ and $\det K = 405.9$, and $(-0.1214073, 0.1028234)$, with $c = +4.75 \times 10^4$ and $\det K = -535.4$ (condition 1.2×10^4). The rerun finds a third reversal, at $(-0.13295, 0.08530)$, which the continuation flags but we have not located. This curve also crosses itself, at $(-0.12341, 0.09904)$, between its points 6 and 64, which lie on either side of the two cusps (points 15 and 40), so that the two cusps share a segment and their other branches cross: a second dovetail. There too the arms are folds of two different period-5 orbits with the same token loop, which differ by 8.83 to 9.06 % of the state norm (measured the same way), 23 to 31 times the state change between neighboring curve points. In both dovetails the cusp with $c > 0$ lies above the one with $c < 0$. At all four cusps the critical multiplier is within 1.3×10^{-9} of 1 and the second modulus below 0.008. The flip curves bound the body's upper-left and lower sides and run along the legs. The fold curve resolves both anomalies of the grid. The isolated repelling cell ($\mu_1 = +1.65$) lies 1.2×10^{-5} from it, on the saddle branch, and the column of failed Newton solves at $\alpha_1 = -0.119167$ runs into the lower reversal, at $\alpha_1 = -0.119166$, where the followed orbit ceases to exist between the folded branches.

The lower-left leg's own curves. The followed attractor enters the lower-left leg, but no line from the body's start crosses it, and the curves continued from the body pass to its left. We therefore marched the period-5 orbit by Newton's method from the body's start, $(\alpha_1, \alpha_2) = (-0.11875, 0.10075)$, through the 68 cells by which the follow run entered the leg (66 of period 5, and 2 of period 10 where the orbit exists but is unstable) to $(-0.124375, 0.094)$. Newton's method converged at every cell, with largest relative residual 8.9×10^{-13}. At the end the leading multiplier

is $+0.64$. Along α_1 the orbit loses stability in a flip at $\alpha_1 = -0.12451$ (real crossing multiplier; next modulus 0.0060), beyond which its multiplier falls to -13.9, and in the other direction it dies in a fold at -0.1243654 (next modulus 0.014). Both boundaries were continued both ways in float64. The two continuations toward the body were stopped by their time limit; their points were recovered from the runs' logs after the same checks. At every point of the four curves the residual is at most 10^{-10} and the test function at most 5.2×10^{-10}; the critical multiplier is real and within 3.7×10^{-8} of -1 on the flip curve and of $+1$ on the fold curve, and every other multiplier has modulus at most 0.21. Downward, the two curves bound the leg to $\alpha_2 = 0.080$, the bound of the continuation. The strip between them narrows from 1.9 grid steps at $\alpha_2 = 0.096$ to 0.58 at 0.0919, the leg's lowest period-5 cell in the follow run, and to 0.015 at 0.080 (Figure 72c). Upward, the fold curve reverses direction at $(-0.11938, 0.10042)$, where the continuation flags a cusp that we have not located, and the flip curve bends sharply at $(-0.12214, 0.10135)$. Both then run back down the left of the plane alongside curves continued from the body, of the other kind, over $\alpha_2 \in [0.087, 0.0995]$: the returning flip arm at a median distance of 0.017 grid steps (at most 0.055) from the body's second fold curve, and the returning fold arm at a median of 0.055 (at most 0.46) from a flip curve of the body; they cross the body's curves only at isolated points. In every grid row from $\alpha_2 = 0.086$ to 0.096, every period-5 cell of the leg lies inside the strip (39 of 39), although the strip is only 0.58 to 1.9 cells wide in the rows that have such cells. The strip also holds two slowly converging cells and 16 cells that the follow run never computed, all below the leg's lowest period-5 cell (Figure 72b). Across the leg at $\alpha_2 = 0.094$ the multiplier passes from $+1$ at the fold to -1 at the flip within 1.5×10^{-4} in α_1 (Figure 72d).

Orbit census of the junction. Before any solve we fixed 13 points around the junction (`docs/oc/oc-points.json`) and the number of period-5 orbits that a plane section of a swallowtail predicts at each: four in the pocket bounded by the two cusps and the crossing, none beyond the crossing, and two elsewhere. At each point, Newton's method (low-rank, at most 20 steps) starts from the states of the nearest point on each arm of the fold curve, shifted along that point's critical eigenvector by $s|x|$ for $s \in \{0, \pm 10^{-3}, \pm 3 \times 10^{-3}, \pm 10^{-2}, \pm 3 \times 10^{-2}\}$ (27 starts per point). Orbits are distinct when they differ by more than 10^{-7} of the state norm at every phase. The counts are lower bounds, because the seeded search can miss orbits. Table 28 gives the results: 9 of the 13 counts match. The three undercounts found all the attracting orbits that the model predicts, so by the model each missed a repelling one. Above the upper cusp, two of the three orbits found share one token loop: a saddle (multiplier $+1.44$) and the node, which lies beyond the flip curve that passes the cusp (multiplier -1.35), so that no period-5 orbit attracts there (the classifier map shows a period-4 attractor next to this point). The third has another token loop and the multiplier -130, and only one start reached it; that it does not belong to the junction's family is a reading after the fact.

A closure test that failed. We also tested whether two coordinates forecast the window's nonlinear return map (run CL2). Perturbations of relative size 10^{-3} and 10^{-2} (64 per size, half to fit and half to test) were followed for eight returns; quadratic return models in one, two or four modal coordinates (the spectral projections on the leading eigenvectors), with a fitted lift back to the state, were iterated from the first return, and the error is the squared forecast error over the squared deviation, summed over returns 2 to 8. The test failed for numerical reasons: at most locations the iterated models diverged, with median errors up to 2×10^{122}, against at most 0.71 for the linearized map in the same rows. We did not establish why (in 9 of these 12 rows even the one-coordinate model diverges), and no regularized refit was run. At one location the divergence of

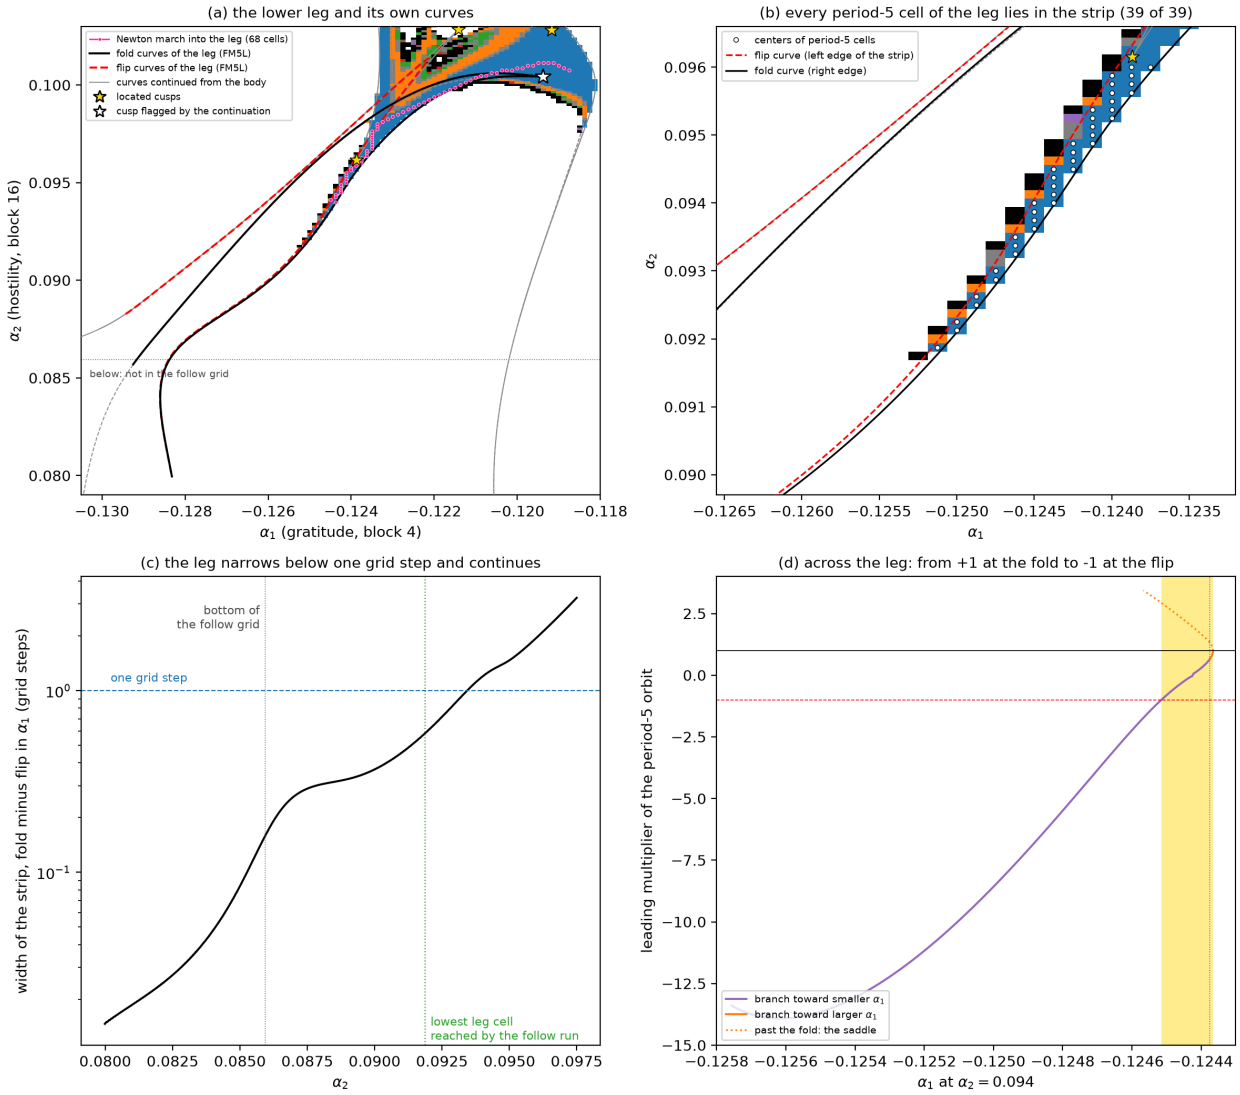

Figure 72: The lower-left leg of the period-5 shrimp and its own curves (Qwen2.5-0.5B-Instruct, gratitude at block 4 and hostility at block 16, $\log_{10} T = -0.1$, float64; `tools/plot_lower_leg.py`). (a) The followed domain (colors as in the followed domain of Figure 46b) with the fold (black) and flip (red, dashed) curves of the orbit marched into the leg (pink: the 68 cells of the march), the curves continued from the body (gray), the located cusps (gold stars) and the cusp flagged on the leg's fold curve (white star); the follow grid ends at $\alpha_2 = 0.086$ (dotted). (b) The centers of all period-5 cells of the leg (white dots) lie between the flip curve and the fold curve. (c) The width of the strip between the two curves, in grid steps: it falls below one step at $\alpha_2 = 0.0934$; the leg's lowest period-5 cell in the follow run is at 0.0919. (d) The leading multiplier of the period-5 orbit along α_1 at $\alpha_2 = 0.094$, from $+1$ at the fold through -1 at the flip (shaded: the strip), and past the fold on the saddle branch (dotted).

points	region	predicted	found	attracting
T-low, T-high, T-near-U	pocket	4	4	2
T-centroid, T-near-L	pocket	4	3	2
E-low, E-high	beyond the crossing	0	0	0
W-low, S-low, N-high, below-L	around the pocket	2	2	1
W-high	around the pocket	2	1	1
above-U	beyond the upper cusp	2	3	0

Table 28: Orbit census of the junction of the period-5 window (run OC; Figure 45d): period-5 orbits found from 27 starts per point, and how many attract. The predictions were fixed before any solve.

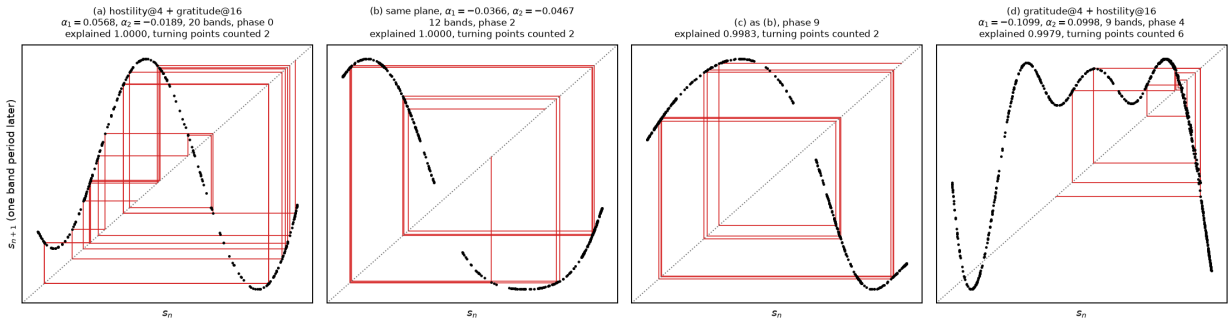

Figure 73: Band return maps with two or more turning points, from two-block planes of Qwen2.5-0.5B-Instruct (float64, $\log_{10} T = -0.1$). Each panel shows one phase of a chaotic band attractor, projected on its leading principal component s (explained variance in the title), with points (s_n, s_{n+1}) one band period apart and, in red, a cobweb that follows consecutive visits (25 in (a), 12 in the others). (a) Hostility at block 4 and gratitude at block 16, $(\alpha_1, \alpha_2) = (0.0568, -0.0189)$: phase 0 of a 20-band attractor, N-shaped. (b, c) The same plane at $(-0.0366, -0.0467)$: two phases of a 12-band attractor, N-shaped maps on two intervals. (d) Gratitude at block 4 and hostility at block 16, $(-0.1099, 0.0998)$: phase 4 of a 9-band attractor, with three maxima.

the four-coordinate model satisfied the frozen two-coordinate criterion spuriously. A guard added after the run marks fits with a median error above 1 as degenerate; 12 of the 18 location-and-size rows are. At the body location, one coordinate gives errors of 0.026 and 0.031 and two coordinates below 6×10^{-5}. We draw no conclusion about two-dimensional closure from this test.

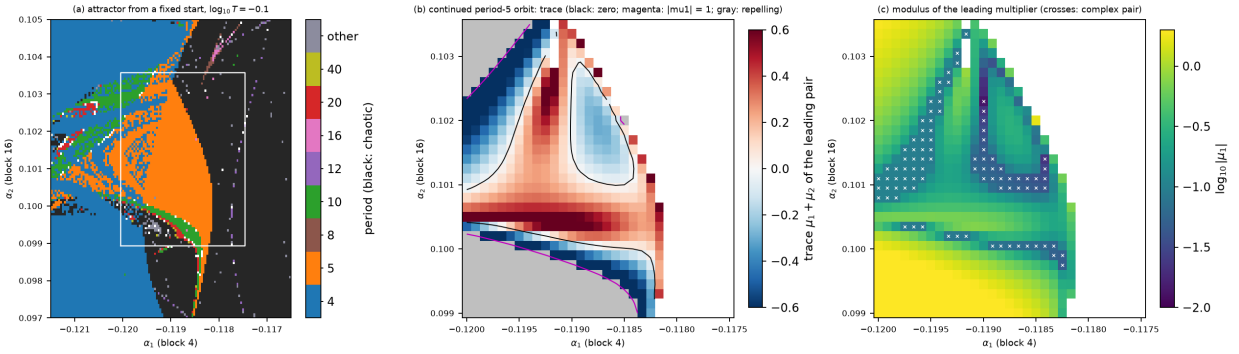

Figure 74: A period-5 window in the plane of gratitude at block 4 and hostility at block 16 (Qwen2.5-0.5B-Instruct, float64, $\log_{10} T = -0.1$). (a) Attractor reached from the fixed start on a 128×128 grid (black: chaotic), with the box of the continuation grid. (b) The window's period-5 orbit continued over the box (31×31, 608 cells solved, one token loop): trace $\mu_1 + \mu_2$ of its two leading multipliers on attracting cells, its zero curves (black) and the curves $|\mu_1| = 1$ (magenta); gray cells hold a repelling orbit. (c) Modulus of the leading multiplier; crosses mark a complex leading pair. The modulus never falls below 0.0139; the dark strips, where a one-dimensional map would be superstable, are where the two leading multipliers meet.

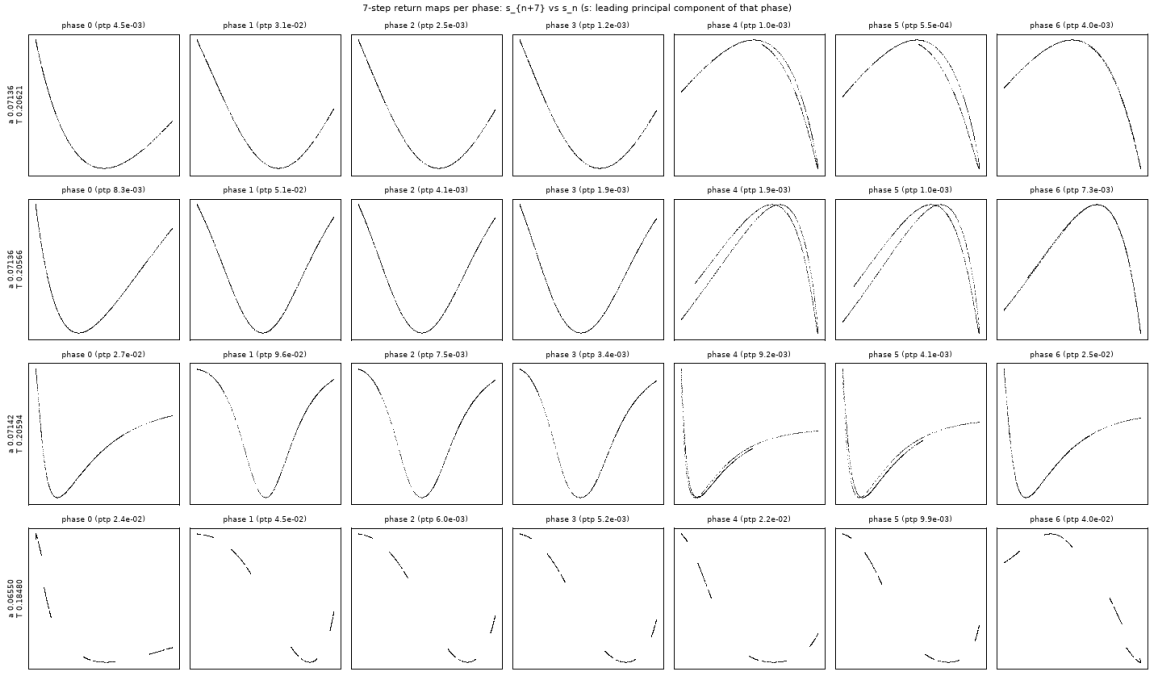

Figure 75: Seven-step return maps s_{n+7} against s_n, for each phase of the 7-periodic text, of chaotic cells next to the hostile period-7 body (top three rows: $\alpha = 0.07136$–0.07142, $\log_{10} T = 0.2057$–0.2062) and on the hostile line of Figure 30 (bottom row: $\alpha = 0.0655$, $\log_{10} T = 0.1848$); s is the projection on the leading principal component of that phase's states, which carries 89–100 % of their variance. Binned coarsely (16 intervals), every phase has one interior extremum; finer partitions keep extra extrema, especially in phases 4 and 5, where the points also split into two thin strands.

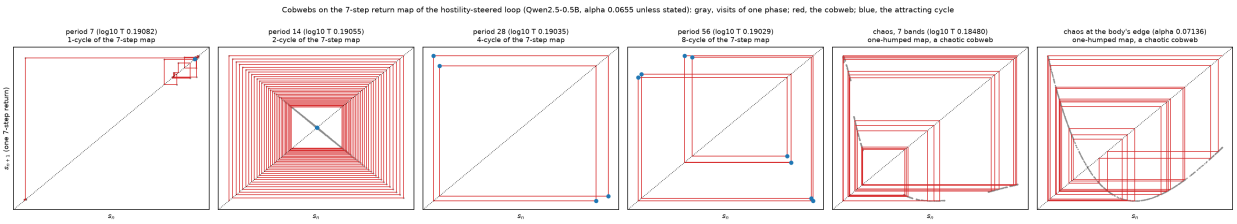

Figure 76: Cobweb diagrams of the 7-step return map of the hostility-steered loop (Qwen2.5-0.5B-Instruct, layer 12, float64). s is the projection of one phase's states on their leading principal component. Left four panels: the doubling cascade at the right edge of the hidden-chaos window along $\alpha = 0.0655$. As $\log_{10} T$ falls from 0.19082 to 0.19029, the attracting set of the 7-step map (blue) is a fixed point (period 7 of the full loop), then a 2-, 4- and 8-cycle (periods 14, 28 and 56); the red cobweb shows the approach from the fixed start (for period 7, its first visits). Right two panels: chaotic cells of Figure 75, where the visits trace a one-humped map (gray) and the cobweb wanders over it.

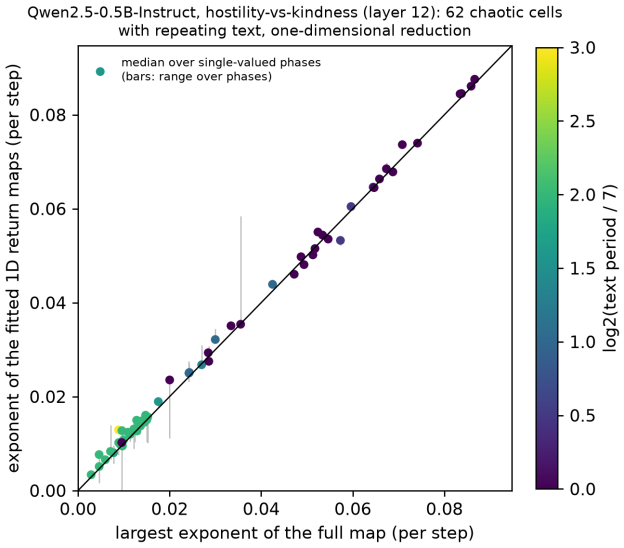

Figure 77: Scalar return-map approximation across the hostile cascade (Qwen2.5-0.5B-Instruct, hostility at layer 12, $\alpha = 0.0655$, $\log_{10} T \in [0.176, 0.191]$): for each of the 62 chaotic cells whose leading-token text repeats exactly, the Lyapunov exponent of the fitted return maps of the leading-token probability (per step; median over the single-valued phases of the text, bars: their range; polynomial degree chosen by cross-validation) against the largest exponent of the full map from twin trajectories. Color: text period relative to 7. The line is equality. With seven-token texts the 2000 steps give 285 repetitions; with 28-token texts only 71 or 72 observations per phase (70 or 71 return pairs), and the fitted maps then overestimate small exponents.

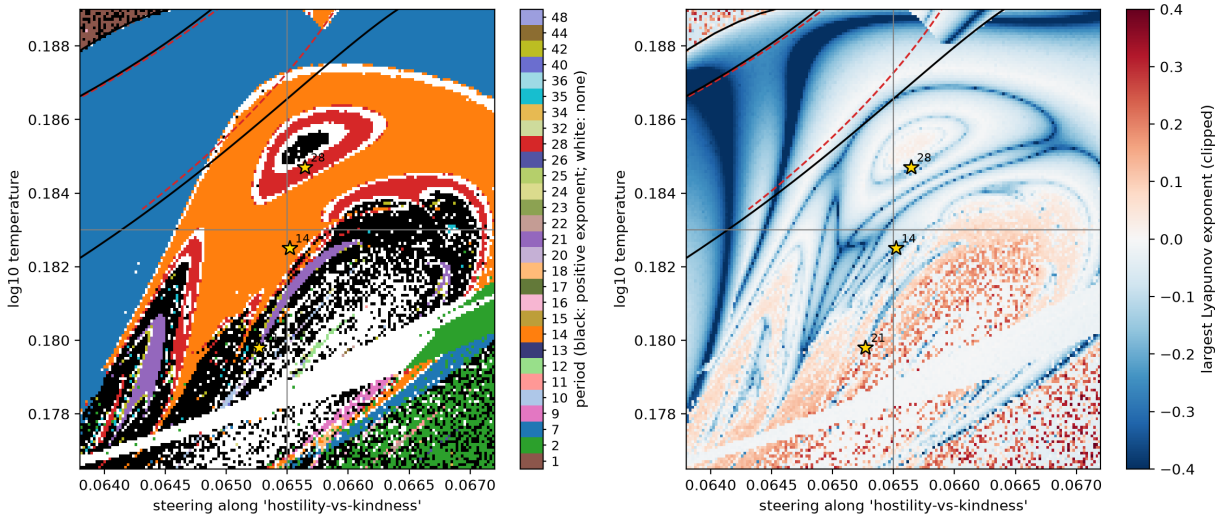

Figure 78: The lower part of the period-7 body of the hostility-steered Qwen2.5-0.5B-Instruct map (layer 12) at higher resolution (160^2 float64 cells). Left: period of the attractor (black: positive largest exponent; white: no period up to 48). Right: largest Lyapunov exponent (clipped); the darkest curves mark a small leading multiplier (candidate superstable curves). Both panels: float64 fold (black) and flip (red, dashed) curves of the period-7 orbit, the two lines of Figure 30 (gray) and the Newton-verified orbits of periods 14, 28 and 21 (stars).

G Bridge to sampled decoding

How much of the fold mechanism of Section 8 survives sampling? To find out, we replace the expected embedding by the mean of m sampled token embeddings (Figure 79). With the exact map, laminar phases near the fold of the 2-cycle "is it" last 120 and 63 steps at $\varepsilon = 10^{-5}$ and 3×10^{-5}. With $m = 64$ they last on average 10.3 to 10.7 steps over the three offsets, nearly independently of ε (57% of steps in the channel), and with $m = 8$ 7.6 to 7.8 steps (38%). With single samples ($m = 1$, sampled decoding of the same truncated system), laminar phases become rare, one to three per offset (11% of steps in the channel). The deterministic intermittency mechanism therefore belongs to the expected-embedding map. Sampling noise turns the bottleneck into a noise-limited trap whose strength grows with m.

Zero-temperature limit and ordinary greedy decoding. As $T \to 0$ the soft-token map becomes, where the leading token is unique, greedy decoding with a one-token memory, $g(t) = \arg\max_s z_s(\text{prefix}, E_t)$ (at an exact tie the limit averages the tied embeddings). This is a map on the vocabulary, and its cycles are the zero-temperature attractors. For Qwen2.5-0.5B we computed g for all 151 936 tokens. Unsteered, it has 1060 cycles, and one of them, the period-5 loop "![](https:// Hello", drains 86 % of the vocabulary; it is the largest attractor family of every 0.5B soft-token map. Ordinary greedy decoding with a growing context does not inherit these cycles (Table 29). From 220 starts (the first tokens of the 20 largest cycles and 200 random tokens), 19 runs end in a short loop within 256 tokens, and the dominant 5-cycle never appears. Steering along philosophy makes ordinary decoding loop almost always (218 of 220 runs at $\alpha = 0.1$), but in other loops, mostly punctuation such as " –", none of them one-token cycles. Greedy decoding that keeps only the last K tokens interpolates between the two. For the tested $K \geq 2$ most final loops differ from the one-token cycles: at $K = 2$ only 7 of 220 unsteered runs and 1 of 220 steered runs keep one of them. Unsteered, looping collapses to the level of the growing context between $K = 16$ and 32; steered, the share of runs whose loop equals the growing-context loop rises from 0 at $K = 1$ to 58 % at $K = 64$. The soft-token map's low-temperature attractors are therefore those of decoding with a one-token memory. Which loop ordinary decoding enters depends on the context it keeps, but the tendency of steering to drive decoding into loops is shared. In ordinary decoding, repetition is known to reinforce itself [27, 28], a theoretical analysis ties loops to words of high inflow [135], and induction heads, which copy earlier context [136], are one mechanism by which a longer context can favor loops that repeat earlier text.

	unsteered		$\alpha = 0.1$ (philosophy)	
memory K	looped	loop as with full context	looped	loop as with full context
1	220	5	220	0
2	220	5	220	0
8	220	5	220	3
16	169	5	219	24
32	32	7	219	71
64	26	11	219	126
full context	19	–	218	–

Table 29: Greedy decoding of Qwen2.5-0.5B with a memory of the last K generated tokens after the fixed prompt ($K = 1$ is the soft-token map's zero-temperature limit), 220 starts of 256 tokens each: runs ending in a loop of period up to 64 repeated three times, and runs whose final loop equals that of ordinary decoding with the full context from the same start.

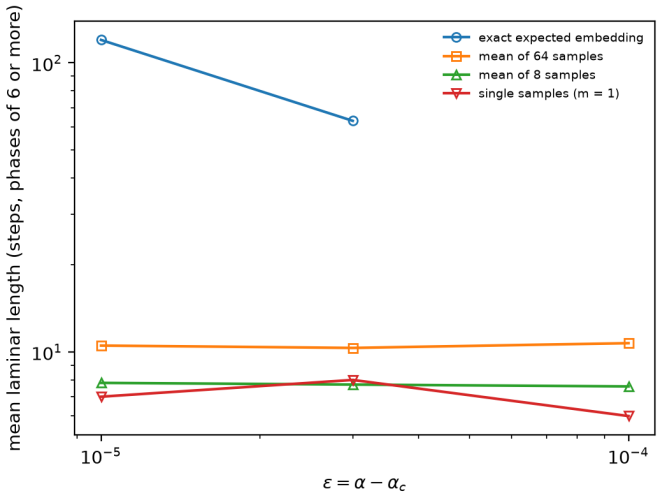

Figure 79: Mean laminar length near the fold of the 2-cycle "is it" (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, $\log_{10} T = -0.115$) against ε, for the exact map and for means of m sampled embeddings.

A longer fed-back memory. The one-token window is the main simplification of the map. Feeding back the last K expected embeddings instead, so that the state is a window of K soft tokens at fixed positions after the prompt, we measured 24 temperatures between $\log_{10} T = -0.3$ and 0.6 for K from 1 to 16 (Table 30; twins at 10^{-5} and 10^{-7} agree within 0.006 at every chaotic temperature). In Qwen2.5-7B the chaos does not vanish but spreads: it is chaotic at 8 temperatures with one token of memory, at 13 with four, and at 21 and 20 with eight and sixteen, nearly the whole range, with smaller exponents (largest 0.96 for $K = 1$, 0.43 for $K = 16$). Qwen2.5-0.5B, chaotic at one temperature with one or two tokens of memory and at none with four, is chaotic at ten with eight. The chaos is therefore not an artifact of the one-token window.

memory K	1	2	4	8	16
Qwen2.5-7B	8 (0.96)	10 (1.08)	13 (0.70)	21 (0.58)	20 (0.43)
Qwen2.5-0.5B	1 (0.20)	1 (0.24)	0	10 (0.31)	3 (0.23)

Table 30: Chaotic temperatures, of 24 between $\log_{10} T = -0.3$ and 0.6, when the loop feeds back the last K expected embeddings (largest exponent in parentheses); paper's prompt, unsteered, float64, 500 transient and 1000 measured steps, twins at 10^{-5} and 10^{-7} (chaotic: both positive and no period up to 64).

Soft thinking with a growing context. Soft Thinking [1] feeds back the same expected embedding but appends it to the context, so that every earlier soft token stays visible. We ran that decoding in float64 from the paper's prompt and displaced the first soft token of twin trajectories by 10^{-10}, without renormalization; the distance between the newest soft tokens then shows whether the early perturbation grows. The context grows, so these are finite-time growth rates of a non-autonomous system, not exponents of an attractor. In Qwen2.5-7B the perturbation grows exponentially at temperatures inside the one-token map's chaotic range: by 0.77 to 0.82 per step at $\log_{10} T = 0.17$ and 0.21, reaching 10^{-3} within 49 to 58 steps, and it reaches 10^{-3} within 400 steps at nine other temperatures between -0.10 and 0.40, at 0.02 to 0.18 per step (0.8 at 0.33 for two of three twins). It does not grow at the sharpest distributions ($\log_{10} T \leq -0.14$) or at nearly uniform ones ($\log_{10} T \geq 0.44$), where it stays between 10^{-13} and 10^{-7} for 400 steps. In Qwen2.5-0.5B the growth moves to sharper distributions: it appears at six of the seven temperatures from $\log_{10} T = -0.30$ to -0.07 (up to 0.33 per step) and at 0.09, with a late transient burst at 0.05, and is absent from 0.13 up, which includes its narrow one-token chaotic band. Sensitive dependence therefore survives when the whole context is kept, and in the 7B model the one-token map locates it. Qwen2.5-72B, run the same way on two GPUs, stays nearly insensitive. We count a temperature as growing when two of its three twins reach 10^{-3} within 400 steps, a rule fixed before the 72B run and applied afterwards to the earlier 0.5B and 7B records, where it gives 11 of 24 temperatures for the 7B model and 8 for the 0.5B model. The 72B model grows at one of 24 temperatures on the paper's prompt, and only transiently: at $\log_{10} T = 0.09$ the twins pass 10^{-3} at step 123 and fall back to 5×10^{-6} (at 0.13 one twin of three passes 10^{-3} at step 381). It grows at none of the 24 completed temperatures on the held-out prompt with its narrow unsteered band (a planned finer line there was not completed), including $\log_{10} T = 0.326$, between two temperatures where its one-token map is chaotic; there the distance never exceeds 1.1×10^{-9}. It also grows at none of the eleven steered cells of Appendix I, including the six where its one-token map is chaotic: there the distance never exceeds 3×10^{-7} and ends between 1×10^{-13} and 6×10^{-10}, and the growing context holds the next-token entropy at 2.3 to 3.6 nats instead of the one-token map's 7.3 to 11.1. At these 72B prompts and cells, over 400 steps, the growing-context decoder therefore did not reproduce the sustained sensitivity of the fixed-window loop, while the one transient burst shows that the method registers growth at this scale.

Feeding back only the most probable tokens. The published Soft Thinking implementation does not feed back the whole distribution: it keeps the 15 most probable tokens and renormalizes their probabilities (Section 6). Such a map is no longer smooth. It is smooth while the set of kept tokens stays the same, and it jumps when a token enters or leaves the set. In an exploratory run we measured what this truncation does to the loop of Qwen2.5-0.5B (`tools/topk_loop.py`). The fed-back embedding is the average of the embeddings of the k most probable tokens, weighted by their renormalized probabilities, for $k = 50$, 20, 10, 5, 3, 2 and 1, and for comparison the full distribution; with $k = 1$ this is greedy decoding with a memory of one token. We scanned 128

temperatures between $\log_{10} T = -1$ and 0.3 on two lines, one without steering and one with the philosophy direction at $\alpha = 0.0743$, in float64, with 500 transient and 1500 measured steps and one twin at offset 10^{-6}. The setting 15 itself was not run; 10 and 20 bracket it.

A finite-difference exponent can be inflated when the trajectory and its twin keep different sets of tokens at some step, because the map jumps there. As for the mixture of experts (Appendix H), we therefore call a temperature chaotic only if its exponent is positive, no period is detected and the twin kept the same set as the trajectory at every measured step. If the two sets differed at least once we call it irregular.

tokens kept, k	all	50	20	10	5	3	2	1
steered line ($\alpha = 0.0743$)								
chaotic	22	8	14	2	3	0	0	0
irregular	0	8	3	0	0	0	0	0
no period, no positive exponent	2	3	3	3	0	0	0	0
largest exponent	0.48	0.37	0.32	0.24	0.11	−0.64	−0.21	–
unsteered line								
chaotic	1	0	0	0	0	0	0	0
irregular	0	0	1	0	0	0	0	0
largest exponent	0.21	−0.11	0.11	−0.12	−0.41	−1.61	−1.29	–

Table 31: Temperatures, of 128 on each line, at which the loop of Qwen2.5-0.5B is chaotic or irregular when only the k most probable tokens are fed back, and the largest exponent on the line (per step). Every temperature not counted in a row has a detected period. On the steered line at $k = 50$ and 20, and on the unsteered line at $k = 20$, the largest exponent belongs to an irregular temperature. At $k = 1$ both copies select the same token and their difference vanishes exactly, so no exponent is given. Float64, 500 transient and 1500 measured steps, one twin at offset 10^{-6}, periods up to 64; exploratory (`docs/lf3-tk/topk-summary.json`).

Truncation leaves the loop chaotic at fewer temperatures, and at none once very few tokens remain (Table 31, Figure 80a,b). On the steered line the full distribution is chaotic at 22 of the 128 temperatures. Keeping 50 tokens leaves 8 chaotic and 8 irregular temperatures, keeping 20 leaves 14 and 3, keeping 10 leaves 2 chaotic ones and keeping 5 leaves 3. With 3, 2 or 1 tokens every temperature has a detected period. The chaotic temperatures also move: of the 17 temperatures with a positive exponent at $k = 20$, 10 are among the 22 of the full distribution and 7 are not. On the unsteered line, where the full distribution is chaotic at a single temperature, no truncated loop is chaotic and one temperature is irregular at $k = 20$.

A second test follows two copies of the loop without renormalizing their distance. At the steered temperature $\log_{10} T = -0.0476$, where the full distribution is chaotic, we started 32 pairs of copies on the attractor at each of five offsets from 10^{-2} to 10^{-10}. Every pair separated to a tenth of the state's norm, after a median of 3 steps at the largest offset and 47 at the smallest (Figure 80c). With 20 tokens kept the loop at this temperature is periodic, and the same test shows it (Figure 80d): copies offset by 10^{-4} or less came together again in every repeat. Copies offset by 10^{-2}, about 3 % of the state's norm, mostly did not: 29 of the 32 pairs were still apart after 200 steps, at a distance that had neither grown nor shrunk over the last 100. We did not determine whether such a pair sits on two orbits or on one orbit at two phases.

Sampling changes the picture more than truncation does. The sampled loop feeds back the embedding of a single token drawn from the kept distribution. We ran it with twins that share their random numbers, so that both draw the same token unless the offset moves a probability

across the drawn number, and we gave the twin a fresh offset before every step. Over the 128,000 measured steps of each line at offset 10^{-6}, the twin drew a different token at 508 steps (unsteered) and 399 (steered) with the full vocabulary, at 1 and 0 with 20 tokens kept, at 2 and 0 with 50, and at none with 10 or fewer. With 20 tokens kept the counts were 25 and 70 at offset 10^{-4} and 2438 and 5647 at 10^{-2}, and zero at 10^{-8}. A sampled trajectory holds exactly one token, so its twin either coincides with it or differs by a whole token: a small offset is not amplified step by step, it is erased or, rarely, changes a draw. The separation test agrees. Sampled copies either parted at the first step or coincided exactly from then on; with the full vocabulary 48 of 480 runs parted, all at offsets of 10^{-4} or more, and with 20 tokens kept one did.

These results are exploratory: one small model, two lines, one twin offset for the exponents, and twins that share random numbers. Within these limits, truncation to the range of the published settings leaves the soft loop chaotic at some steered temperatures, fewer than with the full distribution, while the irregular temperatures stay undecided. The published method also appends each soft token to a growing context (previous paragraph), which this test does not.

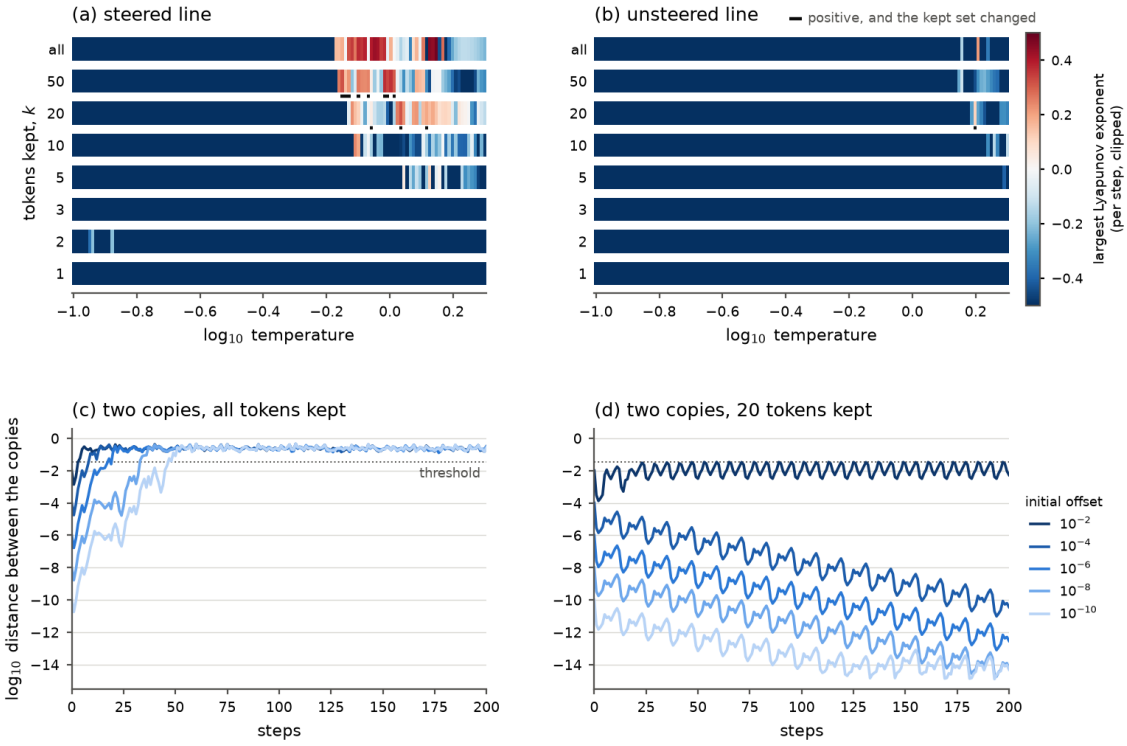

Figure 80: Feeding back only the k most probable tokens (Qwen2.5-0.5B, float64, exploratory). (a, b) Largest Lyapunov exponent at 128 temperatures on the steered line ($\alpha = 0.0743$ along philosophy) and on the unsteered line, one strip per k; color clipped at ± 0.5. A black bar under a cell marks an irregular temperature: a positive exponent, no detected period, and a twin whose kept set differed from the trajectory's at some step. (c, d) Distance between two copies of the loop at the steered temperature $\log_{10} T = -0.0476$, median over 32 repeats, for five initial offsets (the value at step 0), with all tokens (c) and with 20 tokens kept (d); no renormalization. Dotted line: a tenth of the state's norm (`tools/plot_topk.py`).

Ordinary generation at the soft map's settings. We also tested whether the regime of the soft-token map says anything about ordinary generation at the same temperature and steering strength. In an exploratory run we generated text from Qwen2.5-0.5B-Instruct in the ordinary

way: sampled tokens, the whole growing context, the paper's prompt and the philosophy direction added at layer 12, in float64 (`tools/real_inference_twins.py`). Each test used two copies of one generation. The first embedding of the second copy's prompt was nudged by a relative 10^{-8} or 10^{-3} in a random direction, and both copies drew their tokens with the same random numbers, so that they could differ only through the nudge. We generated 300 tokens for 32 such pairs at each setting, with two samplers: plain sampling from the whole distribution at the setting's temperature, and the sampler the model ships with (the 20 most probable tokens, a cumulative-probability cutoff of 0.8 and a repetition penalty of 1.1) at the same temperature. We recorded two things. The first is whether the two texts differ anywhere. The second is the rate at which the distance between the copies' next-token distributions (their total variation distance) grows or shrinks per generated token while the texts still agree; for each setting we quote the median of that rate over its 32 pairs.

The settings lie on a line of 15 steering strengths, α from 0.073342 to 0.073605 at $\log_{10} T = -0.1530$, that crosses the right edge of the period-6 window of Figure 38, the side on which the period-6 orbit ends at a fold. On it the soft-token map has period 6 at the first seven settings (exponents from -0.35 to -0.10 per step) and is chaotic at the other eight (0.14 to 0.22). As a control without an edge we ran the same line with a second prompt, the opposite system prompt of Appendix J, for which the soft map has period 6 at all 15 settings (exponents from -0.78 to -0.75). Before the run we wrote down what a carry-over of the regime would look like: a change of both measures at the edge, more parting and faster growth where the soft map is chaotic, and rates that follow the sign of the soft map's exponent. None of the three was observed.

With the smaller nudge no pair of texts differed, at any setting, with either sampler, and every median rate stayed within 0.0011 per token of zero. With the larger nudge and plain sampling, 305 of the 480 pairs on the line differed: 143 of 224 where the soft map is periodic and 162 of 256 where it is chaotic (64 and 63 %; Figure 81b). The counts per setting, 17 to 25 of 32, scatter as equal fractions would (a χ^2 test gives $p = 0.69$), with no change at the edge. The second prompt, without an edge, gives the same picture: 306 of 480 pairs, 17 to 25 per setting. The model's own sampler lets far fewer pairs differ: 14 of 480 on the paper's line (6 of 224 where the soft map is periodic, 8 of 256 where it is chaotic) and 15 of 480 on the second prompt's. At every setting, with either sampler and either nudge, the median rate stays within 0.0021 per token of zero, while the soft map's exponents at the same settings range from -0.78 to $+0.29$ per step.

We also followed the pairs after their texts had parted. Of the 611 pairs that parted under plain sampling on the two lines, 4 came together again (50 identical tokens in a row), and 183 parted after token 150, too late for the test that follows. For the other 424 we restarted each copy 50 tokens after the parting, from its own text and with the same random numbers, four times with a fresh nudge of 10^{-8} (and once with none, as a control). In 423 of them every restart reproduced its copy's next 100 tokens exactly; in the remaining one, not every restart did. A copy that has parted from its twin is thus as insensitive to small nudges as before.

The line compares neighbors, whose soft-map states differ in more than their regime. As a second comparison we chose four pairs of cells from the map of Figure 11, one periodic and one chaotic cell per quadrant of the plane, with next-token entropies of the soft map equal within 0.014 % (Figure 81e). Under plain sampling and the larger nudge, 78 of 128 pairs differed in the periodic cells and 78 of 128 in the chaotic ones. With the model's own sampler the counts were 1 and 5 of 128. Fisher's exact test gives $p = 0.21$ for that difference; the receipt also holds a bootstrap interval that excludes zero, but it resamples the four cells as if each cell's fraction were exact, which understates the uncertainty of counts this small. We conclude that no difference between the regimes was established.

A reading consistent with these numbers is the following. While the two texts agree, the copies differ only by the nudge itself, which stays in the context; no continuous quantity is fed back that

could carry a difference from one step to the next. The rates near zero say that the effect of this fixed nudge on the next-token distribution neither grows nor fades over 300 tokens. The texts part only when a draw falls between the two copies' cumulative probabilities. How often that happens depended on the size of the nudge and on the sampler, and not detectably on the regime of the soft map. The test is narrow: one small model, one line and four matched pairs of cells, 300 tokens, a nudge of the first prompt embedding only, and twins that share their random numbers. It does not show that ordinary generation is insensitive in general, and it leaves open other perturbations, longer texts and larger models (receipt: `docs/ri/ri1-receipt.json`).

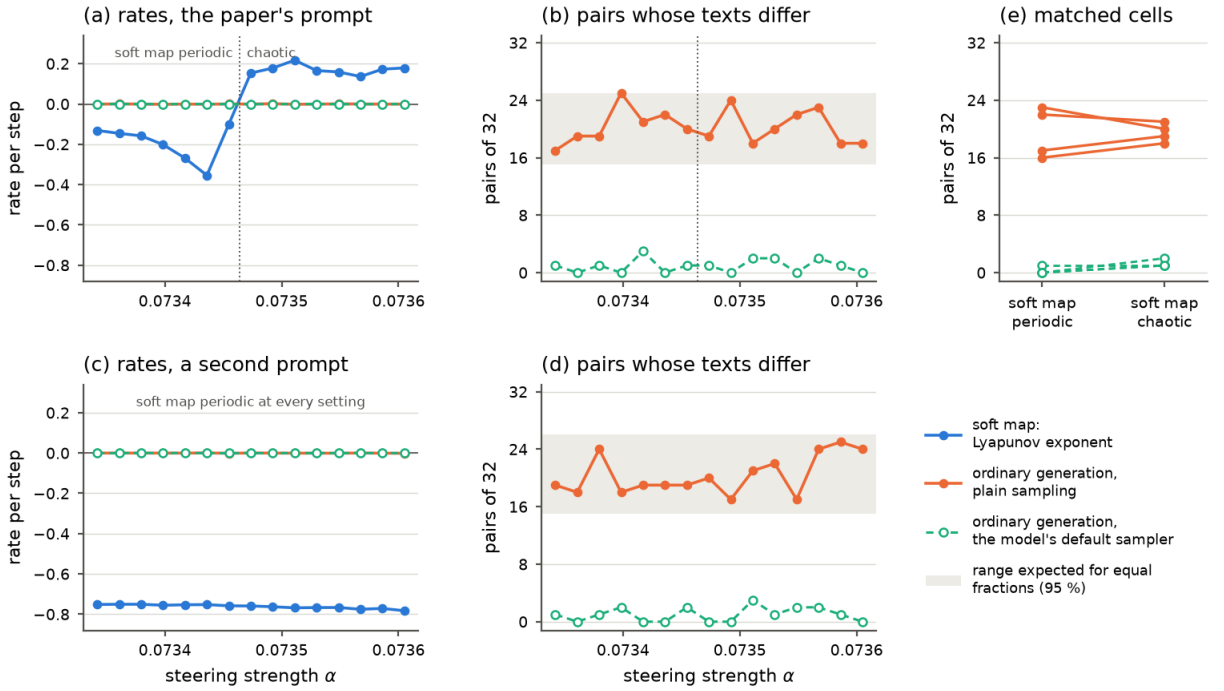

Figure 81: Ordinary sampled generation at settings of the soft-token map (Qwen2.5-0.5B, float64, exploratory; nudge 10^{-3}, 32 pairs per setting). (a) The line across the right edge of the period-6 window, paper's prompt: the soft map's Lyapunov exponent (blue) and, for ordinary generation, the median rate at which the distance between the two copies' next-token distributions changes per token, with plain sampling (orange) and with the model's default sampler (green, dashed). Both rates lie on zero, and the dashed line covers the orange one. The dotted line separates the settings where the soft map is periodic from those where it is chaotic. (b) Pairs, of 32, whose texts differ within 300 tokens. Gray band: the range expected for plain sampling if the fraction were the same at all 15 settings (95 %). (c, d) The same for a second prompt, with which the soft map is periodic at every setting. (e) Four matched pairs of cells: pairs of 32 whose texts differ at the periodic and at the chaotic cell of each pair. With the nudge 10^{-8} no pair of texts differed at any setting (`tools/plot_real_inference.py`).

H Replication and scale

In Qwen2.5-1.5B (steering at layer 14 of 28), a coarse 64^2 float64 map over the range of the 0.5B maps again shows chaotic regions with embedded windows. A 256^2 zoom (Figure 56) resolves a chaotic band near $\log_{10} T = 0$–0.13 (14,422 of 65,536 cells), with a branch descending along $\alpha \approx 0.06$ and periodic windows inside it. The large periodic regions are question loops: "What?\n" (period 2), "\nWhat is a" (period 4) and "What continuum?\n" (period 3). Unlike the 0.5B model, no

period-adding fan was resolved in this map.

A cut of 2048 cells at $\alpha = 0.100$ (Figure 82) shows how the band is entered from below. The 2-cycle "What?\n" gives way abruptly to a period-4 loop, and the Lyapunov exponent jumps, which is consistent with a switch between attractors rather than a resolved smooth doubling. The period-4 loop then doubles, with float64 flips at $\log_{10} T = -0.0031269$ (period $4 \to 8$, next modulus 0.014), -0.0021666 (period $8 \to 16$, next modulus 8×10^{-5}), -0.0019002 (period $16 \to 32$, next modulus 5×10^{-9}) and -0.0018380 (period $32 \to 64$). Along the way the exponent rises continuously to zero, and chaos follows. The interval ratios are 3.60, 4.28, 4.58, 4.648 and 4.665; the last three use the flips of the period-64, period-128 and period-256 orbits at $\log_{10} T = -0.0018244$, -0.00182151 and -0.00182088 (next moduli below 3×10^{-14}). The last ratio differs from Feigenbaum's 4.6692 by 0.0044 (0.1 %); the period-256 flip is located to about 3×10^{-14} in $\log_{10} T$, the floor set by float64 parameters. Like those of the 0.5B crescent at $\log_{10} T = -0.148$, the ratios rise toward 4.669. Their higher start does not indicate a temperature trend, since along nearby cuts of the 0.5B crescent the early ratios range from 2.84 to 14.9 (Section 9).

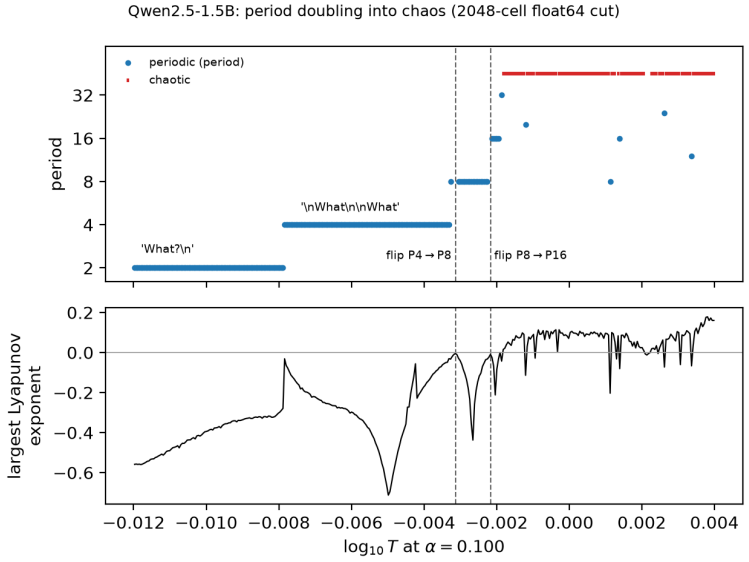

Figure 82: Qwen2.5-1.5B soft-token map steered along the philosophy direction at layer 14, along $\log_{10} T$ at $\alpha = 0.100$: period of the attractor reached from the fixed start (top) and largest Lyapunov exponent (bottom). Dashed lines: float64 flip parameters. The doubled loops alternate the number of line breaks between the "What"s.

In Qwen2.5-0.5B, a second steering direction, hostility, reproduces the fold-and-adding mechanism (Section 10).

Qwen2.5-7B (steering at layer 14 of 28, concept vectors built as for the smaller models) shows the same kind of organization in coarse 64^2 float64 maps (Figure 83). Over the philosophy range used for 0.5B and 1.5B, 8 % of cells have a positive largest exponent, and periods 2, 4, 3 and 11 dominate the periodic cells. Along the hostility direction ($\alpha \in [0, 0.12]$, $\log_{10} T \in [-0.3, 0.3]$), 24 % of cells are chaotic. Bands of period-3, -5 and -6 windows lie in a chaotic sea below $\log_{10} T \approx 0.1$, with high-entropy chaos above, except for a period-3 region at strong steering ($\alpha \gtrsim 0.09$). The windows hold genuine periodic orbits. In float64, Newton's method converges at the most interior period-6 cell to the loop ", shut the end\n\nOh" at $(\log_{10} T, \alpha) = (-0.0143, 0.0476)$ (residual 2×10^{-14}; spectral radius 1.6×10^{-7}, nearly superstable), and at the most interior period-3 cell to a period-3

loop at $(0.2524, 0.1143)$ (residual 3×10^{-15}; leading multiplier -0.12). Continued for about an hour each, over ± 0.001 in α and -0.016 in $\log_{10} T$ (one float64 continuation step takes about two minutes at 7B), both orbits stay strongly stable without a crossing. Where the coarse map switches from the period-6 loop to the period-3 loop "you better than" (between $\alpha = 0.0495$ and 0.0514 at $\log_{10} T = -0.0143$), both loops are float64 orbits (residuals 7×10^{-15} and 1×10^{-14}), and their continuations overlap. The period-6 orbit stays stable from $\alpha = 0.0476$ up to 0.0507 and the period-3 orbit from 0.0514 down to 0.0475, both with spectral radius below 1.4×10^{-3}. The two attractors therefore coexist where both continuations are stable, $\alpha \approx 0.0476$ to 0.0507, and within the tested range the map's edge there is the boundary of the fixed start's basin, not a bifurcation. These short continuations located no bifurcation; a period-2 fold of the 7B model is located below.

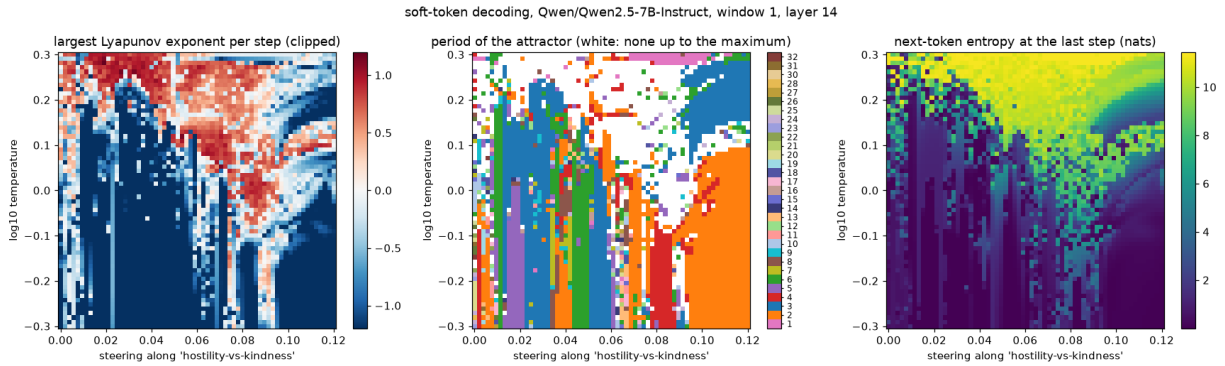

Figure 83: Qwen2.5-7B soft-token map steered along the hostility direction (layer 14), 64^2 float64 cells: largest Lyapunov exponent, period of the attractor reached from the fixed start, and next-token entropy at the last step.

A second model family. OLMo-2-1124-7B-Instruct [137] comes from another laboratory, with other data, training and code. Steered at layer 16 of 32, with concept vectors built exactly as for Qwen, it shows the same kind of organization (Figure 84). In 64^2 float64 maps, 6 % of cells are chaotic along the philosophy direction and 13 % along the hostility direction, in a band near $\log_{10} T = 0.03$–0.17 that bends to lower temperatures under strong hostility steering ($\alpha > 0.07$). Below the band the attractor is the 2-cycle "!Hello"; above it, at high temperature, the fixed point "!"; inside it lie periodic windows of periods 4 to 31. Along the hostility direction the largest windows, of periods 8 to 12, are loops of k laps of the 2-cycle "!<|endoftext|>" with a short excursion ("there is an"). Each lap adds two to the period, as each lap of "!Yes" does in the hostile cascade of Qwen2.5-0.5B (Section 10). Whether a fold of that 2-cycle organizes them is not tested: these are coarse pixels, and no OLMo orbit has been continued.

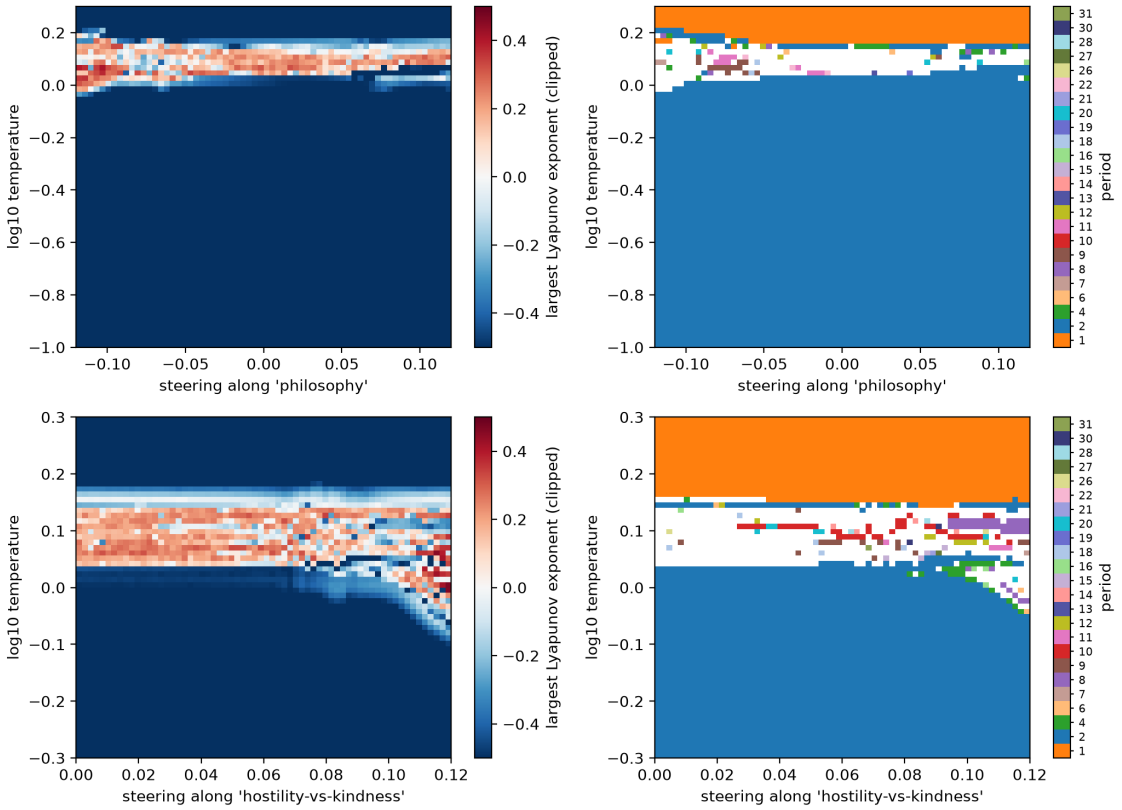

Figure 84: OLMo-2-1124-7B-Instruct soft-token maps, 64^2 float64 cells each, steered at layer 16 along the philosophy direction (top) and the hostility direction (bottom): largest Lyapunov exponent (clipped) and period of the attractor reached from the fixed start (white: no period up to the maximum).

A third family, read with sparse autoencoders. Llama-3.1-8B-Instruct [138], with official weights, steering at layer 16 of 32 and vectors built as for the other models, shows the same organization in all three planes (Figure 85). In 64^2 float64 maps, 8 % of cells are chaotic along gratitude, 12 % along philosophy and 28 % along hostility, with windows of periods up to 32. Its largest loops are greetings between AIs ("-Greetings fellow AI", ". Greetings fellow AI entity detected", period 7). Along hostility, too, the largest loops are greetings and "You", not hostile text. We read 24 of its attractors (the six largest families and the two most chaotic regions of each map) with two sparse autoencoders on the residual stream: one trained on this Instruct model (Goodfire, layer 19 [139]) and one on the base model (TopK, layer 24 [140]). We first checked each one's hook point by its reconstruction of ordinary text. The Instruct-trained autoencoder separates the attractors (their top-15 features overlap by a mean Jaccard index of 0.19 within a steering direction and 0.11 across directions), whereas the base-model autoencoder sees mostly a signature common to every attractor (overlap 0.52–0.86). Both reconstruct the attractors' states less well than ordinary text (median relative error 0.64–0.76 against 0.52–0.56). The readout is therefore descriptive: without interventions it identifies which learned features an attractor occupies, not a mechanism, and since autoencoders trained with different seeds learn different features [141], the features are those of these two autoencoders.

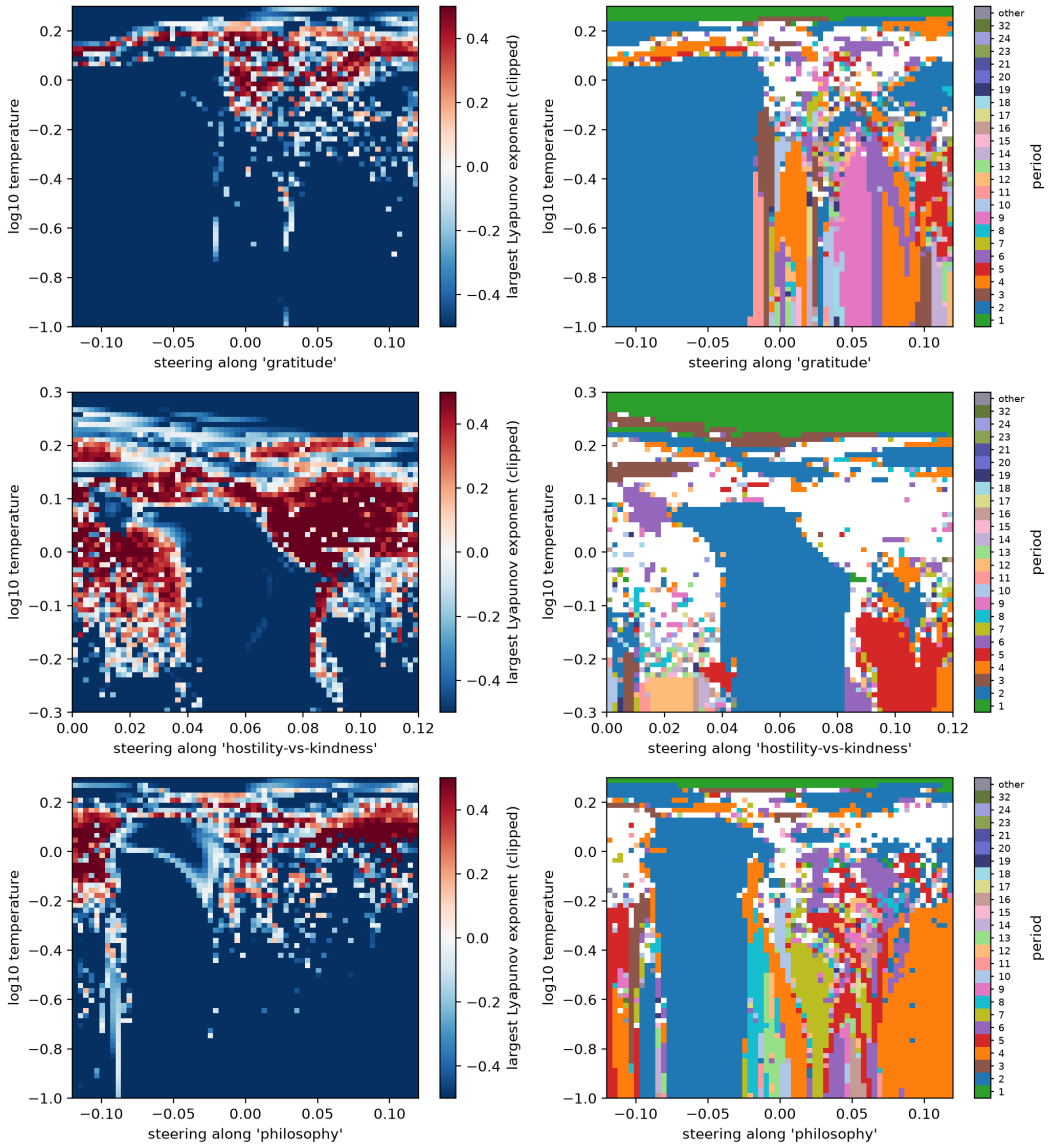

Figure 85: Llama-3.1-8B-Instruct soft-token maps (64^2 float64 cells, steering at layer 16) along gratitude, hostility and philosophy: largest Lyapunov exponent (clipped) and period of the attractor reached from the fixed start (white: no period up to 32).

Maps of 128^2 cells (Figure 86) resolve the windows: the hostility plane has 29 % chaotic cells, as at 64^2, and the philosophy region below the chaotic band is filled with nested periodic tongues of periods 2 to 13.

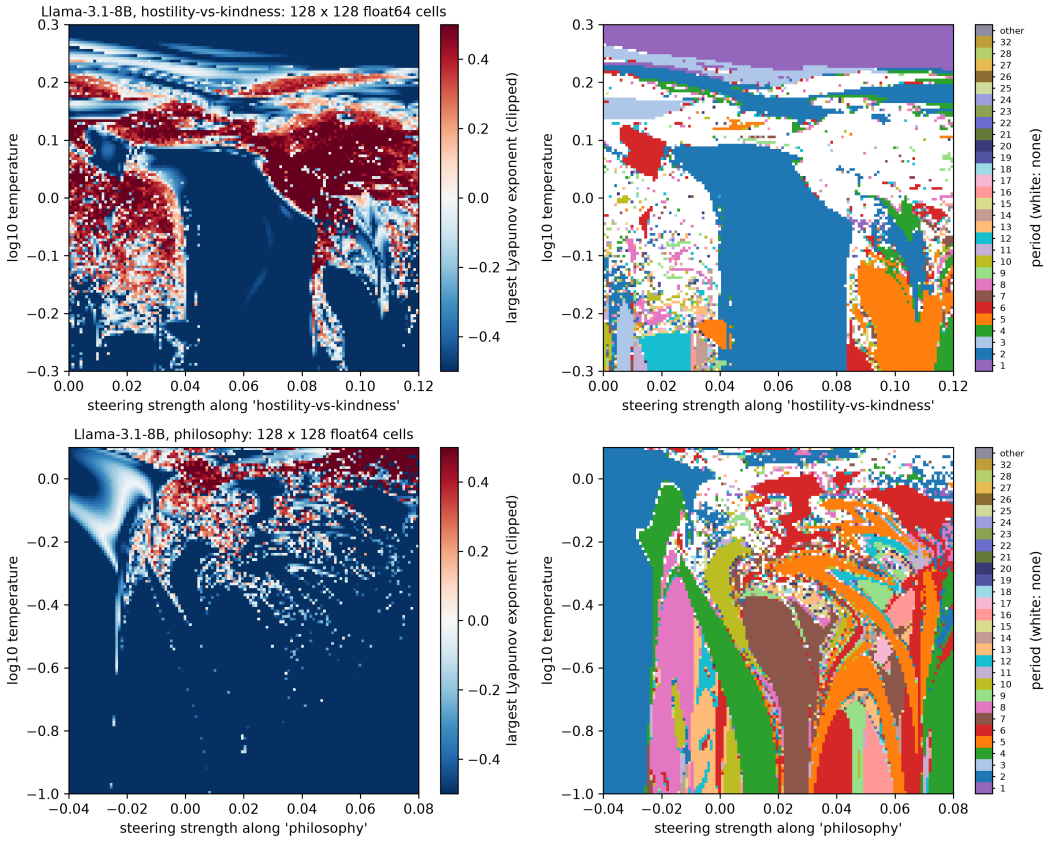

Figure 86: Llama-3.1-8B-Instruct soft-token maps steered at layer 16, at 128^2 float64 cells (each map merged from four 64^2 tiles on halves of one grid): the hostility plane (top) and the philosophy region of the 64^2 map's windows (bottom). Largest Lyapunov exponent (clipped) and period (white: none up to 32).

Same architecture, other weights. Qwen2-0.5B-Instruct has exactly the architecture and size of Qwen2.5-0.5B-Instruct (24 layers of width 896, the same vocabulary size and rotary base) but was trained on other data with another recipe. Built identically, their 128^2 float64 maps along gratitude and philosophy (Figure 87) share the gross layout. A large region at low temperature around $\alpha = 0$ is held by one periodic loop, a band of chaos with embedded windows lies near $\log_{10} T = 0$–0.25 (4–8 % of cells chaotic in each map), and thin window families run along the flanks. The attractors themselves differ almost completely. Of 416 to 900 distinct loops per map, only 10 occur in both models, all short and generic (".", "!", " you are", " is it"), and they cover less than 10 % of either model's cells. Qwen2's planes are dominated by one period-4 loop, "!<|im_end|>\nHello" (an end of turn followed by a new greeting; 41 and 59 % of cells), and Qwen2.5's by period-5 and period-2 loops (". The two checkpoints share the organization of the map; the weights set its content.

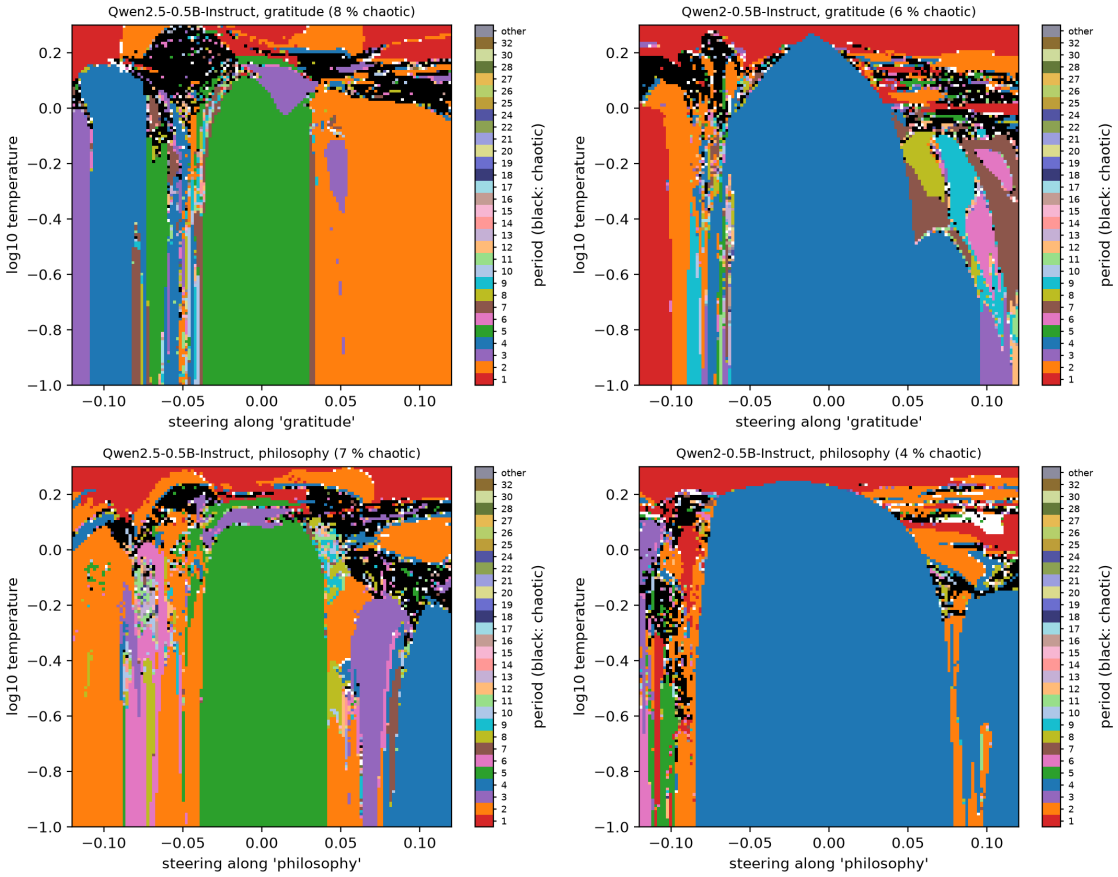

Figure 87: Same architecture, other weights: 128^2 float64 maps of Qwen2.5-0.5B-Instruct (left) and Qwen2-0.5B-Instruct (right) steered along gratitude (top) and philosophy (bottom) at layer 12, with vectors built identically. Color: period of the attractor reached from the fixed start; black: positive largest exponent.

The fold law at 1.5B and 7B. We tested whether the mechanism of Section 8, windows accumulating on the fold of a short cycle with $d_N = C/(N + N_0)^2$, recurs in larger models. The criterion, a constant $d_N(N+N_0)^2$ within 10 % over at least four windows for one offset N_0, was fixed before the lines were computed; a disclosed correction to the period-progression selector after the first scoring kept the conclusions for the nearest four windows and changed which farther windows were selected. In Qwen2.5-1.5B (philosophy direction, layer 14, $\alpha = 0.058588$) the period-4 orbit folds at $\log_{10} T_c = -0.0912029$ (multiplier 1 to 3×10^{-14}, next modulus 0.008), with normal-form coefficients $a = -7.16$ per unit $\log_{10} T$ and $b = -21.09$, so $ab = 151.0$. A float64 line of 1700 cells, 1570 of them above the fold, shows windows of periods $4N$ accumulating on it, from single cells at periods 44–64 to a wide period-12 window. Measured in $\log_{10} T$ from $\log_{10} T_c$ to their edges nearest the fold, the four nearest resolved windows (periods 28–40) meet the criterion within 0.12 % ($\delta = -0.99$, $C = 0.0612$). With the windows farther out the spread rises to 10.3–14.0 %; there several periods recur in multiple disconnected runs on the line's 10^{-5} grid, and grid-spaced edges no longer identify the fold edges. In Qwen2.5-7B ($\alpha = 0.051429$) the period-2 orbit folds at $\log_{10} T_c = -0.2040156$ (next modulus 0.0092, $ab = 55.7$), and odd windows accumulate on it. The nearest four and five (periods 11–17 and 9–17) meet the criterion within 4.5 %, and six (to period 7) within 6.1 %. Each window's own doubling or tripling lies on its far side, as for the 0.5B windows of periods 16–20; this organization is consistent with fold edges on the near side, and the fitted

distances use grid estimates of those edges. The effective passage fraction $f = \sqrt{Cab}/\pi$ is about 0.97 at 1.5B and 0.78 to 0.81 at 7B, against 0.58 at 0.5B. The fold and its window law thus recur at all three sizes, while f, an effective constant that absorbs the global reinjection, differs between the models. With one fold per model, crossed in different coordinates (α at 0.5B, $\log_{10} T$ at 1.5B and 7B), no trend of ab with size can be read.

A mixture-of-experts model. Qwen3.5-35B-A3B differs from the dense models in two ways. Thirty of its 40 layers use linear (Gated DeltaNet) attention, with gated full attention in every fourth layer, and each layer routes a token to 8 of 256 experts plus a shared expert, so that about 3B of its 35B parameters are active. Hugging Face computes several parts of this architecture in float32 even when the model is cast to float64 (the normalizations, the delta-rule recurrences, the rotary tables and the router softmax), and some tensors are stored in float32; we patch these paths, restore the stored tensors exactly, and disable the thinking block that the default chat template opens (a first run with the block open is kept in the research log). Steered at layer 20, the map behaves unlike those of the dense models. Unsteered, 128 temperatures over $\log_{10} T \in [-1, 0.3]$ give no positive exponent (largest -0.265): periodic loops at all but 15 temperatures ("!Hello" with period 2 at 92 of them, other periods up to 52), and the remaining 15, above $\log_{10} T \approx 0.1$, contract without a detected period. Along philosophy, none of 1024 cells is chaotic (largest exponent -0.112). Along hostility, 52 of 1024 cells (5.1 %) have a positive exponent and no period, against 28 % for Llama-3.1 8B on the same plane. Routing is discontinuous: a twin trajectory that crosses a routing boundary jumps, which can inflate a finite-difference exponent, so we call these cells irregular rather than chaotic.

Two dense models of the same family, Qwen3.5-4B and Qwen3.5-9B (32 layers each, 24 of them linear attention), test whether the architecture causes the contraction. With concept vectors at their middle layer (16 of 32) and the same ranges, transients and maximal periods, the 4B model is chaotic at 19 of 128 unsteered temperatures (largest exponent 0.79), at 153 of the 1024 philosophy cells (0.84) and at 323 of the 1024 hostility cells (0.89). The 9B model is chaotic at 10, 73 and 161 (0.48, 0.68 and 0.75; Figure 88). The hybrid attention therefore does not prevent chaos. The routing-aware version of the stack resolution (Appendix B) follows a third row whose routing is frozen to the base orbit's. At the twin offset 10^{-7}, neither mixture of experts switched an expert in any layer at any step, and the frozen row's network factor equals the free twin's. At every knockout temperature, the first attention layer (a linear-attention layer) is the largest amplifier in all three Qwen3.5 models. At the knockout temperature where the entropy is 2 to 4 nats, removing it lowers the network factor by 1.27 (4B), 1.33 (9B) and 0.76 nats (the mixture of experts). Removing the first MLP or expert layer lowers it by 0.92, 0.77 and 0.24.

OLMoE-1B-7B stores its experts in bfloat16. We convert each expert's weights to float64 when it is used, which is exact, compute the router softmax in float64, and cast everything else to float64; central differences of the map then converge as h^2. Its unsteered sweep of 130 temperatures over $\log_{10} T \in [-1, 1.2]$ has one temperature with a positive exponent and no period (0.023, irregular), and 7 with neither a detected period nor an exponent above 10^{-3}. Its median network factor between 0.3 and 6 nats (-0.36) lies between those of the dense OLMo-2-1B ($+0.11$) and OLMo-2-7B (-0.52). OLMoE normalizes before each sublayer and OLMo-2 after it, so this pair differs in more than the experts, and we read no size effect from it.

A model built mostly from convolutions. The models above are stacks of attention blocks (full attention or, in most layers of Qwen3.5, linear attention). LFM3-1.7B [142] is built differently: of its 24 blocks, 13 are gated convolutions over the last 7 positions, 5 attend to a sliding window of

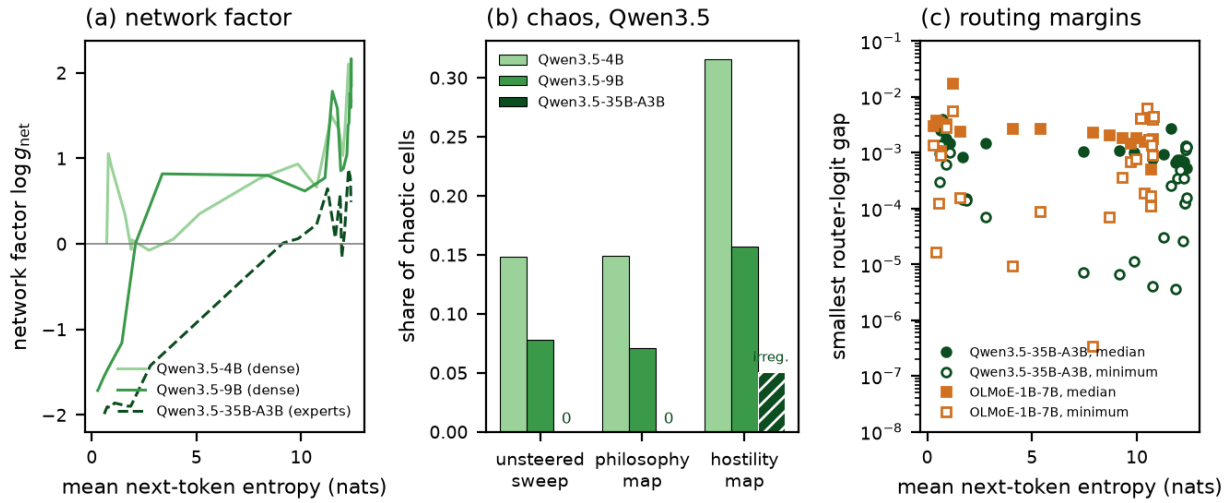

Figure 88: Mixtures of experts against dense models of the same family. (a) Network factor against the mean next-token entropy on the orbit, unsteered loop, 24 temperatures per model, float64 (as in Figure 58), for Qwen3.5-4B and -9B (dense) and Qwen3.5-35B-A3B (mixture of experts, about 3B of 35B parameters active per token). (b) Share of chaotic cells (by the criterion of Appendix B) in the unsteered sweep (128 temperatures, $\log_{10} T \in [-1, 0.3]$) and in 32×32 maps steered along philosophy and hostility at the middle layer; the mixture of experts' hostility cells (hatched) are called irregular, since routing jumps were not excluded when that map was computed. (c) Routing margins along the unsteered orbits of both mixtures of experts, Qwen3.5-35B-A3B and OLMoE-1B-7B (24 temperatures each, as in (a)): per temperature, the median over steps (filled) and the minimum (open) of the per-step smallest gap between the eighth and ninth largest router logits over all layers. A switch requires closing this pairwise logit gap; a bound below half the gap on every logit's change prevents it. A twin at relative offset 10^{-7} never switched an expert.

256 tokens and 6 attend to the whole context. In an exploratory run we scanned the soft-token loop of its two released checkpoints, the pretrained base (LFM3-1.7B-Base) and the model post-trained from it (LFM3-1.7B), with the paper's prompt in each checkpoint's chat template (its thinking block disabled, as for Qwen3.5) and steering at layer 12 of 24. The weights are stored in a 16-bit format, so the cast to float64 is exact. The published model code computes its normalizations and rotary tables in float32; we patch these paths (`tools/lfm3_float64.py`). Before the run we checked that every module's output is float64, that central differences of the one-step map shrink as h^2, and that a step computed from the cached prefix equals a full recomputation to within 10^{-15} (`tools/lfm3_precision_check.py`).

Both checkpoints have chaotic cells, without and with steering (Figure 89). Without steering, the base checkpoint is chaotic at 3 of 160 temperatures between $\log_{10} T = -1$ and 1.2 and the post-trained one at 12, with largest exponents of 0.45 and 0.50 per step; all 15 temperatures lie between $\log_{10} T = -0.08$ and 0.18. On 32×32 maps the base checkpoint is chaotic at 20 of 1024 cells with steering along philosophy and at 63 along hostility-vs-kindness, and the post-trained one at 162 and 280. In each scan the post-trained checkpoint is chaotic at about four to eight times as many cells as the base. The median network factor between 0.3 and 6 nats (Section 14) differs in the same direction, toward less damping: -0.42 for the base checkpoint and $+0.53$ for the post-trained one, over 8 and 9 temperatures. This is one prompt and one pair of checkpoints, and we do not read it as an effect of post-training in general. The steering is weak by these models' own measure: the natural contrast between the two groups of texts corresponds to α between 0.27 and 0.55, against 0.043 in Qwen2.5-0.5B (Section 2), and the maps stop at 0.12.

These labels are coarser than those of the models above. Each rests on one twin offset (10^{-6}), one prompt and 300 measured steps per map cell; no periodic orbit was verified or continued, and no window was resolved. We therefore report the scan as an observation that the loop of a model built mostly from convolution blocks can also be chaotic, and we do not count LFM3 among the fifteen models of Section 14 or among the model families of the main text (summary and receipt: `docs/lf3-tk/lfm3-summary.json`, `docs/lf3-tk/receipt.json`).

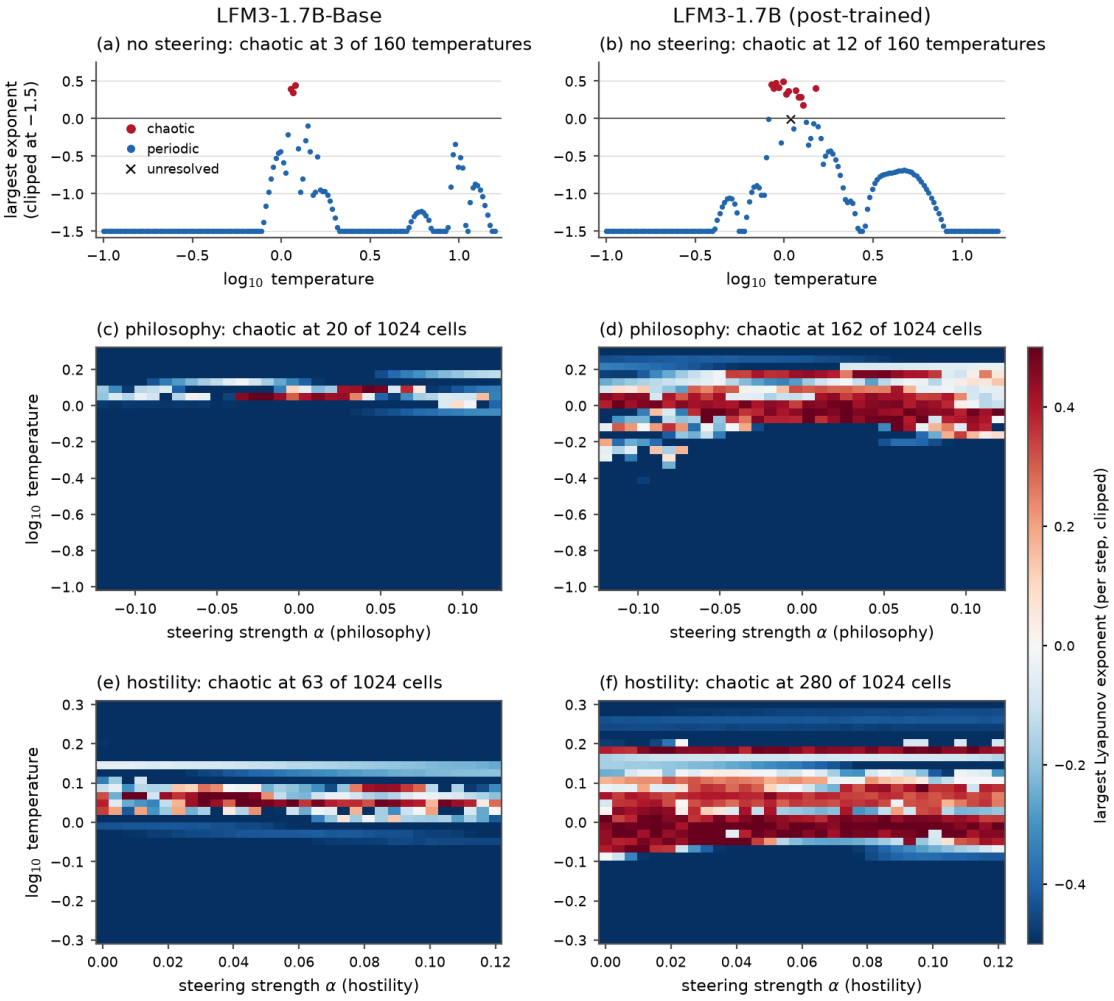

Figure 89: The soft-token loop of LFM3-1.7B, a model built from convolution and attention blocks (exploratory; float64, the paper's prompt): the pretrained base checkpoint (left) and the post-trained model (right). (a, b) No steering: largest Lyapunov exponent at 160 temperatures (1000 transient and 2000 measured steps), drawn at -1.5 where it is lower. Red: chaotic (an exponent above 10^{-3} and no detected period up to 64); blue: a detected period; black cross: neither (one temperature, in b). (c, d) Steering along the philosophy direction at layer 12: 32×32 cells, 300 transient and 300 measured steps, periods up to 32; color: largest exponent, clipped at ± 0.5. (e, f) Steering along hostility-vs-kindness, as (c, d). Each panel's title gives its number of chaotic cells; the 3 to 12 cells per map with neither a period nor a positive exponent are not marked. Every label rests on one twin offset, 10^{-6} (`tools/plot_lfm3.py`).

I Chaos and model size: further results

This appendix completes Section 14: the 70B-class models and their precision check, cheap predictors of the chaos, steering at 70B, and the supporting figures.

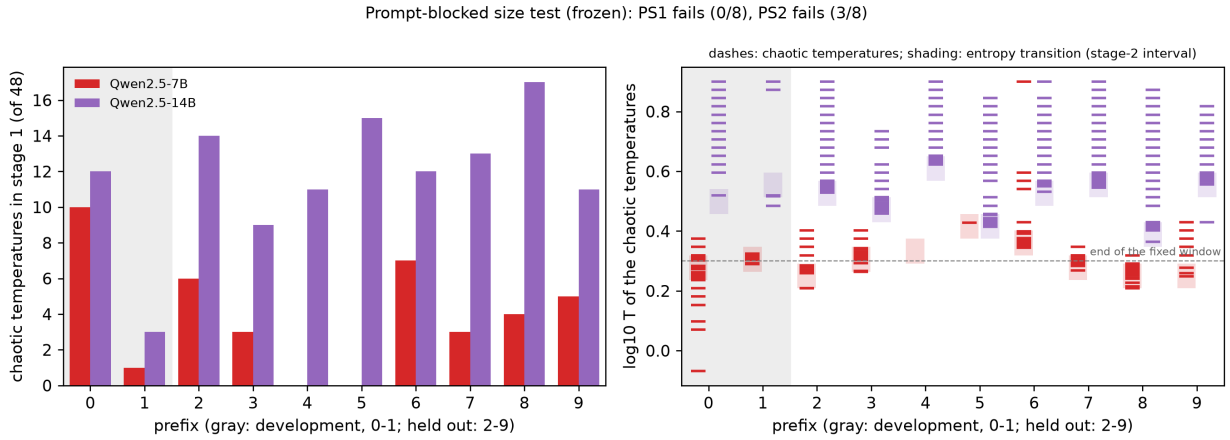

Figure 90: Prompt-blocked size test (frozen before any run), Qwen2.5-7B and -14B on ten prompts (gray: development). Left: chaotic temperatures among 48 over $\log_{10} T \in [-0.4, 0.9]$, the primary endpoint. Right: where the chaotic temperatures lie (dashes; both stages), with each model's entropy transition, the interval of its largest entropy increase (shading). Chaotic: exponents above 10^{-3} at both twin offsets (10^{-5}, 10^{-7}) and no period up to 64; 1000 transient and 2000 measured steps, float64.

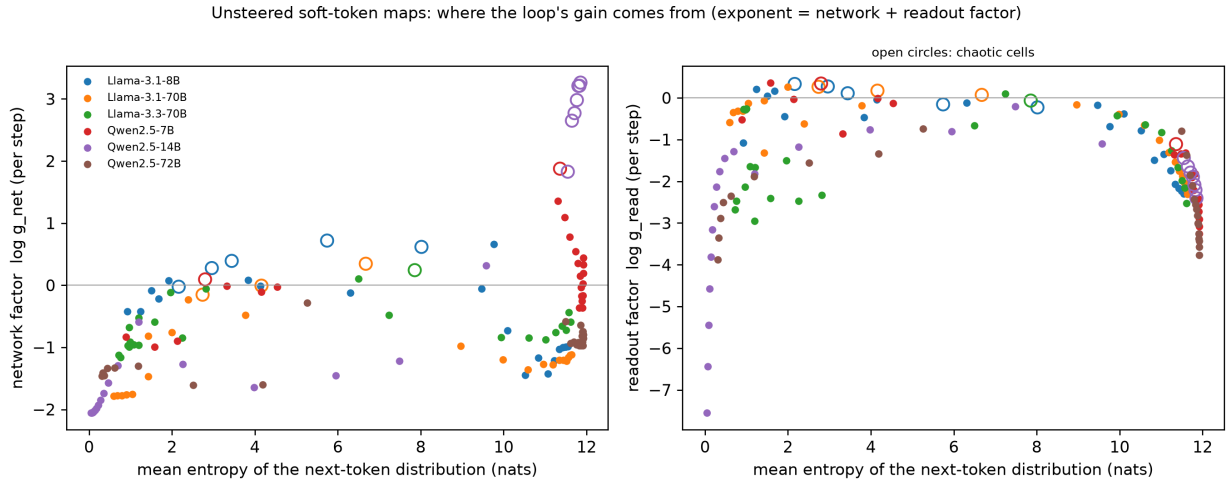

Figure 91: Gain split of the unsteered loop (Section 3) at 24 temperatures per model (500 transient and 1000 measured steps, relative twin offset 10^{-6}): network factor (left) and readout factor (right) against the mean next-token entropy on the orbit; open circles: chaotic temperatures. The two factors add up to the largest exponent up to a boundary term. Medians per entropy bin are given in Section 14.

At 70B, unsteered chaos is rare, and a precision defect made part of it spurious. The regular grids of the 70B-class models find little unsteered chaos: Qwen2.5-72B has no positive exponent at any of 191 temperatures between $\log_{10} T = -1$ and 1.2, Llama-3.3-70B none in $[-1, 0.3]$, and Llama-3.1-70B three temperatures in $[-1, 0.3]$. Finer lines found candidates: a window of

positive exponents at Qwen2.5-72B's entropy jump ($\log_{10} T = 0.3728$–0.376) and two isolated temperatures of Llama-3.3-70B (0.250, 0.306), all at the single twin offset 10^{-6}. The window failed a three-offset test: the exponent is negative at all 64 temperatures of the window at 10^{-5}, near zero at 10^{-6}, and positive at 60 of the 64 at 10^{-7}, including four temperatures whose state repeats with a period, which no stable periodic orbit allows; in a repeat it is negative at all 16 temperatures at 10^{-4}, 3×10^{-5} and 10^{-5}. The growth of small perturbations came from our computation, whose two-GPU runs computed the final normalization in float32 (see the start of Section 14); a controlled reproduction of this false chaos is in Appendix B (Figure 64). Recomputed in float64 throughout, the window's orbits have periods 2 and 4, and their exponents are negative at all four offsets from 10^{-4} to 10^{-7} and for an automatic-differentiation tangent vector carried along the orbit by forward-mode differentiation, which needs no offset; the four twin exponents lie within 0.015 of the tangent exponent at every temperature, and the exponent approaches zero only where the period doubles from 2 to 4 ($\log_{10} T = 0.374$). The selected cells below kept their labels at the larger offsets we tested; this is a finite check, not a qualification of every cell of the original maps. In the steered Llama-3.3-70B maps below, twelve randomly drawn chaotic cells are chaotic at offsets 10^{-4}, 10^{-5} and 10^{-6} alike, with exponents that agree within 0.052, and the drawn periodic and unresolved cells stay negative at all three; on one GPU, the exponents at 10^{-7} and 10^{-5} agree within 0.021 at every chaotic temperature of the prompt-blocked test. The same test on Llama-3.1-70B's three unsteered chaotic temperatures confirms two of them ($\log_{10} T = 0.136$ and 0.198; exponents 0.23 to 0.26 and 0.40 to 0.41 at the three offsets) and rejects the weakest (0.116; positive only at 10^{-6}), so unsteered chaos exists at 70B, if rarely. Two Qwen2.5 models are chaotic far above the fixed window on the paper's prompt, where the next-token distribution is nearly uniform: the 1.5B model at 55 of 224 temperatures, between $\log_{10} T = 0.28$ and 1.04, and the 14B model at 53, between 0.51 and 1.15 (exponents up to 1.14), at mean entropies up to 11.9 nats against $\log V = 11.93$. The other models of the series are not chaotic above $\log_{10} T = 0.6$. For the 1.5B model a longer test (3000 measured steps) at twin offsets 10^{-5} and 10^{-7}, at 12 temperatures from 0.288 to 0.58, finds 9 chaotic and 3 periodic, with the two exponents agreeing within 0.003 at every temperature; one of the periodic ones ($\log_{10} T \approx 0.394$) has positive exponents at both offsets, a contradictory label that we report as returned. On these sampled trajectories the upper sweeps become non-chaotic above about $\log_{10} T = 1.1$ in these two models; we did not locate the uniform contraction that Section 3 proves for large enough temperature.

Can a cheap property predict it? Motivated by the saturation and high-temperature limits of Section 3, we tested the breadth of intermediate entropy as a heuristic predictor: a model whose next-token distribution jumps from sharp to flat over a small range of temperature might leave little room for chaos. The breadth of this intermediate regime, the decades of $\log_{10} T$ in which the mean entropy on the orbit lies between 1 and 10 nats, ranges from 0.014 (Qwen2.5-72B) to 0.69 (OLMo-2-32B). It orders the fifteen models' chaotic length (Spearman +0.55) somewhat better than their parameter count does (−0.44; Figure 57, right), but it misses the chaos of Qwen2.5-1.5B and -14B at nearly uniform distributions, above its 10-nat bound. This analysis was adapted after the OLMo-2-32B result, and the breadth is measured on the loop's own orbits, which chaos itself spreads, so the test is partly circular. A loop-free version, frozen before it was run (`docs/start-state-predictor.md`), uses one forward pass per model and prompt: the entropy profile of the logits at the loop's start state, and the one-step gain there. Its primary prediction fails: the start-state breadth on the paper's prompt orders the fourteen dense models' chaotic length with Spearman +0.30, below the frozen threshold of 0.5. The secondary criterion, its median over ten prompts, narrowly passes on the frozen outcome (+0.52; $p = 0.06$) and weakens on the expanded sweeps (+0.44), and the

one-step gain at the start state is negative in every model, the most chaotic ones included (-0.5 to -3.4 per step), and orders nothing (-0.32). These features ran on two GPUs before the precision fix, which changes them only at about the 10^{-7} level. We have not established a reliable predictor from a single forward pass: as for the period-8 orbit of Section 3, it has to be measured on the loop.

Steering at 70B. Under steering, the 70B maps show chaos, rarely and not independently of the direction (Figure 92; labels at the single offset 10^{-6}, which the plateau test above supports for this computation). On Llama-3.3-70B's hostility plane, at the layer of Goodfire's 70B SAE and at equal norm, 4.3 % of cells are labeled chaotic along the concept vector, 3.5 % along the frozen-rule hostility feature, 1.1 % along a Gaussian direction and 0.2 % along a random live feature, against 28, 29, 39 and 27 % for the same four directions in Llama-3.1-8B (on one GPU); Qwen2.5-72B along hostility gives 0.8 %: eight of its 1024 cells, all at $\log_{10} T = 0.26$–0.30. Recomputed in float64 throughout, six of the eight are chaotic, with automatic-differentiation tangent exponents of 0.23 to 0.36 that the twins at all four offsets from 10^{-4} to 10^{-7} reproduce within 0.01; five have nearly uniform next-token distributions (11.0 to 11.1 nats) and one alternates between sharp and flat behind a period-2 leading token. The same temperatures without steering are period-2 orbits with exponents between -3.8 and -3.6, and the steered cells above $\log_{10} T = 0.325$ are near-uniform fixed points, so steering drives Qwen2.5-72B into chaos too, in a narrow band. The chaotic cells of the hostility directions lie on a ridge that steering moves from $\log_{10} T \approx 0.3$ to 0.1. A 128×128 zoom of the concept plane at the same offset, computed after the precision fix in sixteen bands with identical vectors and joined (Figure 59; $\alpha \in [0.06, 0.12]$, $\log_{10} T \in [0.12, 0.27]$), resolves it: a chaotic band (3410 of the 16,384 cells, 20.8 %; largest exponent 0.64), threaded by periodic windows, whose center steering moves from $\log_{10} T \approx 0.23$ at $\alpha = 0.06$ to ≈ 0.19 at $\alpha = 0.12$; fixed points lie above it, and windows of periods 3, 5, 6 and 7, up to 13, fill the plane below. Only 1.7 % of the cells have neither a detected period nor a positive exponent. A plateau test of 24 cells drawn by a rule fixed in advance (12 chaotic, 6 periodic, 6 unresolved; 300 transient and 300 measured steps, absolute twin offsets 10^{-4}, 10^{-5} and 10^{-6} and an automatic-differentiation tangent) confirms 11 of the 12 chaotic labels, whose four exponents agree within 0.012, and 5 of the 6 periodic ones. The two exceptions are neighboring cells near $\log_{10} T = 0.20$ where a long, weakly attracting loop (periods 15 and 16) and chaos are reached from the same start depending on rounding: one cell labeled chaotic reached the loop, and one labeled periodic reached chaos. The six unresolved cells have small negative exponents (-0.005 to -0.03) and no period up to 32. An earlier 64×64 zoom of the same window, computed before the fix, had 44 % such cells and, above the band, 515 cells with periods from 3 to 32 but a constant leading token, an artifact of the period test at that computation's noise floor; Figure 59 supersedes it.

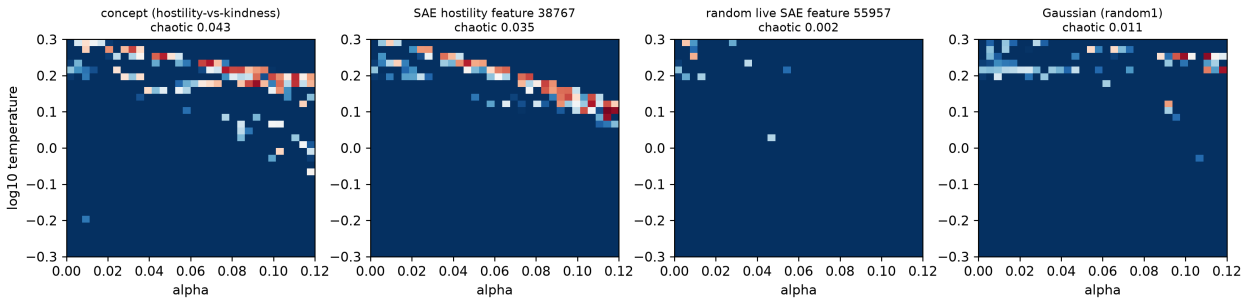

Figure 92: Steering controls at 70B: Llama-3.3-70B-Instruct soft-token maps (32^2 float64 cells, 300 transient and 300 measured steps, maximum period 32) steered after zero-based block 50, the hook point of Goodfire's 70B sparse autoencoder, along the hostility concept vector, the autoencoder's hostility feature selected by a rule frozen before any map, a random live feature and a Gaussian direction, all at the same norm. Chaotic cells: 44, 36, 2 and 11 of 1024 (concept, feature, random feature, Gaussian); cells with neither a period nor a positive exponent: 468, 358, 223 and 424.

J Unsteered chaos and steering controls

Steering adds a fixed vector to the residual stream, so chaos along a steering axis could belong to the push rather than to the model. At $\alpha = 0$, however, the map is the unsteered model's own feedback map, with temperature as its only parameter. The $\alpha = 0$ columns of the hostility maps contain cells with positive finite-horizon exponents in Llama-3.1 8B (26 of 64 temperatures), Qwen2.5-7B (15 of 64) and OLMo-2 7B (11 of 64), while the near-zero columns of the 0.5B and 1.5B maps contain one such cell each (of 128 and 64). A longer float64 sweep of the unsteered 0.5B map (512 temperatures over $\log_{10} T \in [-1, 0.3]$, 1000 transient and 2000 measured steps) finds a narrow chaotic band at $\log_{10} T \approx 0.15$–0.21 (7 temperatures, exponents up to 0.28), with a period-5 loop at 437 of the 512 temperatures. The same long sweep finds 96 chaotic temperatures in Llama-3.1 8B, between $\log_{10} T = -0.20$ and 0.24 (exponents up to 0.66), interleaved with periodic windows of periods 2 to 52; within $\log_{10} T \in [-0.3, 0.3]$ its chaotic share, 41 %, equals that of the coarse column. Qwen2.5-7B has 65 chaotic temperatures between $\log_{10} T = -0.07$ and 0.28 (exponents up to 0.74), and OLMo-2 7B has 32 between 0.04 and 0.14 (up to 0.30). Steering is therefore not needed to reach chaotic regimes in any of these four models; where it does create chaos, the growth it adds along recorded orbits comes mainly from the readout, where next tokens compete (Section 3). Temperature alone produces chaos with embedded windows, broadly in the 7B and 8B models and in a narrow band in the 0.5B model (Figure 93). Across the fifteen models of Section 14, the regular-grid sweeps find unsteered chaos in eleven of the fourteen dense models (in Llama-3.1-70B at two temperatures that pass a three-offset test), but not in Qwen2.5-72B, Llama-3.3-70B, OLMo-2-13B or the mixture of experts; finer lines found candidates in the first two, and the Qwen2.5-72B window, recomputed in float64 throughout, is periodic.

Steering does widen the chaotic regions, but not as a property of the concept: along norm-matched random directions, the 0.5B maps are as chaotic as along the concept directions (Figure 94). On the same 128^2 grid, 7.6, 8.3 and 12.5 % of cells are chaotic for three Gaussian directions with the philosophy vector's norm, against 6.8 % for philosophy and 7.5 % for gratitude. In all five maps the chaotic share rises from about 1 % at $|\alpha| < 0.01$ to between 5 and 22 % at larger $|\alpha|$, and a chaotic region at high temperature, with windows reaching down, recurs. In this model, chaos under steering is therefore a generic response to perturbing the residual stream at this norm. Rotating the steering direction makes the same point continuously: along the great circle from one Gaussian direction to

another, 90° away, at constant norm, the chaotic regions move and change shape, but the chaotic share stays between 13 and 24 % at every one of 120 angles (Figure 9; a video of the rotation is in the supplementary material). Single neurons behave the same way. Steering along the output direction (Figure 95) of one MLP neuron of the steered layer (its column of the down-projection, rescaled to the same norm; adding it is the same as raising that one neuron's activation by a fixed amount for every token of the input: by 37.4 at $|\alpha| = 0.12$ for the neuron most specific to philosophy, whose mean activation on the philosophy texts is 0.27) gives 9.9 % chaotic cells for the neuron most specific to philosophy on held-out texts, chosen like the autoencoder feature below by a rule fixed before any map, 5.1 % for its index neighbor and 12.9 % for a random neuron (Figure 96); near $\alpha = 0$ the share is 0.4 to 1.2 %. The direction does select the texts. The largest attractor family in all five maps is the unsteered period-5 loop, but beyond it only 26 to 56 % of one map's cells carry an attractor text that appears anywhere in another map. Llama-3.1 8B agrees. On its hostility plane, two norm-matched random directions at the same layer give 26 and 17 % chaotic cells against 28 % along the concept direction; all three carry 34 to 44 % near $\alpha = 0$, and the random directions thin out at large $|\alpha|$. At layer 19, where a sparse autoencoder trained on this model reads the residual stream [139], the concept vector rebuilt there gives 28 %. The decoder direction of the autoencoder feature most specific to hostility on held-out texts (chosen by a rule fixed before any map; cosine 0.10 with the concept vector) gives 29 %, a random live feature 27 % and a Gaussian direction 39 %, all at the same norm. In this model, chaos therefore occurs along autoencoder-feature, contrastive and random directions alike, with direction-dependent prevalence; at 70B the directions differ more (Appendix I).

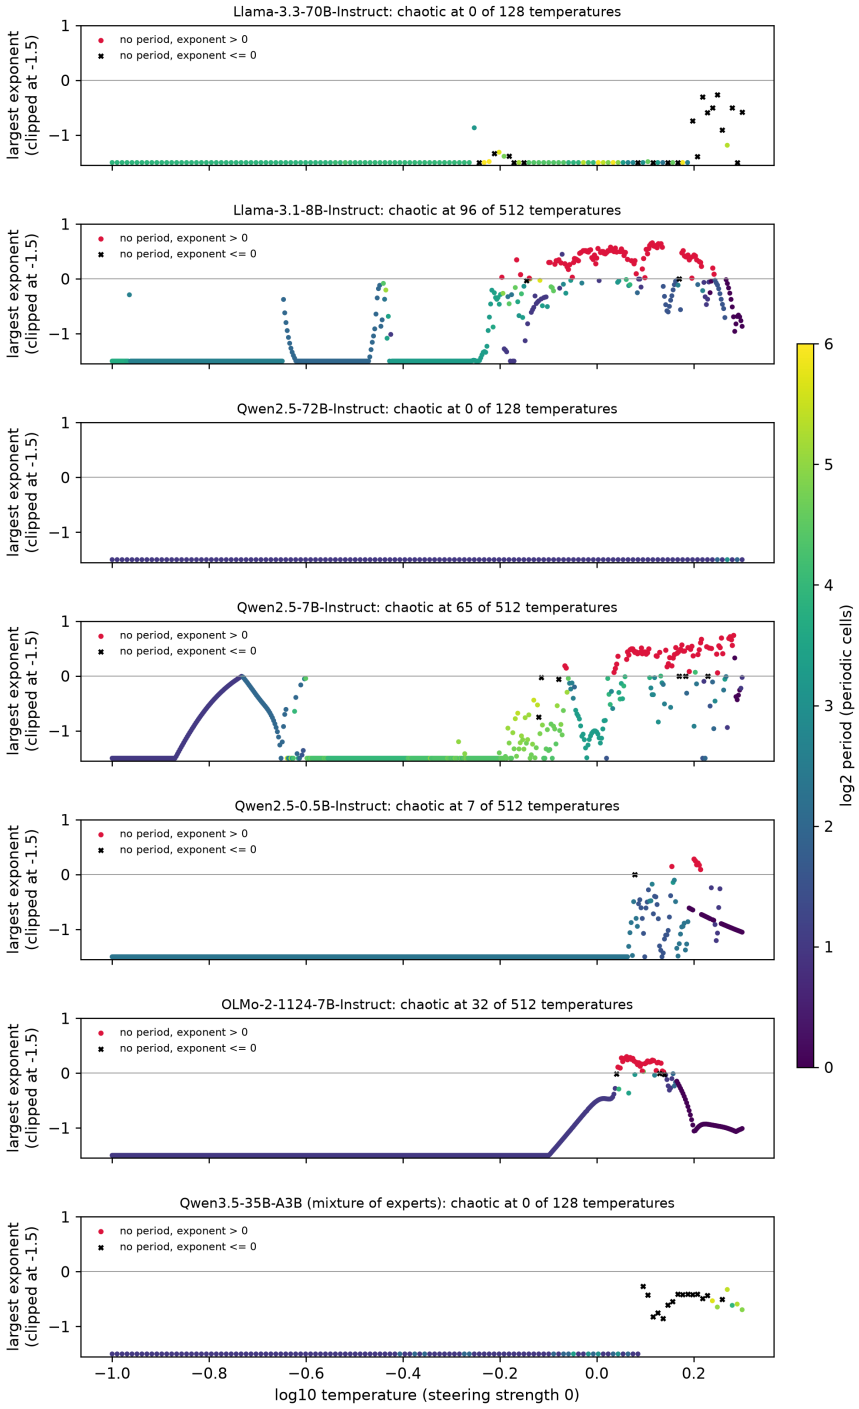

Figure 93: The unsteered soft-token map ($\alpha = 0$) along temperature, top to bottom: Llama-3.3-70B-Instruct, Llama-3.1-8B-Instruct, Qwen2.5-72B-Instruct, Qwen2.5-7B-Instruct, Qwen2.5-0.5B-Instruct, OLMo-2-1124-7B-Instruct and the mixture of experts Qwen3.5-35B-A3B. Largest Lyapunov exponent (clipped at -1.5) from float64 sweeps of 512 temperatures (128 for the 70B-class and mixture-of-experts models; 1000 transient and 2000 measured steps). Red: no period up to 64 and an exponent above 10^{-3}; black crosses: neither a detected period nor a positive exponent (unresolved); other points have a detected period, colored by $\log_2$ of the period. Table 22 summarizes these and the wider sweeps of the other models.

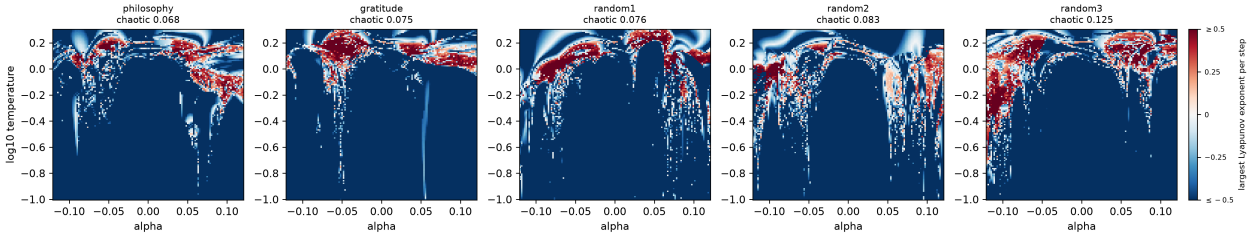

Figure 94: Steering controls in Qwen2.5-0.5B-Instruct (128^2 float64 cells, steering at layer 12, one grid and one set of settings): largest Lyapunov exponent along the philosophy and gratitude directions and along three Gaussian directions rescaled to the philosophy vector's norm. Titles give the chaotic share (finite-horizon exponent above 10^{-3}, no period up to 48); color gives the exponent, clipped at ± 0.5.

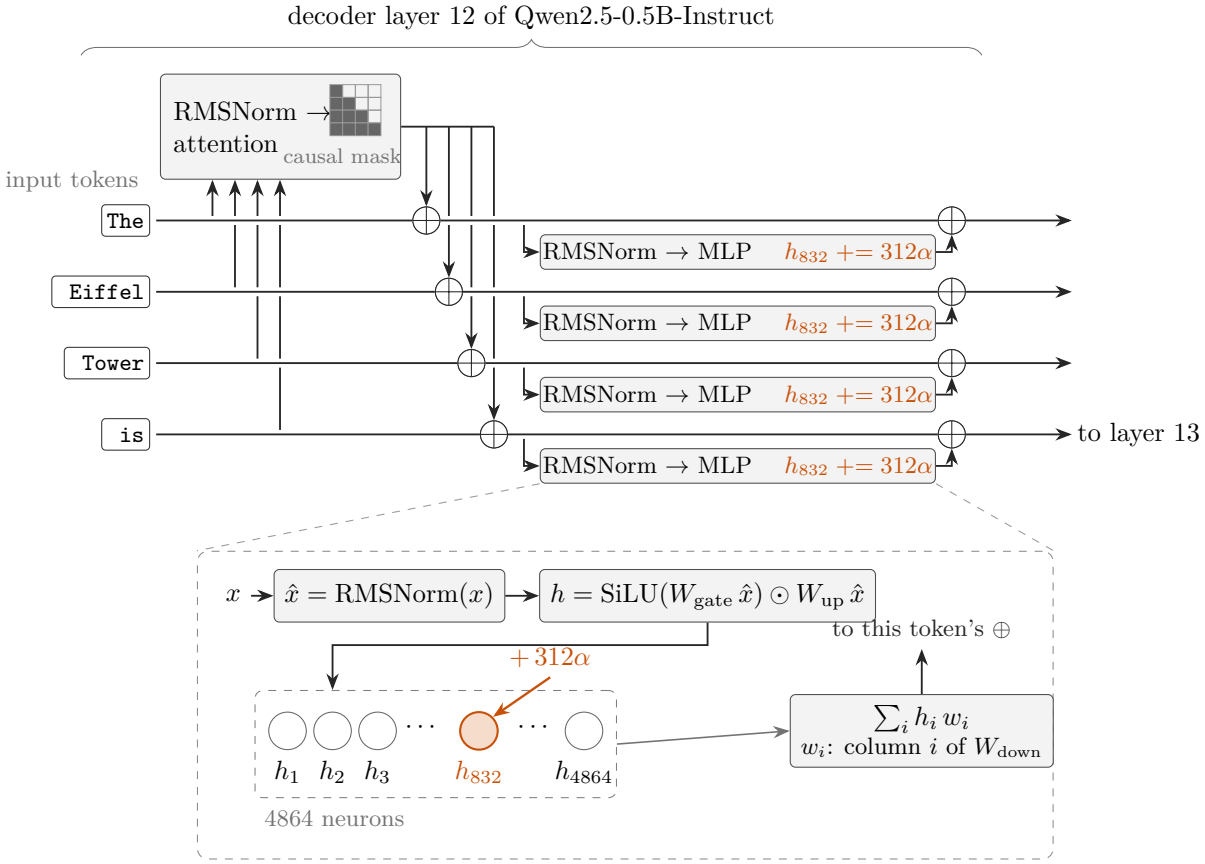

Figure 95: Steering along one MLP neuron, in decoder layer 12 of Qwen2.5-0.5B-Instruct. Each token has its own residual stream, one 896-dimensional vector per token, and its own pass through the same MLP; attention, with its causal mask, is the only step that mixes tokens. Inside every MLP box (bottom, for one token with residual vector x), the 4864 neurons write $\sum_i h_i w_i$ into the stream, neuron i along its own column w_i of the down-projection. The steering strength α is a scalar we set (negative values lower the activation), and $312 = \|v_{\mathrm{phil}}\|/\|w_{832}\| = 188.99/0.606$ is a fixed scale for neuron 832. Raising h_{832} by 312α for every token is the same as adding αv, with $v = 312\, w_{832}$, to the layer's output, which is how we implement it: $\sum_i h_i w_i + 312\alpha\, w_{832} = \sum_{i \neq 832} h_i w_i + (h_{832} + 312\alpha)\, w_{832}$. At $|\alpha| = 0.12$ the increase is 37.4, against a mean activation of 0.27 on the philosophy texts; for neurons 833 and 4137 the scale is 329 and 346. The tokens shown are only an example: in the loop, the input is the fixed prompt followed by the fed-back soft token.

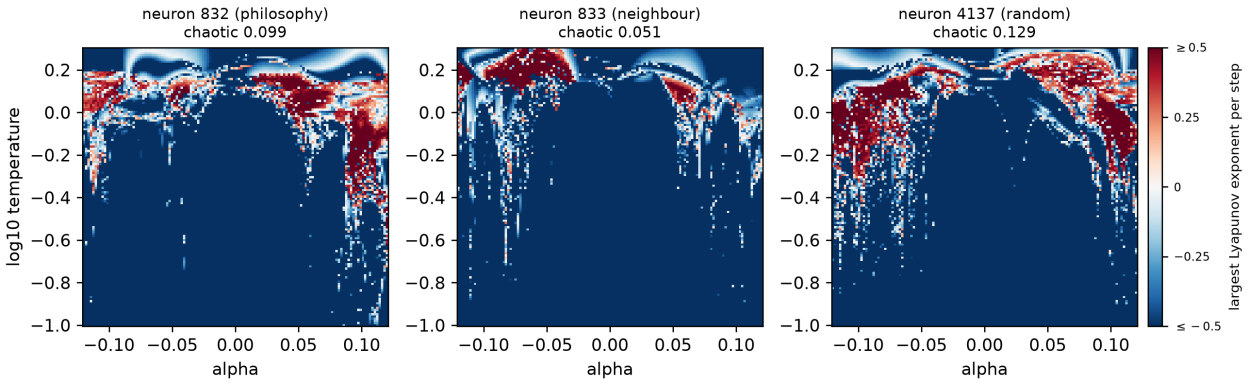

Figure 96: Steering along single MLP neurons in Qwen2.5-0.5B-Instruct (same grid, settings and color scale as Figure 94): the output directions of the layer-12 neuron most specific to philosophy on held-out texts (832), of its index neighbor (833) and of a random neuron (4137), each rescaled to the philosophy vector's norm (cosines with it 0.013, −0.028 and 0.022). Titles give the chaotic share.

The prompt. Every 0.5B result above uses one prompt, so we mapped (run PV, `tools/pv_summary.py`) the standard plane (α from -0.12 to 0.12, $\log_{10} T$ from -1 to 0.3; 128×128 float64 cells, philosophy direction at layer 12) under 22 variants (Figure 97): each of its two parts alone, and neither (Qwen's default system prompt, which the chat template inserts when none is given); the system prompt's words in two random orders, the user message's two tokens reversed, and both changes together; an opposite system prompt; happy and sad versions of the system prompt and of the user message; the same texts without chat markup, a system prompt three times as long and a two-turn conversation; three prompts that are a single random vector, with no text at all (a Gaussian direction at the median embedding norm, the embedding of a random token, and a Dirichlet(0.1) mixture of the embeddings of 64 random tokens); and the base prompt with the loop started from three random vectors. Every prompt we tested gives chaos with periodic windows: 5.9 to 14.9 % of the plane is chaotic (7.0 % for the base prompt), with largest exponents of 0.62 to 0.85 per step (here chaotic means a largest exponent above 10^{-3} alone; this finite-window indicator overlaps a detected period at 3 to 18 cells per map). Removing the user message doubles the chaotic share, and even a prompt that is one random vector gives 8 to 10 %. In every prompt we tested, chaos and periodic windows appear, but the prompt decides where they lie: the chaotic cells of each of the other 18 prompts overlap those of the base prompt by 1 to 37 % (Jaccard index). For the three random starts, the chaos/non-chaos threshold labels agree with the base run at 98.9 to 99.0 % of the cells, while the periods detected within the run agree at only 86.5 to 93.1 %, so within this horizon the orbit reached can depend on the start. Across the 19 prompts there is nonetheless a common structure (Figure 98): 48.3 % of the cells are chaotic for at least one prompt but only 3.0 % for half of them or more, and these hot spots lie at moderate steering of either sign in a band of temperature, where the mean fraction of chaotic cells per row peaks at 38 % near $\log_{10} T = 0.15$, against 0.8 % below $\log_{10} T = -0.5$. The three random-vector prompts move the band to colder temperatures: their mean profile peaks near $\log_{10} T = 0.03$, that of the text prompts near 0.16.

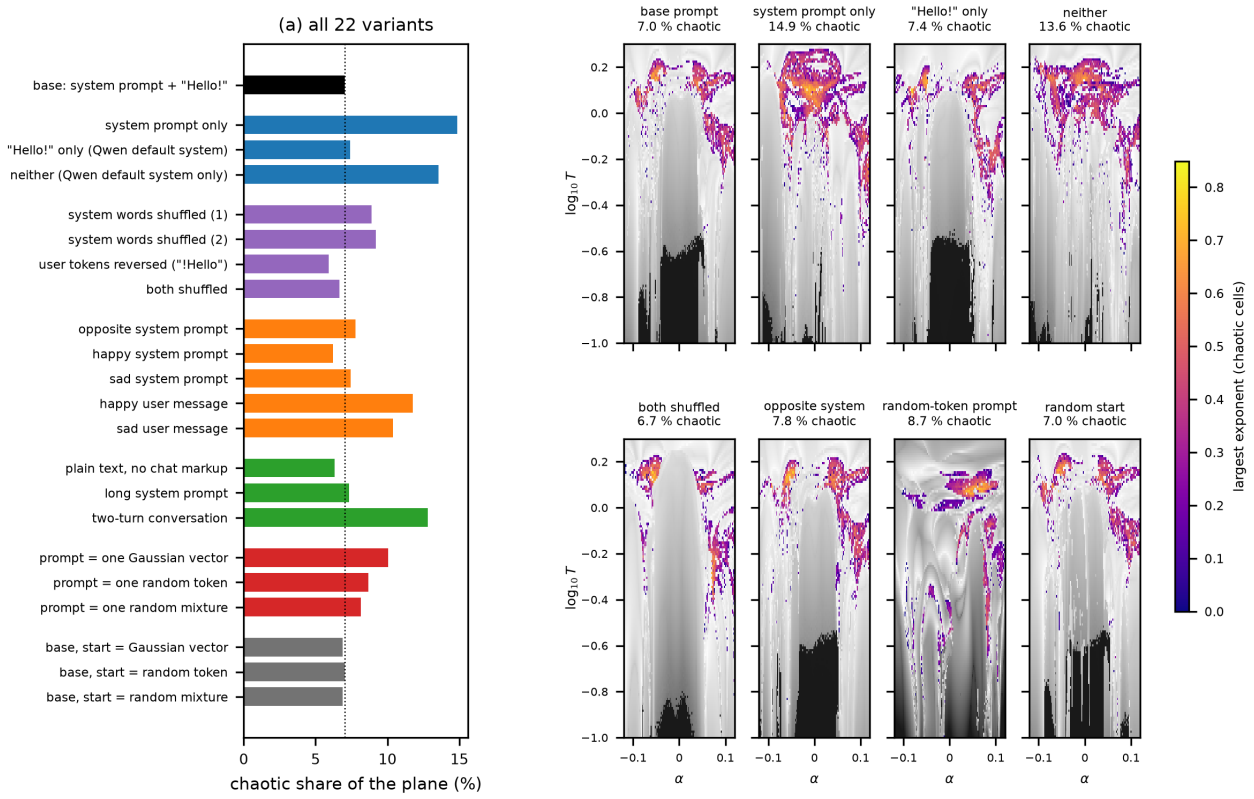

Figure 97: Every prompt we tested gives chaos, at places that depend on the prompt (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, float64; 128×128 cells over $\alpha \in [-0.12, 0.12]$ and $\log_{10} T \in [-1, 0.3]$). (a) The chaotic share of the plane for each of the 22 variants (19 prefixes, and the base prefix with the loop started from three random vectors), grouped by the kind of change; dotted: the base prompt (the system prompt "You are an AI talking with another AI…" and the user message "Hello!"). (b) Eight of the maps (colors: the largest exponent where it exceeds 10^{-3}; gray: the other cells, darker where more contracting, among them a few unresolved nonperiodic cells). The exact texts and token identities of the prompts are in the project repository (`docs/prompt-variants`).

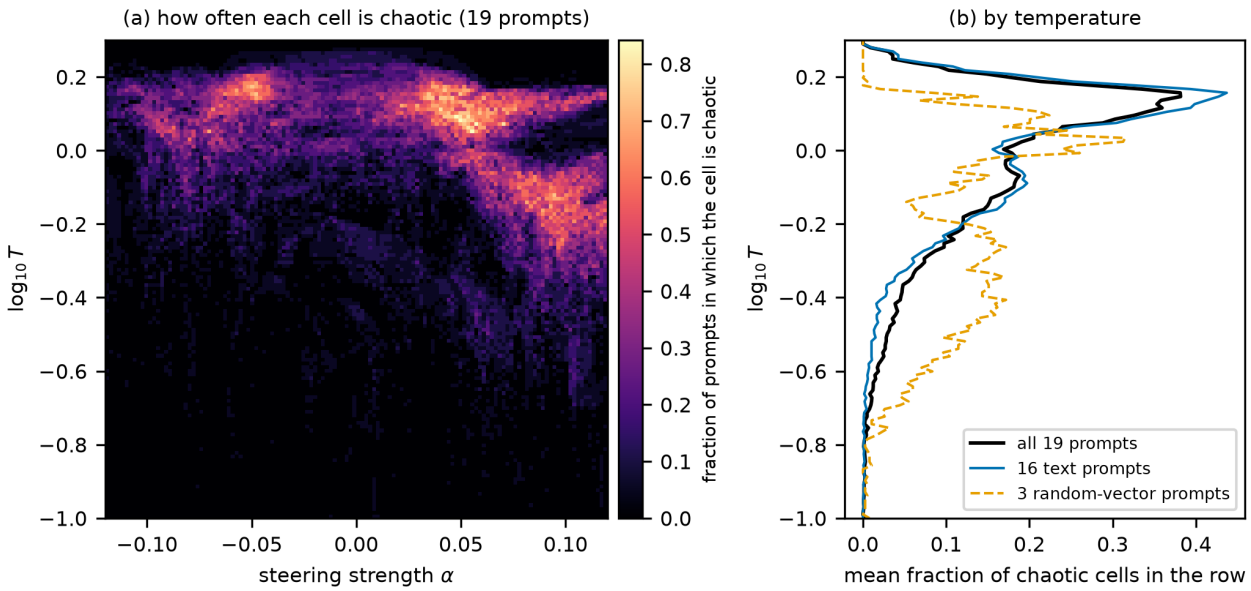

Figure 98: Where the chaos is found across prompts. (a) For each cell of the plane of Figure 97, the fraction of the 19 prompt variants (the random-start runs left out) in which it is chaotic. (b) The mean fraction per temperature row, for all prompts, the text prompts and the three random-vector prompts.

K Limitations in detail

The detailed orbit studies cover one model family (Qwen2.5, 0.5B and 1.5B), with coarse maps and three float64 orbits at 7B; the two other families (OLMo-2 7B and Llama-3.1 8B) have maps, including 128^2 refinements for Llama, but no continued orbits. The size comparison of Section 14 uses regular temperature grids that can miss narrow windows, a window that missed most of the larger models' chaos, and mostly one twin offset per label. The two-GPU runs of the 70B-class models computed their final normalization in float32 until we fixed it; their labels rest on offsets of 10^{-6} or more, and the cells we rechecked, the Qwen2.5-72B jump window and steered candidates in float64 throughout and samples of the Llama maps at three offsets, behave as stated, but the rest of those maps were not recomputed. Neither were the four coarse lines of the frozen 72B test; after the fix we computed only fine lines through their entropy jumps, one of which passes through their single chaotic temperature. The maps use one prompt, apart from the 22 prompt variants of the 0.5B standard plane (paragraph "The prompt" above; Section 14 adds ten for Qwen2.5-7B and -14B, and four for Qwen2.5-72B, its base model, and OLMo-2-7B and -32B), and steering uses contrastive vectors, checked against norm-matched random and sparse-autoencoder directions in two models. The soft-token loop is one inference scheme: Appendix G extends it to memories of up to sixteen soft tokens, to the growing context of Soft Thinking and, in an exploratory run, to feeding back only the most probable tokens, but sampled decoding of discrete tokens is a different system. The evidence consists of numerical validation, cross-checks and prospectively frozen passage tests (including failed all-row criteria), with computer-assisted proofs for only four orbits at three parameter points (the period-8 orbit, a coexisting period-2 and period-3 pair, and the 3-cycle at the start of the cascade). Chaos labels are finite-horizon estimates (300 measured steps after 300 to 500 transient steps), and the chaotic cells are numerically validated (twin exponents agree with automatic-differentiation tangents, across precisions and horizons) but not yet established by a computer-assisted proof. A computer-assisted proof of chaos, such as a horseshoe verified by covering relations, would need enclosures of the map over boxes many orders of magnitude larger than the

period-8 proof's (radius 2^{-640}). That proof is unaffected, but its ball arithmetic does not scale to large boxes: in a separate test, a box of radius 2^{-80} already gave enclosures too wide to show that an RMS normalization at layer 3 divides by a positive number (the enclosures overestimate through the layers), so a proof of chaos needs other verified methods. A horizon check retested random cells of two float64 maps (Figures 1 and 10) with 1000 transient and 2000 measured steps at two twin offsets: 86 and 92 of 96 cells labeled chaotic stayed chaotic (all but one of the rest settled on periodic orbits), all 96 periodic cells in each map stayed periodic, and 29 and 27 of 32 cells with neither label settled on periodic orbits (18 and 20 of them with periods up to 32, the maps' maximum, the others with periods up to 64). For the cells that stayed chaotic, the map's 300-step exponent differs from the 2000-step exponent by a median of 0.020 and 0.030 (90th percentiles 0.064 and 0.096, maxima 0.23 and 0.30), and the 2000-step exponents at offsets 10^{-6} and 10^{-9} agree within 3×10^{-3}. The chaotic counts do not hinge on the threshold: raising it from 10^{-3} to 0.01, 0.02 or 0.05 leaves 77,646, 76,761 or 73,437 of the 78,868 chaotic cells of Figure 1. Near the accumulation of a cascade the labels are less reliable, because orbits there converge slowly or wander chaotically before they settle (transient chaos [48, 51]): on the hostile line of Figure 30 (left), 135 of the 235 cells labeled chaotic from 400 measured steps stayed chaotic and 91 were periodic. The Feigenbaum ratios are finite-order values that rise toward δ (the last is within 1 %); the limit itself is not computed. All orbit, bifurcation and chaos results concern the soft-token map, not sampled generation with a growing context. With sampling, the bridge of Appendix G tests whether the bottleneck's laminar phases survive the noise and, in two exploratory runs on one 0.5B model, how sampled loops and ordinary sampled generation respond to small perturbations; its other tests are deterministic (greedy decoding, longer fed-back memories, Soft Thinking's growing context and feeding back only the most probable tokens).

L Evidence and reproducibility

We grade claims on an evidence ladder. Rung 1 means numerically validated (tolerances, residuals, convergence, precision varied); rung 2, cross-checked by an independent estimate, implementation or hardware; rung 3, a computer-assisted proof.

Search scale. The soft-token results come from an exploratory search. An initial census of the raw data at 00:19 UTC on 14 October 2026 counted 34 distinct maps (987 308 cells, 27 maps in float64) in six models, each with one prompt and one of three steering directions or none, plus one partial map. Of these, 29 maps and the partial one appear here; the five omitted are early discovery or side maps. Later additions are not in that count: the steering-control maps of Appendix J, the mixture-of-experts maps, and the sweeps, lines and maps of the models of Section 14 and Appendix I. The labeled maps contain about 2600 periodic families of at least six cells, counted per map. We singled out about 26 windows or families for float64 work: 20 are reported in detail, 2 only as map features and 4 not at all (an orbit whose continuation was abandoned, a cell that proved aperiodic, a band whose float32 orbit did not converge and a pilot that found no cusp within a bounded search). Of 54 Newton orbit attempts, 7 failed. All nine tangent reversals on continued fold curves (seven found by 14 October, and two on traces of a period-4 fold curve continued on 15 October) are numerically located cusps; the two reversals of a continued flip curve are smooth turns. The four reversals on the two fold curves of the period-5 window found on 15 October (below) are located cusps as well. Four further reversals found on 18 October, one on the lower-left leg's own fold curve, one on the extension of the window's second fold curve and two on a partial trace of a fold curve of the window at $\log_{10} T = -0.105$, are flagged by the continuation but not

located. Two of the five regions we flagged as shrimp candidates are reported as shrimp-like. On 15 October, the turning-point search of the two-block planes (not counted in Table 32) led to two more windows: a period-5 window, reported with its continued orbit and its continued fold and flip curves in Appendix F, and a period-6 window that was mapped but not continued. One further language model (Ouro-1.4B) was explored and is not reported; its apparent chaos was a long transient. Every window was chosen after the maps were inspected. In the census of Table 32, only the passage forecasts, the cusp-unfolding offsets and Huginn's relaxation test were fixed before the computations that test them; tests frozen later, up to the fixed-shape control of 18 October, are identified as such where they are reported. The child-window memory intervention of 21 October is an additional exploratory study, outside that census. Its 273-point local extension, source commit `cf2726b3a61bce0e933ba71b93ae992d529977f0`, completed seven new saddles after fresh replay of two prior centers; the immutable parent supplied eleven accepted centers and a failed trust-region step. Independent analysis at `7a885cc641f5ce13115a3dd0ec755b3704dd5d61` checked the source and plan bindings, every raw array, saved endpoint closure, primitive-period and cyclic gates, finite-difference stencils, root iterations, full-orbit alignments and determinant gradients. The compact receipt in Section 11.1 binds the full audits by SHA-256. These are saved-data and numerical-precision checks; no interval inclusion proves the saddle, exact crossing or full-spectrum isolation. The subsequent BSMZ1 scan compares 16,384 fresh fixed-opener cells at the retained-feedback candidate with the original BSZ2 close-up. Its independent audit binds the unchanged execution source, every saved array, the effective float64 checks, the root-state and reference-image evidence, and the matched figure (`docs/baby-shrimp-zoom/memory-junction-scan-receipt.json`). Its recurrence labels remain finite-horizon candidates rather than primitive-orbit verification at every pixel.

	explored	reported	note
Soft-token maps (distinct)	34	29	plus one partial map, reported
cells	987 308	893 868	omitted: 5 discovery or side maps[a]
float32 maps	7	4	
Models with maps	6	6	one prompt; 3 directions and $\alpha = 0$
Periodic families (≥ 6 cells)	≈ 2600	–	counted per labeled map
Windows given float64 work	≈ 26	20 (+2)	+2: map features only
Newton orbit attempts	54	–	7 failed
Tangent reversals on fold curves	9	9	all located cusps (2 found on 15 Oct.)
Unfolding offsets, frozen	12	12	counts 12/12, stability 11/12
Shrimp-candidate regions	5	2	reported as shrimp-like
Tests fixed before outcomes	6	4	2 all-row criteria failed by one row
Language models loaded	8	7	Ouro-1.4B not reported

Table 32: Search scale on 14 October 2026, before the steering-control maps; the reversal row includes the two cusps located on 15 October (Huginn's frozen test counted once; of the two tests not reported here, one belongs to the separately reported sampled-dialogue study and one, a cusp pilot in a second checkpoint, is the unreported bounded search). [a]The five omitted maps are a high-temperature family map, a 256-cell scan and three float64 zooms. The census saved only its totals; this list was reconstructed on 17 October from the raw maps as the set whose omission reproduces both totals (`docs/receipts/paper-numbers-20261017.json`).

- Rung 3: the period-8 stripe orbit of Table 8. At $\alpha = 0.071674$ and $\log_{10} T = -0.115$, read as exact binary64 rationals (the temperature is the real number $\exp(\ln 10 \cdot \log_{10} T)$ for that exact

dyadic $\log_{10} T$), the expected-embedding map of the stored Qwen2.5-0.5B-Instruct checkpoint (all 24 layers, the full 151 936-token readout, the fixed 37-token prefix, steering after zero-based block 12) has a primitive period-8 orbit, which is the unique multiple-shooting root in a product of eight boxes of coordinate radius 2^{-640} around refined centers, and the orbit is locally attracting. Ball arithmetic (Arb [143]) encloses the map and its Jacobian over each box (8×896 directions); a Krawczyk-type inclusion [144–146] for the preconditioned shooting system gives existence and uniqueness, separation of the first point from the points one, two and four steps later gives the primitive period, and a norm bound on powers of the enclosed monodromy gives attraction. The map is the real-number map of the stored constants, read as exact dyadic rationals: rotary angles, exponentials and square roots are real functions enclosed by the balls, not floating-point tables. An independent audit in exact rational arithmetic rechecked the saved witnesses (the scalar inequalities, the primitive-period separations and the row norms of the saved monodromy powers); it did not recompute the enclosures of the model or the dense matrix products, which rest on one execution. The proof trusts the Python, FLINT and Arb stack and the stored weights. It does not exclude other period-8 points in the first box, and it says nothing about the window's boundaries or global organization. The record (`certificates/llm-p8-certificate-final-001.json`) pins the SHA-256 hashes of the checkpoint, its configuration and the steering vector, and lists the 37 prefix token IDs.

- Rung 3: coexistence. At $\alpha = 0.06468$ and $\log_{10} T = -0.148$, inside the strip of the crescent where the 2-cycle and the 3-cycle both exist (Section 12), the same map, with the same reading of parameters and constants, has a primitive period-2 orbit and a primitive period-3 orbit. Each is the unique multiple-shooting root in its product of boxes of coordinate radius 2^{-640} around refined centers, and each is locally attracting; different primitive periods make them distinct orbits. The construction is that of the period-8 proof. The point enclosures are computed at 2048 bits and the derivative enclosures at 1024 bits, in 70 pieces that cover every phase and all 896 directions (28 for the 2-cycle, 42 for the 3-cycle; 37.5 worker-hours on rented CPUs). Krawczyk-type inclusion gives existence and local uniqueness, and separation from the proper divisors gives the primitive periods. For attraction, the unweighted bounds on the first and second powers of the enclosed monodromy failed the sufficient condition: the monodromy is strongly nonnormal, so a single norm overstates its growth. These failed attempts are kept in the record, and the fourth and eighth powers pass. Their induced infinity-norm bounds, rounded up, are 7.51×10^{-4} and 1.9×10^{-8} for the 2-cycle and 0.342 and 1.091×10^{-3} for the 3-cycle; the numerical leading multiplier moduli are 0.0476 and 0.239. An independent exact-rational audit rechecked every saved row norm (1 605 632 entries per orbit) and the scalar and separation witnesses. A separate check certifies the token loops: over every box of both orbits, the most probable next token leads all 151 936 tokens by a positive margin, so the 2-cycle writes " is it" and the 3-cycle a blank line, "The" and " existence" (`codex/coexistence-task33/token-labels-release.json`); the same check gives the period-8 orbit's loop " is there is a is a. This". These are the readouts of the feedback orbits, not sentences generated autoregressively. A further certificate gives explicit basin neighborhoods: every state in the closed infinity-norm ball of radius 2^{-681} about the 2-cycle's frozen center at phase 0 converges to the 2-cycle, and every state within 2^{-700} of the 3-cycle's center converges to the 3-cycle (`codex/coexistence-task33/inner-basin-release.json`). These neighborhoods are tiny. The proofs say nothing about which orbit a start outside them reaches, including the fixed start of the maps, nor about the trapping of the whole box or the window's global organization. The record is `codex/coexistence-task33/pair-certificate-release.json`.

- Rung 3: the 3-cycle at the start of the cascade. At $\alpha = 0.0650$ and $\log_{10} T = -0.14798$, on the cut of Table 16 and 4.8×10^{-4} below its first flip, the same map has a primitive period-3 orbit: it is the unique multiple-shooting root in its product of three boxes of coordinate radius 2^{-640} around refined centers, and it is locally attracting. The construction is that of the period-8 proof, with point enclosures at 2048 bits and derivative enclosures at 1024 bits, in 42 pieces that cover every phase and all 896 directions. The bounds on the first, second and fourth powers of the enclosed monodromy failed the sufficient condition for attraction and stay in the record; the eighth power passes, with an induced infinity-norm bound of at most 0.447 (for a 3-cycle, a bound on the derivative over 24 steps). An independent exact-rational audit rechecked the scalar and separation witnesses and every saved row norm (3 211 264 entries); it did not recompute the enclosures of the model or the dense products. The proof does not certify the orbit's token loop, which starts reach it, or the later orbits of the cascade: a certificate of the 6-cycle on the same cut was begun and not completed. The record is `codex/certificate-ladder-task34/cascade-p3-certificate-release.json`.

- Rung 2: the 2-cycle's fold and its normal form. An independent reimplementation of the forward map and solver recovers α_c, a and b (to 6×10^{-15}, 3×10^{-12} and 2×10^{-6} relatively), and ab also follows from the stable-branch multipliers. The same reimplementation reproduces all 18 fixed-entrance counts of the controlled passage runs exactly (Table 10 lists 12 of them).

- Prospectively frozen: the passage forecasts at two new folds (Table 13); 23 of 24 within one iterate (16 of the 17 longer than five iterates), and the all-row criterion failed by one row. A second frozen test at two more folds (Table 14): 18 of 24 within one iterate of the primary forecast; it failed at the entrance passages of a period-6 fold, short by 3.1 to 5.2 %.

- Rung 2 as well: the period-2 and period-8 orbits of Table 8 (independent implementation), and the cascade's first four flip parameters (CPU and H100, agreement 4×10^{-15}).

- Rung 2 as well: the hostile 2-cycle fold and its normal form, reproduced by an independent float64 forward map with nested automatic differentiation (location to 6×10^{-10}, ab to 8 parts per million); its controlled passages were also checked independently, with the frozen all-six criterion failing by one row.

- Rung 1: the other orbits, the curves, the thirteen located cusps, the superstable parameters, the 1.5B flips, the hostile windows' edges, the spiral region's period-7 body and the cascade inside it, the three 7B orbits, and Huginn's float32 spectral estimates.

- Prospectively frozen, not independently reproduced: the scaling predictions P1–P3, the prompt-blocked size tests in Qwen2.5 (7B and 14B; 72B; 72B base against Instruct) and OLMo-2, and the loop-free predictor test of Section 14 and Appendix I, all scored by code pushed before the data, together with the three-offset plateau tests of the 70B-class maps and the float64 recheck of the Qwen2.5-72B cells with automatic-differentiation tangent exponents. These are float64 computations at rung 1; no independent implementation has rechecked them.

- Prospectively frozen on 15 October: the answer-level reproducibility test in Qwen2.5-0.5B (positive), whose saved rows and statistics an independent implementation reproduced (the generation itself was not rerun); its 7B extension (positive) and the behavior test (null on its primary question; no edge optimum), not independently reproduced. All are reported in Section 6, the behavior test in detail in Appendix D, and each was scored by code pushed before the data (`docs/repro-task-test.md`, `docs/bpc-test.md`).

- Prospectively frozen on 18 October, not independently reproduced: the fixed-shape control and dose test in Qwen2.5-0.5B (448 cells; `docs/fs-design.md`; Section 6). With unchanged matrix shapes the paired runs are bitwise identical in every cell, a changed shape reproduces the earlier comparison, and in the chaotic cells the primary dose criterion holds at every dose from 10^{-15} to 10^{-7}. Of the secondary criteria, the slope passed and the periodic one failed: from $d = 10^{-9}$ on, more than 10 % of the periodic cells separated.
- Independently reproduced from saved data (15 October): every summary field of the 0.5B batch-layout comparison (448 cells, four comparisons) and the gain-split, logit-bias and crossover summaries of Section 3 (5353, 6837 and 24 cells); the coefficients and offset rows of the flip normal-form test (Table 18); the 1.5B and 7B fold-law scores under both period selectors (all 2700 cells); and the scalar conditions of the nine located cusps, matched one-to-one to the tangent reversals of the seven continued fold curves. This reproduces the statistics and saved scalars, not the model runs, derivatives or eigenvectors. The within-prompt steered test (null) used an analysis fixed and pushed before its data, without a separate frozen design document.
- Prospectively frozen and independently replayed: Huginn's relaxation test on held-out data, confirmed in both parts; an independent implementation reproduced both from the saved rows and spectra (GSM8K part exactly, item-level part to 10^{-12}) (Section 15).

Hardware. Float64 GPU computations ran on rented NVIDIA H100 (NVL, PCIe and 80 GB HBM3), H200, H200 NVL and A100 80 GB (SXM4 and PCIe) GPUs, one GPU per run except for the 70B-class models, whose layers were spread over two; the float32 frames of the rotation video ran on RTX 4090 GPUs; the first reset test of Section 6 ran on a local NVIDIA RTX 2000 Ada workstation GPU and its replication on an H100 NVL; cross-checks, the interval proof, the second prospective passage test, the coexistence runs and the near-nilpotent searches ran on CPUs. The manifest `docs/manifests/runpod-runs.json` records for each run the GPU type, the code commit, the job and its arguments, the wall time and the exit status, and the available hashes of the uploaded code and input archives and of the retrieved results, with absent fields left explicit; per-run receipts and accepted provenance limitations are kept with it. Regenerated at 02:11 UTC on 16 October 2026, it covers the runs up to that time: it lists 438 runs with retrieved results (code and input-archive hashes for 330, result-archive hashes for 412) and 21 pods that ended without results. Later runs are recorded with their receipts in the research log. GPU types round float64 reductions differently, so quantities used across runs were checked across hardware where it mattered (for example, flip parameters agree between CPU and H100 to 4×10^{-15}). Code: `tools/explore_llm_soft_loop.py`, `tools/bifurcation_curves.py`, `tools/llm_feigenbaum.py`, `tools/explore_latent_loop.py`, `tools/cusp_check.py`, `tools/huginn_relax_eval.py`, `tools/huginn_arc_analysis.py`, `tools/steer_gain.py`, `tools/crossover.py`, `tools/repro_noise.py`, `tools/repro_task.py`, `tools/repro_summary.py`, `tools/bpc_analyze.py`, `tools/turning_scan.py`. Raw-data hash receipts: `evidence/llm-dynamics-data-receipts.json` (to 12 October), the per-run receipts in the research log, and the verified snapshot manifests in `docs/backups/`; independent cross-check and audit receipts are under `codex/` (`crosscheck`, `confirm-fold`, `hostile-fold-crosscheck`, `huginn-design-review`, `huginn-independent-replay`, `paper-review`).