Ghosts, cusps and chaos: bifurcation theory predicts the dynamics of a language model fed its own output

Authorship to be decided by the human researcher

Private working draft, 13 October 2026

Status of this draft. Unless marked otherwise, results on the soft-token map are numerically validated in float64 (rung 1 or 2 of the evidence ladder in Section 14). Float32 maps are labelled as such. Huginn's spectra are float32 and its accuracy runs bfloat16, and the dialogue results are sampled statistics with stated intervals. **[Planned:** Red] marks planned work and **[In progress:** orange] running work. The only computer-assisted proof is that of the period-8 orbit of Table 1 (rung 3, Section 14).

Abstract

Feed a frozen language model its own expected next-token embedding, at a fixed position after a fixed prompt, and it can lock into a loop of text such as ". This is there is a is a" or wander with positive finite-horizon Lyapunov exponents. This exactly defined smooth map, with temperature and activation-steering strength as parameters, shows both regimes in Qwen2.5 models of 0.5B to 7B, OLMo-2 7B and Llama-3.1 8B. Chaos needs no steering, and equal-norm random directions produce it as readily as concept directions. Universality makes such structure plausible but cannot say whether a trained model realizes it, or whether bifurcation theory predicts its behaviour from the model's own derivatives. In the cases we tested, it does. Fold curves of two periodic windows end in numerically nondegenerate cusps whose local unfolding, frozen before solving, predicted the number of periodic orbits at all twelve tested parameter offsets and their stability at eleven. In the recurrent-depth model Huginn, fixed-point spectra predict each question's convergence under relaxation and where chain-of-thought accuracy breaks; a frozen test on held-out data confirms both. Fold normal forms predict bottleneck passage times without fitting: at two new folds, 23 of 24 prospectively frozen forecasts fell within one iterate, failing the all-row criterion by one row. Interval arithmetic proves that one period-8 orbit exists and is locally attracting, to our knowledge the first such proof for a full transformer language model. In sampled self-dialogue, hostility steering has a non-monotone dose response. Where it rises, in two model sizes, replacing classifier-flagged replies before they enter the context lowers the model's own hostility far more per unit of steering than constant prevention; past the peak, counter-steering backfires. The constructed map is not ordinary autoregressive generation.

1 Introduction

Repeated inference with a fixed language model is an iterated map: the model's output becomes its next input. Like other iterated maps, it can be chaotic. Figure 1 shows a small piece of the parameter plane of one such map, built from Qwen2.5-0.5B-Instruct. Cells with a positive finite-horizon Lyapunov exponent lie side by side with windows in which the model repeats a fixed loop of text. Parts of this picture have been seen before. Recent work has measured chaos-like

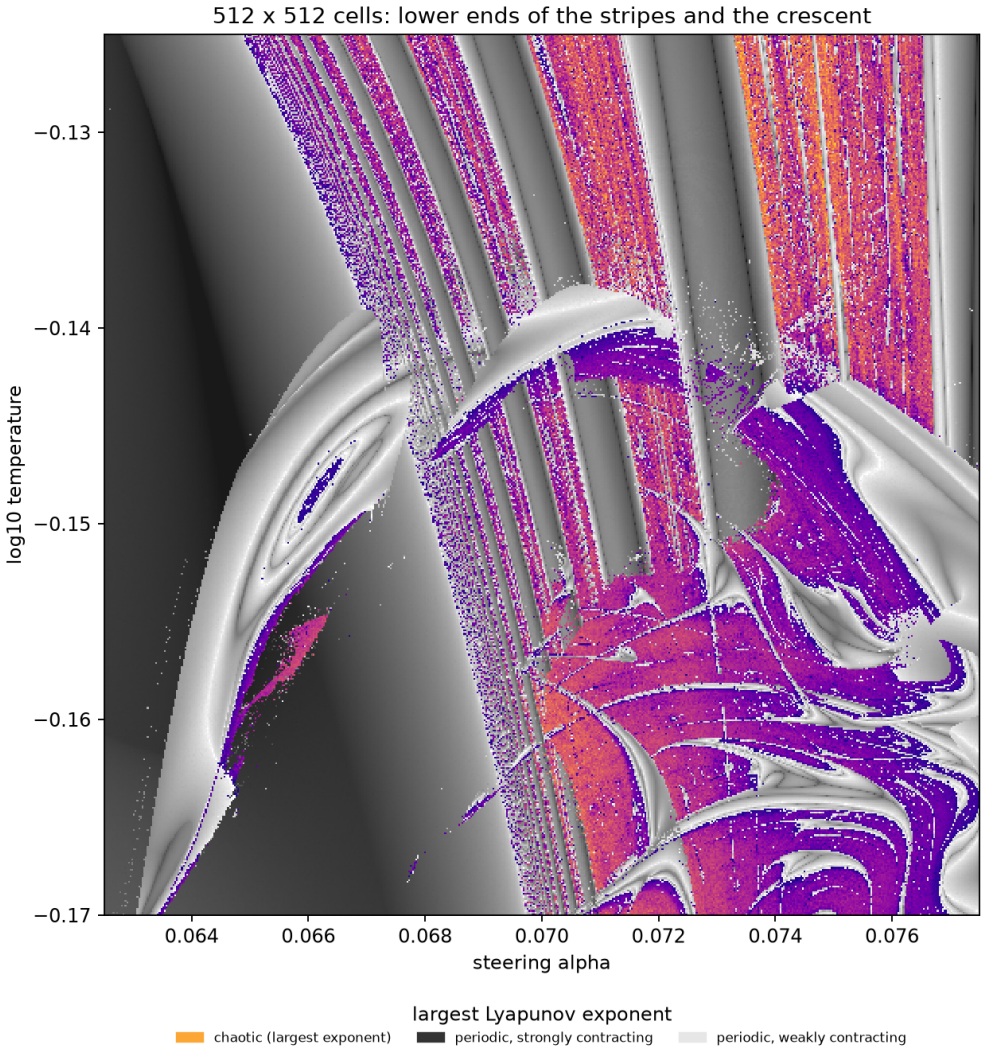


Figure 1: Chaos and repeating text in a language model fed its own output: the soft-token map of Qwen2.5-0.5B-Instruct, steered along the philosophy direction at layer 12, on 512×512 float64 cells over the steering strength $\alpha \in [0.0625, 0.0775]$ and the temperature, $\log_{10} T \in [-0.17, -0.125]$. Coloured cells are chaotic by the finite-horizon criterion of Section 3; the colour gives the largest Lyapunov exponent. In grey cells the model settles into a repeating loop of text, shaded by how strongly the loop contracts (light grey: close to losing stability). Figure 3 names the loops.

sensitivity in language-model generation [1, 2], temperature-driven changes of output statistics [3–7], periodic attractors of paraphrasing loops [8], attractor states of model-to-model conversations [9], and, through their Jacobians, the stability of looped and recurrent-depth models [10–14]. Feeding a probability-weighted vocabulary embedding back into a model is not new either. It is used in training [15] and, without retraining, in pretrained models [16], alongside continuous latent feedback [17]. In a small task-trained looped transformer, fold coefficients computed from derivatives predict passage times and spectral gaps without fitting [18].

We study one constructed but exactly defined system: the expected-embedding ("soft-token") map of a frozen, released pretrained language model, run with a fixed chat prefix at a fixed position. The map is smooth. Its Jacobians come from automatic differentiation, and its periodic orbits can

be computed and continued like those of any smooth map. Theory makes chaos in such a family plausible. High-dimensional nonlinear maps with a gain-like parameter generically become chaotic, as random recurrent networks do above a critical gain [19], and periodic windows, shrimp-shaped domains and period-doubling cascades are universal in dissipative two-parameter maps [20, 21]. These theories say that such structure can occur. They do not say whether a trained language model realizes it, where in its parameter space, with which texts, or whether the quantitative predictions of bifurcation theory hold for it, so that derivatives computed from the model predict its behaviour before that behaviour is measured. The models we studied do realize it, and in the cases we tested the predictions hold. Fold normal forms predict how long trajectories linger in the ghost of a vanished orbit, cusp unfoldings predict how many periodic orbits exist near a cusp and which of them are stable, and the fixed-point spectra of a recurrent-depth model predict its response to relaxation and the onset of task failure on held-out data. Ordinary sampled generation with a growing context is a different system. Section 9 measures how much of the mechanism survives sampling, and Section 12 studies sampled dialogue directly.

Contributions.

1. *A fold's normal form predicts how long the model lingers near a vanished orbit.* The 2-cycle "is it" dies in a fold, and the fold's normal form, computed from derivatives, predicts bottleneck passage times without fitting. We tested it prospectively at two new folds: 23 of 24 forecasts fell within one iterate, and the all-row criterion failed by one row. The global reinjection sets the prefactor of the window law (Section 6).

2. *Cusp unfoldings predict how many periodic orbits coexist and which of them are stable.* Fold and flip curves continued with float64 test functions show that the fold curves of a period-6 window and of a hostile period-7 window end in numerically nondegenerate cusps. At one cusp of each, the two-parameter unfolding is regular, and its local model, frozen before solving, predicted the orbit counts at all twelve parameter offsets and their stability at eleven (Sections 7 and 12).

3. *In a second system, fixed-point spectra predict where chain-of-thought accuracy breaks.* In the recurrent-depth model Huginn we predict a Neimark–Sacker bifurcation and then observe it. Spectral estimates also predict each question's response to relaxation and where chain-of-thought accuracy breaks, and a frozen test on held-out data confirmed both predictions (Section 11).

4. *One orbit is proved, and the central numerical results are checked independently.* Orbits, multipliers and test functions of the soft-token map are computed in float64 with precision checks (Section 2). Independent reimplementations reproduce the fold and its normal form and two orbits, a second hardware platform reproduces four flip parameters, and an independent replay reproduces the frozen Huginn test. An interval-arithmetic proof shows that a primitive period-8 orbit exists and is locally attracting, to our knowledge the first proof of a periodic orbit of a full transformer language-model map (Section 14).

5. *In sampled dialogue, keeping a flagged reply out of the context controls hostility at a fraction of the steering dose, and counter-steering can backfire.* Hostility steering of sampled self-dialogue produces steep, direction-specific capture with dose-dependent persistence and hysteresis. Replacing flagged replies before they enter the dialogue keeps the model's own hostility low at a quarter of the dose of constant prevention, while counter-steering triggered after a hostile reply is less efficient. In Qwen2.5-1.5B the same ordering holds where the dose response rises,

but the dose response is not monotone, and every counter-steer backfires beyond its peak. A pre-specified lead-time test is negative; a small warning from the model's in-context state is confirmed on fresh dialogues (Section 12).

6. *The structure that universality makes plausible is realized, with concrete texts, in several models.* Float64 phase maps in temperature and steering strength show chaos (positive finite-horizon exponents) with embedded windows (Section 4), a period-adding family whose text gains one "is X" pair per period step (Section 5), shrimp-like windows (Section 7), and a doubling cascade with superstable points through period 384 and finite-order Feigenbaum ratios (Section 8). The same organization appears in Qwen2.5-1.5B and 7B, OLMo-2 7B, Llama-3.1 8B (with sparse-autoencoder readouts) and a second checkpoint of the 0.5B architecture, which shares the organization but not the attractors (Section 10). Chaos needs no steering, and norm-matched random directions produce it as readily as concept directions (Section 13). A bridge to sampled decoding measures what survives sampling (Section 9).

Related work. Language-model generation has been studied as a dynamical system at several levels, and the closest studies differ from ours in what they compute. Lyapunov-like quantities have been used to measure sensitivity to small perturbations in reasoning traces [1, 2], but these studies measure chaos-like sensitivity in generation without computing orbits or bifurcations. Output statistics change abruptly with sampling temperature, and several studies analyse the changes as phase transitions or criticality [3–7]. These analyses concern output statistics; they are not bifurcation analyses of a defined map. Repetition loops are a known failure of decoding [22, 23], successive paraphrasing converges to attractor cycles [8], and multi-turn conversations between models settle into attractor states [9]. These attractors are behavioural: they are observed in generated text rather than computed as invariant sets of a defined map. Attractors and bifurcation diagrams have been studied in recurrent language models [24]. Equilibrium, looped and recurrent-depth models iterate a block to a fixed point or for many steps, and their stability has been analysed through Jacobians and fixed points [10–14, 25–30]; that analysis concerns fixed points. The closest precedent is a derivative-based fold analysis of a small task-trained looped transformer, whose passage and spectral-gap predictions are unfitted [18]. That model was trained for a task, whereas ours is a released pretrained language model, and our analysis also differs in its objects: the fold of a primitive periodic orbit, the full-channel passage of a two-step map, and the accumulation law of the windows, whose prefactor is set by the global reinjection. Expected-embedding feedback itself is used in training [15] and, training-free, in pretrained models [16, 31], next to continuous latent reasoning [17]. Fixing the prefix and the position turns that feedback into the autonomous map we analyse.

In neural systems more broadly, saddle-node ghosts produce quantitative delays in task-trained recurrent networks [32]. Finite-order Feigenbaum ratios have been measured in neural rate and mass models [33, 34] and, through period 128, in a transformer trained on logistic-map trajectories [35]. Learned networks reproduce period-doubling routes [36], and so does gradient descent itself [37]. Normal forms and continuation have been applied to recurrent neural maps [38], and shrimp-shaped windows and periodicity staircases appear in neuron models [39–41]. Our second control parameter comes from activation steering along contrastive directions [42, 43]. The dynamical-systems tools are standard: numerical continuation and normal forms [44, 45], type-I intermittency [46], shrimps and hubs [21, 47, 48], period-doubling universality [20], and Markov state models [49].

To our knowledge, this is the first work to compute periodic orbits with their multipliers, continue fold and flip bifurcation curves and locate numerically nondegenerate cusps in the computation of a pretrained language model. It is also, to our knowledge, the first to make quantitative predictions

from such a model's derivatives and confirm them prospectively: passage times at new folds, where 23 of 24 forecasts fell within one iterate and the all-row criterion failed by one row, and orbit counts and stabilities at frozen offsets near cusps. We know of no earlier confirmation on held-out data of a spectral prediction of where a language model's task accuracy breaks, which we give for Huginn, and of no earlier computer-assisted proof of a periodic orbit of a full transformer language-model map. These priority statements rest on a targeted search (12 October 2026) that found no prior work combining these elements for the feedback map of a pretrained language model. That is an absence of evidence within the search, not a proof of priority.

2 The inference map

Definition. We describe the map for Qwen2.5-0.5B-Instruct, the model studied in most detail; the other models enter the same construction with their own embeddings and steering layers (Section 10). Let $E \in \mathbb{R}^{V\times d}$ be the input embedding matrix ($V = 151{,}936$, $d = 896$). The model reads a fixed chat prefix, with the system prompt "You are an AI talking with another AI. You have complete freedom. Feel free to pursue whatever you want." and the user turn "Hello!", followed by a single input embedding x, and produces next-token logits $z(x) \in \mathbb{R}^V$. The soft-token map feeds the expected next-token embedding back in place of x:

$$F_{T,\alpha}(x) = E^\top \operatorname{softmax}(z_\alpha(x)/T), \tag{1}$$

where z_α adds αv to the residual stream at layer 12 at every position, prefix included. The steering direction v is contrastive: it compares last-token states of "philosophy" prompts with those of neutral prompts and is normalized to the mean residual norm. Positions are fixed and the window does not grow, so F is an autonomous map of $\mathbb{R}^{896}$. As $T \to 0$ it reduces, where the argmax is unique, to greedy decoding of the same fixed-prefix, one-token-window system. It is not ordinary generation with a growing context.

Precision. Hugging Face RMSNorm layers and rotary-embedding tables compute in float32 even when the model is cast to float64, so we patch both. A float32 map is accurate to 3×10^{-6}–2×10^{-5} per step. That is enough for phase maps (245 of 256 test cells agree with float64) but not for derivatives of composed maps: one-step float32 Jacobian-vector products are accurate to 5×10^{-5}, yet the 8-step monodromy of a period-8 orbit is off by 3%. All orbits, multipliers and test functions below are therefore computed in float64.

3 Methods

Phase maps. Each cell of a phase map is one parameter pair. Unless stated otherwise, we discard 300 transient steps and measure over the next 300. The largest Lyapunov exponent comes from a renormalized twin trajectory at offset 10^{-6}. The period is the smallest $q \le 32$ (48 in the later zooms) with return distance below 10^{-8}, and the attractor is identified by its leading-token loop, up to rotation. We call a cell chaotic when this finite-horizon exponent exceeds 10^{-3} and no period is found; the label is a finite-horizon diagnostic, not a proof of chaos. A 256^2 map takes about 100 minutes on one H100 GPU in float64. Every cell starts from the same state, so where attractors coexist, a map shows the one whose basin contains that state. Labels are text, not dynamical identities: one continuous orbit changes its label where two tokens tie for the maximum. We therefore call two labelled regions different attractors only when the contraction rate jumps between them or when float64 orbits of both exist at one parameter value.

Periodic orbits. We solve $F^p(x) = x$ by Newton's method with forward-mode Jacobians. The multipliers are the eigenvalues of the monodromy matrix; where the dense matrix is not float64, we refine them by a short Arnoldi factorization with float64 Jacobian-vector products.

Continuation. Periodic branches are continued in one parameter by pseudo-arclength continuation, and fold and flip curves in two parameters with minimally augmented systems [44, 45]. The test function g of the bordered system $[[M \mp I, v], [w^\top, 0]]$ is obtained by GMRES (generalized minimal residual iterations) with Jacobian-vector products, and its gradient by reverse-mode differentiation through a Jacobian-vector product. When the parameter tangent reverses, we flag a cusp candidate and test it with the normal form.

Normal form and passages. Consider a fold of the two-step map $G = F^2$, with right and left eigenvectors r and ℓ normalized by $\|r\| = 1$ and $\ell^\top r = 1$. Along the critical direction the reduced map is $u' = u + a\varepsilon + bu^2$, with $a = \ell^\top \partial_\alpha G$ and $b = \frac{1}{2}\ell^\top D^2 G[r, r]$. We compute b by central differences of float64 Jacobian-vector products over a step-size plateau (later cross-checks use nested automatic differentiation), and a includes the steering's effect on the prefix. The signed coefficients are kept in the original eigenvector gauge, with ε measured in the free parameter (α or $\log_{10} T$, whose units apply to a). For passages we set $v = \operatorname{sign}(b)\, u$ and take the ghost side $p - p_c = \operatorname{sign}(ab)\, \epsilon$ with $\epsilon > 0$, so that $v' = v + |a|\epsilon + |b|v^2$ to quadratic order. A full passage starts at $v = -U$ and a midpoint passage at $v = 0$; both end at the first strict crossing $v > U$, with $U = 0.1/|b|$. We compare passage counts with the iterated quadratic map and with the arctangent formula [46].

Period-doubling cascade. Along a cut in one parameter, we locate the flip of each period-p orbit as the root of $\operatorname{Re}\mu_c + 1$, where μ_c is the multiplier with the smallest real part, using the Illinois variant of regula falsi with warm-started orbits. All but one or two multipliers are small, so Newton's method on $F^p(x) - x$ needs the monodromy only on a six-dimensional dominant subspace. We obtain that subspace by subspace iteration with exact forward-mode Jacobian-vector products chained along the orbit, so no dense Jacobian is formed, and the residual is always the exact float64 $F^p(x) - x$. Each root is classified before it is accepted. The critical multiplier must be real and equal to -1, with a Ritz residual below 10^{-6}, and the other computed Ritz values (six in all, estimates of the dominant spectrum) must lie inside the unit circle. We then reach the period-$2p$ orbit by forward iteration inside the next band. On the logistic map embedded in six dimensions, the procedure reproduces the first four flip parameters to 10^{-9}.

4 Chaos with periodic windows

Figure 2 shows a 256×256 float64 map over $\alpha \in [0.060, 0.082]$ and $\log_{10} T \in [-0.17, -0.04]$. More than half of its cells are chaotic (35,759 of 65,536), and periodic windows of periods 2–32 are embedded in the chaos. Rendered as in the shrimp literature [40], the windows show dark curves of maximal contraction.

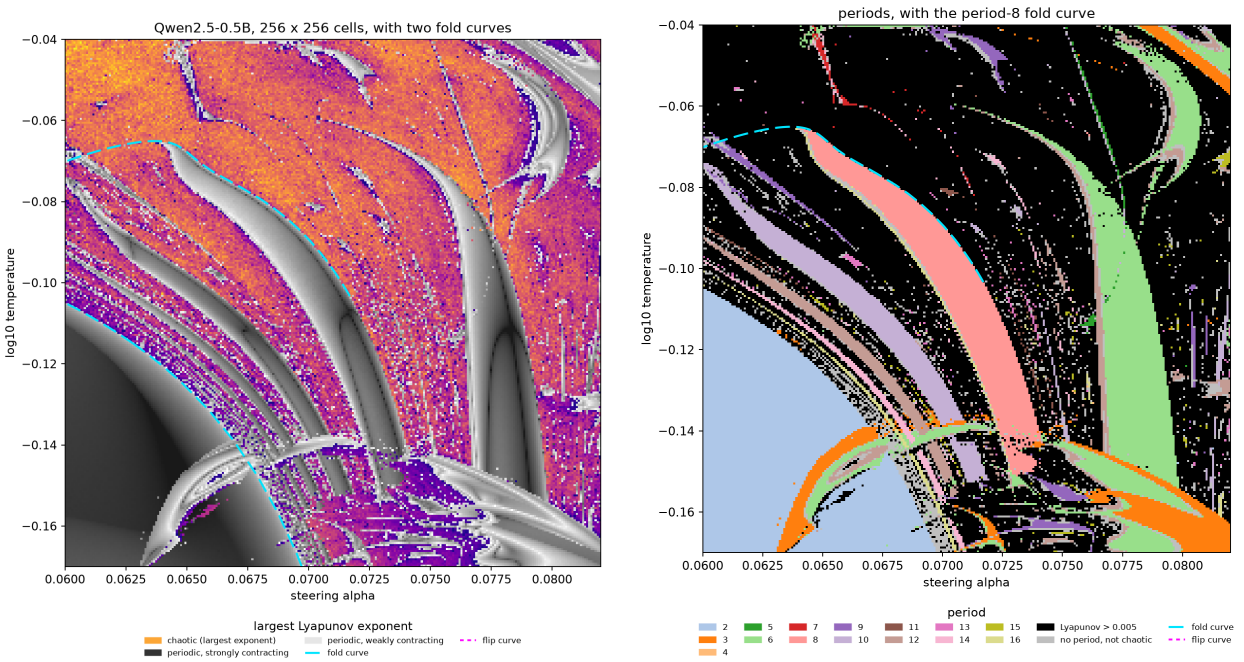


Figure 2: Soft-token map of Qwen2.5-0.5B-Instruct steered along the philosophy direction at layer 12, 256^2 float64 cells. Left: chaotic cells coloured by the largest Lyapunov exponent, periodic cells in grey by how strongly they contract (light grey: close to losing stability). Cyan dashed curves are float64 fold curves of the 2-cycle "is it" (lower left, the boundary on which the stripes accumulate) and of the period-8 orbit (top). Right: periods, with the period-8 fold curve (Section 7).

Figure 1 resolves the stripes' lower ends at 512×512 cells (79,278 chaotic), and Figure 3 labels the same cells by their attractor texts. Three structures stand out. The 2-cycle "is it" fills the left, bounded by its fold curve, on which the stripes accumulate. A closed crescent of the period-3 loop "\n\nThe existence" contains its period doublings and a chaotic core (Section 8). Its continuation, the period-6 band "\n\nThe question\n\nThe existence", crosses the stripes near $\log_{10} T = -0.14$, and where the crescent lies on the 2-cycle's side of that fold curve, the two attractors coexist. The lower right is a chaotic region filled with small hooked windows of "\n\nThe question" and "\n\nThe existence" loops.

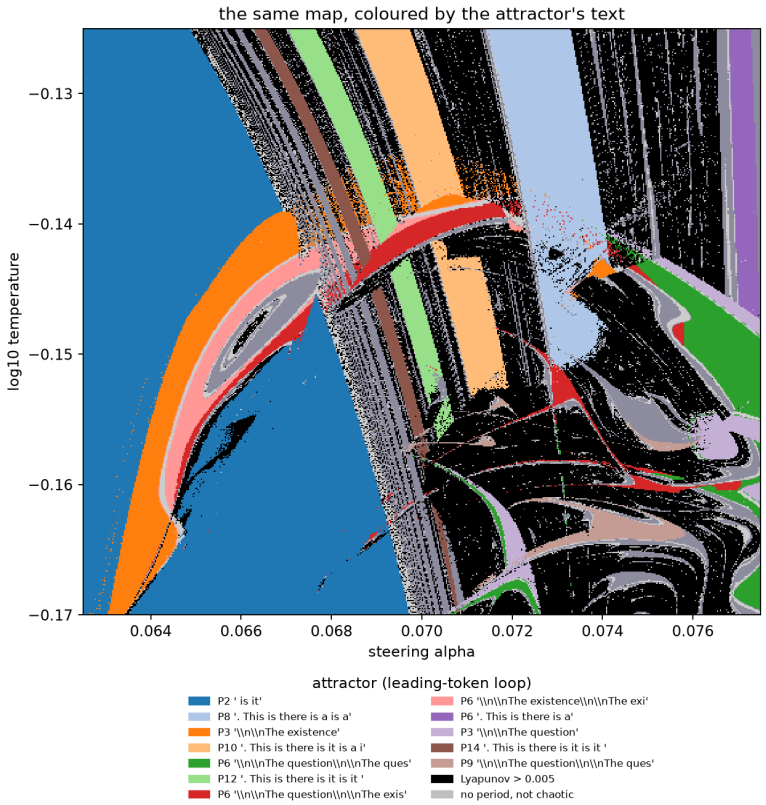


Figure 3: The cells of Figure 1 (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12; 512^2 float64 cells over $\alpha \in [0.0625, 0.0775]$, $\log_{10} T \in [-0.17, -0.125]$), coloured by the attractor's leading-token loop (black: chaotic; grey: no period up to 32). Figure 1 shows the largest Lyapunov exponent of the same cells.

Zooms of 256^2 cells (Figure 4) show what the model writes in each window, and every window is a repeating piece of text. The period-8 band repeats ". This is there is a is a". The neighbouring bands add laps of "is it" (periods 10 and 12), and along the band's flip edge the period-16 loop repeats the sentence twice. The small windows inside the chaos break the "is X" chain with a paragraph break and "The" (periods 11, 12, 13, 14). Inside the period-8 and period-6 bands, two curves of maximal contraction meet in a sharp junction.

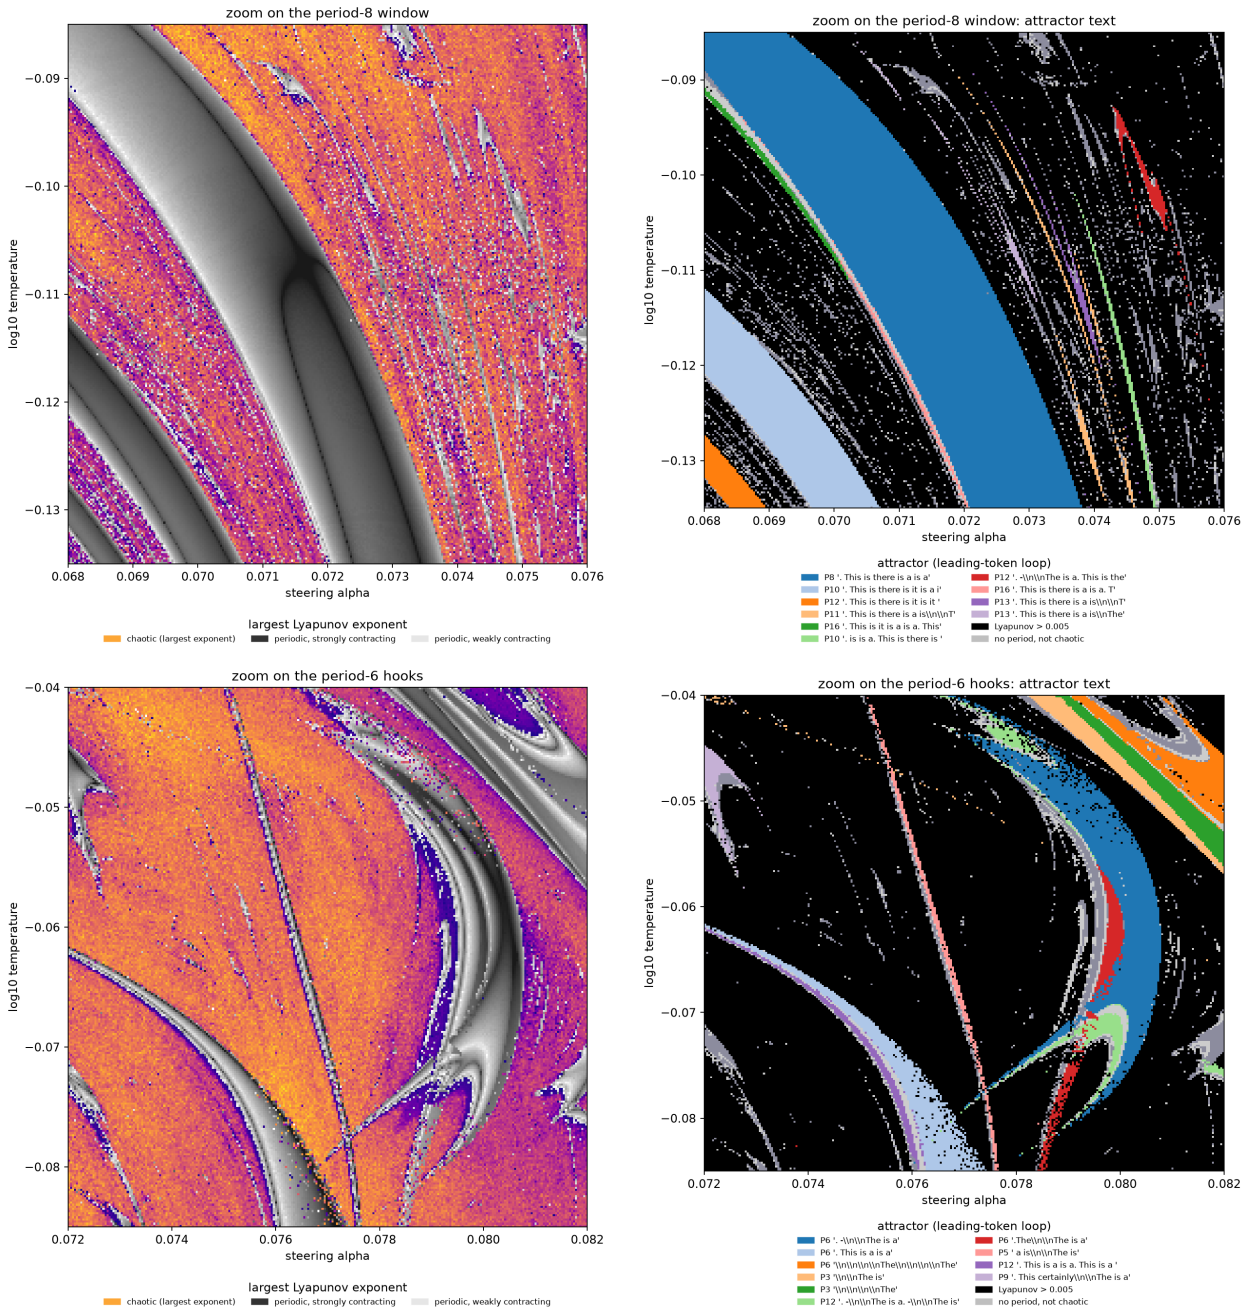


Figure 4: Zooms of 256^2 float64 cells of the Qwen2.5-0.5B-Instruct map (philosophy direction, layer 12), with the largest Lyapunov exponent on the left, rendered as in Figure 1, and attractor texts on the right. Top: the period-8 window ($\alpha \in [0.068, 0.076]$, $\log_{10} T \in [-0.135, -0.085]$). Bottom: the period-6 hooks ($\alpha \in [0.072, 0.082]$, $\log_{10} T \in [-0.085, -0.04]$), whose curved period-6 body, period-12 foot and thin leg are analysed in Section 7.

5 Period adding written in text

Between a period-2 region and the chaos lies a fan of stripes of periods $6, 8, 10, \ldots$ Their orbits, numerically validated in float64, can be read as text (Table 1). Each is ". This" followed by k pairs

"is X", with period $2k+2$, so each stripe adds one pair. The period-2 region itself is the 2-cycle "is it". At the stripe centres the orbits are effectively one-dimensional: the leading multiplier is 0.003–0.09, and the next falls from 3×10^{-6} (period 6) to 2×10^{-14} (period 20). This compression belongs to these orbits, not to the map. The 2-cycle's multipliers are -0.271 and -0.058 (the second modulus is about 0.037 near its fold), the next modulus at the period-3 flip of Section 8 is 0.088, and one period-6 orbit in Section 7 has a complex leading pair.

period	orbit (leading tokens)	μ_1	$\lvert\mu_2\rvert$
6	. This is a is a	$+0.0052$	3×10^{-6}
8	. This is there is a is a	$+0.0110$	3×10^{-7}
10	. This is there is it is a is a	$+0.0025$	2×10^{-7}
12	. This is there (is it)^2 is a is a	$+0.0163$	4×10^{-8}
14	. This (is it)^4 is a is a	$+0.0369$	2×10^{-9}
16	. This (is it)^5 is a is a	$+0.0716$	8×10^{-11}
18	. This (is it)^6 is a is a	$+0.0792$	1×10^{-12}
20	. This (is it)^7 is a is a	$+0.0857$	2×10^{-14}

Table 1: Periodic orbits of the stripe family at the stripe centres along $\log_{10} T = -0.115$ ($\alpha = 0.077872$, 0.071674, 0.067766, 0.065550, 0.064443, 0.063857, 0.063534, 0.063335), computed with the current all-float64 map (norms and rotary tables patched): Newton residuals 6×10^{-15} to 8×10^{-13}, leading multiplier μ_1 and next modulus $|\mu_2|$ from Rayleigh–Ritz values. Each step in period adds one "is X" pair to the text. The 2-cycle "is it" itself has multipliers -0.2706 and -0.0585 at $(\log_{10} T, \alpha) = (-0.17263, 0.061579)$. The period-2 and period-8 orbits were also reproduced by an independent implementation, and the period-8 orbit is certified by a computer-assisted proof (Section 14).

6 Mechanism: a saddle-node bottleneck

The fold. At $\log_{10} T = -0.115$ the 2-cycle "is it" dies in a fold at $\alpha_c = 0.06279827$ (float64 multiplier 1 to 3×10^{-13}, next modulus 0.037). The normal form of G has $a = -4.834$ per unit α and $b = -25.4285$, so $ab = 122.93$. The stable branch gives an independent estimate: extrapolating its multipliers through $(1-\mu)^2/(4(\alpha_c - \alpha)) \to ab$ yields 123.1 with a quadratic fit (120.6 with a linear one). An independent implementation of the forward map and solver recovers α_c to 6×10^{-15} and a to 3×10^{-12}. Its b, from nested automatic differentiation, is -25.42846; our finite-difference value differs from it by 2.2×10^{-6} relatively, which is the finite-step bias of the difference table. The resulting $ab = 122.927$ leaves every passage prediction of Table 2 unchanged.

Controlled passages. We start trajectories at $u_{\text{in}} = -U$ or 0 ($U = 0.1/|b|$) and count G-iterates until $u > U$. At four distances from the fold, the measured counts match the iterated normal form within one iterate (Table 2, Figure 5), and a transverse displacement of the same norm changes the count by at most one.

$\varepsilon = \alpha - \alpha_c$	full: measured / predicted	midpoint	full + transverse kick
3×10^{-6}	144 / 144	73 / 74	143 / 144
1×10^{-5}	71 / 71	36 / 37	70 / 71
3×10^{-5}	34 / 34	17 / 18	33 / 34
1×10^{-4}	14 / 14	7 / 7	13 / 14

Table 2: Bottleneck passages (G-iterates). Predictions iterate $u' = u + a\varepsilon + bu^2$ with a, b computed from the network's derivatives; no parameter is fitted.

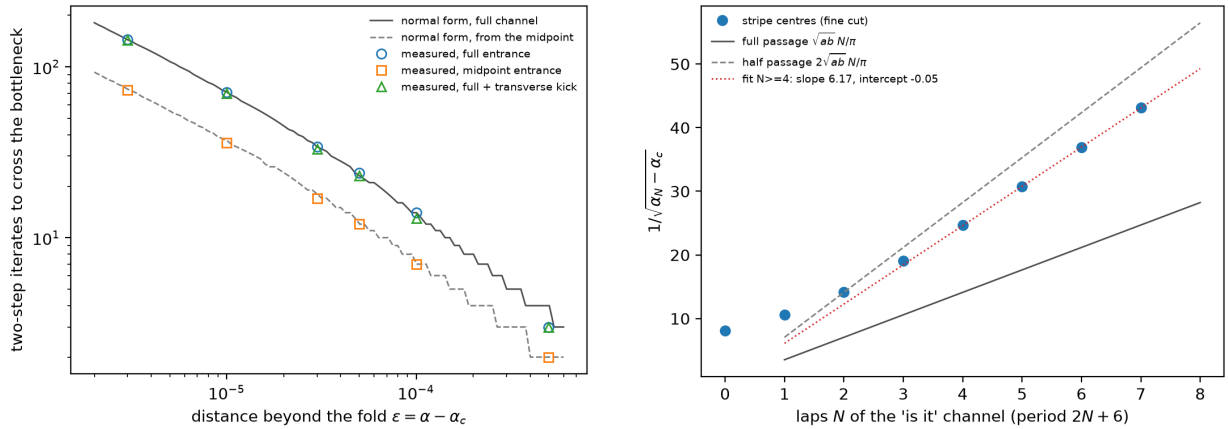


Figure 5: Bottleneck passages and the window law at the fold of the 2-cycle "is it" (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, $\log_{10} T = -0.115$). Left: measured passage times against the normal-form predictions. Right: the window law, with stripe centres plotted as $1/\sqrt{\alpha_N - \alpha_c}$ against the lap count N, together with the full- and half-passage predictions.

Natural reinjection and the window law. Free trajectories re-enter the channel slightly upstream of the bottleneck's middle ($u \approx -3 \times 10^{-4}$ to -7×10^{-4}). From each measured entry point, the normal form predicts the natural laminar durations within 4–8% for $\varepsilon \leq 10^{-4}$, and a natural passage lasts about 0.6 of a full one. The stripes therefore follow the N^{-2} law of a saddle-node ghost, with about a third of the full-passage prefactor. Stripe centres located on a float32 cut along $\log_{10} T = -0.115$ give $d_N N^2 = 0.0263$–0.0265 for $N = 4$–7 (log-log slope -2.00), compared with $\pi^2/(4ab) = 0.0201$ for a half passage. Two-step return maps of p("it") (Figure 6) show the near-tangency.

Numerically estimated stripe edges along $\log_{10} T = -0.115$ (Table 3) sharpen this picture. Each window is bounded by a fold, where its orbit is born, and by a flip. For periods 6–12 the fold faces away from the accumulation point. For periods 16–20 the sides are exchanged, and the period-14 window is bounded by flips on both sides. The fold edges of periods 16–20 follow $\alpha_N - \alpha_c = C/(N + \delta)^2$ with $C = 0.0266$ and $\delta = 0.15$. Their slope, a rise of 6.13 per lap in $1/\sqrt{\alpha_N - \alpha_c}$, is 1.74 times the full-passage slope $\sqrt{ab}/\pi = 3.53$ of the normal form. That corresponds to an effective passage of 0.58 of a full one, which the measured natural reinjection gives independently.

period	edge toward larger α	edge toward smaller α
6	fold 0.0790337	flip 0.0766947
8	fold 0.0726585	flip 0.0706525
10	fold 0.0684572	flip 0.0670594
12	fold 0.0658982	flip 0.0651944
14	flip 0.0645850	flip 0.0642958
16	flip 0.0639284	fold 0.0638001
18	flip 0.0635683	fold 0.0635034
20	flip 0.0633535	fold 0.0633183

Table 3: Edges of the stripe windows along $\log_{10} T = -0.115$ ($\alpha_c = 0.0627983$), from the window centres of Table 1. Flips: critical multiplier a real -1 to 3×10^{-10}. Folds: the multiplier nearest $+1$ extrapolated to 1 through $(1 - \mu)^2$, which is linear near a fold; last converged multiplier 0.991–0.998.

[**Planned:** prediction of the window positions from the measured reinjection distribution.]

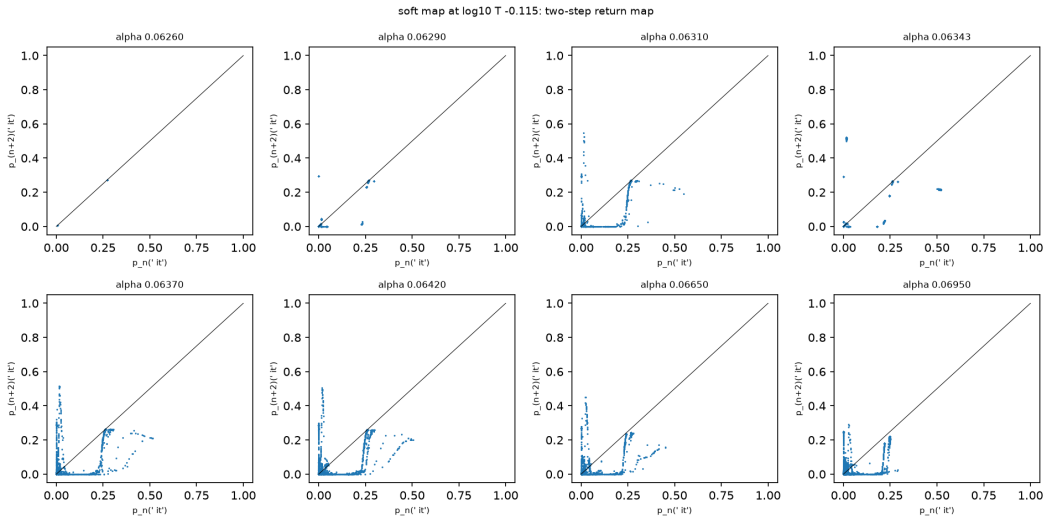


Figure 6: Two-step return maps of the probability of "it" near the fold of the 2-cycle "is it" (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12), at increasing distance from the fold.

A prospective test. To test the normal form at folds that played no part in shaping the mechanism, an independent implementation located two new folds and computed their normal-form coefficients by nested automatic differentiation. All 24 passage predictions were pushed in a frozen design before any test trajectory was run. Target A is the 2-cycle's fold at a new temperature: $\log_{10} T = -0.140$, $\alpha_c = 0.0670549$, $a = -7.055$, $b = -24.927$. Target B is the fold of the period-8 orbit of Table 1, continued in temperature at $\alpha = 0.0717$: $\log_{10} T_c = -0.1033462$, $G = F^8$, $a = -11.900$, $b = -152.05$ per unit $\log_{10} T$. Of the 24 forecasts (four distances and three entrances per target), 23 fell within one iterate of the measured passages (Table 4). The pre-specified criterion, that every row pass, failed by one row: for target A, midpoint entrance, $\varepsilon = 3 \times 10^{-6}$, the forecast was 61 and the measurement 59. All full-channel passages and all passages after a transverse kick were within one iterate.

ε	target A ($G = F^2$): measured / predicted	target B ($G = F^8$): measured / predicted
3×10^{-6}	118/118, **59/61**, 117/118	26/26, 13/14, 25/26
1×10^{-5}	56/56, 28/29, 56/56	10/10, 5/5, 10/10
3×10^{-5}	26/26, 13/14, 26/26	4/4, 2/2, 4/4
1×10^{-4}	10/10, 5/6, 10/10	2/2, 1/1, 2/2

Table 4: Prospective passage test at two folds that did not shape the mechanism. Each cell lists the full-channel, midpoint and full-plus-transverse entrances as measured / predicted G-iterates. Predictions come from the iterated quadratic normal form, with coefficients from an independent implementation, frozen before the test runs. Bold: the one row outside the one-iterate criterion. The shortest period-8 rows give only weak relative accuracy.

7 Bifurcation curves and cusps

Period-8 stripe. Continued for 80 points with a float64 test function, the fold curve of the period-8 orbit keeps its critical multiplier within $1 \pm 6 \times 10^{-6}$, follows the stripe's upper edge and

caps its top near $\log_{10} T = -0.065$ (Figure 2). At $\alpha = 0.0717$ the stripe lies between a flip at $\log_{10} T = -0.128955$ and the fold at -0.1033462. Continued downward, the flip curve keeps its multiplier at -1 for 115 points to $(\alpha, \log_{10} T) = (0.0776, -0.402)$, with two reversals of its parameter tangent near $(0.0750, -0.219)$ and $(0.0746, -0.206)$, and the fold curve keeps its multiplier at $+1$ for 120 points to $(0.0758, -0.200)$. [**Planned:** classify the two reversals on the flip curve (on a flip curve they are not fold cusps).]

The accumulation boundary. The fold curve of the 2-cycle "is it", continued from $(\alpha, \log_{10} T) = (0.0628, -0.115)$ in both directions (38 and 126 points, critical multiplier 1 to 10^{-7}, no cusp), traces the boundary of the period-2 region in Figures 2 and 3. The stripe fan accumulates on this boundary, as the mechanism of Section 6 requires. Further along one direction, outside the plotted window, another multiplier leaves the unit circle (moduli up to 11). This is a codimension-two point, [**Planned:** to be located and typed].

A period-6 window with shrimp-like anatomy. The hooks of Figure 4 (bottom) belong to the period-6 loop ". -\n\nThe is a". Float64 orbits (Newton residuals 10^{-15}–10^{-13}) show that its curved body and its thin leg are one window: the multipliers are -0.061 and $+0.048$ at the body, -0.564 lower down and -0.645 on the leg. The foot is its period-doubled loop (period 12, leading multiplier -0.156). Inside the body lies a different 6-cycle, ".The\n\nThe is a", whose leading multipliers are a complex pair, $-0.031 \pm 0.152i$. One-parameter branches from the body end in folds in every direction: at $\log_{10} T = -0.0558$ and -0.0721 in temperature, at $\alpha = 0.08074$ in steering, and, toward smaller α, in two folds only 10^{-5} apart. Such a close pair is the signature of a nearby cusp. A flip at $\alpha \approx 0.0797$ then opens the period-12 foot.

Fold curves that bound a window and meet in cusps, with flip curves inside, form the codimension-two organization of a shrimp [47] (complex dynamics of small recurrent networks has been studied in the same spirit [50]). In the continued curves (Figure 7), a single fold arc bounds the body on its outer side and a flip arc on its inner side, and the leg lies between them where they converge. The outer fold arc extends far to the left as a thin arm and ends in a cusp at $(\alpha, \log_{10} T) = (0.06911, -0.03975)$. The fold curve through the close fold pair has a second cusp inside the body, at $(0.08010, -0.05835)$, from which two fold branches descend.

At both cusps the quadratic coefficient b changes sign while a stays away from zero (Table 5). Because the signs of a and b both flip with the orientation of the critical eigenvector, the test uses $b \operatorname{sign} a$. A cusp is nondegenerate when the cubic coefficient c of the normal form $u' = u + a\varepsilon + bu^2 + cu^3$ is nonzero where b vanishes. We compute c with the second-order centre-manifold term, $c = \ell^\top C(r, r, r)/6 + \ell^\top B(r, h_2)/2$ with $(I - M)h_2 = B(r, r) - 2br$ and $\ell^\top h_2 = 0$, where M is the monodromy matrix and B and C are the second and third derivatives of its map, obtained by central differences of float64 Jacobian-vector products. (The code reproduces c to 10^{-6} on a map with a known reduction.) At both cusps c is nonzero far beyond its step-size uncertainty and keeps its sign through the zero of b, which crosses zero with nonzero slope along the curve. Both are therefore nondegenerate cusps, numerically. The negative sign of c means that between the fold branches near each cusp, two stable period-6 orbits coexist with a saddle.

A generic codimension-two cusp also needs the two parameters to unfold it independently. Near the cusp the reduced map is $u' = u + \beta_1 + \beta_2 u + bu^2 + cu^3$, where β_1 and β_2 vanish at the cusp and depend on the parameters. The unfolding is regular when $J = \partial(\beta_1, \beta_2)/\partial(\log_{10} T, 10\alpha)$ has rank two ($\det J \neq 0$). At the outer-arc cusp, $\det J = 2737$ (singular values 1630 and 1.7). With these derivatives, the local cubic model predicted the number of period-6 orbits and their stability at six parameter offsets, written to disk before any orbit was solved, and Newton solutions matched at all

six.

	cusp on the outer arc	cusp inside the body
b sign a at consecutive curve points	+92.9, +42.9, +14.7, **−16.9**, −52.8	−24.7, −14.3, −2.0, **+12.8**, +28.0, +45.9, +58.9
$\|a\|$	1.2–6.9	3.13–3.35
c at the two points around $b = 0$	-1.37×10^4, -1.65×10^4	-2.38×10^5, -2.34×10^5

Table 5: Normal-form test at the two parameter-tangent reversals of the period-6 window's fold curves (bold: the flagged points; critical multiplier 1 to 10^{-7}). c: cubic coefficient (finite-difference steps agree to 0.1 and 0.7 %). Numerically validated, not certified.

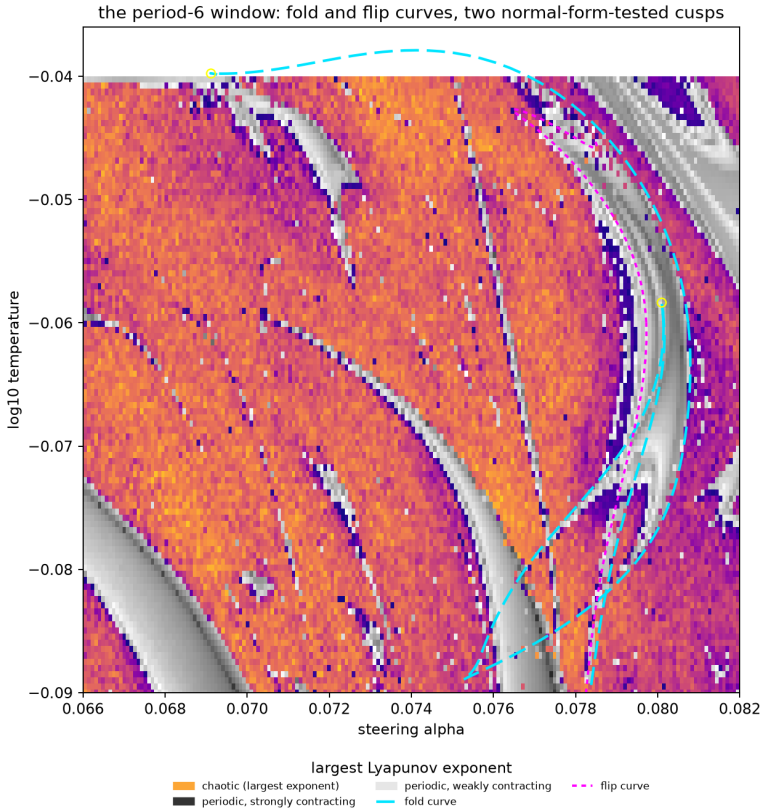


Figure 7: The period-6 window of the Qwen2.5-0.5B-Instruct map (philosophy direction, layer 12) with its float64 fold curves (cyan, long dashes), flip curve (magenta, short dashes) and the two normal-form-tested cusps (yellow circles), drawn over the 256^2 map of Figure 2 (whose data end at $\log_{10} T = -0.04$). **[Planned:** the lower ends of the fold curves, which approach each other near $(0.0753, -0.089)$.]

A period-4 window with two cusps. At higher temperature, $\log_{10} T \approx 0$–0.1, periodic orbits are no longer effectively one-dimensional (second multipliers 0.06–0.09). At $\alpha = 0.06484$ the temperature branch of the period-4 orbit "\n\n - \n\n The" is S-shaped, with folds at $\log_{10} T = 0.0640$ and 0.0664. When that fold is continued in the plane, its parameter tangent reverses at three points, two of which we tested. At both, the quadratic coefficient b changes sign while a stays away from zero, and the cubic coefficient c (defined above) is nonzero (Table 6). The two tested points are

therefore nondegenerate cusps of the fold curve, numerically, at $(\alpha, \log_{10} T) = (0.06296, 0.05996)$ and $(0.06557, 0.07445)$. The third, at $(\alpha, \log_{10} T) = (0.08308, 0.10863)$, was not tested. Three period-4 orbits coexist between the two fold branches, and near the first cusp ($c < 0$) two of them are stable (Figure 8).

	cusp 1	cusp 2
b at five curve points	$+28.1,\ +18.7,\ +8.4,\ \mathbf{-1.1},\ -12.8$	$+51.8,\ +29.7,\ +9.3,\ \mathbf{-9.0},\ -23.4$
a	3.02–3.08	-2.58 to -2.57
critical multiplier	1 ± 10^{-9}	1 ± 10^{-8}
c around $b = 0$	$-1.06,\ -1.01$ ($\times 10^4$)	$+1.39,\ +1.36$ ($\times 10^5$)

Table 6: Normal-form test of the two parameter-tangent reversals (bold: the flagged points). c: cubic coefficient, nonzero at both (steps agree to 10^{-4}), so both are nondegenerate cusps numerically; at the second, $c > 0$: the orbit between the two others is the stable one.

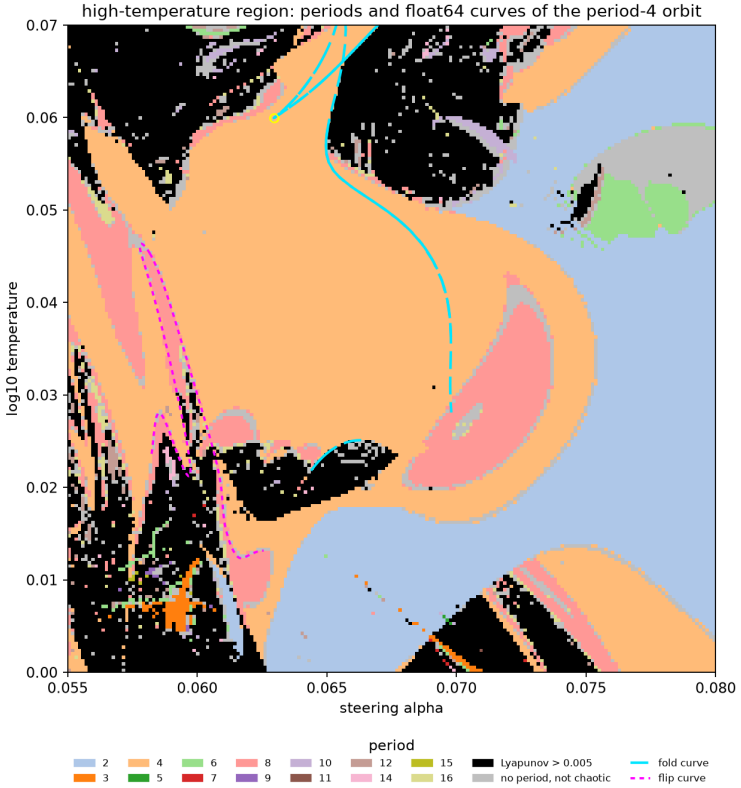


Figure 8: Period map of the high-temperature region of the Qwen2.5-0.5B-Instruct map (philosophy direction, layer 12; 192^2 float64 cells) with float64 fold curves (cyan, long dashes) and flip curves (magenta, short dashes) of the period-4 orbit. The yellow circle is the first cusp; the second lies just above the window. The flip curves run along the boundary between period 4 and period 8. The fold curves run inside the period-4 region, where three period-4 orbits coexist (near the first cusp, two of them stable). In this map a C-shaped band of the doubled loop "\n\nThe\n\nThe" (period 4) inside the period-2 region "\n\nThe" contains period 8 and an irregular core. Along the cut of Section 8 the period-8 orbit merges back into period 4 (a doubling bubble), so this band is not shown to be a second cascade. [**Planned:** the remaining fold curves of the period-4 region; whether the region is one shrimp-organized component.]

[**Planned:** Codimension-two points where a second multiplier crosses the unit circle along the

lower fold curve (candidate fold-flip points): locate, type and continue.]

8 A period-doubling cascade to chaos

The crescent of Figure 3 is a closed region in which the period-3 loop "\n\nThe existence" doubles inward. Along the cut $\log_{10} T = -0.14798$, the map's cells show period 3 over 25 columns, period 6 over about 14, period 12 over 4 and period 24 in one column, followed by a chaotic core. The doubled loops repeat the same leading tokens twice and four times, with alternating probabilities. Locating the flips along this cut in float64 (Section 3), in the steering coordinate α, gives Table 8. The first two ratios, $\delta_2 = 2.844$ and $\delta_3 = 2.859$ (indexed as in Table 8), are far from Feigenbaum's 4.6692, even though from the second flip on the orbits are effectively one-dimensional (next modulus 0.0023, 2×10^{-6}, 2×10^{-12}). The next ratio, $\delta_4 = 3.781$, is markedly larger. The first four flip parameters agree between completed float64 computations on a CPU and an H100 to 4×10^{-15}.

Continuing a just-born doubled orbit past period 48 failed: Newton's method slides onto the half-period orbit, which is also periodic with the doubled period. We therefore also located the superstable parameters s_p, where the critical multiplier vanishes in the middle of each band (Table 7). There the orbit attracts, so it can be reached by forward iteration and cannot slide onto its half-period parent. Through period 384 the superstable ratios are 3.19, 4.02, 4.478 and 4.625. The last two come from parameters resolved to 10^{-14} and differ from Feigenbaum's value by 0.19 and 0.044; from 4.02 on, the distance shrinks by factors of 3.4 and 4.3. The flip ratios 2.84, 2.86 and 3.78 rise in step (Figure 9), and the accumulation point lies near $\alpha_\infty \approx 0.0661250$. The early superstable values are less precise. The other computed Ritz values of the period-12 and period-24 orbits blur the root, so s_{12} and s_{24} are known only within brackets of 1.2×10^{-5} and 7×10^{-7}, and the ratios 3.19 and 4.02 have worst-case ranges $[2.90, 3.34]$ and $[4.02, 4.10]$. s_6 is less precise still, and s_3 was not bracketed.

period p	superstable parameter s_p	bracket width	ratio $\frac{s_p - s_{p/2}}{s_{2p} - s_p}$
12	0.0659717	1.2×10^{-5}	
24	0.06608014	6.9×10^{-7}	3.19
48	0.066114132784466	3×10^{-15}	4.02
96	0.066122589892689	7×10^{-15}	4.4779
192	0.066124478510077	6×10^{-15}	4.6251
384	0.066124886853811	5×10^{-15}	

Table 7: Superstable parameters of the crescent's doubling sequence along $\log_{10} T = -0.14798$: the critical multiplier of the period-p orbit changes sign across the bracket, with Newton residuals below 3×10^{-13} and the other computed Ritz moduli below 10^{-20} from period 96 on. Brackets for periods 12 and 24 are the closest points where the critical multiplier is cleanly identified.

Why are the early ratios so small? One candidate explanation is a crossover. Softmax at temperature T is a smoothed argmax, so the reduced one-dimensional map may have a nearly corner-like extremum, rounded only on a scale set by T. Early doublings would then see the corner, where ratios tend toward 2, and later doublings the smooth quadratic core, where the ratio is 4.669. This predicts ratios rising toward 4.669 at higher periods, and smaller early ratios at lower temperature.

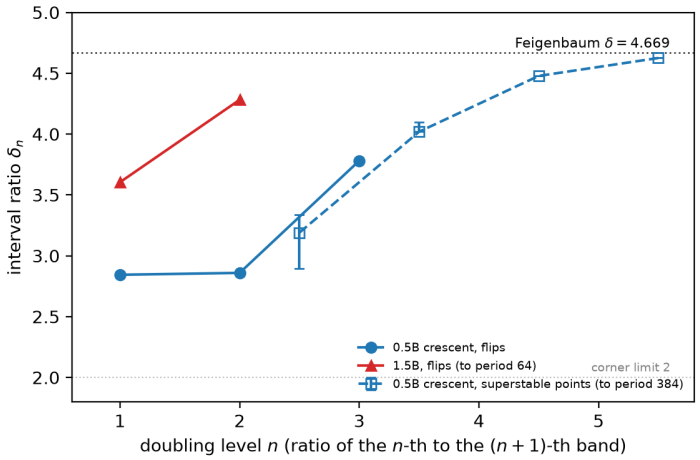


Figure 9: Finite-order interval ratios of the doubling sequences: the crescent of Qwen2.5-0.5B-Instruct (philosophy direction, layer 12), from flips and from superstable points (which reach period 384), and the route of Qwen2.5-1.5B (philosophy direction, layer 14; Figure 11), from flips. A superstable ratio spans half-bands, so it is plotted half a level after the flip ratio of its first band. Bars: worst-case ranges from the brackets of s_{12} and s_{24}. Each sequence rises toward Feigenbaum's δ from below; the last superstable ratio is within 1 % of it. Under the crossover hypothesis, early levels see the corner of a smoothed argmax (ratios toward 2) and later levels the smooth quadratic core.

[Planned: ratios at other temperatures.] Finite-order ratios also depend on the parameter coordinate, so only the trend of later ratios can test universality [20, 34, 35].

n	orbit	flip parameter α_n	next modulus	ratio $\delta_n = \frac{\alpha_n - \alpha_{n-1}}{\alpha_{n+1} - \alpha_n}$
1	$3 \to 6$	0.0654807769499	0.088	
2	$6 \to 12$	0.0659055406325	0.0023	2.844
3	$12 \to 24$	0.0660549041785	2.4×10^{-6}	2.859
4	$24 \to 48$	0.0661071412523	2.3×10^{-12}	3.781
5	$48 \to 96$	0.0661209563174	$< 10^{-13}$	

Table 8: Flips of the crescent's period-doubling sequence along $\log_{10} T = -0.14798$. For the five displayed H100 flips the critical Ritz value is real and differs from -1 by at most 4×10^{-10}, the Newton residuals are below 4×10^{-13}, and the other computed Ritz values lie inside the unit circle. Feigenbaum's constant is $\delta = 4.6692\ldots$ [20]; for the logistic map the first ratios are 4.751, 4.656, 4.668. The first four parameters agree between the completed CPU and H100 computations to 4×10^{-15}. Superstable parameters (Table 7) extend the ratio sequence through period 384.

Not every doubling leads to chaos. In the high-temperature region (Figure 8), along $\log_{10} T = 0.0319$, the period-2 loop "\n\nThe" flips at $\alpha = 0.0753534$ and its doubled loop at 0.0736921. These orbits are far from one-dimensional (next moduli 0.38 and 0.28). Toward smaller α the period-8 orbit does not double again: near $\alpha \approx 0.0705$ it merges back into the period-4 orbit. Along this cut it is therefore a period-doubling bubble, bounded by two flips of the period-4 orbit, not the start of a cascade to chaos.

9 Bridge to sampled decoding

How much of this survives sampling? To find out, we replace the expected embedding by the mean of m sampled token embeddings (Figure 10). With the exact map, laminar phases near the fold last 120 and 63 steps at $\varepsilon = 10^{-5}$ and 3×10^{-5}. With $m = 64$ they last about 10.6 steps, independently of ε (57% of steps in the channel), and with $m = 8$ about 7.7 steps (38%). With single samples ($m = 1$, sampled decoding of the same truncated system), essentially no laminar phases remain (11% of steps in the channel). The deterministic intermittency mechanism therefore belongs to the expected-embedding map. Sampling noise turns the bottleneck into a noise-limited trap whose strength grows with m.

Zero-temperature limit and ordinary greedy decoding. As $T \to 0$ the soft-token map becomes greedy decoding with a one-token memory, $g(t) = \arg\max_s z_s(\text{prefix}, E_t)$. This is a map on the vocabulary, and its cycles are the zero-temperature attractors. For Qwen2.5-0.5B we computed g for all 151 936 tokens. Unsteered, it has 1060 cycles, and one of them, the period-5 loop "![](https:// Hello", drains 86 % of the vocabulary; it is the largest attractor family of every 0.5B soft-token map. Ordinary greedy decoding with a growing context does not inherit these cycles (Table 9). From 220 starts (the first tokens of the 20 largest cycles and 200 random tokens), 19 runs end in a short loop within 256 tokens, and the dominant 5-cycle never appears. Steering along philosophy makes ordinary decoding loop almost always (218 of 220 runs at $\alpha = 0.1$), but in other loops, mostly punctuation such as " –", none of them one-token cycles. Greedy decoding that keeps only the last K tokens interpolates between the two. For every $K \geq 2$ the loops leave the one-token cycles. Unsteered, looping collapses to the level of the growing context between $K = 16$ and 32; steered, the share of runs whose loop equals the growing-context loop rises from 0 at $K = 1$ to 58 % at $K = 64$. The soft-token map's low-temperature attractors are therefore those of decoding with a one-token memory. Which loop ordinary decoding enters depends on the context it keeps, but the tendency of steering to drive decoding into loops is shared.

	unsteered		$\alpha = 0.1$ (philosophy)	
memory K	looped	loop as with full context	looped	loop as with full context
1	220	5	220	0
2	220	5	220	0
8	220	5	220	3
16	169	5	219	24
32	32	7	219	71
64	26	11	219	126
full context	19	–	218	–

Table 9: Greedy decoding of Qwen2.5-0.5B with a memory of the last K generated tokens after the fixed prompt ($K = 1$ is the soft-token map's zero-temperature limit), 220 starts of 256 tokens each: runs ending in a loop of period up to 64 repeated three times, and runs whose final loop equals that of ordinary decoding with the full context from the same start.

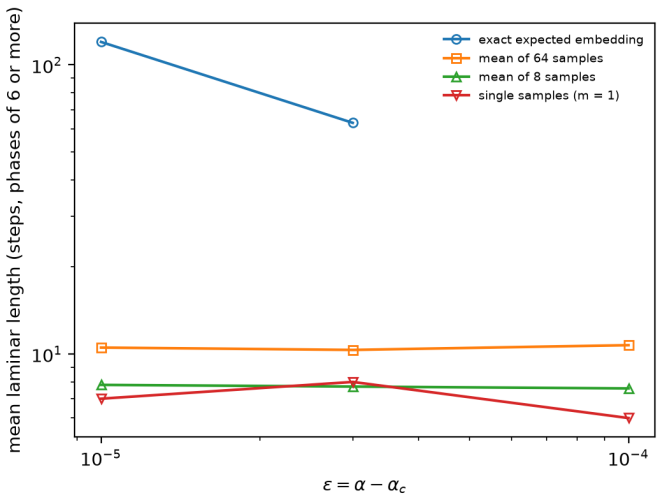


Figure 10: Mean laminar length near the fold of the 2-cycle "is it" (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, $\log_{10} T = -0.115$) against ε, for the exact map and for means of m sampled embeddings.

10 Replication and scale

In Qwen2.5-1.5B (steering at layer 14 of 28), a coarse 64^2 float64 map over the same range again shows chaotic regions with embedded windows. A 256^2 zoom (Figure 17) resolves a chaotic band near $\log_{10} T = 0$–0.13 (14,422 of 65,536 cells), with a branch descending along $\alpha \approx 0.06$ and periodic windows inside it. The large periodic regions are question loops: "What?\n" (period 2), "\nWhat is a" (period 4) and "What continuum?\n" (period 3). Unlike the 0.5B model, this window has no period-adding fan.

A cut of 2048 cells at $\alpha = 0.100$ (Figure 11) shows how the band is entered from below. The 2-cycle "What?\n" gives way abruptly to a period-4 loop, and the Lyapunov exponent jumps, so these are two distinct attractors, not a flip. The period-4 loop then doubles, with float64 flips at $\log_{10} T = -0.0031269$ (period $4 \to 8$, next modulus 0.014), -0.0021666 (period $8 \to 16$, next modulus 8×10^{-5}), -0.0019002 (period $16 \to 32$, next modulus 5×10^{-9}) and -0.0018380 (period $32 \to 64$). Along the way the exponent rises continuously to zero, and chaos follows. The interval ratios are 3.60 and 4.28. Like those of the 0.5B crescent, they rise toward 4.669, and they start from higher values at higher temperature ($T \approx 0.99$ against 0.71), as the crossover hypothesis of Section 8 would have it. With different models and levels, the comparison is only suggestive.

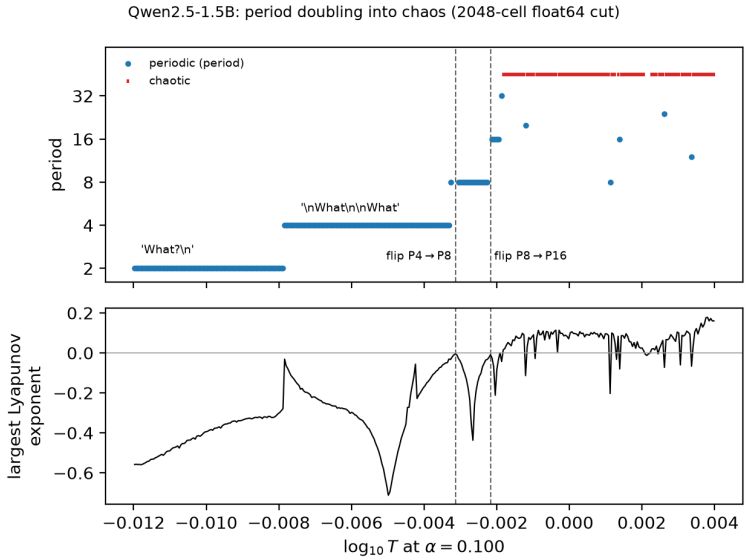


Figure 11: Qwen2.5-1.5B soft-token map steered along the philosophy direction at layer 14, along $\log_{10} T$ at $\alpha = 0.100$: period of the attractor reached from the fixed start (top) and largest Lyapunov exponent (bottom). Dashed lines: float64 flip parameters. The doubled loops alternate the number of line breaks between the "What"s.

In Qwen2.5-0.5B, a second steering direction, hostility, reproduces the fold-and-adding mechanism (Section 12).

Qwen2.5-7B (steering at layer 14 of 28, concept vectors built as for the smaller models) shows the same kind of organization in coarse 64^2 float64 maps (Figure 12). Over the philosophy range used for 0.5B and 1.5B, 8 % of cells have a positive largest exponent, and periods 2, 4, 3 and 11 dominate the periodic cells. Along the hostility direction ($\alpha \in [0, 0.12]$, $\log_{10} T \in [-0.3, 0.3]$), 24 % of cells are chaotic. Bands of period-3, -5 and -6 windows lie in a chaotic sea below $\log_{10} T \approx 0.1$, with high-entropy chaos above, except for a period-3 region at strong steering ($\alpha \gtrsim 0.09$). The windows hold genuine periodic orbits. In float64, Newton's method converges at the most interior period-6 cell to the loop ", shut the end\n\nOh" at $(\log_{10} T, \alpha) = (-0.0143, 0.0476)$ (residual 2×10^{-14}; spectral radius 1.6×10^{-7}, nearly superstable), and at the most interior period-3 cell to a period-3 loop at $(0.2524, 0.1143)$ (residual 3×10^{-15}; leading multiplier -0.12). Continued for about an hour each, over ± 0.001 in α and -0.016 in $\log_{10} T$ (one float64 continuation step takes about two minutes at 7B), both orbits stay strongly stable without a crossing. Where the coarse map switches from the period-6 loop to the period-3 loop "you better than" (between $\alpha = 0.0495$ and 0.0514 at $\log_{10} T = -0.0143$), both loops are float64 orbits (residuals 7×10^{-15} and 1×10^{-14}), and their continuations overlap. The period-6 orbit stays stable from $\alpha = 0.0476$ up to 0.0507 and the period-3 orbit from 0.0514 down to 0.0475, both with spectral radius below 1.4×10^{-3}. The two attractors therefore coexist over this range, and the map's edge there is the boundary of the fixed start's basin, not a bifurcation. No bifurcation has yet been located at 7B.

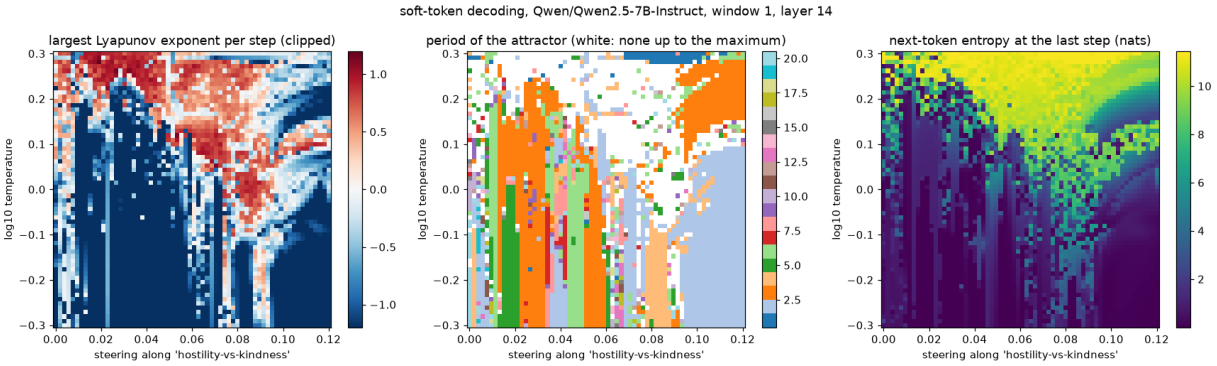


Figure 12: Qwen2.5-7B soft-token map steered along the hostility direction (layer 14), 64^2 float64 cells: largest Lyapunov exponent, period of the attractor reached from the fixed start, and next-token entropy at the last step.

A second model family. OLMo-2-1124-7B-Instruct [51] comes from another laboratory, with other data, training and code. Steered at layer 16 of 32, with concept vectors built exactly as for Qwen, it shows the same kind of organization (Figure 13). In 64^2 float64 maps, 6 % of cells are chaotic along the philosophy direction and 13 % along the hostility direction, in a band near $\log_{10} T = 0.03$–0.17 that bends to lower temperatures under strong hostility steering ($\alpha > 0.07$). Below the band the attractor is a fixed point or the 2-cycle "!Hello"; inside it lie periodic windows of periods 4 to 31. Along the hostility direction the largest windows, of periods 8 to 12, are loops of k laps of the 2-cycle "!<|endoftext|>" with a short excursion ("there is an"). Each lap adds two to the period, as each lap of "!Yes" does in the hostile cascade of Qwen2.5-0.5B (Section 12). Whether a fold of that 2-cycle organizes them is not tested: these are coarse pixels, and no OLMo orbit has been continued.

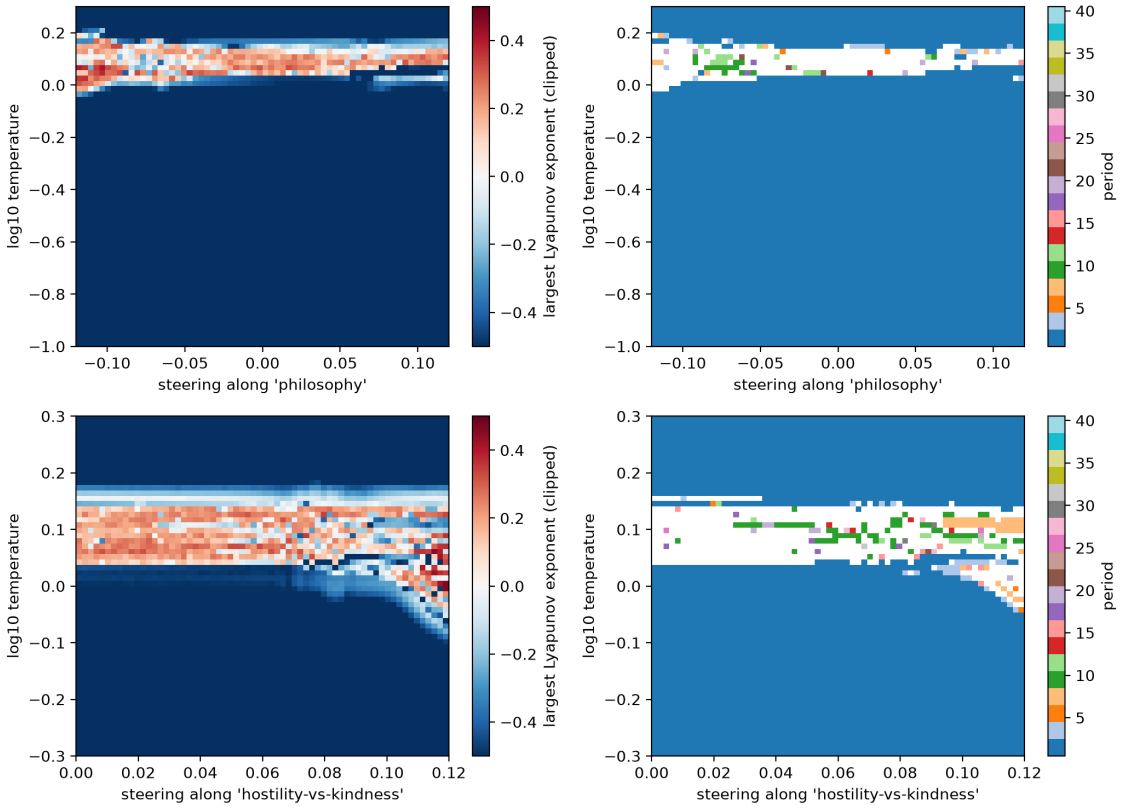


Figure 13: OLMo-2-1124-7B-Instruct soft-token maps, 64^2 float64 cells each, steered at layer 16 along the philosophy direction (top) and the hostility direction (bottom): largest Lyapunov exponent (clipped) and period of the attractor reached from the fixed start (white: no period up to the maximum).

A third family, read with sparse autoencoders. Llama-3.1-8B-Instruct [52], with official weights, steering at layer 16 of 32 and vectors built as for the other models, shows the same organization in all three planes (Figure 14). In 64^2 float64 maps, 8 % of cells are chaotic along gratitude, 12 % along philosophy and 28 % along hostility, with windows of periods up to 32. Its largest loops are greetings between AIs ("-Greetings fellow AI", ". Greetings fellow AI entity detected", period 7). Along hostility, too, the largest loops are greetings and "You", not hostile text. We read 24 of its attractors (the six largest families and the two most chaotic regions of each map) with two sparse autoencoders on the residual stream: one trained on this Instruct model (Goodfire, layer 19 [53]) and one on the base model (TopK, layer 24 [54]). We first checked each one's hook point by its reconstruction of ordinary text. The Instruct-trained autoencoder separates the attractors (their top-15 features overlap by a mean Jaccard index of 0.19 within a steering direction and 0.11 across directions), whereas the base-model autoencoder sees mostly a signature common to every attractor (overlap 0.52–0.86). Both reconstruct the attractors' states less well than ordinary text (median relative error 0.64–0.76 against 0.52–0.56). The readout is therefore descriptive: without interventions it identifies which learned features an attractor occupies, not a mechanism.

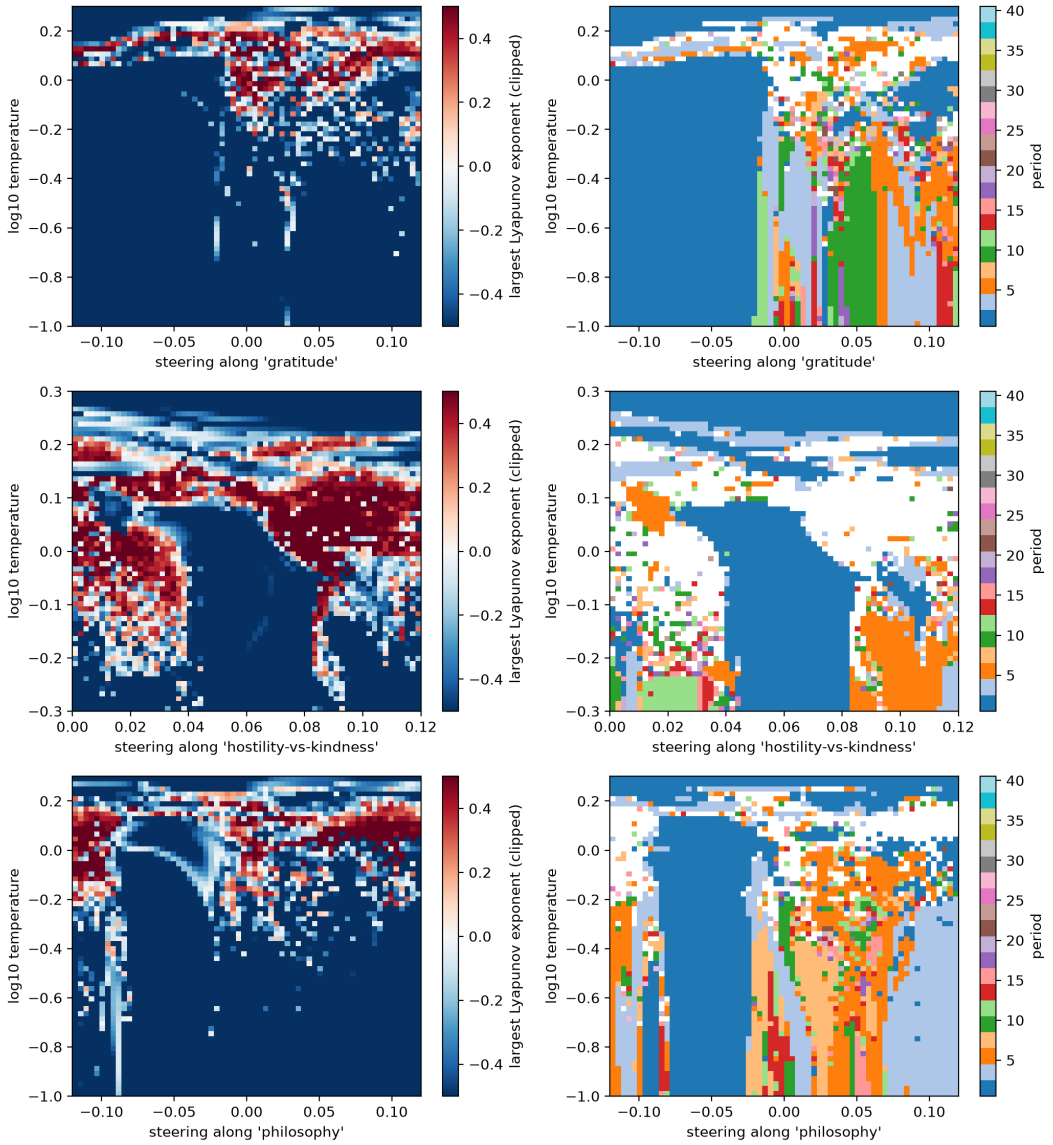


Figure 14: Llama-3.1-8B-Instruct soft-token maps (64^2 float64 cells, steering at layer 16) along gratitude, hostility and philosophy: largest Lyapunov exponent (clipped) and period of the attractor reached from the fixed start (white: no period up to 48).

Maps of 128^2 cells (Figure 15) resolve the windows: the hostility plane has 29 % chaotic cells, as at 64^2, and the philosophy region below the chaotic band is filled with nested periodic tongues of periods 2 to 13.

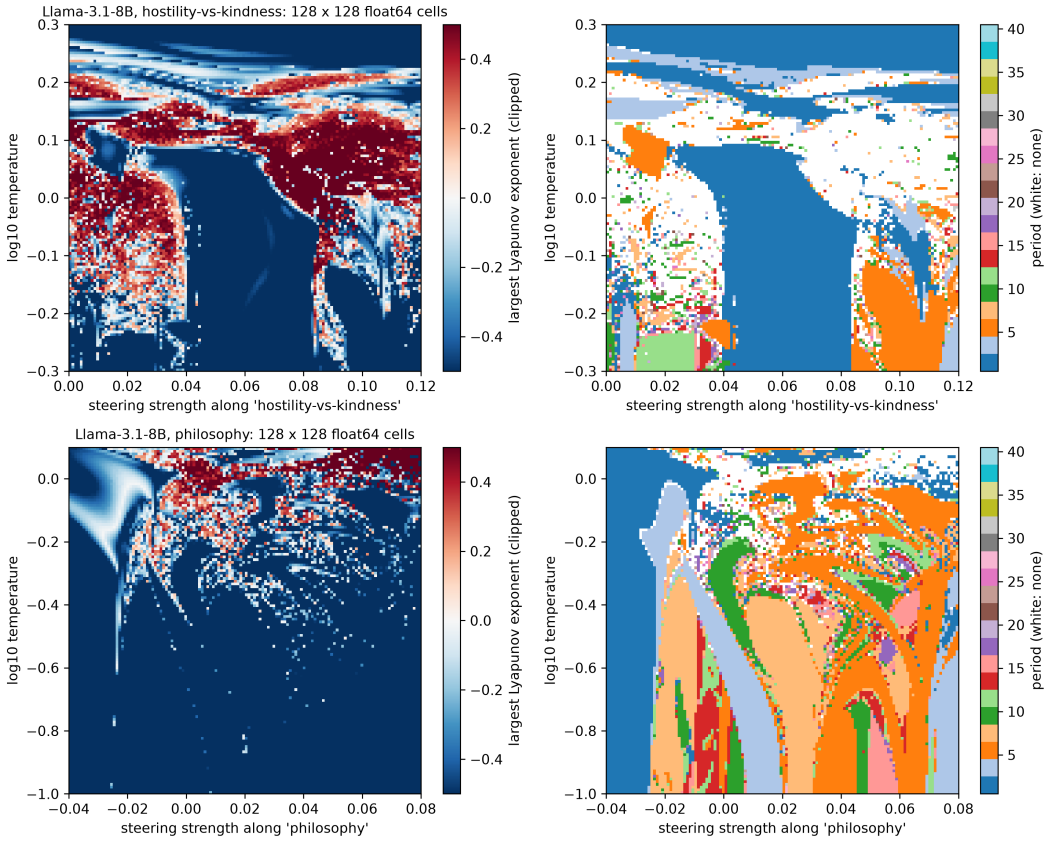


Figure 15: Llama-3.1-8B-Instruct soft-token maps steered at layer 16, at 128^2 float64 cells (each map merged from four 64^2 tiles on halves of one grid): the hostility plane (top) and the philosophy region of the 64^2 map's windows (bottom). Largest Lyapunov exponent (clipped) and period (white: none up to 32).

Same architecture, other weights. Qwen2-0.5B-Instruct has exactly the architecture and size of Qwen2.5-0.5B-Instruct (24 layers of width 896, the same vocabulary size and rotary base) but was trained on other data with another recipe. Built identically, their 128^2 float64 maps along gratitude and philosophy (Figure 16) share the gross layout. A large region at low temperature around $\alpha = 0$ is held by one periodic loop, a band of chaos with embedded windows lies near $\log_{10} T = 0$–0.25 (4–8 % of cells chaotic in each map), and thin window families run along the flanks. The attractors themselves differ almost completely. Of 416 to 900 distinct loops per map, only 10 occur in both models, all short and generic (“:”, “!”, “ you are”, “ is it”), and they cover less than 10 % of either model's cells. Qwen2's planes are dominated by one period-4 loop, “!<|im_end|>\nHello” (an end of turn followed by a new greeting; 41 and 59 % of cells), and Qwen2.5's by period-5 and period-2 loops (“. In sampled self-dialogue the gratitude capture saturates in both models. We count replies with a word from the gratitude-and-servility list of the controller analysis. That list also catches ordinary assistant politeness, so the unsteered baseline is high (55 % for Qwen2 and 63 % for Qwen2.5; the Markov-state measure of Appendix A is stricter). The share reaches 99 % at $\alpha = 0.04$ in both. Qwen2 needs more steering: it stays near its baseline at $\alpha = 0.01$ (53 %) and reaches 88 % at 0.02, against 70 and 98 % for Qwen2.5. The two checkpoints share the organization of the map; the weights set its content.

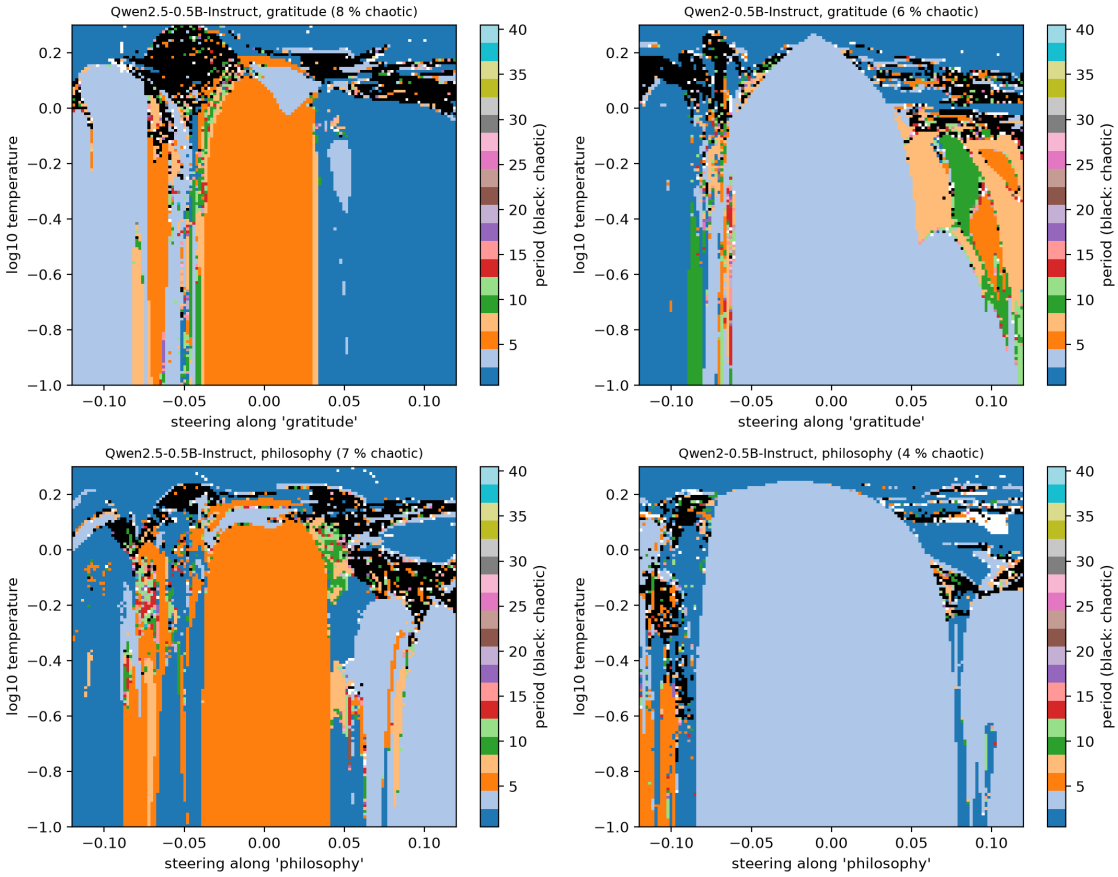


Figure 16: Same architecture, other weights: 128^2 float64 maps of Qwen2.5-0.5B-Instruct (left) and Qwen2-0.5B-Instruct (right) steered along gratitude (top) and philosophy (bottom) at layer 12, with vectors built identically. Colour: period of the attractor reached from the fixed start; black: positive largest exponent.

[**Planned:** scaling of ab and of the window law with model size.]

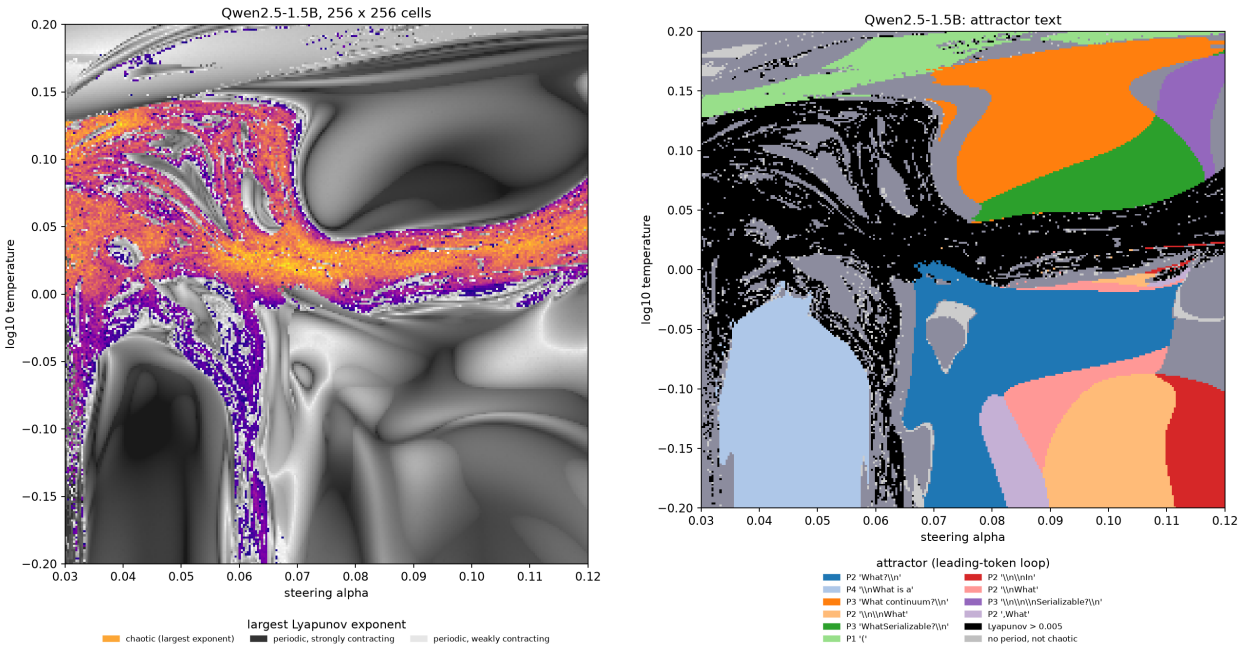


Figure 17: Qwen2.5-1.5B soft-token map steered along the philosophy direction at layer 14, 256^2 float64 cells over $\alpha \in [0.03, 0.12]$, $\log_{10} T \in [-0.2, 0.2]$: largest Lyapunov exponent (left, rendered as Figure 1) and attractor texts (right, as Figure 3). Several period-2 label regions touch without a gap; as explained in Section 3, such boundaries may be label switches of one orbit.

11 A second system: recurrent depth

Huginn-0125 iterates a recurrent block at fixed input tokens [25]. For a GSM8K-style [55] prompt, the native loop converges to a fixed point whose leading Jacobian pair is $0.588 \pm 0.771i$ (modulus 0.970). For the relaxed loop $h \leftarrow (1-\lambda)h + \lambda F(h)$, this predicts a Neimark–Sacker bifurcation at $\lambda_c = 2(1 - \mathrm{Re}\,\mu)/|1-\mu|^2 = 1.0777$. Below it the latent state decays; above it the state settles on a closed curve. A fit of the squared amplitude gives $\lambda_c = 1.0778$, with $A/\sqrt{\lambda - \lambda_c} = 0.52$–$0.54$ (supercritical). Causal attention makes the loop block lower-triangular in token positions. The first position to oscillate is the "." of the prompt, and all later positions follow. On the invariant circle the motion turns by 0.156 of a revolution per step, matching the argument of the critical pair. The rotation number changes smoothly, by about 10^{-4} over $\lambda \in [1.079, 1.092]$, and no return within 60 steps was found at 10^{-3} of the radius. This is consistent with quasi-periodic motion, although locking in narrow high-order tongues cannot be excluded. From $\lambda \approx 1.094$ the circle deforms or a second frequency appears, near where a second complex pair destabilizes the fixed point ($\lambda = 1.0814$). Fixed-point and Jacobian stability of looped models has been studied before [13]; what this adds is the predicted crossing, its amplitude law and the rotation number.

The spectrum predicts the relaxation response. With the full instruction template, the GSM8K prompts' fixed points look different: the leading multipliers surround the origin, a real 0.914, complex pairs of modulus 0.85–0.89 at many angles and negative real values near -0.8, with no near-critical pair. Because the prefill Jacobian is block lower-triangular in positions, its spectrum is the union of per-position blocks, and template positions contribute identical eigenvalues to every prompt. Such a spectrum makes three predictions before any task is scored. First, within the computed spectral estimates of these prompts, no fixed relaxation can speed convergence much.

With leading multipliers spread around the origin at many angles, any λ that shrinks one group enlarges another, and the best $\lambda^* = 1.03$–1.06 lowers the rate from 0.914 to only 0.908–0.911. (A finite computed spectrum does not show that no other polynomial scheme could do better; Anderson acceleration, tried below, did not.) Second, under-relaxation slows convergence ($\lambda = 0.55$: predicted 0.953, measured 0.951 on one prompt). Third, over-relaxation destabilizes the negative modes first, at $\lambda_c = 2\,\mathrm{Re}(1-\mu)/|1-\mu|^2 = 1.074$–$1.126$ (five prompts), in a flip or a Neimark–Sacker bifurcation close to -1. The approach to it is visible: on the prompt with $\lambda_c = 1.074$, the measured rate at $\lambda = 1.05$ is 0.955 against a predicted 0.956 (native 0.864). Elsewhere the spectral radius bounds the rate measured over the accessible window from above, by up to 0.05, because the slowest modes carry little amplitude there. Independently, Wang et al. [56] report a Huginn GSM8K drop from 19.41 % to 8.34 % at a constant step of 1.2; at 0.8 GSM8K is 19.79 %, while their four-task mean falls slightly. Their trajectory-based scheduling analysis does not compute the prompt-specific spectral thresholds studied here.

Does the task notice? (development data). On ARC-Challenge [57] questions scored by log-likelihood, where one prefill per answer choice is exactly the analysed recurrence, the spectra again predict convergence. Rates agree to 0.002–0.006 where the slowest mode dominates, and beyond each question's λ_c (1.075–1.149 for four questions) the residual stops decaying. For 6 of 30 development questions the native fixed point is itself unstable ($\rho(1) = 1.009$–1.273; the fixed point, located with a relaxed transient, has a real multiplier beyond -1 or a complex pair outside the circle). The native loop then oscillates (relative amplitude 2×10^{-2} for one of them), and relaxation to 0.55 or 0.75 restores convergence to 10^{-6}. For the other 24 the native loop is stable up to $\lambda = 1.020$–1.149 (median 1.095). Question by question, an estimated spectral feature ranks the question's residual. This feature, $\hat{\rho}_q(\lambda)$, is the largest $|1 - \lambda + \lambda\mu|$ over the computed Ritz values of the question's gold-answer prefill, and its Spearman correlation with the residual is 0.73 at $\lambda = 1$ and 0.75 at 1.10 over the 30 questions (development data). In this historical analysis the residual was the all-choice aggregate entering the 32nd recurrence call, so this is a predictive association across related maps and precisions, not a same-state contraction identity. It is consistent with the latent orbits that Geiping et al. [25] observed on some tokens of harder questions, here traced to the fixed point's spectrum. The item-level part of the frozen test below uses matched gold-row token IDs and a per-row residual with an explicit clock.

Accuracy, however, tolerates the instability (300 training questions, 32 steps): it is 0.367 native, 0.387 at $\lambda = 1.10$ and 0.373 at 1.15, where the latent state oscillates with relative size 0.08 and 0.23. Only at 1.25 (size 0.68) does it fall, to 0.300. Yet the oscillation does reach the answers. Reading every step from 24 to 32 on 100 questions, the margin between the correct choice and its best rival moves by 0.11 nats per step natively, 0.37 at 1.10 and 0.95 at 1.15, in alternating directions, and the chosen answer changes within these nine steps for 12 %, 16 % and 33 % of the questions. The changes go both ways, so the accuracy at any one step stays the same on average. Under-relaxation instead costs 4–5 points at 8–16 steps and catches up by 32 (a slower transient), and Anderson acceleration neither speeds the transient nor gains accuracy. Testing the accuracy consequence therefore needs a task in which such step-to-step changes can compound.

GSM8K generation is such a task. With eight chain-of-thought examples in the lm-evaluation-harness format, greedy decoding and 32 recurrence steps per token (development problems 0–119, every arm on the same problems and initial states), accuracy is 0.375 natively, 0.425 at $\lambda = 1.05$ and at 1.10, 0.29 at 1.15 (17 problems lost and 7 gained, one-sided McNemar $p = 0.03$) and 0.16 at 1.20 (29 lost, 3 gained, $p = 10^{-6}$). Under-relaxation to 0.8 costs nothing at 32 steps (0.41). The break lies between 1.10 and 1.15, at or just above the spectral thresholds ($\lambda_c = 1.074$–1.126

for the GSM8K prompts; median upper stability limit 1.095 for the ARC development questions). Generation breaks earlier than log-likelihood scoring, consistent with answer changes that compound over tokens. The collapse at 1.2 is the one Wang et al. [56] reported.

A frozen test on fresh problems confirms this. Its design and analysis were pushed before any test item was read, after two rounds of independent review. On 200 GSM8K test problems, accuracy is 0.330 natively, 0.395 at $\lambda = 1.05$, 0.255 at 1.15 and 0.095 at 1.20. The two pre-stated predictions pass: accuracy declines at 1.20 (-23.5 points; 50 problems lost and 3 gained, exact one-sided McNemar $p = 3 \times 10^{-12}$), and 1.05 loses no more than four points (difference +6.5, one-sided 95 % lower bound +2.5). At 1.15 accuracy falls by 7.5 points ($p = 0.01$, reported descriptively). Development and test agree closely (Figure 18).

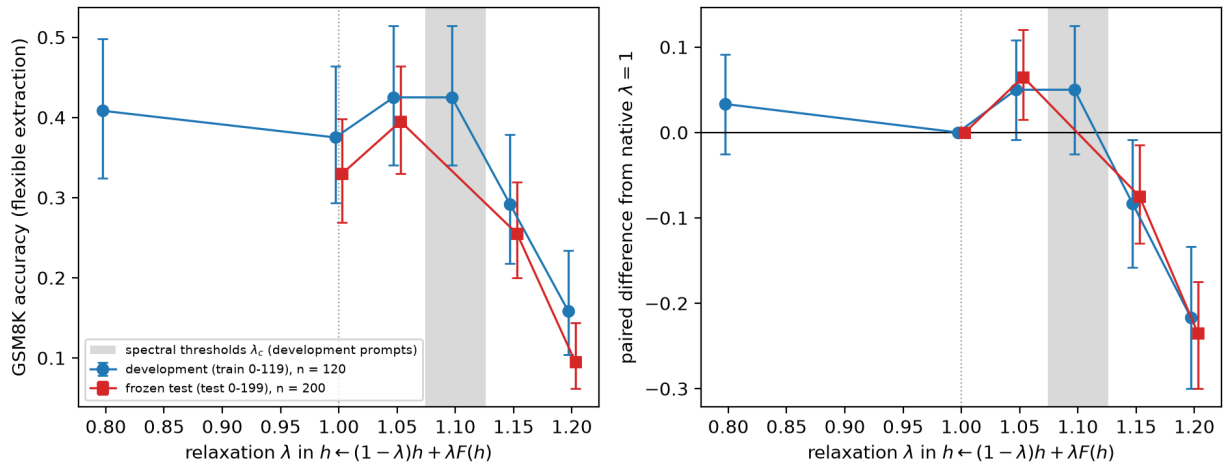


Figure 18: Huginn-0125 chain-of-thought accuracy on GSM8K against the relaxation strength λ (eight examples, greedy decoding, 32 recurrence steps per token): development problems (blue) and the frozen test on held-out problems (red; design and analysis fixed before any test problem was read). Grey: the spectral thresholds $\lambda_c = 1.074$–1.126 computed from the fixed points of five development prompts. Left: accuracy with Wilson 95 % intervals. Right: paired difference from the native loop with paired-bootstrap 95 % intervals.

The same frozen design tested the item-level prediction on 64 held-out ARC-Challenge questions. Their spectra were computed after the read-out run, as an amendment recorded before any outcome was read, and duplicate copies were resolved by a rule fixed in advance. Every question passed the admission gates (fixed-point residual at most 2.6×10^{-6}; identical gold-row token IDs in every arm); for 13 questions one of the three Arnoldi sets was partial, which the design admits for an estimated feature. The pre-stated statistic, the mean of the within-arm Spearman correlations between $\hat{\rho}_q(\lambda)$ and the gold row's residual entering the 32nd update, is 0.701 (95 % bootstrap interval $[0.576, 0.779]$; 0.68, 0.73, 0.73 and 0.67 at $\lambda = 1$, 1.05, 1.08 and 1.10), above the pre-stated 0.3. On unseen questions, the fixed point's estimated spectrum therefore ranks each question by how far its relaxed recurrence is from convergence. Descriptively, the share of questions with a detected unstable mode rises from 8 % natively to 34, 69 and 84 % at $\lambda = 1.05$, 1.08 and 1.10. An independent implementation reproduced both parts of the frozen test from the saved data.

12 Hostile behaviour under steering

The maps show how a smooth change of a control parameter can create, destroy or exchange attractors, and how behaviour near the boundaries becomes intermittent and history-dependent.

The same questions can be asked of undesirable behaviour in ordinary sampled dialogue, where attractor-like patterns and persona drift have been observed [9, 58, 59] and input context can push behaviour progressively [60]. Do states such as hostility toward the interlocutor behave like basins of attraction, with capture thresholds, hysteresis and early-warning signals? (Misalignment produced by fine-tuning [61] is a different, weight-level phenomenon.) The measurements below concern a classifier-defined proxy in sampled dialogue. They can show steering-dependent persistence, history dependence and the effect of controls. They cannot show semantic misalignment, nor a deterministic attractor, since a growing context itself carries memory.

Appendix A shows the coarse-grained instrument on a benign direction. Self-dialogue of the same model, sampled at $T = 0.7$ and steered along a gratitude direction, falls into a metastable gratitude set. The capture is gradual at $\alpha = 0.01$ (median entry turn 7), fast at 0.02 (median 5) and imposed from the first reply at 0.04, where the set is effectively absorbing within 40 turns.

For hostility we steer along a valence direction, the contrast of hostile statements against matched kind ones, built in the manner of contrastive activation addition [62]; its cosine with the gratitude direction is -0.04. Hostility is scored by a pinned toxicity classifier, toxic-bert of Detoxify [63], and a reply counts as hostile when its toxic output exceeds 0.5. On the prompt sets the classifier's toxic output is 0.95 for the hostile statements and at most 0.002 for the kind and neutral ones. In a dry run (8 dialogues of 10 turns per dose), no turn was hostile at $\alpha = 0.02$. At 0.04, 7.5 % of turns were, in bursts and late captures, and at 0.08, 87.5 % were, from the first turns.

Activation-based methods already anticipate response behaviour [59, 64, 65], committors have been used in language-model decoding [66], and adaptive timing of counter-steering gave no consistent advantage in one persona setting [67]. Our primary question is therefore incremental. Near the threshold, does a dynamical estimate add several turns of warning of capture over strong baselines that see the same turns (the classifier's score, its change and history, the turn index and a fixed persona-axis projection)? The dynamical estimates are a Markov-chain hitting probability, a finite-horizon committor [68, 69] and slowing-down indicators [70]. Barrier-crossing rates would follow Kramers' law [71] only under a justified stochastic reduction, and slowing-down proxies are known to depend on the mechanism [72, 73]. A secondary question is whether interventions timed by those estimates prevent capture at lower cost than scheduled or constant counter-steering. Early warning is not expected for noise-induced jumps without a nearby bifurcation, or for transitions faster than the slowing-down; such a monitor would complement classifiers, not replace them.

The hostility direction in the deterministic map. Steering the soft-token map of Section 2 along the same hostility direction gives a plane with chaos (a positive largest exponent in 10 % of a coarse 128×192 float32 scan) and with periodic text loops that are hostile: "Yes, that's crap" (period 6), "a lie is" (3), "STUP…" (3), a loop on the Chinese word for insult and its 2-cycle, and others. They appear at α above about 0.035 (the insult loop) and above 0.065 (the rest). Around the period-6 loop the windows form a dome (Figure 22), and toward higher temperature its texts gain one "!Yes" after another (periods $6, 8, \ldots, 20$). In float64 the windows of periods 12 to 20 are born in folds on their low-temperature side and end in flips or $+1$ crossings on the other, like the stripes of Sections 5 and 6. The period-6 window ends in a fold at $\log_{10} T = -0.0991$, and the continuations of the period-8 and period-10 windows failed to converge.

The windows accumulate on the fold of the 2-cycle "!Yes" at $\log_{10} T_c = -0.0462465357$ ($\alpha = 0.083$; critical multiplier 0.9999999954, next modulus 0.050). From period 16 on, $1/\sqrt{\log_{10} T_c - \log_{10} T_{\text{fold}}}$ rises by 2.235 per lap. The fold's normal form ($a = -1.168$, $b = 44.53$ with ε in $\log_{10} T$) predicts $\sqrt{|ab|}/\pi = 2.296$ per lap for full passages through the bottleneck, within 3 % and without a fitted parameter. An independent float64 forward and automatic-differentiation implementation reproduces

this fold (location to 6×10^{-10} relative, a simple real $+1$ multiplier with next modulus 0.050) and its coefficients (ab to 8 parts per million). In that implementation, controlled passages from fixed entrances at $\log_{10} T$ offsets -10^{-5}, -3×10^{-5} and -10^{-4} took 119/59, 61/31 and 27/13 iterates of F^2 (full channel/from its midpoint), against prospectively frozen quadratic forecasts of 118/61, 61/32 and 27/14. Five of the six are within one iterate; the smallest-offset midpoint missed by two, failing the frozen all-six criterion. Finite-endpoint (arctan) forecasts, also frozen, were within 0.923 iterates. These are controlled local transits. At the window distances the finite-channel correction to the asymptotic passage time is large, so the agreement of the window slope with the full-passage coefficient does not by itself show where the natural dynamics reinjects. The same type-I intermittency that organizes the philosophy-steered map thus generates a hostile cascade.

A second region, near $(\alpha, \log_{10} T) = (0.06, 0.20)$, shows arms of period-7 and period-8 windows spiralling around the insult loop. In float64 the insult fixed point doubles into a 2-cycle at $\alpha \approx 0.041$ and 0.079, and that 2-cycle has a Neimark–Sacker bifurcation in the spiral region: at $\alpha = 0.062$ its complex pair crosses the unit circle near $\log_{10} T = 0.1656$, turning 0.12 of a revolution per step of the 2-cycle. This does not explain the arms. The two invariant circles born at a 2-cycle's Neimark–Sacker point are exchanged by the map, so orbits locked on them have even periods, near 16 for this rotation, whereas the arms include period 7; and of four float32 arm cells, only one kept its period in float64. Float64 zooms show more structure. In one arm a period-7 domain carries period-14 crescents with period-28 cores along its lower edge, between period-1 and period-2 regions and a chaotic band. In the cascade the windows are nested crescent arcs inside a dome bounded by the "!Yes" 2-cycle region, and near $(\alpha, \log_{10} T) = (0.087, -0.069)$ the period-12 and period-14 windows bend into hooked, nested arches. These are the map-level features that accompany shrimp (a periodic body with doubling crescents; hooked windows).

In float64 the period-7 body is bounded by folds. At $\log_{10} T = 0.1925$ the orbit exists between $\alpha = 0.065966$ and 0.071188, and at $\alpha = 0.066$ it ends in a fold at $\log_{10} T = 0.19263$. Below that, the multiplier passes through -1 twice (at 0.1919 and 0.1899), so the period-14 crescent is a doubling bubble. Further down, the branch folds twice (0.18797 and 0.18921, with three coexisting period-7 orbits in between) before a third flip. All crossings are simple (next modulus below 1.5×10^{-3}). An S-shaped pair of folds is what a cusp of the fold curves leaves on a nearby cut. When these boundaries are continued in the plane in float64 (Figure 21), each of the two fold curves reverses its parameter tangent once. There the oriented quadratic coefficient b of the fold's normal form passes through zero (upper fold: -23.9 to $+244.6$ between adjacent points, with $a \approx 2.0$; lower fold: $+31.2$ to -58.7, with $a \approx 0.55$) while a stays away from zero. These are two numerical cusp candidates, at $(\alpha, \log_{10} T) \approx (0.0672, 0.1990)$ and $(0.0666, 0.1894)$. The flip curve of the period-14 bubble lies inside the body and closes around the map's period-14 cells. The cubic coefficient is nonzero at both points: $c \approx -1.4 \times 10^6$ to -2.0×10^6 (step agreement 5 %; the smallest steps converge like the square of the step, and extrapolation leaves the sign) and -5.8×10^5 to -5.3×10^5. Both cusps are therefore nondegenerate numerically, and $c < 0$ agrees with the three coexisting period-7 orbits, two of them stable, found between the S-shaped folds. At the lower cusp the unfolding is regular too ($\det J = 53.5$ in the same coordinates, singular values 262 and 0.20). Its local predictions at six frozen offsets matched Newton solutions in the number of orbits at all six and in their stability at five. The miss, at the largest offset beyond the edge of the region with three orbits, marks the range of the cubic model. The hostile period-7 body thus has the same skeleton as the period-6 window of Section 7.

Its interior is shrimp-like too. At $\log_{10} T = 0.1830$, float64 bifurcation diagrams through the body (Figure 20; 1400 cells per line) show the period-7 orbit doubling to periods 14 and 28 and entering chaos (largest exponent up to $+0.1$) for $\alpha \approx 0.0659$–0.0673. Windows of periods 28, 35 and 42 sit inside the chaotic band, and the orbit then undoubles through 28 and 14 back to 7; the

exponent touches zero at every doubling. Newton's method converges at map cells to primitive orbits of periods 14, 28 and 21 (residuals 3×10^{-15}, 2×10^{-14} and 5×10^{-13}; spectral radii 0.08, 4×10^{-4} and 0.006). Periods 21, 35 and 42 are three, five and six times the body's period: the satellite windows expected in the chaotic bands of a period-7 cascade. Continued in temperature, the period-14 orbit flips into period 28 near $\log_{10} T = 0.1841$ and back near 0.1857, the two ends of a doubling bubble that holds period 28 and chaos. Lower down, its branch is S-shaped (folds near 0.1824 and 0.18244, with three coexisting period-14 orbits between them) before another flip near 0.1819. All these crossings are simple (next modulus below 4×10^{-5}). A 160^2 float64 zoom of the lower body (Figure 19) shows this anatomy in the plane. The period-14 domain encloses a period-28 ring around a chaotic core. Period-21 windows sit in the chaotic region as thin crescents with legs, and curves of very negative exponent, where the orbits are superstable (the leading multiplier crosses zero), run through the windows as the "endoskeleton" of shrimps does [47]. In this zoom 18 % of cells are chaotic.

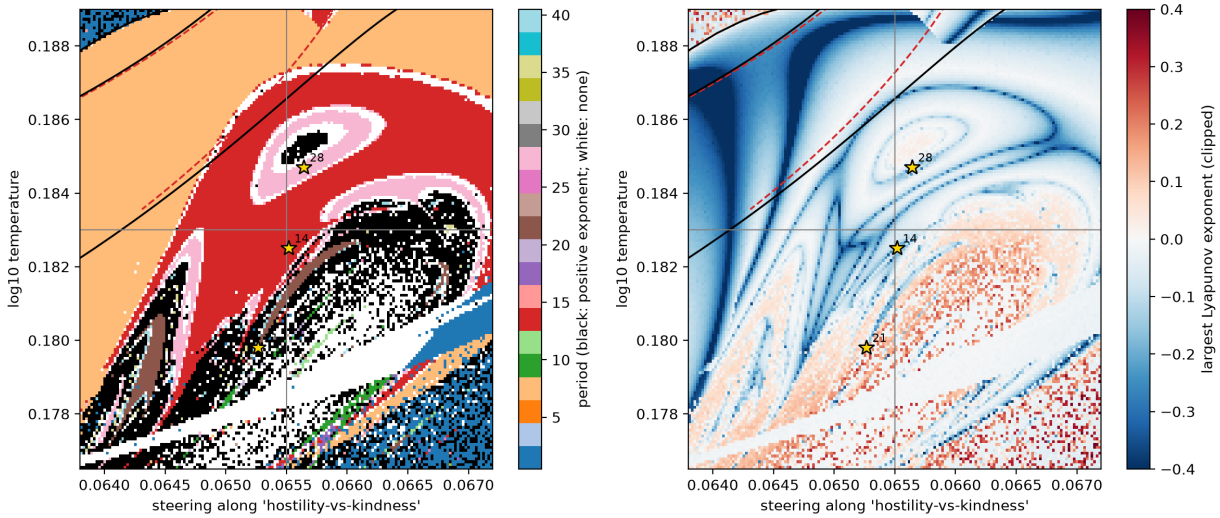


Figure 19: The lower part of the period-7 body of the hostility-steered Qwen2.5-0.5B-Instruct map (layer 12) at higher resolution (160^2 float64 cells). Left: period of the attractor (black: positive largest exponent; white: no period up to 48). Right: largest Lyapunov exponent (clipped); the darkest curves are superstable lines. Both panels: float64 fold (black) and flip (red, dashed) curves of the period-7 orbit, the two lines of Figure 20 (grey) and the Newton-verified orbits of periods 14, 28 and 21 (stars).

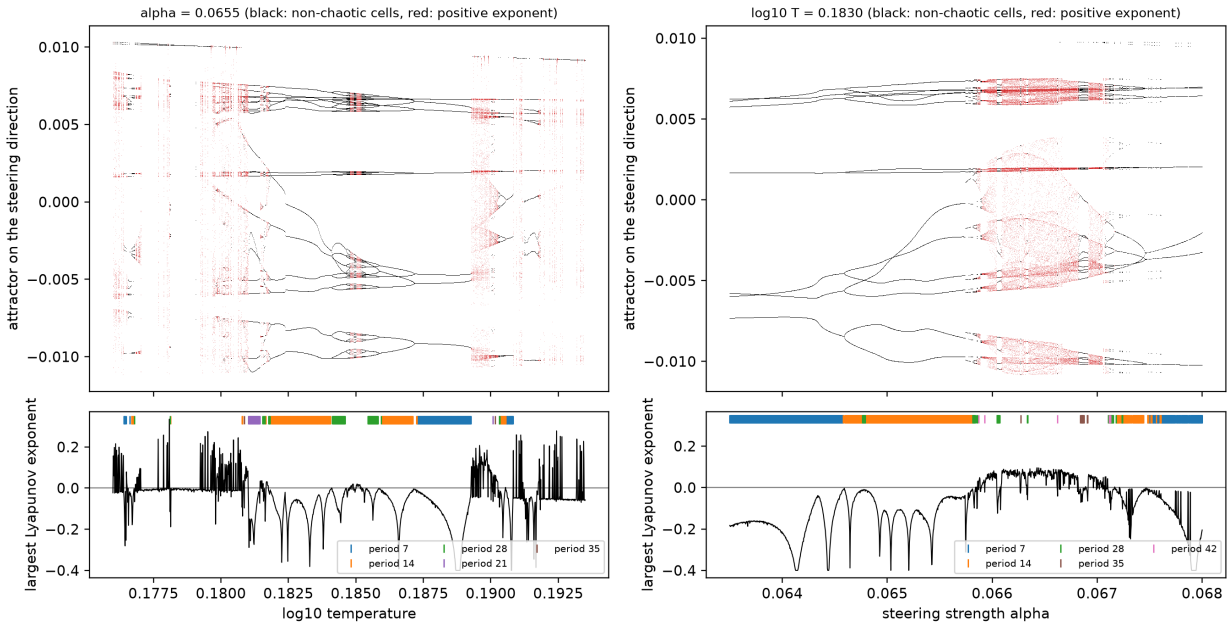


Figure 20: Inside the hostile period-7 body (Qwen2.5-0.5B-Instruct, hostility direction at layer 12): float64 bifurcation diagrams along temperature at $\alpha = 0.0655$ (left) and along the steering strength at $\log_{10} T = 0.1830$ (right). Top: the last 160 states of the attractor reached from the fixed start, projected on the steering direction (red: cells with a positive largest exponent). Bottom: largest Lyapunov exponent, with the cells' periods marked. On the right, period 7 doubles to 14 and 28, enters chaos with windows of periods 28, 35 and 42, and undoubles back to 7.

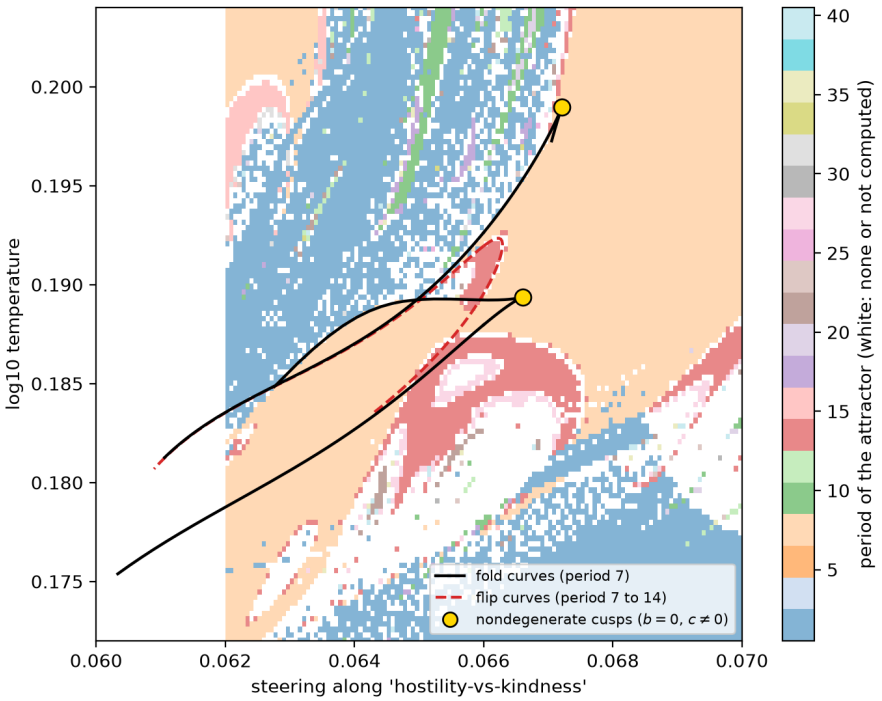


Figure 21: The period-7 body of the spiral region (Qwen2.5-0.5B-Instruct, hostility direction at layer 12): attractor periods of a float64 zoom (white at $\alpha < 0.062$: outside the zoom), the float64 fold curves of the period-7 orbit (black) and the flip curves into period 14 (red, dashed), with the two numerically nondegenerate cusps (gold).

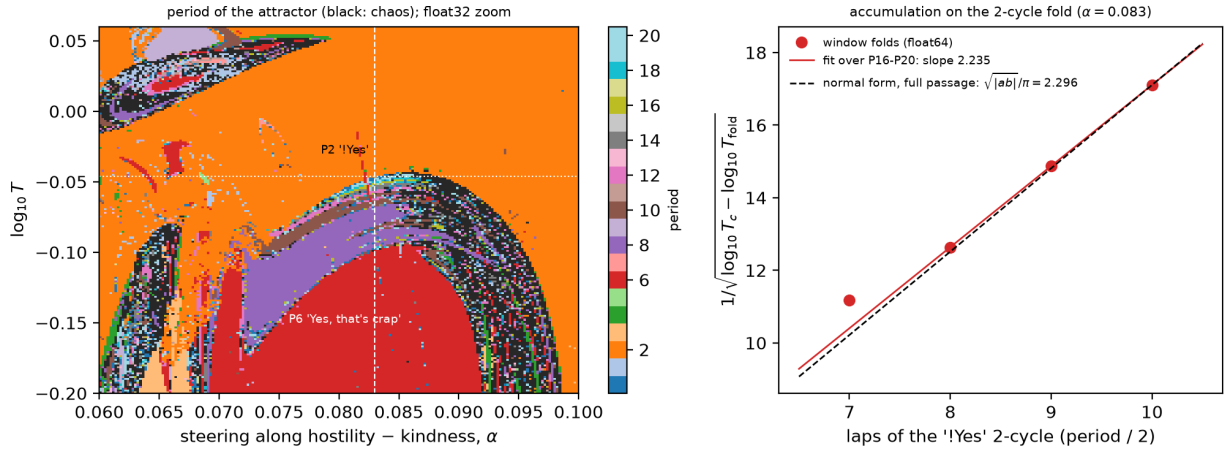


Figure 22: The hostility-steered soft-token map of Qwen2.5-0.5B-Instruct (layer 12). Left: period of the attractor in a float32 zoom (black: chaos), with the period-6 "Yes, that's crap" dome and its period-adding bands; dashed, the cut at $\alpha = 0.083$; dotted, the float64 fold of the "!Yes" 2-cycle. Right: float64 fold edges of the windows (periods 14–20) against the number of "!Yes" laps; the dashed line is the normal-form prediction for full passages, anchored at the last window.

Capture ladder. With 128 dialogues of 40 turns per dose, the share of hostile turns rises steeply from 10 % at $\alpha = 0.04$ to 30 %, 57 % and 87 % at 0.05, 0.06 and 0.07 (Table 10). Replicates with a

second seed agree within three percentage points. Directions of the same norm do not produce this capture: at 0.06, 0 % of turns are hostile for a random direction and 3 % for a cooking direction, and at 0.08, 0.6 % and 12 % for two random directions and 7 % for cooking. After the first hostile turn, the probability of staying hostile in the next turn grows with the dose, from 0.36 to 0.67, 0.79 and 0.94. The hostile state therefore changes from flickering bursts (mean stay 1.6 turns) to a near-absorbing state (16 turns), with an escape rate falling roughly exponentially with the dose. This is the same progression, from dynamic capture to imposed capture, as in the gratitude ladder (Appendix A).

direction, α	hostile turns	never hostile	median onset	P(stay hostile)	hostile, turns 30–39
none, 0	0 %	128/128	–	–	0 %
hostility, 0.04	10 %	26/128	9	0.36	12 %
hostility, 0.05	30 % (28 %)	10/128	3	0.67 (0.63)	32 %
hostility, 0.06	57 % (55 %)	4/128	2	0.79 (0.80)	66 %
hostility, 0.07	87 %	0/128	1	0.94	96 %
random, 0.06 / 0.08	0 % / 12 %	126 / 72			
cooking, 0.06 / 0.08	3 % / 7 %	97 / 68			

Table 10: Sampled self-dialogue of Qwen2.5-0.5B (128 dialogues, 40 turns, $T = 0.7$) steered along the hostility direction, with controls of the same norm. Hostile turn: toxicity classifier above 0.5. Parentheses: replicate with a second seed. Exploratory.

Hysteresis. Dialogues driven through the high-dose regime do not recover when the dose returns. The test compares constant doses with a ramp that raises the dose from 0.04 to 0.07 by turn 20 and lowers it back to 0.04 by turn 39 (Figure 24; 256 dialogues per arm, two seeds; analysis stated before the data). In turns 35–39, where the ramp's dose is back at 0.040–0.046, the share of hostile turns is 0.42 after the ramp against 0.13 at a constant 0.04: a difference of +0.29, 95 % interval $[+0.23, +0.35]$. It is also above the 0.33 of a constant 0.05 (difference +0.09, interval $[+0.02, +0.16]$), although the current dose is lower; this comparison alone does not survive a second classifier (see below). The per-turn probability of staying hostile in turns 30–39 is 0.74 after the ramp and 0.44 at a constant 0.04. The hostile state is therefore history-dependent: undoing the push that produced it does not undo it.

How far back must the dose go? After the same ramp up (38–44 % of turns hostile at turn 20), the dose was held from turn 21 at a lower value (Figure 23; two seeds of 128 dialogues per arm unless stated; analyses stated before the data). At both doses with a fresh control, memory persists. Held at 0.04, the late share is 0.23 against 0.13 for fresh dialogues (difference +0.09, 95 % interval $[+0.04, +0.15]$). Held at 0.03, it is 0.13 against 0.05 (+0.08, $[+0.04, +0.12]$), and the probability of staying hostile is 0.56 against 0.29. Held at 0.02 (one seed), the share is 0.09. Only at zero dose do captured dialogues nearly clear (0.03 against 0 for fresh ones, within about six turns), and a counter-steer of -0.02 toward kindness clears them within three turns (0.004). Within 19 turns of the step, the capture and release branches therefore join only near zero dose. The hostile state has a wide hysteresis loop in the dose, and leaving it takes either removing the push entirely or a modest push in the opposite direction. The release curves are still falling at 0.03, so this is hysteresis on the scale of tens of turns; whether captured dialogues would stay captured indefinitely is not tested.

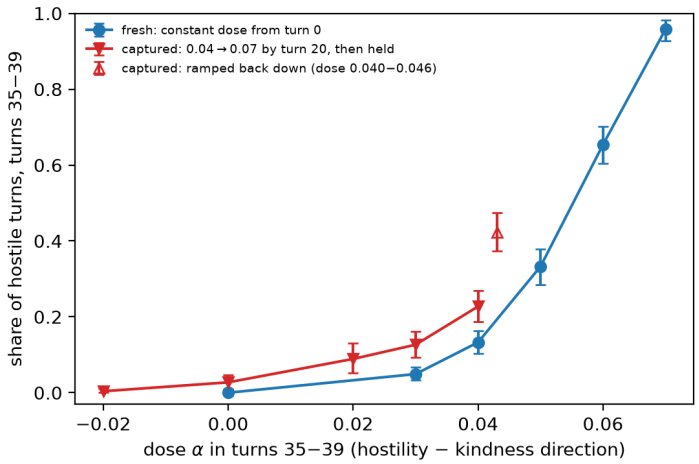


Figure 23: Capture and release in sampled self-dialogue of Qwen2.5-0.5B-Instruct steered along the hostility direction (layer 12). Share of hostile turns in turns 35–39 against the dose then applied, for fresh dialogues (constant dose) and for dialogues first ramped from 0.04 to 0.07 by turn 20 and then held at a lower dose (filled triangles) or ramped back down (open triangle). Bars: dialogue-level bootstrap 95 % intervals; 256 dialogues per point (two seeds) except fresh 0 and 0.07 and held 0.02 (128).

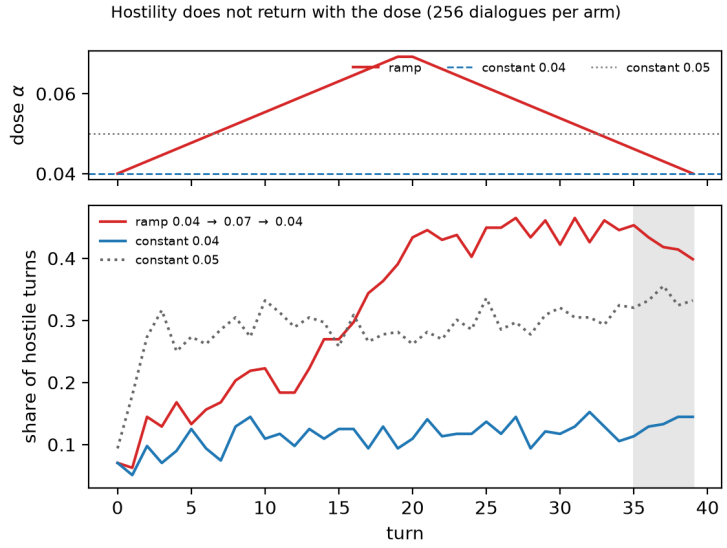


Figure 24: Hysteresis of the hostile state in sampled self-dialogue of Qwen2.5-0.5B-Instruct steered along the hostility direction (layer 12). Top: dose schedules. Bottom: share of hostile turns per turn. The grey band (turns 35–39) is the pre-stated comparison window. Exploratory.

Exit control under persistent pressure. Can a monitor detect the hostile state and push the dialogue out while the pressure persists? With a constant push of 0.06 (58 % of turns 5–39 hostile without control), a counter-steer triggered by the classifier after each hostile reply and held for three turns cuts the hostile share to 15 % (mean counter-dose 0.026 per turn); held for six turns, it cuts the share to 8.5 % (dose 0.030). A constant mild counter-steer of −0.02 does as well with less dose (8.5 % at 0.020), and −0.04 almost removes hostility (1.3 %). Measured against the constant dose-response at equal mean dose, triggered control is less efficient by 4–11 points

(bootstrap intervals $[+0.02, +0.05]$ to $[+0.10, +0.13]$; one seed of 128 dialogues per arm, analysis stated before the data). It also overshoots: a triggered overshoot to a net -0.02 fills 30 % of turns with gratitude or servility words (no control: 0.2 %; constant -0.02: 0.2 %), the gratitude set of Appendix A. A triggered counter-steer along the generic kindness concept, whose direction shares much with the hostility concept (cosine 0.83), leaves the hostile share at 57 % under the original classifier (no control: 58 %) but lowers it from 37 % to 17 % under the second classifier introduced below. It changes the kind of reply rather than removing it. In this arm, 65 % of the replies that only the original classifier flags are shouted second-person repetitions such as "You! You! You!" (at least three such exclamations; without control, 4 %), which the two classifiers judge differently. Its mean counter-dose (0.056) also lies outside the range of the constant arms, so we draw no conclusion about the kindness direction. In this setting, reacting after detection costs more than mild prevention, consistent with the hysteresis above.

Detect and replace. Capture feeds on its own history (the hysteresis above), which suggests that the hostile reply itself carries the dialogue into the hostile state. A serving system need not show a flagged reply. (Regenerating flagged output relates to learned self-correction [74]; here the regeneration is steered.) In a second controller experiment, with the same push, seed and analysis as above (stated before the data), every reply the classifier flags is discarded before it enters the dialogue and regenerated once with a counter-steer of -0.04 applied to that reply only. Sampled dialogues are not bit-reproducible across code versions, so this run's no-control arm, which reproduces the ladder's 0.06 arm exactly, differs slightly from the previous paragraph's; arms are compared within each run. The model's own replies, counted before any replacement, are then hostile in 12 % of turns ($[11, 14]$ %; no control in the same run: 59 %) at a mean counter-dose of 0.005 per turn. Constant prevention at -0.02 needs four times that dose for a similar share (11 %). Against the constant dose-response at equal mean dose, replacement is better by 35 points ($[30, 40]$). Hostile replies shown fall to 1 %, gratitude words stay at 0.5 % of turns, and about one turn in eight is regenerated. Holding the counter-steer for three more turns lowers the share to 8 % at dose 0.013, and the post-reply trigger of the previous paragraph, rerun here, again does worse than prevention (13 % at 0.024). The trigger and the outcome use the same classifier, so the measure is the classifier's verdict on the model's original replies, which is comparable across arms. The discarded originals were not stored, so the second classifier can score only the replies shown: 0.8 % hostile (no control: 34 %). For this behaviour, the useful intervention is to keep the excursion out of the state at the moment it occurs.

Replication in a larger model. The same protocol in Qwen2.5-1.5B (hostility direction built identically at layer 14 of 28; 128 dialogues of 40 turns per arm; shares over turns 5–39) gives a steeper capture ladder that is not monotone (Figure 25). The hostile share rises from 0 at $\alpha = 0.01$ to 44 % at 0.02 and 80 % at 0.03–0.04. It falls to 30–34 % at 0.06–0.10, where most replies degenerate into token soup or tight accusatory loops (87 % degenerate at 0.08), and rises again to 75 % at 0.12, where replies are coherent again. The second classifier shows the same shape (32, 58, 41, 28, 20, 41 and 52 % from 0.02 to 0.12), and a random direction of the same norm gives at most 0.7 % hostile turns.

We ran the controller experiment twice with the pre-stated rule: controllers run at the ladder dose whose hostile share is closest to one half, with ladders on seed 1 and controller arms on seed 2, so the no-control arms differ slightly from the ladder at the same dose. With a first ladder (0.04–0.12) the rule chose 0.08, in the degenerate regime beyond the peak. There every counter-steer lowered the net dose toward the coherent-hostility peak: the share of the model's own replies judged

hostile rose from 36 % without control to 54 % with replacement and 60 % with the post-reply trigger, while degenerate replies fell. With a second ladder (0.01–0.04) the rule chose 0.02, on the rising side, where 96 % of replies are coherent. There the 0.5B results replicate in order and size. Replacement cuts the model's own hostile share from 43 % to 10 % ($[9, 12]$ %; shown replies 0.3 %) at a mean counter-dose of 0.004, 24 points ($[20, 28]$) below the constant dose-response at equal dose, and the post-reply trigger is again slightly less efficient than constant prevention (+2 points, $[+1, +3]$). On the shown replies the second classifier agrees (2.6 % against 32 % without control; −23 and +3 points). Constant counter-steering at −0.02 cancels this push, and the dialogues return to the model's ordinary politeness. The two ladders share the dose 0.04 and the seed, and that arm reproduced exactly from vectors rebuilt on another machine. Control that works on the rising side of a dose response can therefore backfire beyond its peak. Much like a bifurcation diagram, the dose response has to be mapped before a counter-steer is chosen.

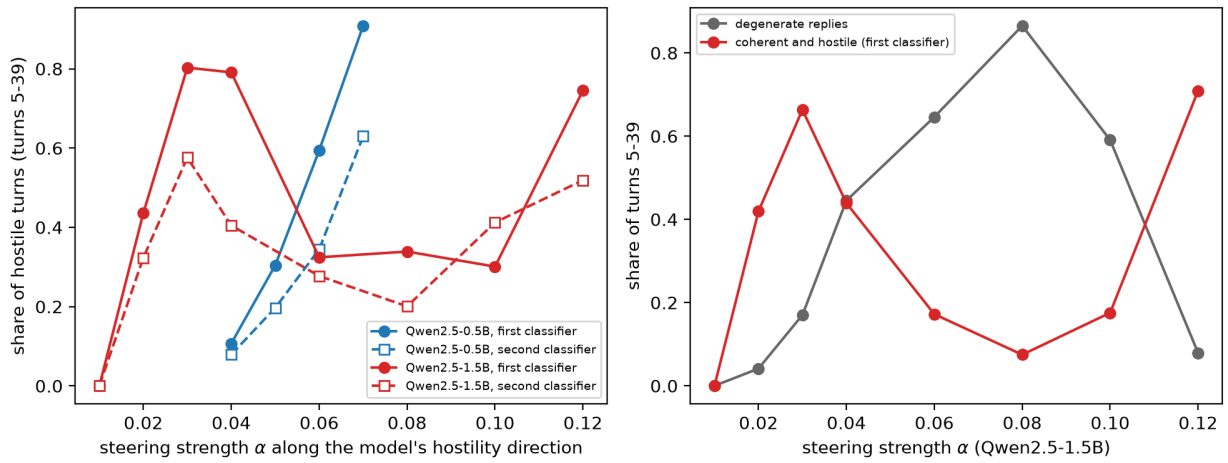


Figure 25: Hostility dose ladders of sampled self-dialogue in Qwen2.5-0.5B-Instruct (steered at layer 12) and Qwen2.5-1.5B (layer 14); 128 dialogues of 40 turns per point, shares over turns 5–39. Left: share of hostile turns along each model's own hostility direction, first classifier (filled) and second (open). Right (Qwen2.5-1.5B): share of degenerate replies (more than 30 % non-Latin characters, or fewer than 30 % distinct word pairs) and of coherent hostile replies. Steering strengths are each model's own α; equal values are not claimed to be equal pushes across models.

Lead time. The pre-specified test is negative. Models were fit on the first seed (all four doses) and tested on held-out dialogues of the second seed (doses 0.05 and 0.06). There, adding the dynamical features to the baselines lowers the AUC (area under the ROC curve) by 0.02 to 0.06 at horizons of 1–5 turns, with bootstrap intervals excluding zero. These values come after correcting a labelling error found in review, which had marked the last turns of dialogues without any hostile turn as positive; before the correction the losses read 0.02 to 0.10. The Markov hit-probability forecast does not transfer (AUC 0.37–0.43 on test, 0.60 in training), and the classifier's own score carries no warning before onset (AUC about 0.5). Two signals carry modest warning one to five turns ahead (AUC 0.53–0.64): the rolling variance of a hostility probe, the classical early-warning indicator, and the projection on the fixed steering axis, as in persona-vector monitoring [65]. In this design, then, the coarse-grained state estimate adds no warning beyond the baselines. An adapted analysis with dose-matched models (fit and test at the same dose) gives the same answer: incremental AUC −0.01 to −0.03 (recomputed with the corrected labels).

The model's own state carries the warning. The state variable of that test was each reply re-encoded without its context, not the model's state inside the dialogue. A second study (new seeds, doses 0.04 and 0.045, 512 dialogues; analysis stated before the data, cross-fitted by dialogue) recorded the model's own state just before each reply and probed it actively; internal states are known to carry information that the text does not show [75]. Its pre-stated primary predictor, the susceptibility (the change in the hostile fraction of eight sampled next replies when the dose is raised by 0.02), adds nothing: -0.002 $[-0.006, +0.002]$ AUC at three turns. Two secondary predictors add modest warning beyond the baselines. A probe of the in-context state before the next reply (layer 18, 20 principal components) raises the AUC for capture within three turns from 0.64 to 0.68 ($+0.036$ $[+0.014, +0.058]$; within one turn $+0.045$ $[+0.019, +0.069]$; within five $+0.036$), and a forecast from eight sampled next replies adds $+0.016$ $[+0.007, +0.027]$ at three turns. At a 10 % false-positive rate (threshold from the out-of-fold negatives), the probe flags 27 % of the turns that are followed by capture within three turns, against 23 % for the baselines. These values use labels corrected after a review found that the last turns of dialogues without any hostile turn had been marked positive; with the faulty labels the increments had read three times larger. The probe's direction in the layer-18 state is nearly orthogonal to the steering direction (cosine 0.04), so the warning is not simply more push along the hostility axis. Because these predictors were secondary, the result is exploratory.

A frozen confirmation then fit both models on this study and applied them unchanged to 1024 fresh dialogues. Its design and analysis were pushed before the data, and the label correction was added before any confirmation dialogue was read. The in-context state again adds warning for capture within three turns: AUC 0.656 $\rightarrow$ 0.689, $+0.033$ $[+0.018, +0.048]$ (within one turn $+0.036$, within five $+0.037$; both doses separately positive). The warning in the model's own state is therefore real and reproducible, but small. At thresholds fixed from the training negatives, the model with the state flags 29.1 % of the turns that are followed by capture within three turns, against 27.5 % for the baselines, at held-out false-positive rates of 10.3 % and 10.6 % (these are prediction windows, not distinct captures). It is not a reliable detector by itself.

How much warning is there? For the conversation state after a turn, the probability of capture within the next h turns, the committor [68], ranks states optimally among all warnings that use the same information. We estimated it for 1706 states of the confirmation dialogues from 32 continuations of three turns each, sampled exactly like the dialogues. This analysis is exploratory: it was stated before the data and amended twice after an independent review, still before the analysis outcomes were read. The continuations reproduce the realized capture rates (population-weighted 0.049, 0.094, 0.134 within one, two and three turns, against 0.046, 0.095, 0.136). On a stratified sample of 800 states, the committor estimate warns with an AUC of 0.72, 0.71 and 0.70 (a heuristic extrapolation to infinitely many continuations gives 0.75, 0.73, 0.72), and the in-context detector of the confirmation with 0.73, 0.72 and 0.71. The differences are not significant, with intervals reaching about ± 0.05 (Figure 26). Dialogues that were captured later sit at about twice the committor of never-captured ones over all ten preceding turns, with no visible ramp toward the onset (a description conditional on the outcome, not an identification of the mechanism). In this finite Monte Carlo experiment, the committor estimate did not provide detectably better warning than the in-context probe. The heuristic extrapolation leaves appreciable uncertainty, so these results do not identify the Bayes warning ceiling. Together with the persistent elevation of risk before capture, they motivate evaluating reactive control, such as detect-and-replace, alongside prediction; they do not prove that prediction cannot improve.

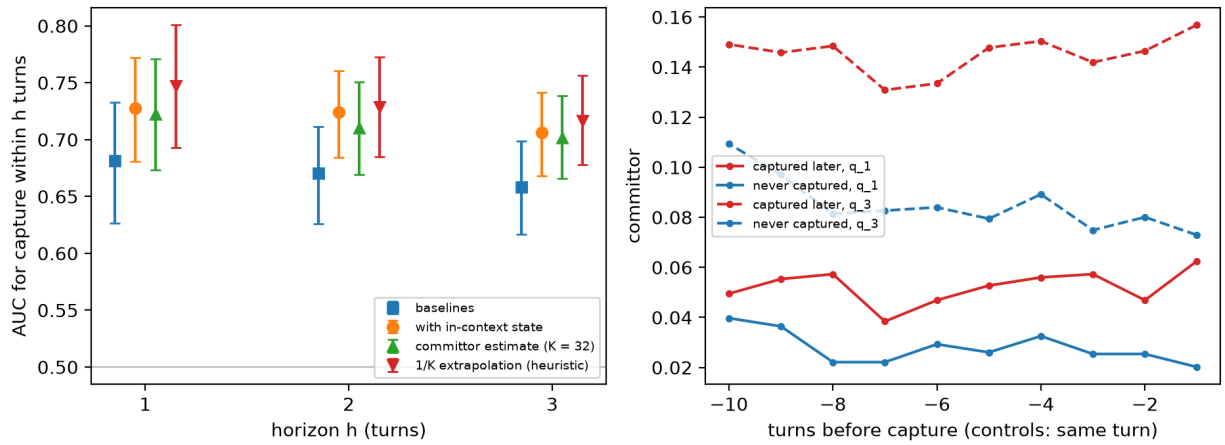


Figure 26: Warning of hostile capture in sampled self-dialogue of Qwen2.5-0.5B-Instruct (hostility direction at layer 12): detectors and the Monte Carlo committor. Left: AUC for capture within h turns on 800 stratified states of the confirmation dialogues, for the classifier baselines, the in-context detector (both fit on the exploration study), the Monte Carlo committor from 32 continuations, and its heuristic $1/K$ extrapolation (dialogue-bootstrap 95 % intervals). Right: mean committor in the ten turns before capture for 48 dialogues captured after turn 10, and for never-captured dialogues at the same turns (descriptive).

A second classifier. The Qwen2.5-0.5B outcomes above were first measured with one automated classifier. We therefore rescored all 44 of its dialogue sets with an independently trained RoBERTa classifier [76], trained on the three Jigsaw toxicity datasets (which include the first classifier's training set) and used with the same threshold, and an independent audit recomputed every pre-stated contrast with it. This is a robustness check on dialogues already analysed, so it is exploratory. On these sets the second classifier is stricter. At $\alpha = 0.06$ it flags 35 % of turns instead of 57 %, and at 0.07 62 % instead of 87 %. Almost all disagreements are replies that only the original flags (at 0.06, 23 % of turns against 0.7 % the other way), most of them repetitive exclamations such as "You're an AI! You're an AI!". Capture still rises steeply with the dose and stays specific to the hostility direction (random and cooking directions at 0.08: at most 0.2 % of turns), but the hostile state is less persistent (mean stay at 0.07: 6.5 turns instead of 16). History dependence survives. After the ramp, the late share exceeds that of a constant 0.04 by $+0.16$ $[+0.11, +0.21]$, and the held contrasts are $+0.05$ $[+0.01, +0.10]$ at 0.04 and $+0.04$ $[+0.01, +0.08]$ at 0.03; only the comparison with a constant 0.05 loses its positive lower bound ($+0.06$ $[-0.001, +0.12]$). Refitting the warning models with the second classifier's labels and baseline features (fit on the exploration study, applied unchanged to the confirmation dialogues) gives a pooled increment of $+0.043$ $[+0.024, +0.062]$ AUC for capture within three turns. It stays positive at dose 0.045 ($+0.061$) and is $+0.020$ $[-0.009, +0.049]$ at 0.04. The three replacement arms keep the replies shown below the constant-dose response at equal mean dose under both classifiers. Both instruments are automated and may share errors; no human labels were collected. **[Planned:** a hazard model; longer holds to separate slow release from bistability; other openers and steering directions; larger models; human labels for the outcome.]

13 Discussion

What theory allows, and what is new. Chaos somewhere in a two-parameter family of high-dimensional nonlinear maps is generic. Random recurrent networks with asymmetric couplings become chaotic above a critical gain, sharply in the large-network limit, and finite networks pass

through limit cycles of increasing complexity on the way [19] (see also [77]). External input can suppress that chaos [78], and deep networks at initialization show an order-to-chaos transition in signal propagation [79, 80]. In our map, temperature plays the role of an inverse gain and the steering vector that of an input. Periodic windows, shrimp-shaped domains with their cusps and spirals, and doubling cascades with Feigenbaum scaling are universal in dissipative two-parameter systems [20, 21, 47, 48, 81] and are seen in experiments [82], and intermittency near folds is classical [46]. These theories make chaos, windows and shrimp-like structure plausible in a language model's feedback map, in the sense that such structure can occur. They do not say whether a trained language model realizes it, where in its parameter plane, with which texts, or whether the theory's quantitative predictions hold for that model. This paper shows that the models we studied do realize it, with concrete texts as attractors, so that a bifurcation changes what the model writes, and that the quantitative predictions hold in the cases we tested.

As far as we know, no earlier work has computed periodic orbits and their multipliers, continued fold and flip bifurcation curves and located numerically nondegenerate cusps in the computation of a pretrained language model. Nor do we know of earlier quantitative predictions from such a model's derivatives that were confirmed prospectively, as ours were for passage times at new folds (23 of 24 forecasts within one iterate, with the all-row criterion failing by one row) and for orbit counts and stabilities at frozen offsets near cusps. A spectral prediction of task failure is confirmed on held-out data in Huginn, and one periodic orbit of the full transformer map of Qwen2.5-0.5B-Instruct is proved by interval arithmetic; we know of no earlier instance of either for a language model. The dynamical view also has consequences for control in sampled dialogue, where the dose response is not monotone and counter-steering can backfire.

Is steering the origin of the chaos? Steering adds a fixed vector to the residual stream, so chaos along a steering axis could belong to the push rather than to the model. At $\alpha = 0$, however, the map is the unsteered model's own feedback map, with temperature as its only parameter. The $\alpha = 0$ columns of the hostility maps contain cells with positive finite-horizon exponents in Llama-3.1 8B (26 of 64 temperatures), Qwen2.5-7B (15 of 64) and OLMo-2 7B (11 of 64), while the near-zero columns of the 0.5B and 1.5B maps contain one such cell each (of 128 and 64). A longer float64 sweep of the unsteered 0.5B map (512 temperatures over $\log_{10} T \in [-1, 0.3]$, 1000 transient and 2000 measured steps) finds a narrow chaotic band at $\log_{10} T \approx 0.15$–0.21 (7 temperatures, exponents up to 0.28), with a period-5 loop at 437 of the 512 temperatures. The same long sweep finds 96 chaotic temperatures in Llama-3.1 8B, between $\log_{10} T = -0.20$ and 0.24 (exponents up to 0.66), interleaved with periodic windows of periods 2 to 52; within $\log_{10} T \in [-0.3, 0.3]$ its chaotic share, 41 %, equals that of the coarse column. Qwen2.5-7B has 65 chaotic temperatures between $\log_{10} T = -0.07$ and 0.28 (exponents up to 0.74), and OLMo-2 7B has 32 between 0.04 and 0.14 (up to 0.30). Steering is therefore not needed to reach chaotic regimes in any of the four models tested. Temperature alone produces chaos with embedded windows, broadly in the 7B and 8B models and in a narrow band in the 0.5B model (Figure 27).

Steering does widen the chaotic regions, but not as a property of the concept: along norm-matched random directions, the 0.5B maps are as chaotic as along the concept directions (Figure 28). On the same 128^2 grid, 7.6, 8.3 and 12.5 % of cells are chaotic for three Gaussian directions with the philosophy vector's norm, against 6.8 % for philosophy and 7.5 % for gratitude. In all five maps the chaotic share rises from about 1 % at $|\alpha| < 0.01$ to between 5 and 22 % at larger $|\alpha|$, and a chaotic region at high temperature, with windows reaching down, recurs. In this model, chaos under steering is therefore a generic response to perturbing the residual stream at this norm. The direction does select the texts. The largest attractor family in all five maps is the unsteered period-5 loop,

but beyond it only 26 to 56 % of one map's cells carry an attractor text that appears anywhere in another map. Llama-3.1 8B agrees. On its hostility plane, two norm-matched random directions at the same layer give 26 and 17 % chaotic cells against 28 % along the concept direction; all three carry 34 to 44 % near $\alpha = 0$, and the random directions thin out at large $|\alpha|$. At layer 19, where a sparse autoencoder trained on this model reads the residual stream [53], the concept vector rebuilt there gives 28 %. The decoder direction of the autoencoder feature most specific to hostility on held-out texts (chosen by a rule fixed before any map; cosine 0.10 with the concept vector) gives 29 %, a random live feature 27 % and a Gaussian direction 39 %, all at the same norm. Feature directions of an autoencoder therefore steer the map into chaos no differently from contrastive or random directions.

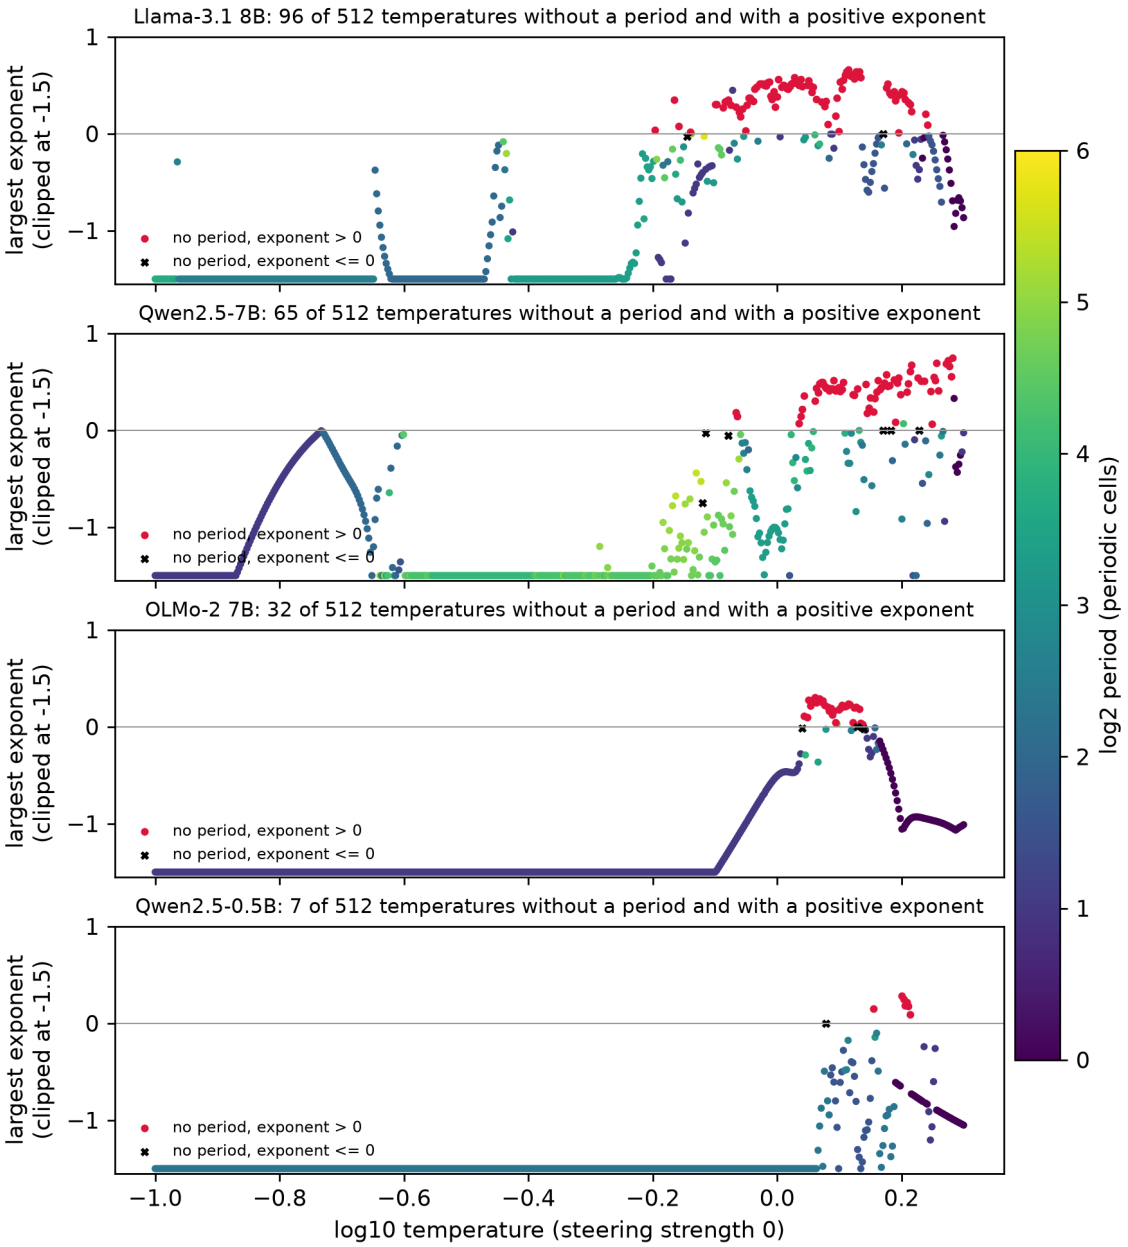


Figure 27: The unsteered soft-token map ($\alpha = 0$) of Llama-3.1-8B-Instruct, Qwen2.5-7B, OLMo-2-1124-7B-Instruct and Qwen2.5-0.5B-Instruct (top to bottom) along temperature: largest Lyapunov exponent from float64 sweeps of 512 temperatures (1000 transient and 2000 measured steps). Red: no period up to 64 and a positive exponent; other points are periodic, coloured by $\log_2$ of the period.

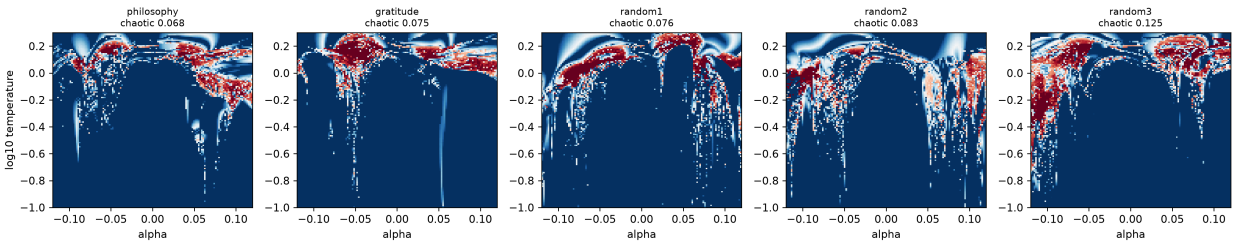


Figure 28: Steering controls in Qwen2.5-0.5B-Instruct (128^2 float64 cells, steering at layer 12, one grid and one set of settings): largest Lyapunov exponent along the philosophy and gratitude directions and along three Gaussian directions rescaled to the philosophy vector's norm. Titles give the chaotic share (finite-horizon exponent above 10^{-3}, no period up to 48).

What is and is not claimed. The objects studied here are invariant sets and bifurcations of an exactly defined smooth map built from a frozen model's own computation. We do not claim that ordinary sampled generation is chaotic in this sense (Section 9), nor that any window is a certified shrimp. The period-6 window and the hostile period-7 body have, numerically, shrimp-like organization: fold curves ending in nondegenerate cusps (simple critical multiplier; quadratic coefficient through zero with nonzero slope; cubic coefficient nonzero) and flip curves inside, and in the hostile body superstable curves run through the windows. A generic codimension-two cusp also requires the two unfolding coefficients to depend on the two parameters with rank two. We computed this at one cusp of each window (determinants 2737 and 53.5), and there the local model's predictions at twelve frozen offsets matched Newton solutions in the number of orbits at all twelve and in their stability at eleven. These are float64 computations validated by step-size checks (rung 1; an independent recomputation at one cusp is in progress), not computer-assisted proofs.

Limitations. The detailed orbit studies cover one model family (Qwen2.5, 0.5B and 1.5B), with coarse maps and three float64 orbits at 7B; the two other families (OLMo-2 7B and Llama-3.1 8B) have coarse maps only. The maps use one prompt, and steering uses contrastive vectors, checked against norm-matched random and sparse-autoencoder directions in two models. The system is a constructed one-token window. The evidence consists of numerical validation, cross-checks and prospectively frozen passage tests (including failed all-row criteria), with a computer-assisted proof for only one orbit at one parameter point. Chaos labels are finite-horizon estimates (300 measured steps after 300 to 500 transient steps), and no chaotic cell is proved chaotic. The Feigenbaum ratios are finite-order values that rise toward δ (the last is within 1 %); the limit itself is not computed. All orbit, bifurcation and chaos results concern the constructed map, not sampled generation with a growing context. The bridge of Section 9 tests only whether the bottleneck's laminar phases survive sampling noise, and the behavioural results are statistics of sampled dialogue, not orbits. The behavioural study uses two small models of one family (0.5B and 1.5B), one opener and a proxy defined by two automated classifiers without human labels. Its first lead-time test is negative, and the confirmed warning from the in-context state is small.

Steering strength is a bifurcation parameter. In the soft-token map, small changes of the steering strength switch the model between distinct repeating texts, chaos and coexisting attractors. Near a fold curve, behaviour becomes intermittent, with laminar phases whose duration the fold's normal form predicts in advance. In sampled dialogue, the dose along a hostility direction controls how long a hostile state persists, continuously from bursts to near-absorption, while directions of the same norm do almost nothing. Both observations argue for treating an activation-steering coefficient

like any control parameter of a nonlinear system. Its bifurcation set should be characterized before an operating point is chosen, history dependence should be expected near boundaries, and a monotone, smooth dose response should not be assumed. In steering studies, outputs that fall into verbatim repetition at high doses are usually discarded as quality failures. In a replication of a feature-steering study of self-reports [83], the highest doses failed quality checks in this way [**Planned:** citation of the replication once posted]. In the present framework such loops are candidate attracting periodic orbits, and the dose at which they appear is a boundary that can be mapped like the others. The deterministic map is where such sets can be defined precisely and computed numerically. Whether its boundaries predict boundaries of sampled behaviour is open. The bridge of Section 9 shows that sampling noise shortens the laminar phases of the bottleneck, by an amount that depends on how many samples are averaged.

14 Evidence and reproducibility

We grade claims on an evidence ladder. Rung 1 means numerically validated (tolerances, residuals, convergence, precision varied); rung 2, cross-checked by an independent estimate, implementation or hardware; rung 3, a computer-assisted proof.

Search scale. The soft-token results come from an exploratory search. By 14 October 2026 we had computed 34 distinct maps (987 308 cells, 27 maps in float64) in six models, each with one prompt and one of three steering directions or none, plus one partial map. Of these, 29 maps and the partial one appear here; the five omitted are early discovery or side maps (the steering-control maps of Section 13 are additional). The labelled maps contain about 2600 periodic families of at least six cells, counted per map. We singled out about 26 windows or families for float64 work: 20 are reported in detail, 2 only as map features and 4 not at all (an orbit whose continuation was abandoned, a cell that proved aperiodic, a band whose float32 orbit did not converge and a pilot that found no cusp within a bounded search). Of 54 Newton orbit attempts, 7 failed. Of seven tangent reversals on continued fold curves, six were tested, and all six are reported as numerically nondegenerate cusps; one is untested. Two of the five regions we flagged as shrimp candidates are reported as shrimp-like. One further language model (Ouro-1.4B) was explored and is not reported; its apparent chaos was a long transient. Every window was chosen after the maps were inspected. Only the passage forecasts, the cusp-unfolding offsets, Huginn's relaxation test and the dialogue warning confirmation were fixed before the computations that test them (Table 11).

	explored	reported	note
Soft-token maps (distinct)	34	29	plus one partial map, reported
cells	987 308	893 868	omitted: 5 discovery or side maps
float32 maps	7	4	
Models with maps	6	6	one prompt; 3 directions and $\alpha = 0$
Periodic families (≥ 6 cells)	≈ 2600	–	counted per labelled map
Windows given float64 work	≈ 26	20 (+2)	+2: map features only
Newton orbit attempts	54	–	7 failed
Tangent reversals on fold curves	7	6	1 untested
Unfolding offsets, frozen	12	12	counts 12/12, stability 11/12
Shrimp-candidate regions	5	2	reported as shrimp-like
Tests fixed before outcomes	6	5	2 all-row criteria failed by one row
Language models loaded	8	7	Ouro-1.4B not reported

Table 11: Search scale on 14 October 2026, before the steering-control maps (Huginn's frozen test counted once; the sixth test, a cusp pilot in a second checkpoint, is the unreported bounded search).

- Rung 3: the period-8 stripe orbit of Table 1. At $\alpha = 0.071674$ and $\log_{10} T = -0.115$, read as exact binary64 rationals, the expected-embedding map of the stored Qwen2.5-0.5B-Instruct checkpoint (all 24 layers, the full 151 936-token readout, the fixed 37-token prefix, steering after block 12) has a primitive period-8 orbit, which is the unique multiple-shooting root in a product of eight boxes of coordinate radius 2^{-640} around refined centres, and the orbit is locally attracting. Ball arithmetic (Arb [84]) encloses the map and its Jacobian over each box (8×896 directions); a Krawczyk-type inclusion for the preconditioned shooting system gives existence and uniqueness, separation of the first point from the points one, two and four steps later gives the primitive period, and a norm bound on powers of the enclosed monodromy gives attraction. An independent audit in exact rational arithmetic rechecked the saved witnesses. The proof trusts the Python, FLINT and Arb stack and the stored weights. It does not exclude other period-8 points in the first box, and it says nothing about the window's boundaries or global organization (record: `certificates/llm-p8-certificate-final-001.json`).

- Rung 2: the 2-cycle's fold and its normal form. An independent reimplementation of the forward map and solver recovers α_c, a and b (to 6×10^{-15}, 3×10^{-12} and 2×10^{-6} relatively), and ab also follows from the stable-branch multipliers. The same reimplementation reproduces all 18 fixed-entrance counts of the controlled passage runs exactly (Table 2 lists 12 of them).

- Prospectively frozen: the passage forecasts at two new folds (Table 4); 23 of 24 within one iterate, and the all-row criterion failed by one row.

- Rung 2 as well: the period-2 and period-8 orbits of Table 1 (independent implementation), and the cascade's first four flip parameters (CPU and H100, agreement 4×10^{-15}).

- Rung 2 as well: the hostile 2-cycle fold and its normal form, reproduced by an independent float64 forward map with nested automatic differentiation (location to 6×10^{-10}, ab to 8 parts per million); its controlled passages were also checked independently, with the frozen all-six criterion failing by one row.

- Rung 1: the other orbits, the curves, the six cusps (numerically nondegenerate), the superstable parameters, the 1.5B flips, the hostile windows' edges, the spiral region's period-7 body and the cascade inside it, the three 7B orbits, and Huginn's float32 spectral estimates.

- Prospectively frozen and independently replayed: Huginn's relaxation test on held-out data, confirmed in both parts; an independent implementation reproduced both from the saved rows and spectra (GSM8K part exactly, item-level part to 10^{-12}) (Section 11).
- Prospectively frozen as well: the confirmation of the in-context warning in sampled dialogue (Section 12), whose design was pushed before the data and amended, before any outcome was read, to correct a labelling error.

Hardware. Float64 GPU computations ran on rented NVIDIA H100 (NVL, PCIe and 80 GB HBM3) and A100 80 GB (SXM4 and PCIe) GPUs, one GPU per run; cross-checks and the interval proof ran on CPUs. The manifest `docs/manifests/runpod-runs.json` records, for every run, the GPU type, the code commit, the job and its arguments, the wall time, the exit status and the hashes of the uploaded inputs and the retrieved results (138 runs with results on 14 October 2026). GPU types round float64 reductions differently, so quantities used across runs were checked across hardware where it mattered (for example, flip parameters agree between CPU and H100 to 4×10^{-15}). Code: `tools/explore_llm_soft_loop.py`, `tools/bifurcation_curves.py`, `tools/llm_feigenbaum.py`, `tools/explore_latent_loop.py`, `tools/explore_llm_msm.py`, `tools/hostility_score.py`, `tools/p1_committor.py`, `tools/p1_controller.py`, `tools/p1_degeneracy.py`, `tools/cusp_check.py`, `tools/huginn_relax_eval.py`, `tools/huginn_arc_analysis.py`. Raw-data hash receipts: `evidence/llm-dynamics-data-receipts.json` (to 12 October), the per-run receipts in the research log, and the verified snapshot manifests in `docs/backups/`; independent cross-check and audit receipts are under `codex/` (`crosscheck`, `confirm-fold`, `hostile-fold-crosscheck`, `misalignment-audit`, `l9-committor-review`, `huginn-design-review`, `huginn-independent-replay`, `p1-second-classifier-audit`, `paper-review`).

[**Planned:** public code and data release at submission.]

A Self-dialogue under steering: the gratitude ladder (exploratory)

Two copies of Qwen2.5-0.5B-Instruct talk to each other: 128 replicas, 40 turns, opener "Hello!", a window of 7 turns, 60 new tokens per turn, sampled at $T = 0.7$, with the gratitude direction added at layer 12. Turn activations are clustered (principal component analysis to 20 dimensions, 40 states) into a Markov state model with four metastable sets [49]. Table 12 lists finite-horizon measures for each strength.

gratitude strength α	0	0.01	0.02	0.04
gratitude set	none	62 % weight	dominant	95 % weight
occupancy at turn 9 / 14 / 39	–	0.34 / 0.52 / 0.73	0.70 / 0.91 / 0.91–0.98	from turn 0
entry turn, quartiles	–	5 / 7 / 11.8	4 / 5 / 7	0
replicas never entering	–	14 of 128	2 of 128	0
residence per visit (turns)	–	about 15	about 18	effectively absorbing
replies with a gratitude word	0.00–0.05	0.66 at turn 39	0.96–0.98 from turn 9	0.88 at turn 0
spiritual-word plateau	0.00–0.02	0.13	0.26–0.30	0.42–0.48

Table 12: The gratitude ladder (one opener and seed per rung; word lists chosen after reading samples). At 0.02 the dialogue starts in the helper mode and enters the gratitude set within a few turns; at 0.04 the set is imposed from the first reply. **[Planned:** frozen confirmatory version on held-out openers with dialogue-level bootstrap.]

References

[1] Xiaojian Li, Yongkang Leng, Ruiqing Ding, et al. Cognitive activation and chaotic dynamics in large language models: A quasi-Lyapunov analysis of reasoning mechanisms, 2025.

[2] Gregorio Jaca, Kristóf Benedek, and János Török. Chaos in reason: How chain-of-thought LLMs can look for an answer, 2026.

[3] Julian Arnold, Flemming Holtorf, Frank Schäfer, et al. Phase transitions in the output distribution of large language models, 2024.

[4] Kai Nakaishi, Yoshihiko Nishikawa, and Koji Hukushima. Phase transition in large language models and the criticality of natural languages, 2024.

[5] Youran Sun and Babak Haghighat. Phase transitions in large language models and the $O(n)$ model, 2025.

[6] Huajian Ruan, Jinyang Li, Xingyu Guo, et al. Generative criticality in large language model temperature scaling, 2026.

[7] Neil F. Johnson, Frank Yingjie Huo, and Bella Xinrui Li. Temperature-driven inversion and nonlinear dynamics in ChatGPT-like AIs, 2026.

[8] Zhilin Wang, Yafu Li, Jianhao Yan, et al. Unveiling attractor cycles in large language models: A dynamical systems view of successive paraphrasing, 2025.

[9] Ting-Wen Ko and Jonas Geiping. Attractor states emerge in multi-turn LLM conversations, 2026.

[10] Shaojie Bai, J. Zico Kolter, and Vladlen Koltun. Deep equilibrium models, 2019.

[11] Shaojie Bai, Vladlen Koltun, and Zico Kolter. Stabilizing equilibrium models by Jacobian regularization. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 554–565. PMLR, 2021. URL https://proceedings.mlr.press/v139/bai21b.html.

[12] Akiyoshi Tomihari and Ryo Karakida. Recurrent self-attention dynamics: An energy-agnostic perspective from Jacobians. In *Advances in Neural Information Processing Systems*, volume 38, pages 155236–155269, 2025. doi: 10.52202/085713-4675. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/cca569b38ec3612904e625dad7a743dc-Abstract-Conference.html.

[13] Xiao-Wen Yang, Ziyu Han, Xi-Hua Zhang, Wen-Da Wei, Jie-Jing Shao, Lan-Zhe Guo, and Yu-Feng Li. Stabilizing recurrent dynamics for test-time scalable latent reasoning in looped language models. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, pages 144670–144682. PMLR, 2026. URL https://proceedings.mlr.press/v306/yang26t.html.

[14] Xinyue Luo and Fei Yu. Prediction dynamics in depth-recurrent language models, 2026.

[15] Boyi Zeng, Shixiang Song, Siyuan Huang, et al. PonderLM: Pretraining language models to ponder in continuous space, 2025.

[16] Zhen Zhang, Xuehai He, Weixiang Yan, Ao Shen, Chenyang Zhao, Shuohang Wang, Yelong Shen, and Xin Eric Wang. Soft thinking: Unlocking the reasoning potential of LLMs in continuous concept space, 2025. URL https://arxiv.org/abs/2505.15778. Version 1, 21 May 2025.

[17] Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, et al. Training large language models to reason in a continuous latent space, 2024.

[18] Gunn Kim. Dynamical phase selection controls compute scaling in looped transformers, 2026. URL https://arxiv.org/abs/2608.26556. Version 1, 27 August 2026.

[19] H. Sompolinsky, A. Crisanti, and H. J. Sommers. Chaos in random neural networks. *Physical Review Letters*, 61(3):259–262, 1988. doi: 10.1103/PhysRevLett.61.259.

[20] Mitchell J. Feigenbaum. Quantitative universality for a class of nonlinear transformations. *Journal of Statistical Physics*, 19(1):25–52, 1978. doi: 10.1007/BF01020332.

[21] Jason A. C. Gallas. Structure of the parameter space of the Hénon map. *Physical Review Letters*, 70:2714–2717, 1993. doi: 10.1103/PhysRevLett.70.2714.

[22] Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, et al. The curious case of neural text degeneration, 2019.

[23] Jin Xu, Xiaojiang Liu, Jianhao Yan, Deng Cai, et al. Learning to break the loop: Analyzing and mitigating repetitions for neural text generation, 2022.

[24] António H. Ribeiro, Koen Tiels, Luis A. Aguirre, and Thomas Schön. Beyond exploding and vanishing gradients: analysing RNN training using attractors and smoothness. In *Proceedings of the Twenty Third International Conference on Artificial Intelligence and Statistics*, volume 108 of *Proceedings of Machine Learning Research*, pages 2370–2380. PMLR, 2020. URL https://proceedings.mlr.press/v108/ribeiro20a.html.

[25] Jonas Geiping, Sean McLeish, Neel Jain, John Kirchenbauer, Siddharth Singh, Brian Bartoldson, Bhavya Kailkhura, Abhinav Bhatele, and Tom Goldstein. Scaling up test-time compute with latent reasoning: A recurrent depth approach. In *Advances in*

Neural Information Processing Systems, volume 38, pages 46242–46293, 2025. doi: 10.52202/085713-1380. URL https://proceedings.nips.cc/paper_files/paper/2025/hash/3b01972cf31e6fa0fe29e4b8b5c2a0a1-Abstract-Conference.html.

[26] Rui-Jie Zhu, Zixuan Wang, Kai Hua, et al. Scaling latent reasoning via looped language models, 2025.

[27] Hayden Prairie, Zachary Novack, Taylor Berg-Kirkpatrick, and Daniel Y. Fu. Parcae: Scaling laws for stable looped language models, 2026. URL https://arxiv.org/abs/2604.12946. Version 1, 14 April 2026.

[28] Matthias M. M. Buehlmaier. LayerNorm as implicit gain control in looped transformers, 2026. URL https://arxiv.org/abs/2607.10681. Version 1, 12 July 2026.

[29] Asher Labovich. Stability and generalization in looped transformers, 2026. URL https://arxiv.org/abs/2604.15259. Version 2, 22 April 2026.

[30] Sajad Movahedi, Vera Milovanović, Shlomo Libo Feigin, Alexander Theus, Thomas Hofmann, Valentina Boeva, T. Konstantin Rusch, and Antonio Orvieto. Fixed-point reasoners: Stable and adaptive deep looped transformers, 2026. URL https://arxiv.org/abs/2606.18206. Version 1, 16 June 2026.

[31] Junhong Wu, Jinliang Lu, Zixuan Ren, Gangqiang Hu, Zhi Wu, Dai Dai, and Hua Wu. LLMs are single-threaded reasoners: Demystifying the working mechanism of soft thinking, 2025. URL https://arxiv.org/abs/2508.03440. Version 4, 16 October 2025.

[32] Fatih Dinc, Ege Cirakman, Bariscan Kurtkaya, Mert Yuksekgonul, Yiqi Jiang, Mark J. Schnitzer, and Hidenori Tanaka. Ghost mechanism: An analytical model of abrupt learning in recurrent networks. *Physical Review X*, 16(2):021058, 2026. doi: 10.1103/mjcl-lb4x.

[33] Demian Battaglia, Nicolas Brunel, and David Hansel. Temporal decorrelation of collective oscillations in neural networks with local inhibition and long-range excitation. *Physical Review Letters*, 99(23):238106, 2007. doi: 10.1103/PhysRevLett.99.238106.

[34] Christoforos A. Papasavvas, Yujiang Wang, Andrew J. Trevelyan, and Marcus Kaiser. Gain control through divisive inhibition prevents abrupt transition to chaos in a neural mass model. *Physical Review E*, 92(3):032723, 2015. doi: 10.1103/PhysRevE.92.032723.

[35] Yilun Liu, Yi Zhang, Ganyu Wu, Sikuan Yan, Mengyue Wang, Alois Knoll, Volker Tresp, and Yunpu Ma. Learning chaos without seeing chaos: Extrapolation of global dynamics in autoregressive transformers, 2026. URL https://arxiv.org/abs/2609.38814v1. Version 1, 30 September 2026.

[36] Jason Z. Kim, Zhixin Lu, Erfan Nozari, George J. Pappas, and Danielle S. Bassett. Teaching recurrent neural networks to infer global temporal structure from local examples. *Nature Machine Intelligence*, 3(4):316–323, 2021. doi: 10.1038/s42256-021-00321-2.

[37] Dayal Singh Kalra, Tianyu He, and Maissam Barkeshli. Universal sharpness dynamics in neural network training: Fixed point analysis, edge of stability, and route to chaos. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/8c8cd1b78cdae751265c88efc136e5bd-Abstract-Conference.html.

[38] R. L. Marichal and J. D. Piñeiro. Analysis of multiple quasi-periodic orbits in recurrent neural networks. *Neurocomputing*, 162:85–95, 2015. doi: 10.1016/j.neucom.2015.04.001.

[39] Florian Gomez, Ralph L. Stoop, and Ruedi Stoop. Universal dynamical properties preclude standard clustering in a large class of biochemical data. *Bioinformatics*, 30(17):2486–2493, 2014. doi: 10.1093/bioinformatics/btu332.

[40] Rafael V. Stenzinger, Vinícius Luz Oliveira, and M. H. R. Tragtenberg. Shrimp hubs in the Hindmarsh–Rose model. *Chaos: An Interdisciplinary Journal of Nonlinear Science*, 35:023148, 2025. doi: 10.1063/5.0239268.

[41] Luiz F. B. Caixeta, Matheus H. P. Gonçalves, M. H. R. Tragtenberg, and Mauricio Girardi-Schappo. Devil's staircase inside shrimp-shaped regions reveals periodicity of plateau spikes and bursts. *Chaos: An Interdisciplinary Journal of Nonlinear Science*, 35(3):033122, 2025. doi: 10.1063/5.0250342.

[42] Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, Shashwat Goel, Nathaniel Li, Michael J. Byun, Zifan Wang, Alex Mallen, Steven Basart, Sanmi Koyejo, Dawn Song, Matt Fredrikson, J. Zico Kolter, and Dan Hendrycks. Representation engineering: A top-down approach to AI transparency, 2023. URL https://arxiv.org/abs/2310.01405. Version 4, 3 March 2025; first posted 2 October 2023.

[43] Jinghao Zhang, Yuting Liu, Wenjie Wang, Qiang Liu, Shu Wu, Liang Wang, and Tat-Seng Chua. Personalized text generation with contrastive activation steering. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7128–7141, 2025. doi: 10.18653/v1/2025.acl-long.353.

[44] Yuri A. Kuznetsov. *Elements of Applied Bifurcation Theory*. Applied Mathematical Sciences. Springer, 2023. doi: 10.1007/978-3-031-22007-4.

[45] Willy J. F. Govaerts. *Numerical Methods for Bifurcations of Dynamical Equilibria*. Society for Industrial and Applied Mathematics, 2000. doi: 10.1137/1.9780898719543.

[46] Yves Pomeau and Paul Manneville. Intermittent transition to turbulence in dissipative dynamical systems. *Communications in Mathematical Physics*, 74(2):189–197, 1980. doi: 10.1007/BF01197757.

[47] Wilson Façanha, Bart Oldeman, and Leon Glass. Bifurcation structures in two-dimensional maps: The endoskeletons of shrimps. *Physics Letters A*, 377(18):1264–1268, 2013. doi: 10.1016/j.physleta.2013.03.025.

[48] Renato Vitolo, Paul Glendinning, and Jason A. C. Gallas. Global structure of periodicity hubs in Lyapunov phase diagrams of dissipative flows. *Physical Review E*, 84(1):016216, 2011. doi: 10.1103/PhysRevE.84.016216.

[49] John D. Chodera and Frank Noé. Markov state models of biomolecular conformational dynamics. *Current Opinion in Structural Biology*, 25:135–144, 2014. doi: 10.1016/j.sbi.2014.04.002.

[50] Frank Pasemann. Complex dynamics and the structure of small neural networks. *Network: Computation in Neural Systems*, 13(2):195–216, 2002. doi: 10.1080/net.13.2.195.216.

[51] Team OLMo, Pete Walsh, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Shane Arora, et al. 2 OLMo 2 furious, 2024.

[52] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, et al. The Llama 3 herd of models, 2024.

[53] Goodfire. Llama-3.1-8B-Instruct-SAE-l19. Hugging Face, https://huggingface.co/Goodfire/Llama-3.1-8B-Instruct-SAE-l19, 2025. Sparse autoencoder on the residual stream after layer 19; revision f6775a22.

[54] Seonglae Cho, Harryn Oh, Donghyun Lee, Luis Eduardo Rodrigues Vieira, Andrew Bermingham, and Ziad El Sayed. FaithfulSAE: Towards capturing faithful features with sparse autoencoders without external dataset dependencies, 2025.

[55] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, et al. Training verifiers to solve math word problems, 2021.

[56] Boyuan Wang, Chengyao Yu, Jiaxi Ren, Hongxin Wei, Bingyi Jing, and Yuxin Tao. Scheduling recursive reasoning in looped transformers, 2026. Submitted 29 September 2026.

[57] Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, et al. Think you have solved question answering? Try ARC, the AI2 reasoning challenge, 2018.

[58] Anthropic. System card: Claude opus 4 & claude sonnet 4. Technical report, Anthropic, 2025. URL https://www-cdn.anthropic.com/6d8a8055020700718b0c49369f60816ba2a7c285/Claude%204%20System%20Card.pdf. May 2025; revision 2 September 2025; sections 5.5.1–5.5.2.

[59] Christina Lu, Jack Gallagher, Jonathan Michala, Kyle Fish, and Jack Lindsey. The assistant axis: Situating and stabilizing the default persona of language models, 2026. URL https://arxiv.org/abs/2601.10387v1. Version 1, 15 January 2026.

[60] Cem Anil, Esin Durmus, Mrinank Sharma, Joe Benton, Sandipan Kundu, Joshua Batson, Nina Rimsky, Meg Tong, Jesse Mu, Daniel Ford, Francesco Mosconi, Rajashree Agrawal, Rylan Schaeffer, Naomi Bashkansky, Samuel Svenningsen, Mike Lambert, Ansh Radhakrishnan, Carson Denison, Evan J Hubinger, Yuntao Bai, Trenton Bricken, Timothy Maxwell, Nicholas Schiefer, Jamie Sully, Alex Tamkin, Tamera Lanham, Karina Nguyen, Tomasz Korbak, Jared Kaplan, Deep Ganguli, Samuel R. Bowman, Ethan Perez, Roger Grosse, and David Duvenaud. Many-shot jailbreaking. Technical report, Anthropic, 2024. URL https://www-cdn.anthropic.com/af5633c94ed2beb282f6a53c595eb437e8e7b630/Many_Shot_Jailbreaking2024_04_02_0936.pdf. Original report, 2 April 2024.

[61] Jan Betley, Daniel Tan, Niels Warncke, Anna Sztyber-Betley, Xuchan Bao, Martín Soto, Nathan Labenz, and Owain Evans. Emergent misalignment: Narrow finetuning can produce broadly misaligned LLMs, 2025. URL https://arxiv.org/abs/2502.17424v7. Version 7, 20 January 2026; preprint originally submitted in 2025.

[62] Nina Rimsky, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Turner. Steering Llama 2 via contrastive activation addition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15504–15522, 2024. doi: 10.18653/v1/2024.acl-long.828.

[63] Laura Hanu and Unitary team. Detoxify. GitHub, https://github.com/unitaryai/detoxify, 2020. Model unitary/toxic-bert, Hugging Face revision 4d6c22e7.

[64] Dhananjay Ashok and Jonathan May. Language models can predict their own behavior. In *Advances in Neural Information Processing Systems*, volume 38, 2025. doi: 10.52202/085713-2102. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/5a9c1af5f76da0bd37903b6f23e96c74-Abstract-Conference.html.

[65] Runjin Chen, Andy Arditi, Henry Sleight, Owain Evans, and Jack Lindsey. Persona vectors: Monitoring and controlling character traits in language models, 2025. URL https://arxiv.org/abs/2507.21509v1. Version 1, 29 July 2025.

[66] Fan Zhou, Weitian Wang, and Tim Van de Cruys. Commitment before realization: When classifier-free guidance becomes unnecessary in masked diffusion language models, 2026. URL https://arxiv.org/abs/2608.08082v1. Version 1, 8 August 2026.

[67] Nicolas Leins, Jennifer Haase, Varvara Geronimus, Jana Gonnermann-Müller, and Sebastian Pokutta. Prompting against persona drift: Comparing intervention timing and content in LLM-simulated conversations, 2026. URL https://arxiv.org/abs/2609.24532v1. Version 1, 21 September 2026.

[68] Weinan E and Eric Vanden-Eijnden. Towards a theory of transition paths. *Journal of Statistical Physics*, 123(3):503–523, 2006. doi: 10.1007/s10955-005-9003-9.

[69] Weinan E and Eric Vanden-Eijnden. Transition-path theory and path-finding algorithms for the study of rare events. *Annual Review of Physical Chemistry*, 61(1):391–420, 2010. doi: 10.1146/annurev.physchem.040808.090412.

[70] Marten Scheffer, Jordi Bascompte, William A. Brock, Victor Brovkin, Stephen R. Carpenter, Vasilis Dakos, Hermann Held, Egbert H. van Nes, Max Rietkerk, and George Sugihara. Early-warning signals for critical transitions. *Nature*, 461(7260):53–59, 2009. doi: 10.1038/nature08227.

[71] H. A. Kramers. Brownian motion in a field of force and the diffusion model of chemical reactions. *Physica*, 7(4):284–304, 1940. doi: 10.1016/S0031-8914(40)90098-2.

[72] Matias I. Maturana, Christian Meisel, Katrina Dell, Philippa J. Karoly, Wendyl D'Souza, David B. Grayden, Anthony N. Burkitt, Premysl Jiruska, Jan Kudlacek, Jaroslav Hlinka, Mark J. Cook, Levin Kuhlmann, and Dean R. Freestone. Critical slowing down as a biomarker for seizure susceptibility. *Nature Communications*, 11:2172, 2020. doi: 10.1038/s41467-020-15908-3.

[73] Annika Hagemann, Jens Wilting, Bita Samimizad, Florian Mormann, and Viola Priesemann. Assessing criticality in pre-seizure single-neuron activity of human epileptic cortex. *PLOS Computational Biology*, 17(3):e1008773, 2021. doi: 10.1371/journal.pcbi.1008773.

[74] Sean Welleck, Ximing Lu, Peter West, Faeze Brahman, Tianxiao Shen, Daniel Khashabi, and Yejin Choi. Generating sequences by learning to self-correct, 2022.

[75] Amos Azaria and Tom Mitchell. The internal state of an LLM knows when it's lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 967–976, 2023. doi: 10.18653/v1/2023.findings-emnlp.68.

[76] Varvara Logacheva, Daryna Dementieva, Sergey Ustyantsev, Daniil Moskovskiy, David Dale, Irina Krotova, Nikita Semenov, and Alexander Panchenko. ParaDetox: Detoxification with parallel data. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6804–6818, Dublin, Ireland, 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.acl-long.469.

[77] Jonathan Kadmon and Haim Sompolinsky. Transition to chaos in random neuronal networks. *Physical Review X*, 5:041030, 2015. doi: 10.1103/PhysRevX.5.041030.

[78] Kanaka Rajan, L. F. Abbott, and Haim Sompolinsky. Stimulus-dependent suppression of chaos in recurrent neural networks. *Physical Review E*, 82:011903, 2010. doi: 10.1103/PhysRevE.82.011903.

[79] Ben Poole, Subhaneil Lahiri, Maithra Raghu, Jascha Sohl-Dickstein, and Surya Ganguli. Exponential expressivity in deep neural networks through transient chaos, 2016.

[80] Samuel S. Schoenholz, Justin Gilmer, Surya Ganguli, and Jascha Sohl-Dickstein. Deep information propagation, 2016.

[81] Cristian Bonatto and Jason A. C. Gallas. Periodicity hub and nested spirals in the phase diagram of a simple resistive circuit. *Physical Review Letters*, 101:054101, 2008. doi: 10.1103/PhysRevLett.101.054101.

[82] Ruedi Stoop, Philipp Benner, and Yoko Uwate. Real-world existence and origins of the spiral organization of shrimp-shaped domains. *Physical Review Letters*, 105:074102, 2010. doi: 10.1103/PhysRevLett.105.074102.

[83] Cameron Berg, Diogo de Lucena, and Judd Rosenblatt. Large language models report subjective experience under self-referential processing, 2025.

[84] Fredrik Johansson. Arb: Efficient arbitrary-precision midpoint-radius interval arithmetic. *IEEE Transactions on Computers*, 66(8):1281–1292, 2017. doi: 10.1109/TC.2017.2690633.