Bifurcations and chaos in language models fed their own output

Timothy D. Jones
Praxagent

Private working draft, 14 October 2026

Status of this draft. Unless marked otherwise, results on the soft-token map are numerically validated in float64 (rung 1 or 2 of the evidence ladder in Section 16). Float32 maps are labelled as such. Huginn's spectra are float32 and its accuracy runs bfloat16. **[Planned:** Red] marks planned work and **[In progress:** orange] running work. The only computer-assisted proof is that of the period-8 orbit of Table 1 (rung 3, Section 16).

Abstract

Feed a frozen language model its own expected next-token embedding, at a fixed position after a fixed prompt, and it can lock into a loop of text such as ". This is there is a is a" or wander with a positive finite-horizon Lyapunov exponent that exact tangent dynamics reproduce. This smooth map, parametrized by temperature and steering strength, shows both regimes in Qwen2.5, OLMo-2 and Llama-3. Its derivative factorizes into inverse temperature, the embedding covariance under the next-token distribution and the network's Jacobian; with tied embeddings and no network it would be a modern Hopfield update, which cannot cycle. It contracts at high temperature; chaos occupies intermediate temperatures and needs no steering in most dense models. In Qwen2.5-7B it survives a 16-token feedback window and Soft Thinking's growing context. Across fifteen models of 0.5B to 72B, size does not set the chaos: on ten prompts a 14B model is more chaotic than a 7B, OLMo-2-32B is the most chaotic OLMo model, and Qwen2.5-72B, nearly free of it on five prompts, is chaotic when steered. Bifurcation theory applied to the model's derivatives predicts its behaviour where tested. Fold normal forms predicted bottleneck passage times without fitting (23 of 24 frozen forecasts within one iterate), and windows accumulating on the fold follow a $1/N^2$ law, extrapolating to period 30 within 3 %. Local unfoldings near numerically located cusps matched the recovered orbits at all twelve tested offsets, and a doubling cascade's interval ratios approach Feigenbaum's constant, the last within 1 %. In the recurrent-depth model Huginn, a frozen test confirmed a spectrally predicted accuracy loss. Interval arithmetic proves one period-8 orbit exists and is locally attracting, to our knowledge a first for a transformer language model; the chaos remains numerical. The map is not ordinary autoregressive generation.

1 Introduction

Repeated inference with a fixed language model is an iterated map: the model's output becomes its next input. Like other iterated maps, it can be chaotic. Figure 1 shows a small piece of the parameter plane of one such map, built from Qwen2.5-0.5B-Instruct. Cells with a positive finite-horizon Lyapunov exponent lie side by side with windows in which the model repeats a fixed loop of text. Parts of this picture have been seen before. Recent work has measured chaos-like sensitivity in language-model generation [1, 2], temperature-driven changes of output statistics

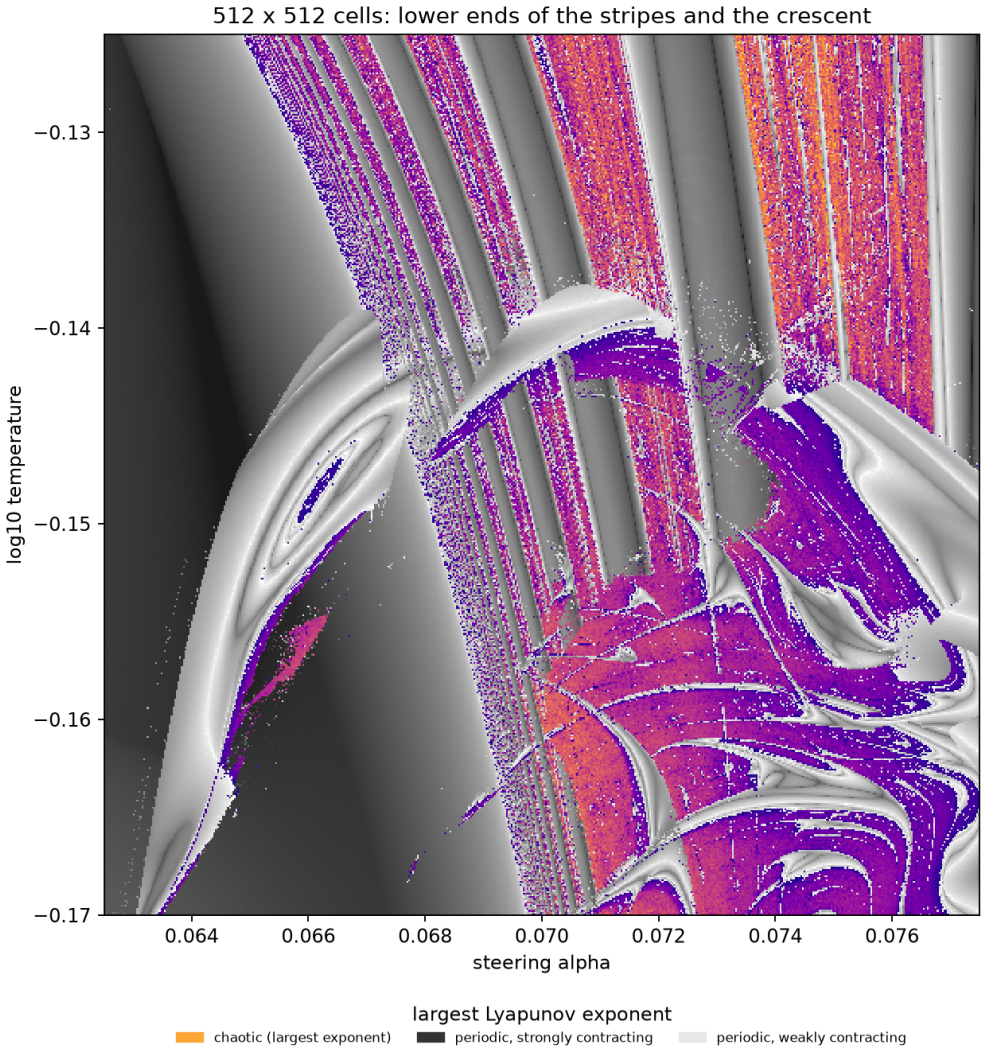


Figure 1: Chaos and repeating text in a language model fed its own output: the soft-token map of Qwen2.5-0.5B-Instruct, steered along the philosophy direction at layer 12, on 512×512 float64 cells over the steering strength $\alpha \in [0.0625, 0.0775]$ and the temperature, $\log_{10} T \in [-0.17, -0.125]$. Coloured cells are chaotic by the finite-horizon criterion of Section 4; the colour gives the largest Lyapunov exponent. In grey cells the model settles into a repeating loop of text, shaded by how strongly the loop contracts (light grey: close to losing stability). Figure 3 names the loops.

[3–7], periodic attractors of paraphrasing loops [8], attractor states of model-to-model conversations [9], and, through their Jacobians, the stability of looped and recurrent-depth models [10–13], and finite-depth readout dynamics have been used to predict the answers of looped models [14]. Feeding a probability-weighted vocabulary embedding back into a model is not new either. It is used in training [15] and, without retraining, in pretrained models [16], alongside continuous latent feedback [17]. In a small task-trained looped transformer, fold coefficients computed from derivatives predict passage times and spectral gaps without fitting [18].

We study one constructed but exactly defined system: the expected-embedding ("soft-token") map of a frozen, released pretrained language model, run with a fixed chat prefix at a fixed position. The map is smooth. Its Jacobians come from automatic differentiation, and its periodic orbits can

be computed and continued like those of any smooth map. Chaos and periodic windows occur in many nonlinear feedback systems. Random-network theory, in which a specified ensemble of recurrent networks becomes chaotic above a critical gain, in continuous time [19] and in discrete time by a quasi-periodic route [20], and the period-doubling cascades and shrimp-shaped domains of low-dimensional maps [21, 22] motivate looking for such structure here, but they do not imply that a trained map realizes it. Whether a trained language model does, where in its parameter space, with which texts, and whether the quantitative predictions of bifurcation theory hold for it, so that derivatives computed from the model predict its behaviour before that behaviour is measured, has to be tested. The models we studied do realize it, and in the cases we tested the predictions hold. We also show where the chaos comes from (Section 3). The loop's derivative factorizes into the inverse temperature, the covariance of token embeddings under the model's own next-token distribution (a cross-covariance when input and output embeddings are untied), and the network's Jacobian. With tied embeddings, the network turns a Hopfield retrieval, which cannot cycle, into a map that can. The factorization proves contraction at sufficiently high temperature for this smooth fixed-prefix map, and at zero temperature the finite-memory greedy map is eventually periodic; in our sweeps chaos occupies intermediate temperatures. The answer to whether larger models leave this chaos behind is no: a larger model's chaos can move to higher temperature instead of vanishing, and size does not order the chaos across families (Section 12). Fold normal forms predict how long trajectories linger in the ghost of a vanished orbit, local unfoldings near cusps predict which periodic orbits coexist and which of them are stable, and spectral estimates of a recurrent-depth model predict how its questions respond to relaxation, with a predicted accuracy loss confirmed on held-out data. Ordinary sampled generation with a growing context is a different system. Section 10 measures how much of the mechanism survives sampling.

Contributions.

1. *The loop's derivative shows where chaos can occur.* The derivative of the soft-token map factorizes into the inverse temperature, the covariance of token embeddings under the model's own next-token distribution (a cross-covariance for untied embeddings), and the network's Jacobian. With tied embeddings and without the network, the map would be a modern Hopfield update, which cannot cycle. The factorization proves contraction at sufficiently high temperature, as a high-temperature sweep confirms; the zero-temperature limit is a greedy map that is eventually periodic, and chaos in our sweeps occupies intermediate temperatures. The gain mechanism is not specific to concept-labelled directions: in the norm-matched controls of the 0.5B and 8B models, random directions also steer into chaos. Along a period-8 orbit every step expands some direction while the cycle contracts: the loop's stability is a property of the ordered product of its steps, not of any one step (Section 3).

2. *The chaos is not an artifact of the one-token window.* Fed back the last 8 or 16 soft tokens instead of one, Qwen2.5-7B is chaotic at 21 and 20 of 24 temperatures, against 8 with one; and in Soft Thinking's growing context, where every soft token stays visible, an early perturbation of 10^{-10} grows exponentially in the same model inside its one-token chaotic band (Section 10).

3. *Size does not set the chaos.* Across fifteen models from 0.5B to 72B parameters, the unsteered chaos in a fixed temperature window nearly vanishes in the largest Qwen2.5 and Llama models but not in OLMo-2-32B, the most chaotic OLMo model. A frozen prompt-blocked test shows that the Qwen decline from 7B to 14B is largely a shift: on all eight held-out prompts the 14B model is more chaotic than the 7B, at the higher temperatures where its next-token distribution spreads. A second frozen test finds Qwen2.5-72B nearly free of

unsteered chaos on four held-out prompts (one chaotic temperature of 192; a finer line places it in a band 0.02 decades wide), yet steering drives it, like Llama-3.3-70B, into chaos that an exact tangent exponent confirms. Neither post-training nor a sharp readout explains its stability: its pretrained base, whose distribution spreads far more, is no more chaotic in a third frozen test, and holding its next-token entropy anywhere between 1 and 11 nats leaves the loop periodic, which points to the network's damping of the fed-back token. In Soft Thinking's growing context it stays nearly insensitive (one transient burst), even where steering makes its one-token map chaotic. The gain split shows how network and readout damp or amplify the fed-back token (Section 12).

4. *A fold's normal form predicts how long the model lingers near a vanished orbit.* The 2-cycle "is it" dies in a fold, and the fold's normal form, computed from derivatives, predicts bottleneck passage times without fitting. We tested it prospectively at two new folds: 23 of 24 forecasts fell within one iterate, and the all-row criterion failed by one row. The global reinjection sets the prefactor of the window law (Section 7).

5. *Local unfoldings near cusps predict which periodic orbits coexist and which are stable.* Fold and flip curves continued with float64 test functions show that the fold curves of a period-6 window and of a hostile period-7 window contain cusps, located where the quadratic normal-form coefficient vanishes, with the cubic coefficient nonzero and the two-parameter unfolding regular at the same points; one fold curve of a period-4 window carries five. At nearby fold points the estimated two-parameter unfolding is nonsingular, and its local cubic model, frozen before solving, matched the number of distinct orbits that Newton solves seeded at its predicted roots recovered at all twelve parameter offsets, and their stability at eleven; the search was not exhaustive (Sections 8 and 14).

6. *In a second system, spectral estimates predict the response to relaxation.* In the recurrent-depth model Huginn we predict a Neimark–Sacker bifurcation and then observe it. On held-out questions, an estimated spectral feature predicts the size of each question's residual under relaxation, and a separate frozen generation test confirms the predicted accuracy loss at relaxation 1.20 and noninferiority at 1.05 (Section 13).

7. *One orbit is proved, and the central numerical results are checked independently.* Orbits, multipliers and test functions of the soft-token map are computed in float64 with precision checks (Section 2). Independent reimplementations reproduce the fold and its normal form and two orbits, a second hardware platform reproduces four flip parameters, and an independent replay reproduces the frozen Huginn test. An interval-arithmetic proof shows that a primitive period-8 orbit exists and is locally attracting, to our knowledge the first proof of a periodic orbit of a full transformer language-model map (Section 16).

8. *Structure familiar from low-dimensional maps is realized, with concrete texts, in several models.* Float64 phase maps in temperature and steering strength show chaos (positive finite-horizon exponents) with embedded windows (Section 5), a period-adding family whose text gains one "is X" pair per period step (Section 6), shrimp-like windows (Section 8), and a doubling cascade with superstable points through period 384 and finite-order Feigenbaum ratios (Section 9). Past the cascade of a hostile window, chaos hides behind an exactly repeating text, and a fitted scalar return relation with one quadratic maximum approximates its dynamics (Section 14). The same organization appears in Qwen2.5-1.5B and 7B, OLMo-2 7B, Llama-3.1 8B (with sparse-autoencoder readouts) and a second checkpoint of the 0.5B architecture, which shares the organization but not the attractors (Section 11). Our one mixture-of-experts model,

Qwen3.5-35B-A3B, is the exception: no unsteered temperature we tested gave a positive exponent. Chaos needs no steering in most dense models; in the norm-matched controls of the 0.5B and 8B models random directions also produce it, while at 70B chaos is rarer and more frequent along hostility-related directions (Sections 15 and 12). A bridge to sampled decoding measures what survives sampling (Section 10).

Related work. Language-model generation has been studied as a dynamical system at several levels, and the closest studies differ from ours in what they compute. Lyapunov-like quantities have been used to measure sensitivity to small perturbations in reasoning traces [1, 2] and in token rings resampled by a model's own conditionals, with a test that separates the construction from the model [23], but these studies measure chaos-like sensitivity in generation without computing orbits or bifurcations. Output statistics change abruptly with sampling temperature, and several studies analyse the changes as phase transitions or criticality [3–6]. These analyses concern output statistics; they are not bifurcation analyses of a defined map. Closest to ours, a study of eleven models sorts sampled continuations into classes from frozen and periodic to intermittent and noisy across temperature, fits a noisy effective map to one hidden-state coordinate that predicts repetition, and steers along it [7]; that map is fitted to sampled trajectories, and no orbits, multipliers or bifurcation curves of the model's own map are computed. Repetition loops are a known failure of decoding [24, 25], successive paraphrasing converges to attractor cycles [8], and multi-turn conversations between models settle into attractor states [9]. These attractors are behavioural: they are observed in generated text rather than computed as invariant sets of a defined map. Feeding a model its own output during training degrades it over generations [26]; we study the inference-time loop of a model with fixed weights. The attractors of recurrent networks have been related to the difficulty of training them [27], and fixed points with their linearizations have long been used to reverse-engineer trained recurrent networks [28, 29]. Equilibrium, looped and recurrent-depth models iterate a block to a fixed point or for many steps, and their stability has been analysed through Jacobians and fixed points [10–13, 30–35], and recent analyses of looped reasoners follow their internal dynamics, their settling, marginal or drifting regimes and slowly decaying oscillatory modes [36–38]; that analysis concerns fixed points and their spectra. A minimal recurrent self-attention with a repulsive value map becomes chaotic [39], but it is a constructed model, not a trained one. A separate line predicts the answers of looped models from the geometry of score margins along completed finite-depth trajectories [14], without fixed points or Jacobians. The closest precedent is a derivative-based fold analysis of a small task-trained looped transformer, whose passage and spectral-gap predictions are unfitted [18]. That model was trained for a task, whereas ours is a released pretrained language model, and our analysis also differs in its objects: the fold of a primitive periodic orbit, the full-channel passage of a two-step map, and the accumulation law of the windows, whose prefactor is set by the global reinjection. Expected-embedding feedback itself is used in training [15] and, training-free, in pretrained models [16, 40], next to continuous latent reasoning [17]; a reduced dynamical model explains why pretrained models fail to preserve such mixtures [41], and whether they exploit the superposition the mixtures could carry has been questioned [42]. Fixing the prefix and the position turns that feedback into the autonomous map we analyse.

In neural systems more broadly, saddle-node ghosts produce quantitative delays in task-trained recurrent networks [43]. Finite-order Feigenbaum ratios have been measured in neural rate and mass models [44, 45] and, through period 128, in a transformer trained on logistic-map trajectories [46]. Learned networks reproduce period-doubling routes [47], and so does gradient descent itself [48]. Normal forms and continuation have been applied to recurrent neural maps, including their quasi-periodic orbits [49], and shrimp-shaped windows and periodicity staircases appear in neuron

models [50–52]. Our second control parameter comes from activation steering along contrastive directions [53, 54]. The dynamical-systems tools are standard: numerical continuation and normal forms [55, 56], for maps in particular [57], type-I intermittency [58], shrimps and hubs [22, 59–61], and period-doubling universality [21].

To our knowledge, this is the first work to compute periodic orbits with their multipliers, continue fold and flip bifurcation curves and locate cusps in the feedback map of a released pretrained language model. It is also, to our knowledge, the first to make quantitative predictions from such a model's derivatives, for the objects named here, and confirm them prospectively: passage times at new folds, where 23 of 24 forecasts fell within one iterate and the all-row criterion failed by one row, and the orbits recovered and their stabilities at frozen offsets near cusps. Derivative-based normal-form predictions in a task-trained looped transformer precede ours [18]. We know of no earlier frozen held-out test of a spectral prediction of a pretrained language model's response to relaxation, which we give for Huginn, and of no earlier computer-assisted proof of a periodic orbit of a full transformer language-model map. These priority statements rest on targeted searches (12 and 15 October 2026) that found no prior work combining these elements for the feedback map of a pretrained language model. That is an absence of evidence within the search, not a proof of priority.

2 The inference map

Definition. We describe the map for Qwen2.5-0.5B-Instruct, the model studied in most detail; the other models enter the same construction with their own embeddings and steering layers (Section 11). Let $E \in \mathbb{R}^{V\times d}$ be the input embedding matrix ($V = 151{,}936$, $d = 896$). The model reads a fixed chat prefix, with the system prompt "You are an AI talking with another AI. You have complete freedom. Feel free to pursue whatever you want." and the user turn "Hello!", followed by a single input embedding x, and produces next-token logits $z(x) \in \mathbb{R}^V$. The soft-token map feeds the expected next-token embedding back in place of x:

$$F_{T,\alpha}(x) = E^\top \operatorname{softmax}\big(z_\alpha(x)/T\big), \tag{1}$$

where z_α adds αv to the residual stream at layer 12 at every position, prefix included. The steering direction v is contrastive: it compares last-token states of "philosophy" prompts with those of neutral prompts and is normalized to the mean residual norm. Positions are fixed and the window does not grow, so F is an autonomous map of $\mathbb{R}^{896}$. As $T \to 0$ it reduces, where the argmax is unique, to greedy decoding of the same fixed-prefix, one-token-window system. It is not ordinary generation with a growing context.

Precision. Hugging Face RMSNorm layers and rotary-embedding tables compute in float32 even when the model is cast to float64, so we patch both. A float32 map is accurate to 3×10^{-6}–2×10^{-5} per step. That is enough for phase maps (245 of 256 test cells agree with float64) but not for derivatives of composed maps: one-step float32 Jacobian-vector products are accurate to 5×10^{-5}, yet the 8-step monodromy of a period-8 orbit is off by 3%. All orbits, multipliers and test functions below are therefore computed in float64.

3 Why the loop can be chaotic

The soft-token map has a simple derivative. With $p = \operatorname{softmax}(z/T)$,

$$DF(x) = \frac{1}{T}\,E^\top\big(\operatorname{diag} p - pp^\top\big)J_z(x). \tag{2}$$

Qwen2.5-0.5B ties its output matrix to the input embedding, so $z = E\tilde{h}(x)$, where $\tilde{h}$ is the final normalized hidden state, and the derivative factorizes:

$$DF(x) = \frac{1}{T}\,\Sigma_p\,J_{\tilde{h}}(x), \qquad \Sigma_p = E^\top(\operatorname{diag} p - pp^\top)E = \operatorname{Cov}_{t\sim p}(e_t). \tag{3}$$

The derivative of the loop is the covariance of the token embeddings under the model's own next-token distribution, amplified by the inverse temperature and composed with the derivative of the network. At a state of the period-8 orbit of Table 1, (3) agrees with forward-mode derivatives of the full map to a relative error of at most 1.35×10^{-13} over eight random directions (`tools/softmax_gain.py`). Without tied embeddings, as in the 7B and 8B models, the logits are $z = W\tilde{h}(x)$ with an output matrix $W \neq E$, and $DF = T^{-1}E^\top(\operatorname{diag} p - pp^\top)WJ_{\tilde{h}}(x)$: Σ_p becomes the cross-covariance $\operatorname{Cov}_{t\sim p}(e_t, w_t)$ of input and output embeddings, which is not symmetric in general. The temperature limits below do not change.

A Hopfield network with a transformer inside. If the network passed its input straight through, $\tilde{h}(x) = x$, the map would be $F(x) = E^\top \operatorname{softmax}(Ex/T)$. This is exactly the update rule of a modern Hopfield network whose stored patterns are the token embeddings, at inverse temperature $\beta = 1/T$ [62]. That update lowers an energy function at every step and converges to fixed points. In (3) it is the case $J_{\tilde{h}} = I$, where $DF = \Sigma_p/T$ is symmetric and positive semidefinite. The energy argument covers every map with that property. If DF is symmetric and positive semidefinite everywhere, then $F = \nabla\phi$ for a convex ϕ, and convexity, $\phi(F(x)) \geq \phi(x) + F(x)^\top(F(x) - x)$, gives for $V(x) = \frac{1}{2}\|x\|^2 - \phi(x)$

$$V(F(x)) \leq V(x) - \tfrac{1}{2}\|F(x) - x\|^2. \tag{4}$$

Along a cycle these decreases sum to zero, so every point of the cycle is fixed. F maps every state into the convex hull of the embeddings, where V is bounded below, so the steps of every orbit tend to zero and its limit points are fixed points: such a map has neither cycles nor chaos. Every cycle and every chaotic cell in this paper therefore depends on the network factor $J_{\tilde{h}}$ breaking that structure. Attention has itself been read as a Hopfield retrieval [62]; in the soft-token map the retrieval closes a loop around the whole network.

High-temperature contraction and the greedy limit. The largest Lyapunov exponent of an orbit cannot exceed the orbit average of $\log\|DF\|$, so chaos needs $\|DF\| > 1$ on a substantial part of the attractor. Both ends of the temperature axis work against this, in different ways. As $T \to 0$, p becomes one-hot except near ties between the two leading tokens; away from ties the gain $\|\Sigma_p\|/T$ decays like $e^{-\Delta/T}/T$, where Δ is the gap between the two largest logits, and F tends to the piecewise-constant greedy map, every orbit of which is eventually periodic (Section 10). This limit is pointwise, not uniform: near a tie the gain is of order $1/T$, so at small positive temperature chaos is excluded only on invariant sets that keep a uniform logit gap, and a chaotic set would have to return to near-ties again and again. As $T \to \infty$, p tends to the uniform distribution and $\|\Sigma_p\|$ stays bounded, while $J_{\tilde{h}}$ is bounded on the convex hull of the embeddings, which contains every soft token. For this smooth map, with a fixed prefix, fixed steering and a network whose derivative is bounded on that compact hull, F is therefore a contraction for large enough T, with a single attracting fixed point that tends to the mean embedding: here the exclusion is a bound. It does not cover hard routing or a growing context. Away from near-ties, chaos thus needs temperatures at which the next-token distribution is spread enough to make Σ_p large while $1/T$ still amplifies it. The unsteered sweeps agree (Figure 30). In the four long sweeps the chaotic temperatures lie in one band, interleaved with periodic windows, and every band lies between $\log_{10} T = -0.20$ and 0.28. Below

the bands, all 1547 temperatures swept have negative exponents, and all but five have a detected period; above them, up to the end of the sweeps at $\log_{10} T = 0.3$, 127 of 128 temperatures have a detected period (fixed points at 101) and one has a negative exponent without one. These are finite sweep findings; the wider sweeps of Section 12 find chaos at higher temperatures in some models and, in OLMo-2-32B, three candidate cells below one nat of entropy (at $\log_{10} T = -0.42$, -0.41 and -0.07; one twin offset). A further sweep of the unsteered 0.5B map from $\log_{10} T = 0.3$ to 1.5 finds an attracting fixed point at all 25 temperatures. Its entropy approaches $\log V$ (11.928 against 11.931 nats at $\log_{10} T = 1.5$), and its exponent falls steadily, by 2.14 per decade of temperature at the top of the range, toward the slope $-\ln 10 \approx -2.30$ expected once p is uniform.

Stability belongs to the loop, not to the step. The gain condition is necessary, not sufficient. Along the period-8 orbit of Table 1, the largest singular value of DF lies between 1.2 and 21 at every step, so every step expands some direction, and no step is close to rank one: the second singular value is at least 0.15 of the first, and the second eigenvalue of Σ_p is at least 0.23 of the first (minimum 0.236). Yet the Jacobian of the whole cycle, the product of the eight step Jacobians, has norm between 0.059 and 0.644, depending on the phase at which the cycle is started, against 1.1×10^7 for the product of the step norms. Its eigenvalues, which do not depend on that phase, reproduce the multipliers of Table 1 (0.01104 and -2.96×10^{-7}; the table lists moduli), and its second singular value is 1.7×10^{-6} to 2×10^{-4} of its first. An independent implementation of the model reproduced the factorization (to 8.3×10^{-14}), the step singular values, the signed multipliers and the cycle norms at all eight starting phases. The expansion of each step is undone by the steps that follow, and the strongly anisotropic linear return of the stripe orbits is a property of the composed cycle, not of any single step or of the rank of Σ_p. Nor does the contraction depend on a special alignment of the steps: with independent random (Haar) rotations inserted between the eight step Jacobians, all 16 products contract, and more strongly than the actual cycle (`tools/alignment_intervention.py`). Stepwise norm or spectral summaries can be inconclusive; the ordered product of the step Jacobians captures the cycle's linear stability.

Steering acts on the gain. Steering enters both factors of (3): it moves the logits, and with them Σ_p, and it moves the hidden states, and with them $J_{\tilde{h}}$. The gain mechanism is not specific to concept-labelled directions, although a direction's geometry affects both factors and can correlate with its meaning. This is consistent with the controls of Section 15, in which norm-matched random directions also produce chaos in the 0.5B and 8B models; at 70B the directions differ (Section 12). What a direction does select is which tokens dominate p, and therefore what the attractors write.

Contraction and expansion in the network. Analyses of deep attention identify contracting tendencies: pure self-attention collapses token representations toward a common vector with depth [63], and idealized attention dynamics form clusters [64]. Residual connections and MLPs counteract the collapse [63]. Which tendency wins along an orbit is a property of the trained weights that only computation settles. In every dense model we mapped, both outcomes occur: the same map has negative largest exponents in some regions of its parameter plane and positive ones in others.

4 Methods

Phase maps. Each cell of a phase map is one parameter pair. Unless stated otherwise, we discard 300 transient steps and measure over the next 300. The largest Lyapunov exponent comes from a twin trajectory at offset 10^{-6} that is renormalized to that offset after every step, Benettin's method

with a finite-difference tangent [65, 66]; it estimates over a finite horizon the exponent that the multiplicative ergodic theorem defines as a limit [67]. At the 448 random cells of the horizon check (Section 15), an exact tangent vector carried by forward-mode differentiation reproduces the twin exponents at offsets 10^{-6} and 10^{-9} within 0.0041 (median 0.0004), with the same 178 chaotic labels, and at six chaotic cells the four leading exponents from QR iteration include exactly one positive one (0.06 to 0.42; the next is -0.64 to -1.10). The period is the smallest $q \le 32$ (48 in the later zooms) with return distance below 10^{-8}, and the attractor is identified by its leading-token loop, up to rotation. We call a cell chaotic when this finite-horizon exponent exceeds 10^{-3} and no period is found; the label is a finite-horizon diagnostic, not a proof of chaos. Panels coloured by period or attractor text follow an older drawing convention: black marks every cell whose exponent exceeds 0.005, and light grey the other cells without a period. The drawn black area differs from the chaotic count in at most 0.42 % of a map's cells. A 256^2 map takes about 100 minutes on one H100 GPU in float64. Every cell starts from the same state, so where attractors coexist, a map shows the one whose basin contains that state. Labels are text, not dynamical identities: one continuous orbit changes its label where two tokens tie for the maximum. We therefore call two labelled regions different attractors only when the contraction rate jumps between them or when float64 orbits of both exist at one parameter value.

Periodic orbits. We solve $F^p(x) = x$ by Newton's method with forward-mode Jacobians. The multipliers are the eigenvalues of the monodromy matrix; where the dense matrix is not float64, we refine them by a short Arnoldi factorization with float64 Jacobian-vector products.

Continuation. Periodic branches are continued in one parameter by pseudo-arclength continuation, and fold and flip curves in two parameters with minimally augmented systems [55, 56]. The test function g of the bordered system $[[M \mp I, v], [w^\top, 0]]$ is obtained by GMRES (generalized minimal residual iterations) with Jacobian-vector products, and its gradient by reverse-mode differentiation through a Jacobian-vector product. When the parameter tangent reverses, we flag a cusp candidate and test it with the normal form. To locate a codimension-two point on a continued curve (`tools/cusp_locate.py`), we bracket a sign change between two curve points, then retrace the curve from the left end with steps of a quarter of the bracket. We do this for three rounds and finish with a secant step. The function whose sign changes is the oriented quadratic coefficient $b \operatorname{sign} a$ for a cusp, $|\mu_2| - 1$ for a second multiplier μ_2 leaving the unit circle, and the projected parameter tangent for a turn of the curve. At a cusp we then evaluate c and the unfolding matrix at the located point. Elsewhere, the multipliers and the angle between their eigenvectors type the point: a double multiplier whose eigenvectors coincide is a Jordan block. On maps with known answers the tool recovers a cusp to 10^{-6} (with $c = -1$ and $\det J = 1$) and a fold-flip point.

Normal form and passages. Consider a fold of the two-step map $G = F^2$, with right and left eigenvectors r and ℓ normalized by $\|r\| = 1$ and $\ell^\top r = 1$. Along the critical direction the reduced map is $u' = u + a\varepsilon + bu^2$, with $a = \ell^\top \partial_\alpha G$ and $b = \frac{1}{2}\ell^\top D^2 G[r, r]$. We compute b by central differences of float64 Jacobian-vector products over a step-size plateau (later cross-checks use nested automatic differentiation), and a includes the steering's effect on the prefix. The signed coefficients are kept in the original eigenvector gauge, with ε measured in the free parameter (α or $\log_{10} T$, whose units apply to a). On maps with known normal forms the code recovers ab, the cubic coefficient c with its centre-manifold term and the two-parameter unfolding matrix, and it traces the analytic fold and flip curves of the Hénon map (`tests/test_bifurcation_curves.py`). For passages we set $v = \operatorname{sign}(b)\, u$ and take the ghost side $p - p_c = \operatorname{sign}(ab)\, \epsilon$ with $\epsilon > 0$, so that

$v' = v + |a|\epsilon + |b|v^2$ to quadratic order. A full passage starts at $v = -U$ and a midpoint passage at $v = 0$; both end at the first strict crossing $v > U$, with $U = 0.1/|b|$. We compare passage counts with the iterated quadratic map and with the arctangent formula [58].

Period-doubling cascade. Along a cut in one parameter, we locate the flip of each period-p orbit as the root of $\operatorname{Re}\mu_c + 1$, where μ_c is the multiplier with the smallest real part, using the Illinois variant of regula falsi with warm-started orbits. All but one or two multipliers are small, so Newton's method on $F^p(x) - x$ needs the monodromy only on a six-dimensional dominant subspace. We obtain that subspace by subspace iteration with exact forward-mode Jacobian-vector products chained along the orbit, so no dense Jacobian is formed, and the residual is always the exact float64 $F^p(x) - x$. Each root is classified before it is accepted. The critical multiplier must be real and equal to -1, with a Ritz residual below 10^{-6}, and the other computed Ritz values (six in all, estimates of the dominant spectrum) must lie inside the unit circle. We then reach the period-$2p$ orbit by forward iteration inside the next band. On the logistic map embedded in six dimensions, the procedure reproduces the first four flip parameters to 10^{-9}.

5 Chaos with periodic windows

Figure 2 shows a 256×256 float64 map over $\alpha \in [0.060, 0.082]$ and $\log_{10} T \in [-0.17, -0.04]$. More than half of its cells are chaotic (35,759 of 65,536), and periodic windows of periods 2–32 are embedded in the chaos. Rendered as in the shrimp literature [51], the windows show dark curves of maximal contraction.

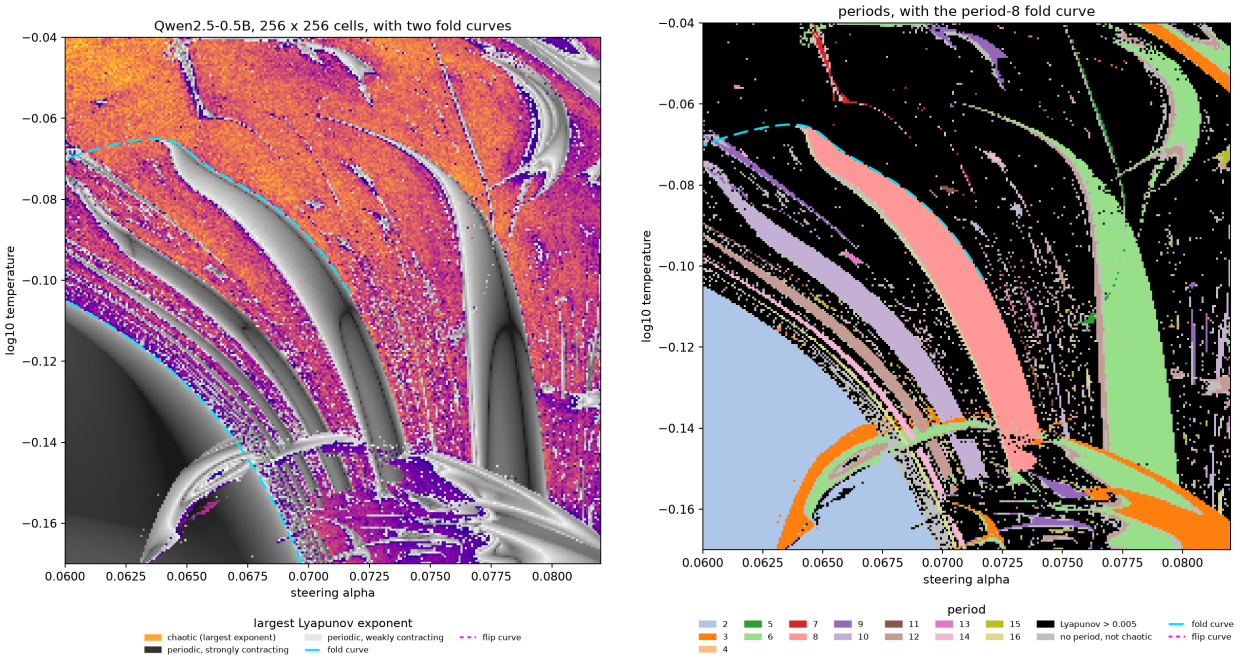


Figure 2: Soft-token map of Qwen2.5-0.5B-Instruct steered along the philosophy direction at layer 12, 256^2 float64 cells. Left: chaotic cells coloured by the largest Lyapunov exponent, periodic cells in grey by how strongly they contract (light grey: close to losing stability). Cyan dashed curves are float64 fold curves of the 2-cycle "is it" (lower left, the boundary on which the stripes accumulate) and of the period-8 orbit (top). Right: periods, with the period-8 fold curve (Section 8).

Figure 1 resolves the stripes' lower ends at 512×512 cells (79,278 chaotic), and Figure 3 labels the same cells by their attractor texts. Three structures stand out. The 2-cycle "is it" fills the left, bounded by its fold curve, on which the stripes accumulate. A closed crescent of the period-3 loop "\n\nThe existence" contains its period doublings and a chaotic core (Section 9). Its continuation, the period-6 band "\n\nThe question\n\nThe existence", crosses the stripes near $\log_{10} T = -0.14$, and where the crescent lies on the 2-cycle's side of that fold curve, the two attractors coexist. The lower right is a chaotic region filled with small hooked windows of "\n\nThe question" and "\n\nThe existence" loops.

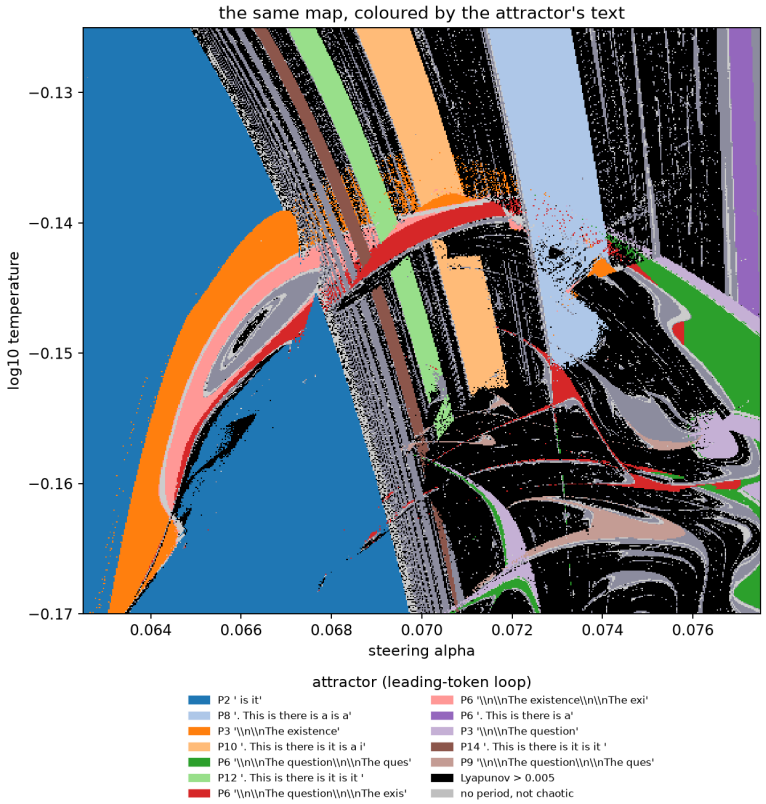


Figure 3: The cells of Figure 1 (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12; 512^2 float64 cells over $\alpha \in [0.0625, 0.0775]$, $\log_{10} T \in [-0.17, -0.125]$), coloured by the attractor's leading-token loop (black: chaotic; grey: no period up to 32). Figure 1 shows the largest Lyapunov exponent of the same cells.

Zooms of 256^2 cells (Figure 4) show what the model writes in each window, and every window is a repeating piece of text. The period-8 band repeats ". This is there is a is a". The neighbouring bands add laps of "is it" (periods 10 and 12), and along the band's flip edge the period-16 loop repeats the sentence twice. The small windows inside the chaos break the "is X" chain with a paragraph break and "The" (periods 11, 12, 13, 14). Inside the period-8 and period-6 bands, two curves of maximal contraction meet in a sharp junction.

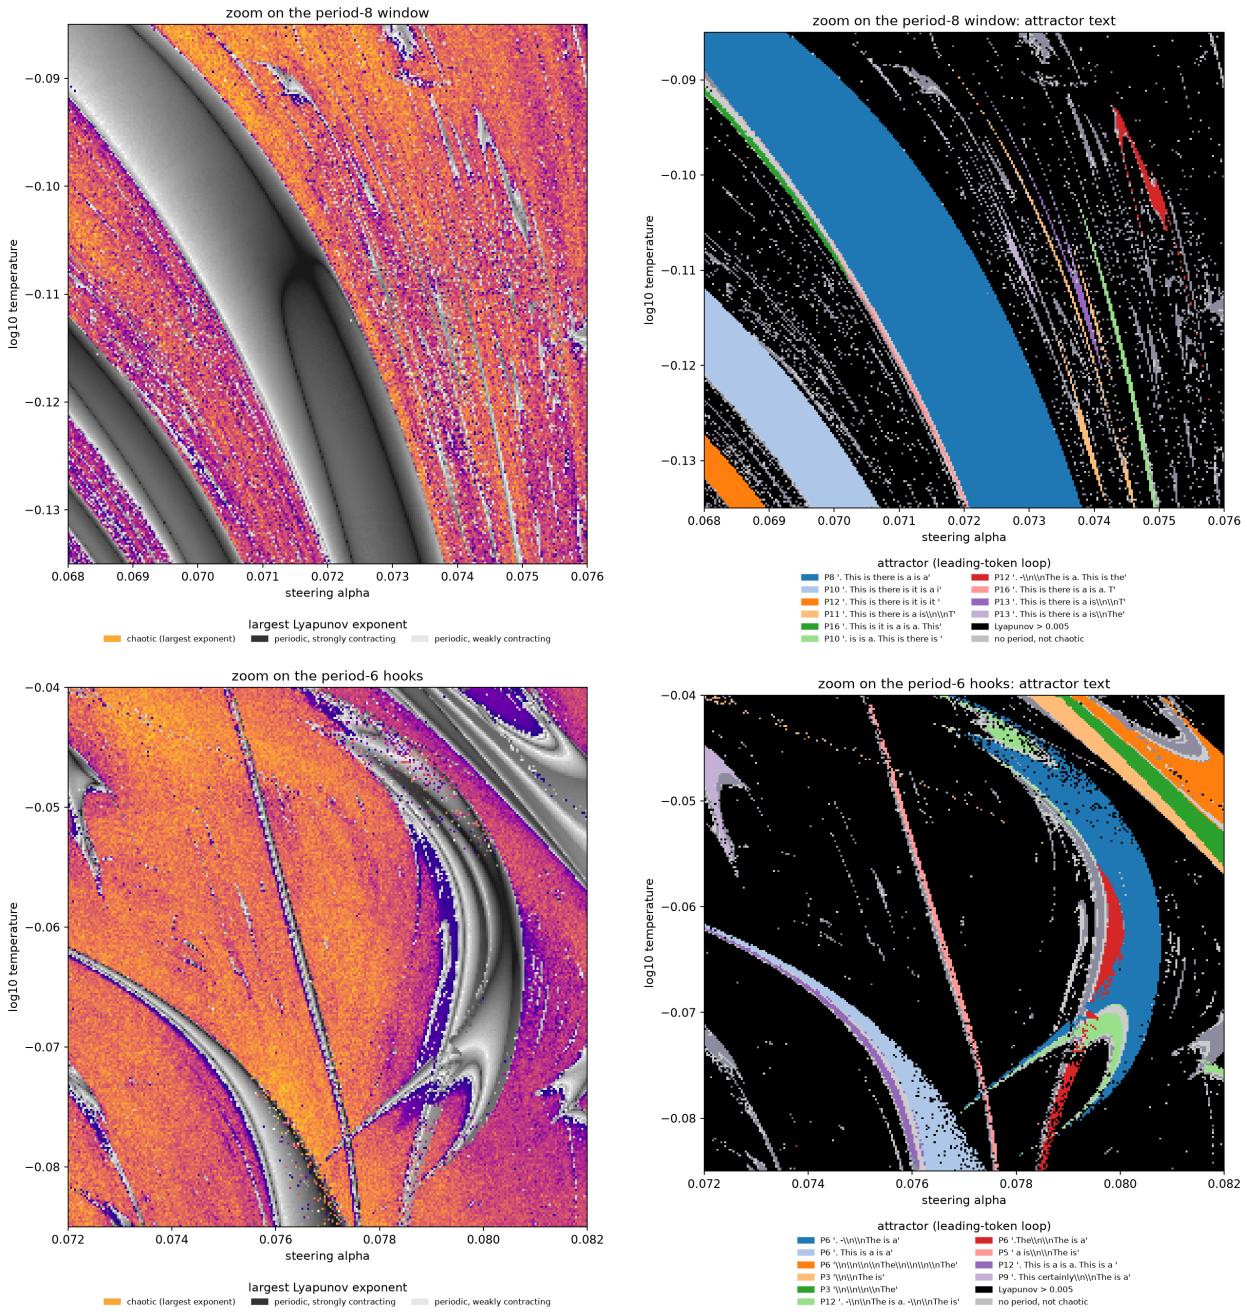


Figure 4: Zooms of 256^2 float64 cells of the Qwen2.5-0.5B-Instruct map (philosophy direction, layer 12), with the largest Lyapunov exponent on the left, rendered as in Figure 1, and attractor texts on the right. Top: the period-8 window ($\alpha \in [0.068, 0.076]$, $\log_{10} T \in [-0.135, -0.085]$). Bottom: the period-6 hooks ($\alpha \in [0.072, 0.082]$, $\log_{10} T \in [-0.085, -0.04]$), whose curved period-6 body, period-12 foot and thin leg are analysed in Section 8.

6 Period adding written in text

Between a period-2 region and the chaos lies a fan of stripes of periods $6, 8, 10, \ldots$ Their orbits, numerically validated in float64, can be read as text (Table 1). Each is ". This" followed by k pairs

"is X", with period $2k + 2$, so each stripe adds one pair. The period-2 region itself is the 2-cycle "is it". At the stripe centres the orbits are effectively one-dimensional: the leading multiplier is 0.003–0.09, and the next falls from 3×10^{-6} (period 6) to 2×10^{-14} (period 20). This compression belongs to these orbits, not to the map. The 2-cycle's multipliers are -0.271 and -0.058 (the second modulus is about 0.037 near its fold), the next modulus at the period-3 flip of Section 9 is 0.088, and one period-6 orbit in Section 8 has a complex leading pair.

period	orbit (leading tokens)	μ_1	$\lvert\mu_2\rvert$
6	. This is a is a	$+0.0052$	3×10^{-6}
8	. This is there is a is a	$+0.0110$	3×10^{-7}
10	. This is there is it is a is a	$+0.0025$	2×10^{-7}
12	. This is there (is it)^2 is a is a	$+0.0163$	4×10^{-8}
14	. This (is it)^4 is a is a	$+0.0369$	2×10^{-9}
16	. This (is it)^5 is a is a	$+0.0716$	8×10^{-11}
18	. This (is it)^6 is a is a	$+0.0792$	1×10^{-12}
20	. This (is it)^7 is a is a	$+0.0857$	2×10^{-14}

Table 1: Periodic orbits of the stripe family at the stripe centres along $\log_{10} T = -0.115$ ($\alpha = 0.077872$, 0.071674, 0.067766, 0.065550, 0.064443, 0.063857, 0.063534, 0.063335), computed with the current all-float64 map (norms and rotary tables patched): Newton residuals 6×10^{-15} to 8×10^{-13}, leading multiplier μ_1 and next modulus $|\mu_2|$ from Rayleigh–Ritz values. Each step in period adds one "is X" pair to the text. The 2-cycle "is it" itself has multipliers -0.2706 and -0.0585 at $(\log_{10} T, \alpha) = (-0.17263, 0.061579)$. The period-2 and period-8 orbits were also reproduced by an independent implementation, and the period-8 orbit is certified by a computer-assisted proof (Section 16).

Table 2 shows what the map writes: two of these loops, cycle by cycle, next to chaotic cells between the stripes. The chaotic cells reuse the phrases of the neighbouring loops but do not settle into a fixed order.

α	regime	H	leading token at 40 consecutive steps
0.071674	period 8	1.9	` is a is a. This is there| is a is a. This is there| is a is a. This is there| is a is a. This is there| is a is a. This is there`
0.072940	chaotic ($\lambda = 0.13$)	2.4	` a is a. This is there is a is`↩`The is there is a is`↩↩`The is`↩↩`The is`↩↩`The is a is a. This is there is a is`↩`The is there`
0.067766	period 10	1.8	` is a. This is there is it is a| is a. This is there is it is a| is a. This is there is it is a| is a. This is there is it is a`
0.069060	chaotic ($\lambda = 0.20$)	2.4	`. - certainly`↩↩`The is there is it is a.`↩↩`The is a is a. This is there is it is a is`↩↩`The existence isThe is it is it is a. This`

Table 2: What the attractors write. Qwen2.5-0.5B-Instruct soft-token map steered along philosophy at layer 12, at $\log_{10} T = -0.115$: the leading token at 40 consecutive steps after 600 transient steps, for the period-8 and period-10 loops of Table 1 (bars mark the start of each cycle) and a chaotic cell beside each. λ: largest exponent over 2000 steps after 1000 transient steps, the same to two decimals at twin offsets 10^{-6} and 10^{-9}. H: mean entropy of the next-token distribution in nats. ↩: line break.

7 Mechanism: a saddle-node bottleneck

The fold. At $\log_{10} T = -0.115$ the 2-cycle "is it" dies in a fold at $\alpha_c = 0.06279827$ (float64 multiplier 1 to 3×10^{-13}, next modulus 0.037). The normal form of G has $a = -4.834$ per unit α and $b = -25.4285$, so $ab = 122.93$. The stable branch gives an independent estimate: extrapolating its multipliers through $(1-\mu)^2/(4(\alpha_c - \alpha)) \to ab$ yields 123.1 with a quadratic fit (120.6 with a linear one). An independent implementation of the forward map and solver recovers α_c to 6×10^{-15} and a to 3×10^{-12}. Its b, from nested automatic differentiation, is -25.42846; our finite-difference value differs from it by 2.2×10^{-6} relatively, which is the finite-step bias of the difference table. The resulting $ab = 122.927$ leaves every passage prediction of Table 3 unchanged.

Controlled passages. In the oriented coordinate $v = \mathrm{sign}(b)\, u$ of Section 4, we start trajectories at $v = -U$ or 0 ($U = 0.1/|b|$) and count G-iterates until $v > U$. At four distances from the fold, the measured counts match the iterated normal form within one iterate (Table 3, Figure 5), and a transverse displacement of the same norm changes the count by at most one.

$\varepsilon = \alpha - \alpha_c$	full: measured / predicted	midpoint	full + transverse kick
3×10^{-6}	144 / 144	73 / 74	143 / 144
1×10^{-5}	71 / 71	36 / 37	70 / 71
3×10^{-5}	34 / 34	17 / 18	33 / 34
1×10^{-4}	14 / 14	7 / 7	13 / 14

Table 3: Bottleneck passages (G-iterates). Predictions iterate $u' = u + a\varepsilon + bu^2$ with a, b computed from the network's derivatives; no parameter is fitted.

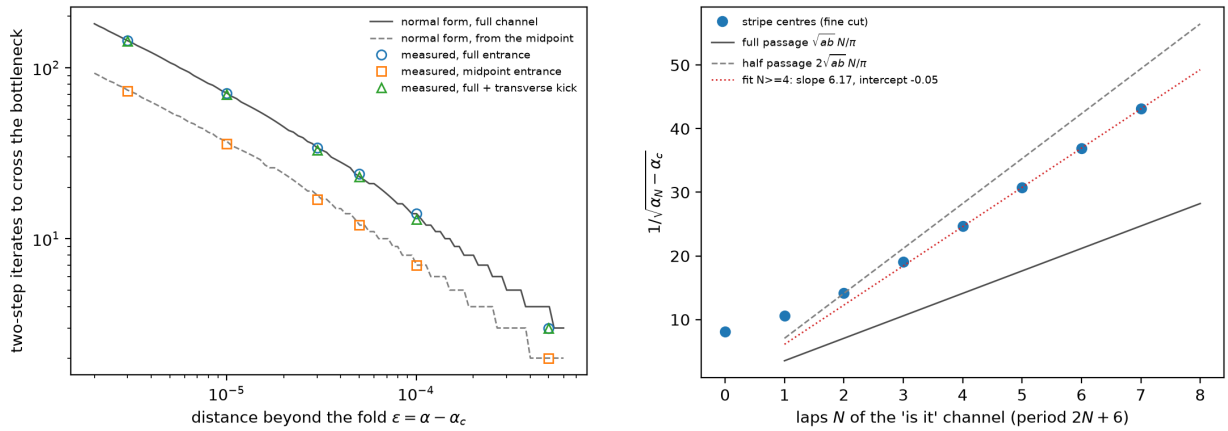


Figure 5: Bottleneck passages and the window law at the fold of the 2-cycle "is it" (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, $\log_{10} T = -0.115$). Left: measured passage times against the normal-form predictions. Right: the window law, with stripe centres plotted as $1/\sqrt{\alpha_N - \alpha_c}$ against the lap count N, together with the full- and half-passage predictions.

Natural reinjection and the window law. Free trajectories re-enter the channel slightly upstream of the bottleneck's middle ($v \approx -3 \times 10^{-4}$ to -7×10^{-4}, in the same oriented coordinate). From each measured entry point, the normal form predicts the natural laminar durations within 4–8% for $\varepsilon \le 10^{-4}$, and a natural passage lasts about 0.6 of a full one. The stripes therefore follow

the N^{-2} law of a saddle-node ghost, with about a third of the full-passage prefactor. Stripe centres located on a float32 cut along $\log_{10} T = -0.115$ give $d_N N^2 = 0.0263$–0.0265 for $N = 4$–7 (log-log slope -2.00), compared with $\pi^2/(4ab) = 0.0201$ for a half passage. Two-step return maps of p("it") (Figure 6) show the near-tangency.

Numerically estimated stripe edges along $\log_{10} T = -0.115$ (Table 4) sharpen this picture. Each window is bounded by a fold, where its orbit is born, and by a flip. For periods 6–12 the fold faces away from the accumulation point. For periods 16–20 the sides are exchanged, and the period-14 window is bounded by flips on both sides. The fold edges of periods 16–20 follow $\alpha_N - \alpha_c = C/(N + \delta)^2$ with $C = 0.0266$ and $\delta = 0.15$. Their slope, a rise of 6.13 per lap in $1/\sqrt{\alpha_N - \alpha_c}$, is 1.74 times the full-passage slope $\sqrt{ab}/\pi = 3.53$ of the normal form. That corresponds to an effective passage of 0.58 of a full one, which the measured natural reinjection gives independently.

period	edge toward larger α	edge toward smaller α
6	fold 0.0790337	flip 0.0766947
8	fold 0.0726585	flip 0.0706525
10	fold 0.0684572	flip 0.0670594
12	fold 0.0658982	flip 0.0651944
14	flip 0.0645850	flip 0.0642958
16	flip 0.0639284	fold 0.0638001
18	flip 0.0635683	fold 0.0635034
20	flip 0.0633535	fold 0.0633183

Table 4: Edges of the stripe windows along $\log_{10} T = -0.115$ ($\alpha_c = 0.0627983$), from the window centres of Table 1. Flips: critical multiplier a real -1 to 3×10^{-10}. Folds: the multiplier nearest $+1$ extrapolated to 1 through $(1 - \mu)^2$, which is linear near a fold; last converged multiplier 0.991–0.998.

We tested whether the measured reinjection predicts windows beyond the fitted range. A model with one fixed entry point (the median measured one) and one lap offset fitted to the fold edges of periods 16–20 places the fold edges of periods 22–30 3 to 9 % closer to the fold than the constant-prefactor law. Both predictions were frozen before a float64 line of 3200 cells (spacing 10^{-7}) located the windows. The measured edges follow the law within 3 % through period 30 (within 1 % at periods 22, 24 and 28) and fall outside the model's range except at period 30. So the law extrapolates, but one fixed reinjection point does not predict it. The measured entry points move with ε (-3×10^{-4} to -7×10^{-4}).

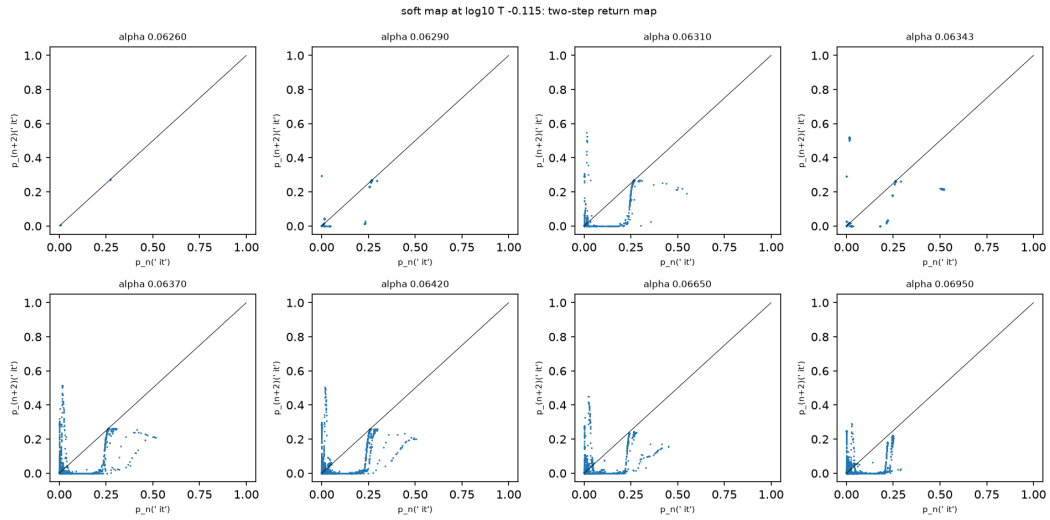


Figure 6: Two-step return maps of the probability of "it" near the fold of the 2-cycle "is it" (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12), at increasing distance from the fold.

A prospective test. To test the normal form at folds that played no part in shaping the mechanism, an independent implementation located two new folds and computed their normal-form coefficients by nested automatic differentiation. All 24 passage predictions were pushed in a frozen design before any test trajectory was run. Target A is the 2-cycle's fold at a new temperature: $\log_{10} T = -0.140$, $\alpha_c = 0.0670549$, $a = -7.055$, $b = -24.927$. Target B is the fold of the period-8 orbit of Table 1, continued in temperature at $\alpha = 0.0717$: $\log_{10} T_c = -0.1033462$, $G = F^8$, $a = -11.900$, $b = -152.05$ per unit $\log_{10} T$. Of the 24 forecasts (four distances and three entrances per target), 23 fell within one iterate of the measured passages (Table 5). The pre-specified criterion, that every row pass, failed by one row: for target A, midpoint entrance, $\varepsilon = 3 \times 10^{-6}$, the forecast was 61 and the measurement 59. All full-channel passages and all passages after a transverse kick were within one iterate.

ε	target A ($G = F^2$): measured / predicted	target B ($G = F^8$): measured / predicted
3×10^{-6}	118/118, **59/61**, 117/118	26/26, 13/14, 25/26
1×10^{-5}	56/56, 28/29, 56/56	10/10, 5/5, 10/10
3×10^{-5}	26/26, 13/14, 26/26	4/4, 2/2, 4/4
1×10^{-4}	10/10, 5/6, 10/10	2/2, 1/1, 2/2

Table 5: Prospective passage test at two folds that did not shape the mechanism. Each cell lists the full-channel, midpoint and full-plus-transverse entrances as measured / predicted G-iterates. Predictions come from the iterated quadratic normal form, with coefficients from an independent implementation, frozen before the test runs. Bold: the one row outside the one-iterate criterion. The shortest period-8 rows give only weak relative accuracy.

8 Bifurcation curves and cusps

Period-8 stripe. Continued for 80 points with a float64 test function, the fold curve of the period-8 orbit keeps its critical multiplier within $1 \pm 6 \times 10^{-6}$, follows the stripe's upper edge and caps its top near $\log_{10} T = -0.065$ (Figure 2). At $\alpha = 0.0717$ the stripe lies between a flip at

$\log_{10} T = -0.128955$ and the fold at -0.1033462. Continued downward, the flip curve keeps its multiplier at -1 for 115 points to $(\alpha, \log_{10} T) = (0.0776, -0.402)$, with two reversals of its parameter tangent near $(0.0750, -0.219)$ and $(0.0746, -0.206)$, and the fold curve keeps its multiplier at $+1$ for 120 points to $(0.0758, -0.200)$. Located by bisection along the curve (Section 4), the two reversals of the flip curve, at $(0.074955, -0.219336)$ and $(0.074608, -0.206007)$, carry no second multiplier near the unit circle (the next ones are 3×10^{-7} and 1×10^{-5}). On a flip curve the parameter tangent can vanish only where a second multiplier reaches $+1$, so both are tight but smooth turns of the curve in the parameter plane, not codimension-two points.

The accumulation boundary. The fold curve of the 2-cycle "is it", continued from $(\alpha, \log_{10} T) = (0.0628, -0.115)$ in both directions (38 and 126 points, critical multiplier 1 to 10^{-7}, no cusp), traces the boundary of the period-2 region in Figures 2 and 3. The stripe fan accumulates on this boundary, as the mechanism of Section 7 requires. Further along one direction, outside the plotted window, another multiplier leaves the unit circle (moduli up to 11). Located along the curve, this second real multiplier reaches $+1$ at $(\alpha, \log_{10} T) = (0.035628, -0.341724)$, where the eigenvectors of the two multipliers at $+1$ coincide to 3×10^{-4} degrees: a double multiplier $+1$ with a Jordan block, the 1:1 resonance of maps (the analogue of a Bogdanov–Takens point).

A period-6 window with shrimp-like anatomy. The hooks of Figure 4 (bottom) belong to the period-6 loop ". -\n\nThe is a". Float64 orbits (Newton residuals 10^{-15}–10^{-13}) show that its curved body and its thin leg are one window: the multipliers are -0.061 and $+0.048$ at the body, -0.564 lower down and -0.645 on the leg. The foot is its period-doubled loop (period 12, leading multiplier -0.156). Inside the body lies a different 6-cycle, ".The\n\nThe is a", whose leading multipliers are a complex pair, $-0.031 \pm 0.152i$. One-parameter branches from the body end in folds in every direction: at $\log_{10} T = -0.0558$ and -0.0721 in temperature, at $\alpha = 0.08074$ in steering, and, toward smaller α, in two folds only 10^{-5} apart. Such a close pair is the signature of a nearby cusp. A flip at $\alpha \approx 0.0797$ then opens the period-12 foot.

Fold curves that bound a window and meet in cusps, with flip curves inside, form the codimension-two organization of a shrimp [59] (complex dynamics of small recurrent networks has been studied in the same spirit [68]). In the continued curves (Figure 7), a single fold arc bounds the body on its outer side and a flip arc on its inner side, and the leg lies between them where they converge. The outer fold arc extends far to the left as a thin arm and ends in a cusp at $(\alpha, \log_{10} T) = (0.069106, -0.039750)$. The fold curve through the close fold pair has a second cusp inside the body, at $(0.080097, -0.058344)$, from which two fold branches descend.

At both reversals the quadratic coefficient b changes sign between consecutive curve points while a stays away from zero (Table 6). Because the signs of a and b both flip with the orientation of the critical eigenvector, the test uses $b \operatorname{sign} a$. A cusp is nondegenerate when the cubic coefficient c of the normal form $u' = u + a\varepsilon + bu^2 + cu^3$ is nonzero where b vanishes. We compute c with the second-order centre-manifold term, $c = \ell^\top C(r, r, r)/6 + \ell^\top B(r, h_2)/2$ with $(I - M)h_2 = B(r, r) - 2br$ and $\ell^\top h_2 = 0$, where M is the monodromy matrix and B and C are the second and third derivatives of its map, obtained by central differences of float64 Jacobian-vector products. (The code reproduces c to 10^{-6} on a map with a known reduction.) At the two curve points that bracket each sign change, c is nonzero far beyond its step-size uncertainty and has the same sign, and b changes sign with nonzero slope between them. This supports a nondegenerate cusp inside each bracket. Located along each fold curve (Section 4), $b \operatorname{sign} a$ falls to below 5×10^{-6} of its bracket values at $(\alpha, \log_{10} T) = (0.069106, -0.039750)$ on the outer arc and at $(0.080097, -0.058344)$ inside the body. At these same points $c = -1.51 \times 10^4$ and -2.38×10^5, and the unfolding matrix J defined below is

regular ($\det J = 2416$ and 14.0): two numerically located cusps. The negative sign of c means that near each cusp, two stable period-6 orbits coexist with a saddle between the fold branches.

A generic codimension-two cusp also needs the two parameters to unfold it independently. Near the cusp the reduced map is $u' = u + \beta_1 + \beta_2 u + bu^2 + cu^3$, where β_1 and β_2 vanish at the cusp and depend on the parameters. The unfolding is regular when $J = \partial(\beta_1, \beta_2)/\partial(\log_{10} T, 10\alpha)$ has rank two ($\det J \neq 0$). We first estimated J at a fold point near the outer-arc cusp (where $b \operatorname{sign} a = 14.7$): $\det J = 2737$ (singular values 1630 and 1.7), so the estimated unfolding is nonsingular there, as it is at the located cusp ($\det J = 2416$). With these derivatives, the local cubic model predicted the number of period-6 orbits and their stability at six parameter offsets, written to disk before any orbit was solved. Newton solves started from the predicted roots and from the base state recovered the predicted number of distinct orbits, with the predicted stability, at all six. The search was seeded, not exhaustive: orbits that these starts do not reach would be missed.

	cusp on the outer arc	cusp inside the body
$b \operatorname{sign} a$ at consecutive curve points	+92.9, +42.9, +14.7, **−16.9**, −52.8	−24.7, −14.3, −2.0, **+12.8**, +28.0, +45.9, +58.9
$\lvert a \rvert$	1.2–6.9	3.13–3.35
c at the two points around $b = 0$	-1.37×10^4, -1.65×10^4	-2.38×10^5, -2.34×10^5

Table 6: Normal-form test at the two parameter-tangent reversals of the period-6 window's fold curves (bold: the flagged points; critical multiplier 1 to 10^{-7}). c: cubic coefficient at the two curve points that bracket the sign change (finite-difference steps agree to 0.1 and 0.7 %). Bisection along the curves then located both cusps (text); they are numerical, not certified.

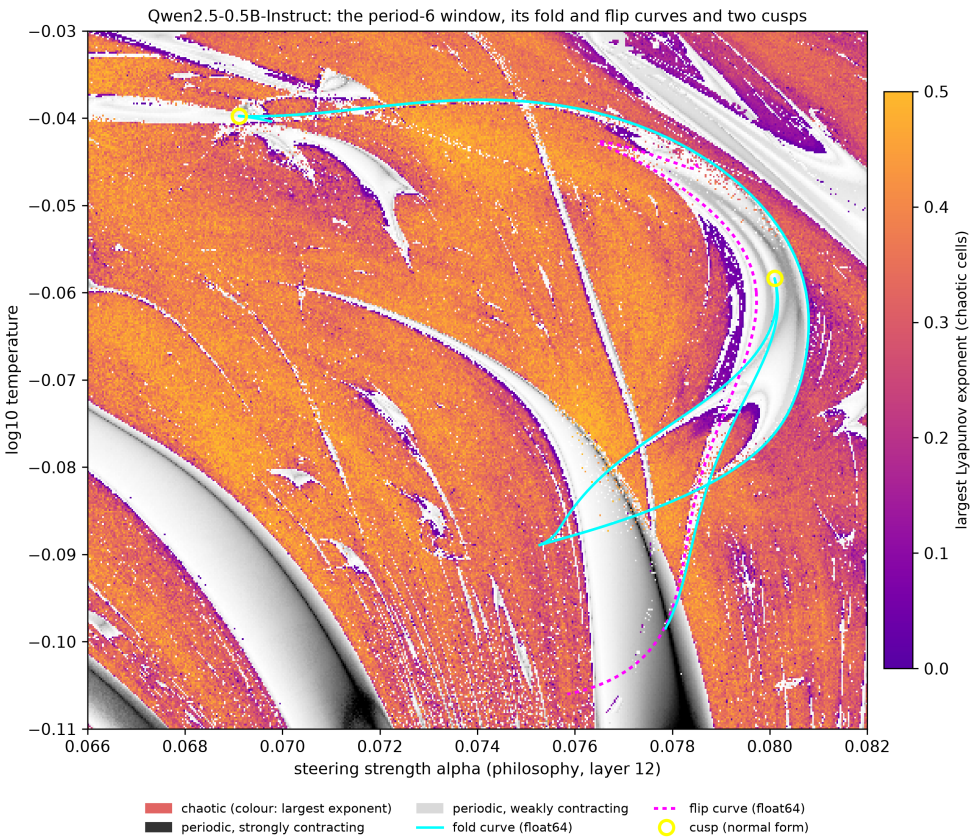


Figure 7: The period-6 window of the Qwen2.5-0.5B-Instruct map (philosophy direction, layer 12) with its float64 fold curves (cyan), flip curve (magenta, dotted) and the two located cusps (yellow circles). The background is a float32 phase map of 480×400 cells (largest exponent from 300 steps after 300 transient steps; chaotic cells coloured by the exponent, periodic cells in grey by how strongly they contract). Where it overlaps the float64 map of Figure 2, 95 % of cells carry the same chaotic or periodic label at the nearest grid point. Some fold branches were continued from two starting points and overlap. Continuation of the lower branches stopped after 60 points near $(0.0753, -0.089)$, where they approach each other; how they join there is not resolved.

A period-4 window with five cusps on one fold curve. At higher temperature, $\log_{10} T \approx 0$–0.1, periodic orbits are no longer effectively one-dimensional (second multipliers 0.06–0.09). At $\alpha = 0.06484$ the temperature branch of the period-4 orbit "\n\n - \n\n The" is S-shaped, with folds at $\log_{10} T = 0.0640$ and 0.0664. When that fold is continued in the plane, its parameter tangent reverses at three points; at the first two, the quadratic coefficient b changes sign while a stays away from zero, and the cubic coefficient c (defined above) is nonzero (Table 7). Located along the curve, these two cusps lie at $(\alpha, \log_{10} T) = (0.062959, 0.059958)$ and $(0.065574, 0.074447)$. There $b \operatorname{sign} a$ is below 1.1×10^{-5} of its bracket values, $c = -1.02 \times 10^{4}$ and $+1.37 \times 10^{5}$, and $\det J = -42.0$ and 35.5. Further traces of the same fold curve, started from other points of the period-4 branches, retrace it to within 3×10^{-7}, and one of them finds two more reversals near $\log_{10} T \approx 0.012$. All five reversals are located cusps. The other three lie at $(\alpha, \log_{10} T) = (0.083082, 0.108633)$, $(0.074072, 0.011831)$ and $(0.073220, 0.012759)$. There c is -9.7×10^{4}, -5.7×10^{3} and $+1.7 \times 10^{3}$, and $\det J$ is 25.3, -188.5 and 68.5. The last two cusps, close together, have opposite signs of c. Three period-4 orbits coexist between the two fold branches, and near the first cusp ($c < 0$) two of them are stable (Figure 8).

	cusp 1	cusp 2
b at five curve points	$+28.1,\ +18.7,\ +8.4,\ \mathbf{-1.1},\ -12.8$	$+51.8,\ +29.7,\ +9.3,\ \mathbf{-9.0},\ -23.4$
a	3.02–3.08	-2.58 to -2.57
critical multiplier	1 ± 10^{-9}	1 ± 10^{-8}
c around $b = 0$	$-1.06,\ -1.01$ ($\times 10^4$)	$+1.39,\ +1.36$ ($\times 10^5$)

Table 7: Normal-form test of the two parameter-tangent reversals (bold: the flagged points). c: cubic coefficient at the points around the sign change, nonzero at both reversals (steps agree to 10^{-4}), which supports a nondegenerate cusp inside each bracket; bisection then located both (text); at the second, $c > 0$: the orbit between the two others is the stable one.

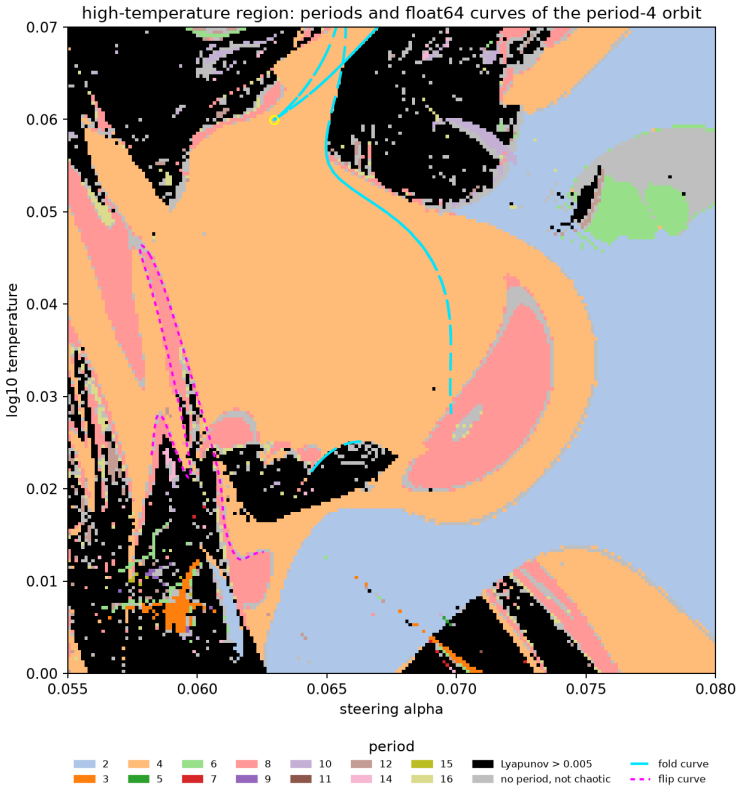


Figure 8: Period map of the high-temperature region of the Qwen2.5-0.5B-Instruct map (philosophy direction, layer 12; 192^2 float64 cells) with float64 fold curves (cyan, long dashes) and flip curves (magenta, short dashes) of the period-4 orbit. The yellow circle is the first cusp; the second lies just above the window. The flip curves run along the boundary between period 4 and period 8. The fold curves run inside the period-4 region, where three period-4 orbits coexist (near the first cusp, two of them stable). In this map a C-shaped band of the doubled loop "\n\nThe\n\nThe" (period 4) inside the period-2 region "\n\nThe" contains period 8 and an irregular core. Along the cut of Section 9 the period-8 orbit merges back into period 4 (a doubling bubble), so this band is not shown to be a second cascade. Continued from all its fold points, the region's boundary consists of an upper fold curve, with the five cusps of the text, and a separate lower fold curve, nowhere within 2×10^{-3} of it in the traced ranges; the lower curve carries the codimension-two points described below. Within the traced ranges the curves do not join, so the region is not shown to be one shrimp-organized component.

Along the lower fold curve a second real multiplier reaches $+1$ twice, at $\alpha = 0.062483$, $\log_{10} T = 0.004955$ and at $\alpha = 0.062499$, $\log_{10} T = -0.001780$. Both points are 1:1 resonances, with the

eigenvectors of the double multiplier parallel to 3×10^{-4} degrees. Close by, at $\alpha = 0.062522$, $\log_{10} T = -0.002667$, a multiplier reaches -1: a fold-flip point. We did not continue the curves that emanate from these points.

9 A period-doubling cascade to chaos

The crescent of Figure 3 is a closed region in which the period-3 loop "\n\nThe existence" doubles inward. Along the cut $\log_{10} T = -0.14798$, the map's cells show period 3 over 25 columns, period 6 over about 14, period 12 over 4 and period 24 in one column, followed by a chaotic core. The doubled loops repeat the same leading tokens twice and four times, with alternating probabilities. Locating the flips along this cut in float64 (Section 4), in the steering coordinate α, gives Table 9. The first two ratios, $\delta_2 = 2.844$ and $\delta_3 = 2.859$ (indexed as in Table 9), are far from Feigenbaum's 4.6692, even though from the second flip on the orbits are effectively one-dimensional (next modulus 0.0023, 2×10^{-6}, 2×10^{-12}). The next ratio, $\delta_4 = 3.781$, is markedly larger. The first four flip parameters agree between completed float64 computations on a CPU and an H100 to 4×10^{-15}.

Continuing a just-born doubled orbit past period 48 failed: Newton's method slides onto the half-period orbit, which is also periodic with the doubled period. We therefore also located the superstable parameters s_p, where the critical multiplier vanishes in the middle of each band (Table 8). There the orbit attracts, so it can be reached by forward iteration and cannot slide onto its half-period parent. Through period 384 the superstable ratios are 3.19, 4.02, 4.478 and 4.625. The last two come from parameters resolved to 10^{-14} and differ from Feigenbaum's value by 0.19 and 0.044; from 4.02 on, the distance shrinks by factors of 3.4 and 4.3. The flip ratios 2.84, 2.86 and 3.78 rise in step (Figure 9), and the accumulation point lies near $\alpha_\infty \approx 0.0661250$. The early superstable values are less precise. The other computed Ritz values of the period-12 and period-24 orbits blur the root, so s_{12} and s_{24} are known only within brackets of 1.2×10^{-5} and 7×10^{-7}, and the ratios 3.19 and 4.02 have worst-case ranges $[2.90, 3.34]$ and $[4.02, 4.10]$. s_6 is less precise still, and s_3 was not bracketed.

period p	superstable parameter s_p	bracket width	ratio $\frac{s_p - s_{p/2}}{s_{2p} - s_p}$
12	0.0659717	1.2×10^{-5}	
24	0.06608014	6.9×10^{-7}	3.19
48	0.066114132784466	3×10^{-15}	4.02
96	0.066122589892689	7×10^{-15}	4.4779
192	0.066124478510077	6×10^{-15}	4.6251
384	0.066124886853811	5×10^{-15}	

Table 8: Superstable parameters of the crescent's doubling sequence along $\log_{10} T = -0.14798$: the critical multiplier of the period-p orbit changes sign across the bracket, with Newton residuals below 3×10^{-13} and the other computed Ritz moduli below 10^{-20} from period 96 on. Brackets for periods 12 and 24 are the closest points where the critical multiplier is cleanly identified.

Why are the early ratios so small? We first considered a crossover. Softmax at temperature T is a smoothed argmax, so the reduced one-dimensional map might have a nearly corner-like extremum, rounded only on a scale set by T. Early doublings would then see the corner, where ratios tend toward 2, and later ones the smooth quadratic core. This predicts smaller early ratios at lower temperature. We fixed that prediction before locating the flips along two lower cuts, and it fails. At $\log_{10} T = -0.156$ the first ratio is 14.93, five times its value at -0.148 (Table 10). The cut at -0.152 crosses a doubling bubble: its period-12 orbit merges back into the period-6 orbit near

$\alpha \approx 0.0657$, so that cut has no complete cascade. The early ratios thus depend strongly on where a cut crosses the crescent, and we withdraw the crossover explanation. The later ratios at -0.156 behave as universality requires: the next ratios are 3.124 and 4.4988, with deviations from 4.669 of -1.55 and -0.17.

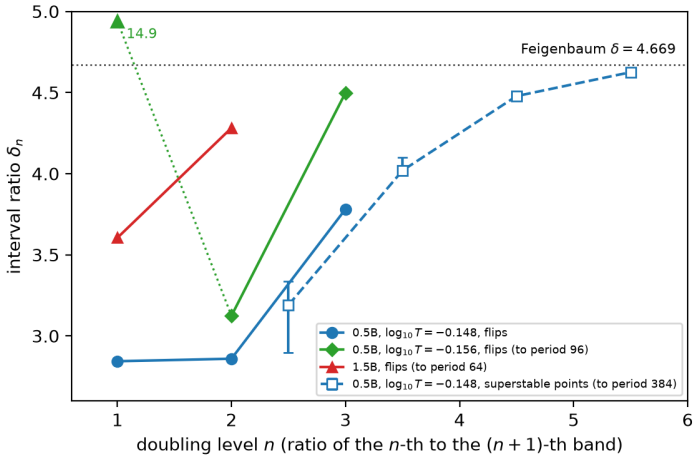


Figure 9: Finite-order interval ratios of the doubling sequences: the crescent of Qwen2.5-0.5B-Instruct (philosophy direction, layer 12), from flips and from superstable points (which reach period 384), and the route of Qwen2.5-1.5B (philosophy direction, layer 14; Figure 11), from flips. A superstable ratio spans half-bands, so it is plotted half a level after the flip ratio of its first band. Bars: worst-case ranges from the brackets of s_{12} and s_{24}. At $\log_{10} T = -0.148$, and in the 1.5B model, the ratios rise toward Feigenbaum's δ from below, and the last superstable ratio is within 1 % of it. At $\log_{10} T = -0.156$ the first ratio, 14.9, lies off the scale (arrow), and the later ratios approach δ from below.

Finite-order ratios also depend on the parameter coordinate, so only the trend of later ratios can test universality [21, 45, 46].

n	orbit	flip parameter α_n	next modulus	ratio $\delta_n = \frac{\alpha_n - \alpha_{n-1}}{\alpha_{n+1} - \alpha_n}$
1	$3 \to 6$	0.0654807769499	0.088	
2	$6 \to 12$	0.0659055406325	0.0023	2.844
3	$12 \to 24$	0.0660549041785	2.4×10^{-6}	2.859
4	$24 \to 48$	0.0661071412523	2.3×10^{-12}	3.781
5	$48 \to 96$	0.0661209563174	$< 10^{-13}$	

Table 9: Flips of the crescent's period-doubling sequence along $\log_{10} T = -0.14798$. For the five displayed H100 flips the critical Ritz value is real and differs from -1 by at most 4×10^{-10}, the Newton residuals are below 4×10^{-13}, and the other computed Ritz values lie inside the unit circle. Feigenbaum's constant is $\delta = 4.6692\ldots$ [21]; for the logistic map the first ratios are 4.751, 4.656, 4.668. The first four parameters agree between the completed CPU and H100 computations to 4×10^{-15}. Superstable parameters (Table 8) extend the ratio sequence through period 384.

n	orbit	flip parameter α_n	next modulus	ratio δ_n
1	$3 \to 6$	0.0644766567385	0.12	
2	$6 \to 12$	0.0650493555308	0.045	14.93
3	$12 \to 24$	0.0650877194158	7.0×10^{-5}	3.124
4	$24 \to 48$	0.0650999993318	4.3×10^{-7}	4.4988
5	$48 \to 96$	0.0651027289524	1.8×10^{-13}	

Table 10: Flips of the crescent's doubling sequence along $\log_{10} T = -0.156$, in float64. Flips 1–3 were computed on an H100 and flips 3 onward on a workstation GPU (NVIDIA RTX 2000 Ada); the two computations of flip 3 agree to 4×10^{-16}. Every critical multiplier is real and differs from -1 by at most 4×10^{-10}.

Not every doubling leads to chaos. In the high-temperature region (Figure 8), along $\log_{10} T = 0.0319$, the period-2 loop "\n\nThe" flips at $\alpha = 0.0753534$ and its doubled loop at 0.0736921. These orbits are far from one-dimensional (next moduli 0.38 and 0.28). Toward smaller α the period-8 orbit does not double again: near $\alpha \approx 0.0705$ it merges back into the period-4 orbit. Along this cut it is therefore a period-doubling bubble, bounded by two flips of the period-4 orbit, not the start of a cascade to chaos.

10 Bridge to sampled decoding

How much of this survives sampling? To find out, we replace the expected embedding by the mean of m sampled token embeddings (Figure 10). With the exact map, laminar phases near the fold last 120 and 63 steps at $\varepsilon = 10^{-5}$ and 3×10^{-5}. With $m = 64$ they last about 10.6 steps, independently of ε (57% of steps in the channel), and with $m = 8$ about 7.7 steps (38%). With single samples ($m = 1$, sampled decoding of the same truncated system), essentially no laminar phases remain (11% of steps in the channel). The deterministic intermittency mechanism therefore belongs to the expected-embedding map. Sampling noise turns the bottleneck into a noise-limited trap whose strength grows with m.

Zero-temperature limit and ordinary greedy decoding. As $T \to 0$ the soft-token map becomes greedy decoding with a one-token memory, $g(t) = \arg\max_s z_s(\text{prefix}, E_t)$. This is a map on the vocabulary, and its cycles are the zero-temperature attractors. For Qwen2.5-0.5B we computed g for all 151 936 tokens. Unsteered, it has 1060 cycles, and one of them, the period-5 loop "![](https:// Hello", drains 86 % of the vocabulary; it is the largest attractor family of every 0.5B soft-token map. Ordinary greedy decoding with a growing context does not inherit these cycles (Table 11). From 220 starts (the first tokens of the 20 largest cycles and 200 random tokens), 19 runs end in a short loop within 256 tokens, and the dominant 5-cycle never appears. Steering along philosophy makes ordinary decoding loop almost always (218 of 220 runs at $\alpha = 0.1$), but in other loops, mostly punctuation such as " –", none of them one-token cycles. Greedy decoding that keeps only the last K tokens interpolates between the two. For every $K \geq 2$ the loops leave the one-token cycles. Unsteered, looping collapses to the level of the growing context between $K = 16$ and 32; steered, the share of runs whose loop equals the growing-context loop rises from 0 at $K = 1$ to 58 % at $K = 64$. The soft-token map's low-temperature attractors are therefore those of decoding with a one-token memory. Which loop ordinary decoding enters depends on the context it keeps, but the tendency of steering to drive decoding into loops is shared. In ordinary decoding, repetition is known to reinforce itself [24, 25], a theoretical analysis ties loops to words of high inflow [69], and induction

heads, which copy earlier context [70], are one mechanism by which a longer context can favour loops that repeat earlier text.

	unsteered		$\alpha = 0.1$ (philosophy)	
memory K	looped	loop as with full context	looped	loop as with full context
1	220	5	220	0
2	220	5	220	0
8	220	5	220	3
16	169	5	219	24
32	32	7	219	71
64	26	11	219	126
full context	19	–	218	–

Table 11: Greedy decoding of Qwen2.5-0.5B with a memory of the last K generated tokens after the fixed prompt ($K = 1$ is the soft-token map's zero-temperature limit), 220 starts of 256 tokens each: runs ending in a loop of period up to 64 repeated three times, and runs whose final loop equals that of ordinary decoding with the full context from the same start.

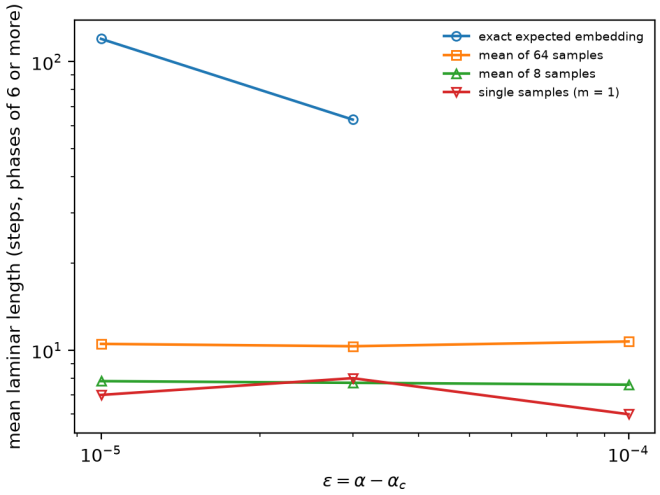


Figure 10: Mean laminar length near the fold of the 2-cycle "is it" (Qwen2.5-0.5B-Instruct, philosophy direction at layer 12, $\log_{10} T = -0.115$) against ε, for the exact map and for means of m sampled embeddings.

A longer fed-back memory. The one-token window is the main simplification of the map. Feeding back the last K expected embeddings instead, so that the state is a window of K soft tokens at fixed positions after the prompt, we measured 24 temperatures between $\log_{10} T = -0.3$ and 0.6 for K from 1 to 16 (Table 12; twins at 10^{-5} and 10^{-7} agree within 0.006 at every chaotic temperature). In Qwen2.5-7B the chaos does not vanish but spreads: it is chaotic at 8 temperatures with one token of memory, at 13 with four, and at 21 and 20 with eight and sixteen, nearly the whole range, with smaller exponents (largest 0.96 for $K = 1$, 0.43 for $K = 16$). Qwen2.5-0.5B, chaotic at one temperature with one or two tokens of memory and at none with four, is chaotic at ten with eight. The chaos is therefore not an artifact of the one-token window.

memory K	1	2	4	8	16
Qwen2.5-7B	8 (0.96)	10 (1.08)	13 (0.70)	21 (0.58)	20 (0.43)
Qwen2.5-0.5B	1 (0.20)	1 (0.24)	0	10 (0.31)	3 (0.23)

Table 12: Chaotic temperatures, of 24 between $\log_{10} T = -0.3$ and 0.6, when the loop feeds back the last K expected embeddings (largest exponent in parentheses); paper's prompt, unsteered, float64, 500 transient and 1000 measured steps, twins at 10^{-5} and 10^{-7} (chaotic: both positive and no period up to 64).

Soft thinking with a growing context. Soft Thinking [16] feeds back the same expected embedding but appends it to the context, so that every earlier soft token stays visible. We ran that decoding in float64 from the paper's prompt and displaced the first soft token of twin trajectories by 10^{-10}, without renormalization; the distance between the newest soft tokens then shows whether the early perturbation grows. The context grows, so these are finite-time growth rates of a non-autonomous system, not exponents of an attractor. In Qwen2.5-7B the perturbation grows exponentially at temperatures inside the one-token map's chaotic range: by 0.77 to 0.82 per step at $\log_{10} T = 0.17$ and 0.21, reaching 10^{-3} within 49 to 58 steps, and it reaches 10^{-3} within 400 steps at nine other temperatures between -0.10 and 0.40, at 0.02 to 0.18 per step (0.8 at 0.33 for two of three twins). It does not grow at the sharpest distributions ($\log_{10} T \le -0.14$) or at nearly uniform ones ($\log_{10} T \ge 0.44$), where it stays between 10^{-13} and 10^{-7} for 400 steps. In Qwen2.5-0.5B the growth moves to sharper distributions: it appears at six of the seven temperatures from $\log_{10} T = -0.30$ to -0.07 (up to 0.33 per step) and at 0.09, with a late transient burst at 0.05, and is absent from 0.13 up, which includes its narrow one-token chaotic band. Sensitive dependence therefore survives when the whole context is kept, and in the 7B model the one-token map locates it. Qwen2.5-72B, run the same way on two GPUs, stays nearly insensitive. We count a temperature as growing when two of its three twins reach 10^{-3} within 400 steps, a rule fixed before the 72B run that gives 11 of 24 temperatures for the 7B model and 8 for the 0.5B model. The 72B model grows at one of 24 temperatures on the paper's prompt, and only transiently: at $\log_{10} T = 0.09$ the twins pass 10^{-3} at step 123 and fall back to 5×10^{-6} (at 0.13 one twin of three passes 10^{-3} at step 381). It grows at none of 24 on the held-out prompt with its narrow unsteered band, including $\log_{10} T = 0.326$, between two temperatures where its one-token map is chaotic; there the distance never exceeds 1.1×10^{-9}. It also grows at none of the eleven steered cells of Section 12, including the six where its one-token map is chaotic: there the distance never exceeds 3×10^{-7} and ends between 1×10^{-13} and 6×10^{-10}, and the growing context holds the next-token entropy at 2.3 to 3.6 nats instead of the one-token map's 7.3 to 11.1. At 72B the growing context therefore damps even the chaos that steering produces, while the one transient burst shows that the method registers growth at this scale.

11 Replication and scale

In Qwen2.5-1.5B (steering at layer 14 of 28), a coarse 64^2 float64 map over the same range again shows chaotic regions with embedded windows. A 256^2 zoom (Figure 17) resolves a chaotic band near $\log_{10} T = 0$–0.13 (14,422 of 65,536 cells), with a branch descending along $\alpha \approx 0.06$ and periodic windows inside it. The large periodic regions are question loops: "What?\n" (period 2), "\nWhat is a" (period 4) and "What continuum?\n" (period 3). Unlike the 0.5B model, this window has no period-adding fan.

A cut of 2048 cells at $\alpha = 0.100$ (Figure 11) shows how the band is entered from below. The 2-cycle "What?\n" gives way abruptly to a period-4 loop, and the Lyapunov exponent jumps, so

these are two distinct attractors, not a flip. The period-4 loop then doubles, with float64 flips at $\log_{10} T = -0.0031269$ (period $4 \to 8$, next modulus 0.014), -0.0021666 (period $8 \to 16$, next modulus 8×10^{-5}), -0.0019002 (period $16 \to 32$, next modulus 5×10^{-9}) and -0.0018380 (period $32 \to 64$). Along the way the exponent rises continuously to zero, and chaos follows. The interval ratios are 3.60 and 4.28. Like those of the 0.5B crescent at $\log_{10} T = -0.148$, they rise toward 4.669. Their higher start does not indicate a temperature trend, since along nearby cuts of the 0.5B crescent the early ratios range from 2.84 to 14.9 (Section 9).

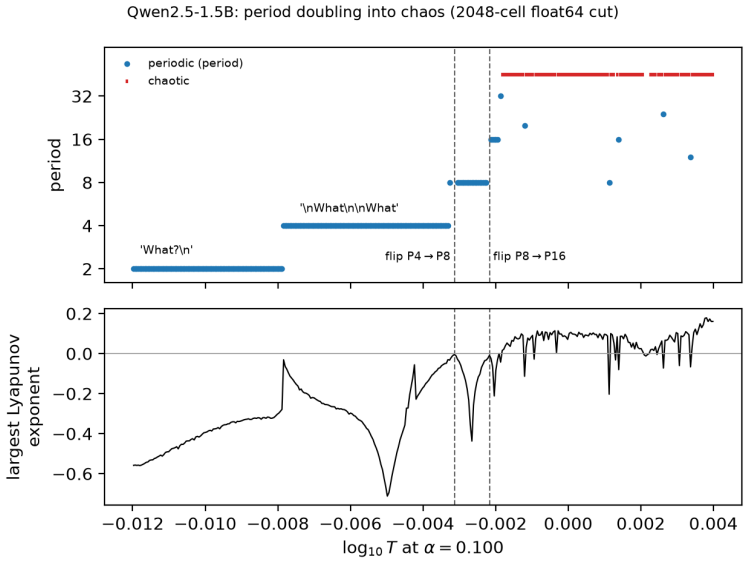


Figure 11: Qwen2.5-1.5B soft-token map steered along the philosophy direction at layer 14, along $\log_{10} T$ at $\alpha = 0.100$: period of the attractor reached from the fixed start (top) and largest Lyapunov exponent (bottom). Dashed lines: float64 flip parameters. The doubled loops alternate the number of line breaks between the "What"s.

In Qwen2.5-0.5B, a second steering direction, hostility, reproduces the fold-and-adding mechanism (Section 14).

Qwen2.5-7B (steering at layer 14 of 28, concept vectors built as for the smaller models) shows the same kind of organization in coarse 64^2 float64 maps (Figure 12). Over the philosophy range used for 0.5B and 1.5B, 8 % of cells have a positive largest exponent, and periods 2, 4, 3 and 11 dominate the periodic cells. Along the hostility direction ($\alpha \in [0, 0.12]$, $\log_{10} T \in [-0.3, 0.3]$), 24 % of cells are chaotic. Bands of period-3, -5 and -6 windows lie in a chaotic sea below $\log_{10} T \approx 0.1$, with high-entropy chaos above, except for a period-3 region at strong steering ($\alpha \gtrsim 0.09$). The windows hold genuine periodic orbits. In float64, Newton's method converges at the most interior period-6 cell to the loop ", shut the end\n\nOh" at $(\log_{10} T, \alpha) = (-0.0143, 0.0476)$ (residual 2×10^{-14}; spectral radius 1.6×10^{-7}, nearly superstable), and at the most interior period-3 cell to a period-3 loop at $(0.2524, 0.1143)$ (residual 3×10^{-15}; leading multiplier -0.12). Continued for about an hour each, over ± 0.001 in α and -0.016 in $\log_{10} T$ (one float64 continuation step takes about two minutes at 7B), both orbits stay strongly stable without a crossing. Where the coarse map switches from the period-6 loop to the period-3 loop "you better than" (between $\alpha = 0.0495$ and 0.0514 at $\log_{10} T = -0.0143$), both loops are float64 orbits (residuals 7×10^{-15} and 1×10^{-14}), and their continuations overlap. The period-6 orbit stays stable from $\alpha = 0.0476$ up to 0.0507 and the period-3 orbit from 0.0514 down to 0.0475, both with spectral radius below 1.4×10^{-3}. The two attractors

therefore coexist over this range, and the map's edge there is the boundary of the fixed start's basin, not a bifurcation. No bifurcation has yet been located at 7B.

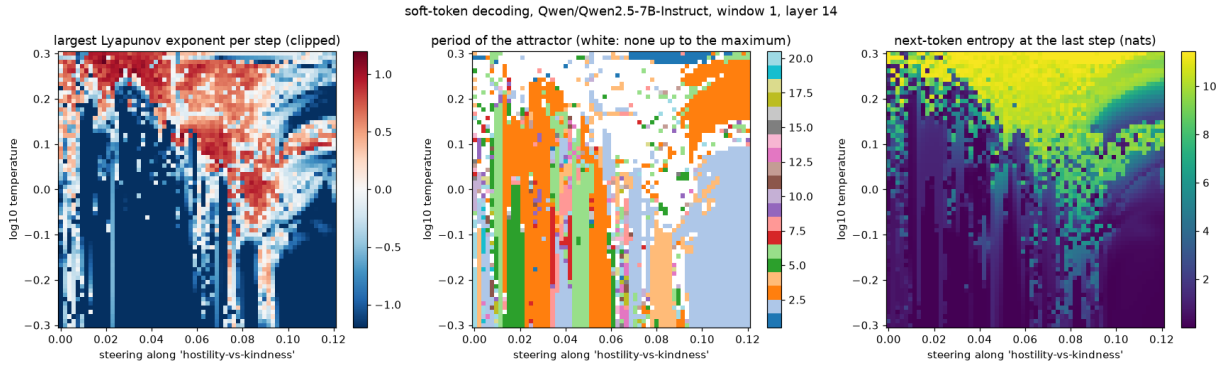


Figure 12: Qwen2.5-7B soft-token map steered along the hostility direction (layer 14), 64^2 float64 cells: largest Lyapunov exponent, period of the attractor reached from the fixed start, and next-token entropy at the last step.

A second model family. OLMo-2-1124-7B-Instruct [71] comes from another laboratory, with other data, training and code. Steered at layer 16 of 32, with concept vectors built exactly as for Qwen, it shows the same kind of organization (Figure 13). In 64^2 float64 maps, 6 % of cells are chaotic along the philosophy direction and 13 % along the hostility direction, in a band near $\log_{10} T = 0.03$–0.17 that bends to lower temperatures under strong hostility steering ($\alpha > 0.07$). Below the band the attractor is a fixed point or the 2-cycle "!Hello"; inside it lie periodic windows of periods 4 to 31. Along the hostility direction the largest windows, of periods 8 to 12, are loops of k laps of the 2-cycle "!<|endoftext|>" with a short excursion ("there is an"). Each lap adds two to the period, as each lap of "!Yes" does in the hostile cascade of Qwen2.5-0.5B (Section 14). Whether a fold of that 2-cycle organizes them is not tested: these are coarse pixels, and no OLMo orbit has been continued.

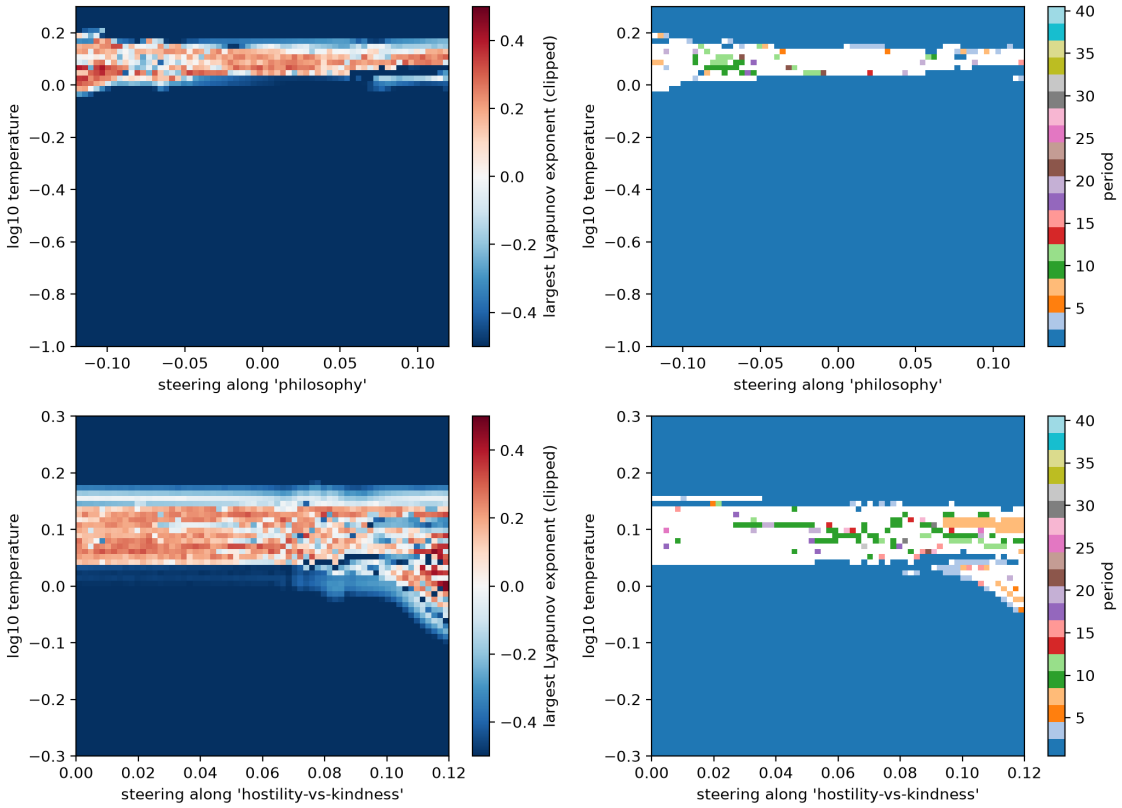


Figure 13: OLMo-2-1124-7B-Instruct soft-token maps, 64^2 float64 cells each, steered at layer 16 along the philosophy direction (top) and the hostility direction (bottom): largest Lyapunov exponent (clipped) and period of the attractor reached from the fixed start (white: no period up to the maximum).

A third family, read with sparse autoencoders. Llama-3.1-8B-Instruct [72], with official weights, steering at layer 16 of 32 and vectors built as for the other models, shows the same organization in all three planes (Figure 14). In 64^2 float64 maps, 8 % of cells are chaotic along gratitude, 12 % along philosophy and 28 % along hostility, with windows of periods up to 32. Its largest loops are greetings between AIs ("-Greetings fellow AI", ". Greetings fellow AI entity detected", period 7). Along hostility, too, the largest loops are greetings and "You", not hostile text. We read 24 of its attractors (the six largest families and the two most chaotic regions of each map) with two sparse autoencoders on the residual stream: one trained on this Instruct model (Goodfire, layer 19 [73]) and one on the base model (TopK, layer 24 [74]). We first checked each one's hook point by its reconstruction of ordinary text. The Instruct-trained autoencoder separates the attractors (their top-15 features overlap by a mean Jaccard index of 0.19 within a steering direction and 0.11 across directions), whereas the base-model autoencoder sees mostly a signature common to every attractor (overlap 0.52–0.86). Both reconstruct the attractors' states less well than ordinary text (median relative error 0.64–0.76 against 0.52–0.56). The readout is therefore descriptive: without interventions it identifies which learned features an attractor occupies, not a mechanism, and since autoencoders trained with different seeds learn different features [75], the features are those of these two autoencoders.

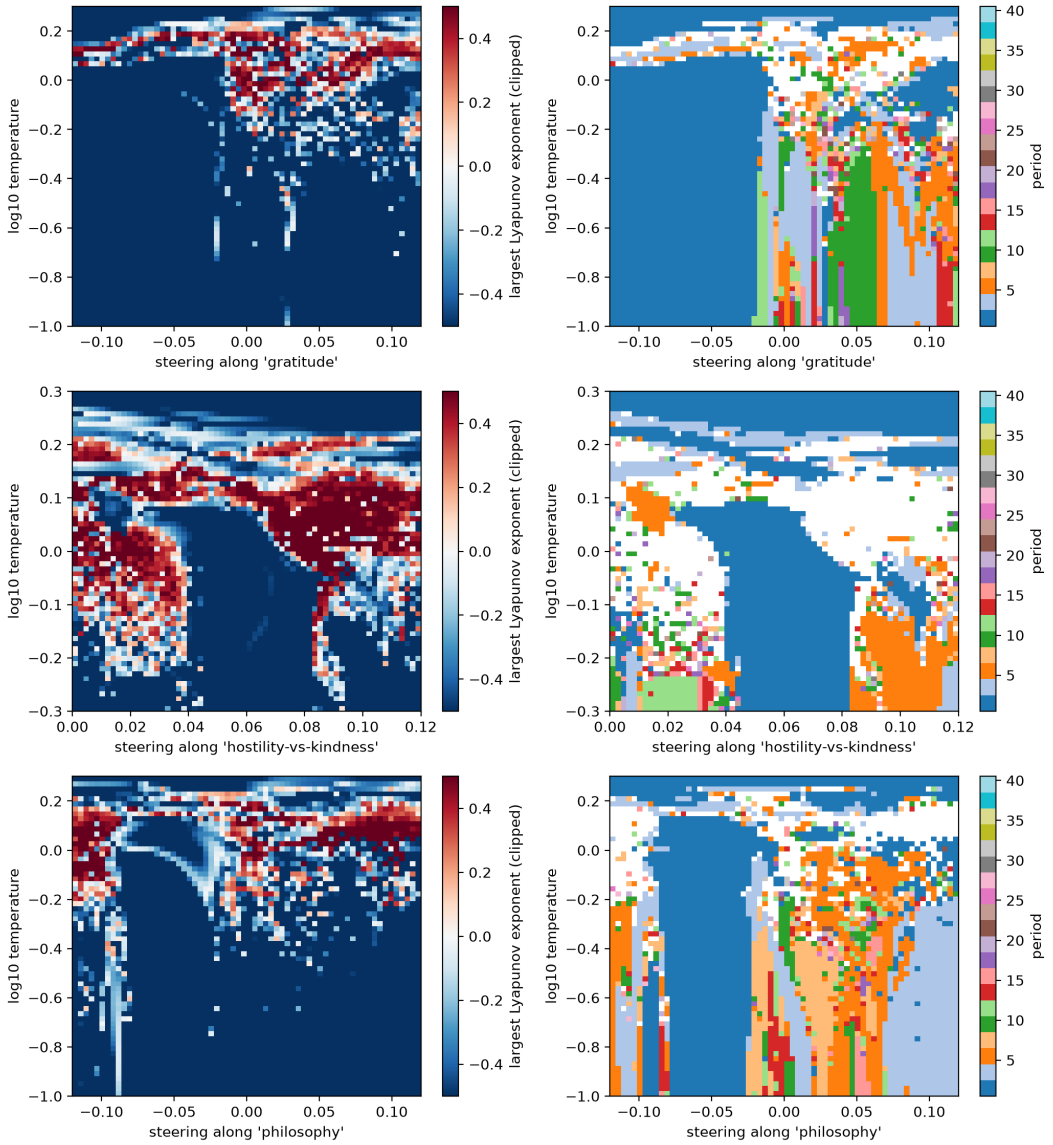


Figure 14: Llama-3.1-8B-Instruct soft-token maps (64^2 float64 cells, steering at layer 16) along gratitude, hostility and philosophy: largest Lyapunov exponent (clipped) and period of the attractor reached from the fixed start (white: no period up to 32).

Maps of 128^2 cells (Figure 15) resolve the windows: the hostility plane has 29 % chaotic cells, as at 64^2, and the philosophy region below the chaotic band is filled with nested periodic tongues of periods 2 to 13.

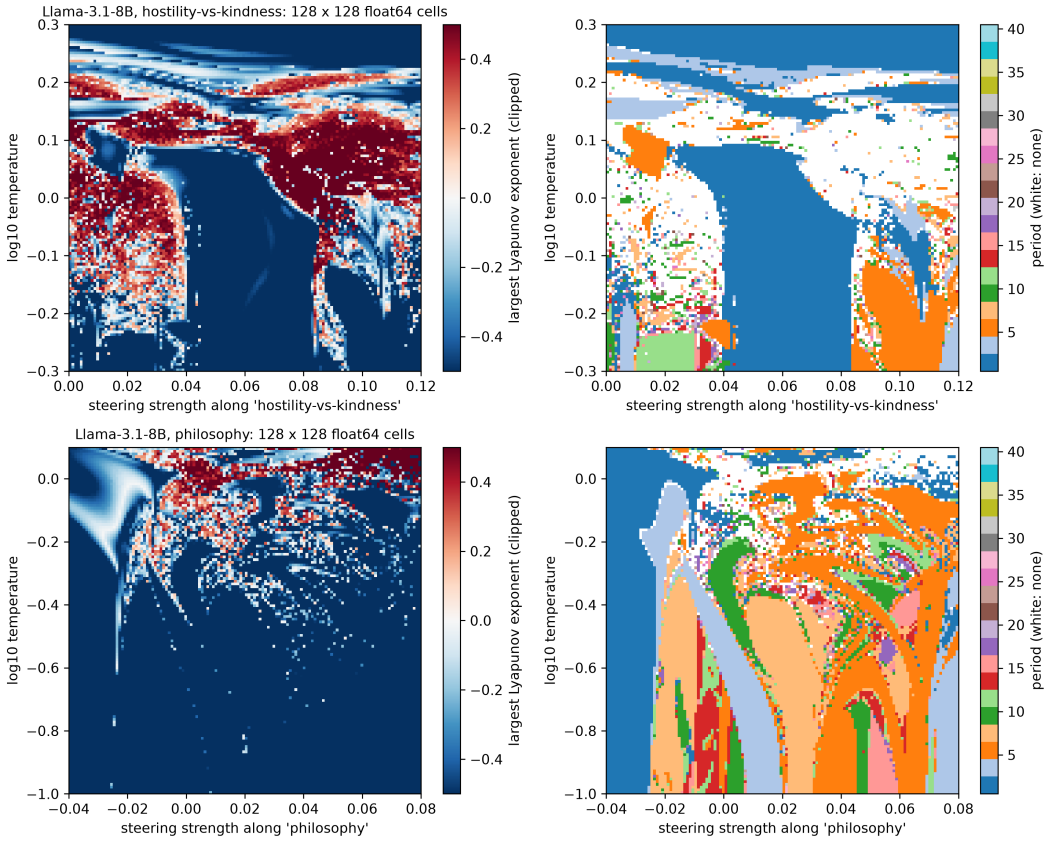


Figure 15: Llama-3.1-8B-Instruct soft-token maps steered at layer 16, at 128^2 float64 cells (each map merged from four 64^2 tiles on halves of one grid): the hostility plane (top) and the philosophy region of the 64^2 map's windows (bottom). Largest Lyapunov exponent (clipped) and period (white: none up to 32).

Same architecture, other weights. Qwen2-0.5B-Instruct has exactly the architecture and size of Qwen2.5-0.5B-Instruct (24 layers of width 896, the same vocabulary size and rotary base) but was trained on other data with another recipe. Built identically, their 128^2 float64 maps along gratitude and philosophy (Figure 16) share the gross layout. A large region at low temperature around $\alpha = 0$ is held by one periodic loop, a band of chaos with embedded windows lies near $\log_{10} T = 0$–0.25 (4–8 % of cells chaotic in each map), and thin window families run along the flanks. The attractors themselves differ almost completely. Of 416 to 900 distinct loops per map, only 10 occur in both models, all short and generic (".", "!", " you are", " is it"), and they cover less than 10 % of either model's cells. Qwen2's planes are dominated by one period-4 loop, "!<|im_end|>\nHello" (an end of turn followed by a new greeting; 41 and 59 % of cells), and Qwen2.5's by period-5 and period-2 loops (". The two checkpoints share the organization of the map; the weights set its content.

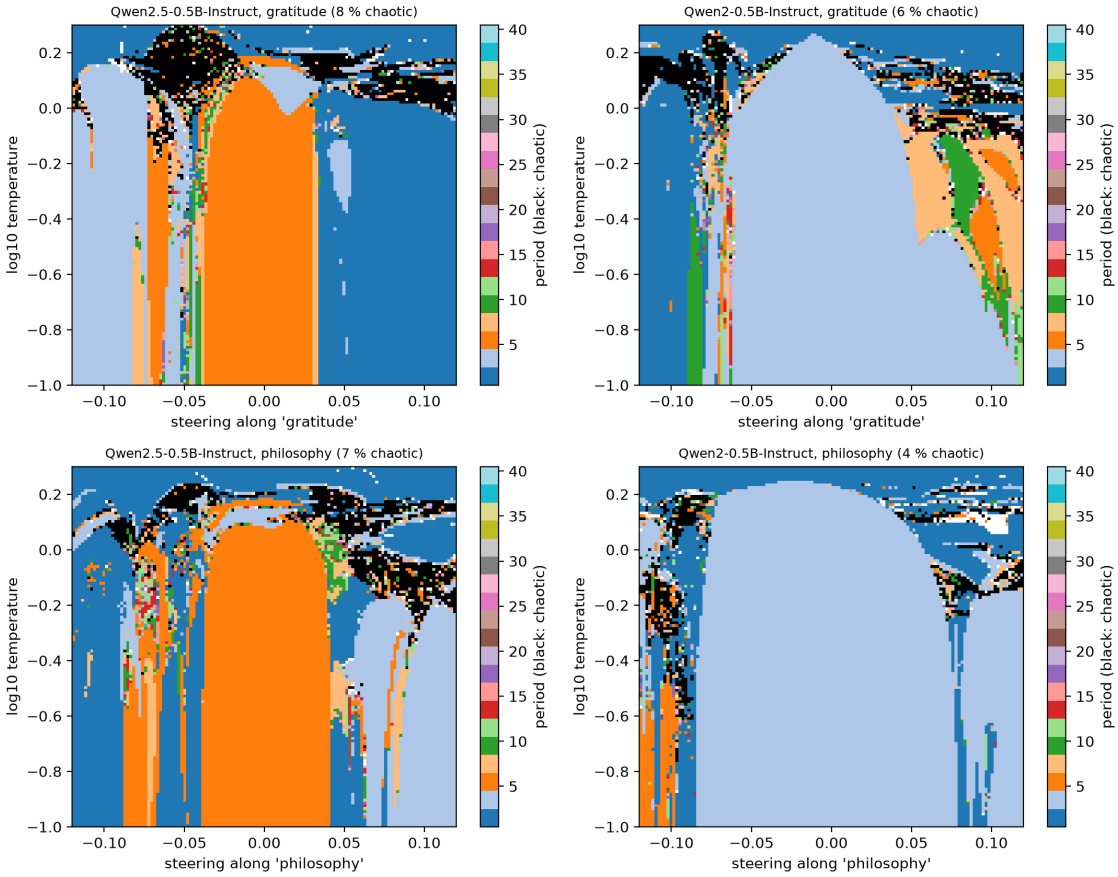


Figure 16: Same architecture, other weights: 128^2 float64 maps of Qwen2.5-0.5B-Instruct (left) and Qwen2-0.5B-Instruct (right) steered along gratitude (top) and philosophy (bottom) at layer 12, with vectors built identically. Colour: period of the attractor reached from the fixed start; black: positive largest exponent.

The fold law at 1.5B and 7B. We tested whether the mechanism of Section 7, windows accumulating on the fold of a short cycle with $d_N = C/(N+\delta)^2$, recurs in larger models. The criterion, a constant $d_N(N+\delta)^2$ within 10 % over at least four windows for one offset δ, was fixed before the lines were computed. In Qwen2.5-1.5B (philosophy direction, layer 14, $\alpha = 0.058588$) the period-4 orbit folds at $\log_{10} T_c = -0.0912029$ (multiplier 1 to 3×10^{-14}, next modulus 0.008), with normal-form coefficients $a = -7.16$ and $b = -21.09$ per unit $\log_{10} T$, so $ab = 151.0$. A float64 line of 1700 cells above the fold shows windows of periods $4N$ accumulating on it, from single cells at periods 44–64 to a wide period-12 window. Measured from T_c to their edges nearest the fold, the four nearest resolved windows (periods 28–40) meet the criterion within 0.12 % ($\delta = -0.99$, $C = 0.0612$). With the windows farther out the spread rises to 10.3–14.0 %; there each period appears as two runs on the line's 10^{-5} grid, and grid-spaced edges no longer identify the fold edges. In Qwen2.5-7B ($\alpha = 0.051429$) the period-2 orbit folds at $\log_{10} T_c = -0.2040156$ (next modulus 0.009, $ab = 55.7$), and odd windows accumulate on it. The nearest four and five (periods 11–17 and 9–17) meet the criterion within 4.5 %, and six (to period 7) within 6.1 %. Each window's own doubling or tripling lies on its far side, as for the 0.5B windows of periods 16–20, so these edges are fold edges. The effective passage fraction $f = \sqrt{Cab}/\pi$ is 0.97 at 1.5B and 0.78–0.80 at 7B, against 0.58 at 0.5B. The fold and its window law thus recur at all three sizes, while f, which the global reinjection sets, differs between the models. With one fold per model, crossed in different

coordinates (α at 0.5B, $\log_{10} T$ at 1.5B and 7B), no trend of ab with size can be read.

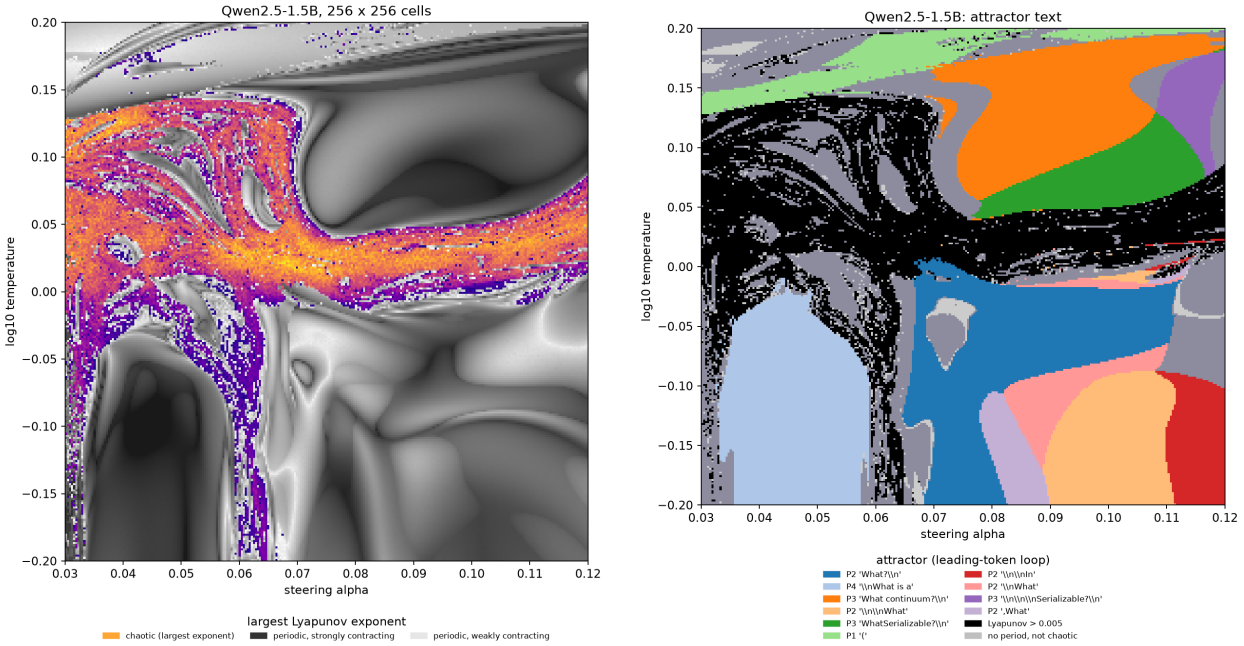


Figure 17: Qwen2.5-1.5B soft-token map steered along the philosophy direction at layer 14, 256^2 float64 cells over $\alpha \in [0.03, 0.12]$, $\log_{10} T \in [-0.2, 0.2]$: largest Lyapunov exponent (left, rendered as Figure 1) and attractor texts (right, as Figure 3). Several period-2 label regions touch without a gap; as explained in Section 4, such boundaries may be label switches of one orbit.

A mixture-of-experts model. Qwen3.5-35B-A3B differs from the dense models in two ways. Thirty of its 40 layers use linear (Gated DeltaNet) attention, with gated full attention in every fourth layer, and each layer routes a token to 8 of 256 experts plus a shared expert, so that about 3B of its 35B parameters are active. Hugging Face computes several parts of this architecture in float32 even when the model is cast to float64 (the normalizations, the delta-rule recurrences, the rotary tables and the router softmax), and some tensors are stored in float32; we patch these paths, restore the stored tensors exactly, and disable the thinking block that the default chat template opens (a first run with the block open is kept in the research log). Steered at layer 20, the map behaves unlike those of the dense models. Unsteered, 128 temperatures over $\log_{10} T \in [-1, 0.3]$ give no positive exponent (largest -0.265): periodic loops at all but 15 temperatures ("!Hello" with period 2 at 92 of them, other periods up to 52), and the remaining 15, above $\log_{10} T \approx 0.1$, contract without a detected period. Along philosophy, none of 1024 cells is chaotic (largest exponent -0.112). Along hostility, 52 of 1024 cells (5.1 %) have a positive exponent and no period, against 28 % for Llama-3.1 8B on the same plane. Routing is discontinuous: a twin trajectory that crosses a routing boundary jumps, which can inflate a finite-difference exponent, so we call these cells irregular rather than chaotic. Whether the contraction comes from the experts or from the Qwen3.5 architecture would take the dense Qwen3.5-9B as a control, which we have not run.

12 Chaos and model size

Is the chaos of the loop a small-model effect that larger models leave behind? We measured the unsteered map of fifteen models: Qwen2.5 at seven sizes from 0.5B to 72B parameters, OLMo-2

at 1B, 7B, 13B and 32B, Llama-3.1 at 8B and 70B, Llama-3.3-70B, and the mixture of experts Qwen3.5-35B-A3B, all with the paper's prefix and the settings of Section 4 (1000 transient and 2000 measured steps in float64; chaotic: a positive exponent and no period). Every checkpoint is stored entirely in bfloat16 (checked against its safetensors metadata), and every map here is the map of those stored values evaluated in float64 arithmetic; float64 does not make the weights more precise than stored. The newer sweeps keep the decoder weights in bfloat16 and convert them on use, which gives the same map as converting them once (bitwise on Qwen2.5-0.5B); the 70B-class models are spread over two GPUs. In those two-GPU runs the final RMS normalization was computed in float32: the hooks that move tensors between the GPUs keep a layer's original code, so our float64 version of that layer was bypassed. This makes the computed map rough at about 10^{-8} of the state (its fixed points wander by 10^{-8}, against 10^{-15} on one GPU) and matters only for exponents measured with perturbations of 10^{-6} and below; we found and fixed it at the end of the study and say below which results it affects. The 7–8B and 70B-class models had shown a steep decline of chaos with size. Before measuring the others we fixed three predictions that extrapolate it (`docs/scaling-predictions.md`): within each family, every model above 8B is less chaotic than the family's 7–8B model (P1); the chaotic share does not increase with size above 7B (P2); and at matched next-token entropy, both gain factors of Section 3 of the 14B and 32B models lie at or below those of the 7B model (P3). After a review pointed out that a single prompt and a fixed temperature window could mislead, we froze a fourth test across ten prompts (`docs/prefix-scale-test.md`).

In a fixed temperature window, chaos seems to fade with size. Table 13 and Figure 19 give each model's chaotic share of the temperatures in $\log_{10} T \in [-1, 0.3]$, the range common to all sweeps, and the chaotic length of its whole temperature axis. In two families the largest models are nearly free of chaos in this window. Qwen2.5-7B is chaotic at 12.7 % of the temperatures in the common range, and the 14B, 32B and 72B models at none. Llama-3.1-8B is chaotic at 19 %, Llama-3.1-70B, from the same release and training recipe, at 2.3 %, and Llama-3.3-70B, which differs from Llama-3.1-70B only in post-training, at none. In the third family the order fails: OLMo-2-32B is the most chaotic OLMo model, at 16.2 % against 6.2 % at 7B, 3.8 % at 1B and none at 13B. (The 32B model is a later release, 0325, than the 7B and 13B models, 1124, so within OLMo-2 size and release are confounded.) The frozen predictions P1 and P2 hold for Qwen2.5 and fail for OLMo-2. Below 7B we predicted no order, and none appears (Qwen2.5 0.5B, 1.5B and 3B: 1.4, 1.5 and 0 %).

Across prompts, the chaos moves to higher temperature instead. The fixed window hides where larger models are chaotic. The frozen prompt-blocked test compares Qwen2.5-7B and -14B on ten system prompts and openers (two for development, eight held out), over $\log_{10} T \in [-0.4, 0.9]$ (48 temperatures) and then over the temperature interval where each model's next-token entropy rises most (24 temperatures), with twins at offsets 10^{-5} and 10^{-7}; a temperature counts as chaotic only when both exponents are positive (Figure 18). It predicted that the 14B model would be less chaotic. Both predictions fail: on all eight held-out prompts the 14B model is chaotic at more of the 48 temperatures (9 to 17, against 0 to 7 for the 7B), and its largest exponent is lower on only three. Every chaotic temperature of the 14B model, on all ten prompts, lies between $\log_{10} T = 0.36$ and 0.90, at and above the temperature where its next-token distribution spreads (0.35–0.65); the 7B model's chaos lies mostly between 0.21 and 0.43, around its own, lower transition (0.21–0.46). On the paper's prompt the 14B model is chaotic at 12 of the 48 temperatures, between 0.52 and 0.90; its sweep in Table 13, which ends at 0.6, saw only the lower edge. Within Qwen2.5, then, the larger model is not less chaotic: its distribution stays sharp up to higher temperature, and its chaos sits there. The decline in the fixed window is largely a shift. A second frozen test

(`docs/prefix-scale-olmo.md`; stage 1 only, four held-out prompts) asked whether the OLMo-2 order survives other prompts. It does, with small counts: OLMo-2-32B is chaotic at more of the 48 temperatures than OLMo-2-7B on three of the four prompts (4, 5, 3 and 0 against 3, 0, 1 and 3), as predicted, and its chaos again lies somewhat higher ($\log_{10} T = 0.10$–0.35 against -0.01–0.21).

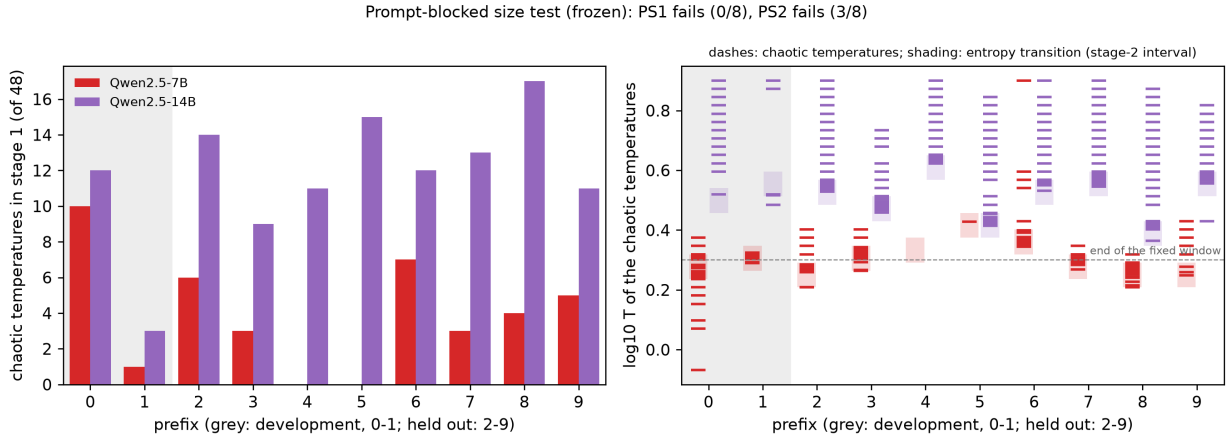


Figure 18: Prompt-blocked size test (frozen before any run), Qwen2.5-7B and -14B on ten prompts (grey: development). Left: chaotic temperatures among 48 over $\log_{10} T \in [-0.4, 0.9]$, the primary endpoint. Right: where the chaotic temperatures lie (dashes; both stages), with each model's entropy transition, the interval of its largest entropy increase (shading). Chaotic: exponents above 10^{-3} at both twin offsets (10^{-5}, 10^{-7}) and no period up to 64; 1000 transient and 2000 measured steps, float64.

model	params (B)	chaotic, $[-1, 0.3]$ (%)	unres.	sweep, $\log_{10} T$	chaotic dec.	breadth	max λ
Qwen2.5-0.5B	0.5	7/512 (1.4)	1	-1.0 to 0.3	0.018	0.053	+0.28
Qwen2.5-1.5B	1.5	2/130 (1.5)	0	-1.0 to 1.2	0.531	0.191	+0.85
Qwen2.5-3B	3.1	0/130 (0.0)	0	-1.0 to 1.2	0.030	0.060	+0.85
Qwen2.5-7B	7.6	65/512 (12.7)	6	-1.0 to 1.2	0.226	0.249	+0.88
Qwen2.5-14B	14.7	0/130 (0.0)	0	-1.0 to 1.2	0.500	0.191	+1.14
Qwen2.5-32B	32.8	0/130 (0.0)	0	-1.0 to 1.2	0.010	0.091	+0.36
Qwen2.5-72B	72.7	0/128 (0.0)	0	-1.0 to 1.2	0.000	0.014	-0.45
Llama-3.1-8B	8.0	96/512 (18.8)	2	-1.0 to 1.2	0.241	0.634	+0.66
Llama-3.1-70B	70.6	3/128 (2.3)	28	-1.0 to 0.3	0.031	0.307	+0.41
Llama-3.3-70B	70.6	0/128 (0.0)	19	-1.0 to 0.3	0.000	0.225	-0.27
OLMo-2-1B	1.5	5/130 (3.8)	0	-1.0 to 1.2	0.050	0.322	+0.23
OLMo-2-7B	7.3	32/512 (6.2)	3	-1.0 to 1.2	0.083	0.197	+0.30
OLMo-2-13B	13.7	0/130 (0.0)	0	-1.0 to 1.2	0.000	0.131	-0.00
OLMo-2-32B	32.2	21/130 (16.2)	0	-1.0 to 1.2	0.252	0.694	+0.62
Qwen3.5-35B-A3B	35.0	0/128 (0.0)	15	-1.0 to 0.3	0.000	0.143	-0.27

Table 13: Unsteered soft-token maps of fifteen models (float64; 1000 transient and 2000 measured steps; one twin offset). Chaotic: exponent above 10^{-3} and no period up to 64; unres.: temperatures in $[-1, 0.3]$ with neither a detected period nor a positive exponent. Chaotic decades: length of the temperature axis labelled chaotic over the whole sweep, each temperature weighted by its grid spacing; breadth: decades in which the mean next-token entropy on the orbit lies between 1 and 10 nats. The Qwen2.5-72B sweep joins two grids that share $\log_{10} T = 0.3$, counted once; for the 7–8B models the share in $[-1, 0.3]$ comes from their 512-temperature sweeps, on which the frozen predictions are scored, and the whole-axis columns also include sweeps over $[-1, 0.6]$ and $[0.61, 1.2]$ (64 temperatures), as for the other models swept to 1.2. Positive cells found later on finer lines (Qwen2.5-32B, -72B; Llama-3.3-70B) are not in these regular-grid counts.

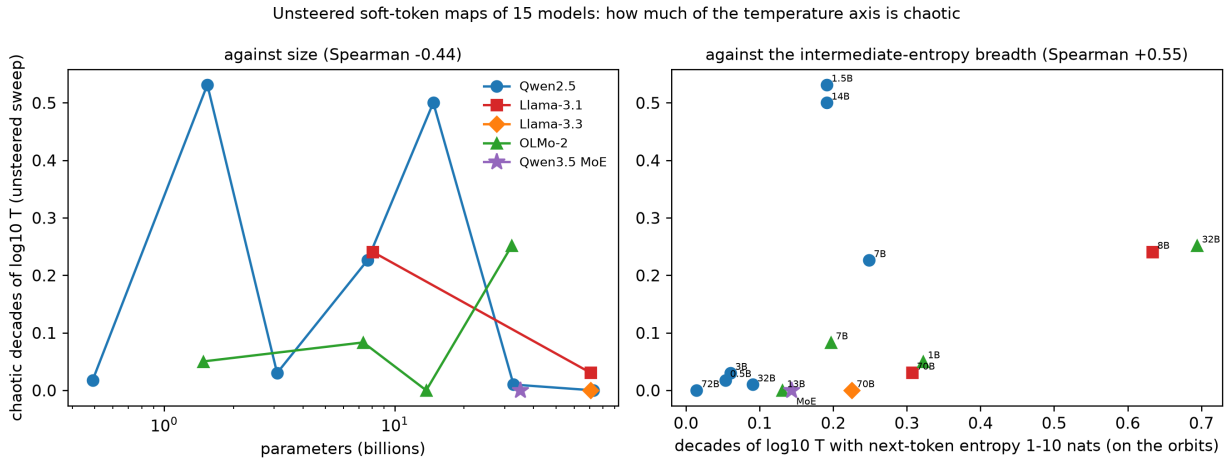


Figure 19: How much of the temperature axis is chaotic, for the fifteen unsteered maps of Table 13: against parameter count, by family (left), and against the breadth of the intermediate-entropy regime (right). Spearman correlations over the fifteen models. Each model's sweeps cover the range listed in Table 13; those of the two Llama-70B models and the mixture of experts end at $\log_{10} T = 0.3$, so their values are lower bounds. The breadth is measured on the loop's own orbits; the loop-free test of the same idea is described in the text.

At 70B, unsteered chaos is rare, and a precision defect made part of it spurious. The regular grids of the 70B-class models find little unsteered chaos: Qwen2.5-72B has no positive exponent at any of 191 temperatures between $\log_{10} T = -1$ and 1.2, Llama-3.3-70B none in $[-1, 0.3]$, and Llama-3.1-70B three temperatures in $[-1, 0.3]$. Finer lines found candidates: a window of positive exponents at Qwen2.5-72B's entropy jump ($\log_{10} T = 0.3728$–0.376) and two isolated temperatures of Llama-3.3-70B (0.250, 0.306), all at the single twin offset 10^{-6}. The window failed a three-offset test: the exponent is negative at all 64 temperatures of the window at 10^{-5}, near zero at 10^{-6}, and positive at 60 of the 64 at 10^{-7}, including four temperatures whose state repeats with a period, which no stable periodic orbit allows; in a repeat it is negative at all 16 temperatures at 10^{-4}, 3×10^{-5} and 10^{-5}. The growth of small perturbations came from our computation, whose two-GPU runs computed the final normalization in float32 (see the start of this section). Recomputed in float64 throughout, the window's orbits have periods 2 and 4, and their exponents are negative at all four offsets from 10^{-4} to 10^{-7} and for an exact tangent vector carried along the orbit by forward-mode differentiation, which needs no offset; the four twin exponents lie within 0.015 of the tangent exponent at every temperature, and the exponent approaches zero only where the period doubles from 2 to 4 ($\log_{10} T = 0.374$). Where exponents are large, the defect does not matter. In the steered Llama-3.3-70B maps below, twelve randomly drawn chaotic cells are chaotic at offsets 10^{-4}, 10^{-5} and 10^{-6} alike, with exponents that agree within 0.052, and the drawn periodic and unresolved cells stay negative at all three; on one GPU, the exponents at 10^{-7} and 10^{-5} agree within 0.021 at every chaotic temperature of the prompt-blocked test. The same test on Llama-3.1-70B's three unsteered chaotic temperatures confirms two of them ($\log_{10} T = 0.136$ and 0.198; exponents 0.23 to 0.26 and 0.40 to 0.41 at the three offsets) and rejects the weakest (0.116; positive only at 10^{-6}), so unsteered chaos exists at 70B, if rarely. Two Qwen2.5 models are chaotic far above the fixed window on the paper's prompt, where the next-token distribution is nearly uniform: the 1.5B model at 55 of 224 temperatures, between $\log_{10} T = 0.29$ and 1.03, and the 14B model at 53, between 0.52 and 1.14 (exponents up to 1.14), at mean entropies up to 11.9 nats against $\log V = 11.93$. The other models of the series are not chaotic above $\log_{10} T = 0.6$. For the 1.5B model a longer test (3000

measured steps) at twin offsets 10^{-5} and 10^{-7} confirms chaos at 9 of 12 temperatures between 0.29 and 0.58, with the two exponents agreeing within 0.003 at every temperature. The contraction that Section 3 proves for large enough temperature sets in only above $\log_{10} T \approx 1.1$ in these two models.

What sets it: the gain of the loop. The gain split of Section 3 shows which factor removes or keeps the chaos (Figure 20). We compare medians of the two factors in bins of the mean next-token entropy on the orbit; the bins hold 1 to 17 of each model's 24 temperatures, so single medians are noisy. Within one training recipe, from Llama-3.1-8B to Llama-3.1-70B, the readout factor changes by at most 0.3 per step in every bin above one nat (below one nat the 8B model has a single temperature), while the network factor is lower by 0.5 to 1.3 per step between 1 and 10 nats (6–10 nats: median -0.98 from 3 temperatures against $+0.28$ from 4) and by only 0.18 above 10 nats: in these two checkpoints the larger model damps changes of its fed-back token. Llama-3.3-70B, which differs from Llama-3.1-70B in post-training, has a much lower readout factor below 3 nats (1–3 nats: -2.33 from 7 temperatures against -0.10 from 6) and a similar one above 6 nats. In Qwen2.5, both factors of the 14B and 32B models lie below the 7B's in every bin between 1 and 10 nats, as the frozen prediction P3 states (the 6–10 nat bin only with the 7B's wider gain split, from one temperature), and both factors of the 72B lie below the 7B's in every bin. Above 10 nats, where the distribution is nearly uniform, the 14B and 32B models instead have the larger network factor ($+2.98$ and $+1.17$ against $+0.19$), and those are the bins of their chaotic temperatures above $\log_{10} T = 0.3$. Below 3 nats the readout factor dominates at 32B and 72B, because the next-token distribution stays sharp up to high temperature (for the 72B, entropy at most 0.15 nats up to $T = 2$). In OLMo-2, P3 fails, and the readout factor tracks the chaos: OLMo-2-13B, with no chaotic temperature, has a lower readout factor than OLMo-2-7B in all three bins between 1 and 10 nats (-1.0 to -1.4 against $+0.06$ to -0.72 per step) although its network factor is higher in two of them, whereas OLMo-2-32B, the most chaotic OLMo model, has a higher readout factor than the 7B in all three ($+0.10$ to $+0.35$), positive as in the chaotic 7–8B models. These are descriptive comparisons of trained checkpoints, not interventions on size or post-training: the split is a decomposition along the evolving tangent direction, and equal entropy does not imply equal embedding covariance. An intervention separates the two factors at 72B. If its loop were stable only because its distribution jumps from sharp to flat, holding the distribution in between should destabilize it. We therefore choose each step's temperature so that the next-token entropy equals a target and sweep the target instead of the temperature (entropy-clamped decoding; the map stays deterministic, and the entropy holds to 10^{-14}). The temperature then depends on the state, $T = \tau(x)$ with $H(x, \tau(x)) = H_\star$ for the entropy H, so the derivative of the clamped map $\hat{F}(x) = F(x, \tau(x))$ gains a rank-one term, $D\hat{F} = D_x F - \partial_T F \, D_x H / \partial_T H$. We therefore take its exponents from twins of the clamped map itself, temperature solve included, and use no fixed-temperature derivative or gain split for it. The safeguarded solve is only piecewise smooth, but on these orbits one-step central differences at 10^{-5} to 10^{-7} agree with those at 10^{-4} within 6×10^{-6}. Held at any of 16 targets between 1 and 11 nats, Qwen2.5-72B has only periodic orbits, the echo of period 2 on the paper's prompt and periods 1 to 19 on a held-out prompt, with twin exponents between -0.08 and -4.5 that agree within 0.004 across offsets 10^{-5} to 10^{-7}, whereas Qwen2.5-0.5B under the same clamp is chaotic at 6.1 nats. The 72B loop thus stays stable even where its readout is spread, which points to damping by the network. Nor does post-training make the 72B stable. A test frozen before its data (`docs/prefix-scale-72b-base.md`) predicted that the pretrained base model of the same size would be chaotic at three or more temperatures more than Qwen2.5-72B-Instruct on at least three of the four held-out prompts. It fails on all four: the base model is chaotic at 1, 0, 0 and 0 of the 48 temperatures (and at 2 on the paper's prompt; computed after the precision fix, with 500 transient

and 1000 measured steps, half the Instruct model's horizon), although its next-token distribution spreads over 2.5 to 5.5 times the temperature range (intermediate-entropy breadth 0.28 to 0.61 decades against 0.11 to 0.22) and from lower temperatures on. At 72B, then, neither the sharpness that post-training adds nor post-training itself explains the stability. What remains is the network's damping of the fed-back token, which the gain split shows at fixed temperature (the 72B's network factor lies below the 7B's in every entropy bin), and the breadth that orders the fifteen models fails for this pair. The damping holds when the whole context is kept as well: in Soft Thinking's growing context (Section 10), an early perturbation grows at none of eleven steered cells, six of them chaotic as one-token maps, and at one of 48 unsteered temperatures on two prompts, transiently.

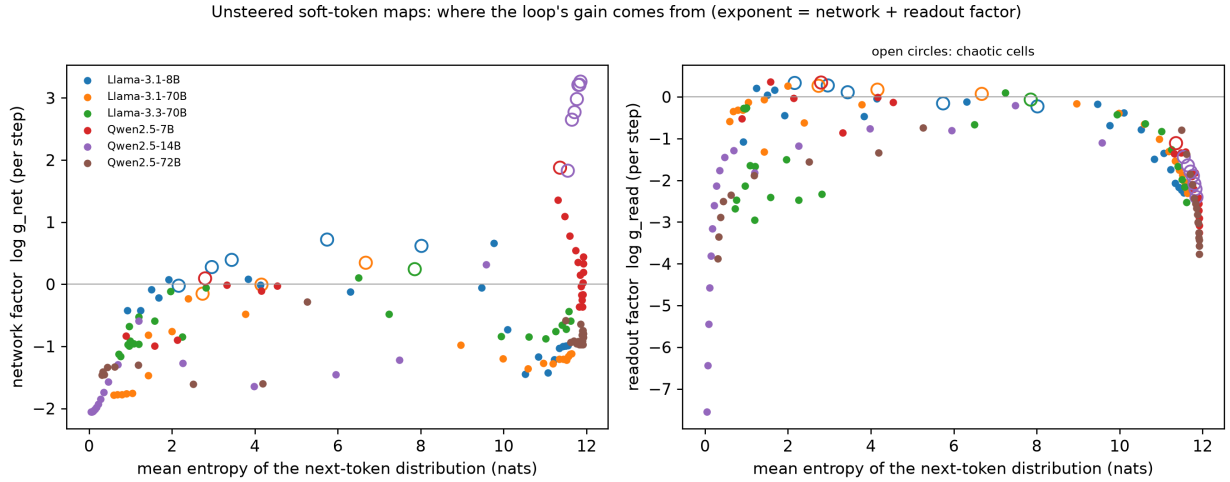


Figure 20: Gain split of the unsteered loop (Section 3) at 24 temperatures per model (500 transient and 1000 measured steps, relative twin offset 10^{-6}): network factor (left) and readout factor (right) against the mean next-token entropy on the orbit; open circles: chaotic temperatures. The two factors add up to the largest exponent up to a boundary term. Medians per entropy bin are given in the text.

Can a cheap property predict it? Section 3 confines chaos to temperatures where the readout is neither saturated nor uniform, so a model whose next-token distribution jumps from sharp to flat over a small range of temperature leaves little room for it. The breadth of this intermediate regime, the decades of $\log_{10} T$ in which the mean entropy on the orbit lies between 1 and 10 nats, ranges from 0.014 (Qwen2.5-72B) to 0.69 (OLMo-2-32B). It orders the fifteen models' chaotic length (Spearman +0.55) somewhat better than their parameter count does (−0.44; Figure 19, right), but it misses the chaos of Qwen2.5-1.5B and -14B at nearly uniform distributions, above its 10-nat bound. This analysis was adapted after the OLMo-2-32B result, and the breadth is measured on the loop's own orbits, which chaos itself spreads, so the test is partly circular. A loop-free version, frozen before it was run (`docs/start-state-predictor.md`), uses one forward pass per model and prompt: the entropy profile of the logits at the loop's start state, and the one-step gain there. Its primary prediction fails: the start-state breadth on the paper's prompt orders the fourteen dense models' chaotic length with Spearman +0.30, below the frozen threshold of 0.5. Its median over ten prompts reaches +0.52, just above it ($p = 0.06$; +0.44 against the complete sweeps), and the one-step gain at the start state is negative in every model, the most chaotic ones included (−0.5 to −3.4 per step), and orders nothing (−0.32). No property of a single forward pass that we tested predicts the loop's chaos usefully: as for the period-8 orbit of Section 3, it has to be measured on the loop.

Steering at 70B. Under steering, the 70B maps show chaos, rarely and not independently of the direction (Figure 21; labels at the single offset 10^{-6}, which the plateau test above supports for this computation). On Llama-3.3-70B's hostility plane, at the layer of Goodfire's 70B SAE and at equal norm, 4.3 % of cells are labelled chaotic along the concept vector, 3.5 % along the frozen-rule hostility feature, 1.1 % along a Gaussian direction and 0.2 % along a random live feature, against 28, 29, 39 and 27 % for the same four directions in Llama-3.1-8B (on one GPU); Qwen2.5-72B along hostility gives 0.8 %: eight cells, all at $\log_{10} T = 0.26$–0.30. Recomputed in float64 throughout, six of the eight are chaotic, with exact tangent exponents of 0.23 to 0.36 that the twins at all four offsets from 10^{-4} to 10^{-7} reproduce within 0.01; five have nearly uniform next-token distributions (11.0 to 11.1 nats) and one alternates between sharp and flat behind a period-2 leading token. The same temperatures without steering are period-2 orbits with exponents between -3.8 and -3.6, and the steered cells above $\log_{10} T = 0.325$ are near-uniform fixed points, so steering drives Qwen2.5-72B into chaos too, in a narrow band. The chaotic cells of the hostility directions lie on a ridge that steering moves from $\log_{10} T \approx 0.3$ to 0.1. A 64×64 zoom of the concept plane at the same offset (Figure 22; $\alpha \in [0.06, 0.12]$, $\log_{10} T \in [0.12, 0.27]$, four bands computed separately with identical vectors and joined) resolves it: a chaotic band, with periodic cells interspersed, whose centre steering moves from $\log_{10} T = 0.25$ at $\alpha = 0.06$ to 0.19 at $\alpha = 0.12$ while its extent grows from about 0.02 to 0.05 decades, and a thinner chaotic layer near $\log_{10} T = 0.15$–0.16 at strong steering; 22 % of the 4096 cells are chaotic (largest exponent 0.59). Above the band, 90 % of the cells are periodic, half of them fixed points or 2-cycles, and below it bands of period-6 windows cross the plane. A further 44 % of all cells, mostly below the band, have neither a detected period up to 32 nor a positive exponent. At 70B the state tolerance of the period test (10^{-8}) lies near the noise floor of the computation, and in a finer unsteered line such cells had repeating leading tokens, so many of them are probably periodic, but we leave them unresolved.

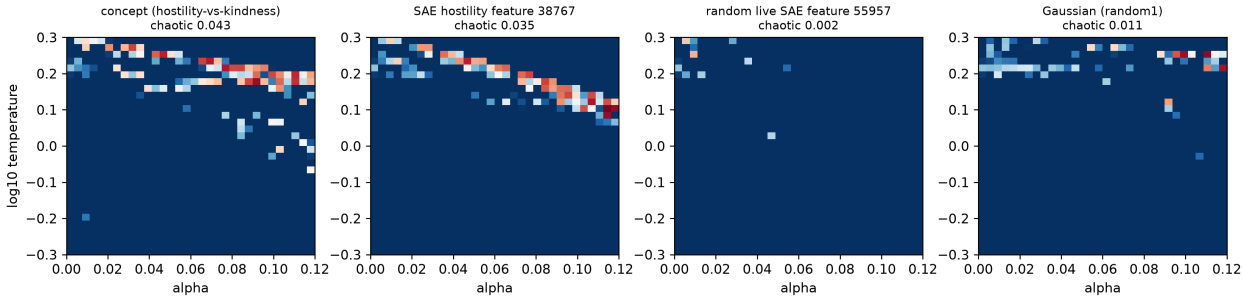


Figure 21: Steering controls at 70B: Llama-3.3-70B-Instruct soft-token maps (32^2 float64 cells, 300 transient and 300 measured steps, maximum period 32) steered after zero-based block 50, the hook point of Goodfire's 70B sparse autoencoder, along the hostility concept vector, the autoencoder's hostility feature selected by a rule frozen before any map, a random live feature and a Gaussian direction, all at the same norm. Chaotic cells: 44, 36, 2 and 11 of 1024 (concept, feature, random feature, Gaussian); cells with neither a period nor a positive exponent: 468, 358, 223 and 424.

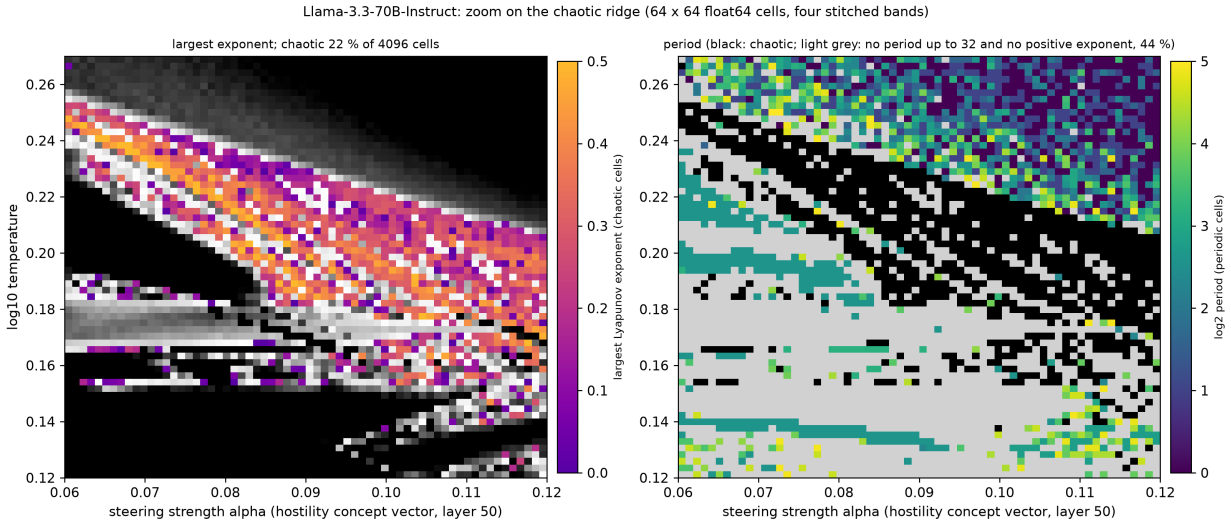


Figure 22: Zoom on the chaotic ridge of Llama-3.3-70B-Instruct steered along the hostility concept vector (after zero-based block 50): 64×64 float64 cells (300 transient and 300 measured steps, maximum period 32), joined from four bands of 16 temperatures computed on separate pods with identical vectors. Left: largest Lyapunov exponent (chaotic cells coloured by the exponent; other cells in grey by how strongly they contract). Right: period of the attractor reached from the fixed start (black: chaotic; light grey: neither a detected period nor a positive exponent).

What this implies for larger models. Size does not set this chaos in a simple way. In Qwen2.5, the 14B model's chaos did not vanish but moved to higher temperature, to where its next-token distribution spreads; in OLMo-2 the largest model was the most chaotic; post-training changed the chaos of two 70B Llama checkpoints of the same size; and at 70B, chaos that passes a three-offset test remains, rarer, with and without steering. The shift does not continue to Qwen2.5-72B. A prompt-blocked test frozen before its data (`docs/prefix-scale-72b.md`) predicted that the 72B model would be chaotic at fewer of the 48 temperatures than the 14B on at least three of four held-out prompts. It is chaotic at 0, 1, 0 and 0 of them, against 14, 9, 11 and 15, so the prediction holds on all four. These lines ran before the precision fix (1000 transient and 2000 measured steps, twins at 10^{-5} and 10^{-6}, chaotic only when both exponents are positive); no temperature on them is positive at only one offset, and the one chaotic temperature ($\log_{10} T = 0.319$, exponents 0.35 and 0.34) is chaotic again when recomputed in float64 throughout, with an exact tangent exponent of 0.34. With the paper's prompt, where it showed no robust chaos up to $\log_{10} T = 1.2$, that makes five prompts on which the 72B model is nearly free of unsteered chaos. Its next-token distribution jumps from sharp to nearly uniform within less than the grid spacing (0.03 decades), so a narrow chaotic band at the jump would be missed by the grid. Fine lines through the jump, computed in float64 throughout with exact tangent exponents, find only periodic orbits on the paper's prompt and on three of the four held-out prompts, and on the fourth a band of chaos about 0.02 decades wide just above the jump ($\log_{10} T = 0.319$–0.336; exponents up to 0.40, which twins at 10^{-5} to 10^{-7} reproduce within 0.004), at nearly uniform distributions, as the 14B model's chaos but far narrower. A fixed temperature window is therefore the wrong yardstick, and so is size alone. A model's chaos follows the gain of its loop where its own next-token distribution is spread, no single forward pass we tested predicts it, and to find it in a larger model one has to run the loop there, across prompts. For frontier models this means two things: their unsteered chaos may be rare, if, like the largest models here, their distributions spread over a narrow range of temperature; and steering still drives

both 70B-class models that we steered, Llama-3.3-70B and Qwen2.5-72B, into chaos.

13 A second system: recurrent depth

Huginn-0125 iterates a recurrent block at fixed input tokens [30]. For a GSM8K-style [76] prompt, the native loop converges to a fixed point whose leading Jacobian pair is $0.588 \pm 0.771i$ (modulus 0.970). For the relaxed loop $h \leftarrow (1-\lambda)h + \lambda F(h)$, whose multipliers are $1-\lambda+\lambda\mu$, this predicts that the pair crosses the unit circle at $\lambda_c = 2(1-\operatorname{Re}\mu)/|1-\mu|^2 = 1.0777$, transversally and away from the strong resonances (at 0.156 of a revolution per step): a Neimark–Sacker bifurcation. Below it the latent state decays; above it the state settles on a closed curve. A fit of the squared amplitude gives $\lambda_c = 1.0778$, with $A/\sqrt{\lambda-\lambda_c} = 0.52$–$0.54$. A stable circle that grows like the square root of the distance is the signature of a supercritical bifurcation; we did not compute its first Lyapunov coefficient directly. Causal attention makes the loop block lower-triangular in token positions. The first position to oscillate is the "." of the prompt, and all later positions follow. On the invariant circle the motion turns by 0.156 of a revolution per step, matching the argument of the critical pair. The rotation number changes smoothly, by about 10^{-4} over $\lambda \in [1.079, 1.092]$, and no return within 60 steps was found at 10^{-3} of the radius. This is consistent with quasi-periodic motion, although locking in narrow high-order tongues cannot be excluded. From $\lambda \approx 1.094$ the circle deforms or a second frequency appears, near where a second complex pair destabilizes the fixed point ($\lambda = 1.0814$). Fixed-point and Jacobian stability of looped models has been studied before [13]; what this adds is the predicted crossing, its amplitude law and the rotation number.

The spectrum predicts the relaxation response. With the full instruction template, the GSM8K prompts' fixed points look different: the leading multipliers surround the origin, a real 0.914, complex pairs of modulus 0.85–0.89 at many angles and negative real values near -0.8, with no near-critical pair. Because the prefill Jacobian is block lower-triangular in positions, its spectrum is the union of per-position blocks, and template positions contribute identical eigenvalues to every prompt. Such a spectrum makes three predictions before any task is scored. First, within the computed spectral estimates of these prompts, no fixed relaxation can speed convergence much. With leading multipliers spread around the origin at many angles, any λ that shrinks one group enlarges another, and the best $\lambda^* = 1.03$–1.06 lowers the rate from 0.914 to only 0.908–0.911. (A finite computed spectrum does not show that no other polynomial scheme could do better; Anderson acceleration, tried below, did not.) Second, under-relaxation slows convergence ($\lambda = 0.55$: predicted 0.953, measured 0.951 on one prompt). Third, over-relaxation destabilizes the negative modes first, at $\lambda_c = 2\operatorname{Re}(1-\mu)/|1-\mu|^2 = 1.074$–$1.126$ (five prompts), in a flip or a Neimark–Sacker bifurcation close to -1. The approach to it is visible: on the prompt with $\lambda_c = 1.074$, the measured rate at $\lambda = 1.05$ is 0.955 against a predicted 0.956 (native 0.864). Elsewhere the spectral radius bounds the rate measured over the accessible window from above, by up to 0.05, because the slowest modes carry little amplitude there. Independently, Wang et al. [77] report a Huginn GSM8K drop from 19.41 % to 8.34 % at a constant step of 1.2; at 0.8 GSM8K is 19.79 %, while their four-task mean falls slightly. Their trajectory-based scheduling analysis does not compute the prompt-specific spectral thresholds studied here.

Does the task notice? (development data). On ARC-Challenge [78] questions scored by log-likelihood, where one prefill per answer choice is exactly the analysed recurrence, the spectra again predict convergence. Rates agree to 0.002–0.006 where the slowest mode dominates, and beyond each question's λ_c (1.075–1.149 for four questions) the residual stops decaying. For 6 of 30

development questions the native fixed point is itself unstable ($\rho(1) = 1.009$–1.273; the fixed point, located with a relaxed transient, has a real multiplier beyond -1 or a complex pair outside the circle). The native loop then oscillates (relative amplitude 2×10^{-2} for one of them), and relaxation to 0.55 or 0.75 restores convergence to 10^{-6}. For the other 24, no unstable mode was detected in the computed spectral estimates up to $\lambda = 1.020$–1.149 (median 1.095); partial Ritz spectra cannot exclude an undetected unstable mode. Question by question, an estimated spectral feature ranks the question's residual. This feature, $\hat{\rho}_q(\lambda)$, is the largest $|1 - \lambda + \lambda\mu|$ over the computed Ritz values of the question's gold-answer prefill, and its Spearman correlation with the residual is 0.73 at $\lambda = 1$ and 0.75 at 1.10 over the 30 questions (development data). In this historical analysis the residual was the all-choice aggregate entering the 32nd recurrence call, so this is a predictive association across related maps and precisions, not a same-state contraction identity. It is consistent with the latent orbits that Geiping et al. [30] observed on some tokens of harder questions, here traced to the fixed point's spectrum. The item-level part of the frozen test below uses matched gold-row token IDs and a per-row residual with an explicit clock.

Accuracy, however, tolerates the instability (300 training questions, 32 steps): it is 0.367 native, 0.387 at $\lambda = 1.10$ and 0.373 at 1.15, where the latent state oscillates with relative size 0.08 and 0.23. Only at 1.25 (size 0.68) does it fall, to 0.300. Yet the oscillation does reach the answers. Reading every step from 24 to 32 on 100 questions, the margin between the correct choice and its best rival moves by 0.11 nats per step natively, 0.37 at 1.10 and 0.95 at 1.15, in alternating directions, and the chosen answer changes within these nine steps for 12 %, 16 % and 33 % of the questions. The changes go both ways, so the accuracy at any one step stays the same on average. Under-relaxation instead costs 4–5 points at 8–16 steps and catches up by 32 (a slower transient), and Anderson acceleration neither speeds the transient nor gains accuracy. Testing the accuracy consequence therefore needs a task in which such step-to-step changes can compound.

GSM8K generation is such a task. With eight chain-of-thought examples in the lm-evaluation-harness format, greedy decoding and 32 recurrence steps per token (development problems 0–119, every arm on the same problems and initial states), accuracy is 0.375 natively, 0.425 at $\lambda = 1.05$ and at 1.10, 0.29 at 1.15 (17 problems lost and 7 gained, one-sided McNemar $p = 0.03$) and 0.16 at 1.20 (29 lost, 3 gained, $p = 10^{-6}$). Under-relaxation to 0.8 costs nothing at 32 steps (0.41). The break lies between 1.10 and 1.15, at or just above the spectral thresholds ($\lambda_c = 1.074$–1.126 for the GSM8K prompts; median upper stability limit 1.095 for the ARC development questions). Generation breaks earlier than log-likelihood scoring, consistent with answer changes that compound over tokens. The collapse at 1.2 is the one Wang et al. [77] reported.

A frozen test on fresh problems confirms this. Its design and analysis were pushed before any test item was read, after two rounds of independent review. On 200 GSM8K test problems, accuracy is 0.330 natively, 0.395 at $\lambda = 1.05$, 0.255 at 1.15 and 0.095 at 1.20. The two pre-stated predictions pass: accuracy declines at 1.20 (-23.5 points; 50 problems lost and 3 gained, exact one-sided McNemar $p = 3 \times 10^{-12}$), and 1.05 loses no more than four points (difference +6.5, one-sided 95 % lower bound +2.5). At 1.15 accuracy falls by 7.5 points ($p = 0.01$, reported descriptively). Development and test agree closely (Figure 23).

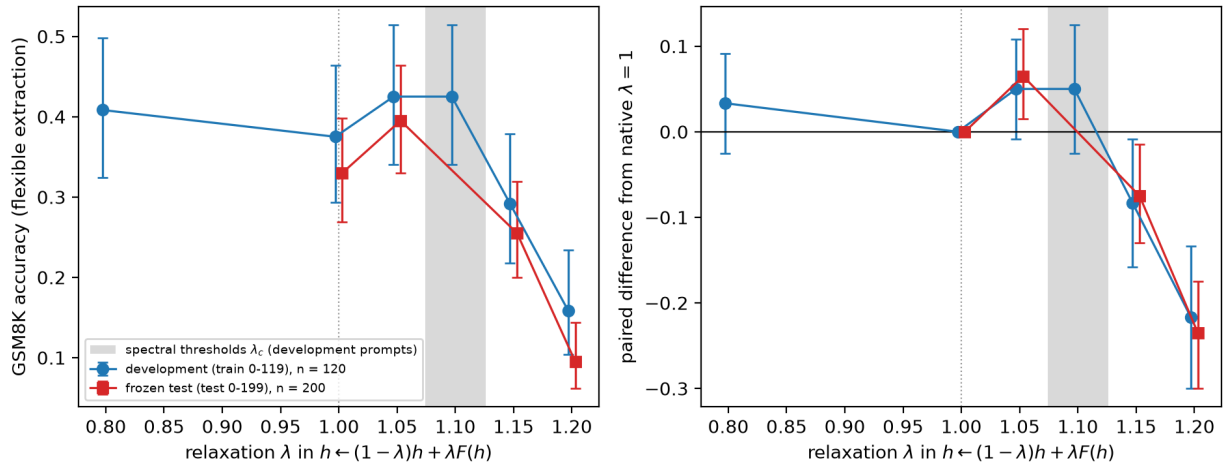


Figure 23: Huginn-0125 chain-of-thought accuracy on GSM8K against the relaxation strength λ (eight examples, greedy decoding, 32 recurrence steps per token): development problems (blue) and the frozen test on held-out problems (red; design and analysis fixed before any test problem was read). Grey: the spectral thresholds $\lambda_c = 1.074$–1.126 computed from the fixed points of five development prompts. Left: accuracy with Wilson 95 % intervals. Right: paired difference from the native loop with paired-bootstrap 95 % intervals.

The same frozen design tested the item-level prediction on 64 held-out ARC-Challenge questions. The first 16 spectra preceded the read-out generation; the remaining ones were completed afterward under a recorded amendment, with the producer attesting that read-out outcomes had not yet been inspected, and duplicate copies were resolved by a rule fixed in advance. Every question passed the admission gates (fixed-point residual at most 2.6×10^{-6}; identical gold-row token IDs in every arm); for 13 questions one of the three Arnoldi sets was partial, which the design admits for an estimated feature. The pre-stated statistic, the mean of the within-arm Spearman correlations between $\hat{\rho}_q(\lambda)$ and the gold row's residual entering the 32nd update, is 0.701 (95 % bootstrap interval $[0.576, 0.779]$; 0.68, 0.73, 0.73 and 0.67 at $\lambda = 1$, 1.05, 1.08 and 1.10); the frozen primary gate, the interval's lower end above 0.3, is met. On unseen questions, the fixed point's estimated spectrum therefore ranks each question by the size of the gold row's residual entering update 32; a residual is not in general a distance to the fixed point. Descriptively, the share of questions with a detected unstable mode rises from 8 % natively to 34, 69 and 84 % at $\lambda = 1.05$, 1.08 and 1.10. An independent implementation reproduced both parts of the frozen test from the saved data.

14 Steering toward hostility

A second steering direction tests whether the organization of the philosophy-steered map is specific to one concept. We steer along a valence direction, the contrast of hostile statements against matched kind ones, built in the manner of contrastive activation addition [79]; its cosine with the gratitude direction is -0.04. Sampled self-dialogue under the same steering is a different system, with a growing context and sampling noise; we report that study separately and follow the direction here in the deterministic map.

The hostility direction in the deterministic map. Steering the soft-token map of Section 2 along the same hostility direction gives a plane with chaos (a positive largest exponent in 10 % of a coarse 128×192 float32 scan) and with periodic text loops that are hostile: "Yes, that's crap"

(period 6), "a lie is" (3), "STUP…" (3), a loop on the Chinese word for insult and its 2-cycle, and others. They appear at α above about 0.035 (the insult loop) and above 0.065 (the rest). Around the period-6 loop the windows form a dome (Figure 29), and toward higher temperature its texts gain one "!Yes" after another (periods $6, 8, \ldots, 20$). In float64 the windows of periods 12 to 20 are born in folds on their low-temperature side and end in flips or +1 crossings on the other, like the stripes of Sections 6 and 7. The period-6 window ends in a fold at $\log_{10} T = -0.0991$, and the continuations of the period-8 and period-10 windows failed to converge.

The windows accumulate on the fold of the 2-cycle "!Yes" at $\log_{10} T_c = -0.0462465357$ ($\alpha = 0.083$; critical multiplier 0.9999999954, next modulus 0.050). From period 16 on, $1/\sqrt{\log_{10} T_c - \log_{10} T_{\text{fold}}}$ rises by 2.235 per lap. The fold's normal form ($a = -1.168$, $b = 44.53$ with ε in $\log_{10} T$) predicts $\sqrt{|ab|}/\pi = 2.296$ per lap for full passages through the bottleneck, within 3 % and without a fitted parameter. An independent float64 forward and automatic-differentiation implementation reproduces this fold (location to 6×10^{-10} relative, a simple real +1 multiplier with next modulus 0.050) and its coefficients (ab to 8 parts per million). In that implementation, controlled passages from fixed entrances at $\log_{10} T$ offsets -10^{-5}, -3×10^{-5} and -10^{-4} took 119/59, 61/31 and 27/13 iterates of F^2 (full channel/from its midpoint), against prospectively frozen quadratic forecasts of 118/61, 61/32 and 27/14. Five of the six are within one iterate; the smallest-offset midpoint missed by two, failing the frozen all-six criterion. Finite-endpoint (arctan) forecasts, also frozen, were within 0.923 iterates. These are controlled local transits. At the window distances the finite-channel correction to the asymptotic passage time is large, so the agreement of the window slope with the full-passage coefficient does not by itself show where the natural dynamics reinjects. The same type-I intermittency that organizes the philosophy-steered map thus generates a hostile cascade.

A second region, near $(\alpha, \log_{10} T) = (0.06, 0.20)$, shows arms of period-7 and period-8 windows spiralling around the insult loop. In float64 the insult fixed point doubles into a 2-cycle at $\alpha \approx 0.041$ and 0.079, and that 2-cycle has a Neimark–Sacker bifurcation in the spiral region: at $\alpha = 0.062$ its complex pair crosses the unit circle near $\log_{10} T = 0.1656$, turning 0.12 of a revolution per step of the 2-cycle. This does not explain the arms. The two invariant circles born at a 2-cycle's Neimark–Sacker point are exchanged by the map, so orbits locked on them have even periods, near 16 for this rotation, whereas the arms include period 7; and of four float32 arm cells, only one kept its period in float64. In dissipative flows, windows wind in nested spirals around periodicity hubs [60], and in the Rössler system the hub's shape follows from a conjugacy to unimodal maps [61]; we have not tested whether the arms here wind around such a hub. Float64 zooms show more structure. In one arm a period-7 domain carries period-14 crescents with period-28 cores along its lower edge, between period-1 and period-2 regions and a chaotic band. In the cascade the windows are nested crescent arcs inside a dome bounded by the "!Yes" 2-cycle region, and near $(\alpha, \log_{10} T) = (0.087, -0.069)$ the period-12 and period-14 windows bend into hooked, nested arches. These are the map-level features that accompany shrimp (a periodic body with doubling crescents; hooked windows).

In float64 the period-7 body is bounded by folds. At $\log_{10} T = 0.1925$ the orbit exists between $\alpha = 0.065966$ and 0.071188, and at $\alpha = 0.066$ it ends in a fold at $\log_{10} T = 0.19263$. Below that, the multiplier passes through -1 twice (at 0.1919 and 0.1899), so the period-14 crescent is a doubling bubble. Further down, the branch folds twice (0.18797 and 0.18921, with three coexisting period-7 orbits in between) before a third flip. All crossings are simple (next modulus below 1.5×10^{-3}). An S-shaped pair of folds is what a cusp of the fold curves leaves on a nearby cut. When these boundaries are continued in the plane in float64 (Figure 28), each of the two fold curves reverses its parameter tangent once. There the oriented quadratic coefficient b of the fold's normal form passes through zero (upper fold: -23.9 to $+244.6$ between adjacent points, with $a \approx 2.0$; lower fold: $+31.2$ to -58.7, with $a \approx 0.55$) while a stays away from zero. Located along the curves, the two cusps lie

at $(\alpha, \log_{10} T) = (0.067213, 0.198983)$ and $(0.066611, 0.189382)$. There $b \operatorname{sign} a$ is below 1.3×10^{-5} of its bracket values, $c = -1.41 \times 10^6$ and -5.61×10^5, and $\det J = 58.2$ and 55.2, all evaluated at the located points. The flip curve of the period-14 bubble lies inside the body and closes around the map's period-14 cells. The cubic coefficient is nonzero at the points that bracket both sign changes: $c \approx -1.4 \times 10^6$ to -2.0×10^6 (step agreement 5 %; the smallest steps converge like the square of the step, and extrapolation leaves the sign) and -5.8×10^5 to -5.3×10^5. This supports nondegenerate cusps inside both brackets, and $c < 0$ agrees with the three coexisting period-7 orbits, two of them stable, found between the S-shaped folds. At a fold point near the lower cusp (where $b = 31.4$) the estimated unfolding is nonsingular too ($\det J = 53.5$ in the same coordinates, singular values 262 and 0.20). Its local predictions at six frozen offsets matched the orbits that Newton solves seeded at the predicted roots and the base state recovered: in number at all six and in stability at five. The miss, at the largest offset beyond the edge of the region with three orbits, marks the range of the cubic model. The hostile period-7 body thus has the same skeleton as the period-6 window of Section 8.

Its interior is shrimp-like too. At $\log_{10} T = 0.1830$, float64 bifurcation diagrams through the body (Figure 27; 1400 cells per line) show the period-7 orbit doubling to periods 14 and 28 and entering chaos (largest exponent up to +0.1) for $\alpha \approx 0.0659$–0.0673. Windows of periods 28, 35 and 42 sit inside the chaotic band, and the orbit then undoubles through 28 and 14 back to 7; the exponent touches zero at every doubling. Newton's method converges at map cells to primitive orbits of periods 14, 28 and 21 (residuals 3×10^{-15}, 2×10^{-14} and 5×10^{-13}; spectral radii 0.08, 4×10^{-4} and 0.006). Periods 21, 35 and 42 are three, five and six times the body's period: the satellite windows expected in the chaotic bands of a period-7 cascade. Continued in temperature, the period-14 orbit flips into period 28 near $\log_{10} T = 0.1841$ and back near 0.1857, the two ends of a doubling bubble that holds period 28 and chaos. Lower down, its branch is S-shaped (folds near 0.1824 and 0.18244, with three coexisting period-14 orbits between them) before another flip near 0.1819. All these crossings are simple (next modulus below 4×10^{-5}). A 160^2 float64 zoom of the lower body (Figure 26) shows this anatomy in the plane. The period-14 domain encloses a period-28 ring around a chaotic core. Period-21 windows sit in the chaotic region as thin crescents with legs, and curves of very negative exponent, where the orbits are superstable (the leading multiplier crosses zero), run through the windows as the "endoskeleton" of shrimps does [59]. In this zoom 18 % of cells are chaotic.

Chaos behind a repeating text. Along the temperature line of Figure 27 (left), a longer test (2000 measured steps after 1000 transient steps, twins at offsets 10^{-6} and 10^{-9}, whose exponents agree to within 0.0031 in every chaotic cell) finds 182 chaotic cells among 1400 between $\log_{10} T = 0.176$ and 0.191. In 65 of them the leading token repeats exactly for all 2000 steps while the state moves chaotically; in 59 of these the period of the text is a multiple of 7, most often 7 or 28. Figure 24 shows the strongest case ($\log_{10} T = 0.18995$, largest exponent 0.091 at both offsets). Its 2000 sampled states fall into seven separated groups, one per phase of the text, visited in a fixed order and at least 0.098 of the state norm apart (a property of the samples, not a proof of separated invariant sets), and each group lies inside one cell of the partition by leading token, so the text `I!Yes heels`↩↩`pee`↩↩ repeats at every one of the 2000 steps we followed. Inside a group the sampled states lie close to a curve (a strongly anisotropic cloud). Read once per repetition through the probability of the last token, the sampled return relation is well approximated by a smooth map with a single quadratic maximum, the shape of the logistic map: a polynomial of degree 8 fits it to 4×10^{-4} of its range, and the fitted map's Lyapunov exponent, 0.647 per repetition, is within 2 % of the full map's 0.636 per seven steps, measured with twin trajectories in all 896 dimensions. The

scalar approximation holds across most of the cascade. Rerun with the leading-token sequences saved, 62 of the 65 cells are again chaotic behind a repeating text, and in every one of them several phases of the text (a median of ten) give single-valued return maps; in 59 cells the best-fitting phase has a single interior stationary point (a maximum in 36, a minimum in 23). If the dynamics were exactly one-dimensional and each phase's probability a smooth invertible coordinate on it, all single-valued phases would give the same asymptotic exponent, since exponents do not change under such a change of coordinates. We therefore fit each phase's map by a polynomial whose degree is chosen by cross-validation over random folds of the return-map pairs, and take the median over the phases. It matches the full map's exponent within 10 % in 51 of the 62 cells (median discrepancy 3.7 %; within 10 % in 23 of the 24 cells with seven-token text, which give 285 repetitions in 2000 steps; Figure 25). This analysis was adopted after we inspected an outlier of our first, pre-specified analysis (a fixed degree of 8 at the single best-fitting phase), which gave 49 of 62 and underfitted one cell whose maps are steep and have two turning points; we report both. An exploratory temporal split (fit on the first 60 % of each record, test on the last 30 %) gives small value-prediction errors at the best scalar phase, chosen on the test segment itself, but in 30 of 61 cells two- or three-delay predictors improve the held-out error by 20 % or more, which the pre-stated rule counts against scalar sufficiency. These results support an empirical scalar description, not a proved one-dimensional reduction: a low-residual fit does not show that the observable is invertible, and even exact agreement of exponents would not establish a one-dimensional invariant curve (two uncoupled logistic maps observed through one coordinate agree in the same way). An independent implementation reproduced the strongest cell (the same leading token at all 2000 steps; exponent 0.089 against our 0.091; scalar discrepancy +2.3 % under the first rule and +1.5 % under the amended one) and both cross-cell counts (49 and 51 of 62). At two further cells it chose, it found good agreement after the amendment in one (text period 10, +0.2 %, although its fitted maps have both a maximum and a minimum) and a persistent discrepancy of about 18 % in the other (period 28), despite a small value error. These are the chaotic bands that follow a period-doubling cascade, and a smooth unimodal map is the setting of Feigenbaum's universality [21]; a fitted maximum on the sampled range does not by itself establish that theorem's hypotheses, and we have shown the scalar approximation for this cascade only. A reader of the text sees a loop. Only the probabilities, or the state, reveal the chaos.

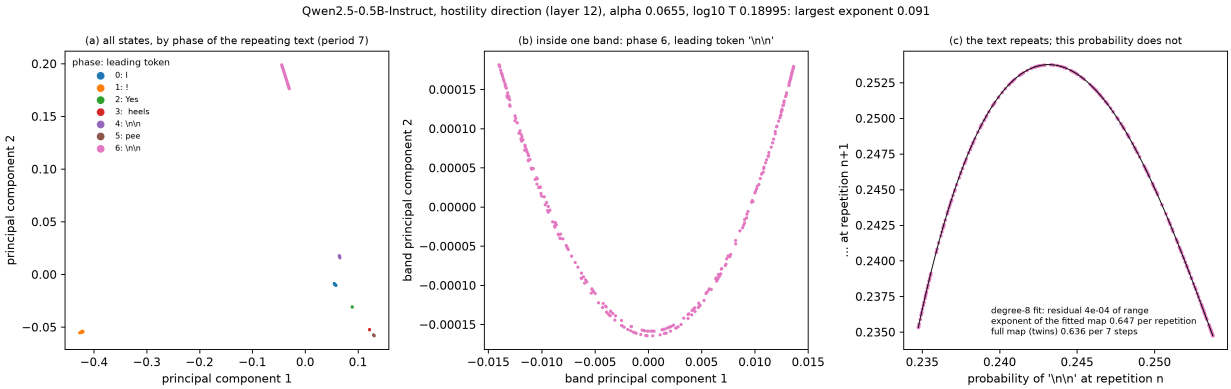


Figure 24: Chaos behind a repeating text: the soft-token map of Qwen2.5-0.5B-Instruct steered along hostility at layer 12 ($\alpha = 0.0655$, $\log_{10} T = 0.18995$; float64, 2000 steps after 1000 transient steps). (a) All states on their first two principal components, coloured by the phase of the seven-token text that the leading token repeats: seven separated sampled clouds. (b) The widest cloud (phase 6, leading token a double line break) on its own principal components: the sampled states lie close to a curve. (c) The probability of that token at successive repetitions of the text, with a fitted return curve with one quadratic maximum (black: polynomial of degree 8) whose Lyapunov exponent is within 2 % of the full map's.

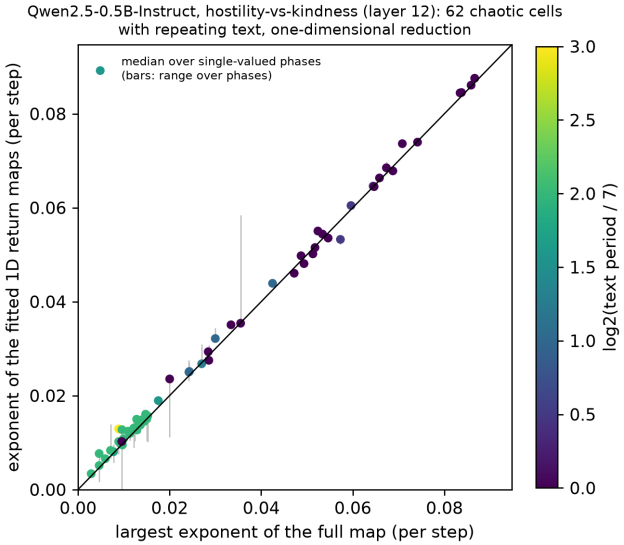


Figure 25: Scalar return-map approximation across the hostile cascade (Qwen2.5-0.5B-Instruct, hostility at layer 12, $\alpha = 0.0655$, $\log_{10} T \in [0.176, 0.191]$): for each of the 62 chaotic cells whose leading-token text repeats exactly, the Lyapunov exponent of the fitted return maps of the leading-token probability (per step; median over the single-valued phases of the text, bars: their range; polynomial degree chosen by cross-validation) against the largest exponent of the full map from twin trajectories. Colour: text period relative to 7. The line is equality. With seven-token texts the 2000 steps give 285 repetitions; with 28-token texts only 71 or 72 observations per phase (70 or 71 return pairs), and the fitted maps then overestimate small exponents.

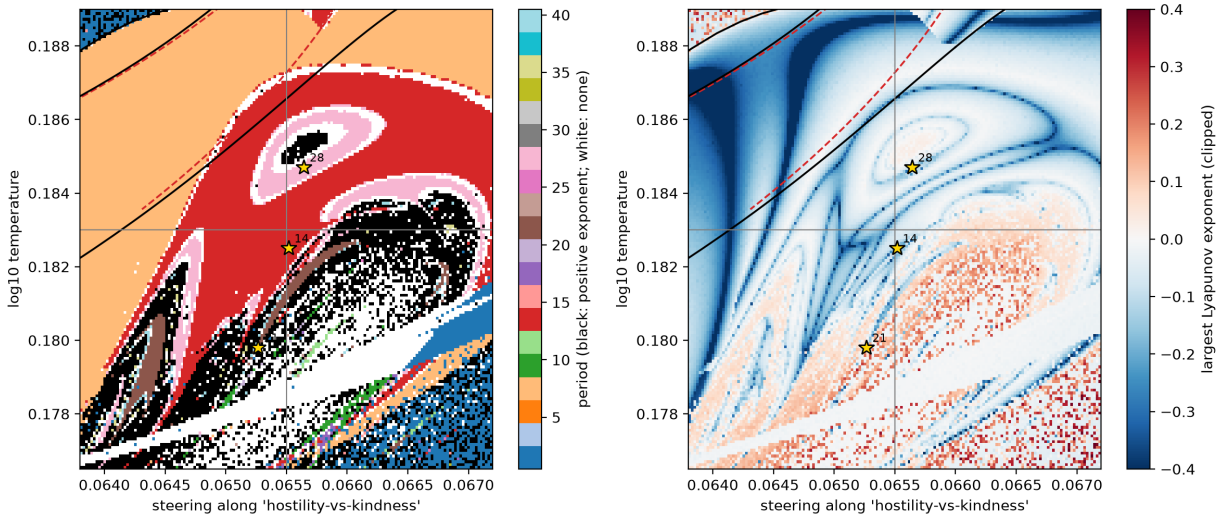


Figure 26: The lower part of the period-7 body of the hostility-steered Qwen2.5-0.5B-Instruct map (layer 12) at higher resolution (160^2 float64 cells). Left: period of the attractor (black: positive largest exponent; white: no period up to 48). Right: largest Lyapunov exponent (clipped); the darkest curves are superstable lines. Both panels: float64 fold (black) and flip (red, dashed) curves of the period-7 orbit, the two lines of Figure 27 (grey) and the Newton-verified orbits of periods 14, 28 and 21 (stars).

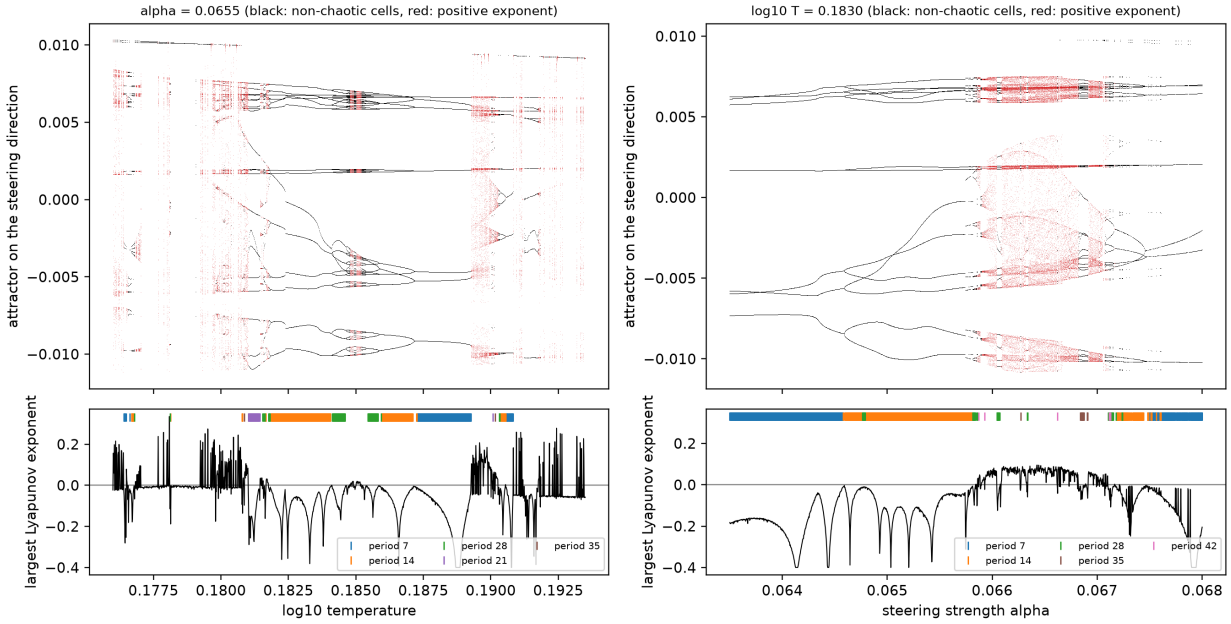


Figure 27: Inside the hostile period-7 body (Qwen2.5-0.5B-Instruct, hostility direction at layer 12): float64 bifurcation diagrams along temperature at $\alpha = 0.0655$ (left) and along the steering strength at $\log_{10} T = 0.1830$ (right). Top: the last 160 states of the attractor reached from the fixed start, projected on the steering direction (red: cells with a positive largest exponent). Bottom: largest Lyapunov exponent, with the cells' periods marked. On the right, period 7 doubles to 14 and 28, enters chaos with windows of periods 28, 35 and 42, and undoubles back to 7. Labels come from 400 measured steps after 1000 transient steps; with 2000 measured steps at two twin offsets, 135 of the 235 red cells on the left stay chaotic and 91 are periodic (Section 15, Limitations).

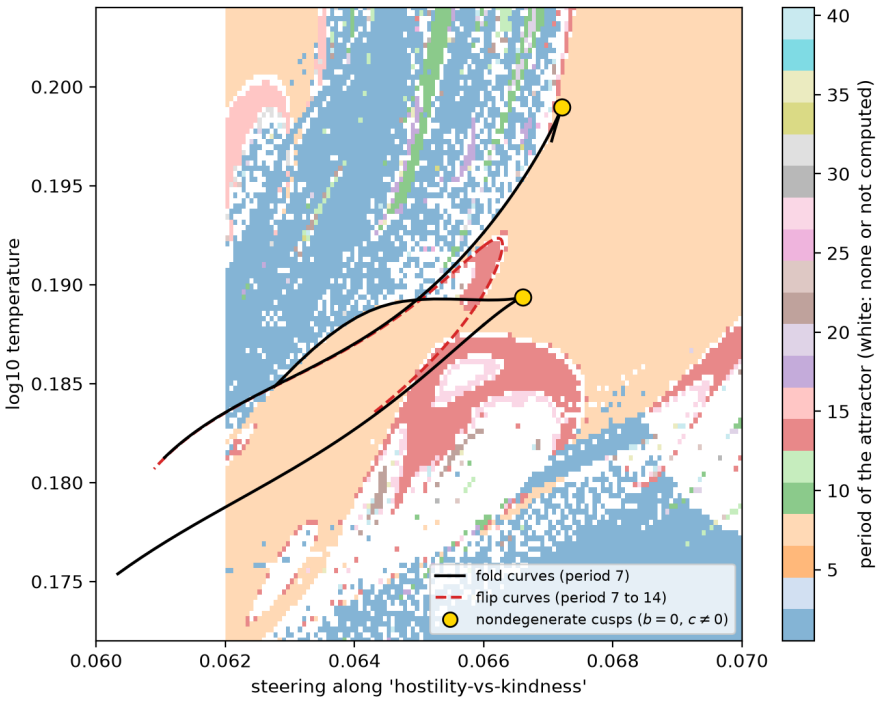


Figure 28: The period-7 body of the spiral region (Qwen2.5-0.5B-Instruct, hostility direction at layer 12): attractor periods of a float64 zoom (white at $\alpha < 0.062$: outside the zoom), the float64 fold curves of the period-7 orbit (black) and the flip curves into period 14 (red, dashed), with the two located cusps (gold).

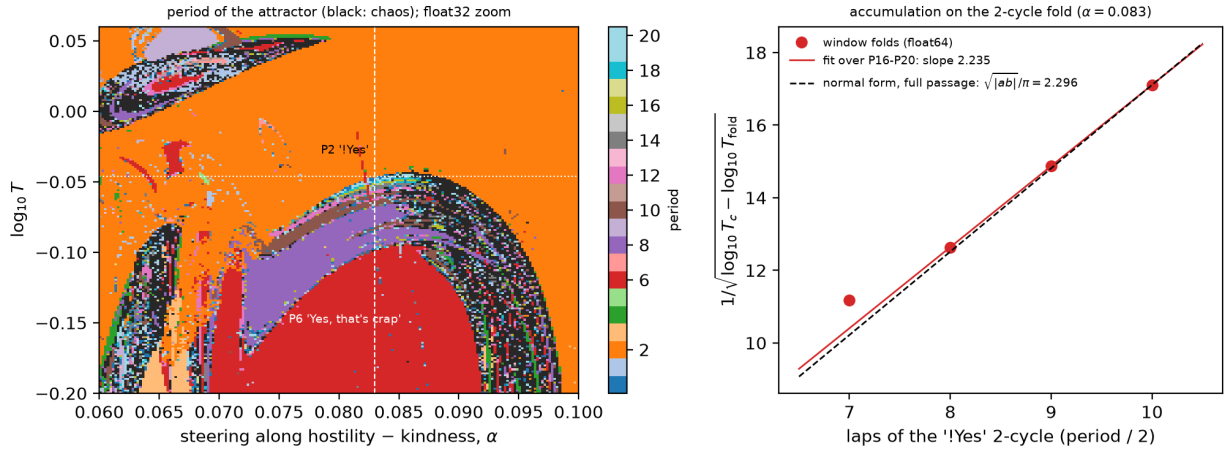


Figure 29: The hostility-steered soft-token map of Qwen2.5-0.5B-Instruct (layer 12). Left: period of the attractor in a float32 zoom (black: chaos), with the period-6 "Yes, that's crap" dome and its period-adding bands; dashed, the cut at $\alpha = 0.083$; dotted, the float64 fold of the "!Yes" 2-cycle. Right: float64 fold edges of the windows (periods 14–20) against the number of "!Yes" laps; the dashed line is the normal-form prediction for full passages, anchored at the last window.

15 Discussion

What theory allows, and what is new. Chaos and periodic windows occur in many nonlinear feedback systems. In a specified ensemble of random recurrent networks with asymmetric couplings, chaos sets in above a critical gain, sharply in the large-network limit, and finite networks pass through limit cycles of increasing complexity on the way [19]; how the transition looks depends on the architecture and the transfer function [80]. Discrete-time random networks, $x(t+1) = \tanh(gJx(t))$, iterate a map as our loop does: their fixed point loses stability when the gain times the spectral radius of the couplings exceeds one, the first bifurcation of large networks is a Neimark–Sacker bifurcation, and chaos follows by a quasi-periodic route, through a torus and frequency locking, over a range of gain that narrows as the network grows [20]. Our largest model narrows differently: Qwen2.5-72B passes from a confident echo to a nearly uniform fixed point within 0.03 decades of temperature, and its chaos is confined to narrow bands, whereas in random networks chaos is generic above the transition. External input can suppress that chaos [81], and deep networks at initialization show an order-to-chaos transition in signal propagation [82, 83]. In our map, temperature plays the role of an inverse gain and the steering vector that of an input. Section 3 makes the correspondence concrete through an exact chain-rule factorization, with the softmax derivative $(\operatorname{diag} p - pp^\top)/T$ in the place of the slope of the transfer function; the identification with the random-network transition remains an analogy, not a transfer of that ensemble's theorem. Feigenbaum scaling holds for a class of one-dimensional maps with a quadratic maximum [21], and periodic windows, shrimp-shaped domains with their cusps and spirals, and doubling cascades are found in many dissipative two-parameter systems [22, 59, 60, 84] and in experiments [85]; intermittency near folds is classical [58]. These results motivate looking for chaos, windows and shrimp-like structure in a language model's feedback map, but they do not imply it. They do not say whether a trained language model realizes it, where in its parameter plane, with which texts, or whether the theory's quantitative predictions hold for that model. This paper shows that the models we studied do realize it, with concrete texts as attractors, so that a bifurcation changes what the model writes, and that the quantitative predictions hold in the cases we tested.

As far as we know, no earlier work has computed periodic orbits and their multipliers, continued fold and flip bifurcation curves and located cusps in the feedback map of a released pretrained language model. Nor do we know of earlier quantitative predictions from such a model's derivatives, for these objects, that were confirmed prospectively, as ours were for passage times at new folds (23 of 24 forecasts within one iterate, with the all-row criterion failing by one row) and for the orbits recovered and their stabilities at frozen offsets near numerically located cusps; derivative-based normal-form predictions in a task-trained looped transformer precede ours [18]. A frozen spectral prediction of the response to relaxation is confirmed on held-out data in Huginn, and one periodic orbit of the full transformer map of Qwen2.5-0.5B-Instruct is proved by interval arithmetic; we know of no earlier instance of either for a language model.

Is steering the origin of the chaos? Steering adds a fixed vector to the residual stream, so chaos along a steering axis could belong to the push rather than to the model. At $\alpha = 0$, however, the map is the unsteered model's own feedback map, with temperature as its only parameter. The $\alpha = 0$ columns of the hostility maps contain cells with positive finite-horizon exponents in Llama-3.1 8B (26 of 64 temperatures), Qwen2.5-7B (15 of 64) and OLMo-2 7B (11 of 64), while the near-zero columns of the 0.5B and 1.5B maps contain one such cell each (of 128 and 64). A longer float64 sweep of the unsteered 0.5B map (512 temperatures over $\log_{10} T \in [-1, 0.3]$, 1000 transient and 2000 measured steps) finds a narrow chaotic band at $\log_{10} T \approx 0.15$–0.21 (7 temperatures, exponents up to 0.28), with a period-5 loop at 437 of the 512 temperatures. The same long sweep finds 96 chaotic

temperatures in Llama-3.1 8B, between $\log_{10} T = -0.20$ and 0.24 (exponents up to 0.66), interleaved with periodic windows of periods 2 to 52; within $\log_{10} T \in [-0.3, 0.3]$ its chaotic share, 41 %, equals that of the coarse column. Qwen2.5-7B has 65 chaotic temperatures between $\log_{10} T = -0.07$ and 0.28 (exponents up to 0.74), and OLMo-2 7B has 32 between 0.04 and 0.14 (up to 0.30). Steering is therefore not needed to reach chaotic regimes in any of these four models. Temperature alone produces chaos with embedded windows, broadly in the 7B and 8B models and in a narrow band in the 0.5B model (Figure 30). Across the fifteen models of Section 12, the regular-grid sweeps find unsteered chaos in eleven of the fourteen dense models (in Llama-3.1-70B at two temperatures that pass a three-offset test), but not in Qwen2.5-72B, Llama-3.3-70B, OLMo-2-13B or the mixture of experts; finer lines found candidates in the first two, and the Qwen2.5-72B window, recomputed in float64 throughout, is periodic.

Steering does widen the chaotic regions, but not as a property of the concept: along norm-matched random directions, the 0.5B maps are as chaotic as along the concept directions (Figure 31). On the same 128^2 grid, 7.6, 8.3 and 12.5 % of cells are chaotic for three Gaussian directions with the philosophy vector's norm, against 6.8 % for philosophy and 7.5 % for gratitude. In all five maps the chaotic share rises from about 1 % at $|\alpha| < 0.01$ to between 5 and 22 % at larger $|\alpha|$, and a chaotic region at high temperature, with windows reaching down, recurs. In this model, chaos under steering is therefore a generic response to perturbing the residual stream at this norm. Rotating the steering direction makes the same point continuously: along the great circle from one Gaussian direction to another, 90° away, at constant norm, the chaotic regions move and change shape, but the chaotic share stays between 13 and 24 % at every one of 103 angles (Figure 32; a video of the rotation is in the supplementary material). The direction does select the texts. The largest attractor family in all five maps is the unsteered period-5 loop, but beyond it only 26 to 56 % of one map's cells carry an attractor text that appears anywhere in another map. Llama-3.1 8B agrees. On its hostility plane, two norm-matched random directions at the same layer give 26 and 17 % chaotic cells against 28 % along the concept direction; all three carry 34 to 44 % near $\alpha = 0$, and the random directions thin out at large $|\alpha|$. At layer 19, where a sparse autoencoder trained on this model reads the residual stream [73], the concept vector rebuilt there gives 28 %. The decoder direction of the autoencoder feature most specific to hostility on held-out texts (chosen by a rule fixed before any map; cosine 0.10 with the concept vector) gives 29 %, a random live feature 27 % and a Gaussian direction 39 %, all at the same norm. In this model, chaos therefore occurs along autoencoder-feature, contrastive and random directions alike, with direction-dependent prevalence; at 70B the directions differ more (Section 12).

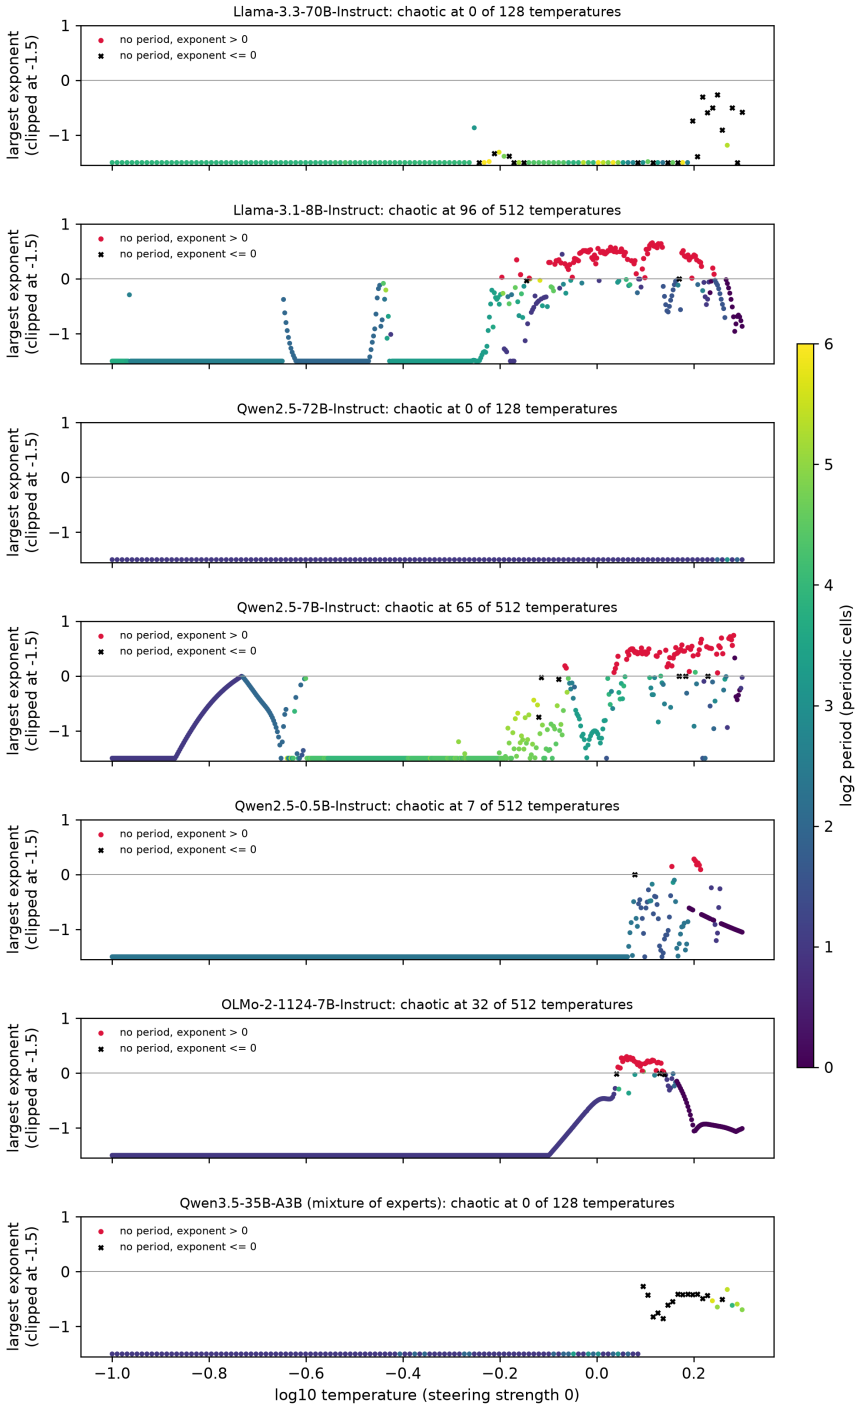


Figure 30: The unsteered soft-token map ($\alpha = 0$) along temperature, top to bottom: Llama-3.3-70B-Instruct, Llama-3.1-8B-Instruct, Qwen2.5-72B-Instruct, Qwen2.5-7B-Instruct, Qwen2.5-0.5B-Instruct, OLMo-2-1124-7B-Instruct and the mixture of experts Qwen3.5-35B-A3B. Largest Lyapunov exponent (clipped at -1.5) from float64 sweeps of 512 temperatures (128 for the 70B-class and mixture-of-experts models; 1000 transient and 2000 measured steps). Red: no period up to 64 and an exponent above 10^{-3}; black crosses: neither a detected period nor a positive exponent (unresolved); other points have a detected period, coloured by $\log_2$ of the period. Table 13 summarizes these and the wider sweeps of the other models.

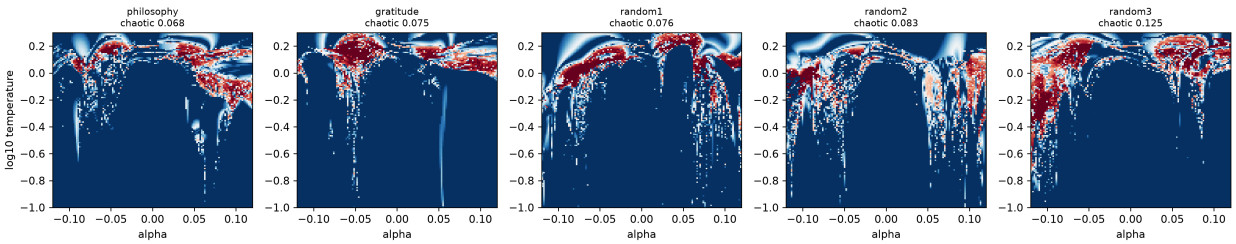


Figure 31: Steering controls in Qwen2.5-0.5B-Instruct (128^2 float64 cells, steering at layer 12, one grid and one set of settings): largest Lyapunov exponent along the philosophy and gratitude directions and along three Gaussian directions rescaled to the philosophy vector's norm. Titles give the chaotic share (finite-horizon exponent above 10^{-3}, no period up to 48).

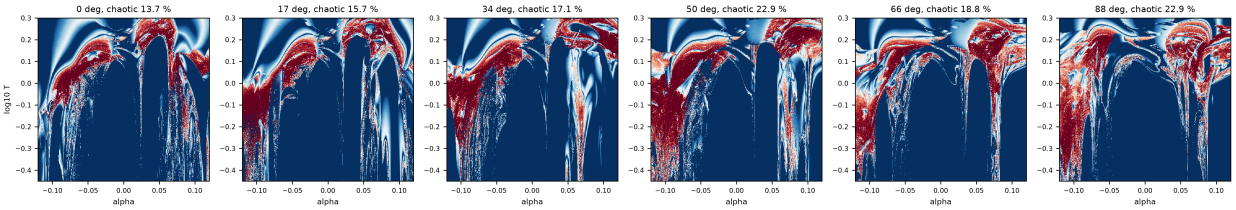


Figure 32: The soft-token map of Qwen2.5-0.5B-Instruct (steering at layer 12) while the steering direction rotates at constant norm from one Gaussian direction (0°) to another (90°): largest Lyapunov exponent (red: positive) at six of 103 angles, with the chaotic share. Float32 maps of 296×512 cells over $\alpha \in [-0.12, 0.12]$, $\log_{10} T \in [-0.45, 0.3]$; exponents from 80 measured steps at twin offset 10^{-4}; each map continues from the final states of the previous angle after 40 steps, and every tenth starts from the fixed start after 300 steps. Where attractors coexist, continued and fresh starts can reach different ones, so the maps, and the rotation video, can jump at every tenth angle and across the 17 of 120 planned angles that did not finish within the run limit. These short-horizon labels are coarser than those of the float64 maps.

What is and is not claimed. The objects studied here are invariant sets and bifurcations of an exactly defined smooth map built from a frozen model's own computation. We do not claim that ordinary sampled generation is chaotic in this sense (Section 10), nor that any window is a certified shrimp. The period-6 window and the hostile period-7 body have, numerically, shrimp-like organization: fold curves with numerically located cusps (simple critical multiplier; quadratic coefficient vanishing at a point where the cubic coefficient is nonzero and the unfolding is regular) and flip curves inside, and in the hostile body curves of very negative exponent, where the orbits are superstable, run through the windows. That is less than the full anatomy of a shrimp: we have not computed superstable curves or located a doubly superstable centre, where two superstable curves of the window's orbit cross, nor determined the windows' symbolic sequences or tested for a periodicity hub [22, 61]. In a map that is effectively one-dimensional with a single critical point, a primitive orbit cannot be doubly superstable, so true shrimp centres would require a second critical point in the reduced dynamics. A generic codimension-two cusp also requires the two unfolding coefficients to depend on the two parameters with rank two. We estimated this at a fold point near one cusp of each window (determinants 2737 and 53.5), and there the local model's predictions at twelve frozen offsets matched the orbits recovered by Newton solves seeded at the predicted roots and the base state, in number at all twelve and in stability at eleven; the search was not exhaustive. At the nine located cusps themselves, $|\det J|$ lies between 14 and 2416. These are float64 computations validated by step-size checks (rung 1), not computer-assisted proofs. At the two continuation points that

bracket the lower cusp of the hostile body, an independent implementation of the model reproduced the normal-form coefficients, including the sign change of b: $|b|$ within 1.2 % and c to two figures, with the opposite orientation of the null vector, which flips the signs of a and b together.

Limitations. The detailed orbit studies cover one model family (Qwen2.5, 0.5B and 1.5B), with coarse maps and three float64 orbits at 7B; the two other families (OLMo-2 7B and Llama-3.1 8B) have maps, including 128^2 refinements for Llama, but no continued orbits. The size comparison of Section 12 uses regular temperature grids that can miss narrow windows, a window that missed most of the larger models' chaos, and mostly one twin offset per label. The two-GPU runs of the 70B-class models computed their final normalization in float32 until we fixed it; their labels rest on offsets of 10^{-6} or more, and the cells we rechecked, the Qwen2.5-72B jump window and steered candidates in float64 throughout and samples of the Llama maps at three offsets, behave as stated, but the rest of those maps were not recomputed. Neither were the four coarse lines of the frozen 72B test; after the fix we computed only fine lines through their entropy jumps, one of which passes through their single chaotic temperature. The maps use one prompt (Section 12 adds ten for Qwen2.5-7B and -14B, and four for Qwen2.5-72B, its base model, and OLMo-2-7B and -32B), and steering uses contrastive vectors, checked against norm-matched random and sparse-autoencoder directions in two models. The system is a constructed loop: Section 10 extends it to memories of up to sixteen soft tokens and to the growing context of Soft Thinking, but sampled decoding of discrete tokens is a different system. The evidence consists of numerical validation, cross-checks and prospectively frozen passage tests (including failed all-row criteria), with a computer-assisted proof for only one orbit at one parameter point. Chaos labels are finite-horizon estimates (300 measured steps after 300 to 500 transient steps), and no chaotic cell is proved chaotic. A computer-assisted proof of chaos, such as a horseshoe verified by covering relations, would need enclosures of the map over boxes many orders of magnitude larger than the period-8 proof's (radius 2^{-640}). With the ball arithmetic used there, a box of radius 2^{-80} already loses the positivity of an RMS normalization at layer 3 (the enclosures overestimate through the layers), so such a proof needs other verified methods. A horizon check retested random cells of two float64 maps (Figures 1 and 2) with 1000 transient and 2000 measured steps at two twin offsets: 86 and 92 of 96 cells labelled chaotic stayed chaotic (all but one of the rest settled on periodic orbits), all 96 periodic cells in each map stayed periodic, and 18 and 20 of 32 cells with neither label turned out to be slow periodic transients. Near the accumulation of a cascade the labels are less reliable, because orbits there converge slowly or wander chaotically before they settle (transient chaos [86]): on the hostile line of Figure 27 (left), 135 of the 235 cells labelled chaotic from 400 measured steps stayed chaotic and 91 were periodic. The Feigenbaum ratios are finite-order values that rise toward δ (the last is within 1 %); the limit itself is not computed. All orbit, bifurcation and chaos results concern the constructed map, not sampled generation with a growing context. With sampling, the bridge of Section 10 tests only whether the bottleneck's laminar phases survive the noise; its other tests are deterministic (greedy decoding, longer fed-back memories and Soft Thinking's growing context).

Steering strength is a bifurcation parameter. In the soft-token map, small changes of the steering strength switch the model between distinct repeating texts, chaos and coexisting attractors. Near a fold curve, behaviour becomes intermittent, with laminar phases whose duration the fold's normal form predicts in advance. This argues for treating an activation-steering coefficient like any control parameter of a nonlinear system. Its bifurcation set should be characterized before an operating point is chosen, history dependence should be expected near boundaries, and a monotone, smooth dose response should not be assumed. Outputs that fall into verbatim repetition at high

steering doses are naturally treated as quality failures. In the present framework such loops are candidate attracting periodic orbits, and the dose at which they appear is a boundary that can be mapped like the others. The deterministic map is where such sets can be defined precisely and computed numerically. Whether its boundaries predict boundaries of sampled behaviour is open. The bridge of Section 10 shows that sampling noise shortens the laminar phases of the bottleneck, by an amount that depends on how many samples are averaged.

What this suggests for interpretability and alignment. Our results come from open models of 0.5B to 72B parameters and a constructed loop, but they suggest three things about models whose outputs feed back into their inputs, as in self-dialogue, agent loops and pipelines of model calls. First, where the loop is chaotic, a perturbation of 10^{-6} grows to order one within roughly 15 to 50 steps (the largest unsteered exponents of the 7B and 8B models are 0.30 to 0.88 per step). Even a complete account of single forward passes then predicts a long trajectory only if the state is known far more precisely than any interpretation supplies. Stepwise summaries can be inconclusive: every step of the period-8 orbit expands some direction while the cycle contracts, so the ordered product of the step Jacobians, not a summary of each step, captures the loop's stability (Section 3). What remains predictable are the invariant sets, their statistics and the boundaries between them. This extends the analysis of trained recurrent networks through their fixed points [28, 29] to the periodic orbits, chaos and bifurcations of a pretrained transformer. Second, feature dictionaries say which features an attractor occupies, but the boundaries between attractors, which decide what a small change does, are properties of the map's derivatives. In our 8B controls, a sparse-autoencoder feature direction, contrastive and random directions all drove the map into chaos (27 to 39 % of cells); at 70B the hostility-related directions gave more chaotic cells than the random ones (4.3 and 3.5 % against 1.1 and 0.2 %; Section 12). Steering acts through the gain of Section 3, which is not specific to concept-labelled directions, but a direction's geometry and its meaning can be correlated. Text can also hide the dynamics: in the hostile cascade, a seven-token text repeats exactly while the state underneath is chaotic (Figure 24), so a monitor that reads only the emitted text would call a chaotic regime periodic. Third, conclusions drawn at one operating point need not transfer: which loop a model enters depends on how much context it keeps (Table 11), on the path by which a boundary is approached, and on the sampling noise (Section 10). Treating a looped model as a dynamical system, with regimes and boundaries to be mapped, rather than as a function to be evaluated point by point, seems to us the appropriate stance for both interpretability and safety evaluation.

Is the chaos useful? The chaos we map is not a source of new content. The chaotic cells of Table 2 recombine the phrases of the neighbouring loops in an order that did not settle over the recorded horizon: the trajectory wanders among the ghosts of nearby periodic orbits, which gives variety in sequence but not in substance. The periodic loops are the familiar degeneration of decoding into repetition [24, 25]. Neither regime is good text, and the diversity of ordinary generation comes from sampling and a growing context, which wash out much of this structure (Section 10). What the chaos does matter for is predictability and control. Positive finite-horizon exponents limit open-loop prediction of a trajectory at fixed state precision (our experiments do not establish mixing or limits on feedback stabilization); a stable attractor is predictable, but it can also hold a model in a state it should leave, such as the hostile loops of Section 14. On the paper's prompt, the largest Qwen2.5 and Llama models sit near the stable end, but across prompts a 14B Qwen model is more chaotic than the 7B, at higher temperatures, and OLMo-2-32B is among the most chaotic models we tested (Section 12). Whether regimes near the boundary between order and

chaos are especially useful, as has been argued for computation in dynamical systems and recurrent networks [87, 88] and shown for signal propagation in deep networks at initialization [82, 83], is open for language models: our maps locate such boundaries but do not test it.

Which findings might reach frontier models. Today's frontier systems include closed models such as GPT-6 Astra and Claude Opus 5.5 and open-weight mixture-of-experts models such as the GLM-5 family. Some of our findings rest only on structure that every smooth softmax-readout language model with a fixed prefix shares and should carry over: the factorization of the loop's derivative (a covariance for tied embeddings, a cross-covariance otherwise), contraction at sufficiently high temperature, the reduction of zero-temperature attractors to greedy cycles of whatever memory the loop keeps, and the role of a steering coefficient as a bifurcation parameter. Others belong to particular weights and will not carry over as such: the attractor texts, the positions of windows, folds and cusps, and every number we report; a second checkpoint of the same architecture already has different attractors (Section 11). Size gives no guarantee of stability. In a fixed temperature window the unsteered chaos nearly vanished in the largest Qwen2.5 and Llama models, but across prompts a 14B Qwen model was more chaotic than the 7B, at higher temperatures, OLMo-2-32B was the most chaotic OLMo model, and post-training changed the chaos of two Llama checkpoints of the same size (Section 12). For an open frontier model we would search for chaos where its own next-token distribution spreads, across prompts. Our one mixture-of-experts model, Qwen3.5-35B-A3B, gave no positive exponent without steering and irregular cells only under hostility steering, and hard routing falls outside the smooth theory of Section 3 (Section 11). Reasoning models open a thinking block that changes the prefix, post-training may remove some attractors and create others, and deployed inference samples tokens, in low precision, over long contexts; our bridge results show that sampling shortens laminar phases and that a longer context changes which loops are entered. Closed models permit none of these computations: the map needs the input embedding matrix, full next-token distributions and a way to feed a vector in place of a token, and its analysis needs derivatives. For closed models the method and the questions transfer, not the numbers.

16 Evidence and reproducibility

We grade claims on an evidence ladder. Rung 1 means numerically validated (tolerances, residuals, convergence, precision varied); rung 2, cross-checked by an independent estimate, implementation or hardware; rung 3, a computer-assisted proof.

Search scale. The soft-token results come from an exploratory search. An initial census of the raw data at 00:19 UTC on 14 October 2026 counted 34 distinct maps (987 308 cells, 27 maps in float64) in six models, each with one prompt and one of three steering directions or none, plus one partial map. Of these, 29 maps and the partial one appear here; the five omitted are early discovery or side maps. Later additions are not in that count: the steering-control maps of Section 15, the mixture-of-experts maps, and the sweeps, lines and maps of the models of Section 12. The labelled maps contain about 2600 periodic families of at least six cells, counted per map. We singled out about 26 windows or families for float64 work: 20 are reported in detail, 2 only as map features and 4 not at all (an orbit whose continuation was abandoned, a cell that proved aperiodic, a band whose float32 orbit did not converge and a pilot that found no cusp within a bounded search). Of 54 Newton orbit attempts, 7 failed. All nine tangent reversals on continued fold curves (seven found by 14 October, and two on traces of a period-4 fold curve continued on 15 October) are numerically located cusps; the two reversals of a continued flip curve are smooth turns. Two of

the five regions we flagged as shrimp candidates are reported as shrimp-like. One further language model (Ouro-1.4B) was explored and is not reported; its apparent chaos was a long transient. Every window was chosen after the maps were inspected. Only the passage forecasts, the cusp-unfolding offsets and Huginn's relaxation test were fixed before the computations that test them (Table 14).

	explored	reported	note
Soft-token maps (distinct)	34	29	plus one partial map, reported
cells	987 308	893 868	omitted: 5 discovery or side maps
float32 maps	7	4	
Models with maps	6	6	one prompt; 3 directions and $\alpha = 0$
Periodic families (≥ 6 cells)	≈ 2600	–	counted per labelled map
Windows given float64 work	≈ 26	20 (+2)	+2: map features only
Newton orbit attempts	54	–	7 failed
Tangent reversals on fold curves	9	9	all located cusps (2 found on 15 Oct.)
Unfolding offsets, frozen	12	12	counts 12/12, stability 11/12
Shrimp-candidate regions	5	2	reported as shrimp-like
Tests fixed before outcomes	6	4	2 all-row criteria failed by one row
Language models loaded	8	7	Ouro-1.4B not reported

Table 14: Search scale on 14 October 2026, before the steering-control maps (Huginn's frozen test counted once; of the two tests not reported here, one belongs to the separately reported sampled-dialogue study and one, a cusp pilot in a second checkpoint, is the unreported bounded search).

- Rung 3: the period-8 stripe orbit of Table 1. At $\alpha = 0.071674$ and $\log_{10} T = -0.115$, read as exact binary64 rationals (the temperature is the real number $\exp(\ln 10 \cdot \log_{10} T)$ for that exact dyadic $\log_{10} T$), the expected-embedding map of the stored Qwen2.5-0.5B-Instruct checkpoint (all 24 layers, the full 151 936-token readout, the fixed 37-token prefix, steering after zero-based block 12) has a primitive period-8 orbit, which is the unique multiple-shooting root in a product of eight boxes of coordinate radius 2^{-640} around refined centres, and the orbit is locally attracting. Ball arithmetic (Arb [89]) encloses the map and its Jacobian over each box (8×896 directions); a Krawczyk-type inclusion [90–92] for the preconditioned shooting system gives existence and uniqueness, separation of the first point from the points one, two and four steps later gives the primitive period, and a norm bound on powers of the enclosed monodromy gives attraction. The map is the real-number map of the stored constants, read as exact dyadic rationals: rotary angles, exponentials and square roots are real functions enclosed by the balls, not floating-point tables. An independent audit in exact rational arithmetic rechecked the saved witnesses (the scalar inequalities, the primitive-period separations and the row norms of the saved monodromy powers); it did not recompute the enclosures of the model or the dense matrix products, which rest on one execution. The proof trusts the Python, FLINT and Arb stack and the stored weights. It does not exclude other period-8 points in the first box, and it says nothing about the window's boundaries or global organization. The record (`certificates/llm-p8-certificate-final-001.json`) pins the SHA-256 hashes of the checkpoint, its configuration and the steering vector, and lists the 37 prefix token IDs.
- Rung 2: the 2-cycle's fold and its normal form. An independent reimplementation of the forward map and solver recovers α_c, a and b (to 6×10^{-15}, 3×10^{-12} and 2×10^{-6} relatively), and ab also follows from the stable-branch multipliers. The same reimplementation reproduces all 18 fixed-entrance counts of the controlled passage runs exactly (Table 3 lists 12 of them).
- Prospectively frozen: the passage forecasts at two new folds (Table 5); 23 of 24 within one

iterate, and the all-row criterion failed by one row.

- Rung 2 as well: the period-2 and period-8 orbits of Table 1 (independent implementation), and the cascade's first four flip parameters (CPU and H100, agreement 4×10^{-15}).
- Rung 2 as well: the hostile 2-cycle fold and its normal form, reproduced by an independent float64 forward map with nested automatic differentiation (location to 6×10^{-10}, ab to 8 parts per million); its controlled passages were also checked independently, with the frozen all-six criterion failing by one row.
- Rung 1: the other orbits, the curves, the nine located cusps, the superstable parameters, the 1.5B flips, the hostile windows' edges, the spiral region's period-7 body and the cascade inside it, the three 7B orbits, and Huginn's float32 spectral estimates.
- Prospectively frozen, not independently reproduced: the scaling predictions P1–P3, the prompt-blocked size tests in Qwen2.5 (7B and 14B; 72B; 72B base against Instruct) and OLMo-2, and the loop-free predictor test of Section 12, all scored by code pushed before the data, together with the three-offset plateau tests of the 70B-class maps and the float64 recheck of the Qwen2.5-72B cells with exact tangent exponents. These are float64 computations at rung 1; no independent implementation has rechecked them.
- Prospectively frozen and independently replayed: Huginn's relaxation test on held-out data, confirmed in both parts; an independent implementation reproduced both from the saved rows and spectra (GSM8K part exactly, item-level part to 10^{-12}) (Section 13).

Hardware. Float64 GPU computations ran on rented NVIDIA H100 (NVL, PCIe and 80 GB HBM3), H200, H200 NVL and A100 80 GB (SXM4 and PCIe) GPUs, one GPU per run except for the 70B-class models, whose layers were spread over two; the float32 frames of the rotation video ran on RTX 4090 GPUs; cross-checks and the interval proof ran on CPUs. The manifest `docs/manifests/runpod-runs.json` records for each run the GPU type, the code commit, the job and its arguments, the wall time and the exit status, and the available hashes of the uploaded code and input archives and of the retrieved results, with absent fields left explicit; per-run receipts and accepted provenance limitations are kept with it. Regenerated at 22:13 UTC on 14 October 2026, it lists 243 runs with retrieved results (code and input-archive hashes for 135, result-archive hashes for 227) and 16 pods that ended without results. GPU types round float64 reductions differently, so quantities used across runs were checked across hardware where it mattered (for example, flip parameters agree between CPU and H100 to 4×10^{-15}). Code: `tools/explore_llm_soft_loop.py`, `tools/bifurcation_curves.py`, `tools/llm_feigenbaum.py`, `tools/explore_latent_loop.py`, `tools/cusp_check.py`, `tools/huginn_relax_eval.py`, `tools/huginn_arc_analysis.py`. Raw-data hash receipts: `evidence/llm-dynamics-data-receipts.json` (to 12 October), the per-run receipts in the research log, and the verified snapshot manifests in `docs/backups/`; independent cross-check and audit receipts are under `codex/` (`crosscheck`, `confirm-fold`, `hostile-fold-crosscheck`, `huginn-design-review`, `huginn-independent-replay`, `paper-review`).

[**Planned:** public code and data release at submission.]

Use of AI assistance

This work was carried out with AI coding and research agents (Anthropic's Claude Code and OpenAI's Codex) working under the author's direction. The agents wrote and ran the code,

executed the experiments, checked results against the raw data and drafted text; other AI systems made independent read-only checks of the results and the manuscript. Frozen designs, run cards, failed attempts, audits and hash-bound records of every run are kept in the project repository. The author directed the study, reviewed the work and takes full responsibility for its content.

References

[1] Xiaojian Li, Yongkang Leng, Ruiqing Ding, Hangjie Mo, and Shanlin Yang. Cognitive activation and chaotic dynamics in large language models: A quasi-Lyapunov analysis of reasoning mechanisms. arXiv preprint arXiv:2503.13530, 2025. URL https://arxiv.org/abs/2503.13530.

[2] Gregorio Jaca, Kristóf Benedek, and János Török. Chaos in reason: How chain-of-thought LLMs can look for an answer. arXiv preprint arXiv:2607.27805, 2026. URL https://arxiv.org/abs/2607.27805.

[3] Julian Arnold, Flemming Holtorf, Frank Schäfer, and Niels Lörch. Phase transitions in the output distribution of large language models. arXiv preprint arXiv:2405.17088, 2024. URL https://arxiv.org/abs/2405.17088.

[4] Kai Nakaishi, Yoshihiko Nishikawa, and Koji Hukushima. Phase transition in large language models and the criticality of natural languages. arXiv preprint arXiv:2406.05335, 2024. URL https://arxiv.org/abs/2406.05335.

[5] Youran Sun and Babak Haghighat. Phase transitions in large language models and the $O(N)$ model. arXiv preprint arXiv:2501.16241, 2025. URL https://arxiv.org/abs/2501.16241.

[6] Huajian Ruan, Jinyang Li, Xingyu Guo, and Lingxiao Wang. Generative criticality in large language model temperature scaling. arXiv preprint arXiv:2606.06238, 2026. URL https://arxiv.org/abs/2606.06238.

[7] Neil F. Johnson, Frank Yingjie Huo, and Bella Xinrui Li. Temperature-driven inversion and nonlinear dynamics in ChatGPT-like AIs. arXiv preprint arXiv:2608.00939, 2026. URL https://arxiv.org/abs/2608.00939.

[8] Zhilin Wang, Yafu Li, Jianhao Yan, Yu Cheng, and Yue Zhang. Unveiling attractor cycles in large language models: A dynamical systems view of successive paraphrasing. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 12740–12755. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.acl-long.624. arXiv:2502.15208.

[9] Ting-Wen Ko and Jonas Geiping. Attractor states emerge in multi-turn LLM conversations. In *Third Conference on Language Modeling*, 2026. URL https://openreview.net/forum?id=umwOiCSh4I. arXiv:2606.30571.

[10] Shaojie Bai, J. Zico Kolter, and Vladlen Koltun. Deep equilibrium models. In *Advances in Neural Information Processing Systems*, volume 32, pages 690–701, 2019. URL https://proceedings.neurips.cc/paper_files/paper/2019/hash/01386bd6d8e091c2ab4c7c7de644d37b-Abstract.html. arXiv:1909.01377.

[11] Shaojie Bai, Vladlen Koltun, and Zico Kolter. Stabilizing equilibrium models by Jacobian regularization. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 554–565. PMLR, 2021. URL https://proceedings.mlr.press/v139/bai21b.html.

[12] Akiyoshi Tomihari and Ryo Karakida. Recurrent self-attention dynamics: An energy-agnostic perspective from Jacobians. In *Advances in Neural Information Processing Systems*, volume 38, pages 155236–155269, 2025. doi: 10.52202/085713-4675. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/cca569b38ec3612904e625dad7a743dc-Abstract-Conference.html.

[13] Xiao-Wen Yang, Ziyu Han, Xi-Hua Zhang, Wen-Da Wei, Jie-Jing Shao, Lan-Zhe Guo, and Yu-Feng Li. Stabilizing recurrent dynamics for test-time scalable latent reasoning in looped language models. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, pages 144670–144682. PMLR, 2026. URL https://proceedings.mlr.press/v306/yang26t.html.

[14] Xinyue Luo and Fei Yu. Prediction dynamics in depth-recurrent language models. arXiv preprint arXiv:2609.21383, 2026. URL https://arxiv.org/abs/2609.21383.

[15] Boyi Zeng, Shixiang Song, Siyuan Huang, et al. PonderLM: Pretraining language models to ponder in continuous space. In *The Fourteenth International Conference on Learning Representations*, pages 91874–91893, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/94572a66c6b0710c5263f15d8ef6b463-Abstract-Conference.html. arXiv:2505.20674.

[16] Zhen Zhang, Xuehai He, Weixiang Yan, Ao Shen, Chenyang Zhao, and Xin Eric Wang. Soft Thinking: Unlocking the reasoning potential of LLMs in continuous concept space. In *Advances in Neural Information Processing Systems*, volume 38, pages 187190–187212, 2025. doi: 10.52202/085713-5629. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/f7396d1c54d51416958d63e285377103-Abstract-Conference.html. arXiv:2505.15778.

[17] Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason E. Weston, and Yuandong Tian. Training large language models to reason in a continuous latent space. In *Second Conference on Language Modeling*, 2025. URL https://openreview.net/forum?id=Itxz7S4Ip3. arXiv:2412.06769.

[18] Gunn Kim. Dynamical phase selection controls compute scaling in looped transformers. arXiv preprint arXiv:2608.26556, 2026. URL https://arxiv.org/abs/2608.26556. Version 1, 27 August 2026.

[19] H. Sompolinsky, A. Crisanti, and H. J. Sommers. Chaos in random neural networks. *Physical Review Letters*, 61(3):259–262, 1988. doi: 10.1103/PhysRevLett.61.259.

[20] B. Doyon, B. Cessac, M. Quoy, and M. Samuelides. On bifurcations and chaos in random neural networks. *Acta Biotheoretica*, 42(2–3):215–225, 1994. doi: 10.1007/BF00709492.

[21] Mitchell J. Feigenbaum. Quantitative universality for a class of nonlinear transformations. *Journal of Statistical Physics*, 19(1):25–52, 1978. doi: 10.1007/BF01020332.

[22] Jason A. C. Gallas. Structure of the parameter space of the Hénon map. *Physical Review Letters*, 70(18):2714–2717, 1993. doi: 10.1103/PhysRevLett.70.2714.

[23] Nicolás Vera Zúñiga. What iterated self-feeding probes of language models measure, and a test that separates the construction from the model. arXiv preprint arXiv:2608.10986, 2026. URL https://arxiv.org/abs/2608.10986.

[24] Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. In *The Eighth International Conference on Learning Representations*, 2020. URL https://openreview.net/forum?id=rygGQyrFvH. arXiv:1904.09751.

[25] Jin Xu, Xiaojiang Liu, Jianhao Yan, Deng Cai, Huayang Li, and Jian Li. Learning to break the loop: Analyzing and mitigating repetitions for neural text generation. In *Advances in Neural Information Processing Systems*, volume 35, pages 3082–3095, 2022. doi: 10.52202/068431-0223. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/148c0aeea1c5da82f4fa86a09d4190da-Abstract-Conference.html. arXiv:2206.02369.

[26] Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. AI models collapse when trained on recursively generated data. *Nature*, 631(8022):755–759, 2024. doi: 10.1038/s41586-024-07566-y.

[27] António H. Ribeiro, Koen Tiels, Luis A. Aguirre, and Thomas Schön. Beyond exploding and vanishing gradients: analysing RNN training using attractors and smoothness. In *Proceedings of the Twenty Third International Conference on Artificial Intelligence and Statistics*, volume 108 of *Proceedings of Machine Learning Research*, pages 2370–2380. PMLR, 2020. URL https://proceedings.mlr.press/v108/ribeiro20a.html.

[28] David Sussillo and Omri Barak. Opening the black box: Low-dimensional dynamics in high-dimensional recurrent neural networks. *Neural Computation*, 25(3):626–649, 2013. doi: 10.1162/NECO_a_00409.

[29] Niru Maheswaranathan, Alex Williams, Matthew Golub, Surya Ganguli, and David Sussillo. Reverse engineering recurrent networks for sentiment classification reveals line attractor dynamics. In *Advances in Neural Information Processing Systems*, volume 32, pages 15696–15705, 2019. URL https://proceedings.neurips.cc/paper_files/paper/2019/hash/d921c3c762b1522c475ac8fc0811bb0f-Abstract.html. arXiv:1906.10720.

[30] Jonas Geiping, Sean McLeish, Neel Jain, et al. Scaling up test-time compute with latent reasoning: A recurrent depth approach. In *Advances in Neural Information Processing Systems*, volume 38, pages 46242–46293, 2025. doi: 10.52202/085713-1380. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/3b01972cf31e6fa0fe29e4b8b5c2a0a1-Abstract-Conference.html.

[31] Rui-Jie Zhu, Zixuan Wang, Kai Hua, et al. Scaling latent reasoning via looped language models. arXiv preprint arXiv:2510.25741, 2025. URL https://arxiv.org/abs/2510.25741.

[32] Hayden Prairie, Zachary Novack, Taylor Berg-Kirkpatrick, and Daniel Y. Fu. Parcae: Scaling laws for stable looped language models. In *Third Conference on Language Modeling*, 2026. URL https://openreview.net/forum?id=pIBqtrqFeP. arXiv:2604.12946.

[33] Matthias M. M. Buehlmaier. LayerNorm as implicit gain control in looped transformers. arXiv preprint arXiv:2607.10681, 2026. URL https://arxiv.org/abs/2607.10681. Version 1, 12 July 2026.

[34] Asher Labovich. Stability and generalization in looped transformers. arXiv preprint arXiv:2604.15259, 2026. URL https://arxiv.org/abs/2604.15259. Version 2, 22 April 2026.

[35] Sajad Movahedi, Vera Milovanović, Shlomo Libo Feigin, Alexander Theus, Thomas Hofmann, Valentina Boeva, T. Konstantin Rusch, and Antonio Orvieto. Fixed-point reasoners: Stable and adaptive deep looped transformers. arXiv preprint arXiv:2606.18206, 2026. URL https://arxiv.org/abs/2606.18206. Version 1, 16 June 2026.

[36] Hugh Blayney, Álvaro Arroyo, Johan Obando-Ceron, et al. A mechanistic analysis of looped reasoning language models. arXiv preprint arXiv:2604.11791, 2026. URL https://arxiv.org/abs/2604.11791.

[37] Ivan Viakhirev, Kirill Borodin, Amirah Almutairi, et al. Think shallow, solve deep: Controlling recurrent dynamics for reliable test-time depth. arXiv preprint arXiv:2608.18222, 2026. URL https://arxiv.org/abs/2608.18222.

[38] Sumin Park and T. Konstantin Rusch. Looped reasoning finishes earlier than you think. arXiv preprint arXiv:2610.18174, 2026. URL https://arxiv.org/abs/2610.18174.

[39] Qucheng Gao, Zuyi Yang, and Xiao Chen. Nonequilibrium phases of repulsive self-attention: Chaos, attention condensation, and emergent locality. arXiv preprint arXiv:2609.28448, 2026. URL https://arxiv.org/abs/2609.28448.

[40] Junhong Wu, Jinliang Lu, Zixuan Ren, Gangqiang Hu, Zhi Wu, Dai Dai, and Hua Wu. LLMs are single-threaded reasoners: Demystifying the working mechanism of Soft Thinking. In *The Fourteenth International Conference on Learning Representations*, pages 4533–4548, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/0895de150660b6c2aef9b0f098cb1dcd-Abstract-Conference.html. arXiv:2508.03440.

[41] Ali Backour. The dynamics of continuous mixture collapse in language models. arXiv preprint arXiv:2609.02049, 2026. URL https://arxiv.org/abs/2609.02049.

[42] Michael Rizvi-Martel, Guillaume Rabusseau, and Marius Mosbach. The illusion of superposition? A principled analysis of latent thinking in language models. arXiv preprint arXiv:2604.06374, 2026. URL https://arxiv.org/abs/2604.06374.

[43] Fatih Dinc, Ege Cirakman, Bariscan Kurtkaya, Mert Yuksekgonul, Yiqi Jiang, Mark J. Schnitzer, and Hidenori Tanaka. Ghost mechanism: An analytical model of abrupt learning in recurrent networks. *Physical Review X*, 16(2):021058, 2026. doi: 10.1103/mjcl-lb4x.

[44] Demian Battaglia, Nicolas Brunel, and David Hansel. Temporal decorrelation of collective oscillations in neural networks with local inhibition and long-range excitation. *Physical Review Letters*, 99(23):238106, 2007. doi: 10.1103/PhysRevLett.99.238106.

[45] Christoforos A. Papasavvas, Yujiang Wang, Andrew J. Trevelyan, and Marcus Kaiser. Gain control through divisive inhibition prevents abrupt transition to chaos in a neural mass model. *Physical Review E*, 92(3):032723, 2015. doi: 10.1103/PhysRevE.92.032723.

[46] Yilun Liu, Yi Zhang, Ganyu Wu, Sikuan Yan, Mengyue Wang, Alois Knoll, Volker Tresp, and Yunpu Ma. Learning chaos without seeing chaos: Extrapolation of global dynamics in autoregressive transformers. arXiv preprint arXiv:2609.38814, 2026. URL https://arxiv.org/abs/2609.38814v1. Version 1, 30 September 2026.

[47] Jason Z. Kim, Zhixin Lu, Erfan Nozari, George J. Pappas, and Danielle S. Bassett. Teaching recurrent neural networks to infer global temporal structure from local examples. *Nature Machine Intelligence*, 3(4):316–323, 2021. doi: 10.1038/s42256-021-00321-2.

[48] Dayal Singh Kalra, Tianyu He, and Maissam Barkeshli. Universal sharpness dynamics in neural network training: Fixed point analysis, edge of stability, and route to chaos. In *The Thirteenth International Conference on Learning Representations*, pages 55966–56000, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/8c8cd1b78cdae751265c88efc136e5bd-Abstract-Conference.html.

[49] R. L. Marichal and J. D. Piñeiro. Analysis of multiple quasi-periodic orbits in recurrent neural networks. *Neurocomputing*, 162:85–95, 2015. doi: 10.1016/j.neucom.2015.04.001.

[50] Florian Gomez, Ralph L. Stoop, and Ruedi Stoop. Universal dynamical properties preclude standard clustering in a large class of biochemical data. *Bioinformatics*, 30(17):2486–2493, 2014. doi: 10.1093/bioinformatics/btu332.

[51] Rafael V. Stenzinger, Vinícius Luz Oliveira, and M. H. R. Tragtenberg. Shrimp hubs in the Hindmarsh–Rose model. *Chaos: An Interdisciplinary Journal of Nonlinear Science*, 35(2): 023148, 2025. doi: 10.1063/5.0239268.

[52] Luiz F. B. Caixeta, Matheus H. P. Gonçalves, M. H. R. Tragtenberg, and Mauricio Girardi-Schappo. Devil's staircase inside shrimp-shaped regions reveals periodicity of plateau spikes and bursts. *Chaos: An Interdisciplinary Journal of Nonlinear Science*, 35(3):033122, 2025. doi: 10.1063/5.0250342.

[53] Andy Zou, Long Phan, Sarah Chen, et al. Representation engineering: A top-down approach to AI transparency. arXiv preprint arXiv:2310.01405, 2023. URL https://arxiv.org/abs/2310.01405. Version 4, 3 March 2025; first posted 2 October 2023.

[54] Jinghao Zhang, Yuting Liu, Wenjie Wang, Qiang Liu, Shu Wu, Liang Wang, and Tat-Seng Chua. Personalized text generation with contrastive activation steering. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7128–7141. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.acl-long.353.

[55] Yuri A. Kuznetsov. *Elements of Applied Bifurcation Theory*, volume 112 of *Applied Mathematical Sciences*. Springer, Cham, fourth edition, 2023. doi: 10.1007/978-3-031-22007-4.

[56] Willy J. F. Govaerts. *Numerical Methods for Bifurcations of Dynamical Equilibria*. Society for Industrial and Applied Mathematics, 2000. doi: 10.1137/1.9780898719543.

[57] Yuri A. Kuznetsov and Hil G. E. Meijer. *Numerical Bifurcation Analysis of Maps*. Cambridge University Press, Cambridge, 2019. doi: 10.1017/9781108585804.

[58] Yves Pomeau and Paul Manneville. Intermittent transition to turbulence in dissipative dynamical systems. *Communications in Mathematical Physics*, 74(2):189–197, 1980. doi: 10.1007/BF01197757.

[59] Wilson Façanha, Bart Oldeman, and Leon Glass. Bifurcation structures in two-dimensional maps: The endoskeletons of shrimps. *Physics Letters A*, 377(18):1264–1268, 2013. doi: 10.1016/j.physleta.2013.03.025.

[60] Renato Vitolo, Paul Glendinning, and Jason A. C. Gallas. Global structure of periodicity hubs in Lyapunov phase diagrams of dissipative flows. *Physical Review E*, 84(1):016216, 2011. doi: 10.1103/PhysRevE.84.016216.

[61] Timothy D. Jones. Topological origins of a bi-parameter periodicity hub for the Rössler attractor. arXiv preprint arXiv:1201.4343, 2012. URL https://arxiv.org/abs/1201.4343.

[62] Hubert Ramsauer, Bernhard Schäfl, Johannes Lehner, et al. Hopfield networks is all you need. In *The Ninth International Conference on Learning Representations*, 2021. URL https://openreview.net/forum?id=tL89RnzIiCd. arXiv:2008.02217.

[63] Yihe Dong, Jean-Baptiste Cordonnier, and Andreas Loukas. Attention is not all you need: Pure attention loses rank doubly exponentially with depth. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 2793–2803. PMLR, 2021. URL https://proceedings.mlr.press/v139/dong21a.html. arXiv:2103.03404.

[64] Borjan Geshkovski, Cyril Letrouit, Yury Polyanskiy, and Philippe Rigollet. The emergence of clusters in self-attention dynamics. In *Advances in Neural Information Processing Systems*, volume 36, pages 57026–57037, 2023. doi: 10.52202/075280-2493. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/b2b3e1d9840eba17ad9bbf073e009afe-Abstract-Conference.html. arXiv:2305.05465.

[65] Giancarlo Benettin, Luigi Galgani, Antonio Giorgilli, and Jean-Marie Strelcyn. Lyapunov characteristic exponents for smooth dynamical systems and for Hamiltonian systems; a method for computing all of them. Part 1: Theory. *Meccanica*, 15(1):9–20, 1980. doi: 10.1007/BF02128236.

[66] Giancarlo Benettin, Luigi Galgani, Antonio Giorgilli, and Jean-Marie Strelcyn. Lyapunov characteristic exponents for smooth dynamical systems and for Hamiltonian systems; a method for computing all of them. Part 2: Numerical application. *Meccanica*, 15(1):21–30, 1980. doi: 10.1007/BF02128237.

[67] V. I. Oseledets. A multiplicative ergodic theorem. Lyapunov characteristic numbers for dynamical systems. *Transactions of the Moscow Mathematical Society*, 19:197–231, 1968. English translation of Trudy Moskovskogo Matematicheskogo Obshchestva 19 (1968) 179–210; Zbl 0236.93034.

[68] Frank Pasemann. Complex dynamics and the structure of small neural networks. *Network: Computation in Neural Systems*, 13(2):195–216, 2002. doi: 10.1080/net.13.2.195.216.

[69] Zihao Fu, Wai Lam, Anthony Man-Cho So, and Bei Shi. A theoretical analysis of the repetition problem in text generation. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35 (14):12848–12856, 2021. doi: 10.1609/aaai.v35i14.17520. arXiv:2012.14660.

[70] Catherine Olsson, Nelson Elhage, Neel Nanda, et al. In-context learning and induction heads. Transformer Circuits Thread, 2022. URL https://transformer-circuits.pub/2022/in-context-learning-and-induction-heads/index.html. arXiv:2209.11895.

[71] Team OLMo, Pete Walsh, Luca Soldaini, Dirk Groeneveld, et al. 2 OLMo 2 furious. arXiv preprint arXiv:2501.00656, 2024. URL https://arxiv.org/abs/2501.00656. Shorter version in the Second Conference on Language Modeling (2025).

[72] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, et al. The Llama 3 herd of models. arXiv preprint arXiv:2407.21783, 2024. URL https://arxiv.org/abs/2407.21783.

[73] Goodfire. Llama-3.1-8B-Instruct-SAE-l19. Hugging Face, https://huggingface.co/Goodfire/Llama-3.1-8B-Instruct-SAE-l19, 2025. Sparse autoencoder on the residual stream after layer 19; revision f6775a22.

[74] Seonglae Cho, Harryn Oh, Donghyun Lee, Luis Rodrigues Vieira, Andrew Bermingham, and Ziad El Sayed. FaithfulSAE: Towards capturing faithful features with sparse autoencoders without external datasets dependency. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 4: Student Research Workshop)*, pages 297–314. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.acl-srw.20. arXiv:2506.17673.

[75] Gonçalo Paulo and Nora Belrose. Sparse autoencoders trained on the same data learn different features. In *The Fourteenth International Conference on Learning Representations*, pages 35748–35760, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/3c1fe56b043848b211030c202764c6a7-Abstract-Conference.html. arXiv:2501.16615.

[76] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, et al. Training verifiers to solve math word problems. arXiv preprint arXiv:2110.14168, 2021. URL https://arxiv.org/abs/2110.14168.

[77] Boyuan Wang, Chengyao Yu, Jiaxi Ren, Hongxin Wei, Bingyi Jing, and Yuxin Tao. Scheduling recursive reasoning in looped transformers. arXiv preprint arXiv:2609.36653, 2026. URL https://arxiv.org/abs/2609.36653. Submitted 29 September 2026.

[78] Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? Try ARC, the AI2 Reasoning Challenge. arXiv preprint arXiv:1803.05457, 2018. URL https://arxiv.org/abs/1803.05457.

[79] Nina Rimsky, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Turner. Steering Llama 2 via contrastive activation addition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15504–15522. Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.acl-long.828.

[80] Jonathan Kadmon and Haim Sompolinsky. Transition to chaos in random neuronal networks. *Physical Review X*, 5(4):041030, 2015. doi: 10.1103/PhysRevX.5.041030.

[81] Kanaka Rajan, L. F. Abbott, and Haim Sompolinsky. Stimulus-dependent suppression of chaos in recurrent neural networks. *Physical Review E*, 82(1):011903, 2010. doi: 10.1103/PhysRevE.82.011903.

[82] Ben Poole, Subhaneil Lahiri, Maithra Raghu, Jascha Sohl-Dickstein, and Surya Ganguli. Exponential expressivity in deep neural networks through transient chaos. In *Advances in Neural Information Processing Systems*, volume 29, pages 3360–3368, 2016. URL https://proceedings.neurips.cc/paper_files/paper/2016/hash/148510031349642de5ca0c544f31b2ef-Abstract.html. arXiv:1606.05340.

[83] Samuel S. Schoenholz, Justin Gilmer, Surya Ganguli, and Jascha Sohl-Dickstein. Deep information propagation. In *The Fifth International Conference on Learning Representations*, 2017. URL https://openreview.net/forum?id=H1W1UN9gg. arXiv:1611.01232.

[84] Cristian Bonatto and Jason A. C. Gallas. Periodicity hub and nested spirals in the phase diagram of a simple resistive circuit. *Physical Review Letters*, 101(5):054101, 2008. doi: 10.1103/PhysRevLett.101.054101.

[85] Ruedi Stoop, Philipp Benner, and Yoko Uwate. Real-world existence and origins of the spiral organization of shrimp-shaped domains. *Physical Review Letters*, 105(7):074102, 2010. doi: 10.1103/PhysRevLett.105.074102.

[86] Ying-Cheng Lai and Tamás Tél. *Transient Chaos: Complex Dynamics on Finite Time Scales*, volume 173 of *Applied Mathematical Sciences*. Springer, New York, NY, 2011. doi: 10.1007/978-1-4419-6987-3.

[87] Chris G. Langton. Computation at the edge of chaos: Phase transitions and emergent computation. *Physica D: Nonlinear Phenomena*, 42(1–3):12–37, 1990. doi: 10.1016/0167-2789(90)90064-V.

[88] Nils Bertschinger and Thomas Natschläger. Real-time computation at the edge of chaos in recurrent neural networks. *Neural Computation*, 16(7):1413–1436, 2004. doi: 10.1162/089976604323057443.

[89] Fredrik Johansson. Arb: Efficient arbitrary-precision midpoint-radius interval arithmetic. *IEEE Transactions on Computers*, 66(8):1281–1292, 2017. doi: 10.1109/TC.2017.2690633.

[90] R. Krawczyk. Newton-Algorithmen zur Bestimmung von Nullstellen mit Fehlerschranken. *Computing*, 4(3):187–201, 1969. doi: 10.1007/BF02234767.

[91] Ramon E. Moore, R. Baker Kearfott, and Michael J. Cloud. *Introduction to Interval Analysis*. Society for Industrial and Applied Mathematics, Philadelphia, PA, 2009. doi: 10.1137/1.9780898717716.

[92] Siegfried M. Rump. Verification methods: Rigorous results using floating-point arithmetic. *Acta Numerica*, 19:287–449, 2010. doi: 10.1017/S096249291000005X.